

# AWS EKS Kubernetes - Masterclass | DevOps, Microservices

Kalyan Reddy Daida



# AWS EKS Course Outline



# STACKSIMPLIFY

## Kubernetes On Cloud Roadmap

**Kubernetes for Beginners On Cloud**

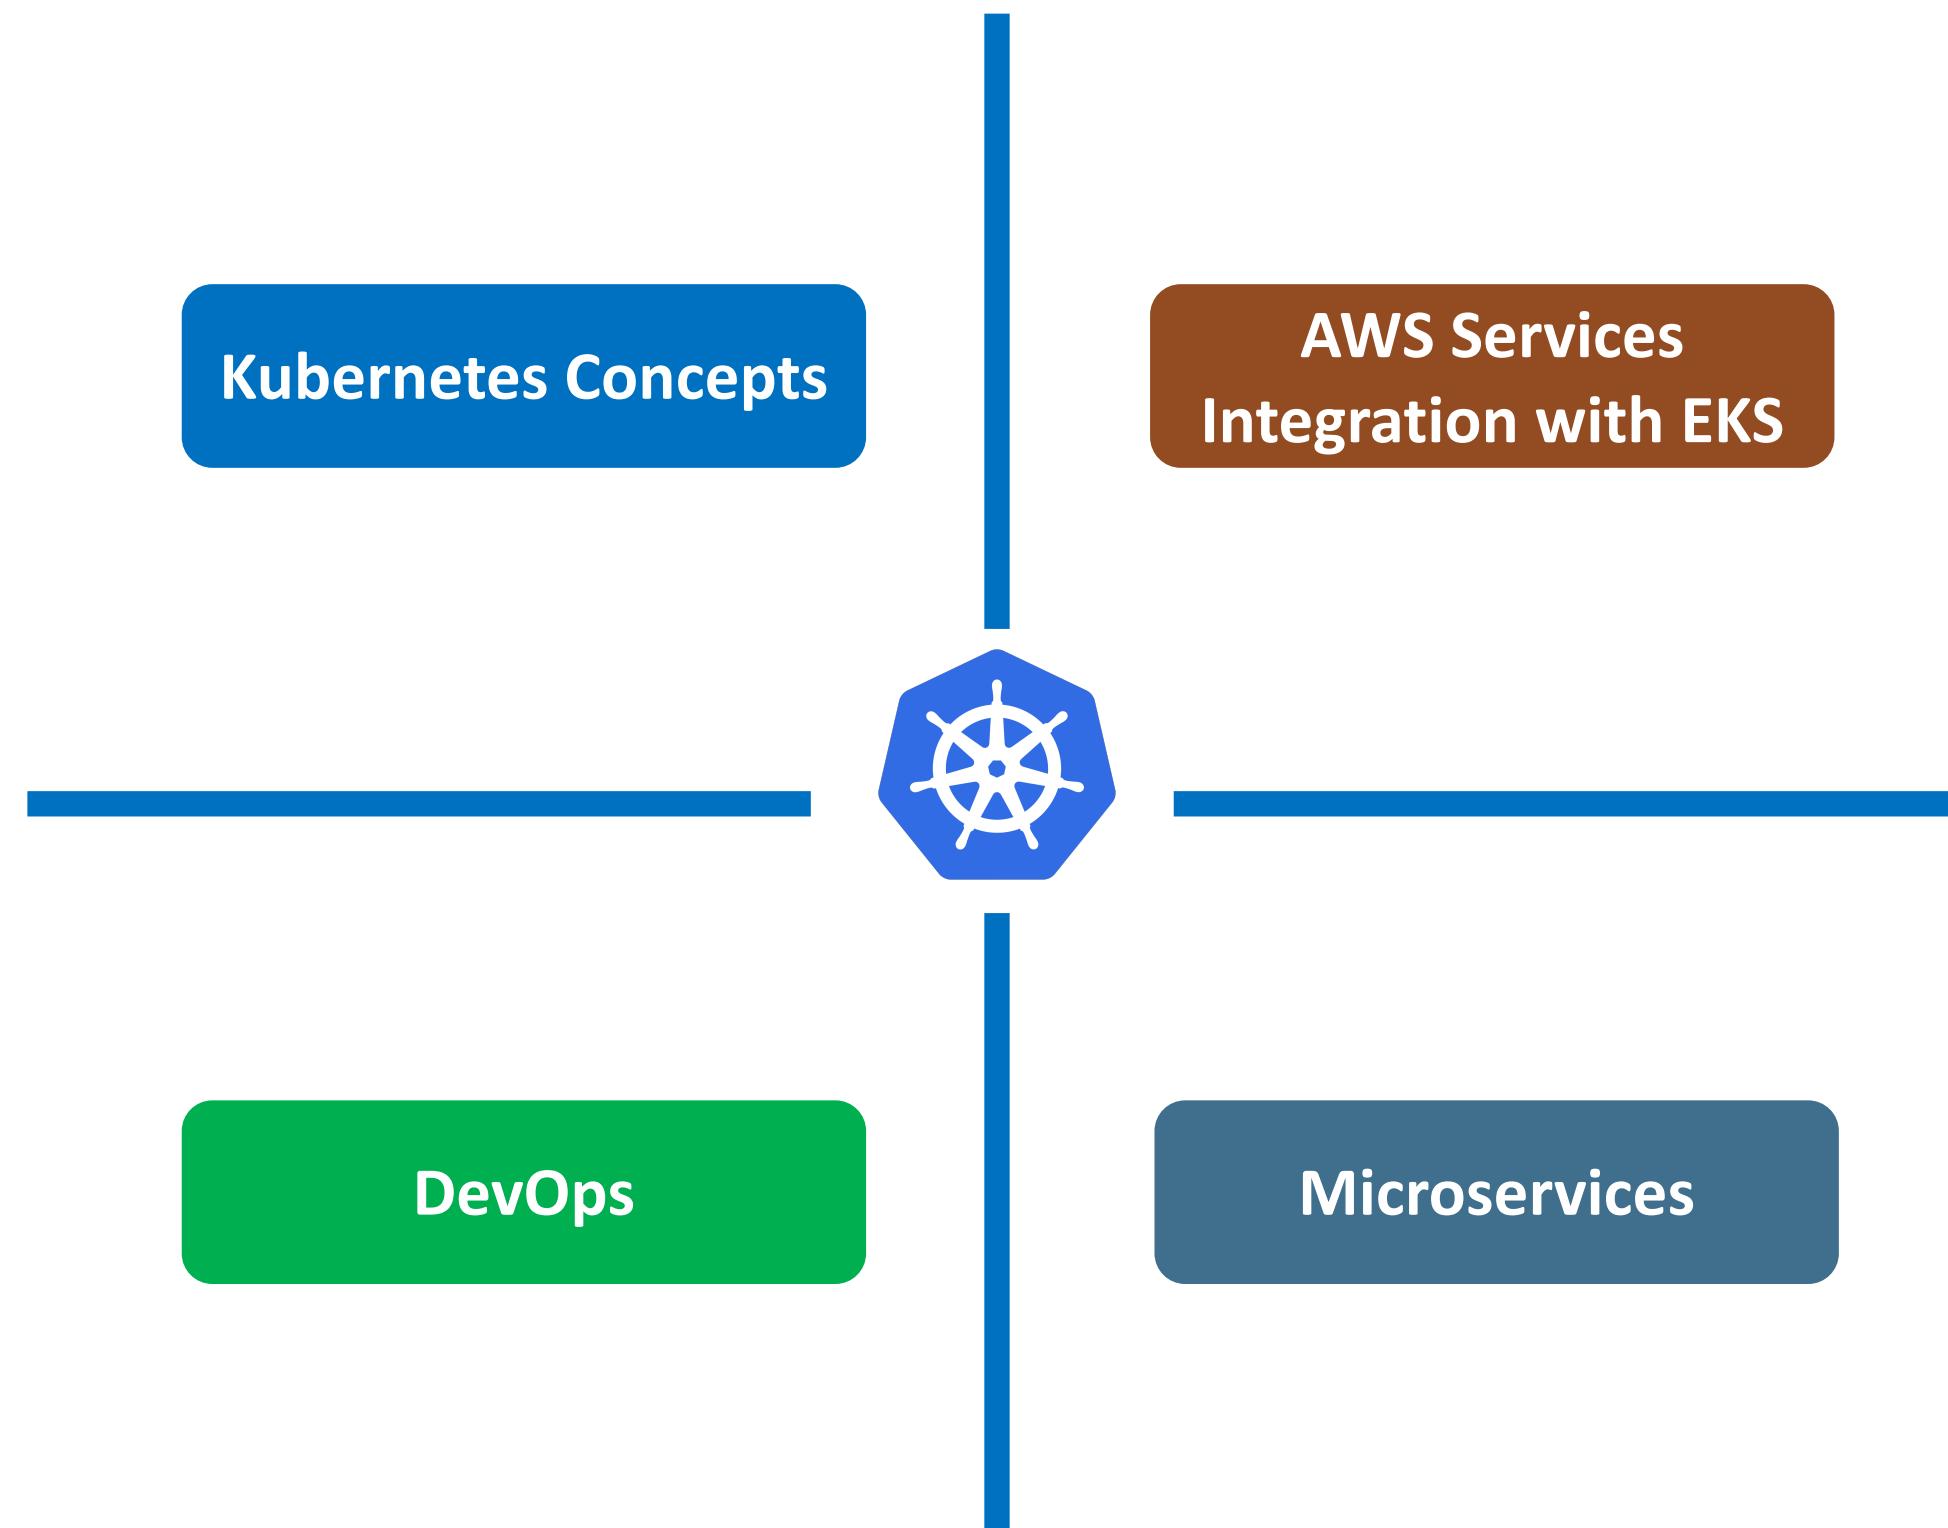
**AWS EKS Kubernetes - Masterclass |  
DevOps, Microservices**

**Azure AKS Kubernetes - Masterclass |  
DevOps, Microservices**

**Google GKE Kubernetes - Masterclass |  
DevOps, Microservices**

**Docker & Kubernetes for Java Spring Boot  
Developers on AWS**

**Master HELM3 with Kubernetes on AWS &  
Azure**



Architecture	kubectl - Imperative	Persistent Volumes
Pods	Declarative with YAML	PVC
ReplicaSets	Secrets	Load Balancers
Deployments	Init Containers	Annotations
Node Port Service	Probes	Canary Deployments
Cluster IP Service	Requests & Limits	HPA
External Name Service	Namespaces	VPA
Ingress Service	Limit Range	DaemonSets
Ingress SSL	Resource Quota	Fluentd for logs
Ingress & External DNS	Storage Classes	ConfigMaps

# K U B E R N E T E S



# A W S S E R V I C E S



**DevOps**

**AWS Developer Tools**



AWS  
CodeCommit



AWS CodeBuild



AWS CodePipeline

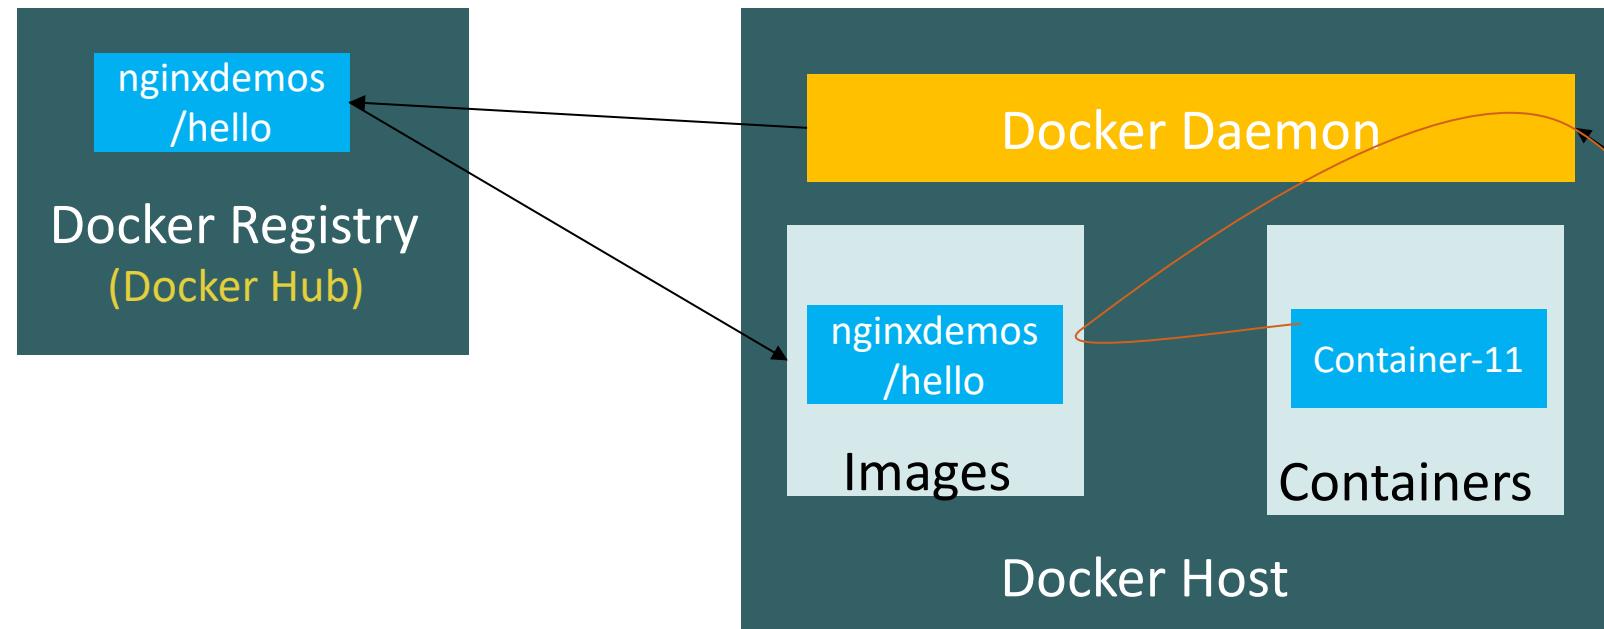
**Microservices**

**Service Discovery**

**Distributed Tracing**

**Canary Deployments**

# Docker - Fundamentals



- **Docker Registry or Docker Hub**
  - A Docker *registry* **stores** Docker images.
  - **Docker Hub** is a public registry that anyone can use, and Docker is configured to look for images on Docker Hub by default.
  - We can even run our own **private registry**.
  - When we use the **docker pull** or **docker run** commands, the required images are pulled from our configured registry.
  - When we use the **docker push** command, our image is pushed to our configured registry.

```
docker pull nginxdemos/hello  
docker run -p 82:80 -d nginxdemos/hello
```

Docker Client (My Desktop or Docker Host)

# Kubernetes - Imperative & Declarative

## Kubernetes Fundamentals

Imperative

Declarative

kubectl

Pod

ReplicaSet

Deployment

Service

YAML & kubectl

Pod

ReplicaSet

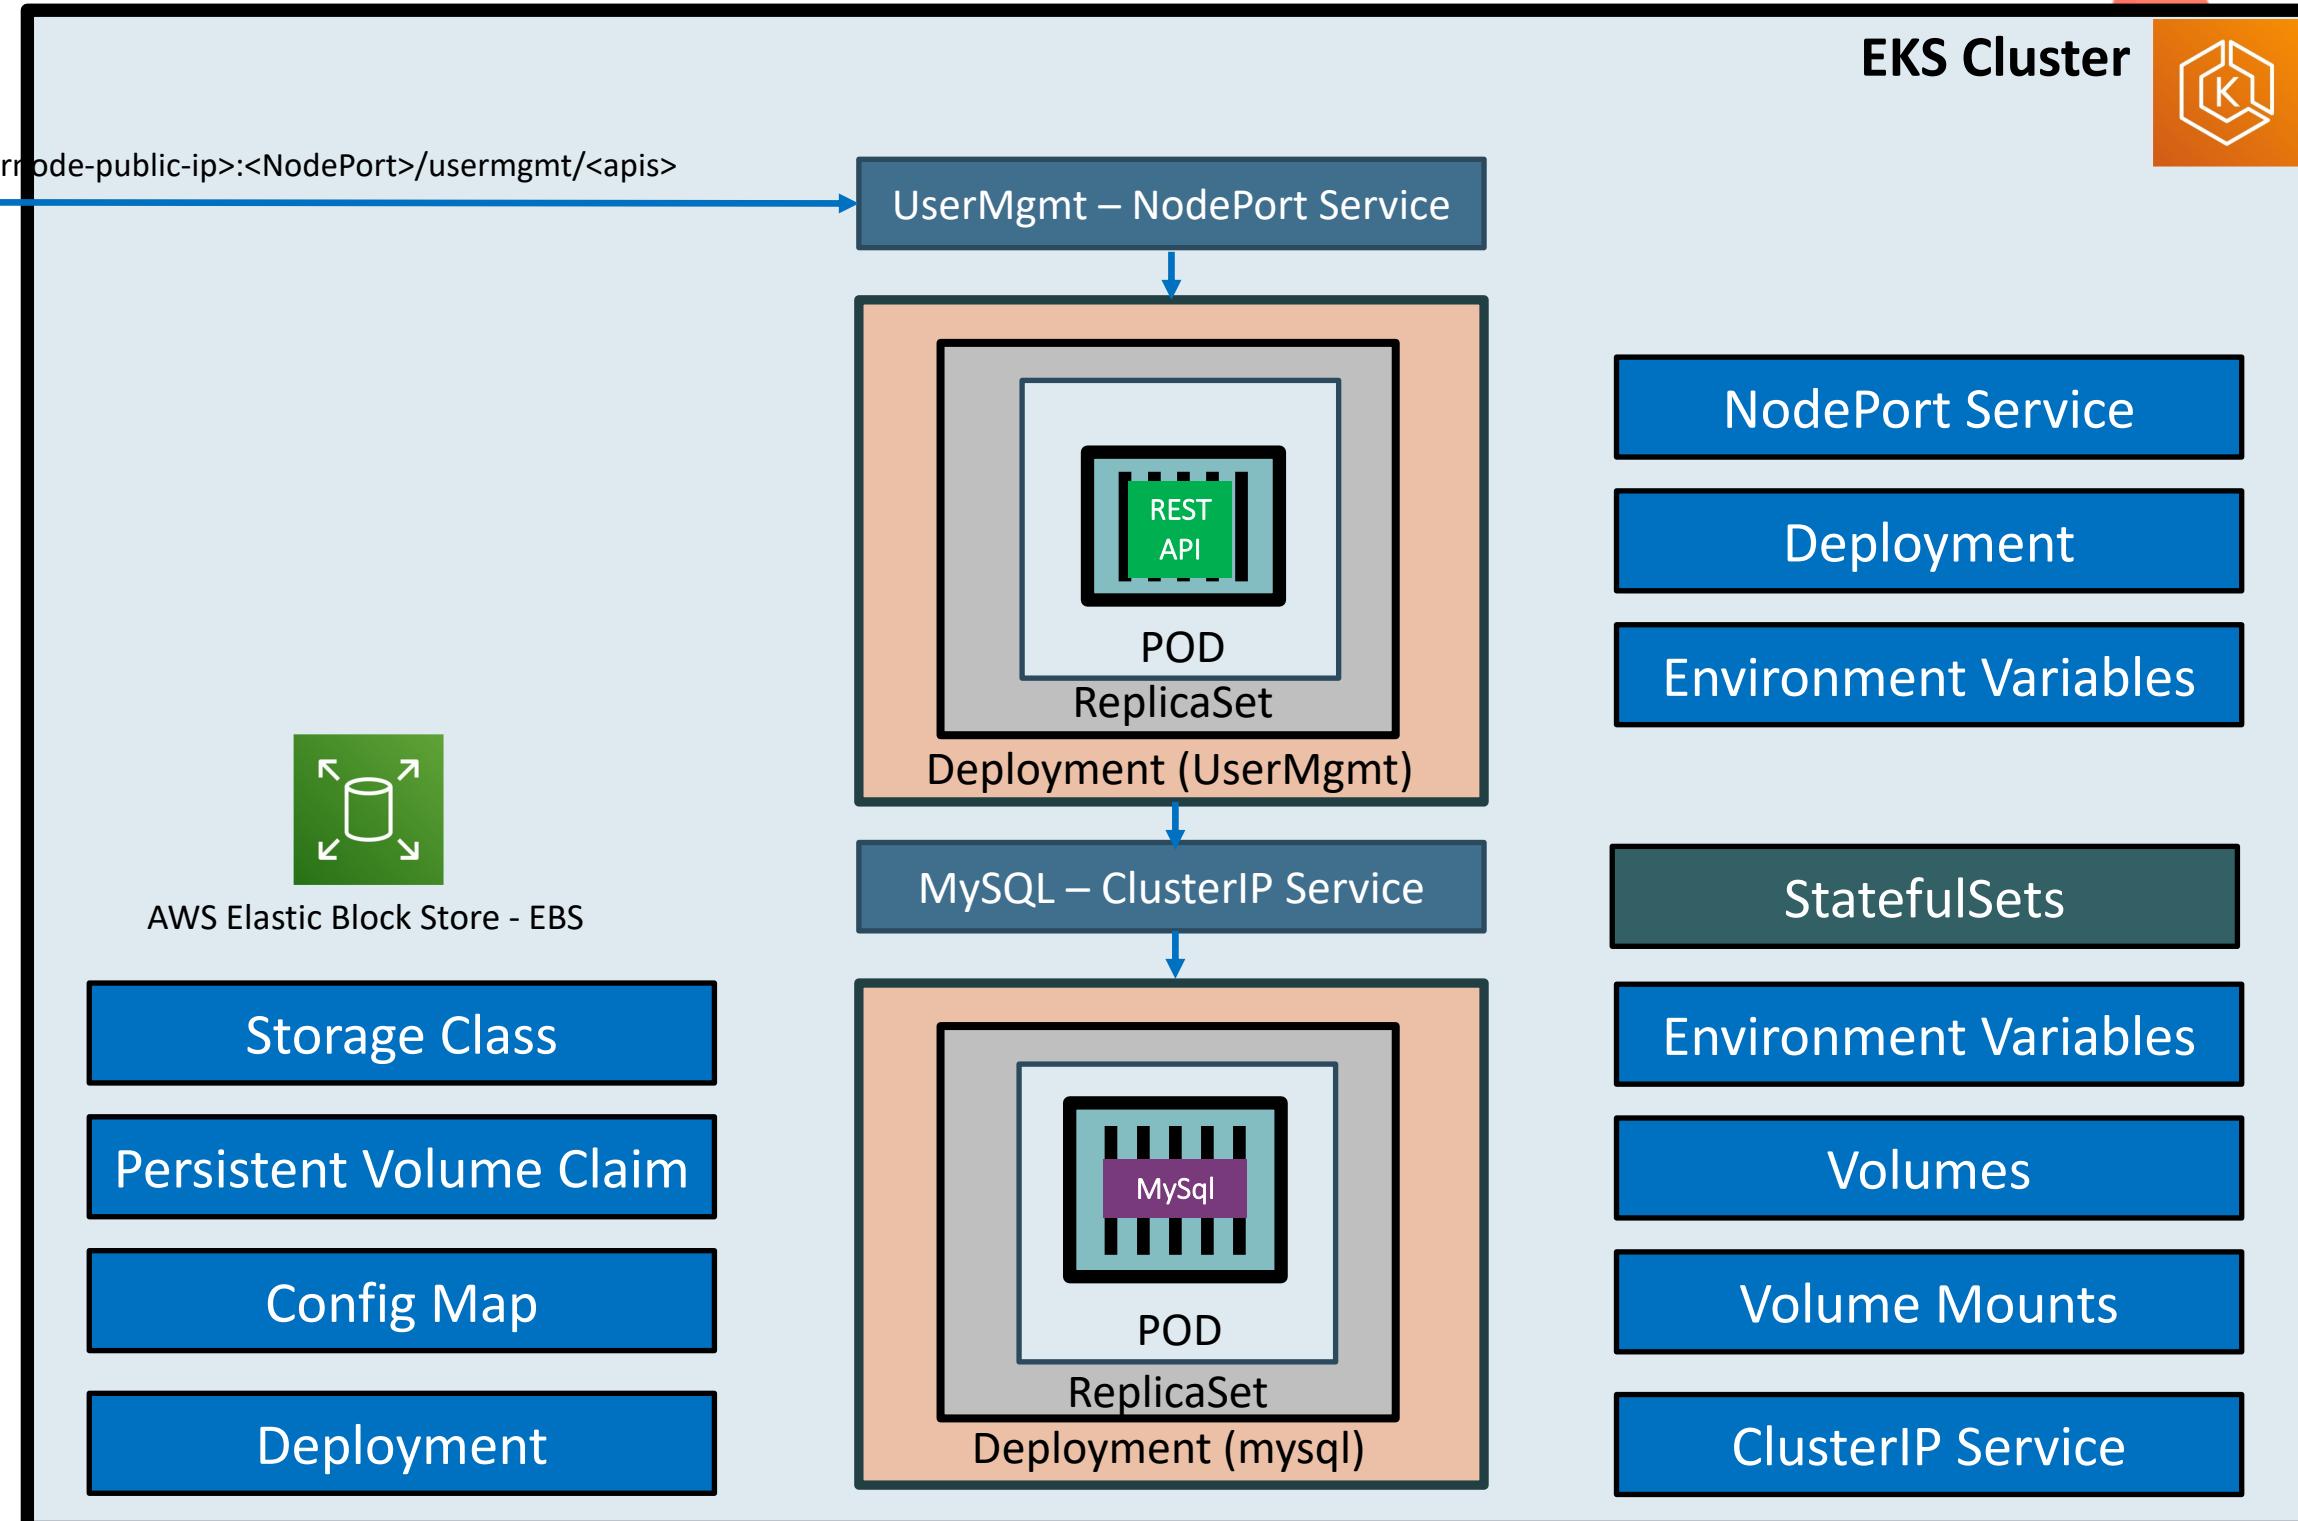
Deployment

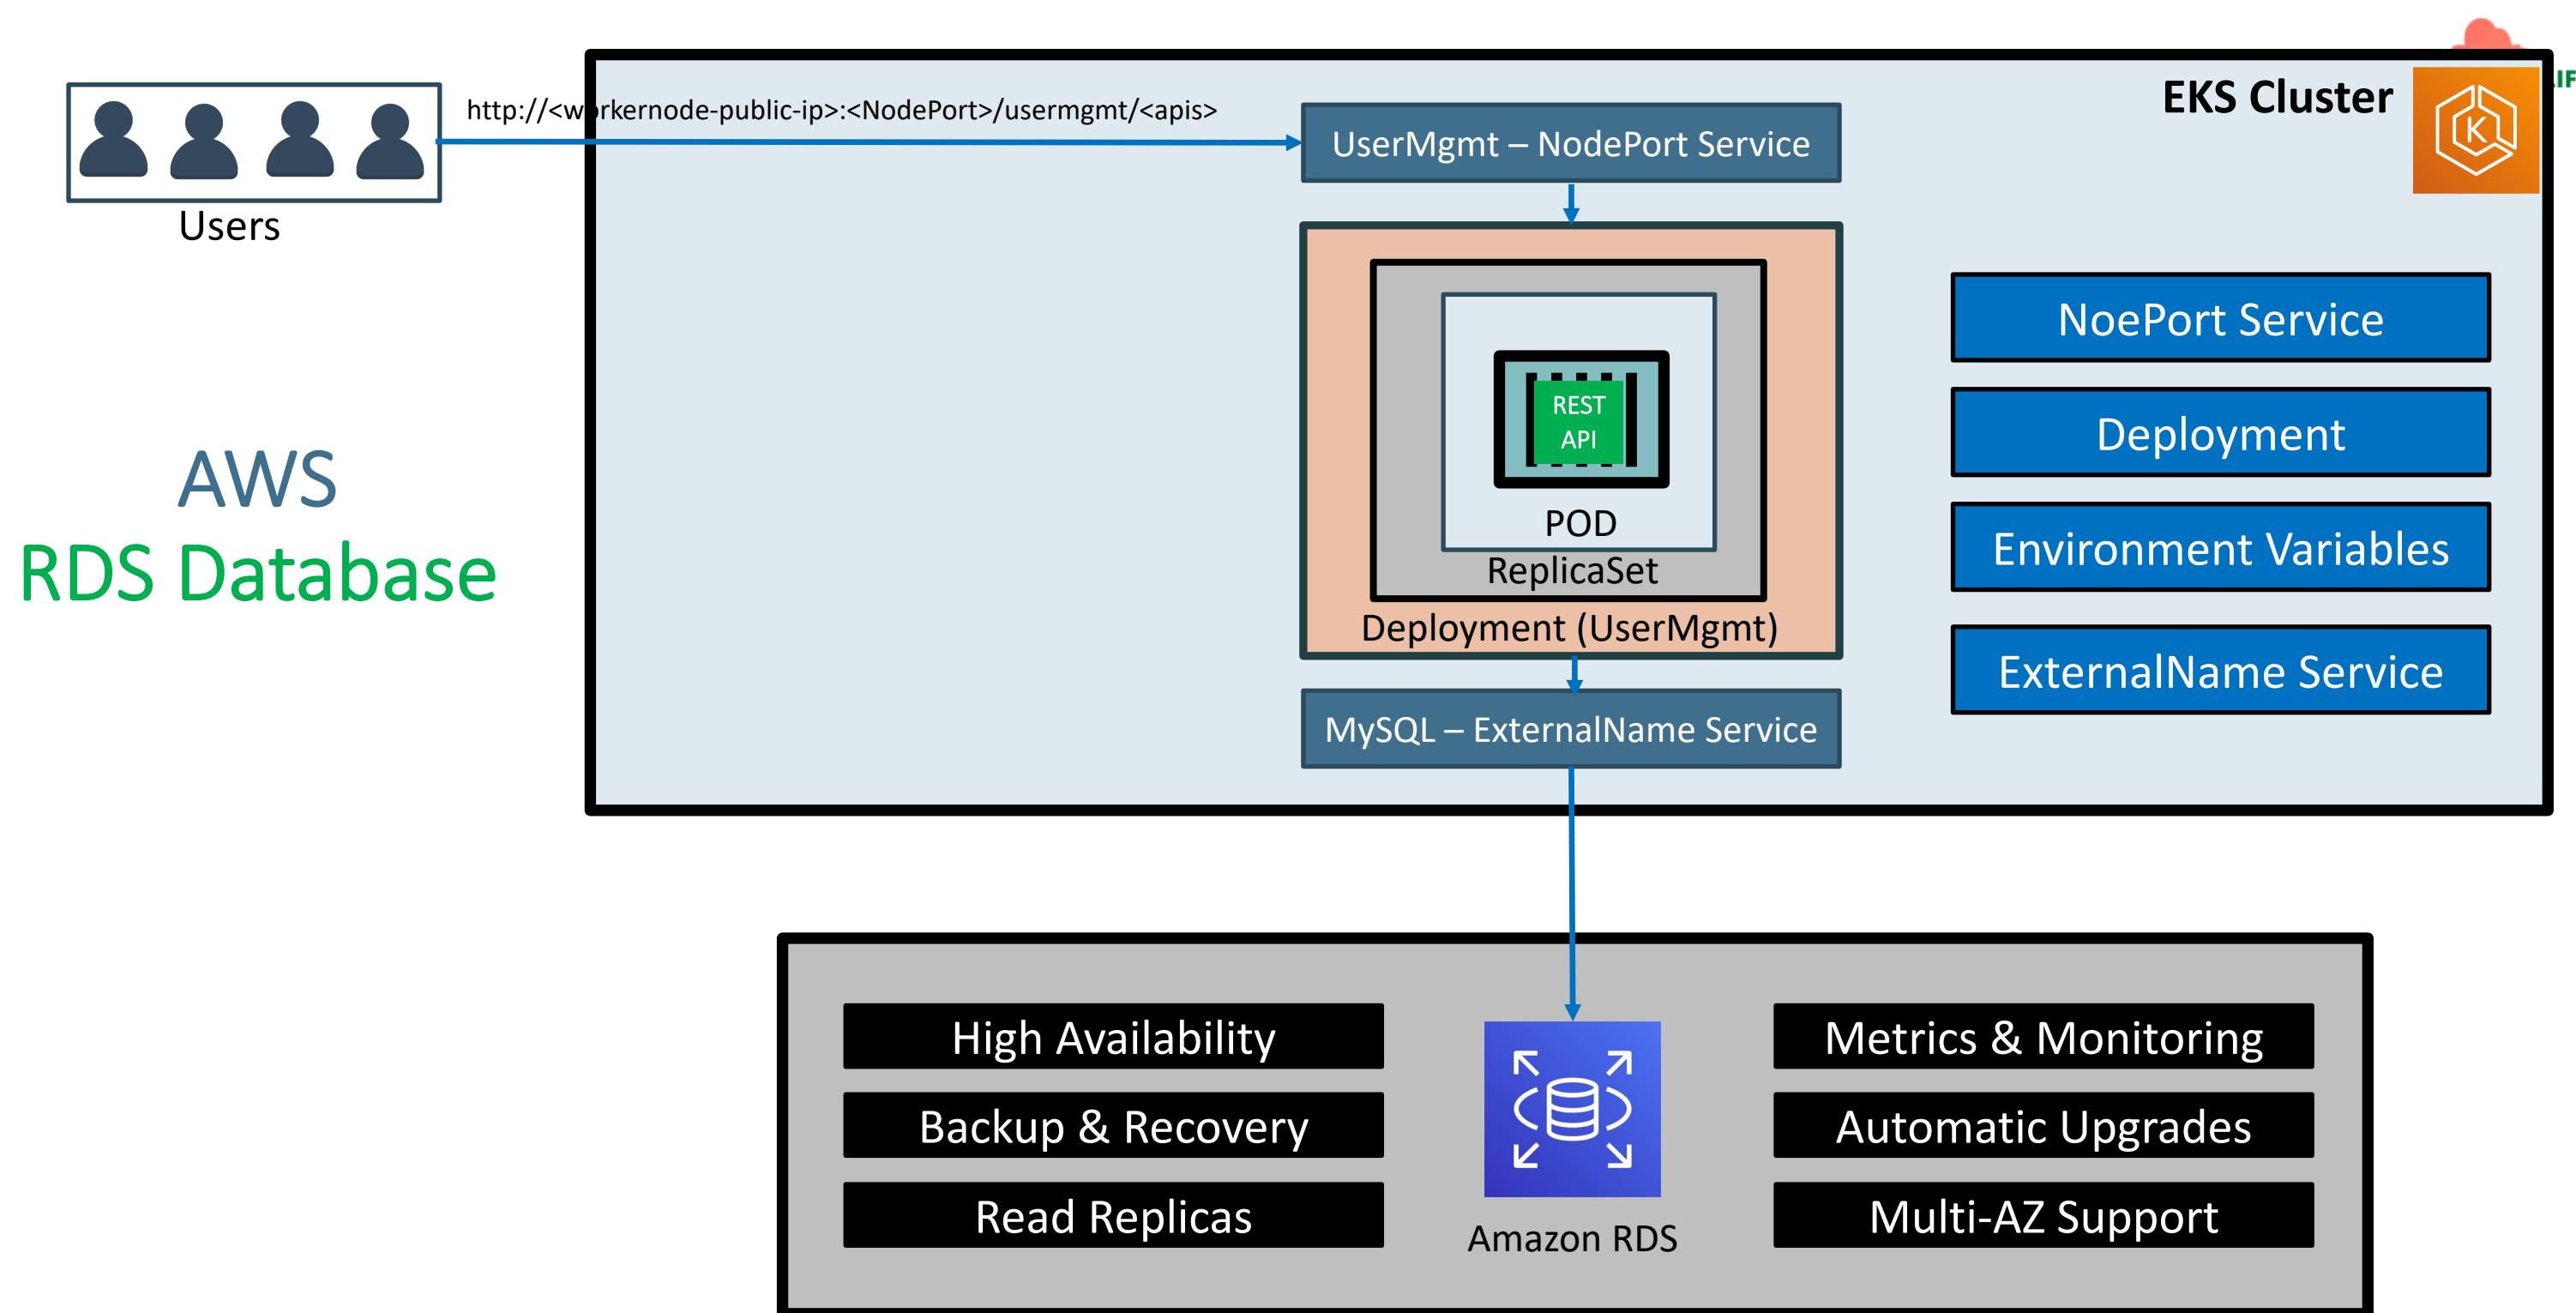
Service

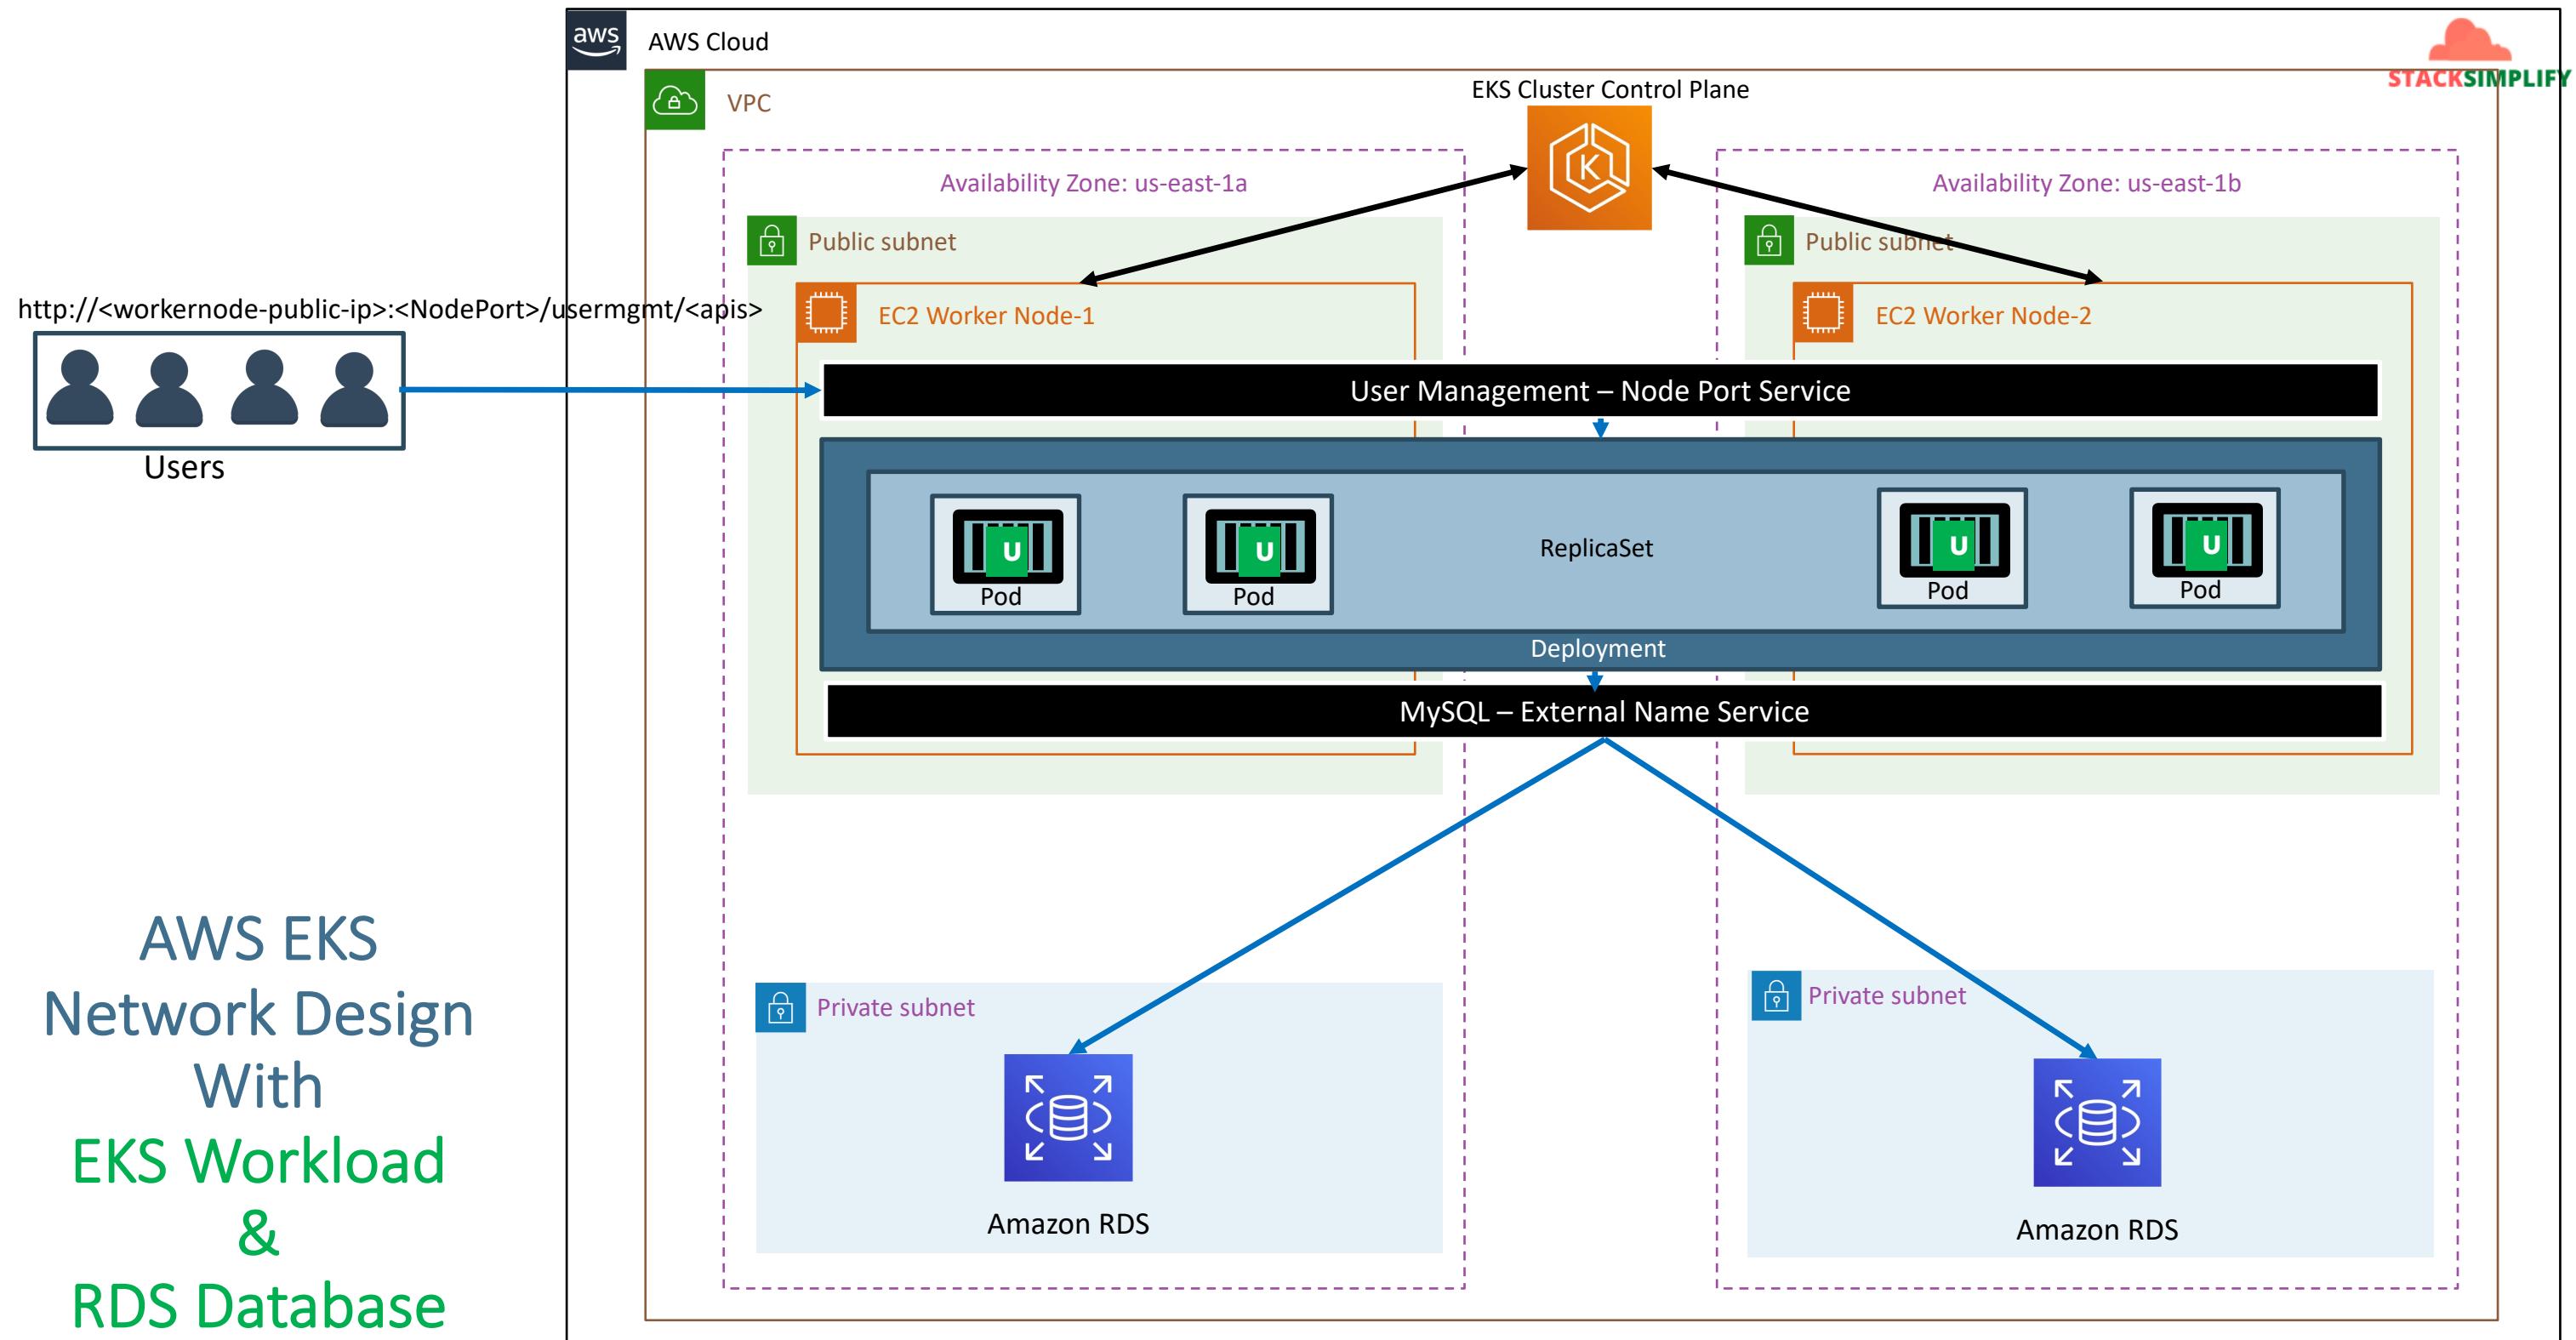


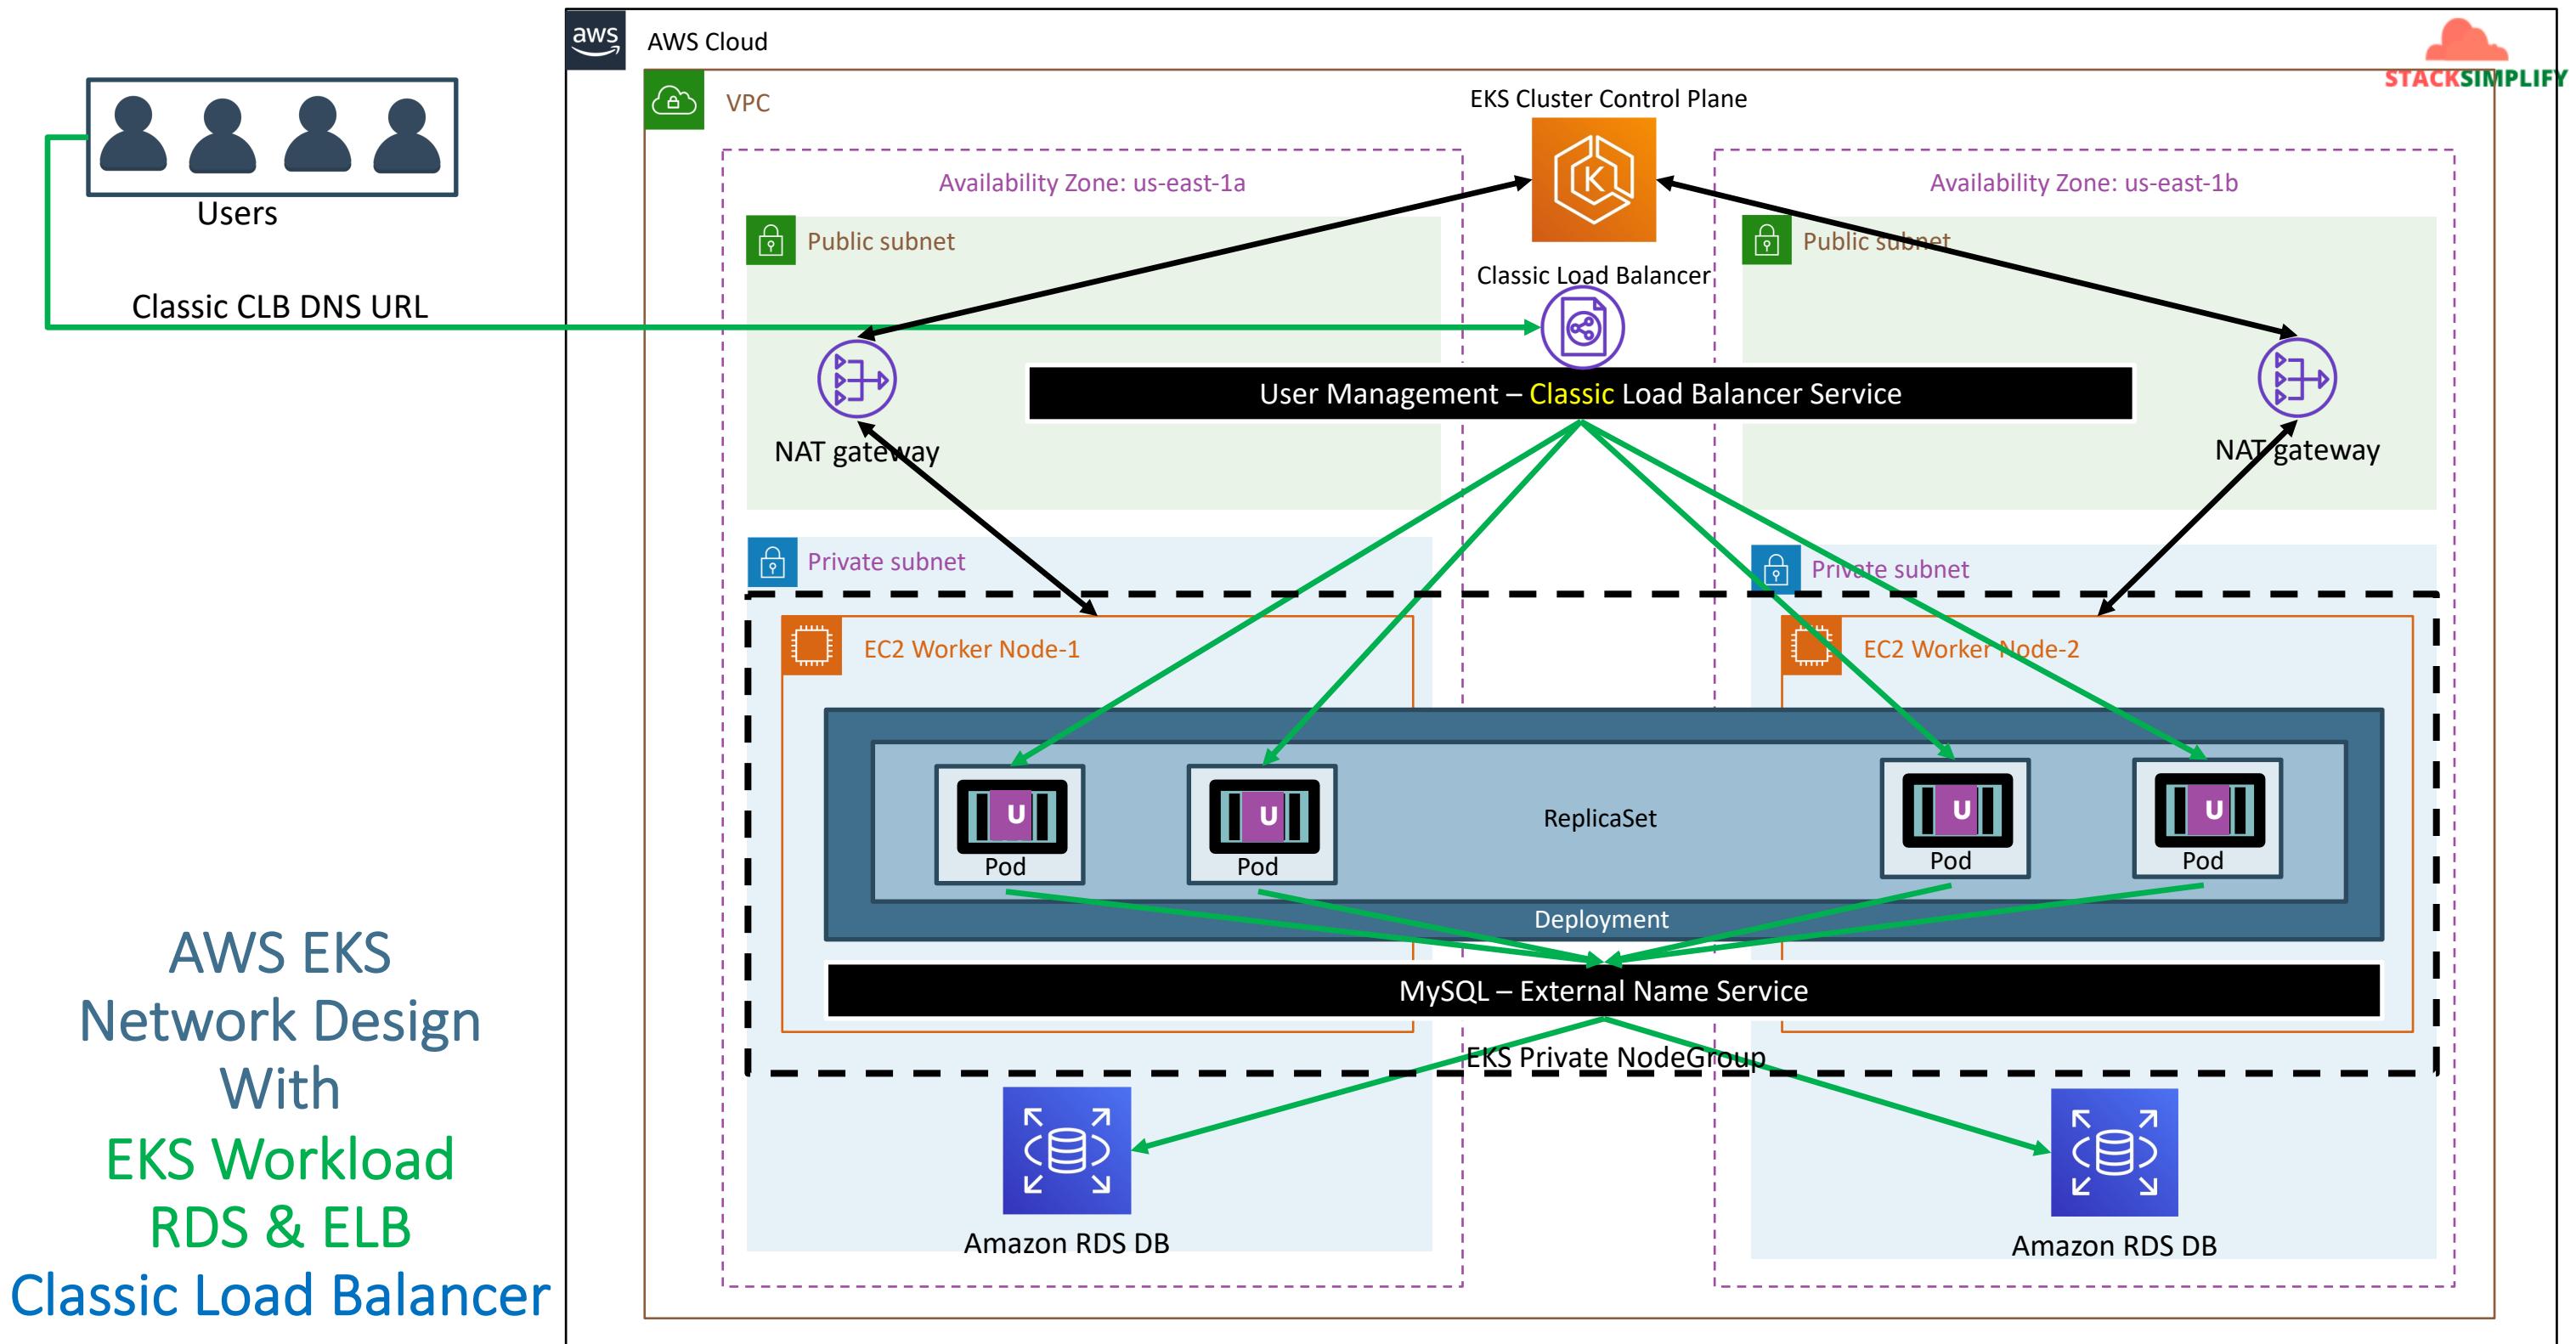
# EKS Storage EBS CSI Driver

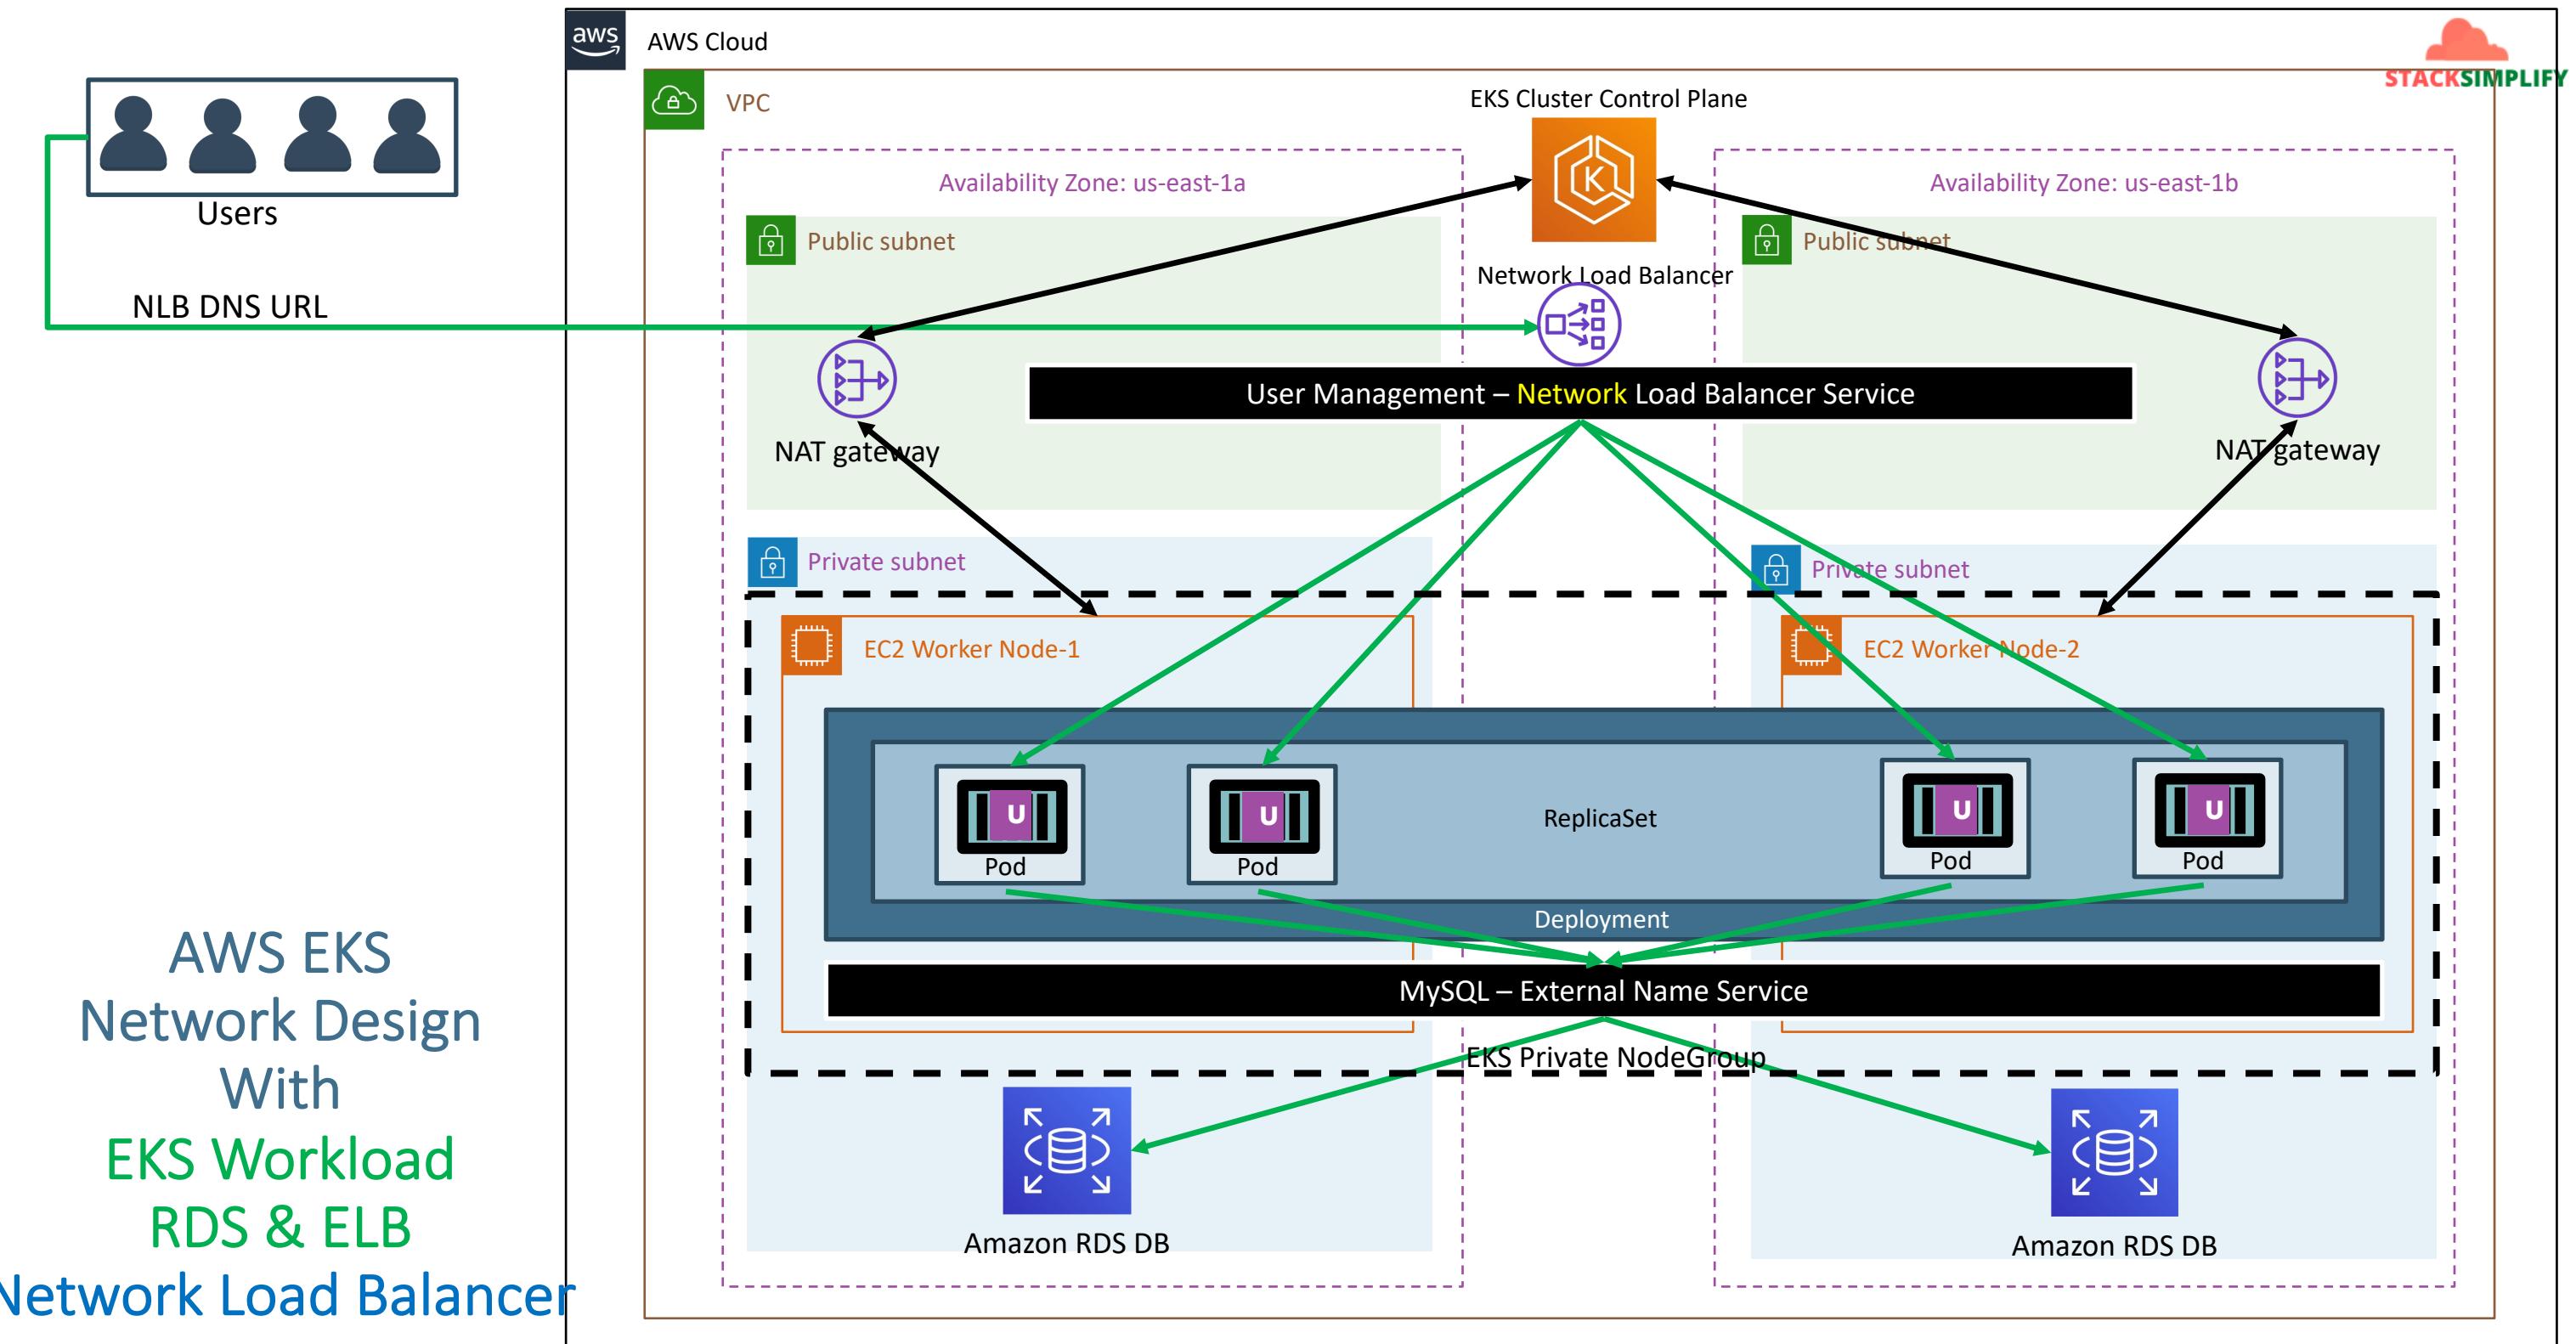
- Drawbacks of EBS CSI for MySQL DB
- Complex setup to achieve HA
- Complex Multi-Az support for EBS
- Complex Master-Master MySQL setup
- Complex Master-Slave MySQL setup
- No Automatic Backup & Recovery
- No Auto-Upgrade MySQL

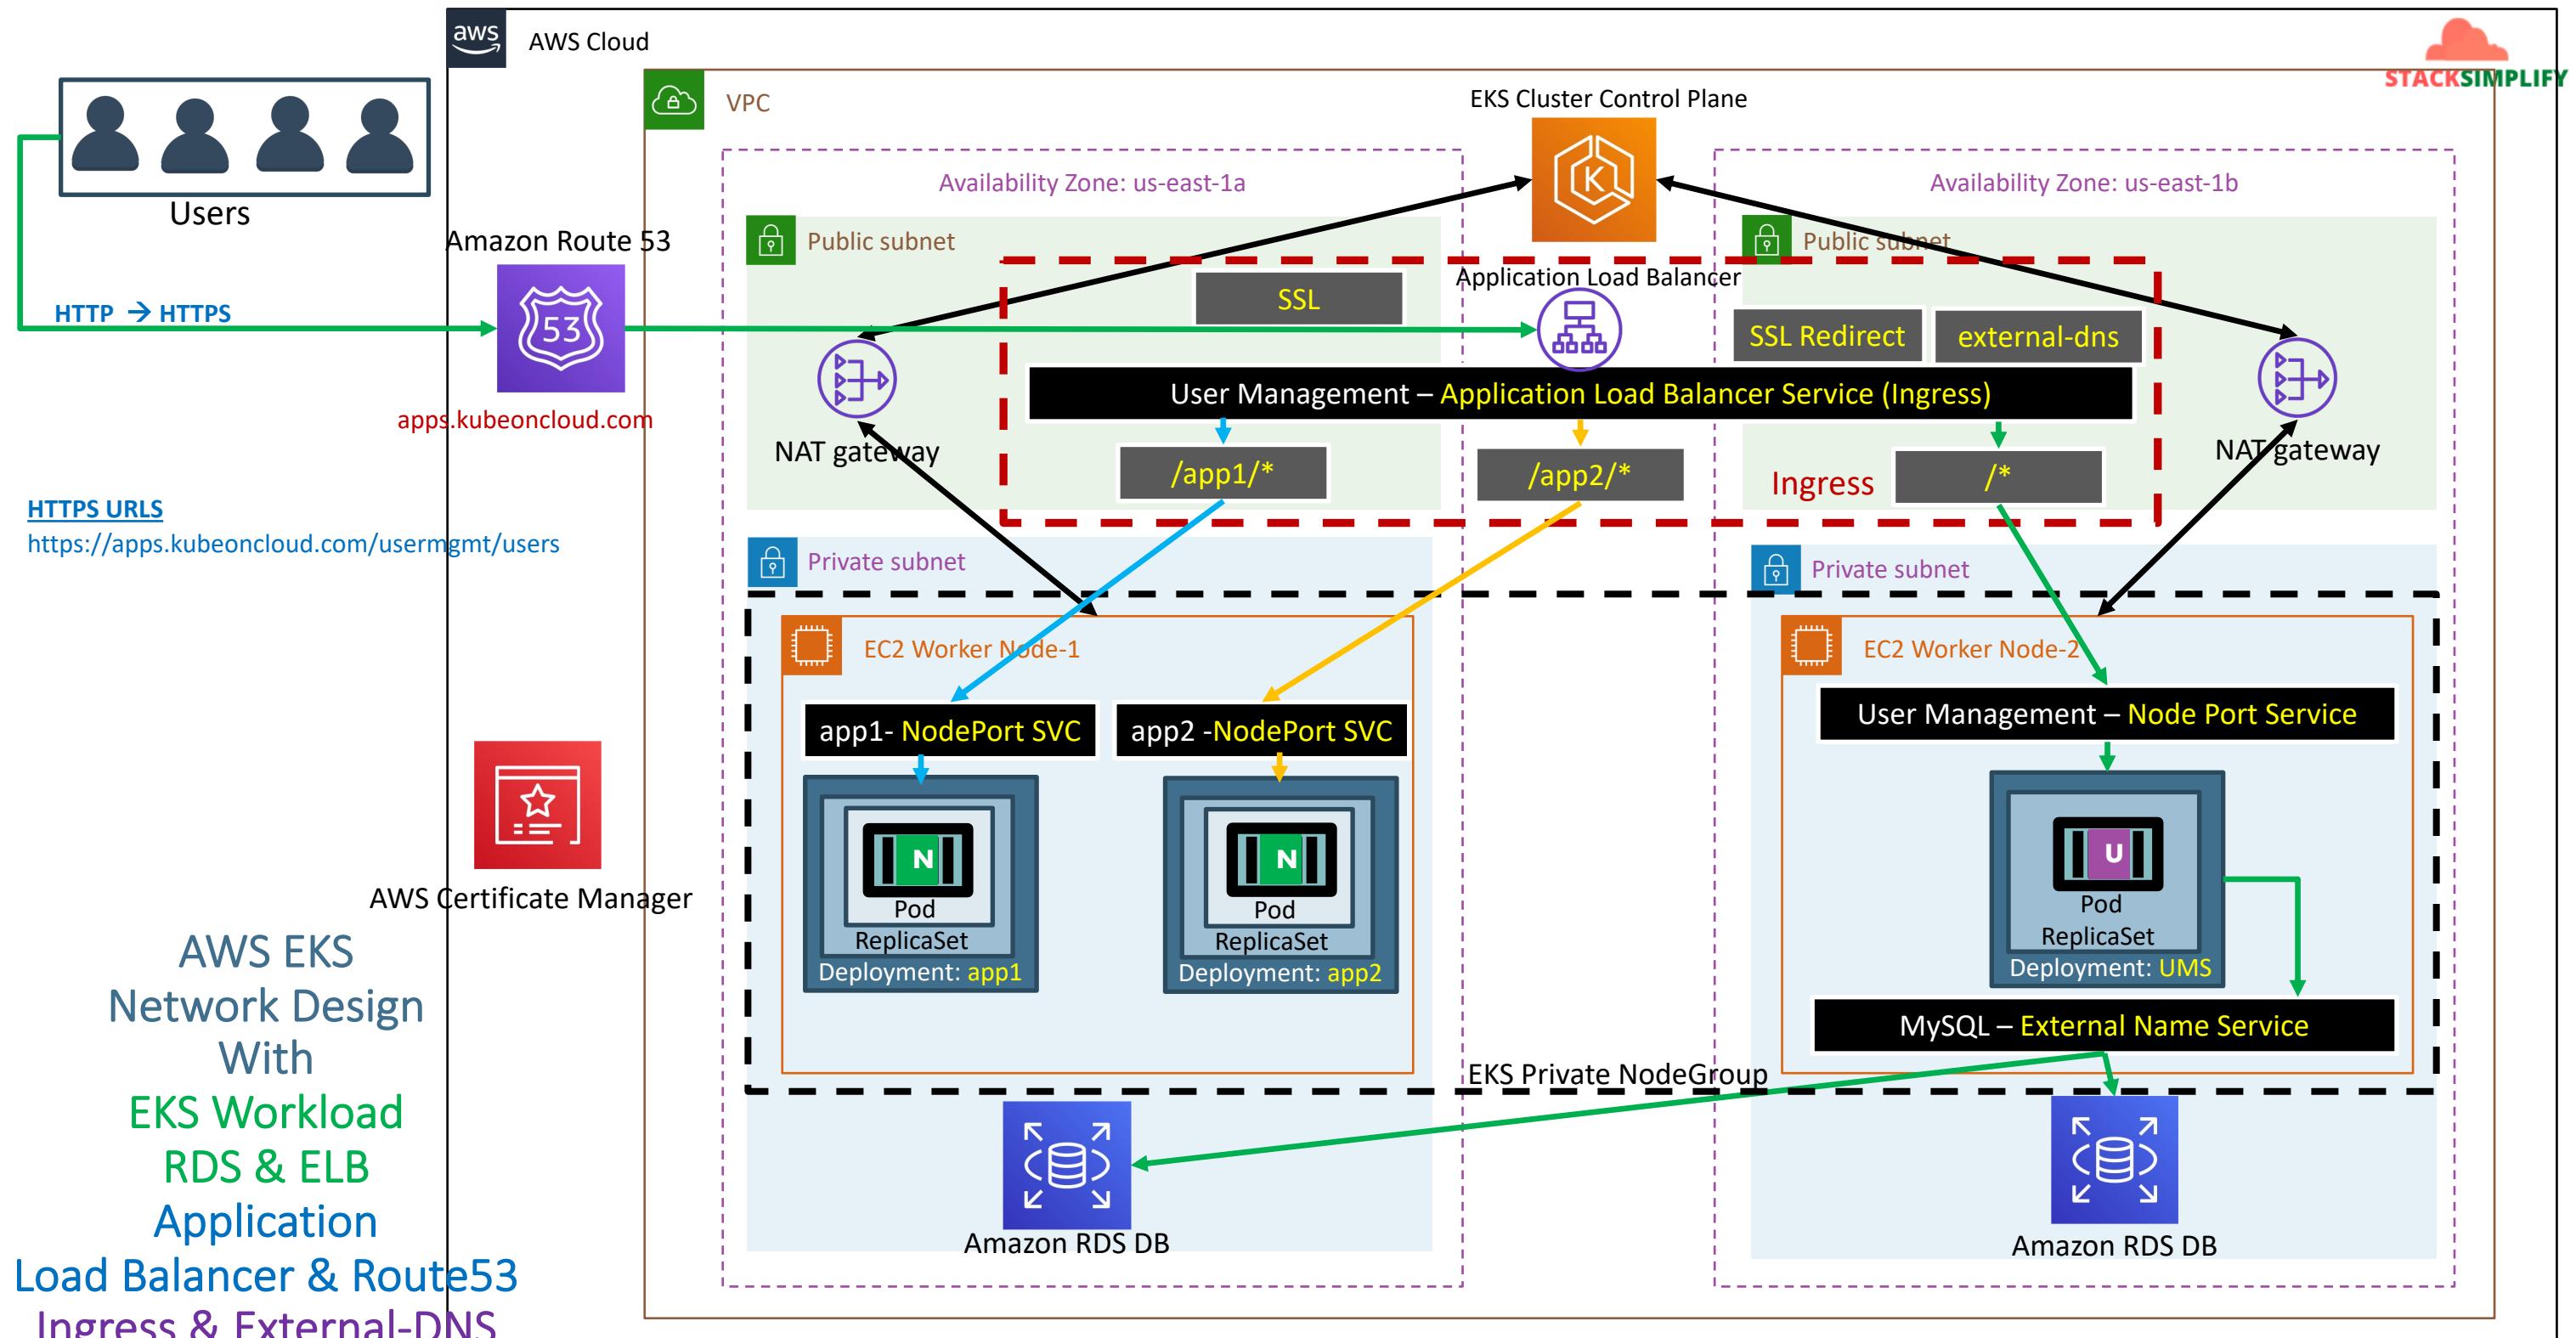




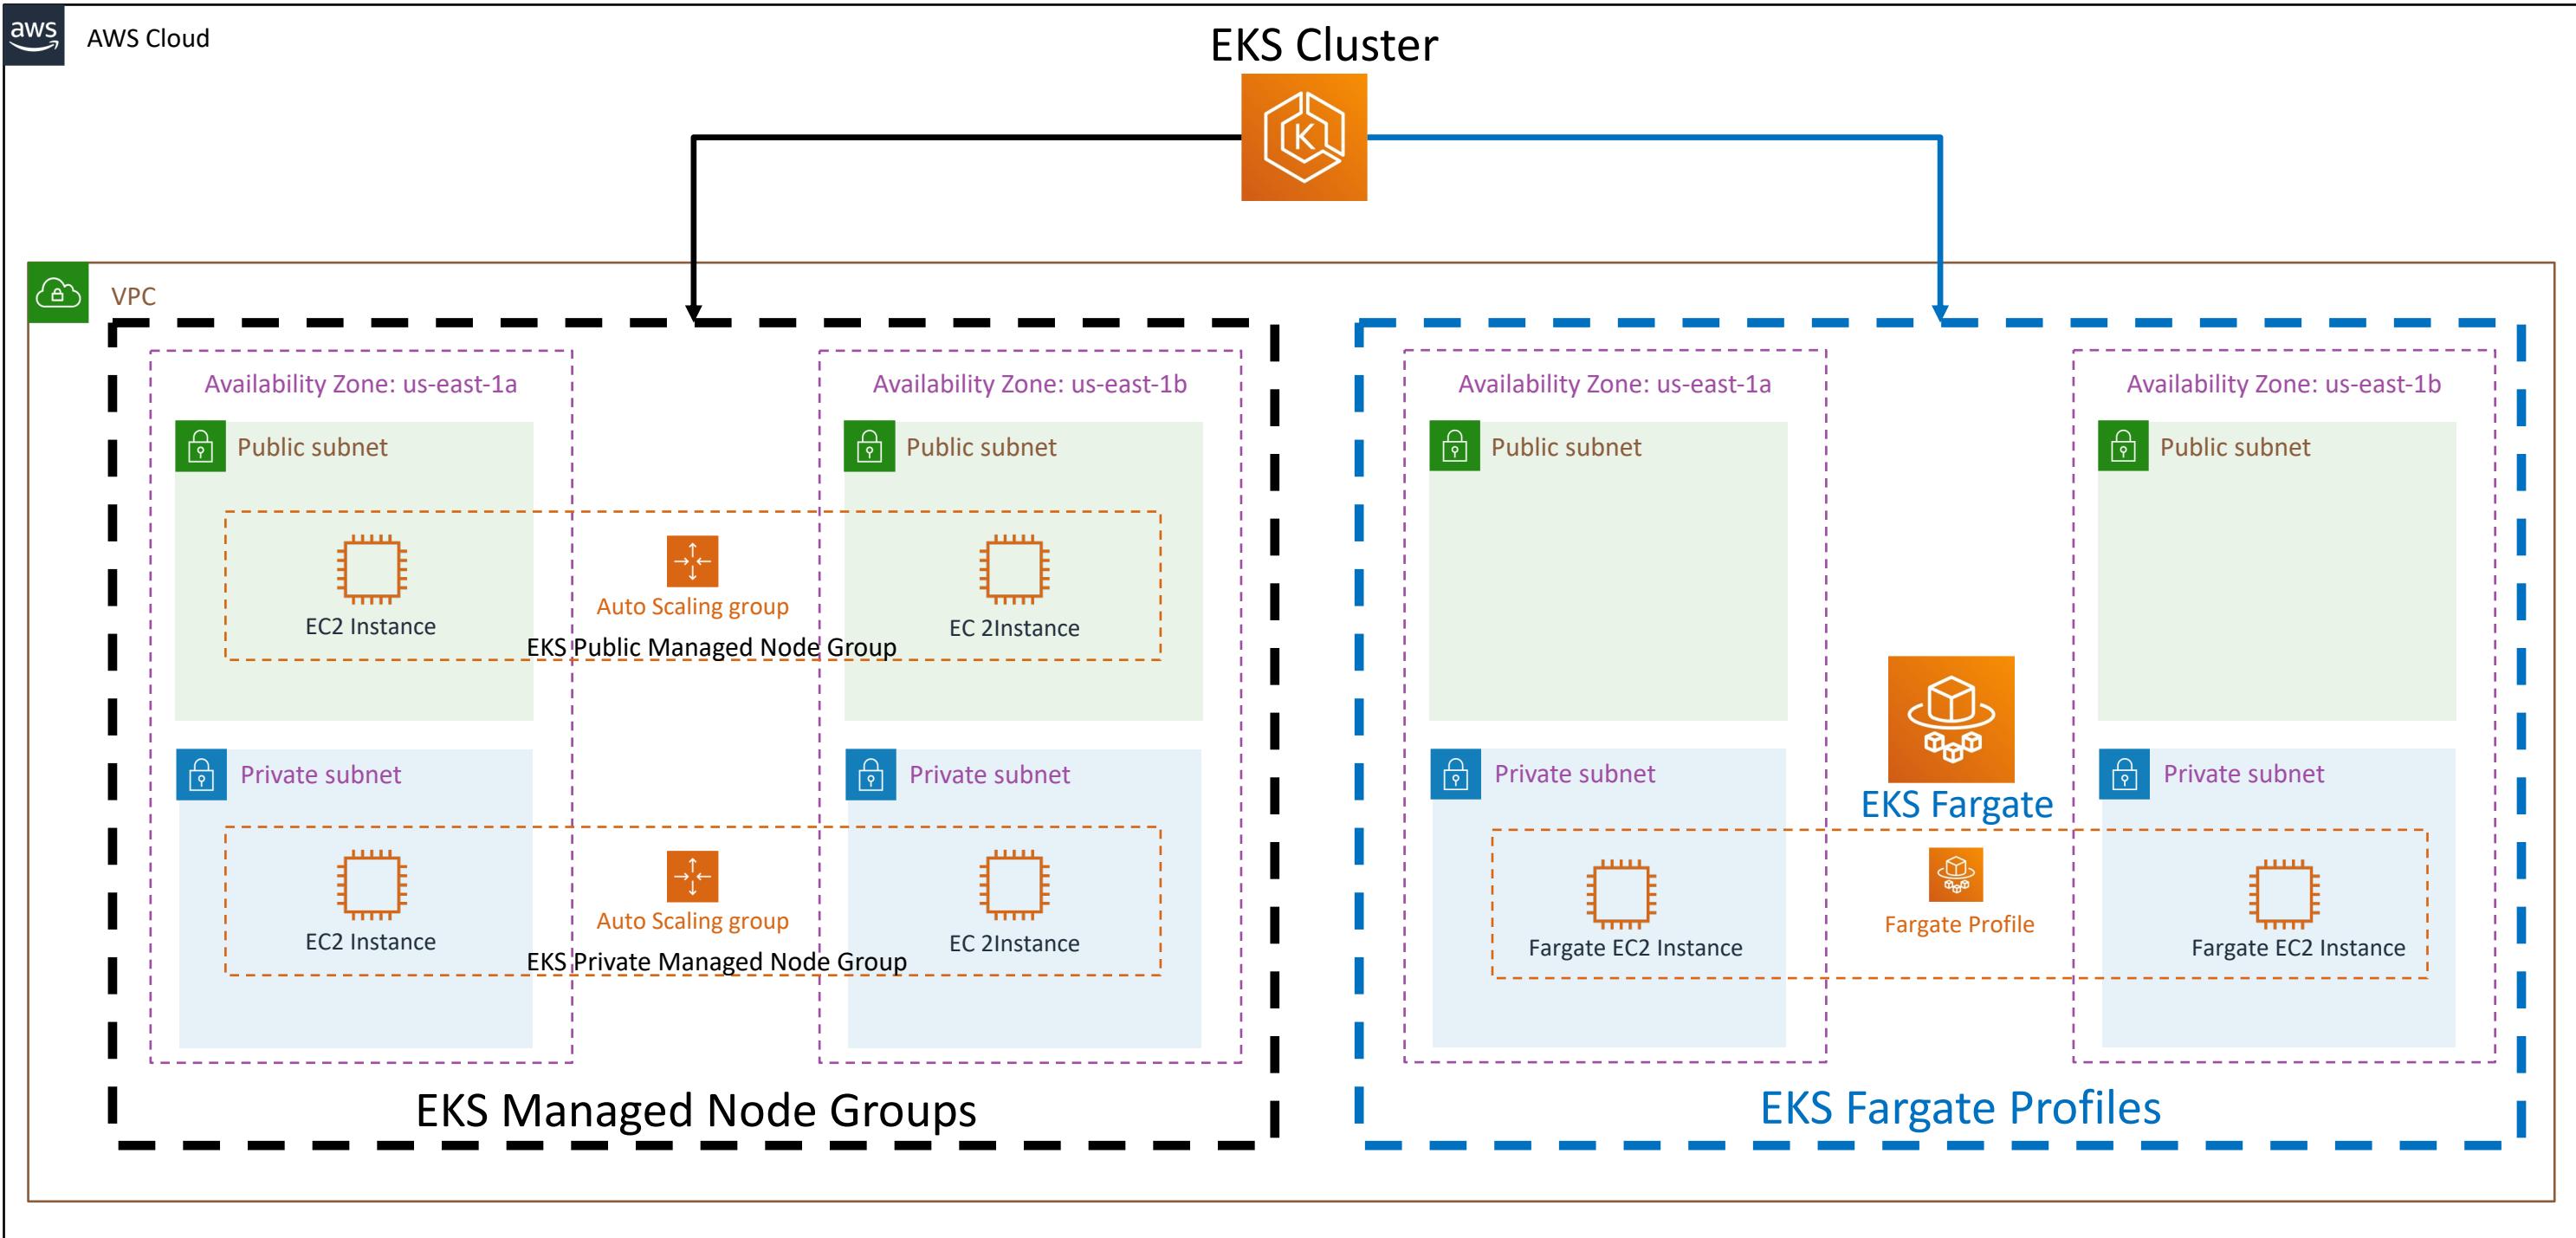




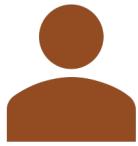




# EKS Deployment Options - Mixed

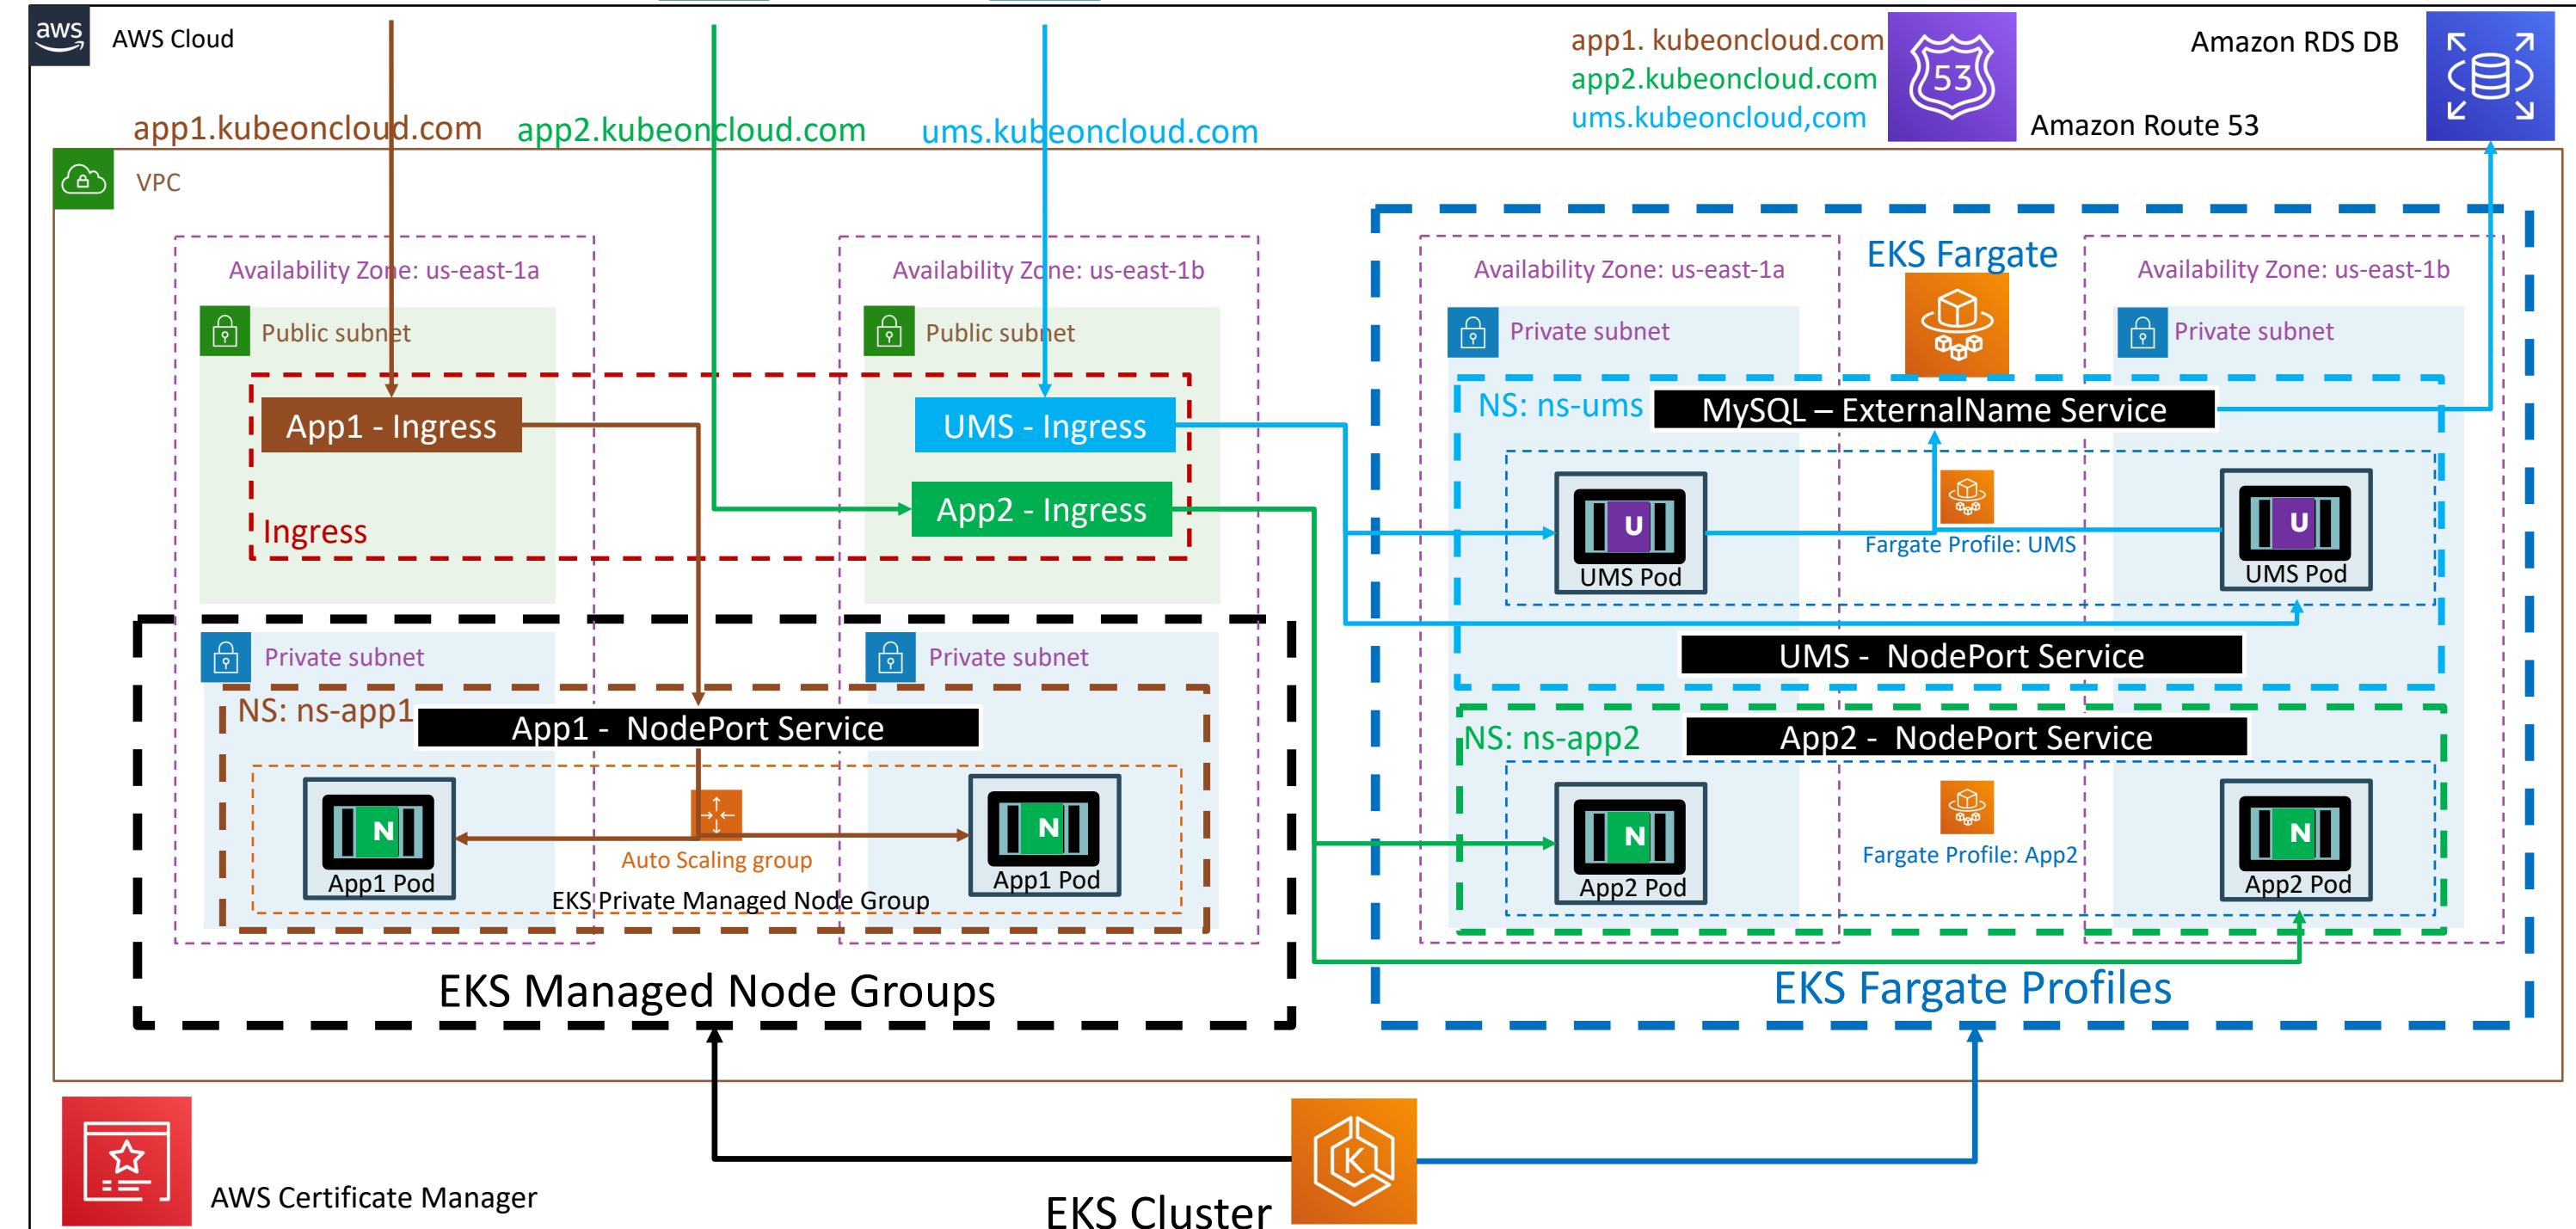


Users



# EKS Deployment Options – Mixed Mode - 3 Apps

StackSimplify

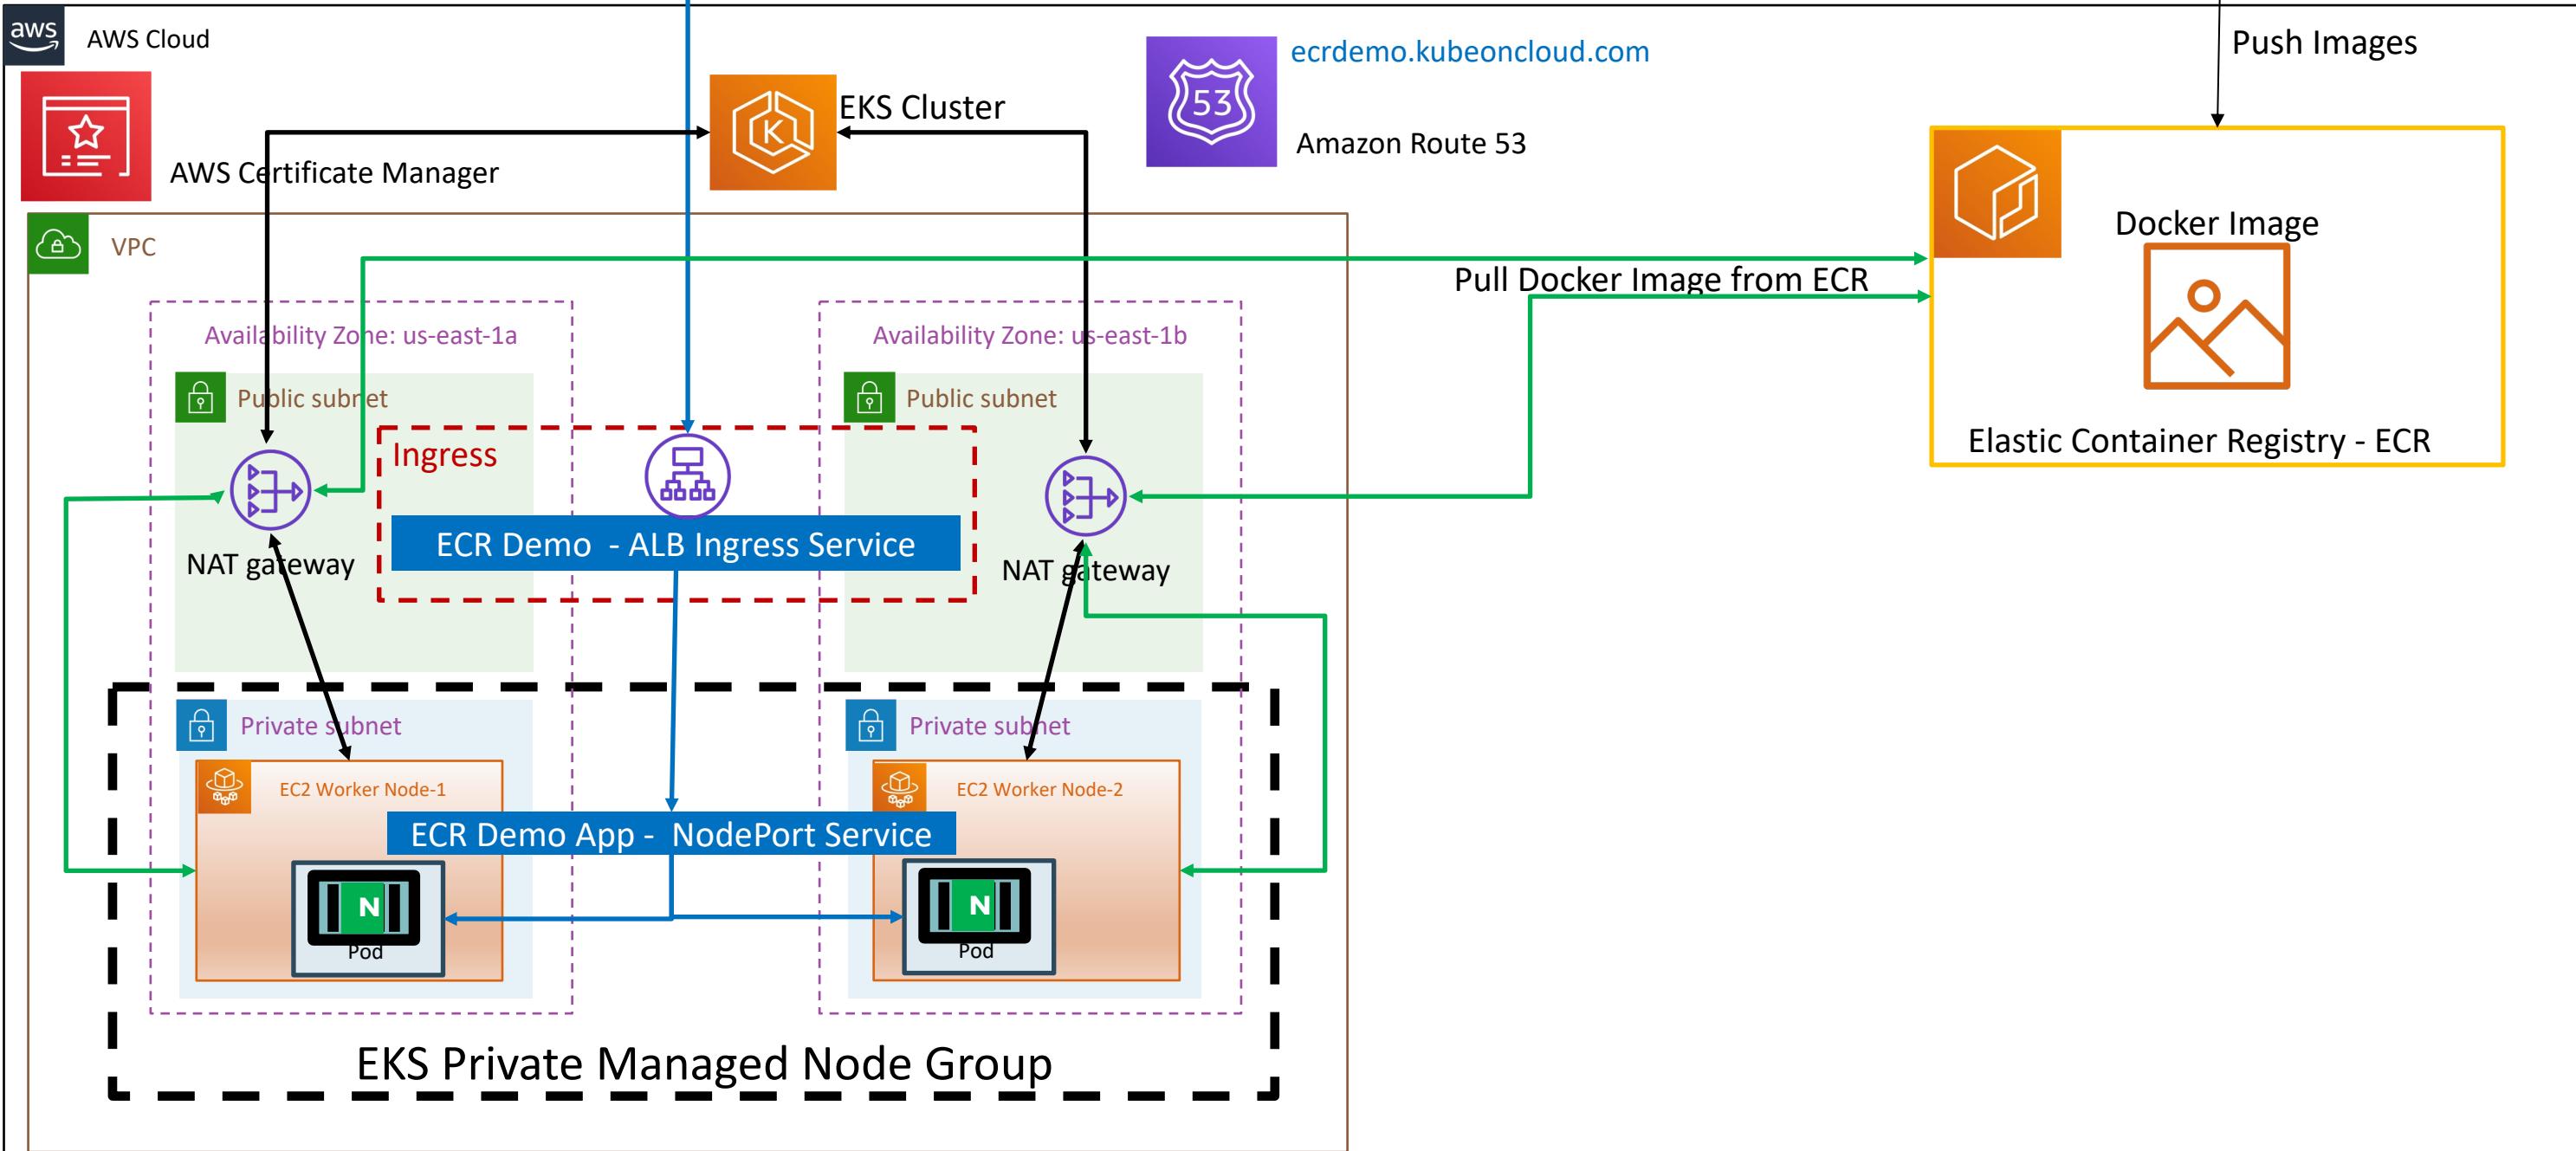


# EKS & ECR

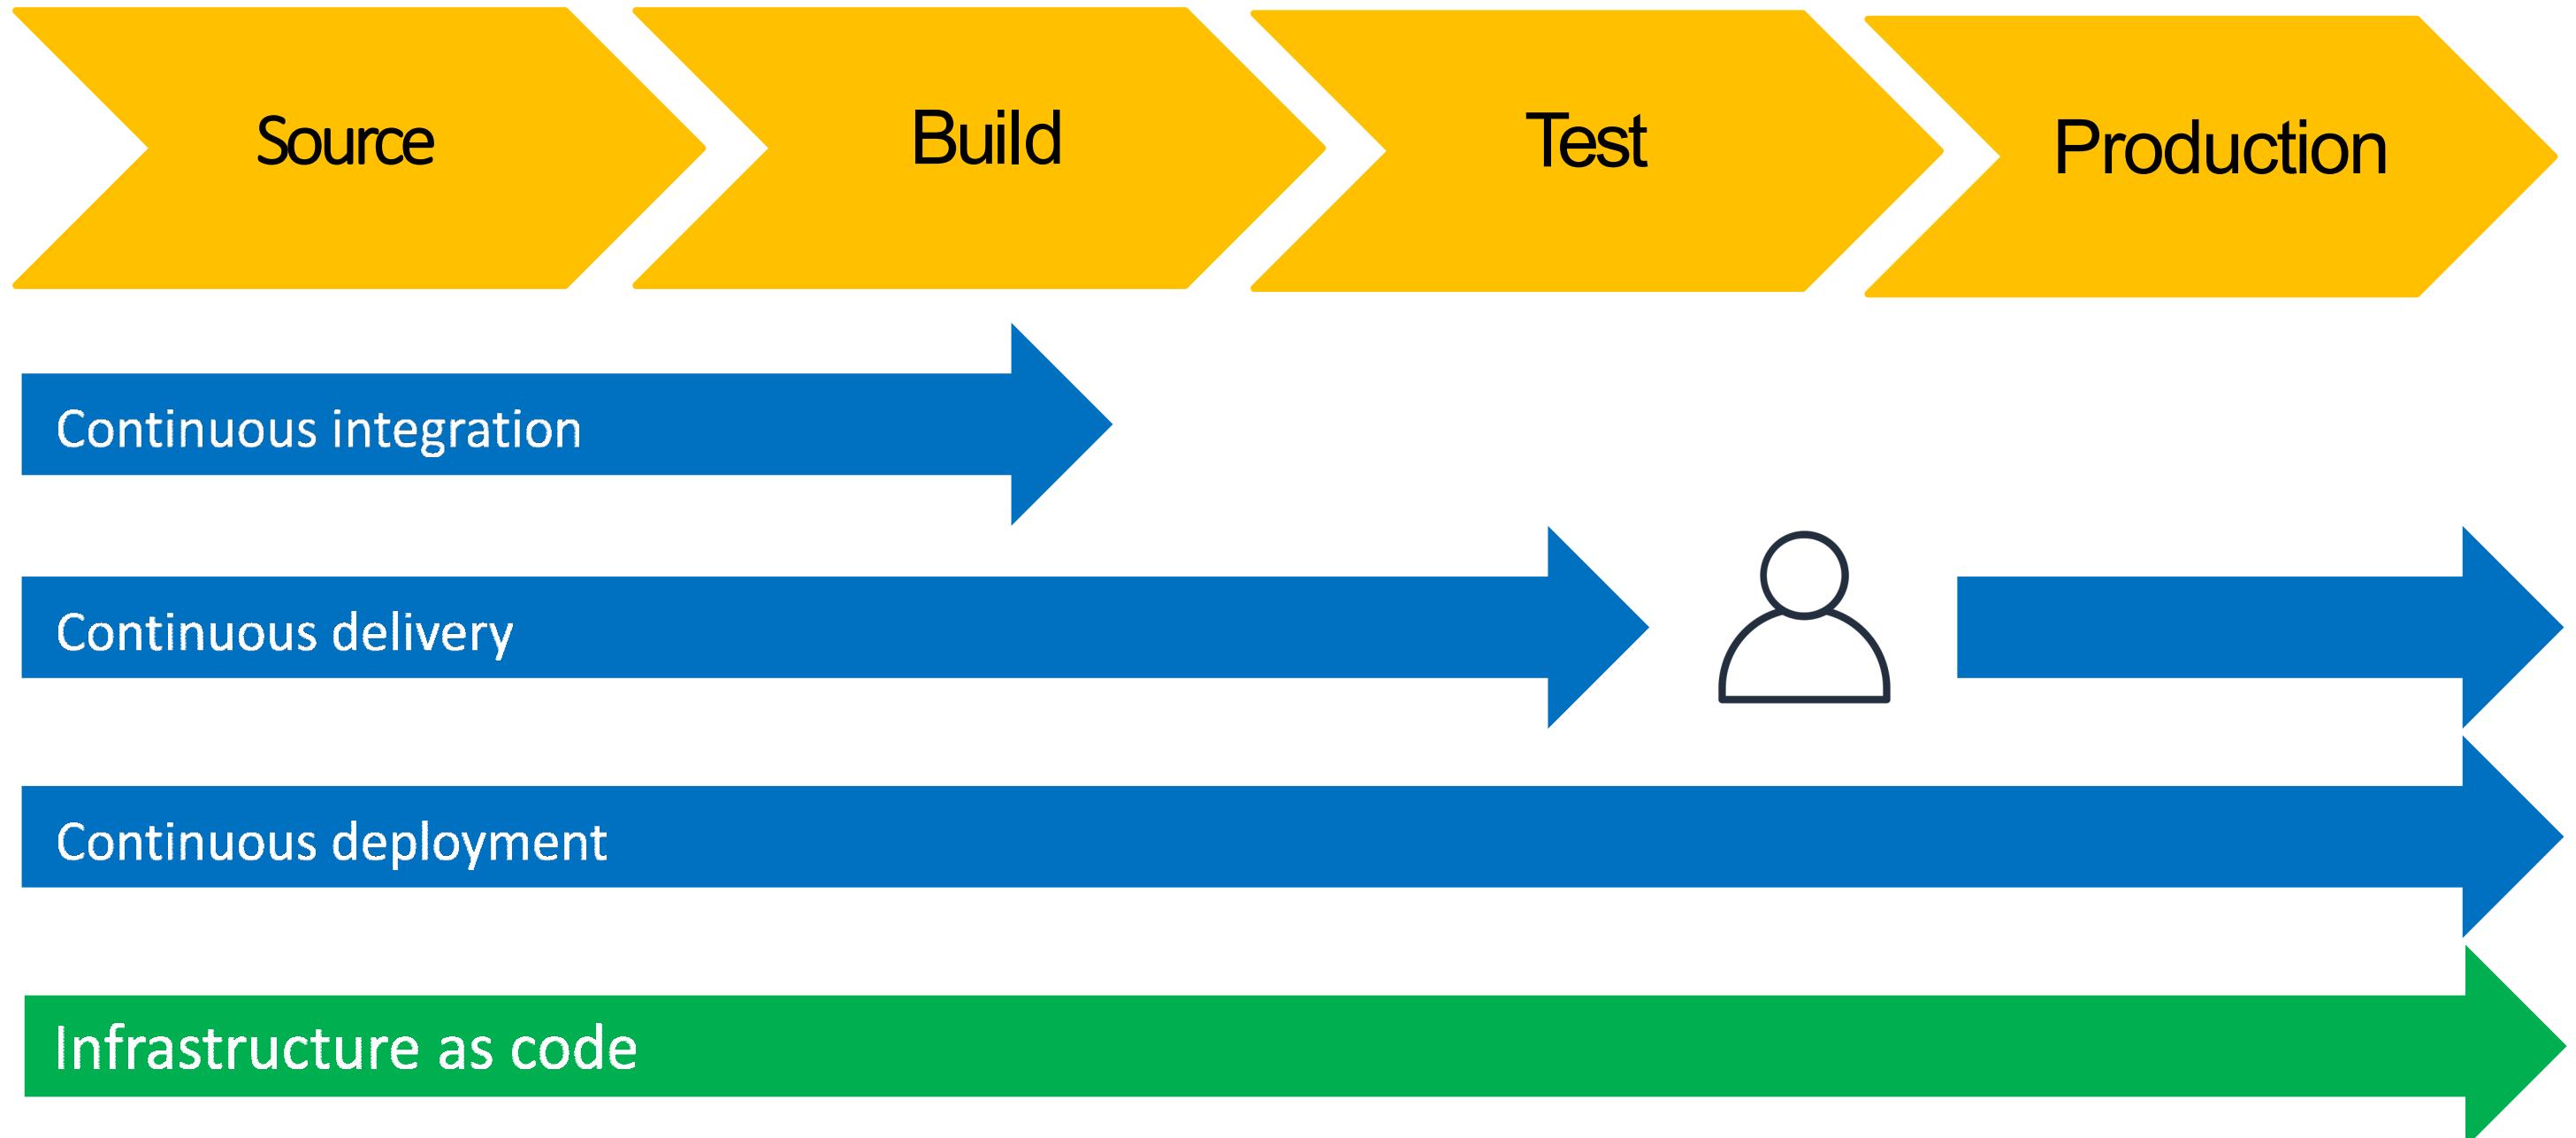
Developer  
StackSimplify



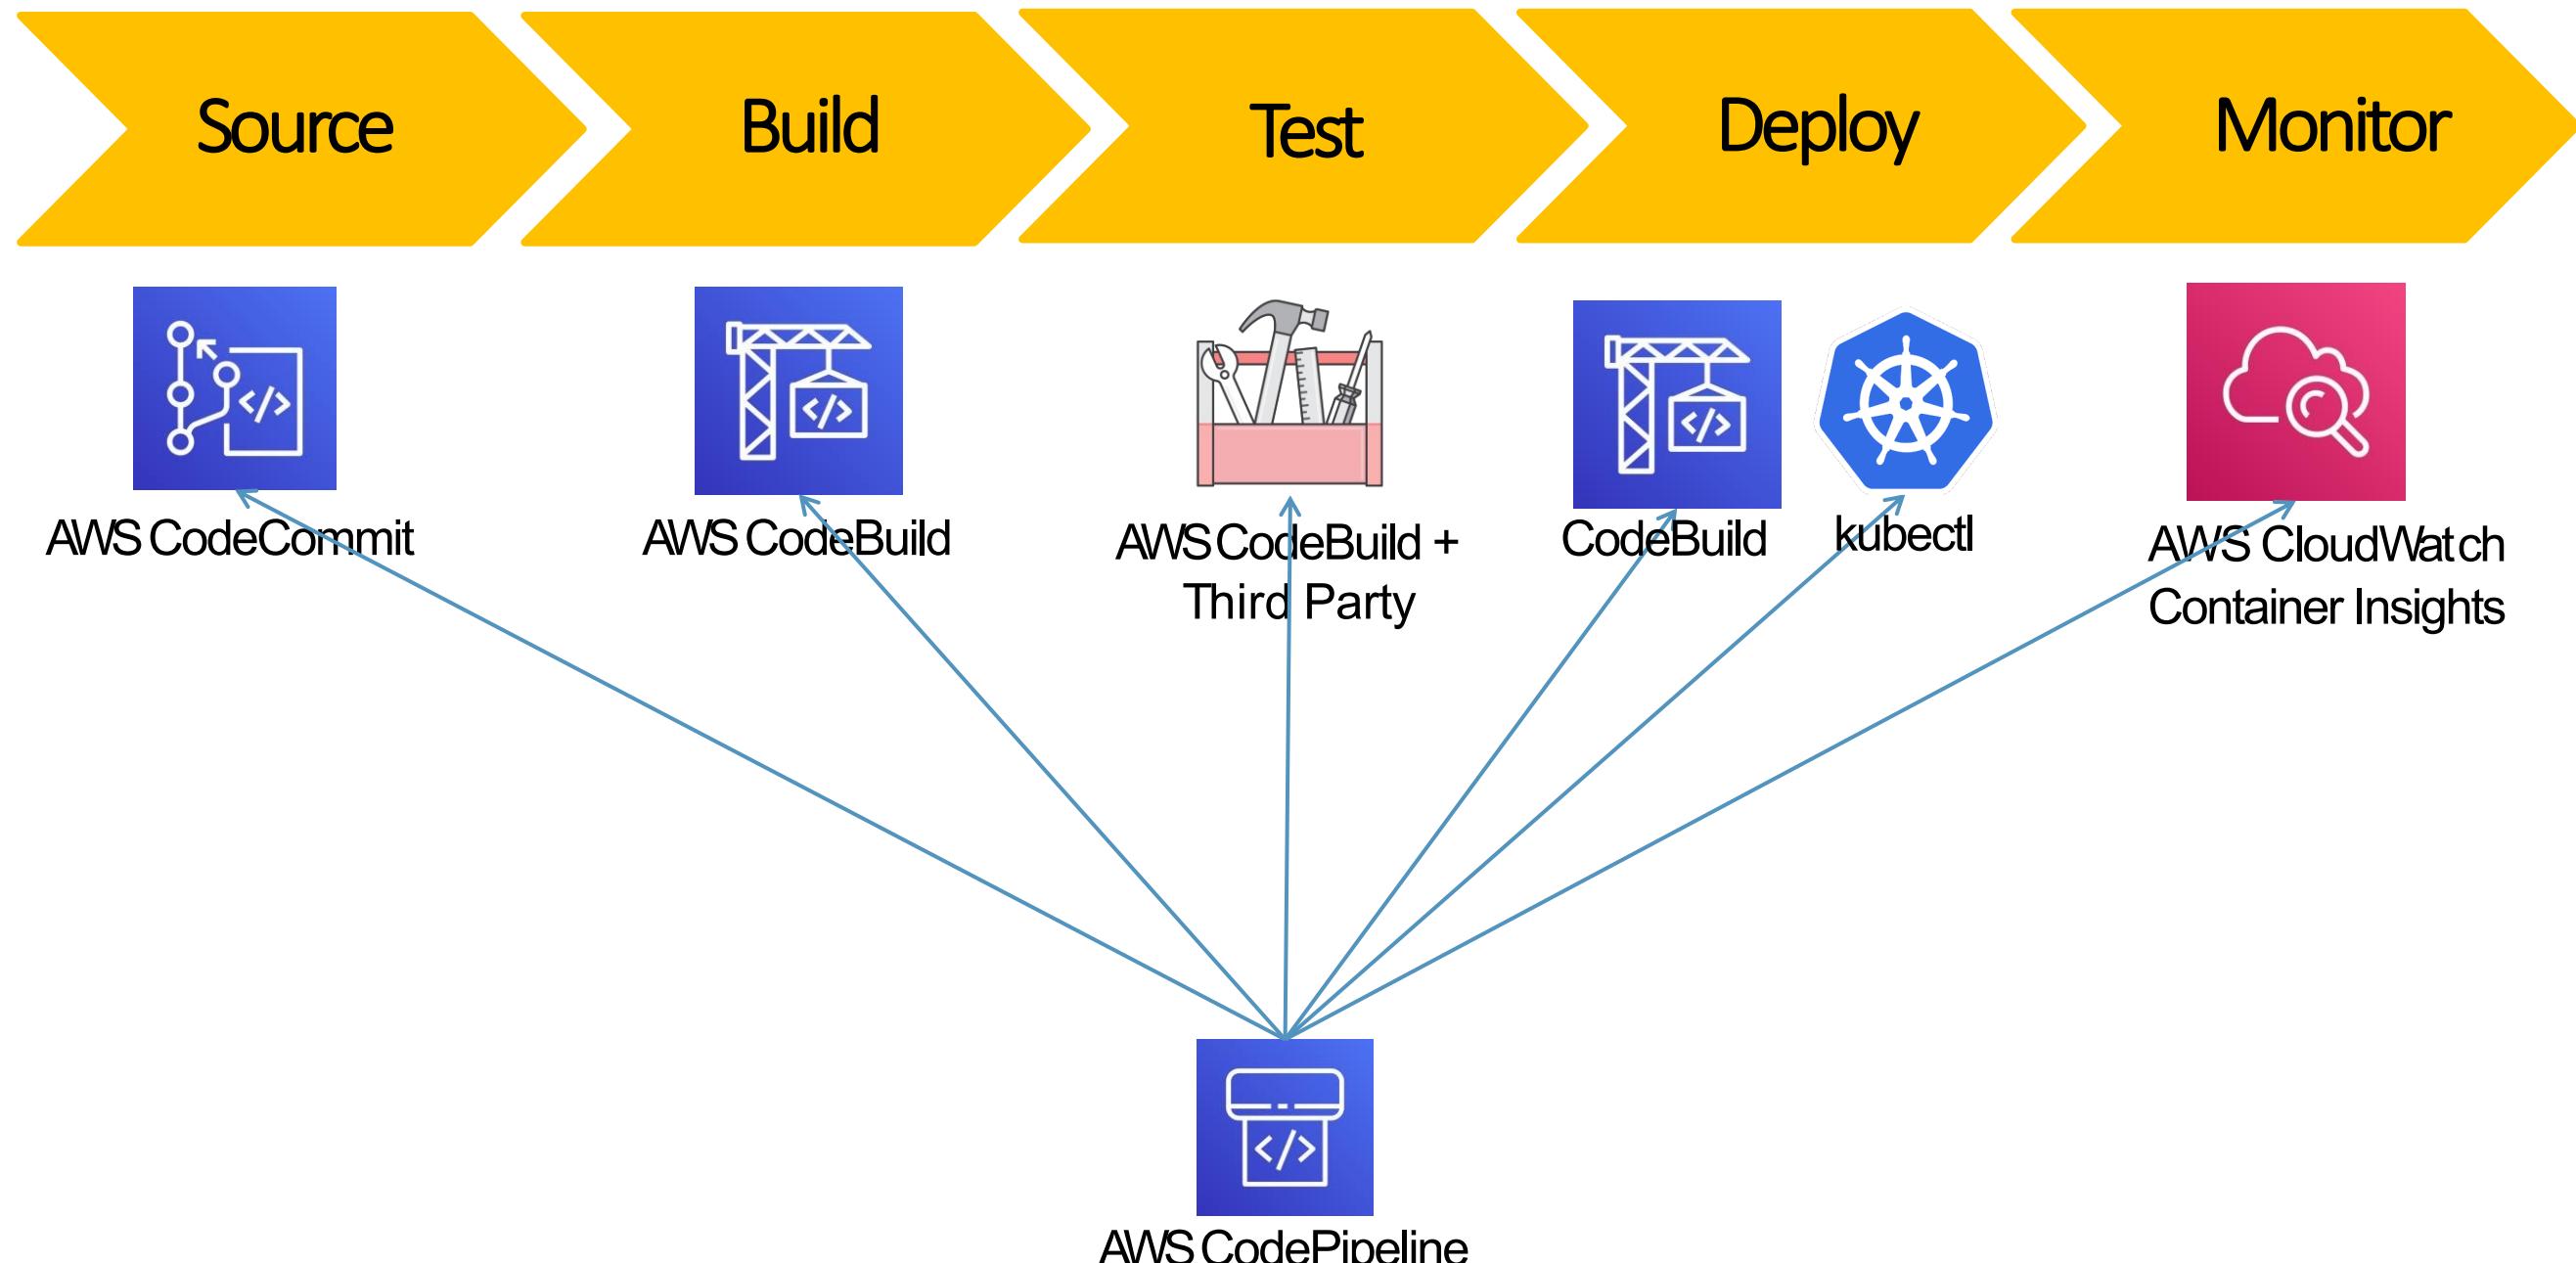
http://ecrdemo.kubeoncloud.com



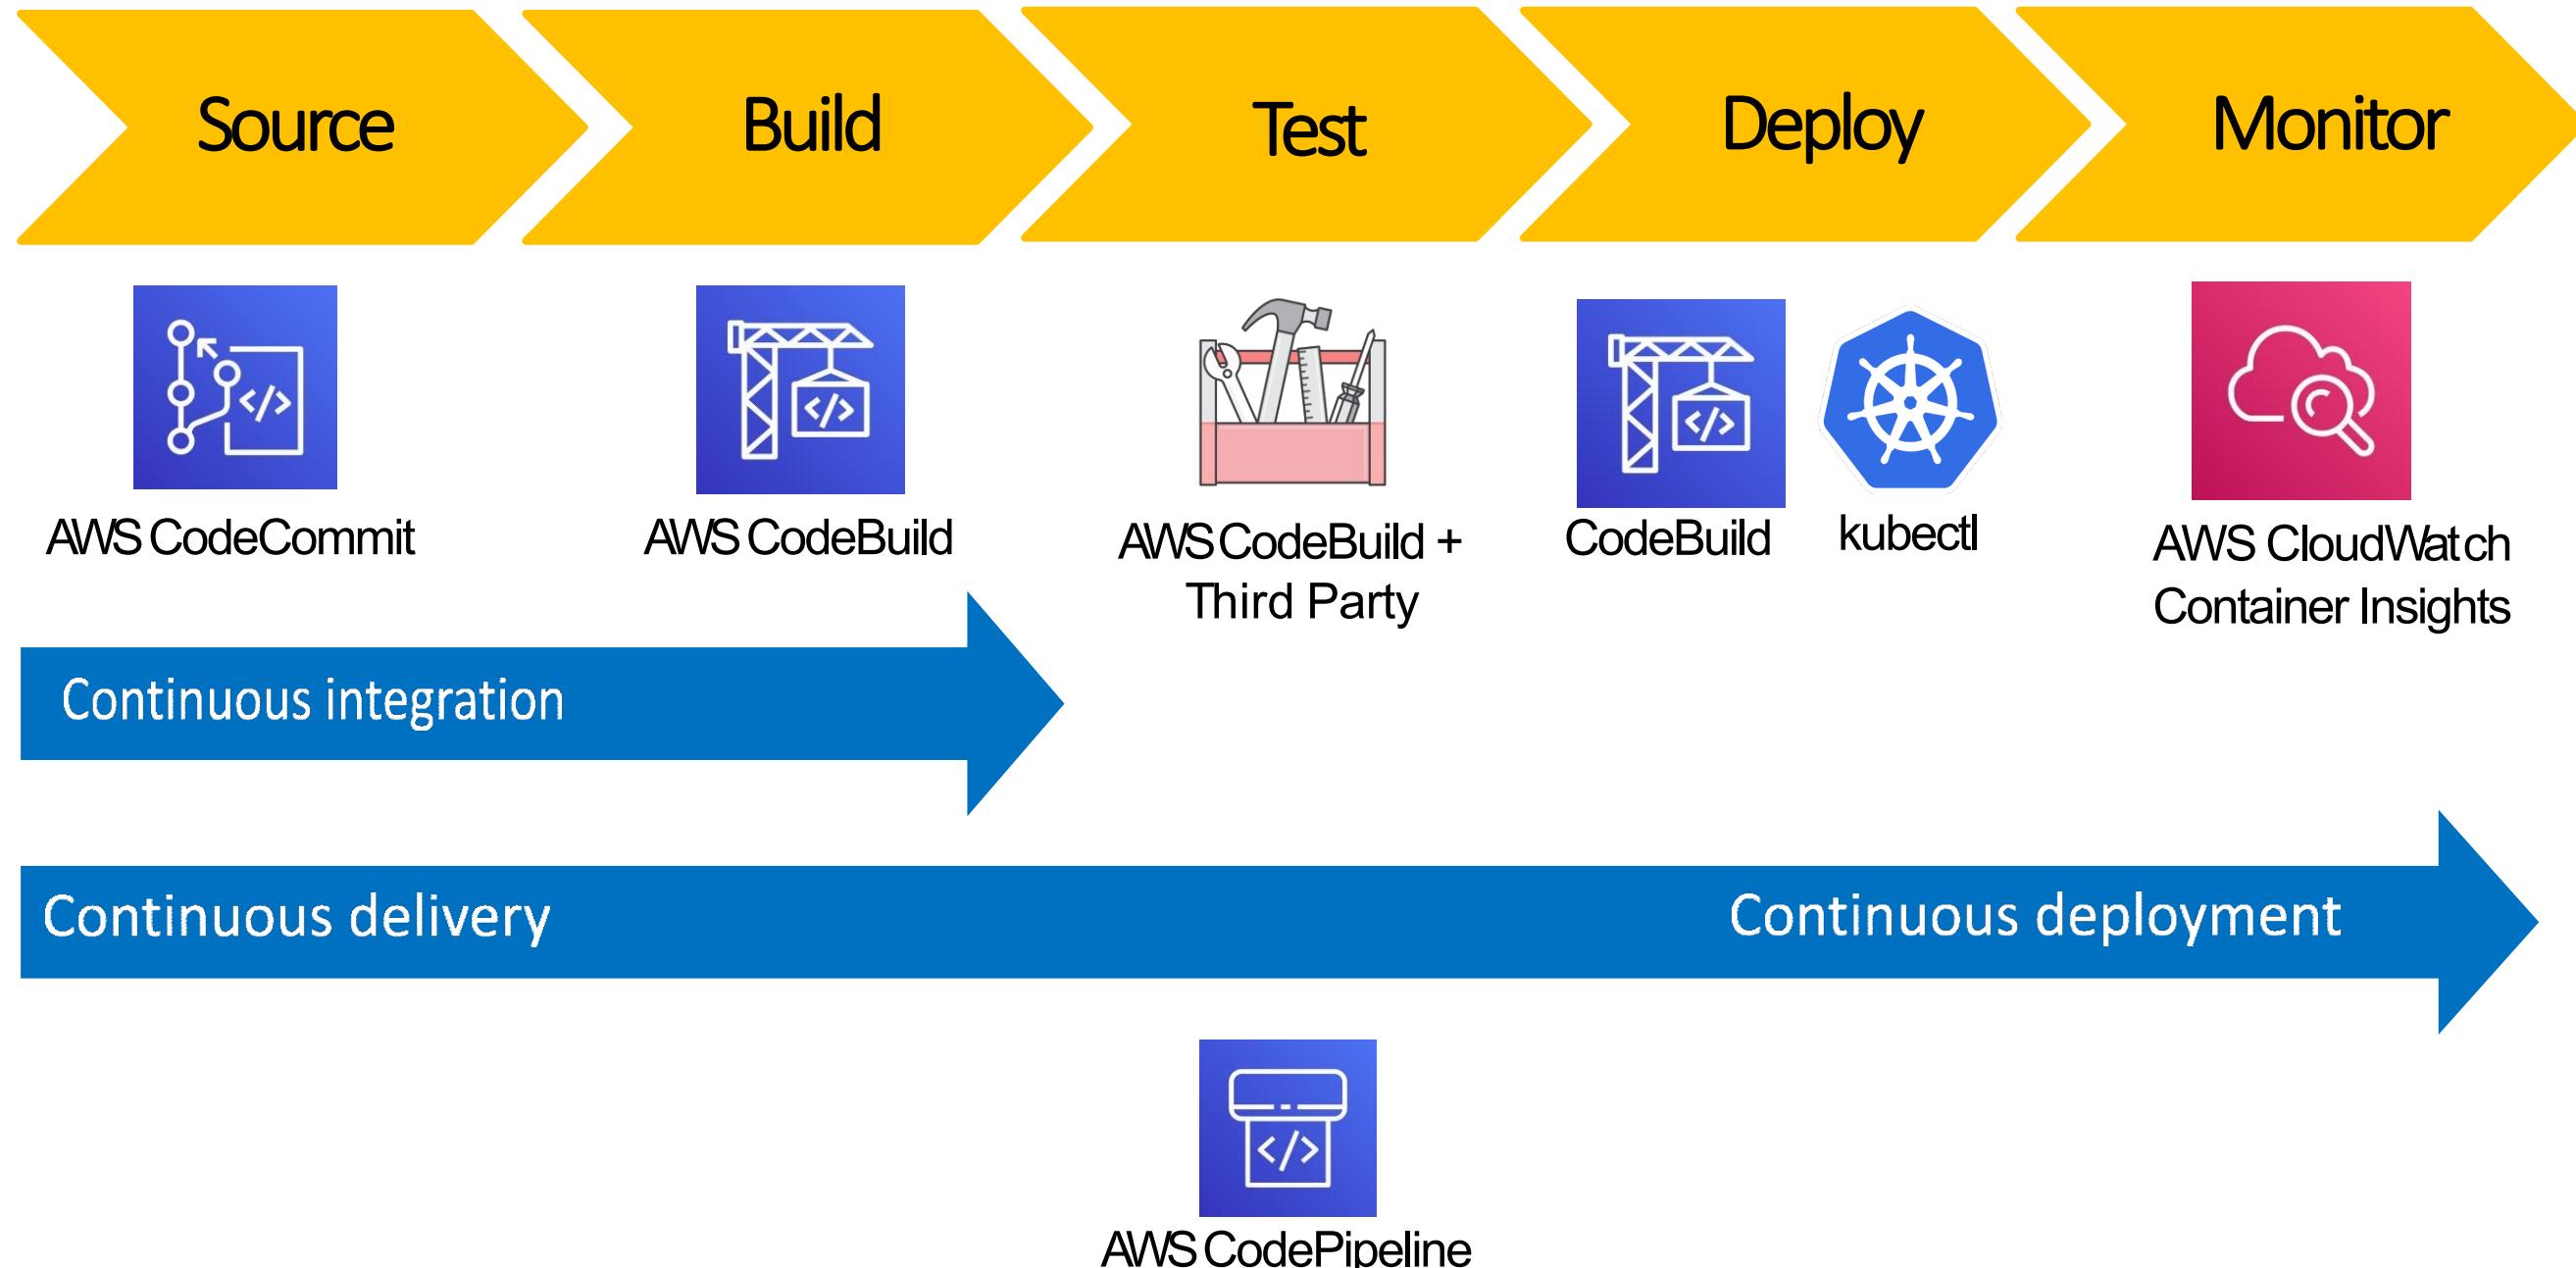
# Stages in Release Process



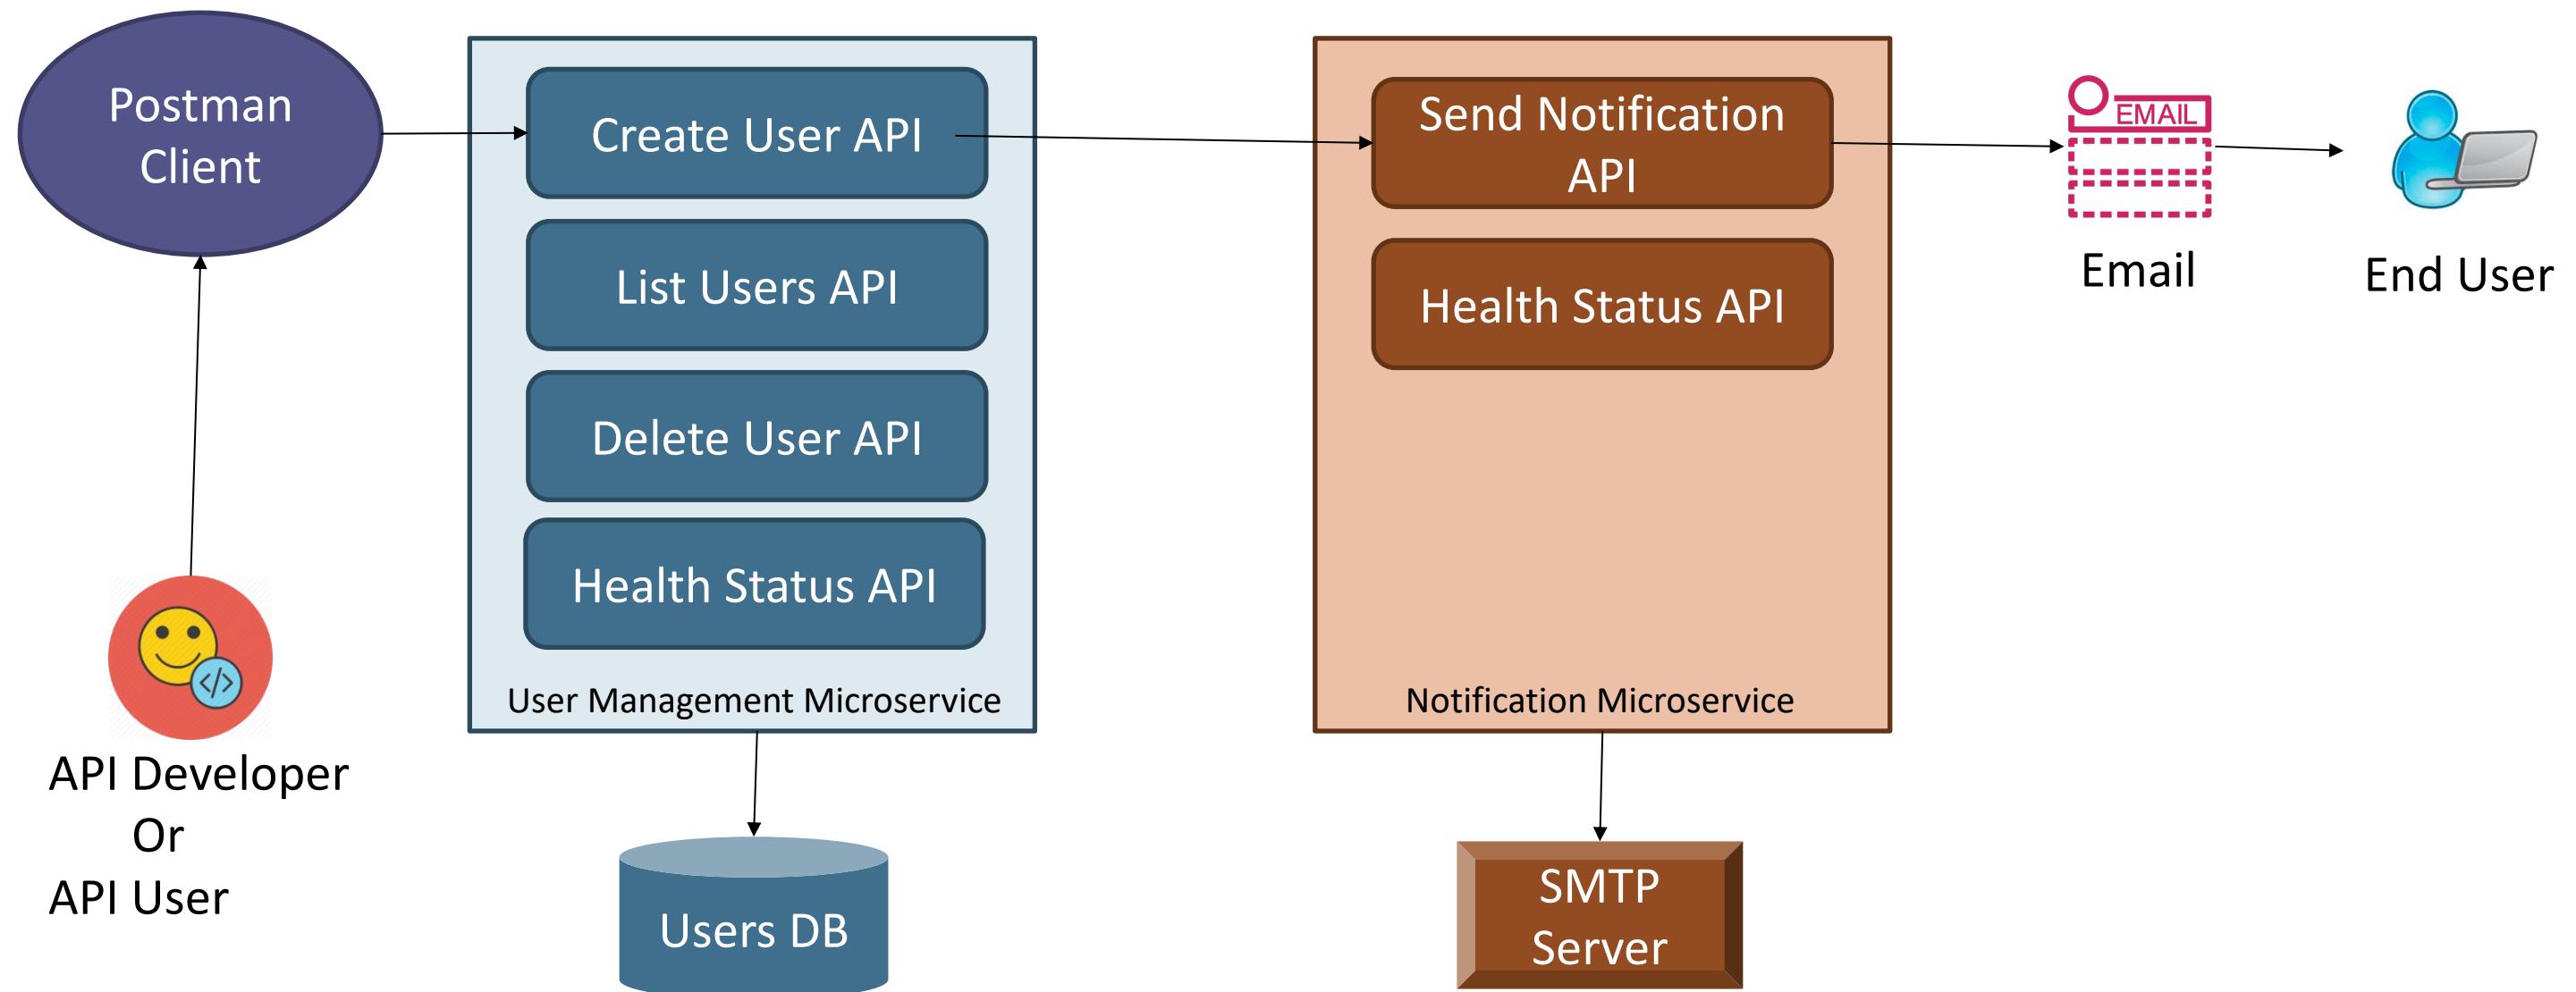
# AWS Developer Tools or Code Services



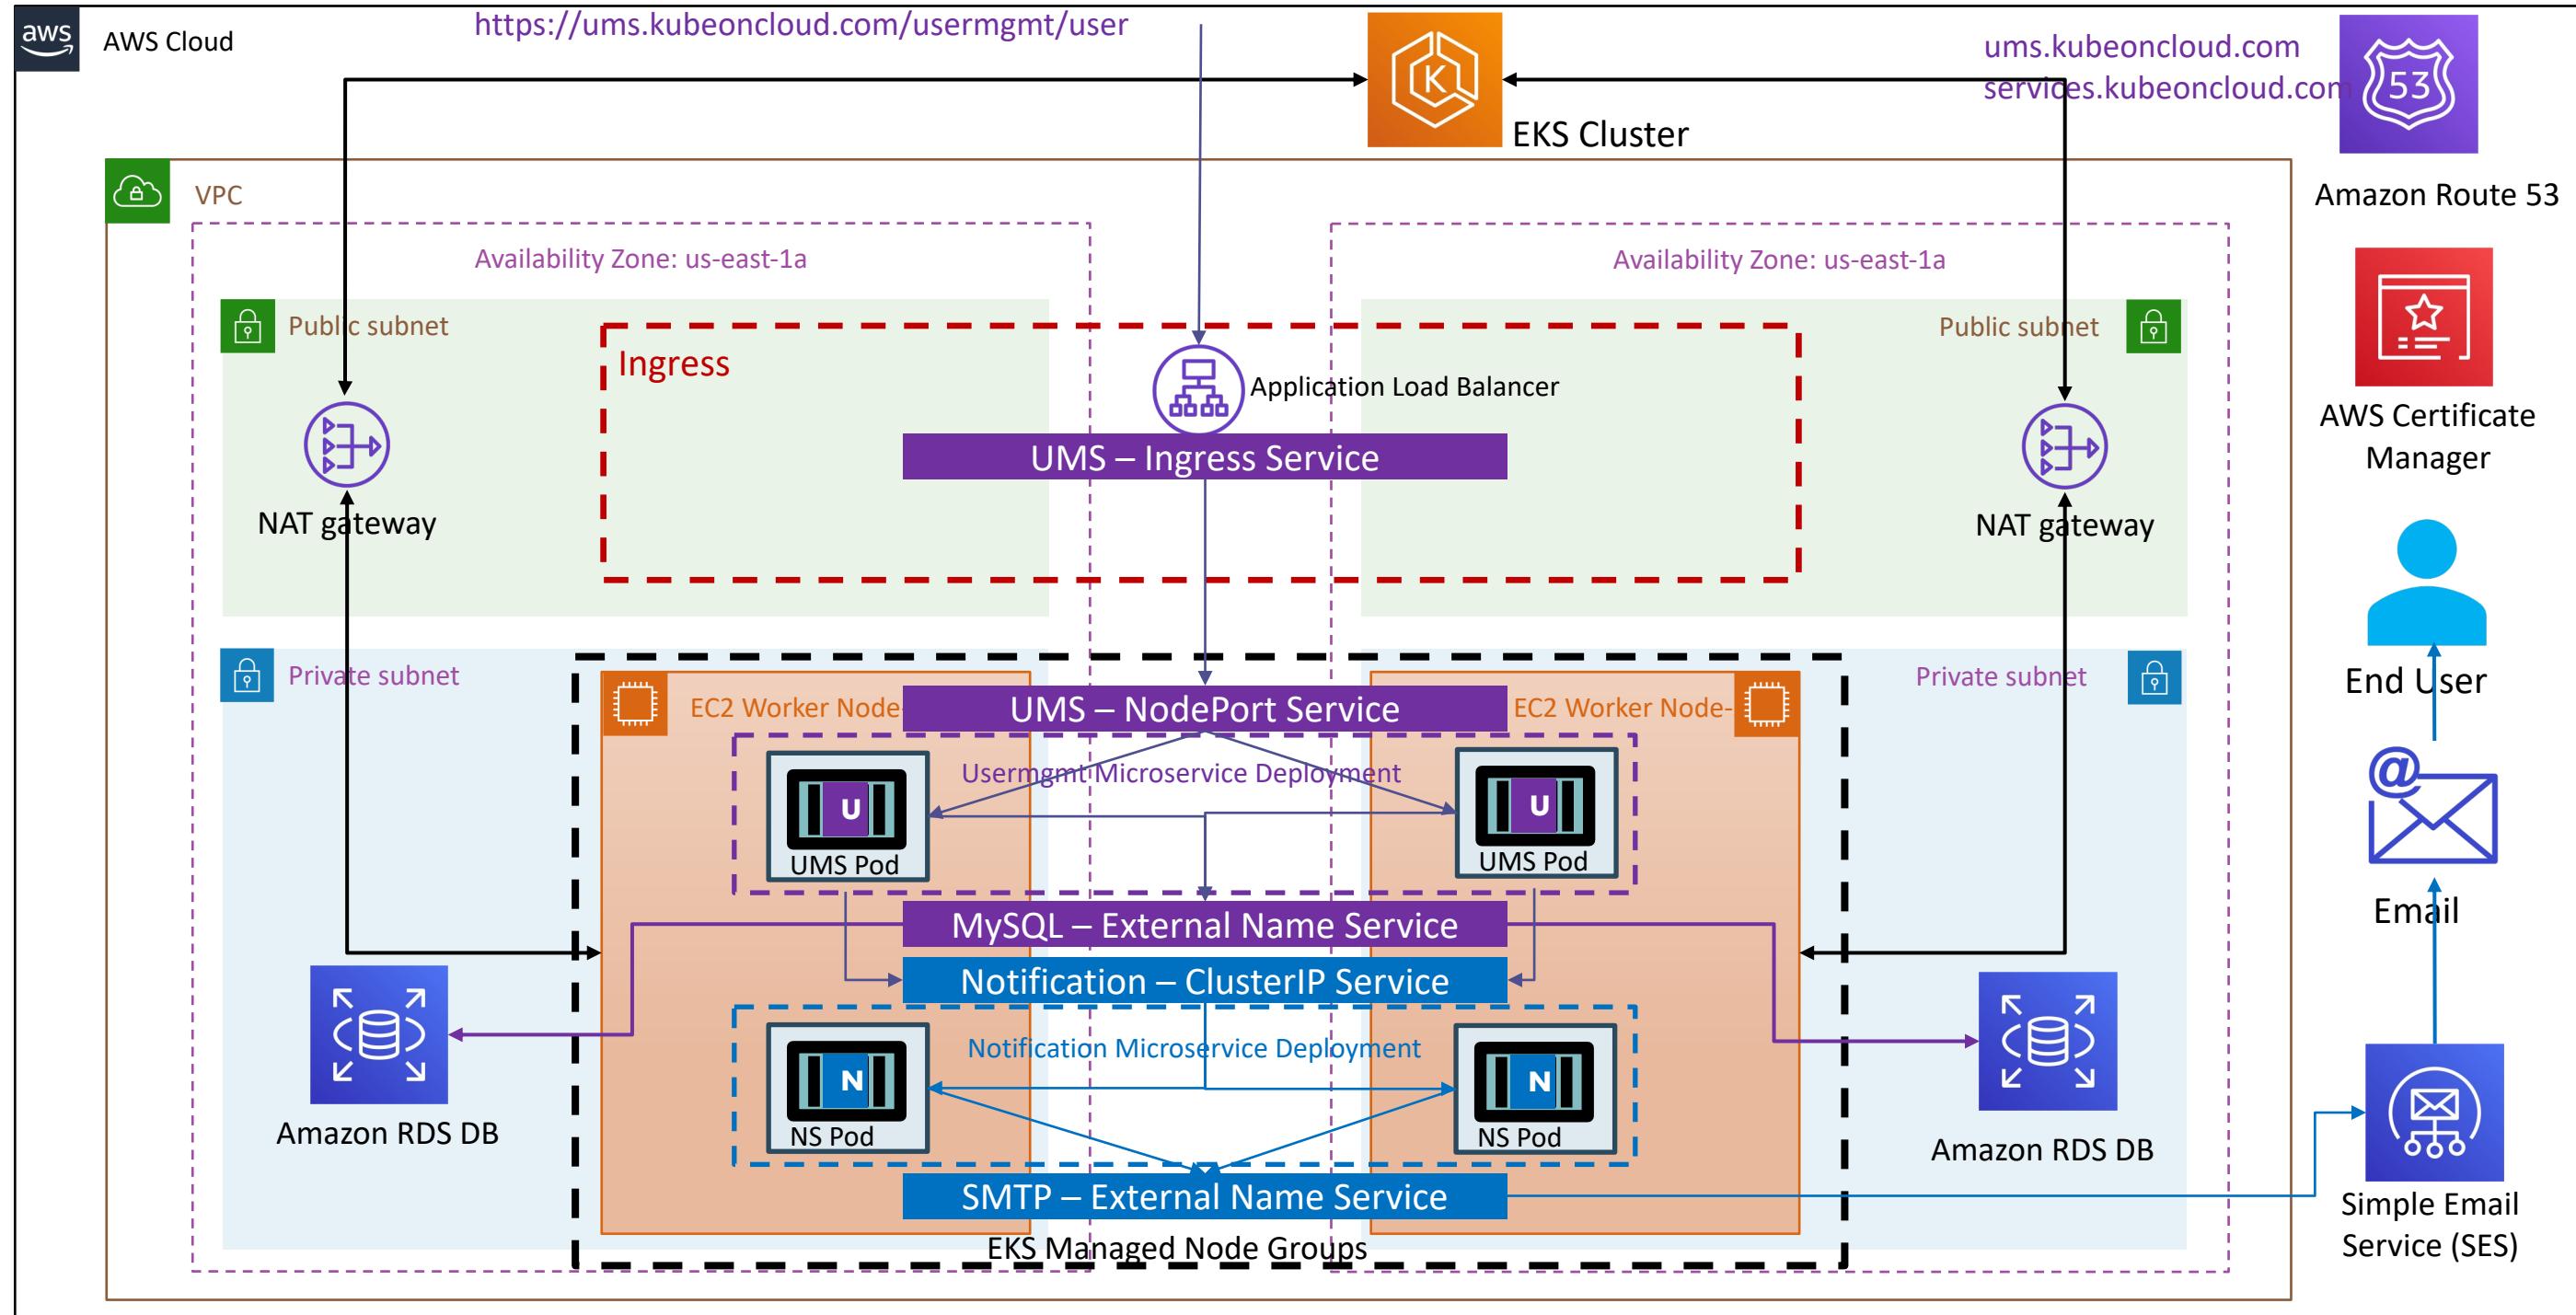
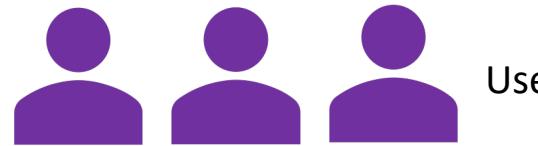
# AWS Developer Tools or Code Services



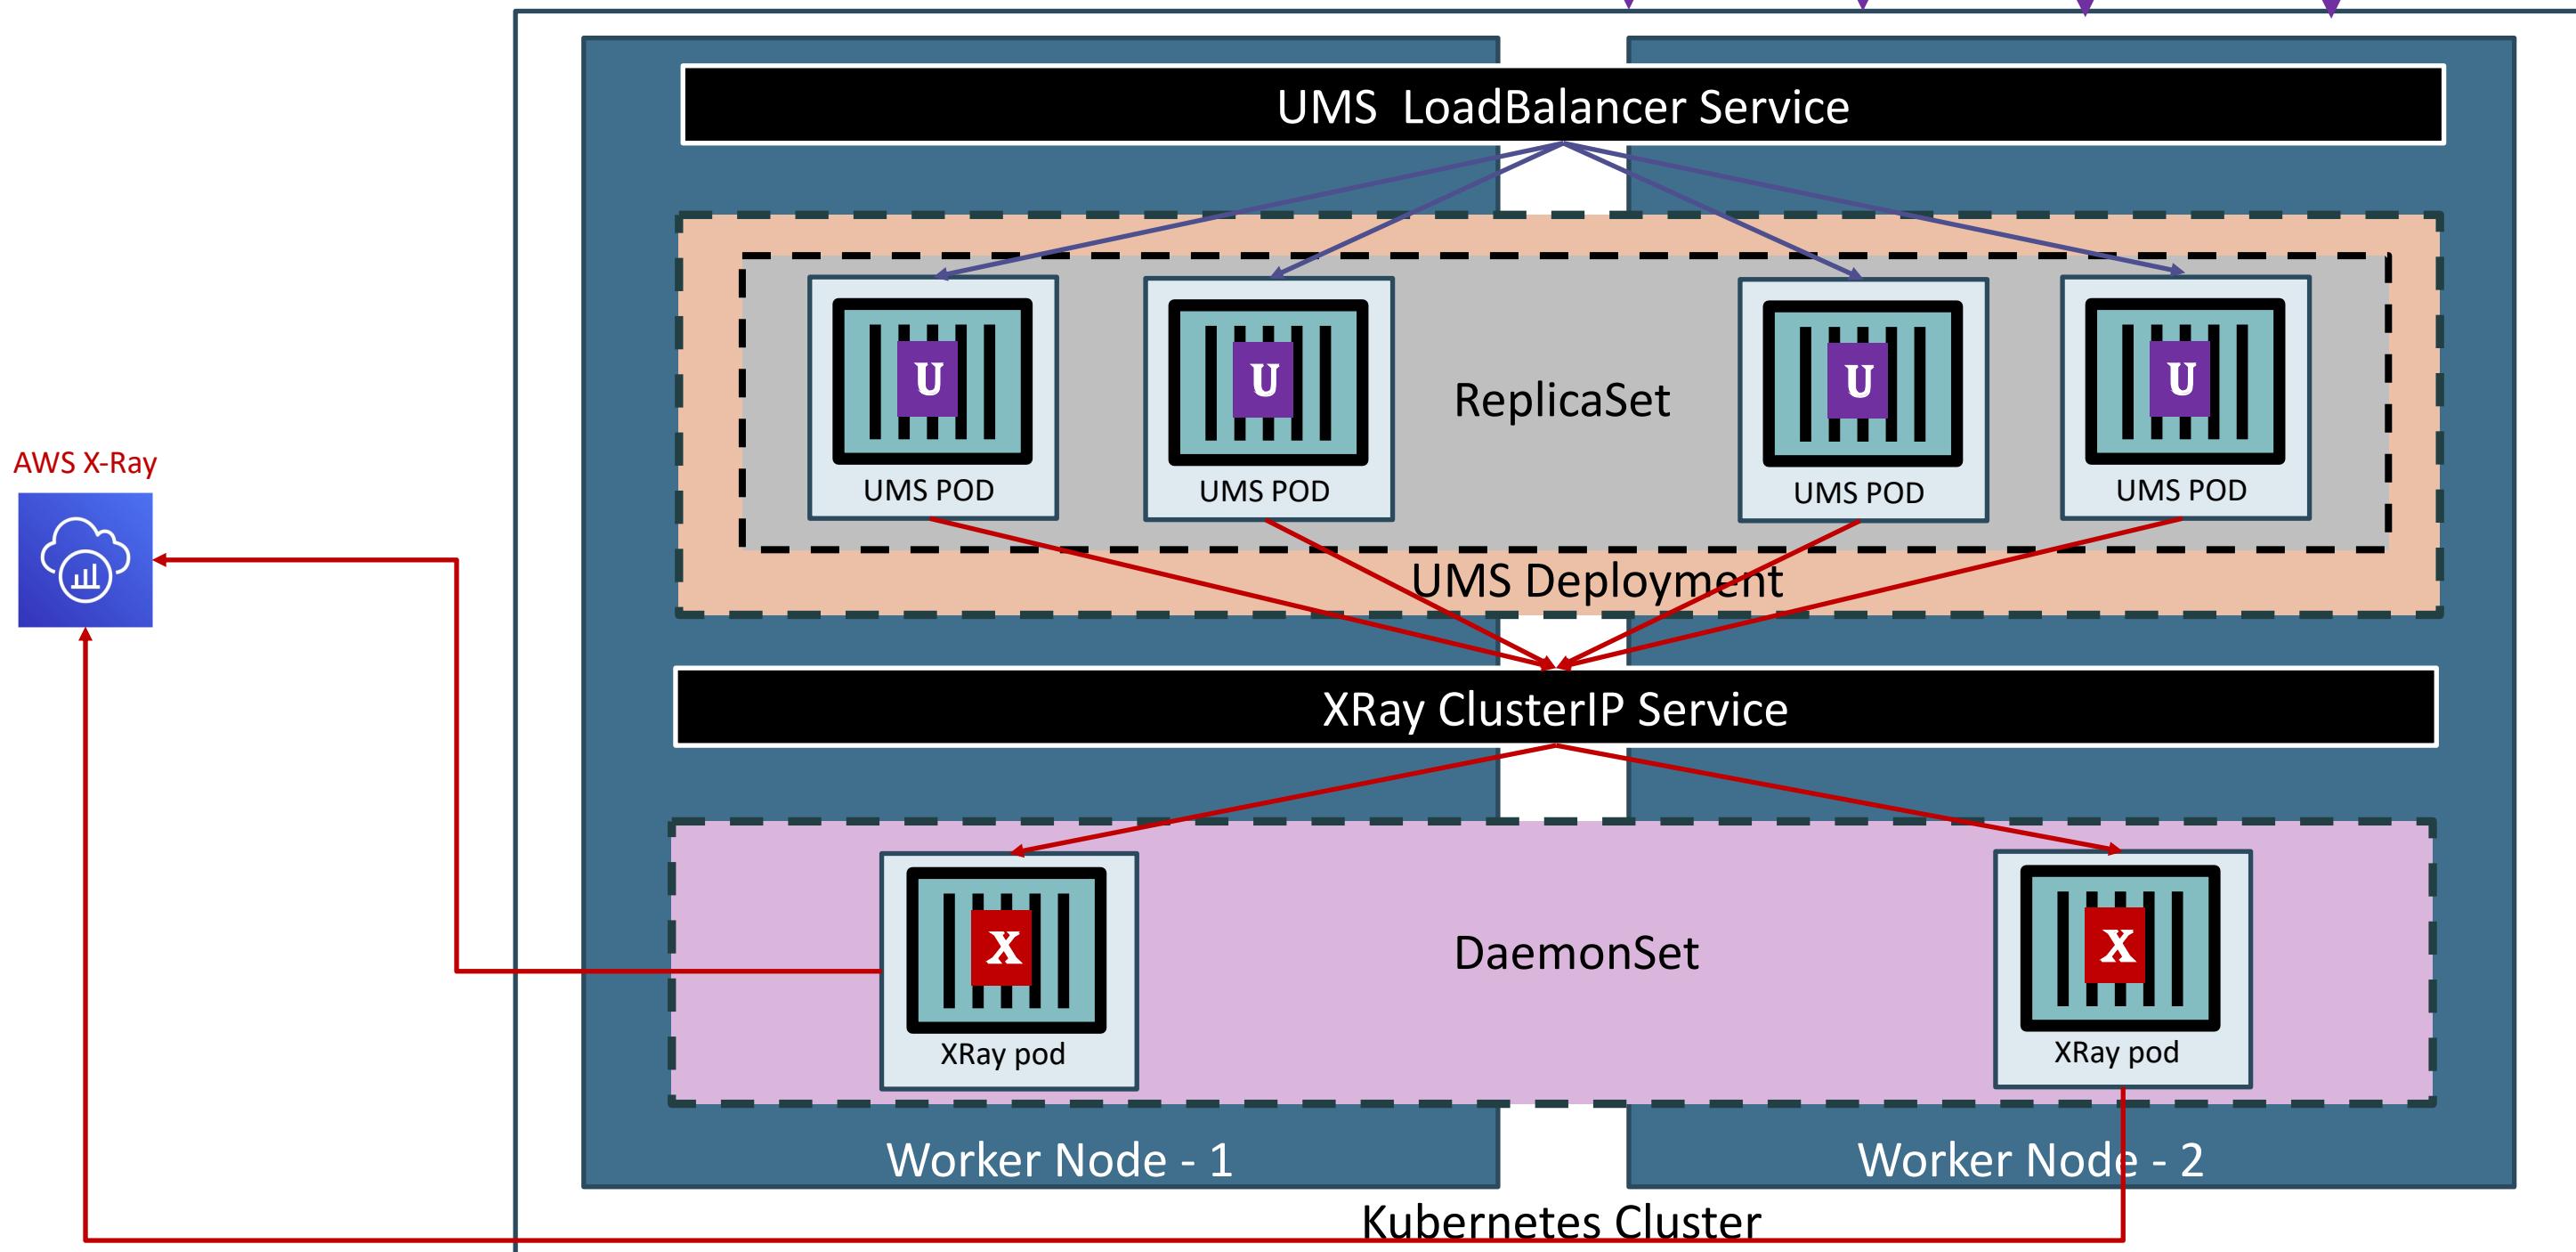
# Microservices



# Microservices Deployment on AWS EKS



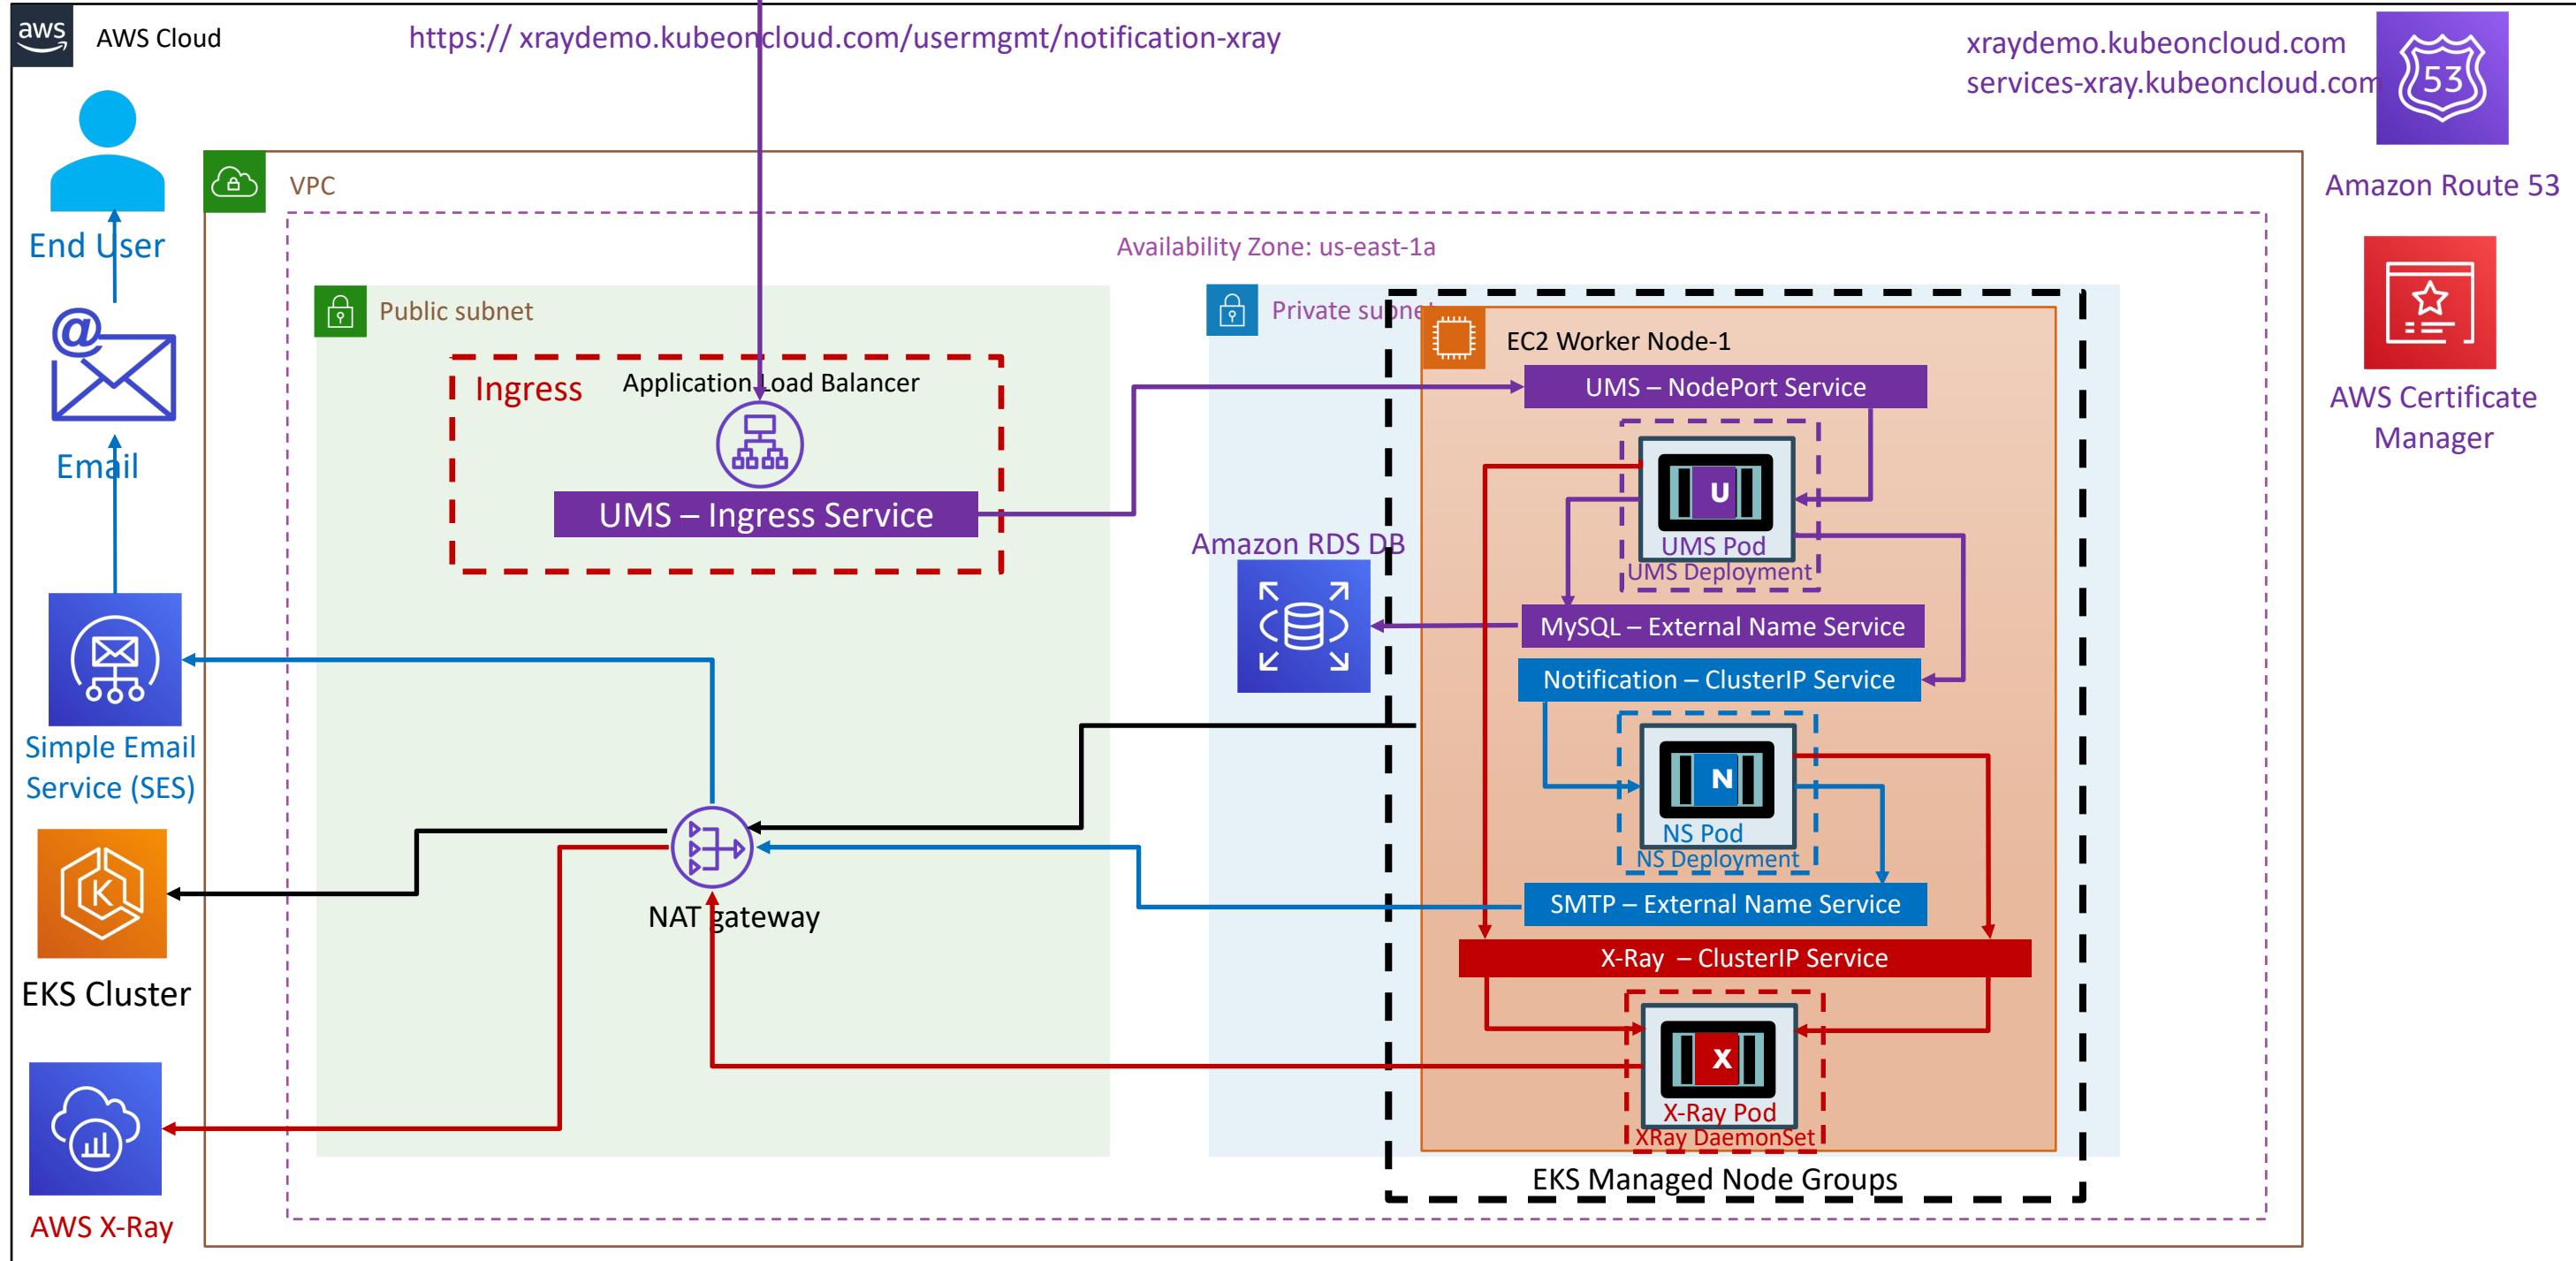
# Kubernetes – DaemonSets





# Microservices Distributed Tracing with AWS X-Ray

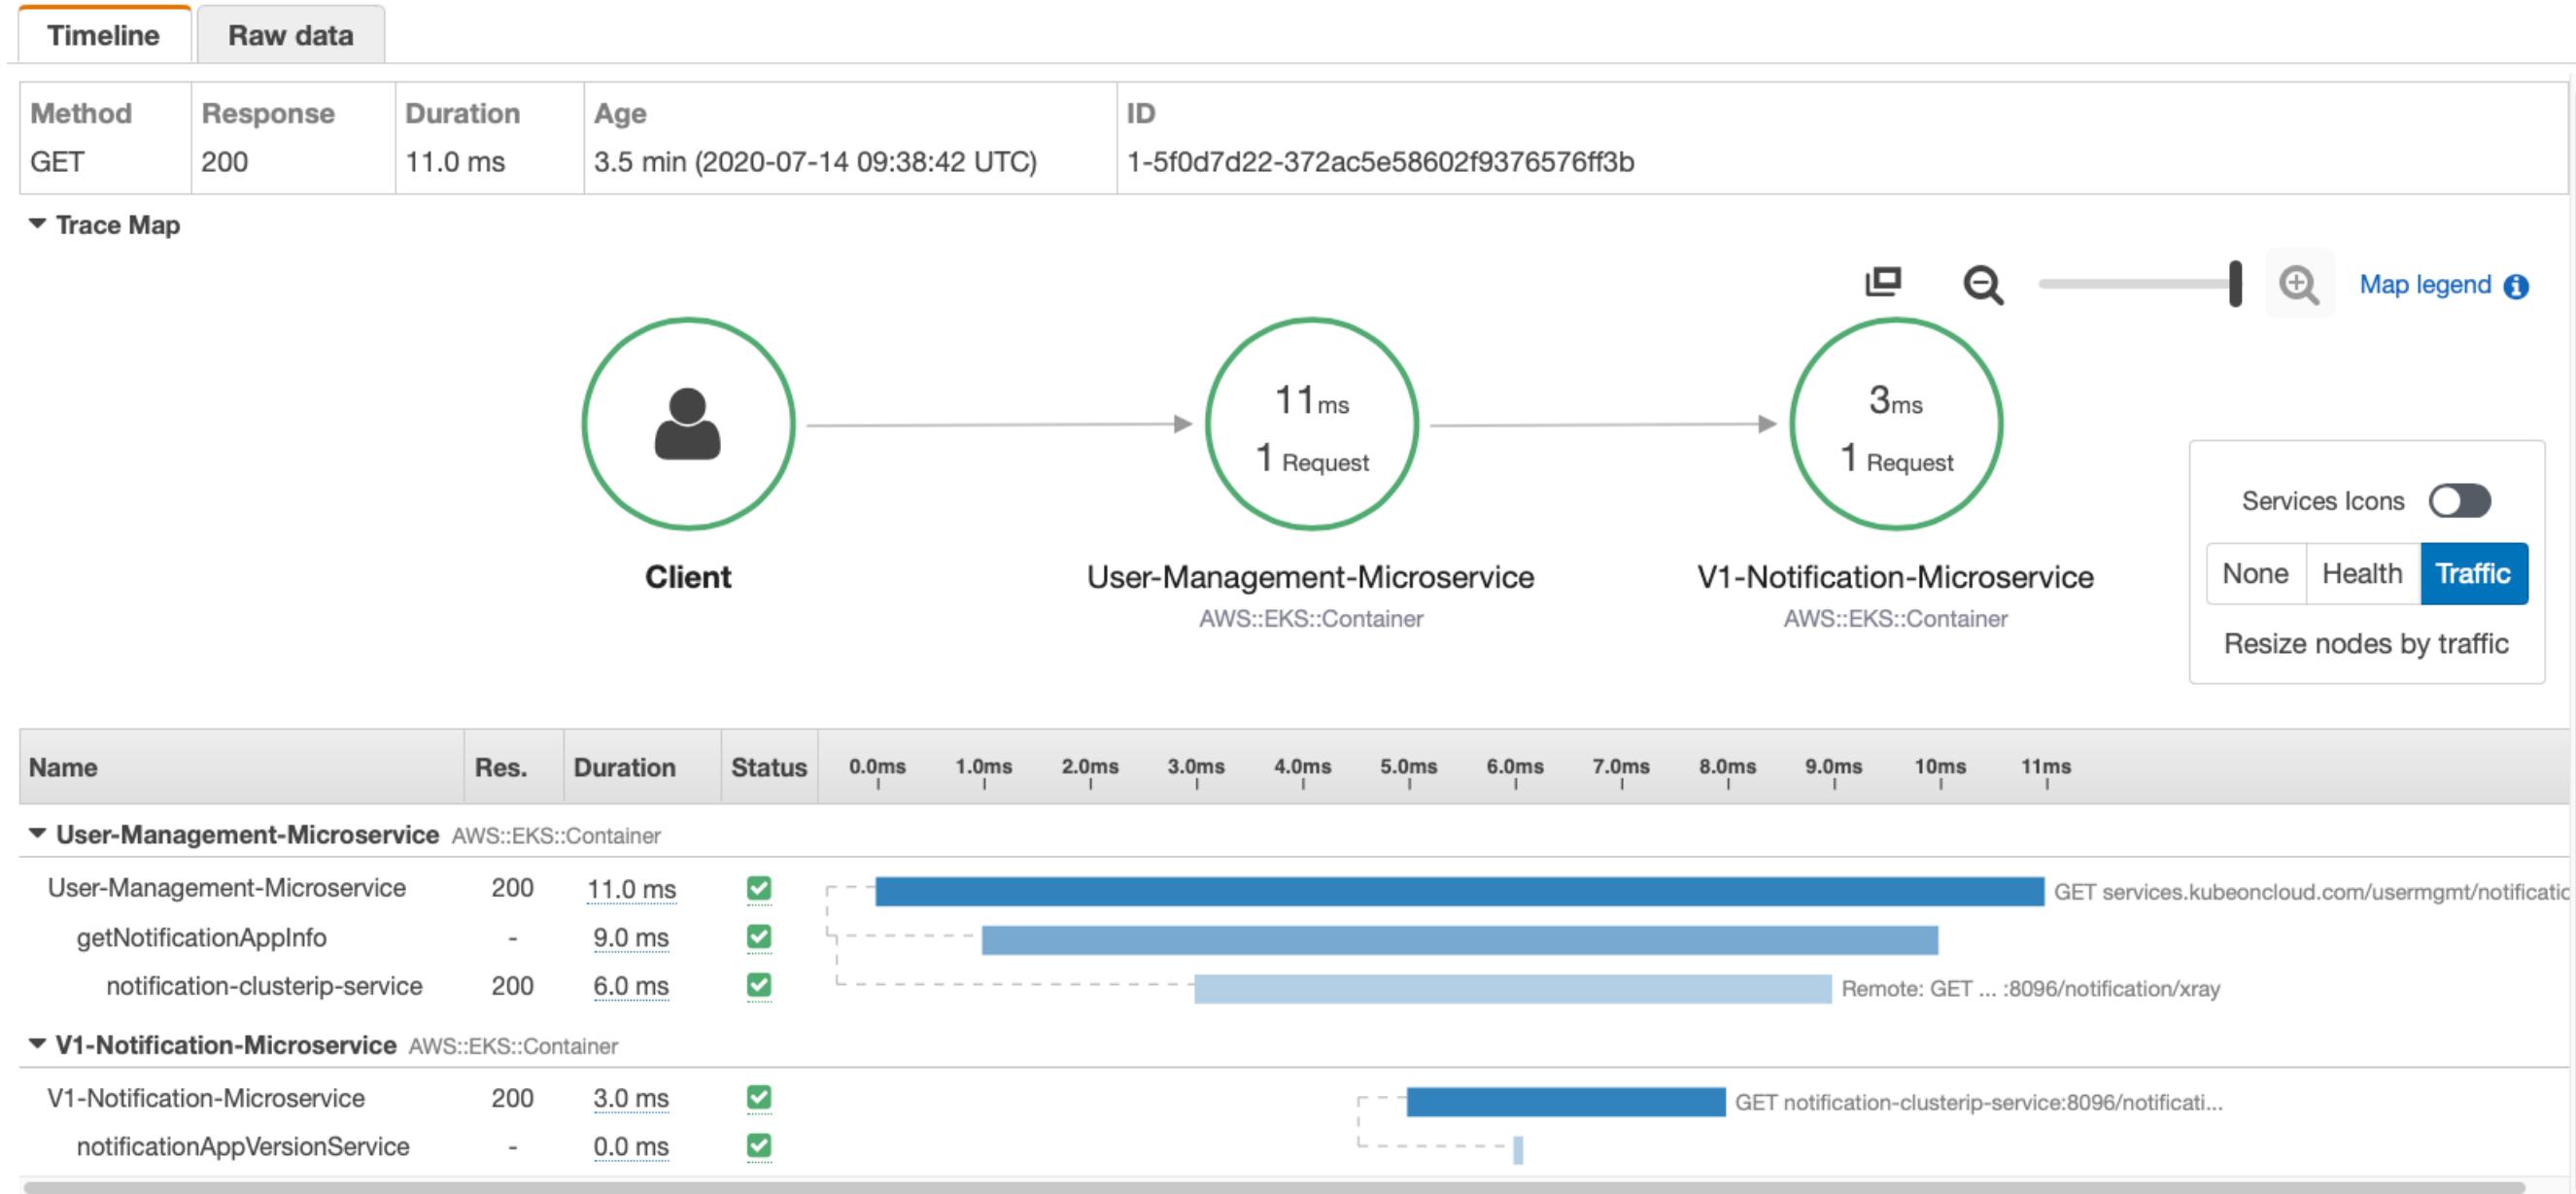
STACKSIMPLIFY



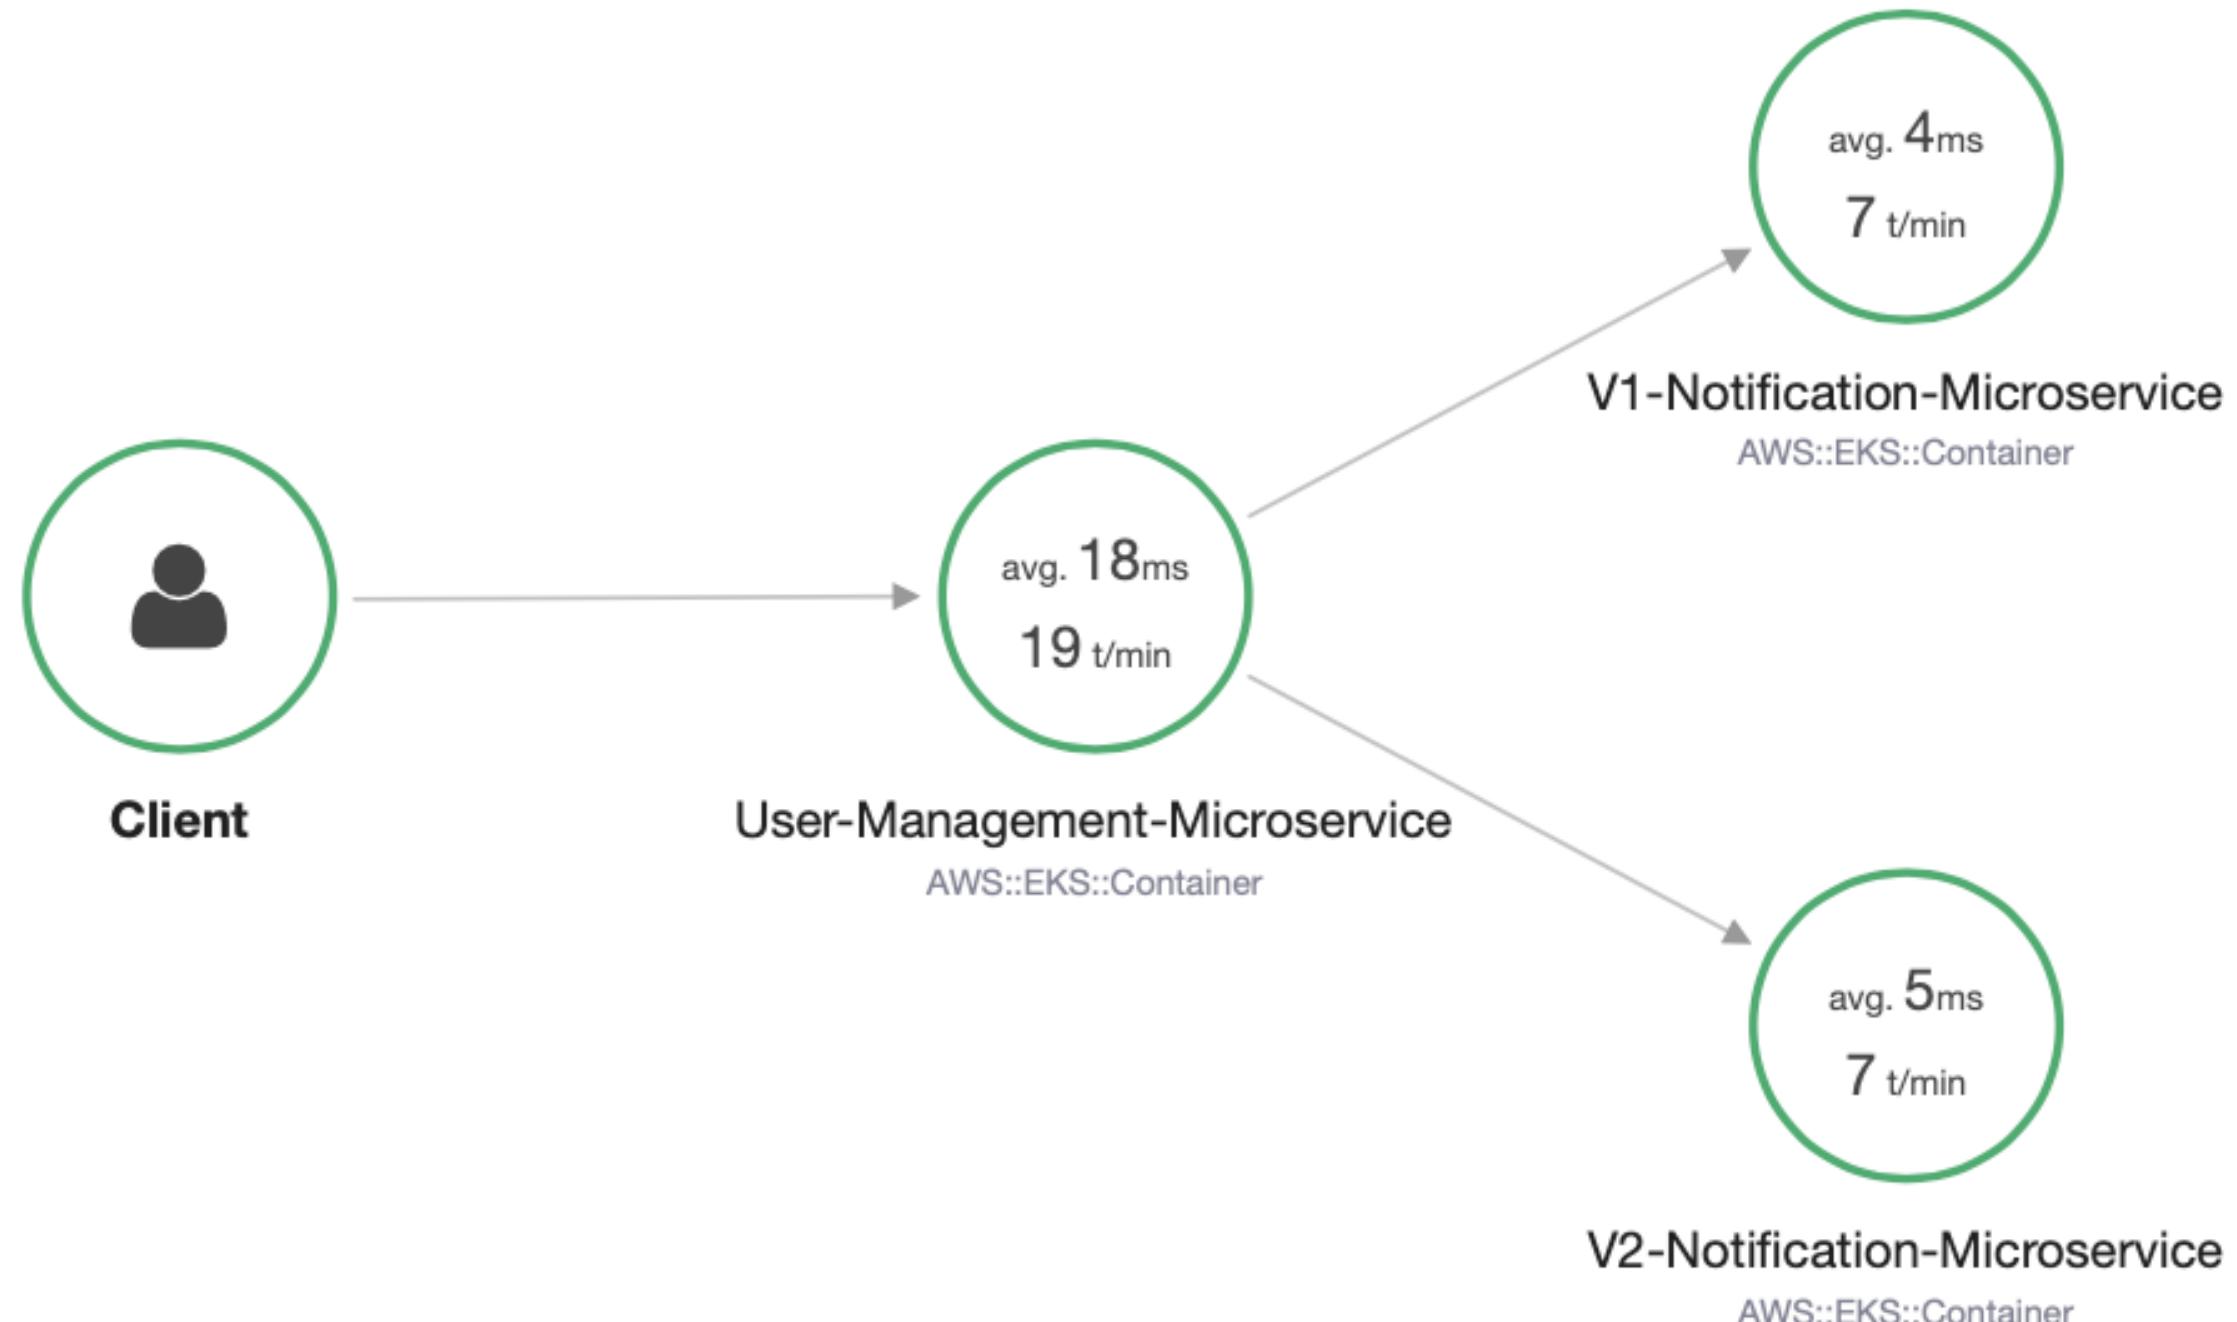
# AWS X-Ray – Service Map



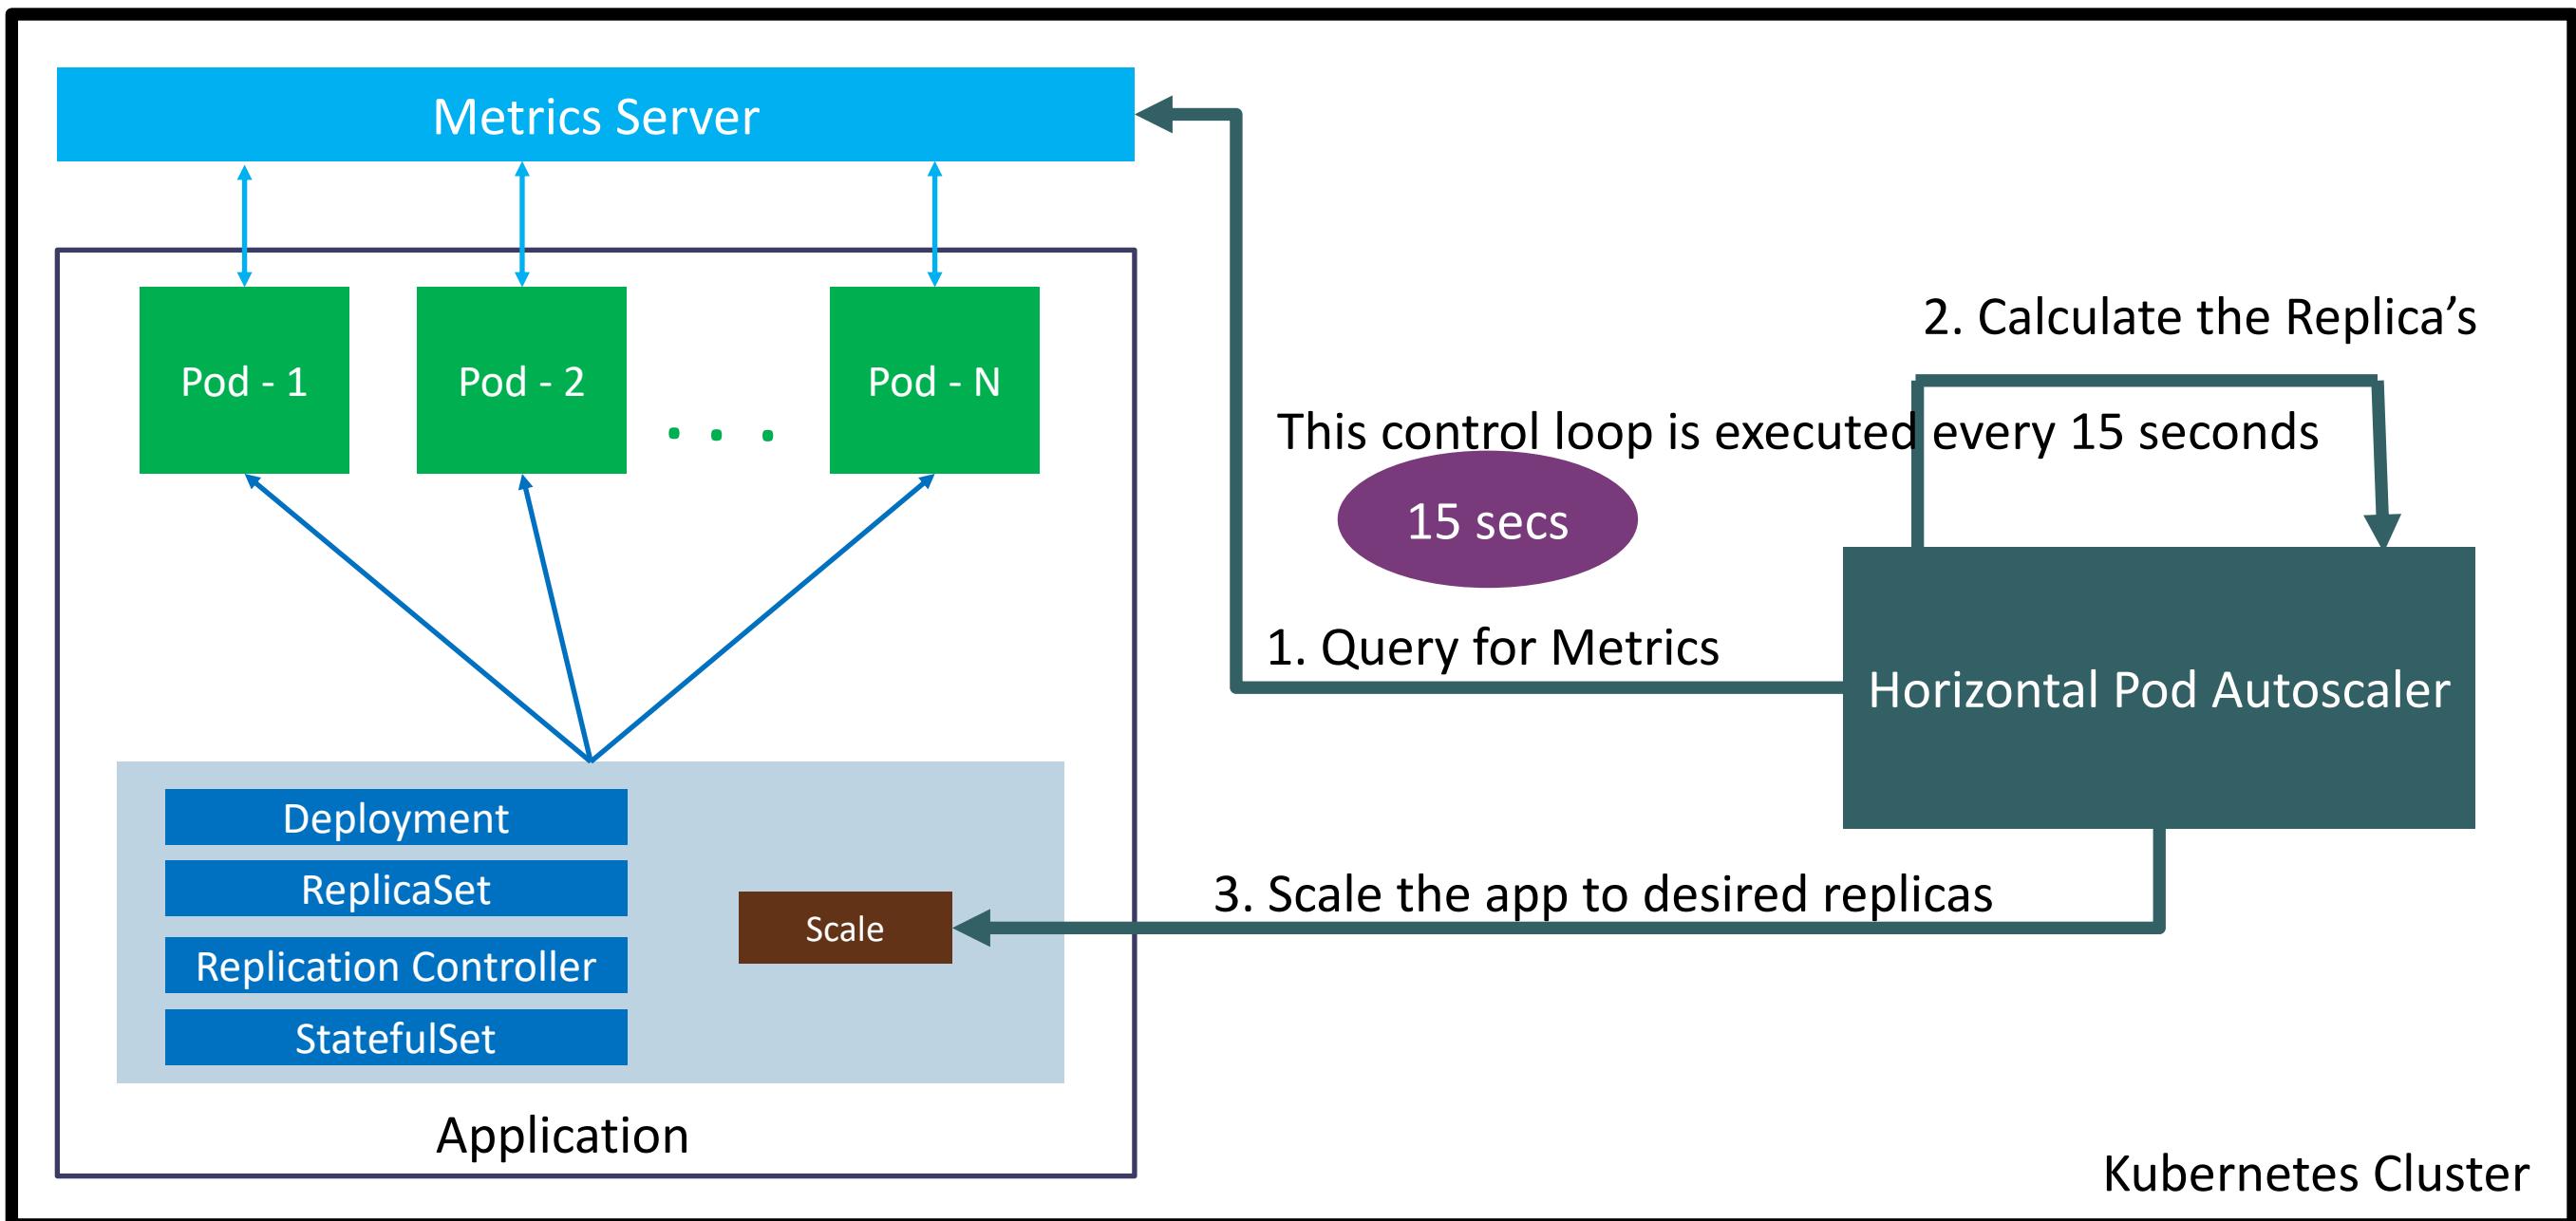
# AWS X-Ray - Traces



# Microservices – Canary Deployments



# How HPA works?



- Container Map
- Container Resources
- Performance Dashboards
- Log Groups
- Log Insights
- Alarms

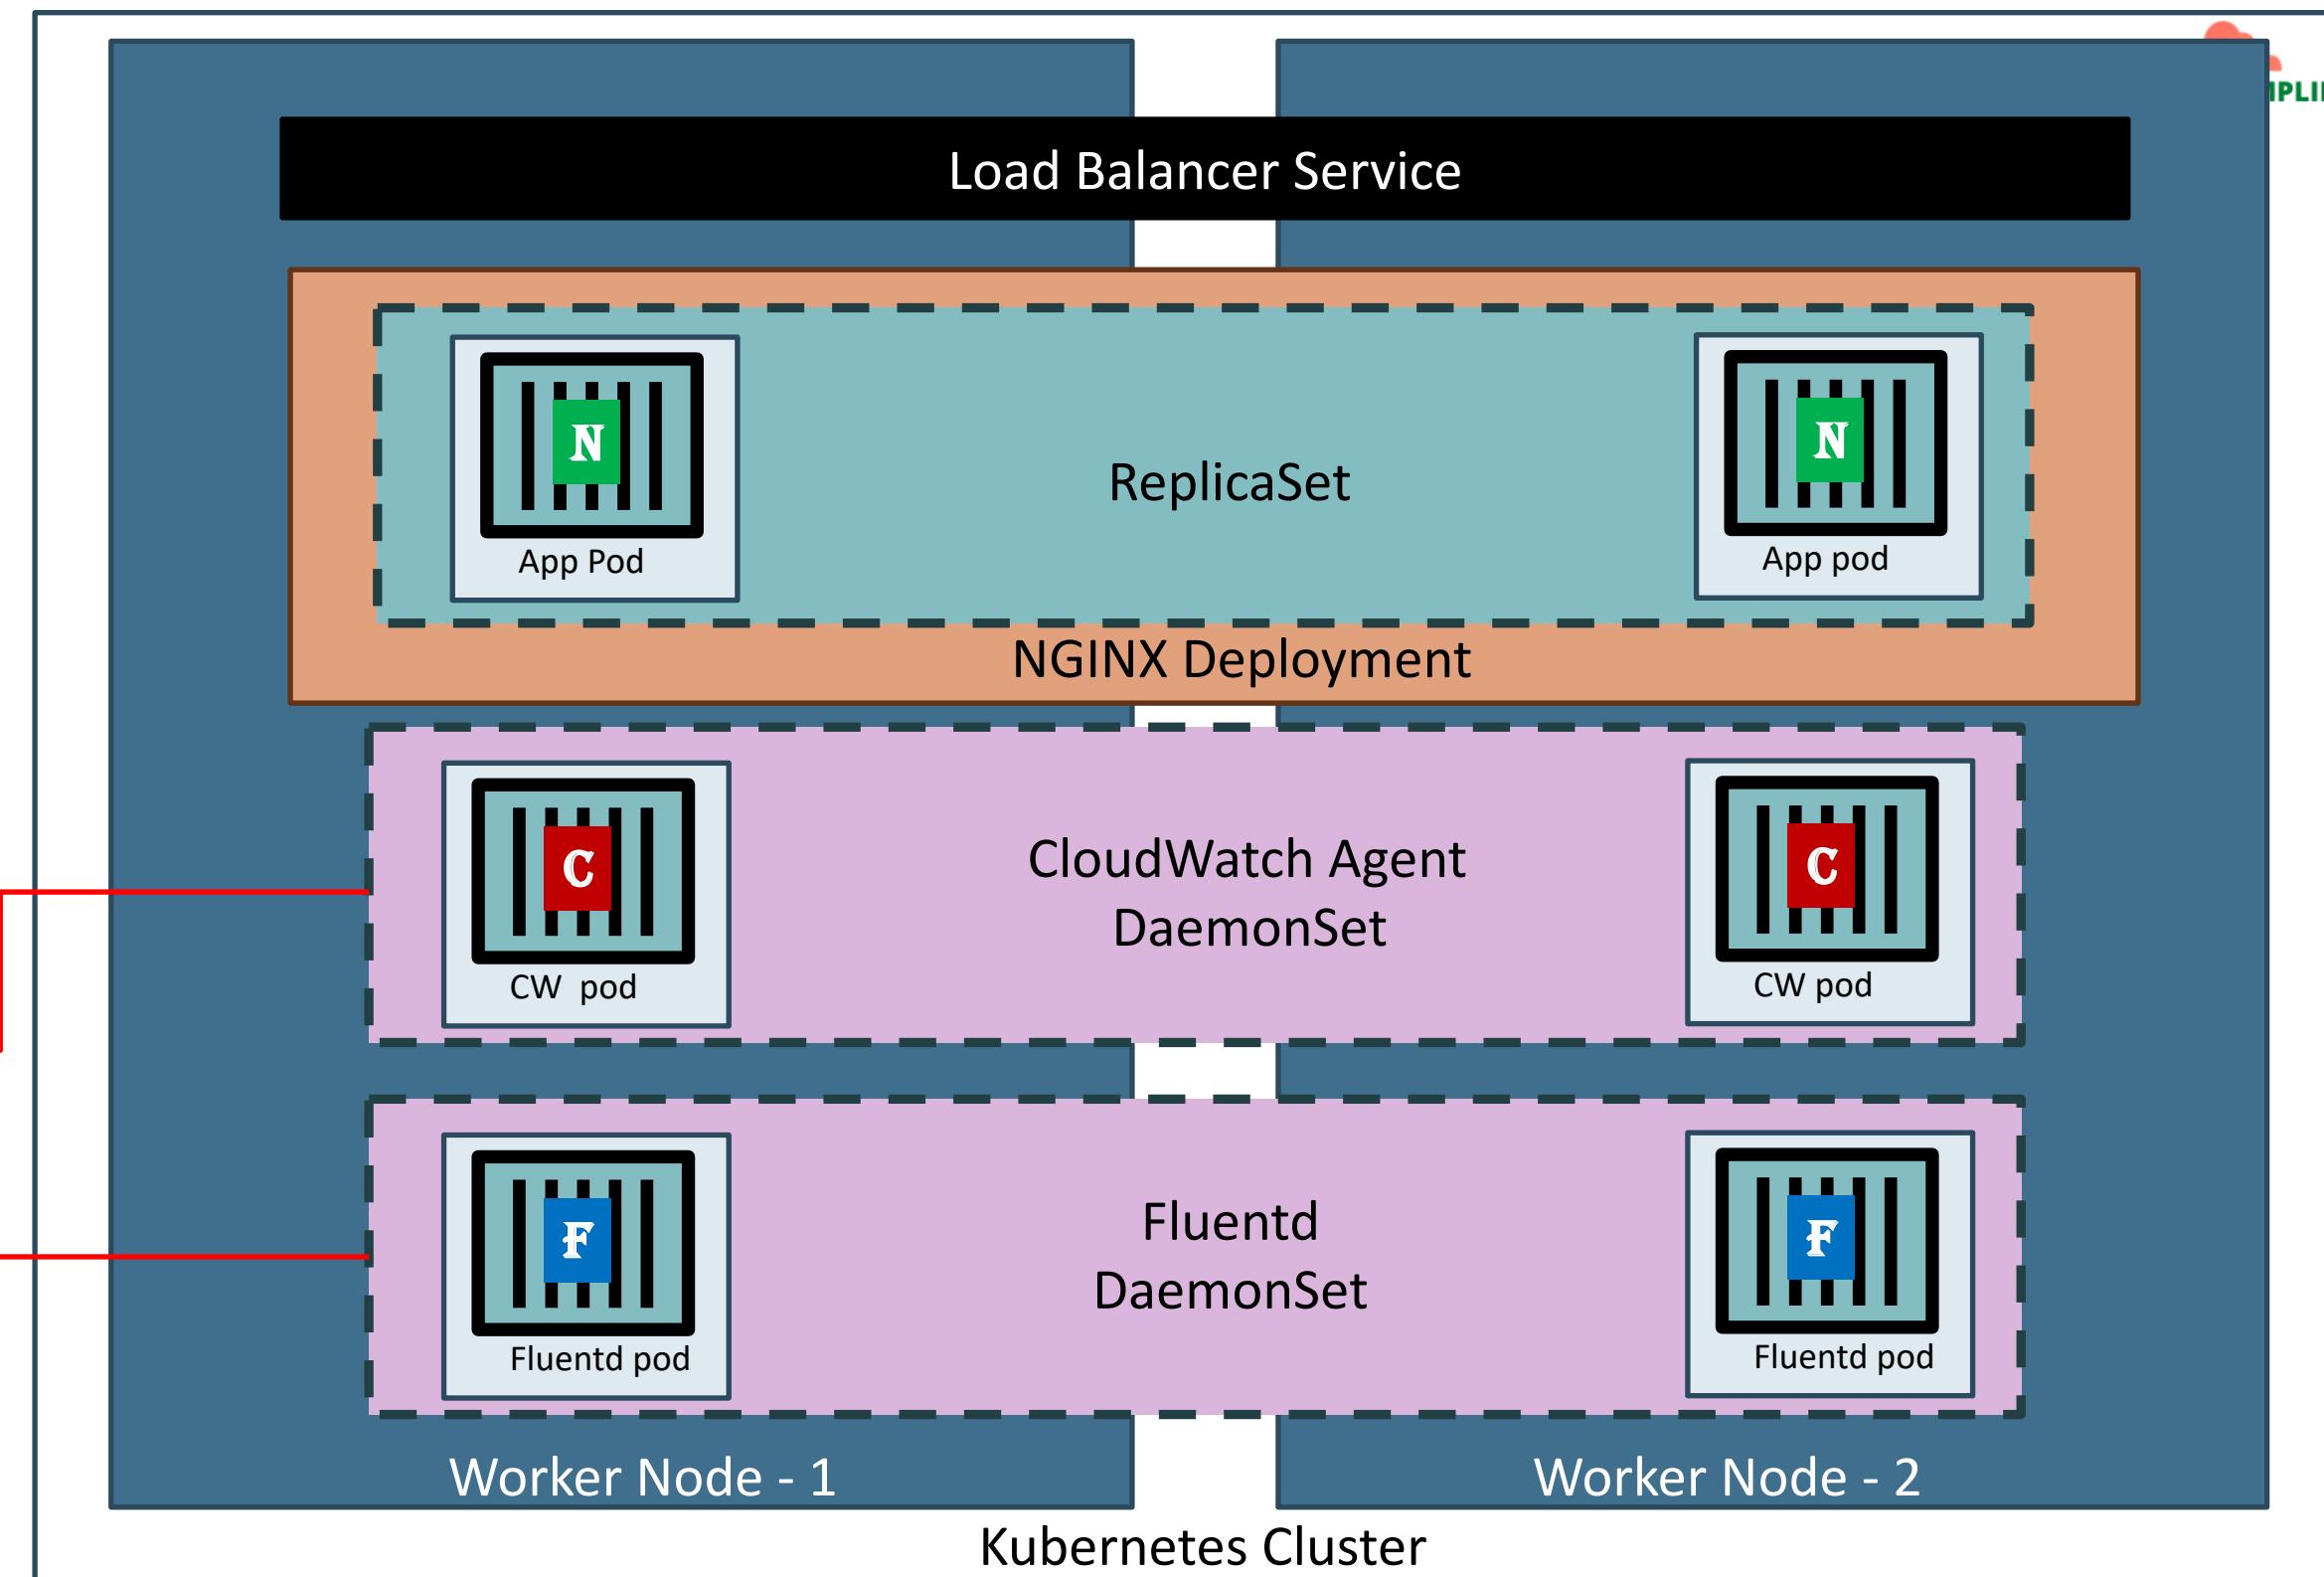


Developer or Operations User

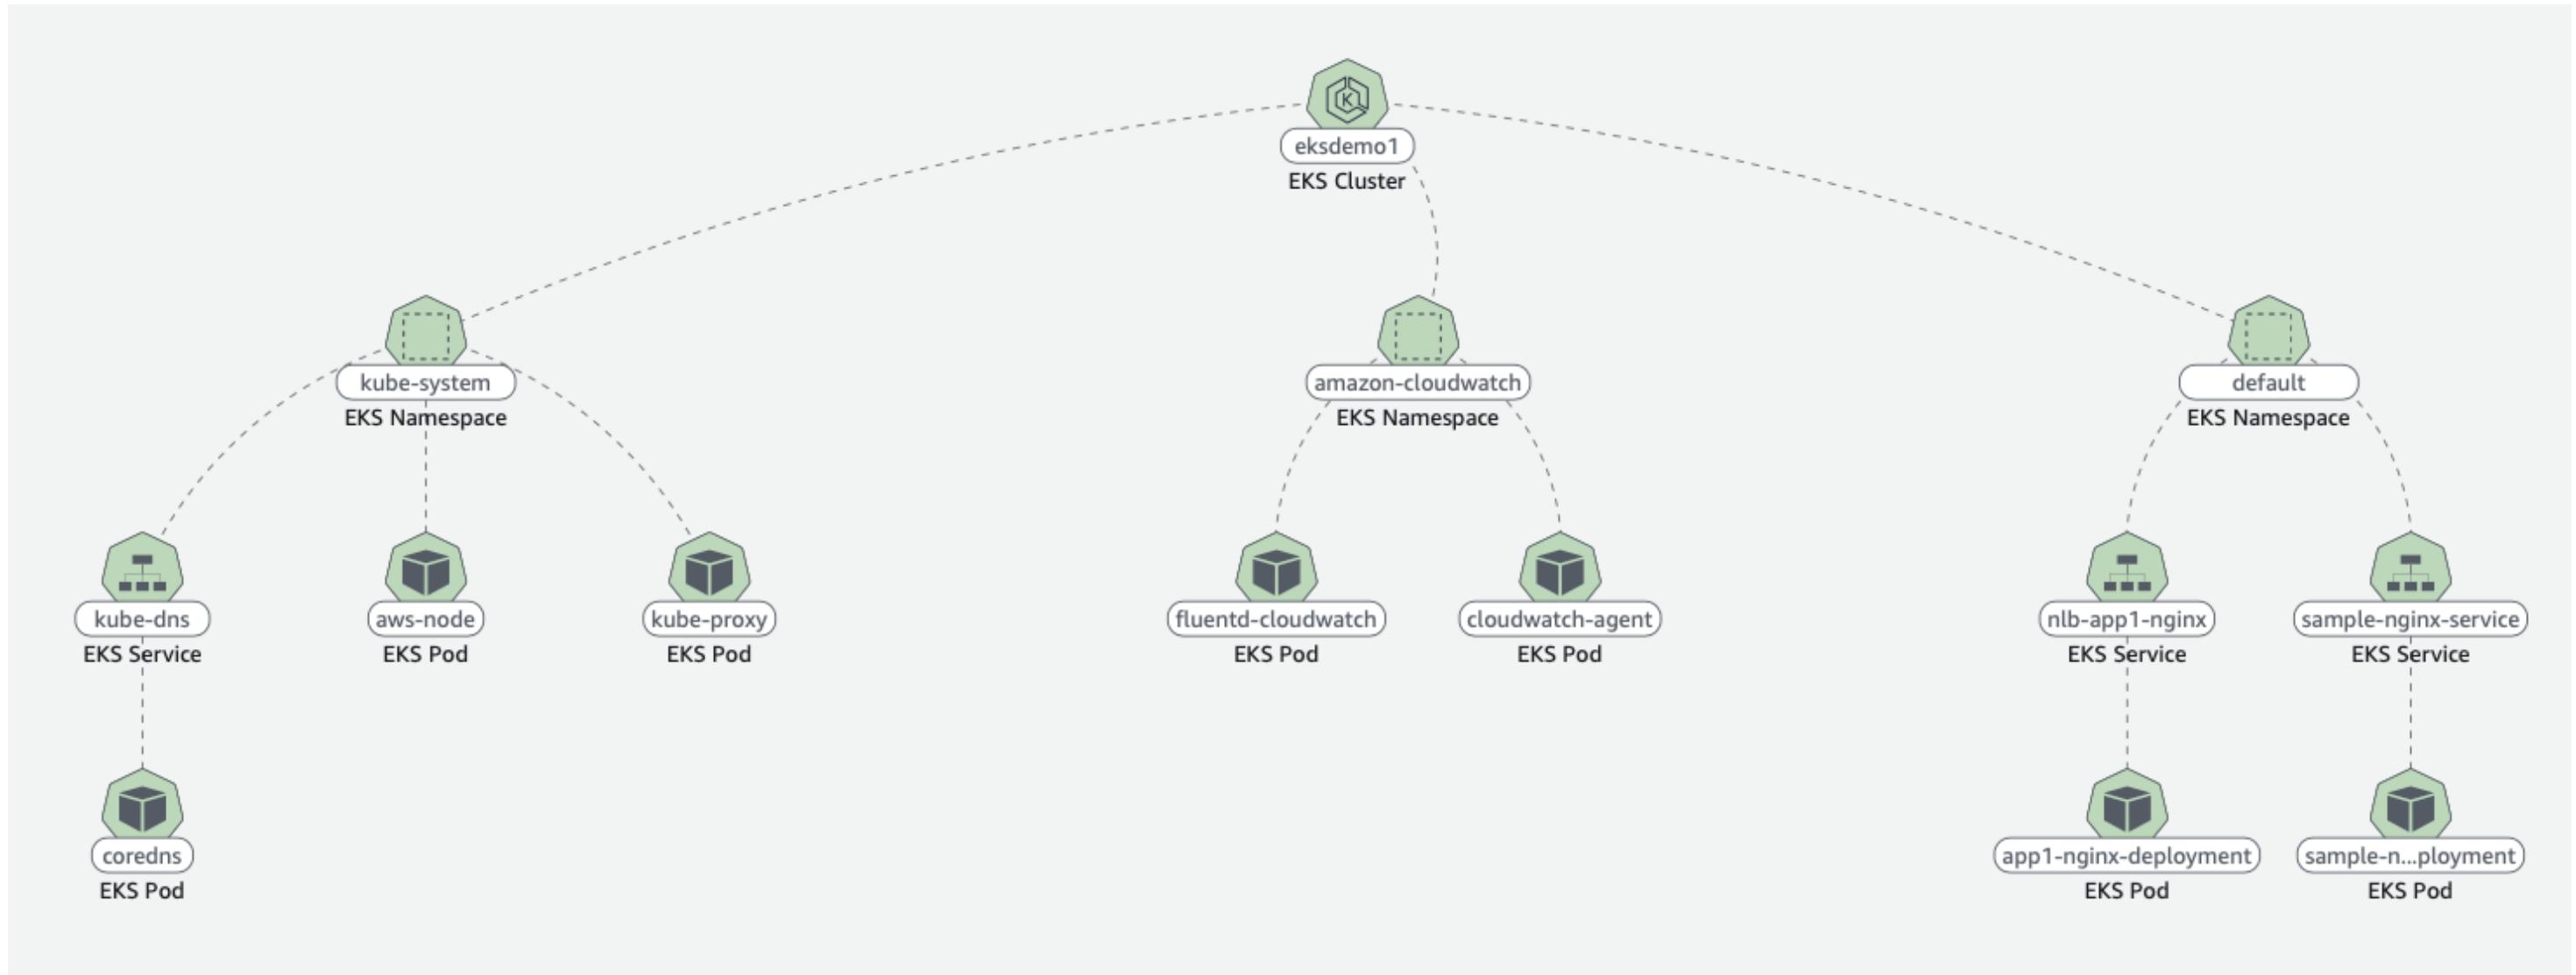


CloudWatch

# Container Insights



# CloudWatch Container Insights Map

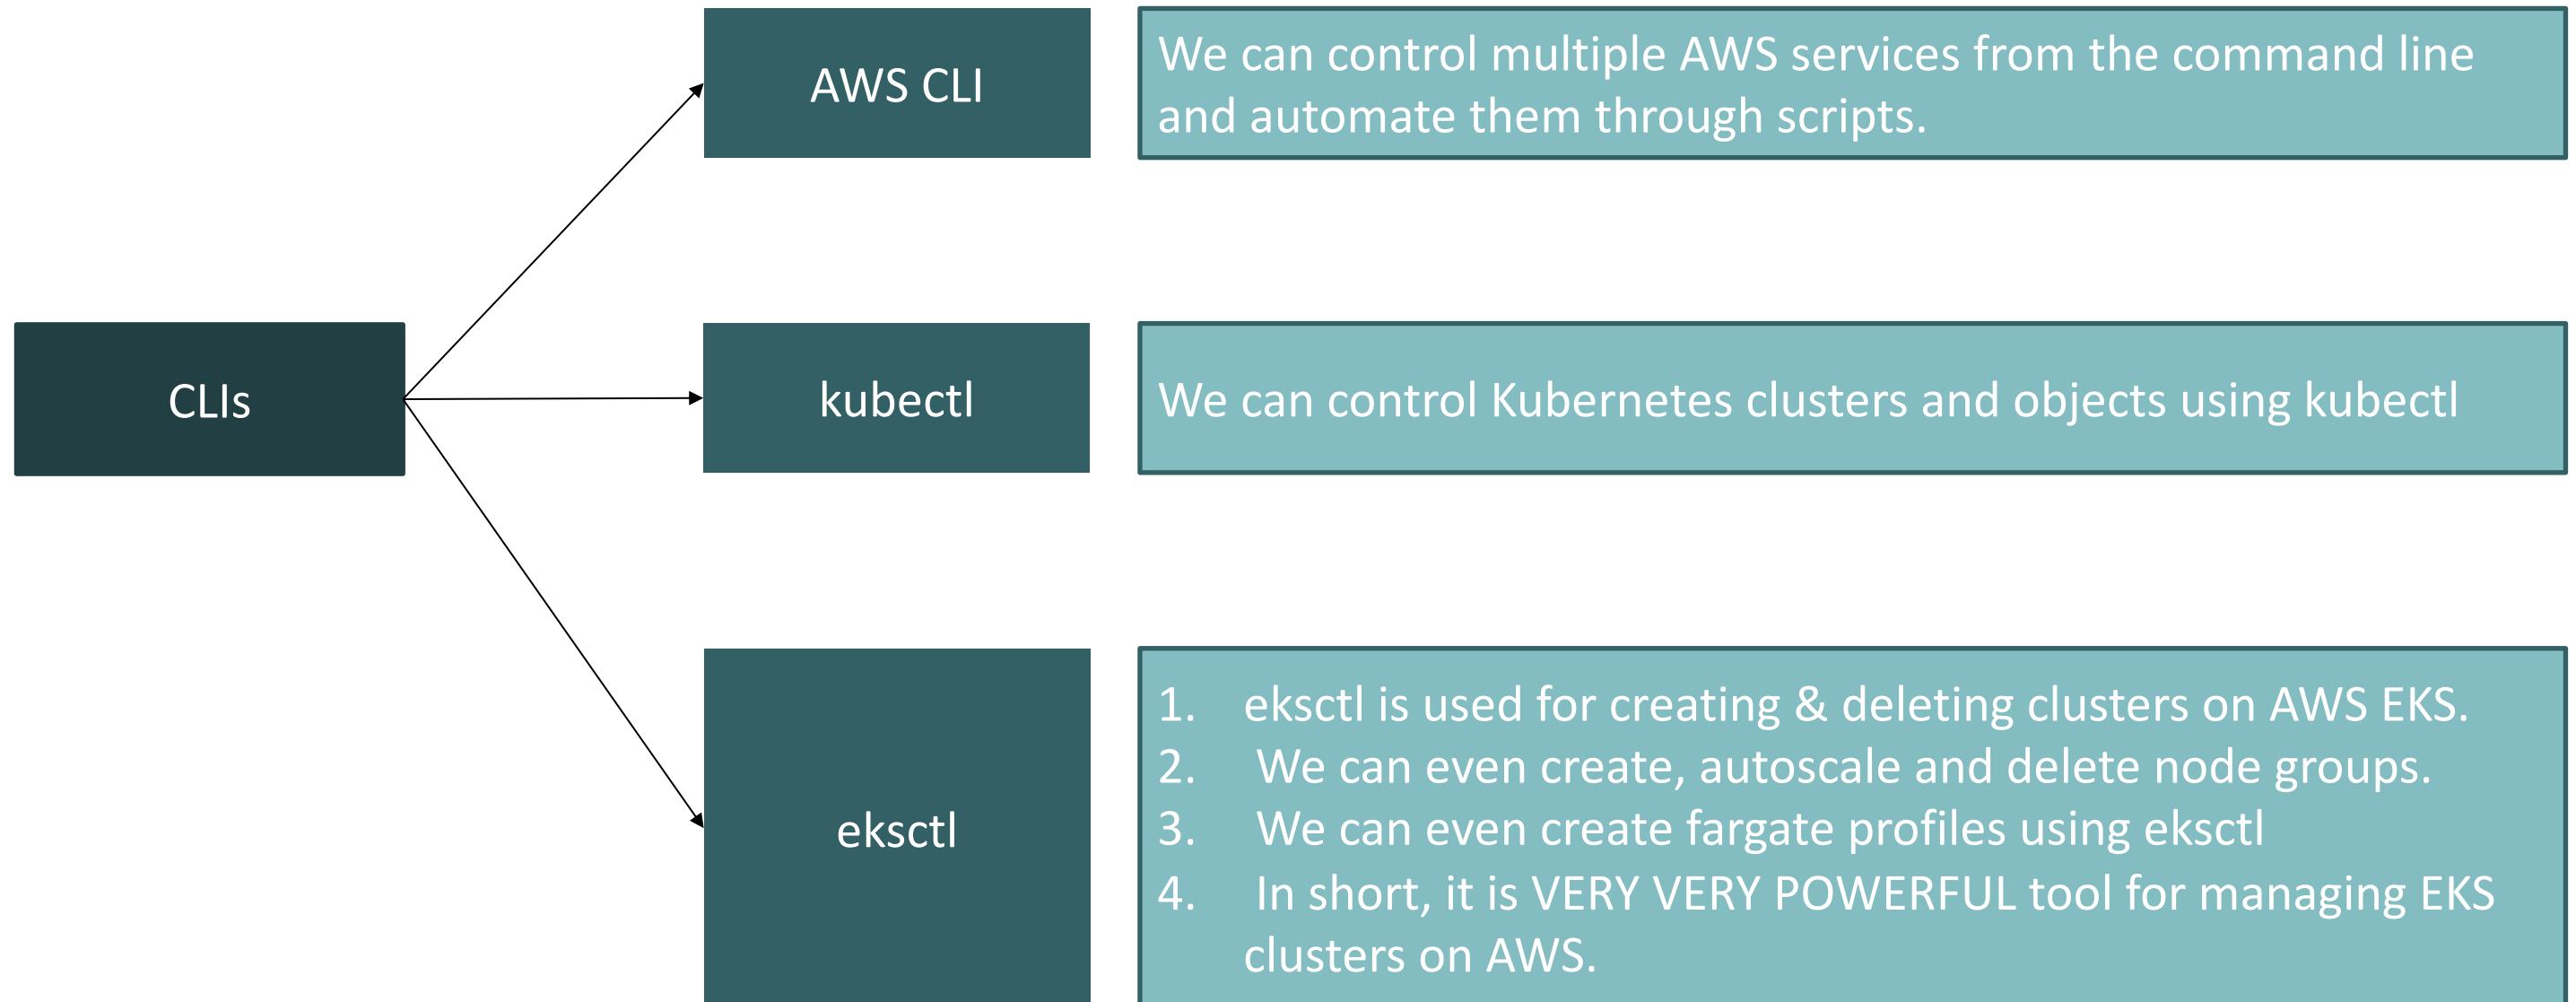




# AWS EKS CLIs



# AWS EKS Cluster - CLIs

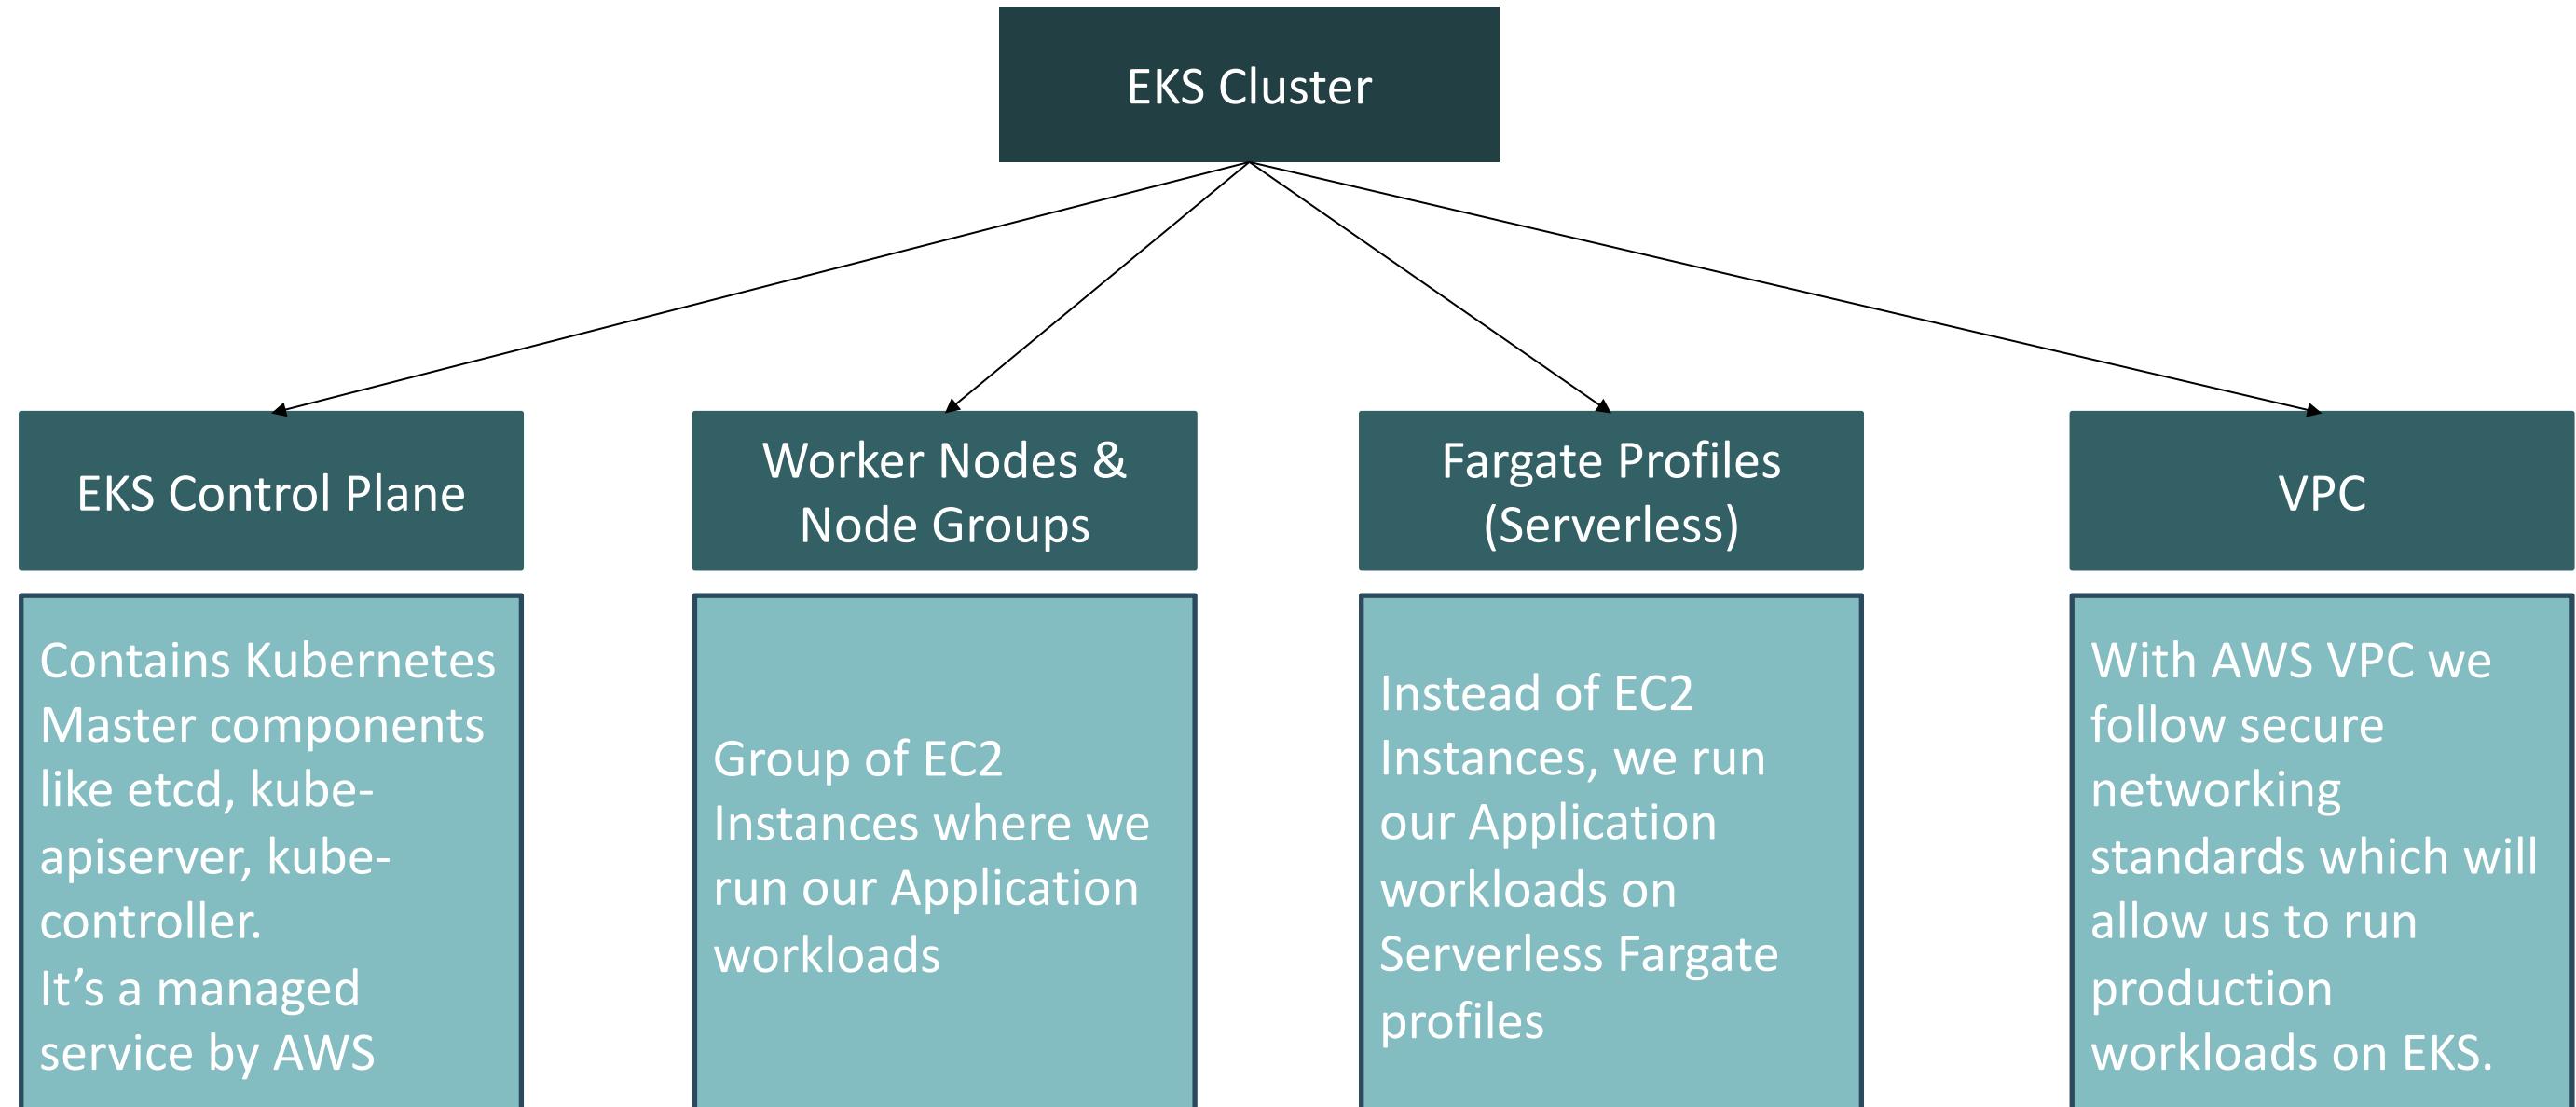




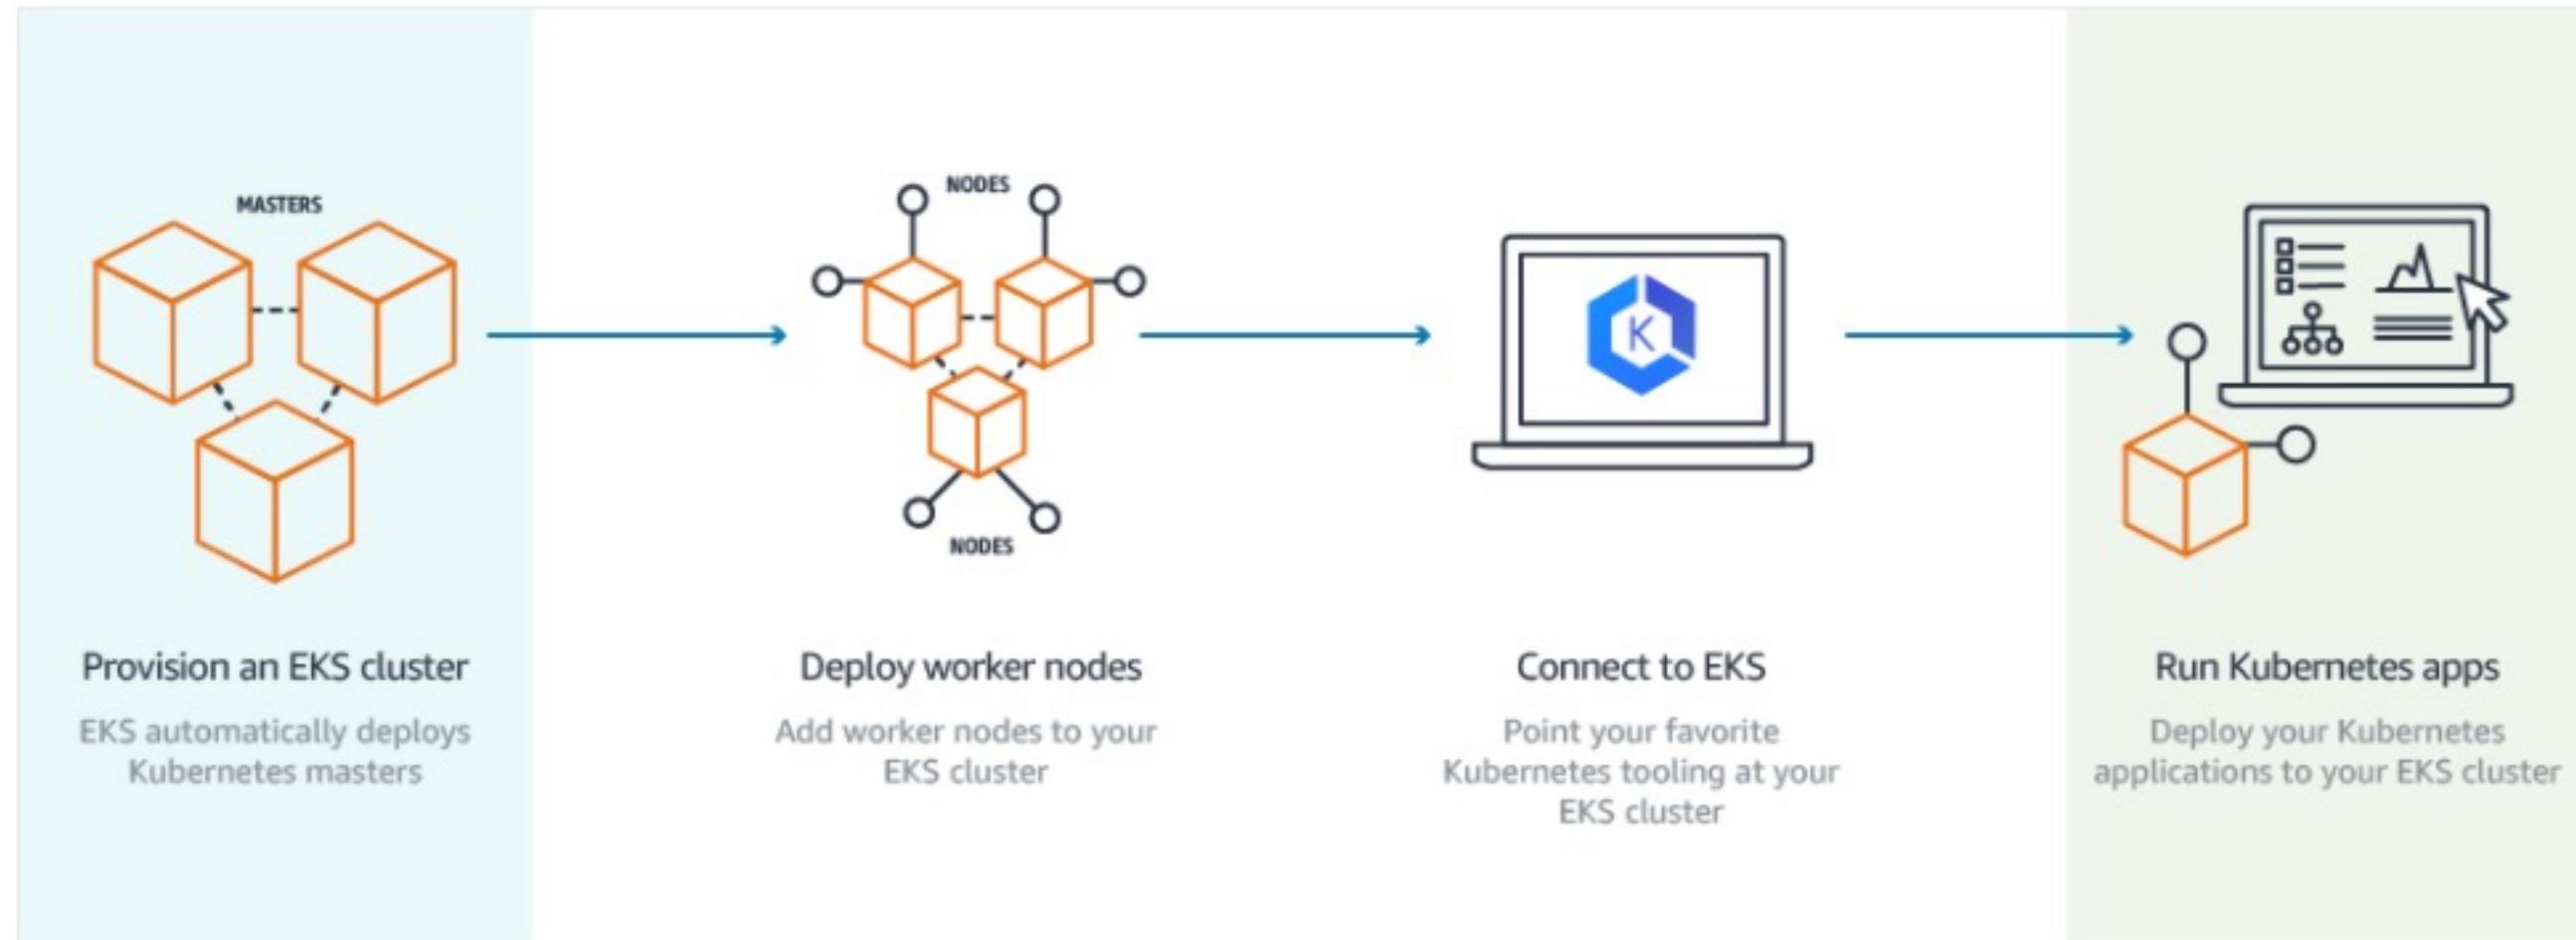
# AWS EKS Cluster



# AWS EKS – Core Objects



# How does EKS work?



# EKS Cluster – Core Objects Detailed

## EKS Control Plane

1. EKS runs a single tenant Kubernetes control plane for each cluster, and control plane infrastructure is **not shared** across clusters or AWS accounts.
2. This control plane consists of at least two API server nodes and three etcd nodes that run across **three Availability Zones within a Region**
3. EKS automatically detects and replaces **unhealthy** control plane instances, restarting them across the Availability Zones within the Region as needed.

## Worker Nodes & Node Groups

1. Worker machines in Kubernetes are called nodes. These are EC2 Instances
2. EKS worker nodes run in our AWS account and connect to our cluster's control plane via the **cluster API server endpoint**.
3. A node group is **one or more** EC2 instances that are deployed in an EC2 Autoscaling group.
4. All instances in a node group must
  1. Be the **same instance type**
  2. Be **running the same AMI**
  3. Use the **same EKS worker node IAM role**

# EKS Cluster – Core Objects Detailed

## Fargate Profiles

1. AWS Fargate is a technology that provides **on-demand, right-sized compute capacity** for containers
2. With Fargate, we **no longer have to provision, configure, or scale groups of virtual machines** to run containers.
3. Each pod running on Fargate has its **own isolation boundary** and does not share the underlying kernel, CPU resources, memory resources, or elastic network interface with another pod.
4. AWS specially built **Fargate controllers** that recognizes the pods belonging to fargate and schedules them on Fargate profiles.
5. We will see more in our Fargate learning section.

## VPC

1. EKS uses AWS VPC network policies **to restrict traffic between control plane components** to within a single cluster.
2. Control plane components for a EKS cluster **cannot view or receive communication** from other clusters or other AWS accounts, except as authorized with Kubernetes RBAC policies.
3. This **secure and highly-available configuration** makes EKS reliable and recommended for **production workloads**.

# Kubernetes Architecture

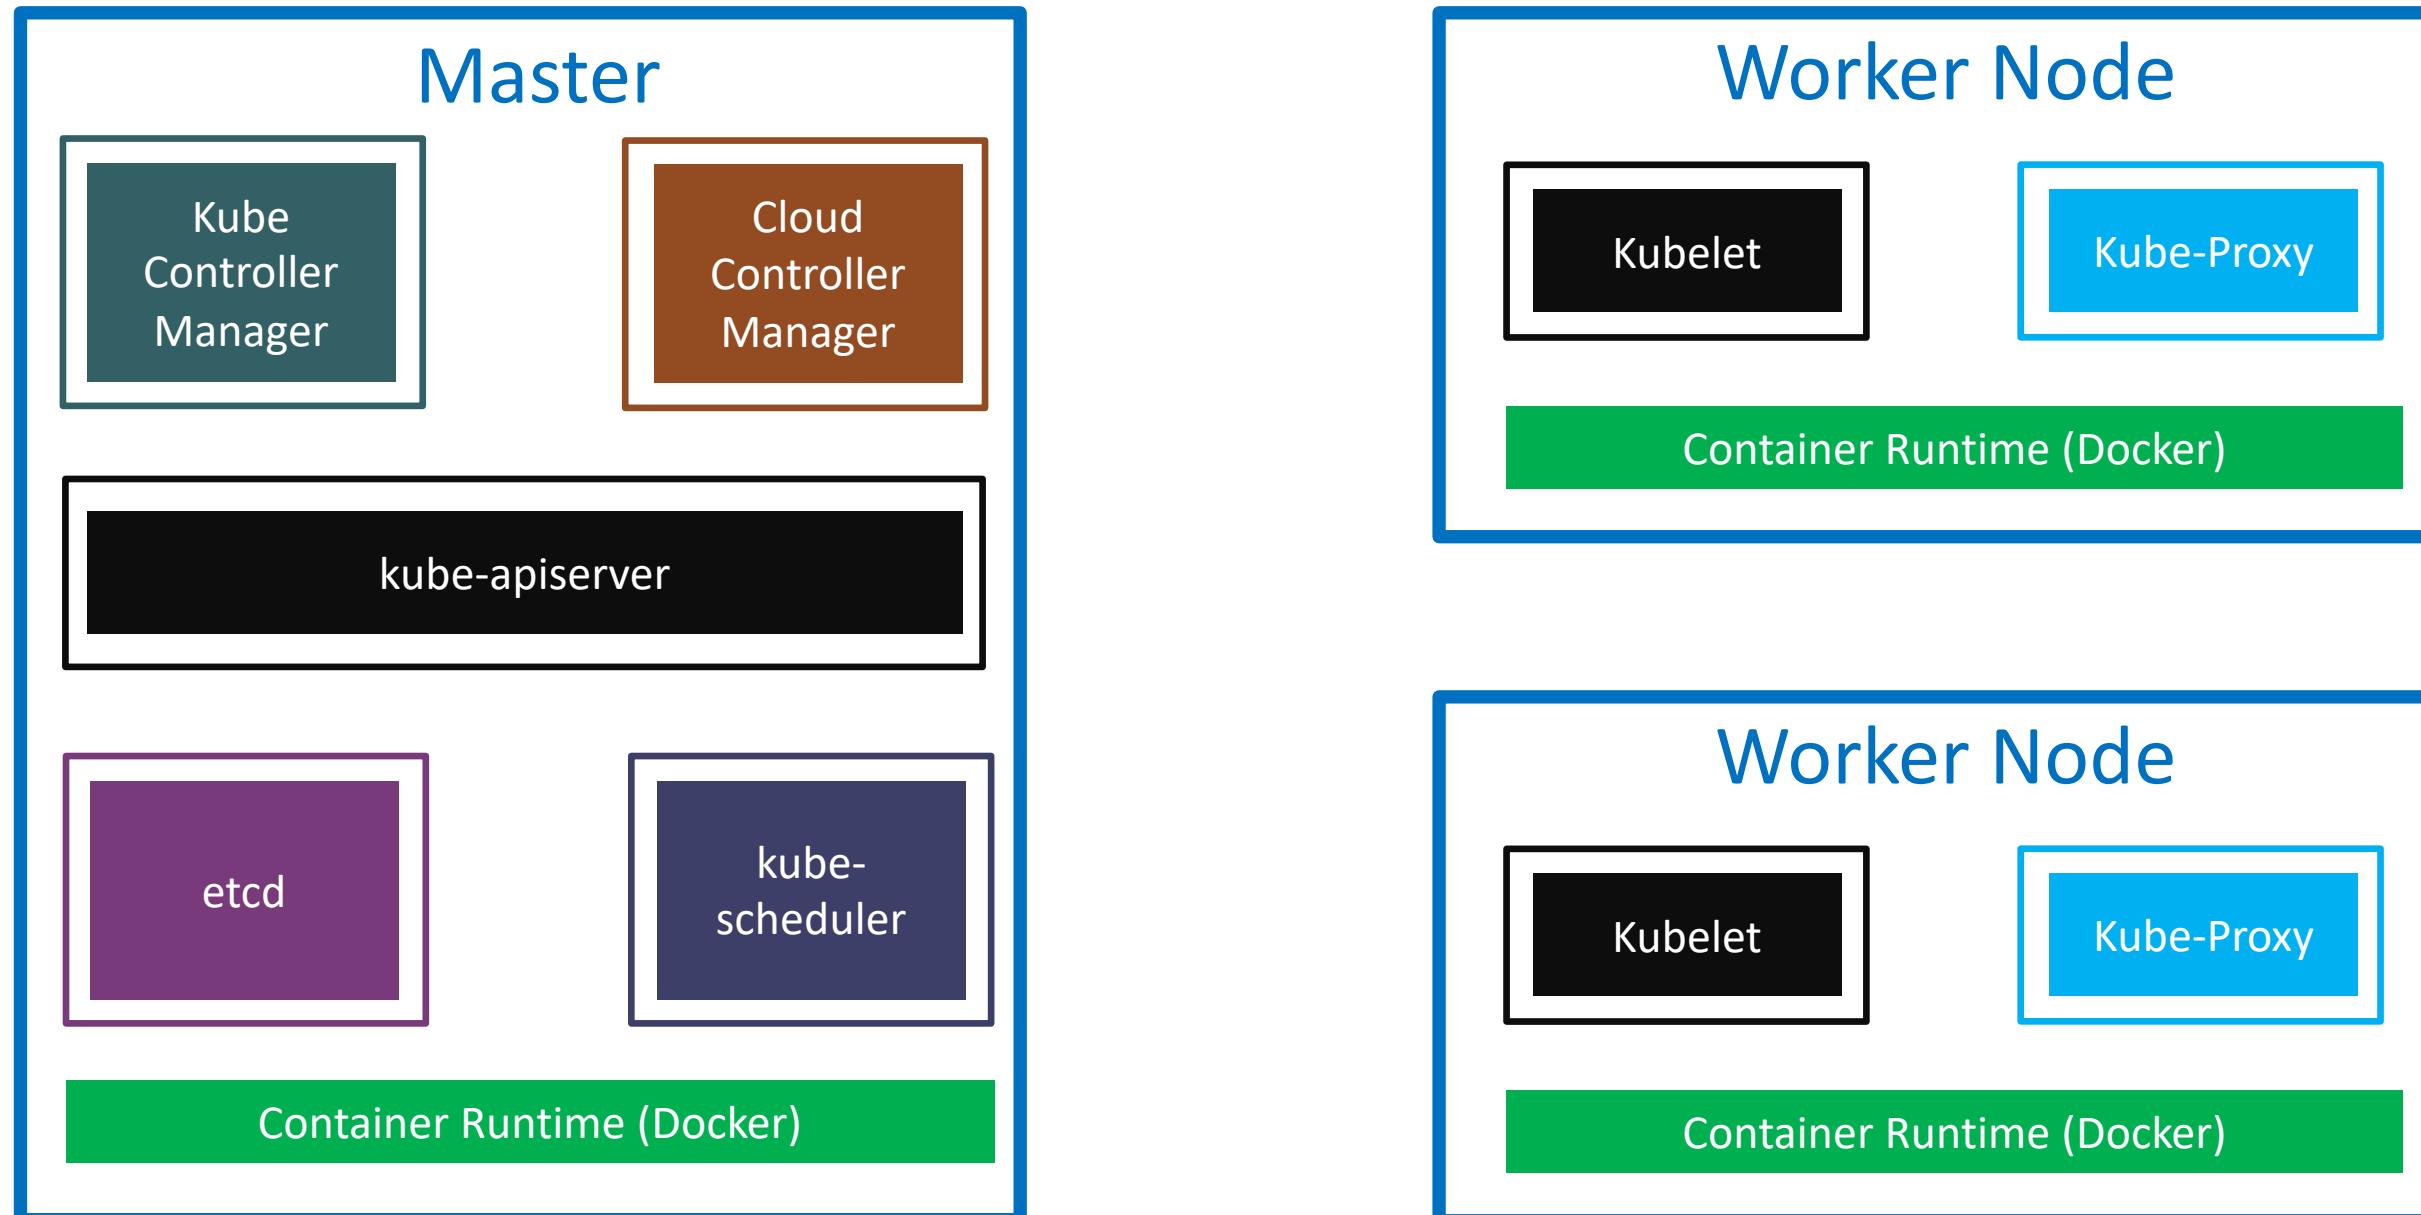




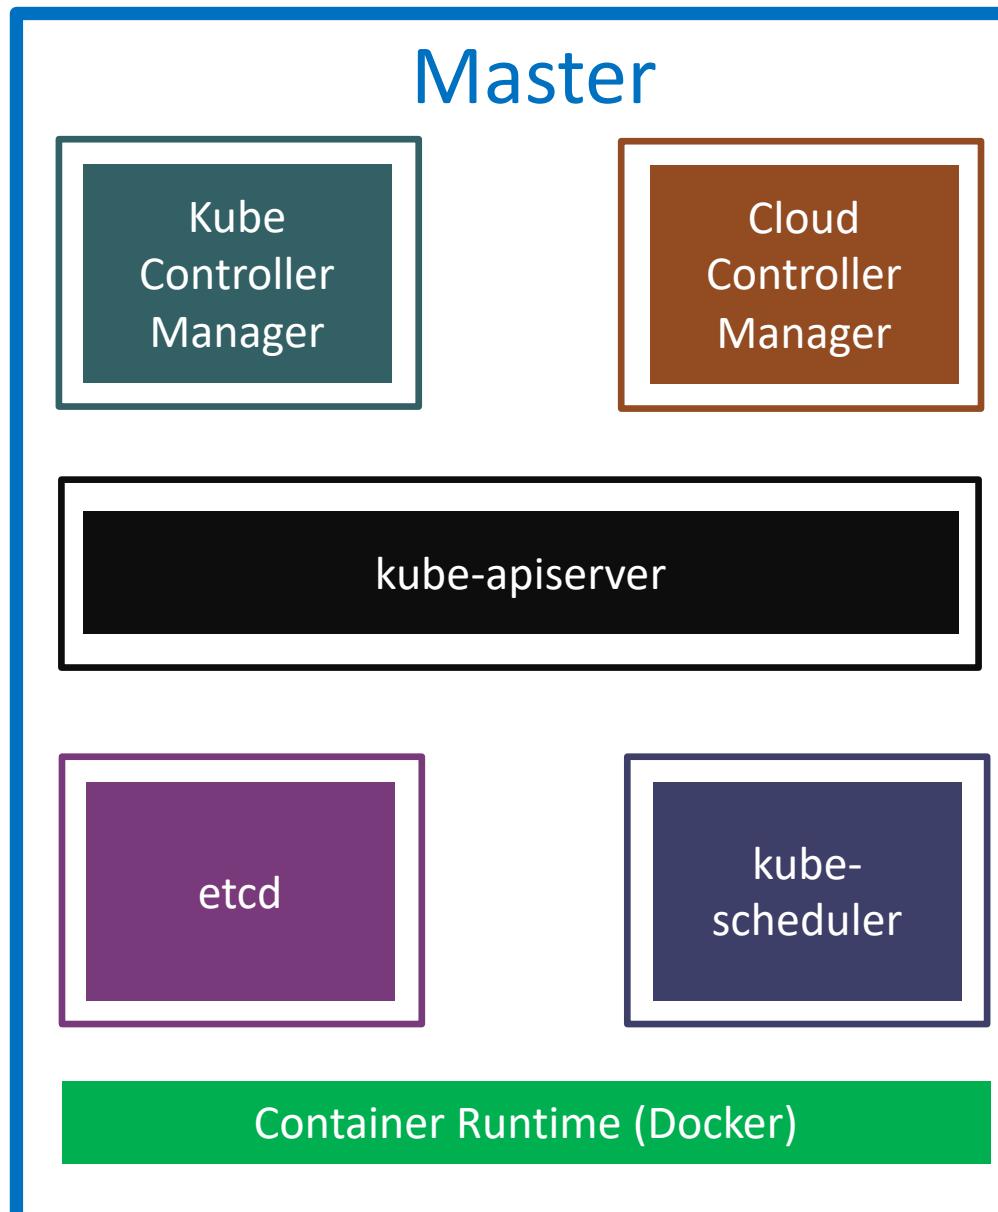
# Kubernetes Architecture



# Kubernetes - Architecture



# Kubernetes Architecture - Master



- **kube-apiserver**

- It acts as **front end** for the Kubernetes control plane. It **exposes** the Kubernetes API
- Command line tools (like kubectl), Users and even Master components (scheduler, controller manager, etcd) and Worker node components like (Kubelet) **everything talk** with API Server.

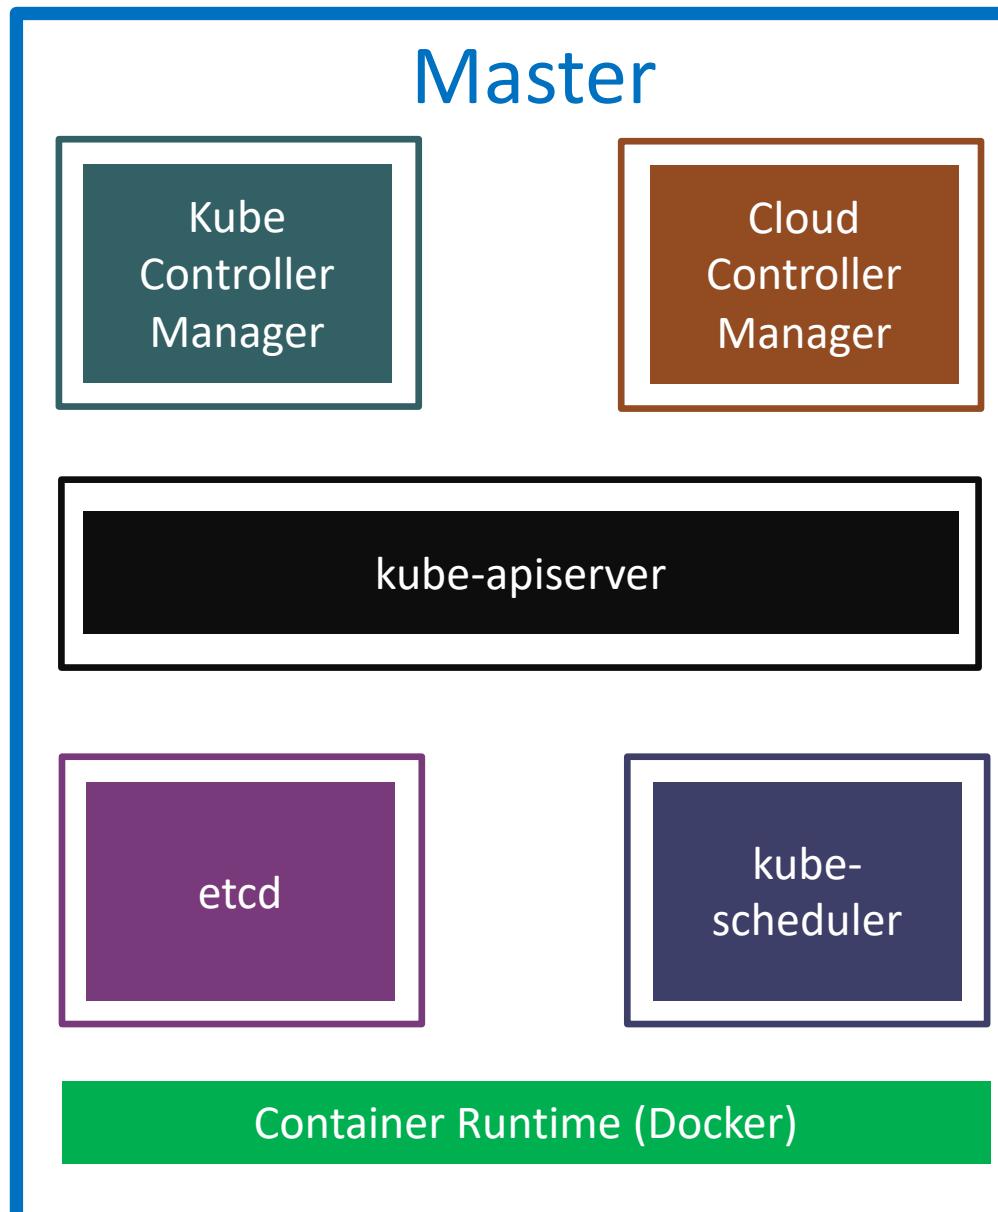
- **etcd**

- Consistent and highly-available **key value store** used as Kubernetes' **backing store** for all cluster data.
- It **stores** all the masters and worker node information.

- **kube-scheduler**

- Scheduler is responsible for distributing containers across multiple nodes.
- It watches for newly created Pods with no assigned node, and selects a node for them to run on.

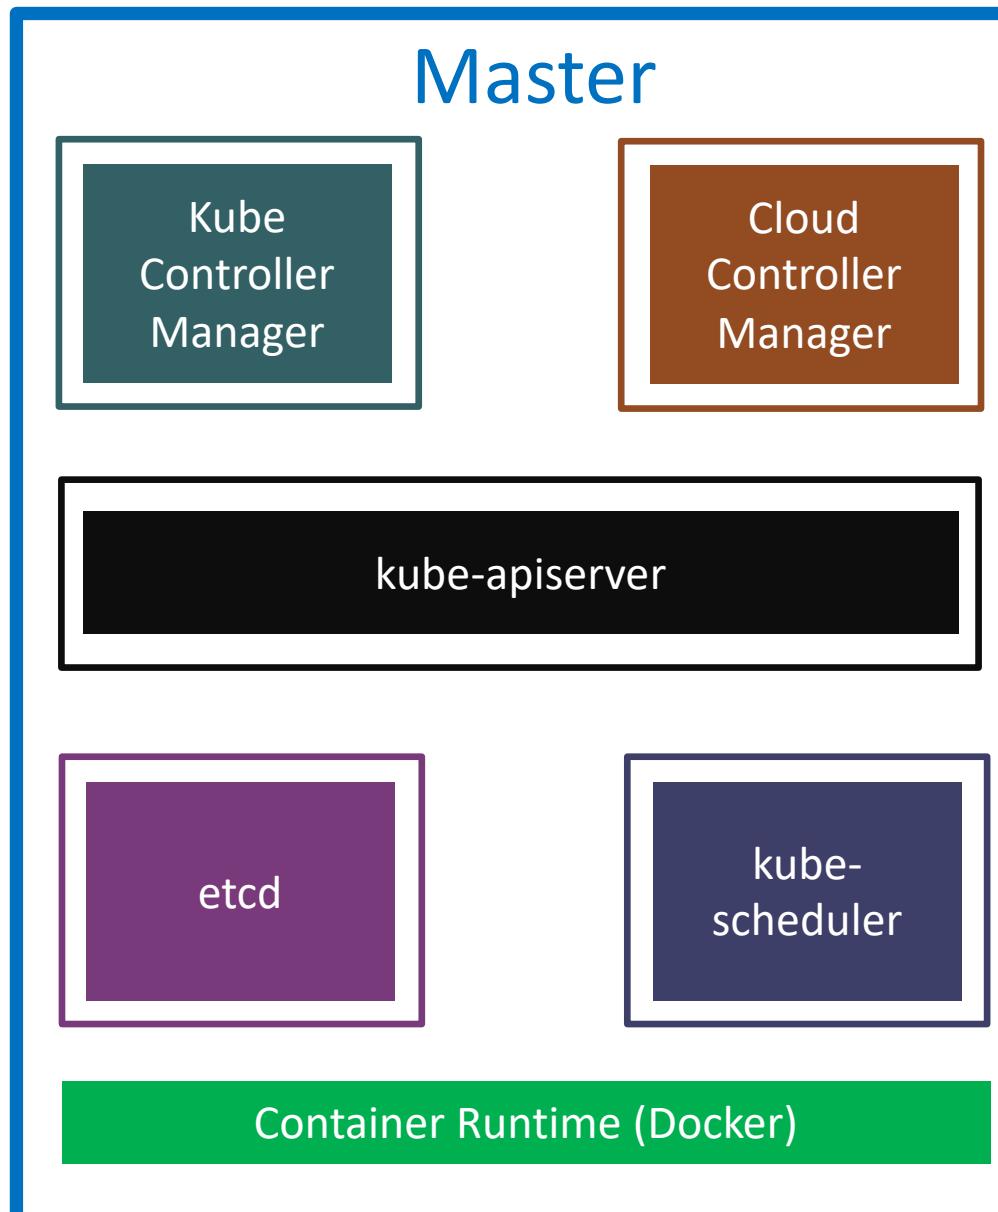
# Kubernetes Architecture - Master



- **kube-controller-manager**

- Controllers are responsible for noticing and responding when nodes, containers or endpoints go down. They make decisions to bring up new containers in such cases.
- **Node Controller**: Responsible for noticing and responding when **nodes go down**.
- **Replication Controller**: Responsible for maintaining the **correct number of pods** for every replication controller object in the system.
- **Endpoints Controller**: **Populates** the Endpoints object (that is, joins Services & Pods)
- **Service Account & Token Controller**: Creates default accounts and API Access for **new namespaces**.

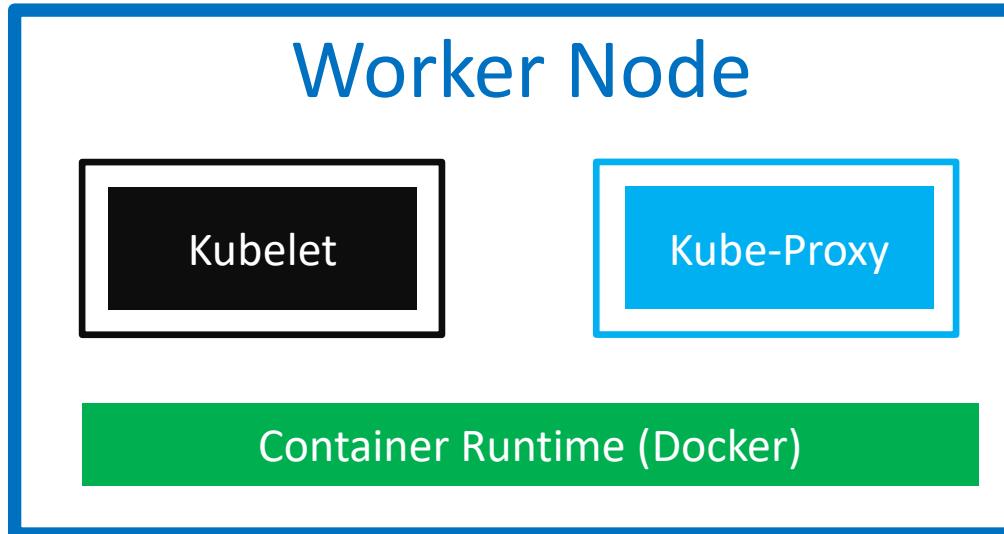
# Kubernetes Architecture - Master



- **cloud-controller-manager**

- A Kubernetes control plane component that embeds **cloud-specific control logic**.
- It only runs controllers that are **specific to your cloud provider**.
- **On-Premise** Kubernetes clusters will not have this component.
- **Node controller**: For **checking** the cloud provider to determine if a node has been deleted in the cloud after it stops responding
- **Route controller**: For setting up **routes** in the underlying cloud infrastructure
- **Service controller**: For creating, updating and deleting cloud provider **load balancer**

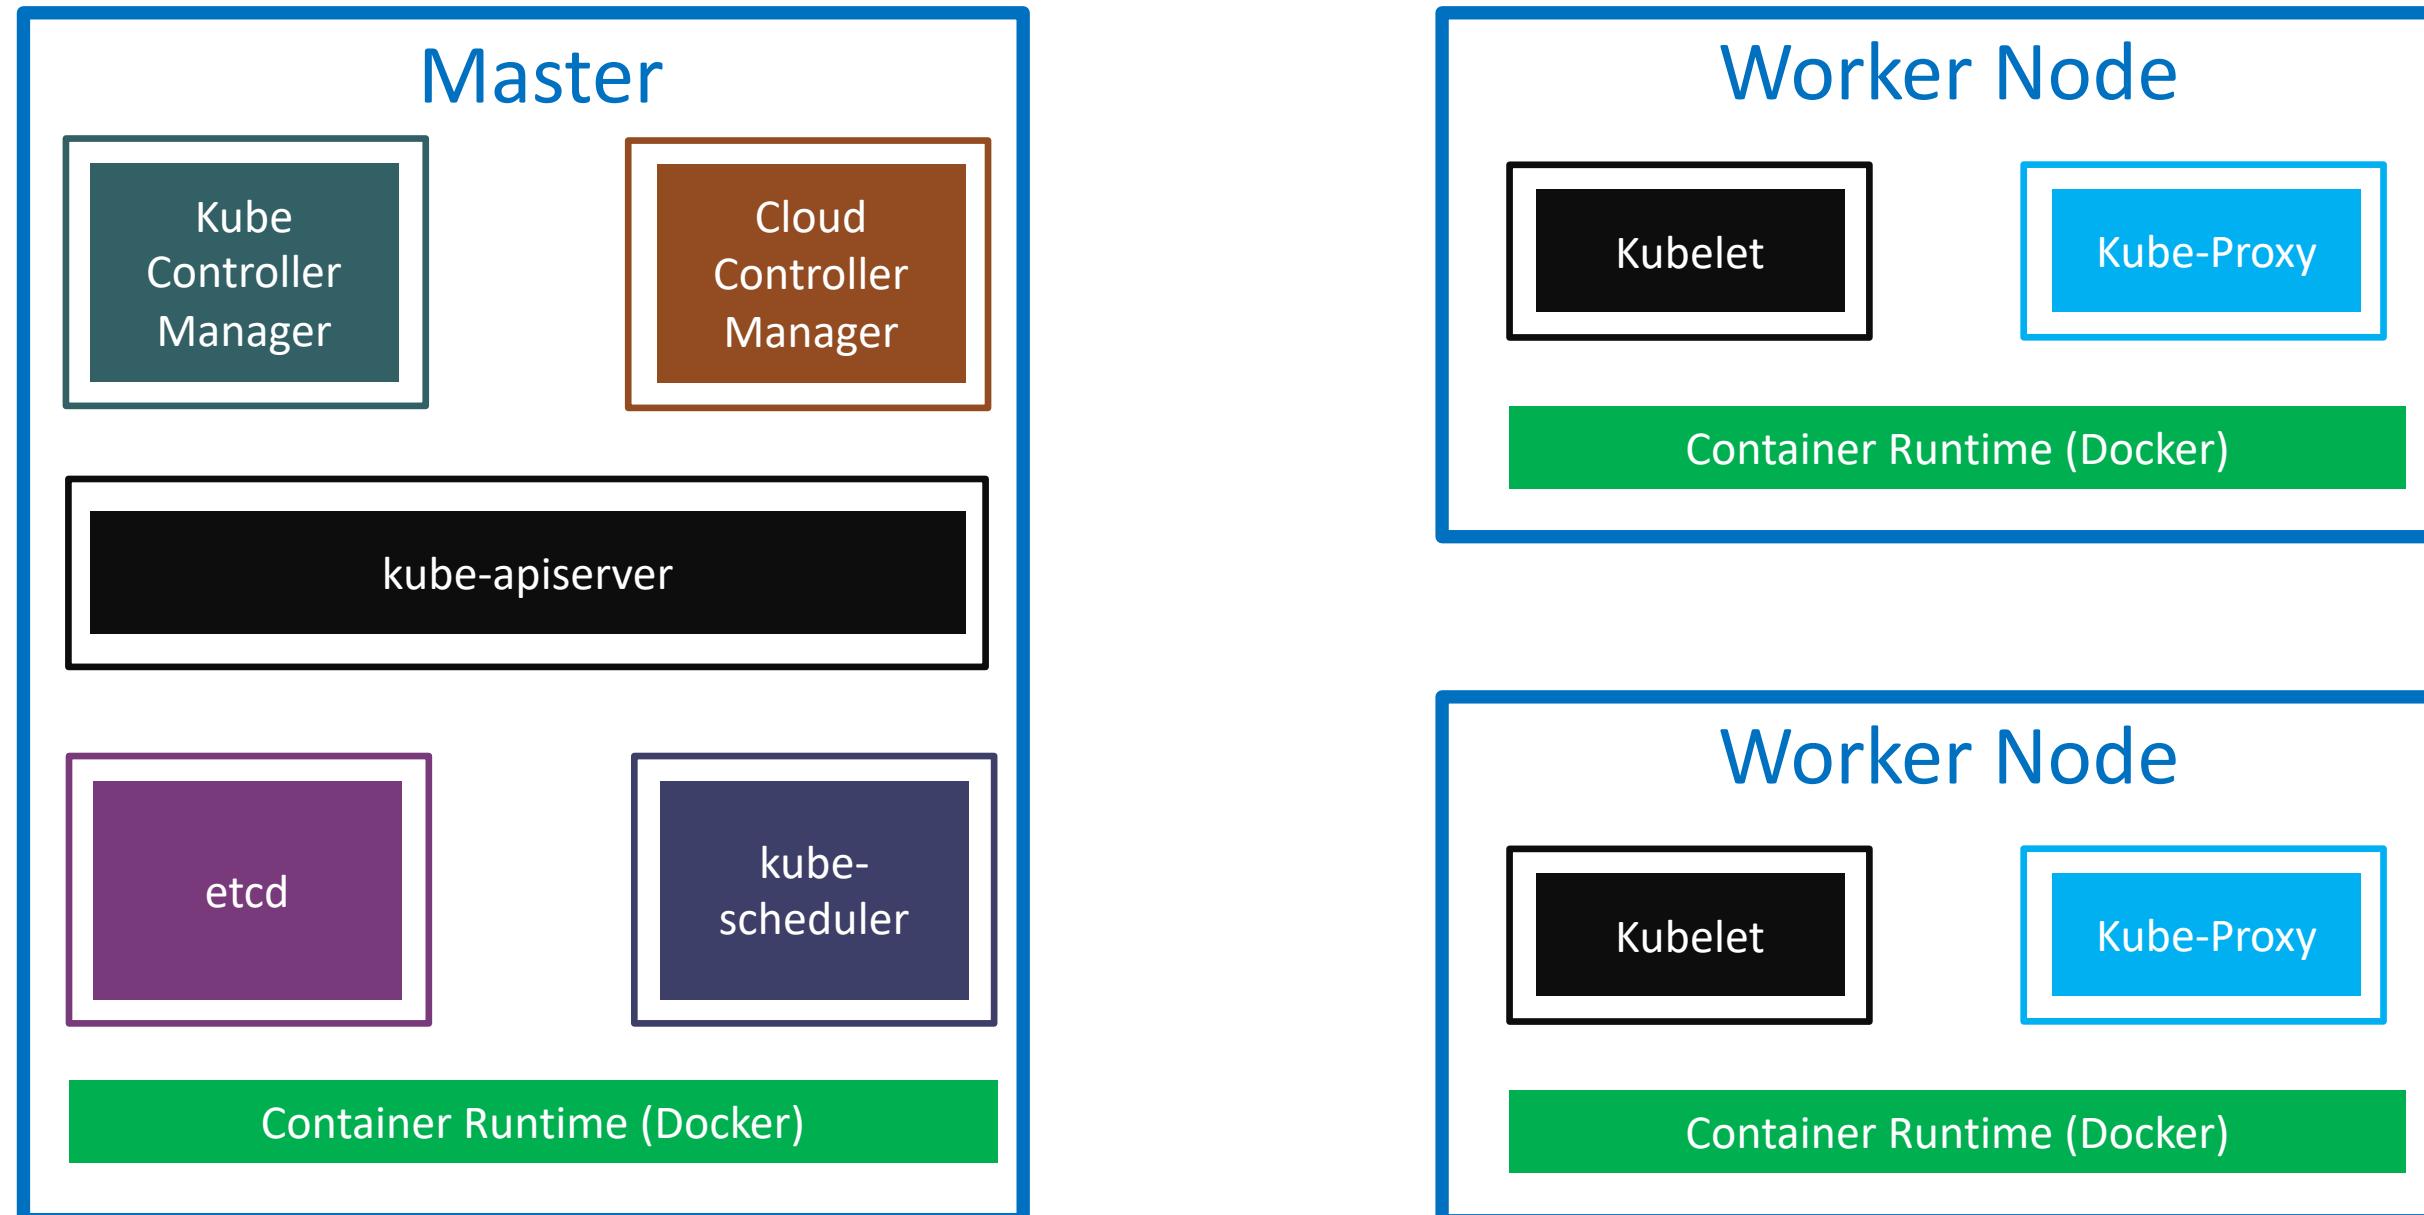
# Kubernetes Architecture – Worker Nodes



- Container Runtime
  - Container Runtime is the **underlying software** where we run all these Kubernetes components.
  - We are using Docker, but we have other runtime options like rkt, container-d etc.

- Kubelet
  - Kubelet is the **agent** that runs on every node in the cluster
  - This agent is **responsible** for making sure that containers are running in a Pod on a node.
- Kube-Proxy
  - It is a **network proxy** that runs on each node in your cluster.
  - It maintains **network rules** on nodes
  - In short, these network rules allow network communication to your Pods from network sessions inside or outside of your cluster.

# Kubernetes - Architecture

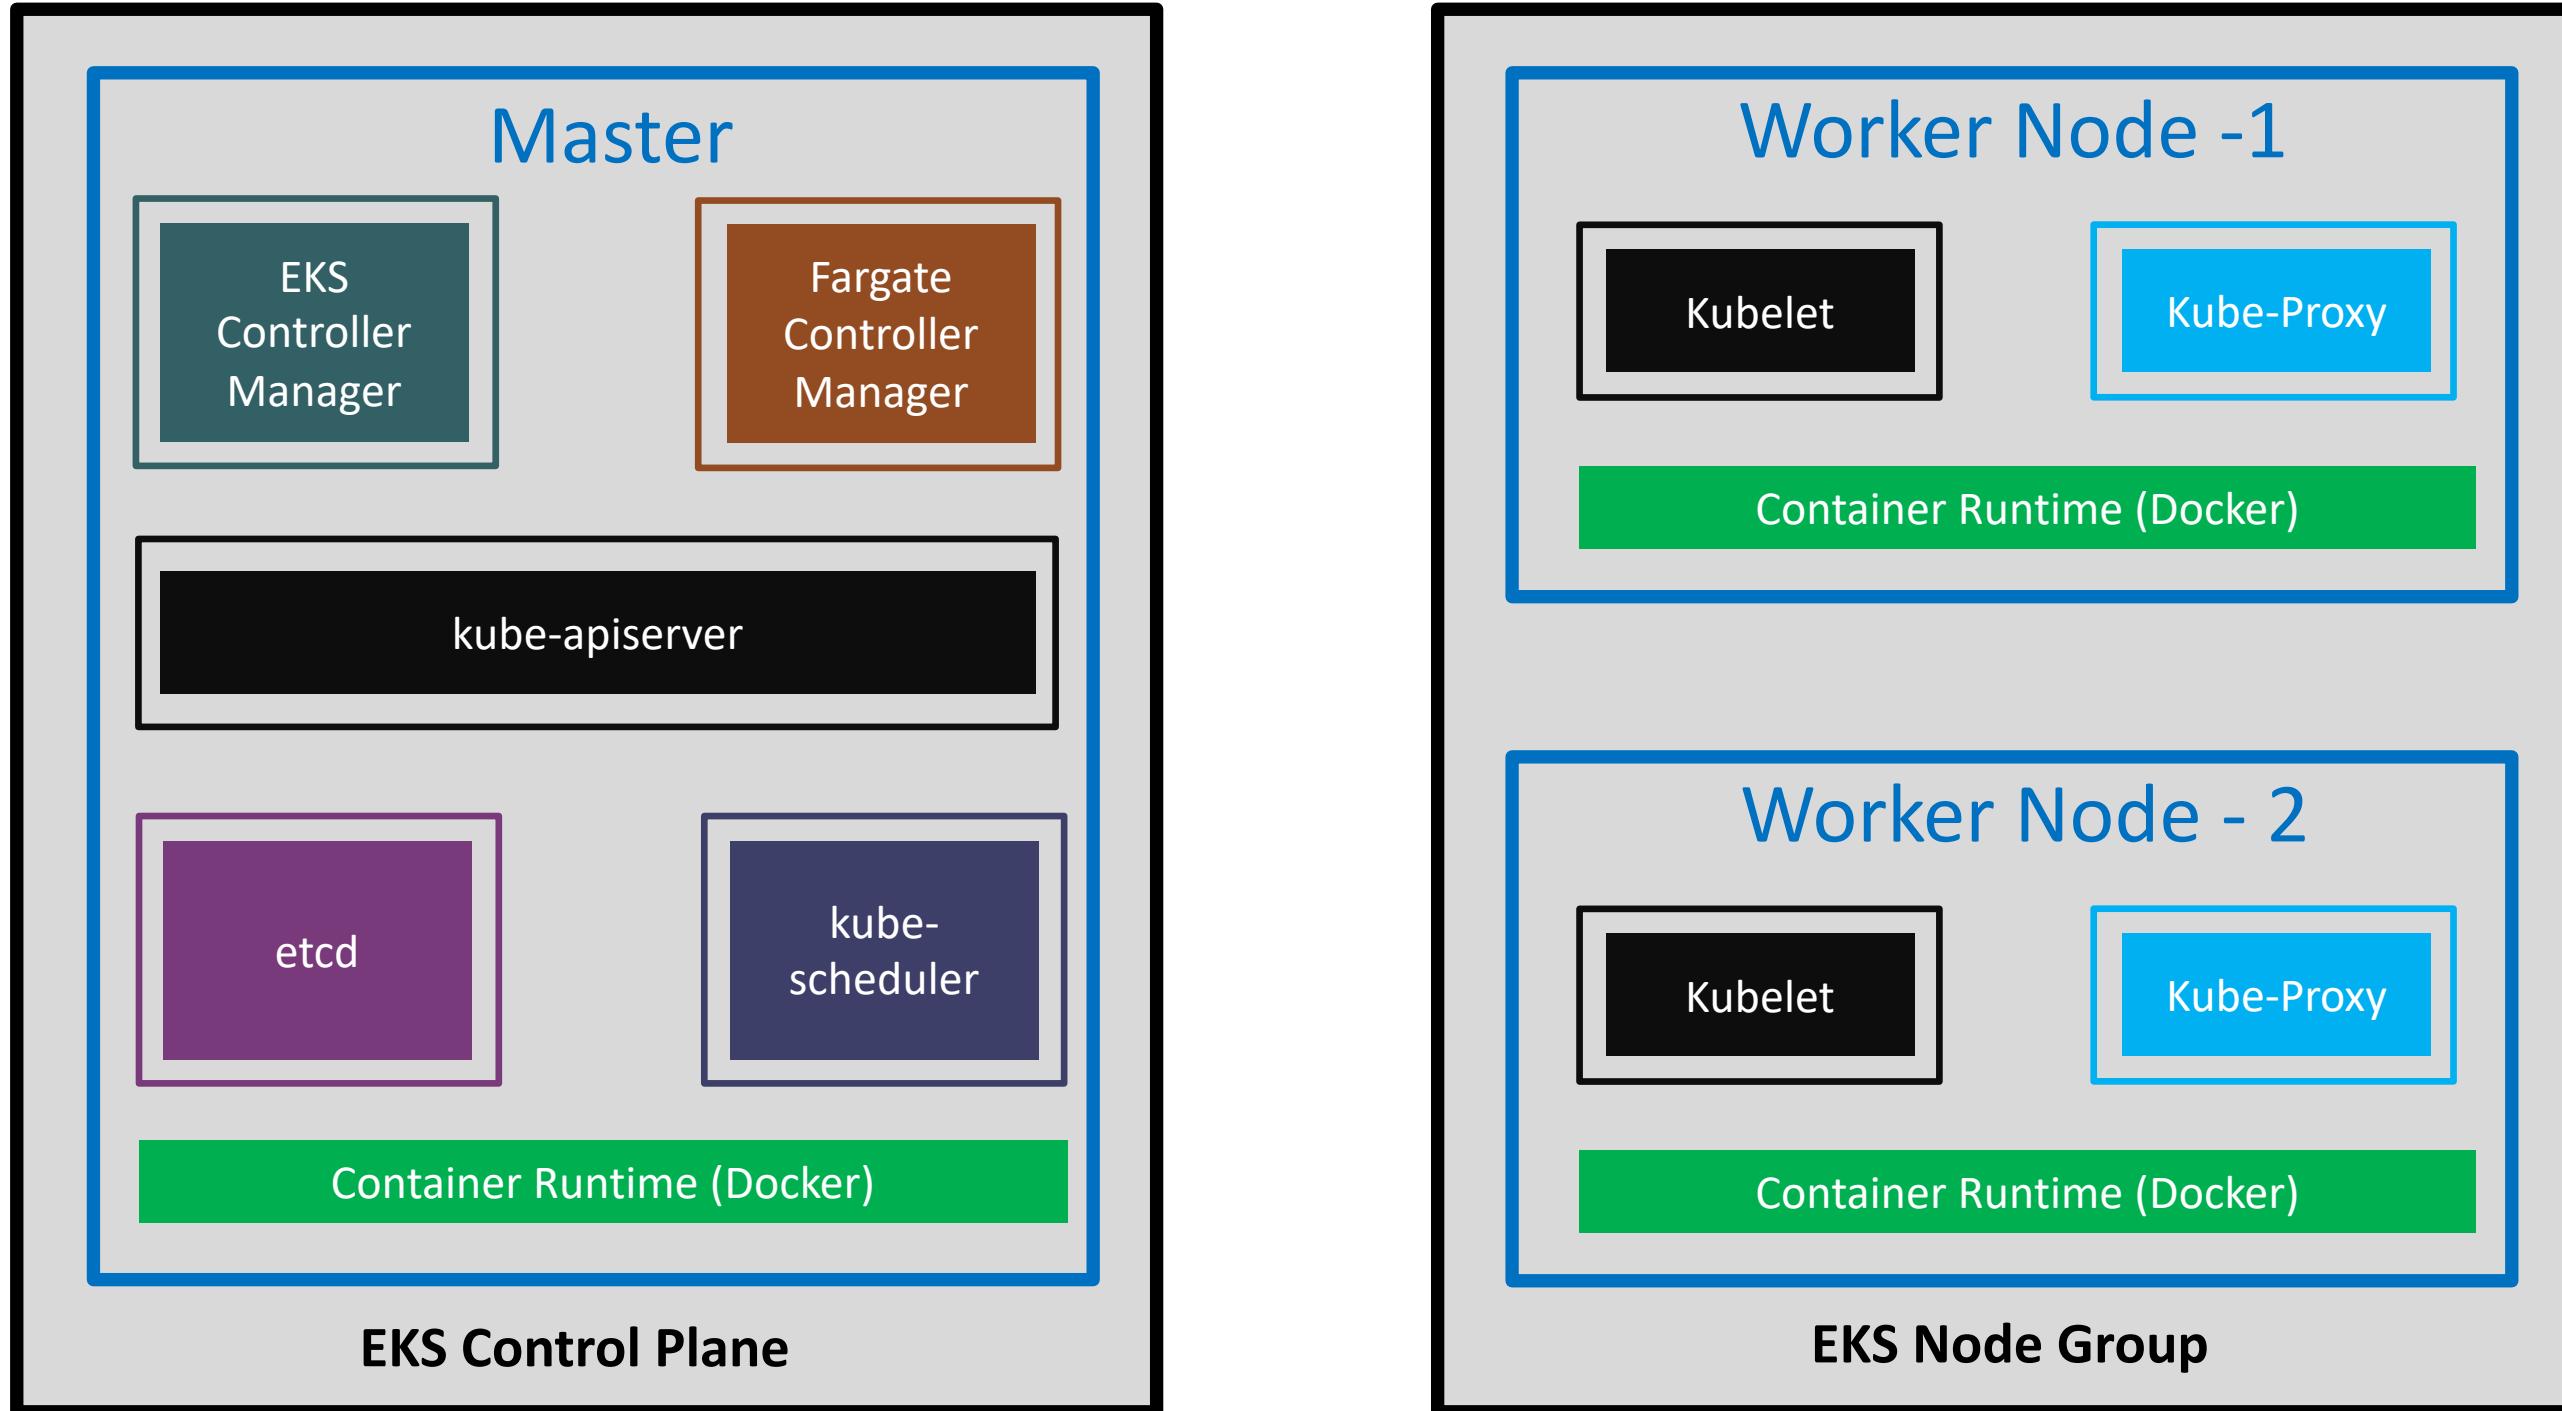




# AWS EKS Cluster



# EKS Kubernetes - Architecture

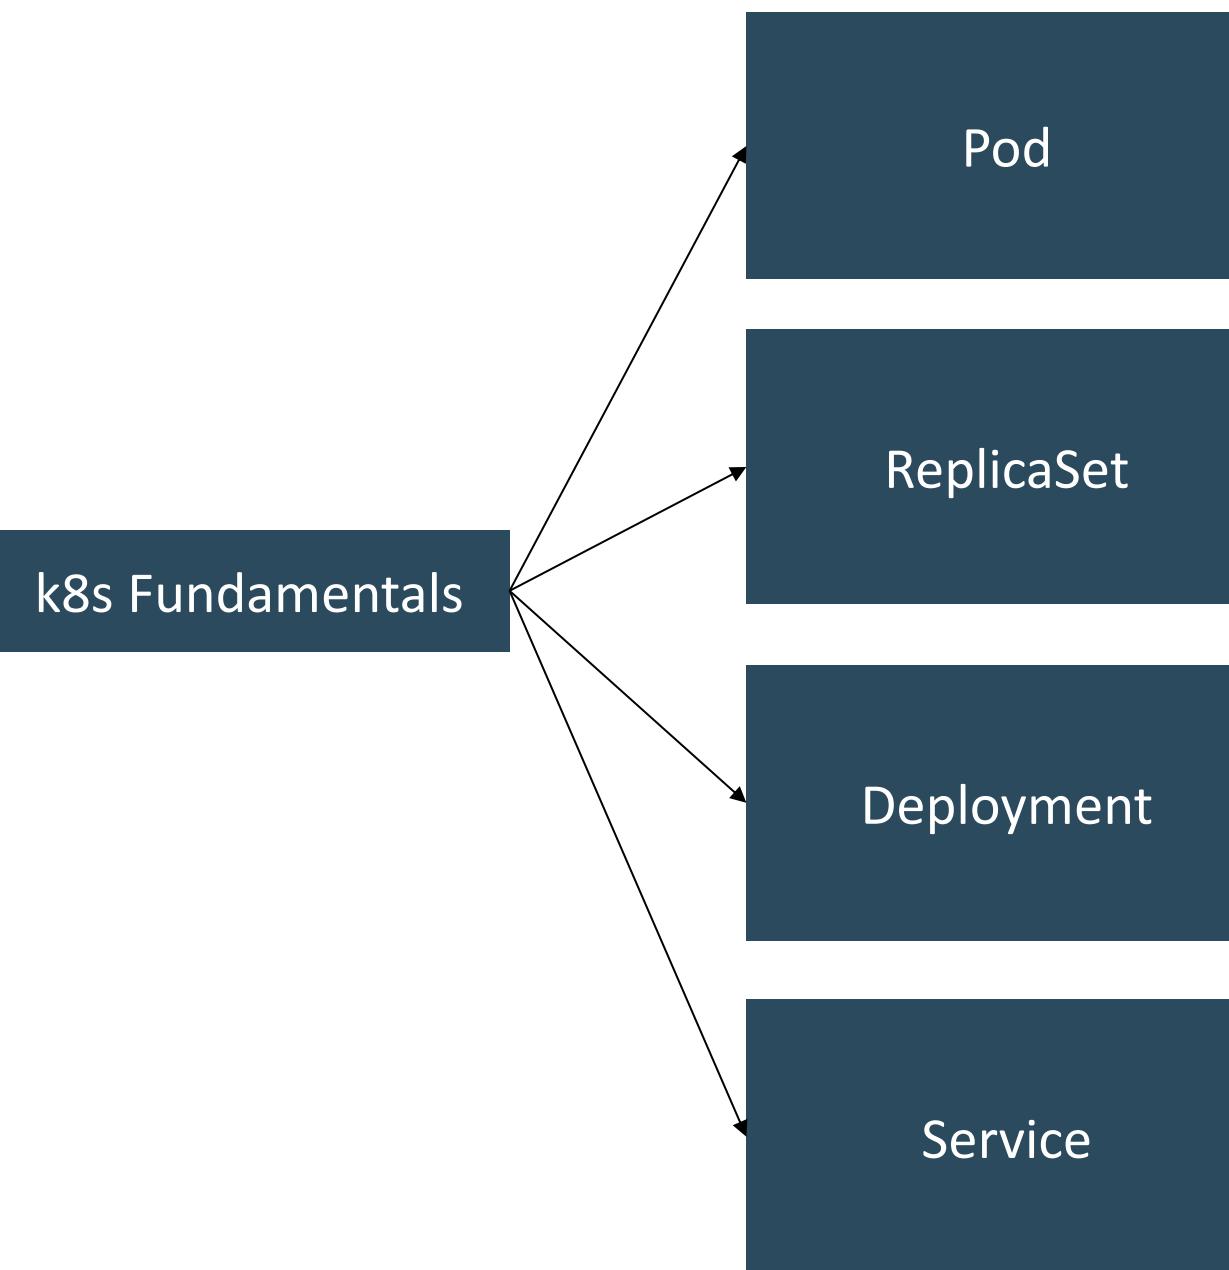


# Kubernetes Fundamentals

## Pod, ReplicaSet, Deployment & Service



# Kubernetes - Fundamentals



A POD is a single instance of an Application.  
A POD is the smallest object, that you can create in Kubernetes.

A ReplicaSet will maintain a stable set of replica Pods running at any given time.  
In short, it is often used to guarantee the availability of a specified number of identical Pods

A Deployment runs multiple replicas of your application and automatically replaces any instances that fail or become unresponsive.  
Rollout & rollback changes to applications. Deployments are well-suited for stateless applications.

A service is an abstraction for pods, providing a stable, so called virtual IP (VIP) address.  
In simple terms, service sits Infront of a POD and acts as a load balancer.

# Kubernetes - Imperative & Declarative

## Kubernetes Fundamentals

Imperative

Declarative

kubectl

Pod

ReplicaSet

Deployment

Service

YAML & kubectl

Pod

ReplicaSet

Deployment

Service

# Kubernetes POD

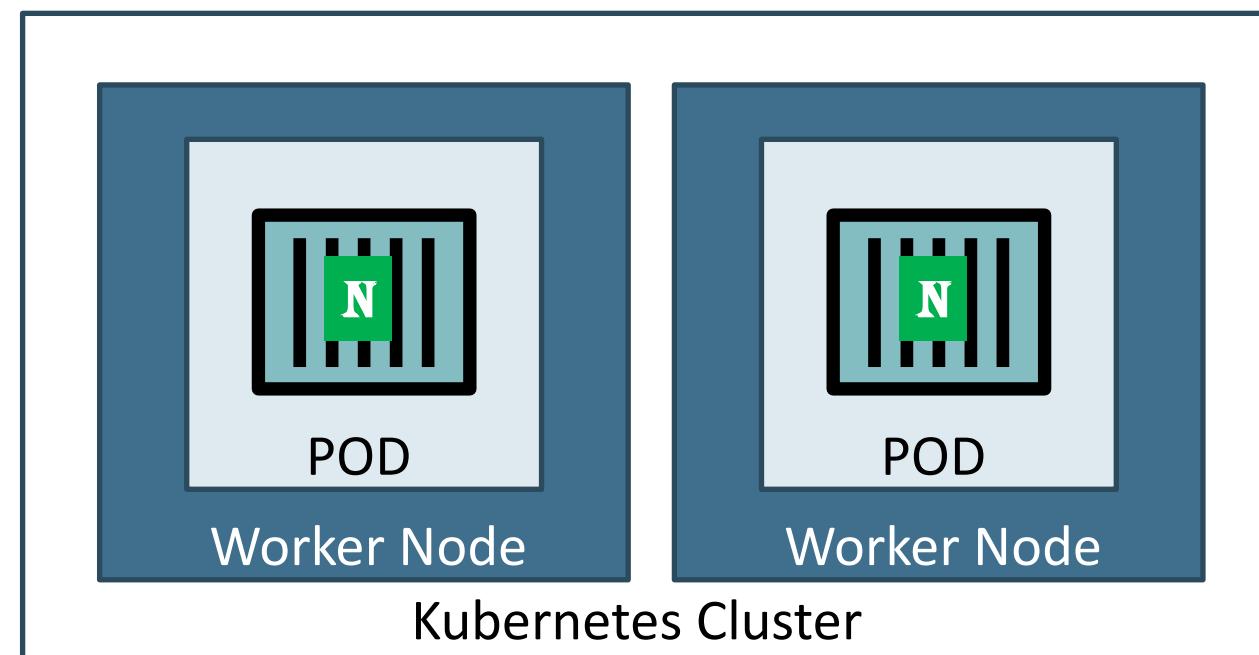


# Kubernetes - POD

- With Kubernetes our core goal will be to deploy our applications in the form of containers on worker nodes in a k8s cluster.
- Kubernetes does not deploy containers directly on the worker nodes.
- Container is encapsulated in to a Kubernetes Object named POD.
- A POD is a single instance of an application.
- A POD is the smallest object that we can create in Kubernetes.

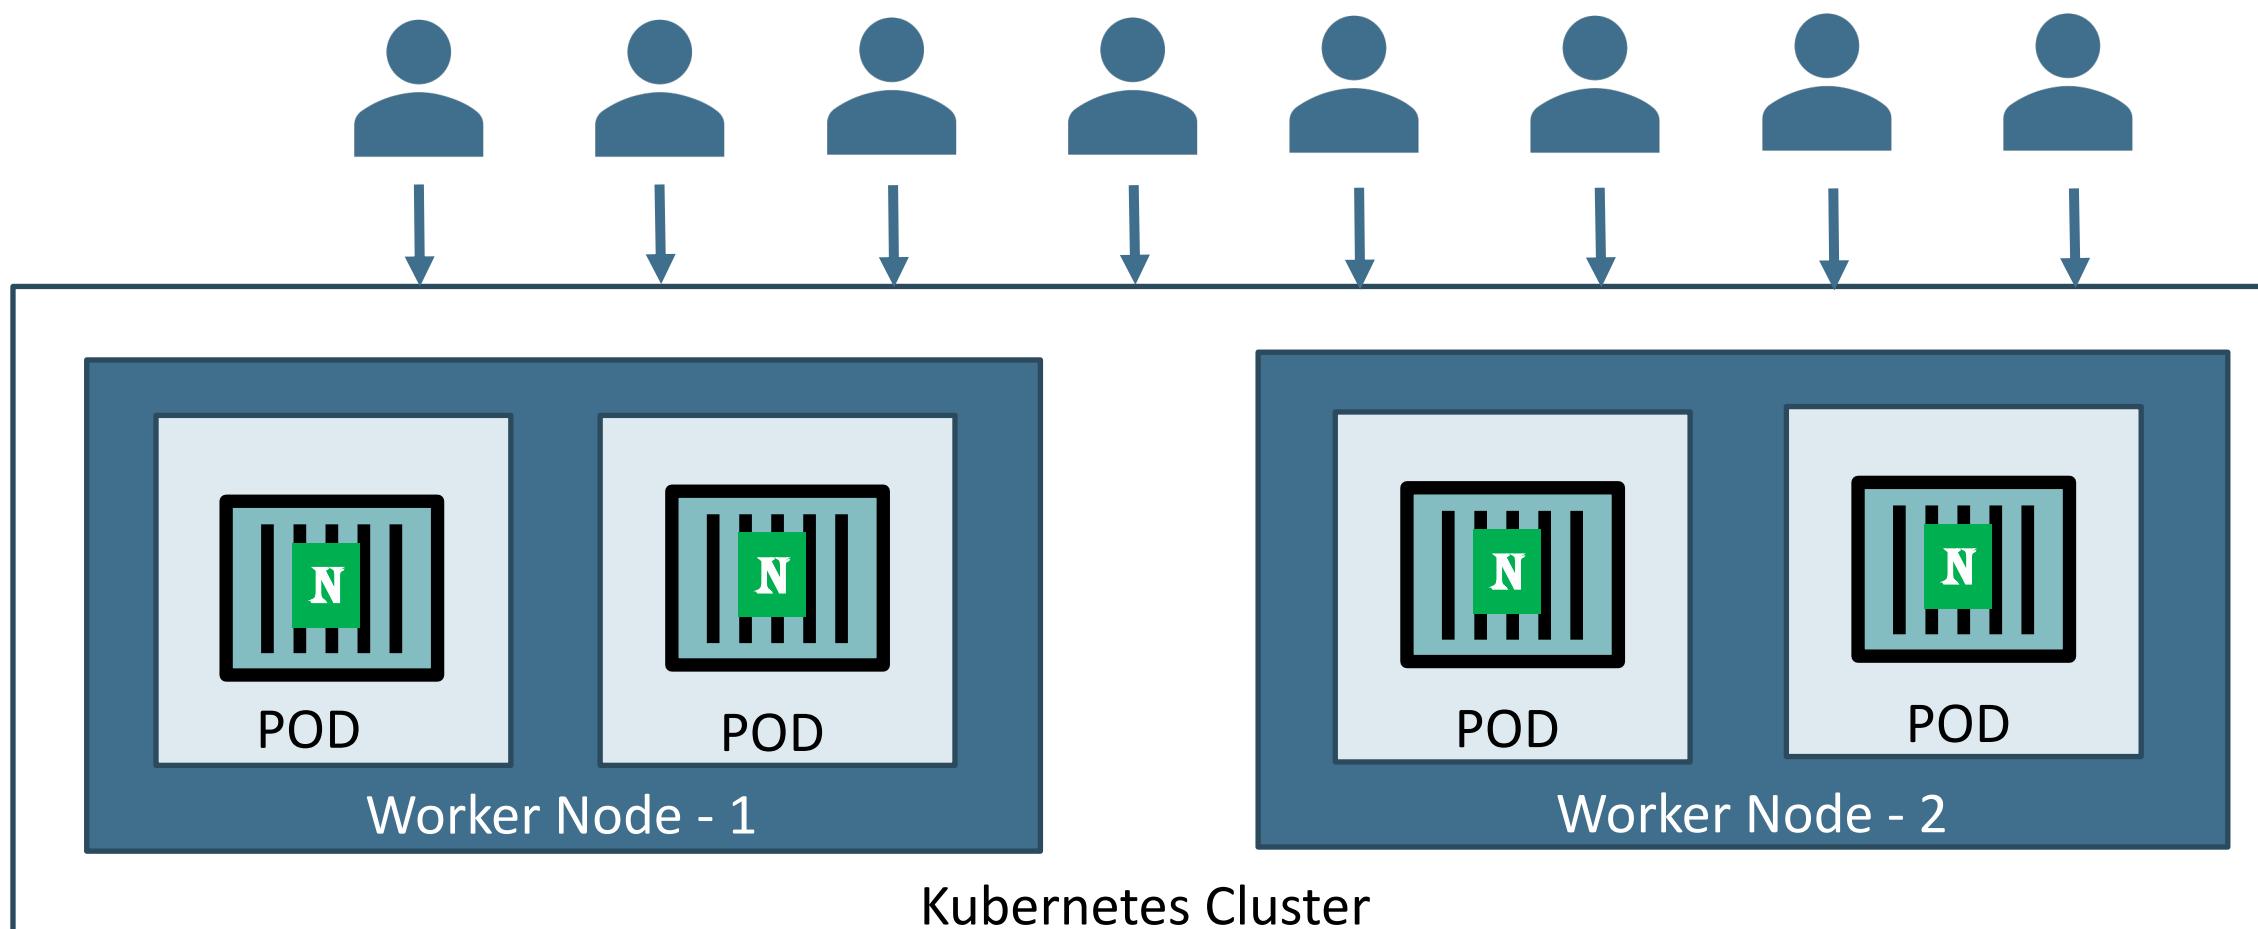


Nginx Container Image



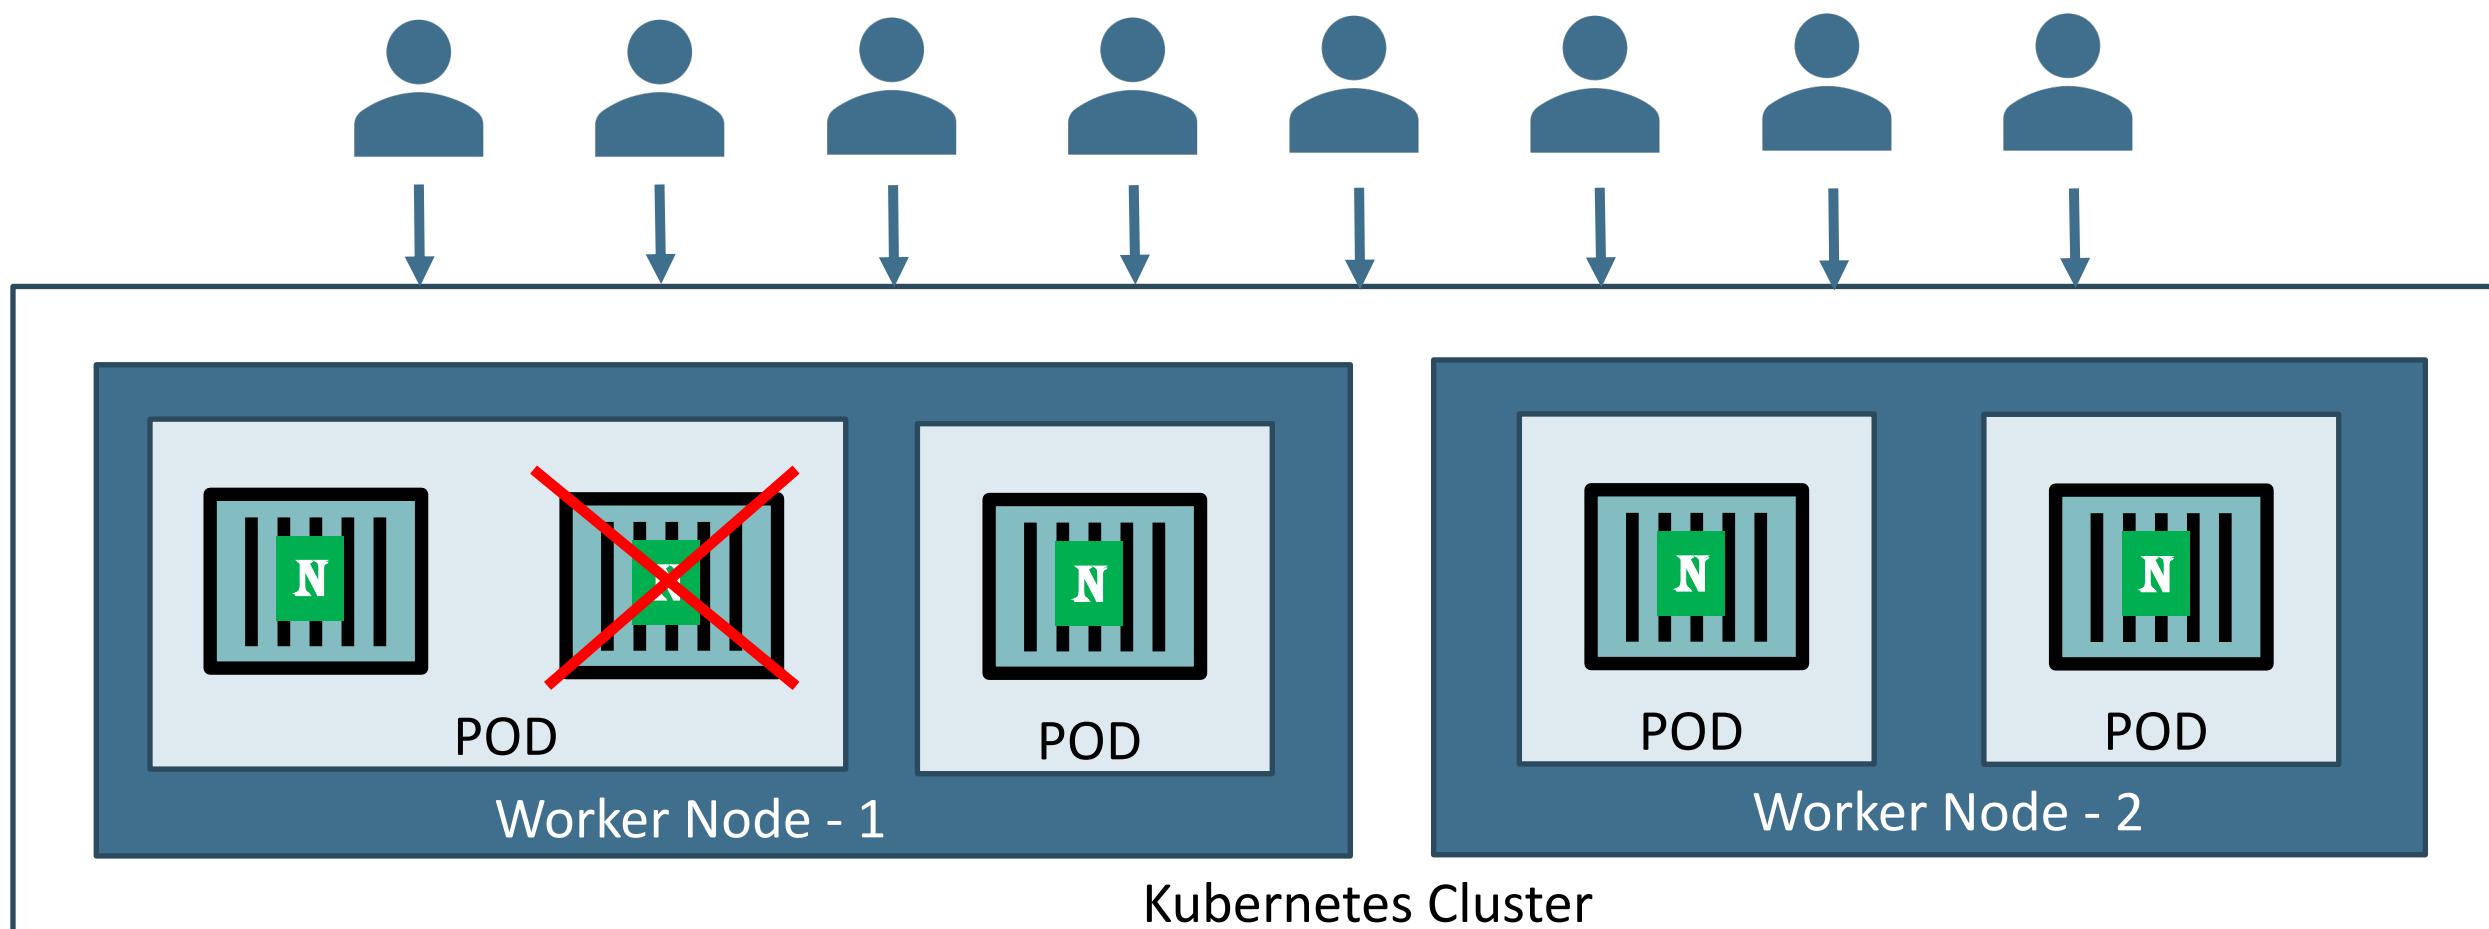
# Kubernetes - POD

- PODs generally have **one to one** relationship with containers.
- To scale up we **create** new POD and to scale down we **delete** the POD.



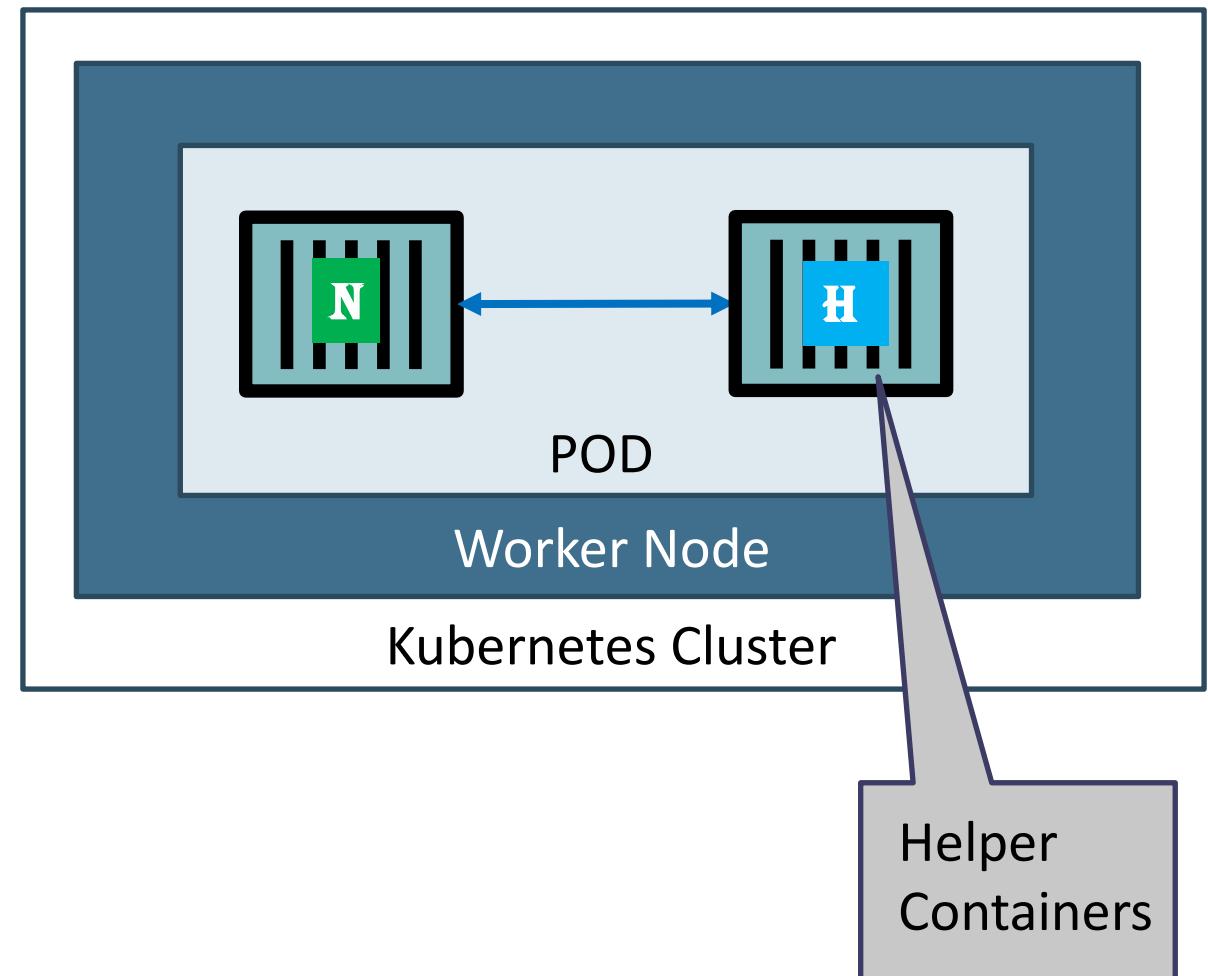
# Kubernetes – PODs

- We **cannot have** multiple containers of **same kind** in a single POD.
- Example: Two NGINX containers in single POD serving same purpose is **not recommended**.



# Kubernetes – Multi-Container Pods

- We can have multiple containers in a single POD, provided **they are not of same kind**.
- Helper Containers (Side-car)
  - Data Pullers: Pull data required by Main Container
  - Data pushers: Push data by collecting from main container (logs)
  - Proxies: Writes static data to html files using Helper container and Reads using Main Container.
- Communication
  - The two containers can easily communicate with each other easily as they share same **network space**.
  - They can also easily share **same storage space**.
- Multi-Container Pods is a **rare use-case** and we will try to focus on core fundamentals.



# Kubernetes PODs Demo



# Kubernetes Services - NodePort



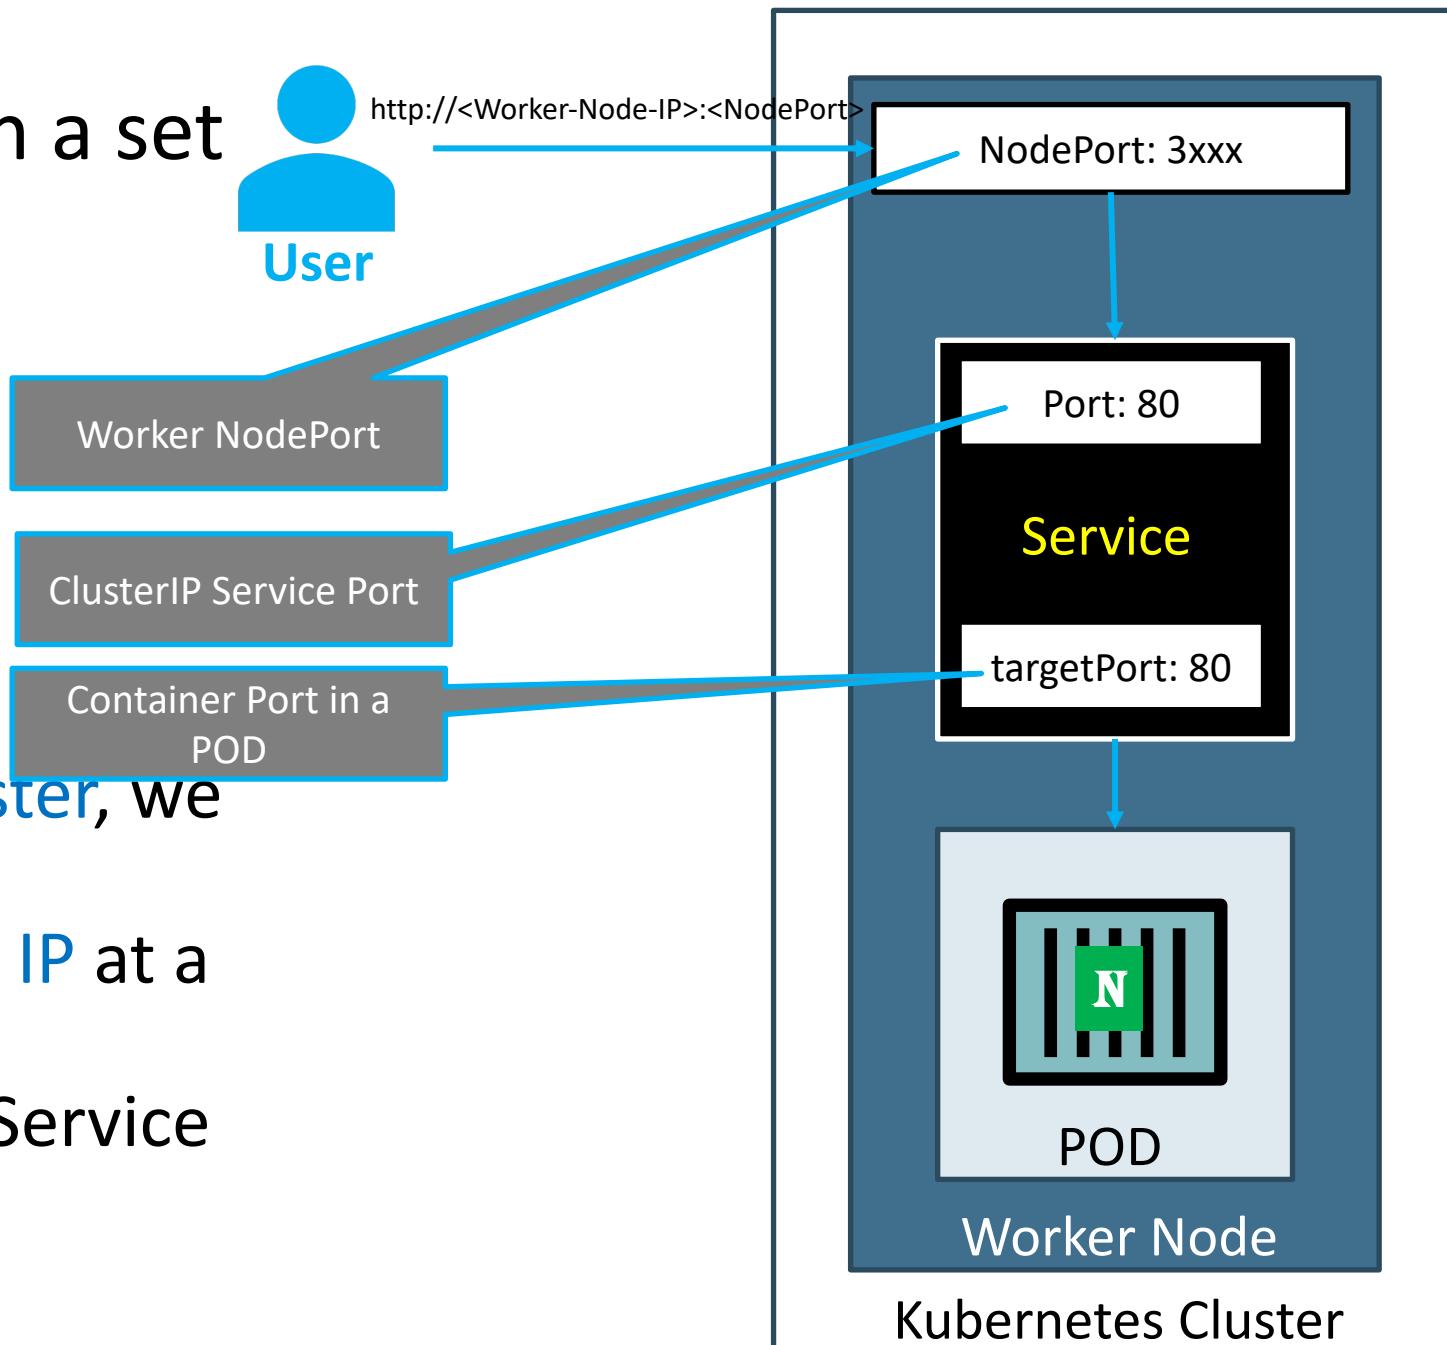
# Kubernetes – Service - NodePort

- We can **expose an application** running on a set of **PODs** using different types of Services available in k8s.

- ClusterIP
- NodePort
- LoadBalancer

## • NodePort Service

- To access our application **outside of k8s cluster**, we can use NodePort service.
- Exposes the Service on each **Worker Node's IP** at a static port (nothing but NodePort).
- A **ClusterIP Service**, to which the **NodePort Service** routes, is **automatically** created.
- Port Range **30000-32767**



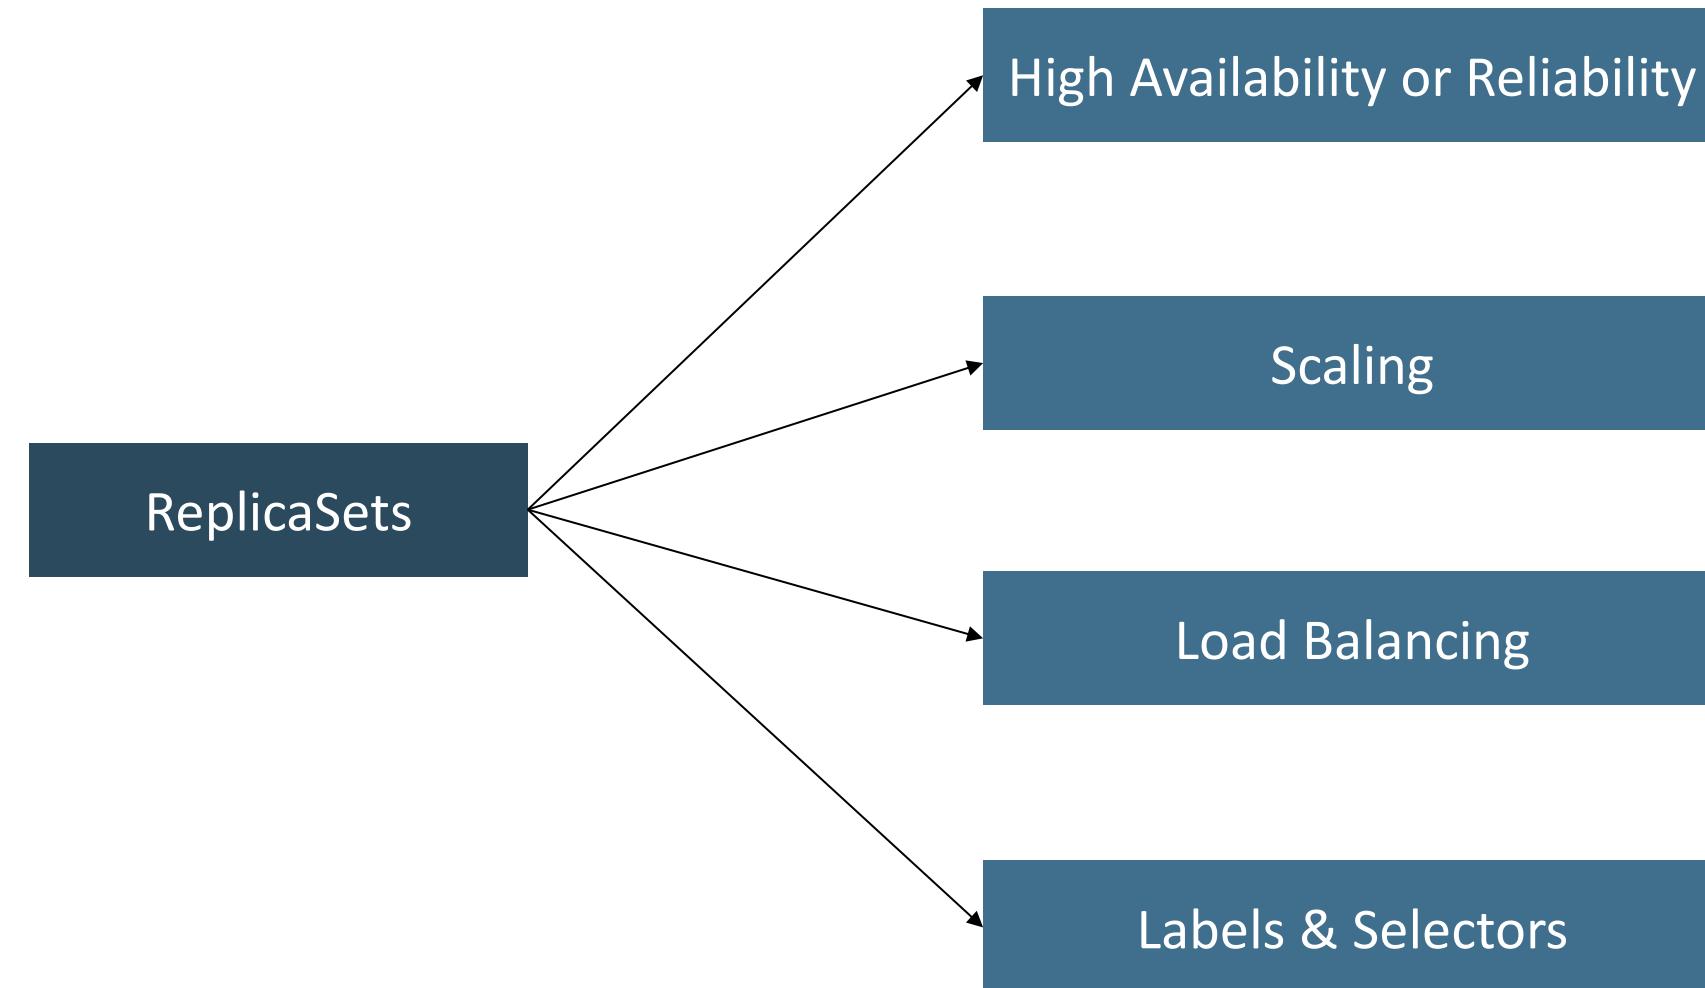
# Kubernetes POD & NodePort Service Demo



# Kubernetes ReplicaSets



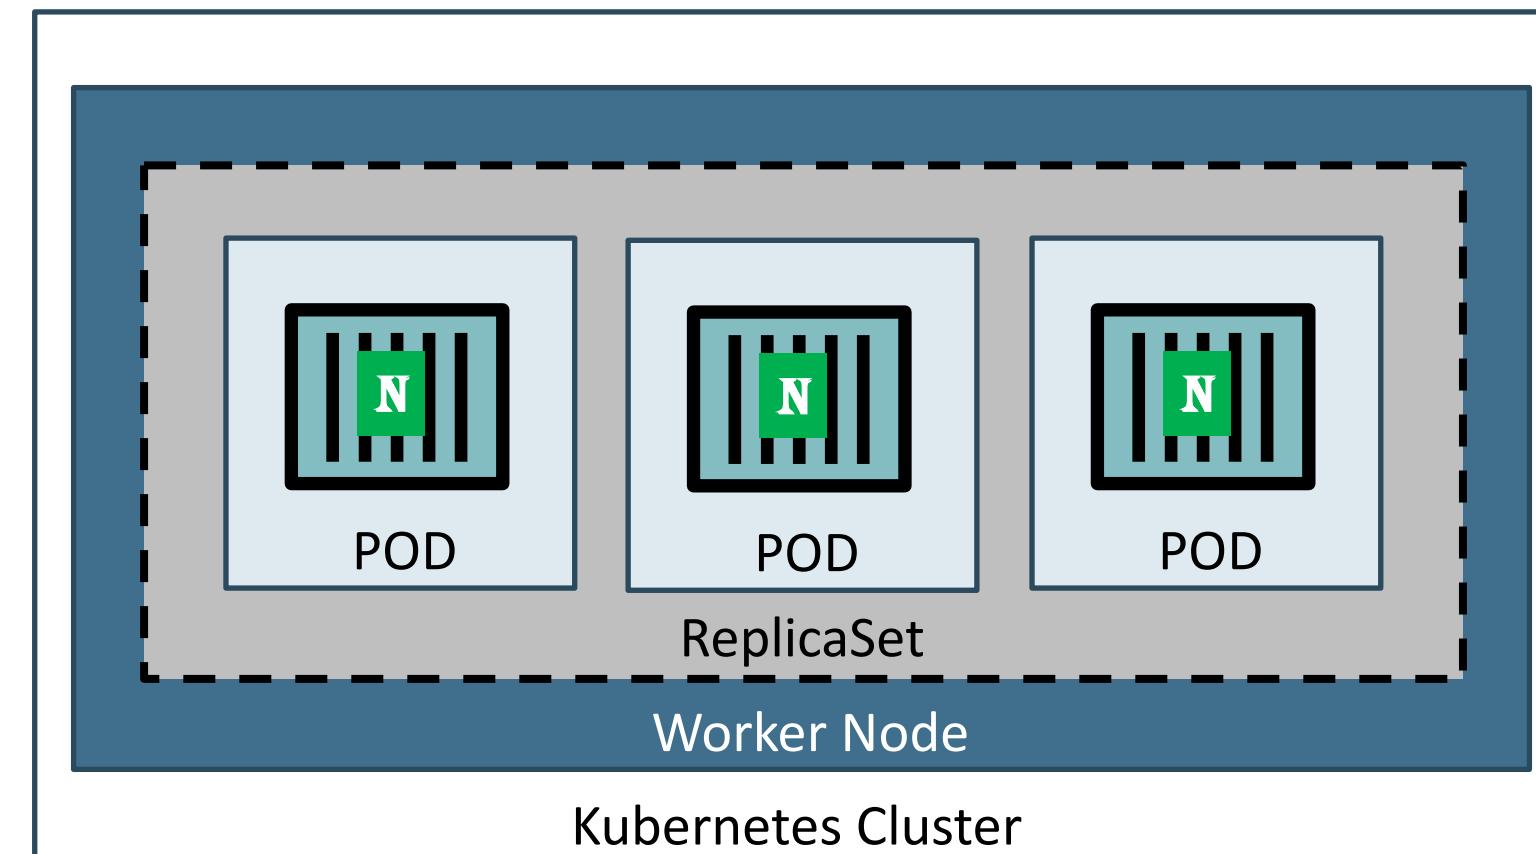
# Kubernetes - ReplicaSets



# Kubernetes – ReplicaSet

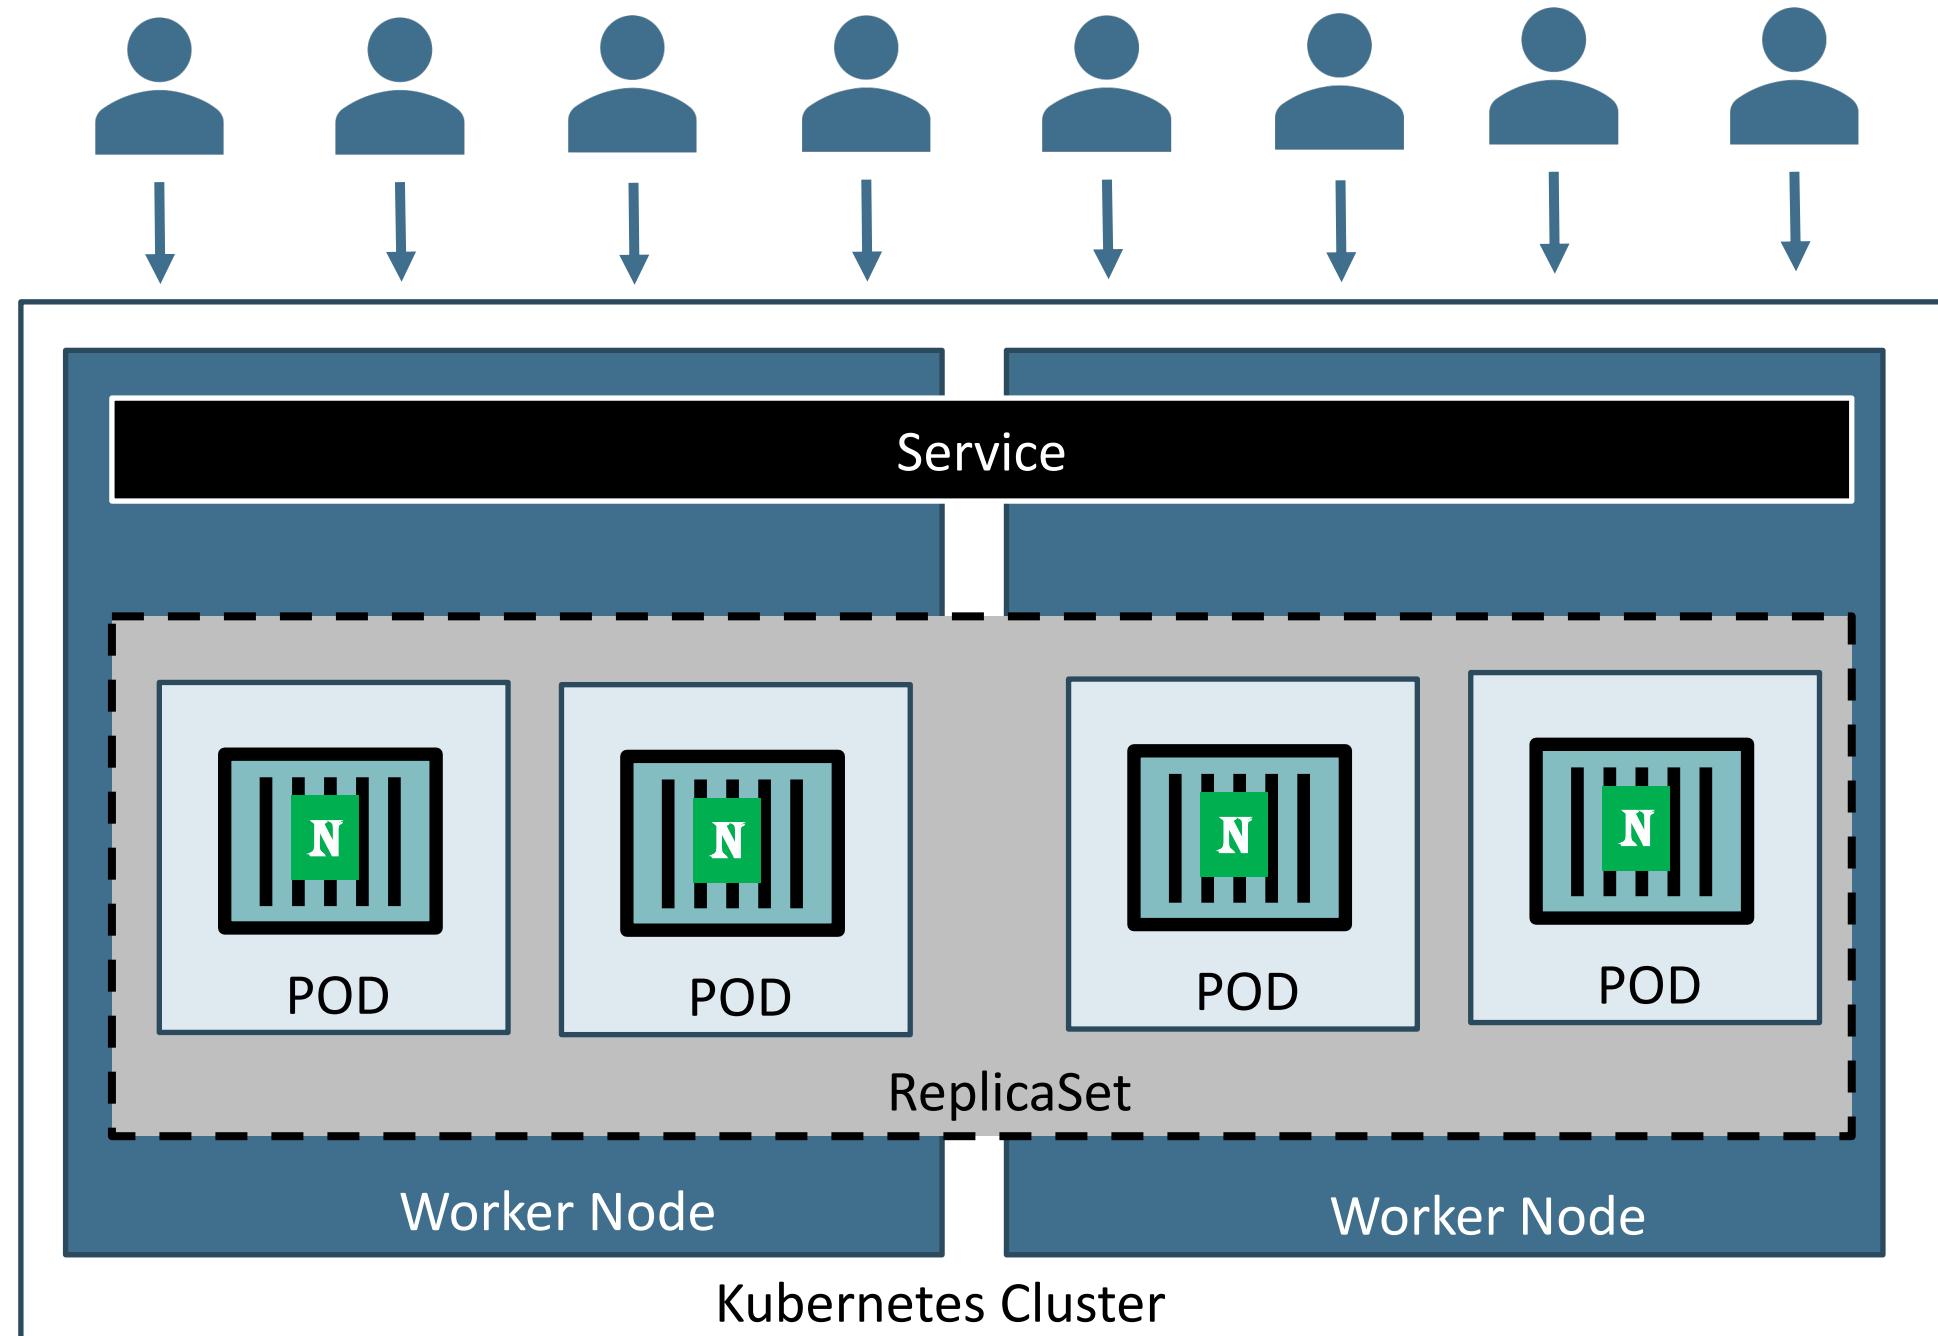
- A ReplicaSet's purpose is to maintain a **stable set of replica Pods** running at any given time.
- If our application crashes (any pod dies), replicaset will **recreate** the pod immediately to ensure the configured number of pods running at any given time.

Reliability  
Or  
High Availability



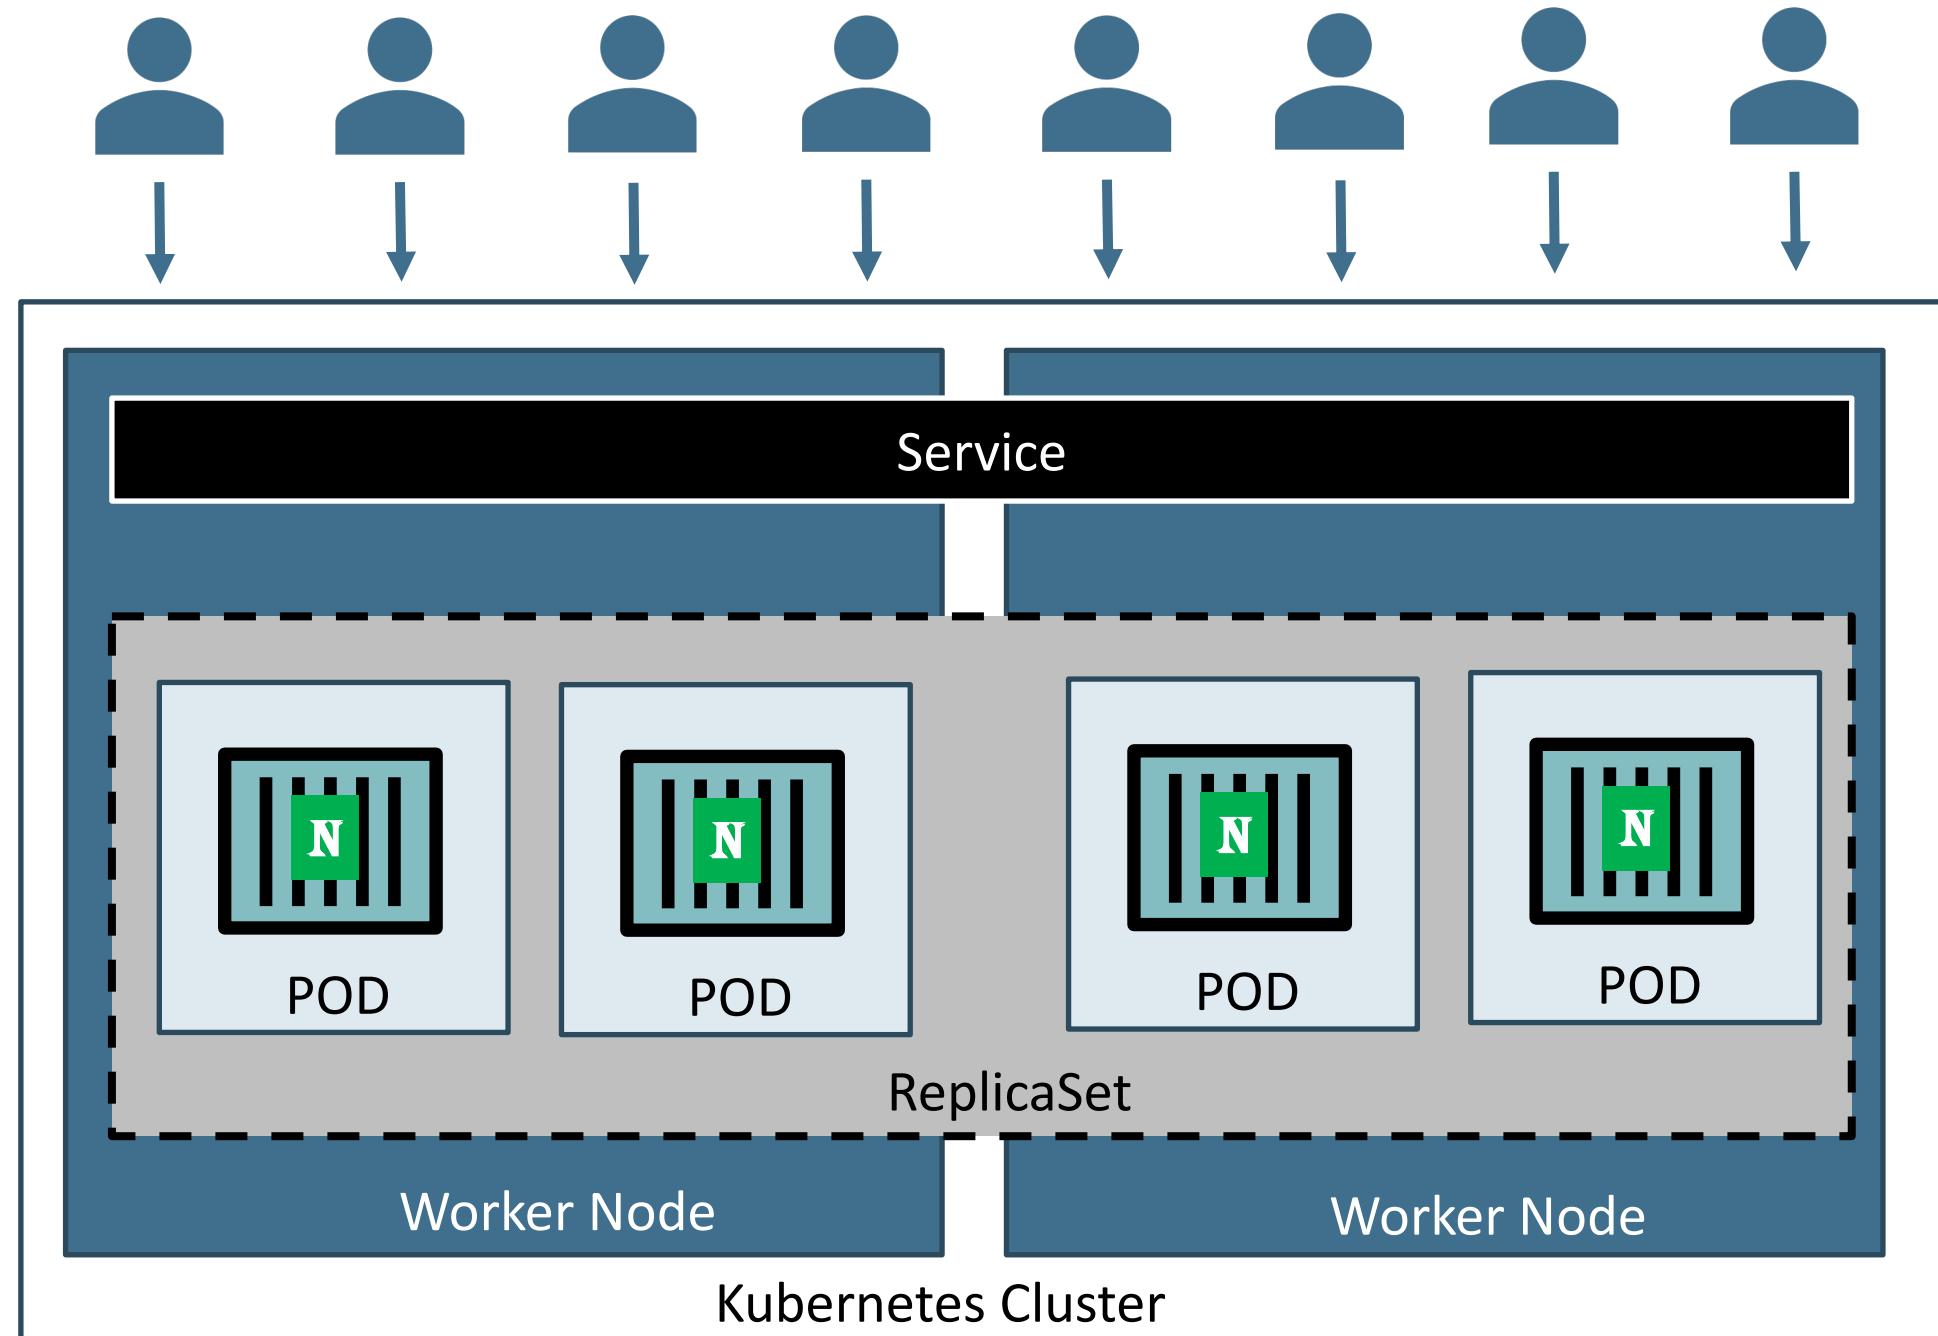
# Kubernetes – ReplicaSet

- Load Balancing
- To avoid overloading of traffic to single pod we can use **load balancing**.
- Kubernetes provides pod load balancing **out of the box** using **Services** for the pods which are part of a ReplicaSet
- **Labels & Selectors** are the **key items** which **ties** all 3 together (Pod, ReplicaSet & Service), we will know in detail when we are writing YAML manifests for these objects



# Kubernetes – ReplicaSet

- Scaling
- When load become too much for the number of existing pods, Kubernetes enables us to easily **scale** up our application, adding additional pods as needed.
- This is going to be **seamless and super quick**.



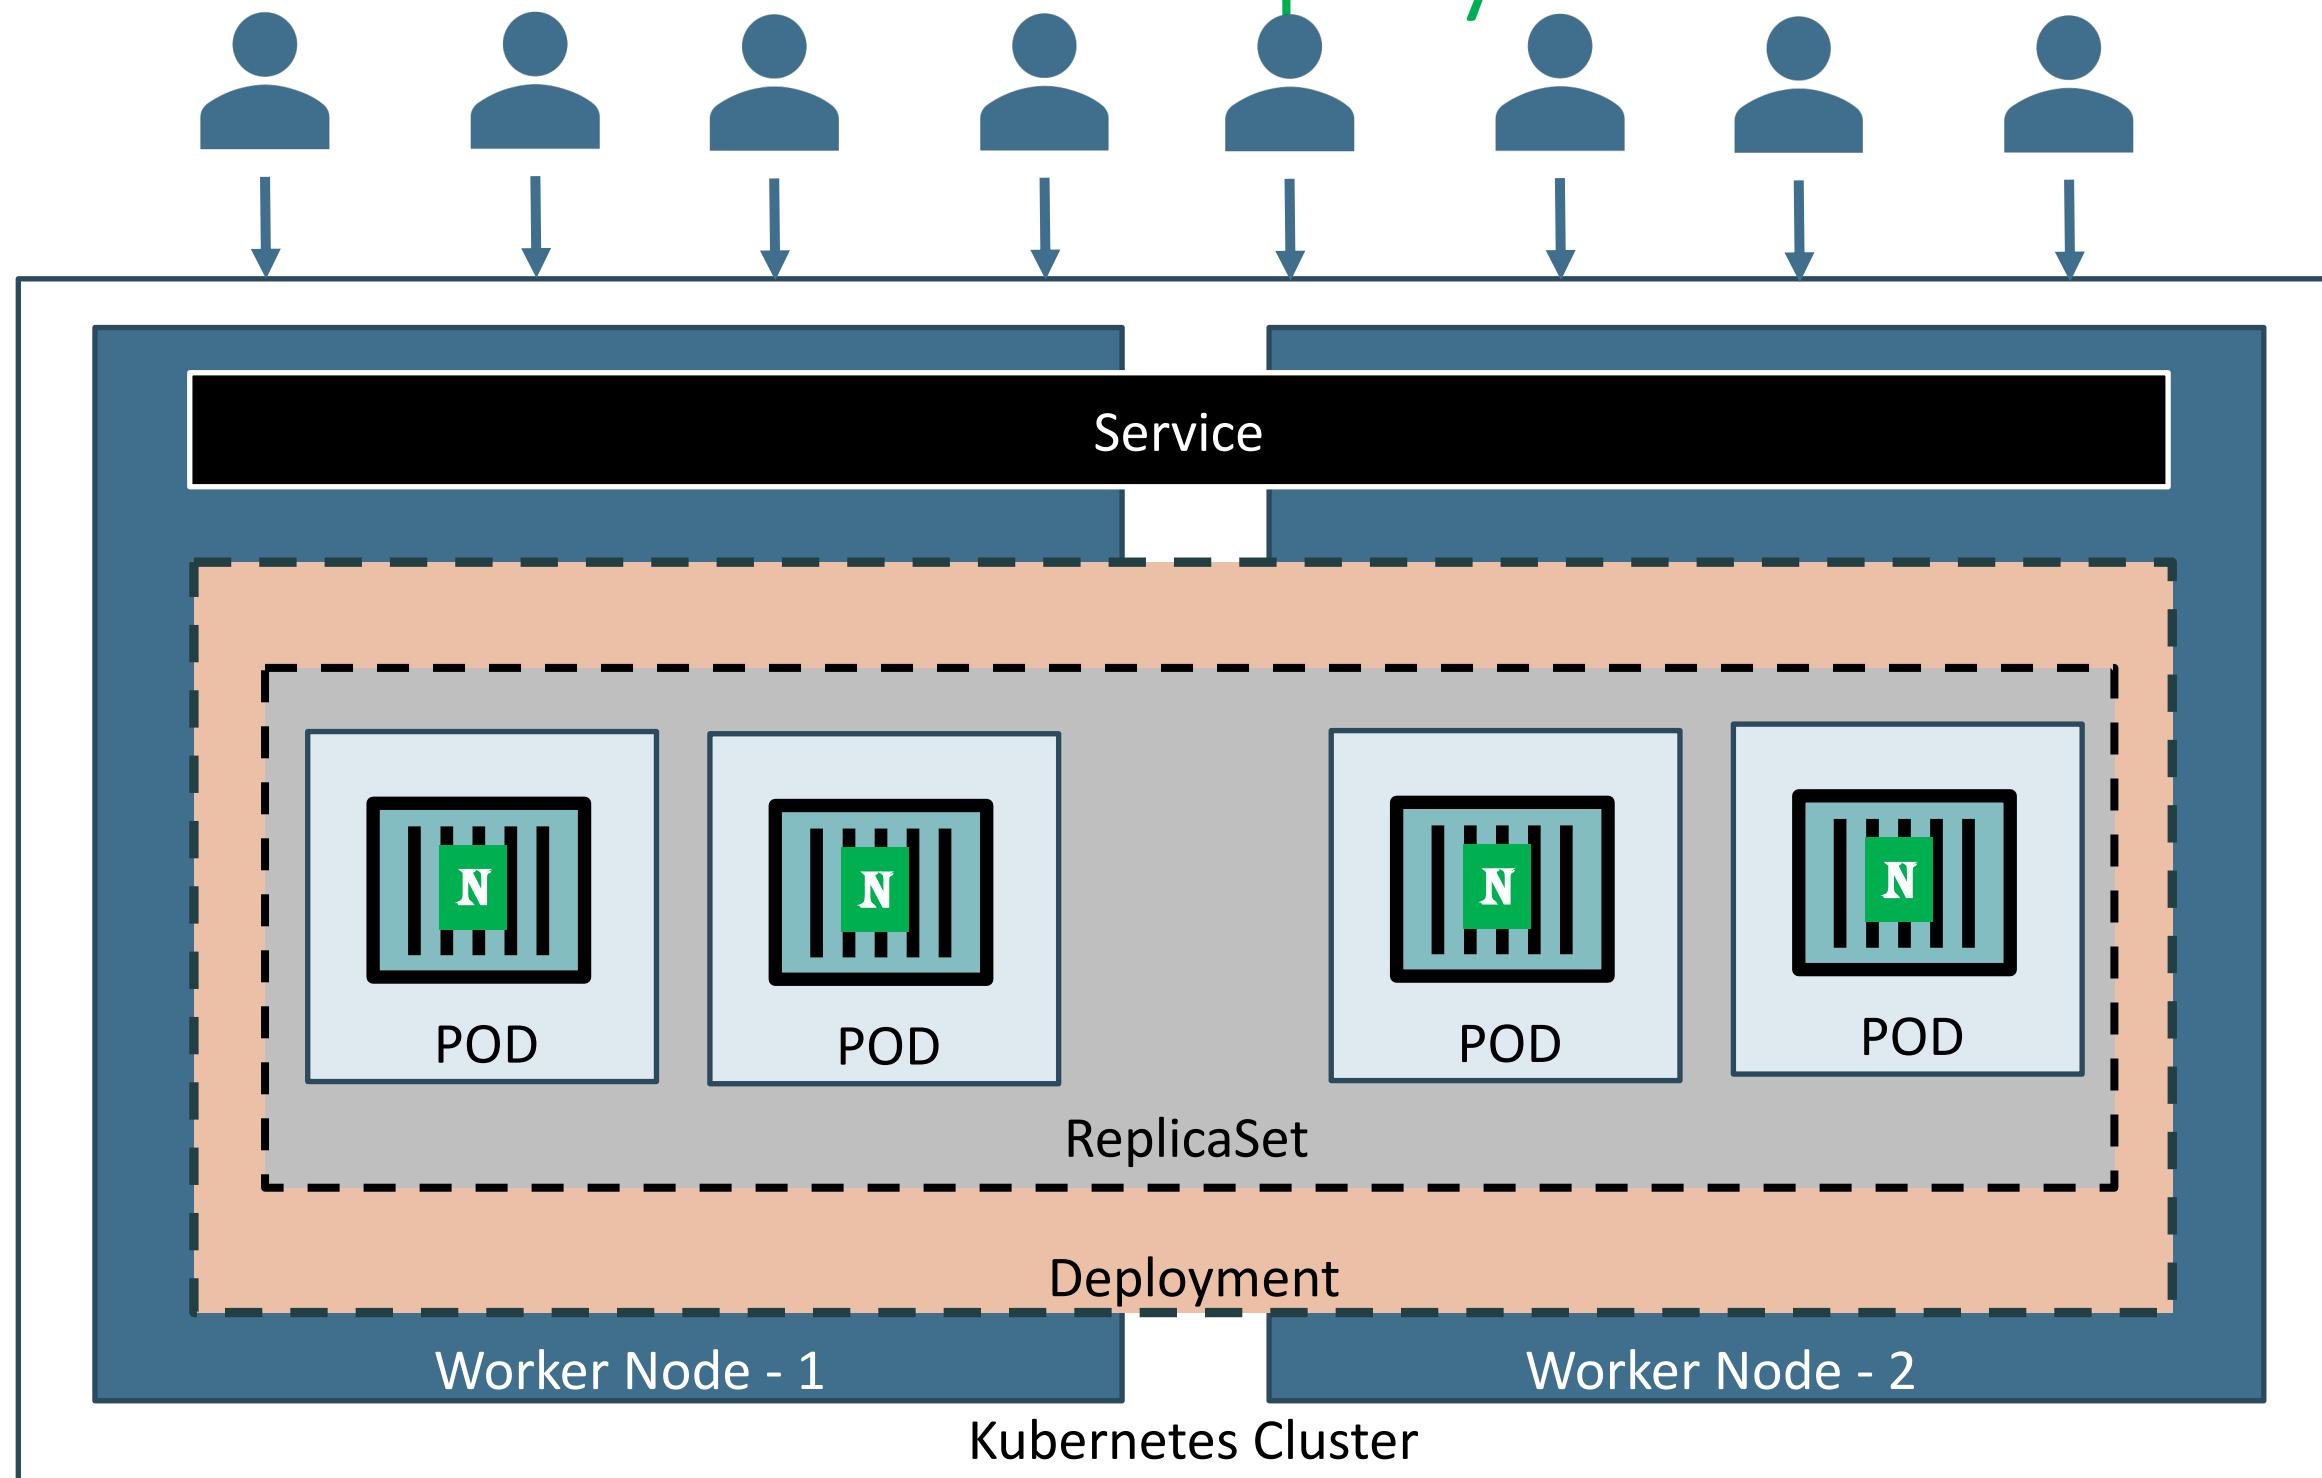
# Kubernetes ReplicaSets Demo



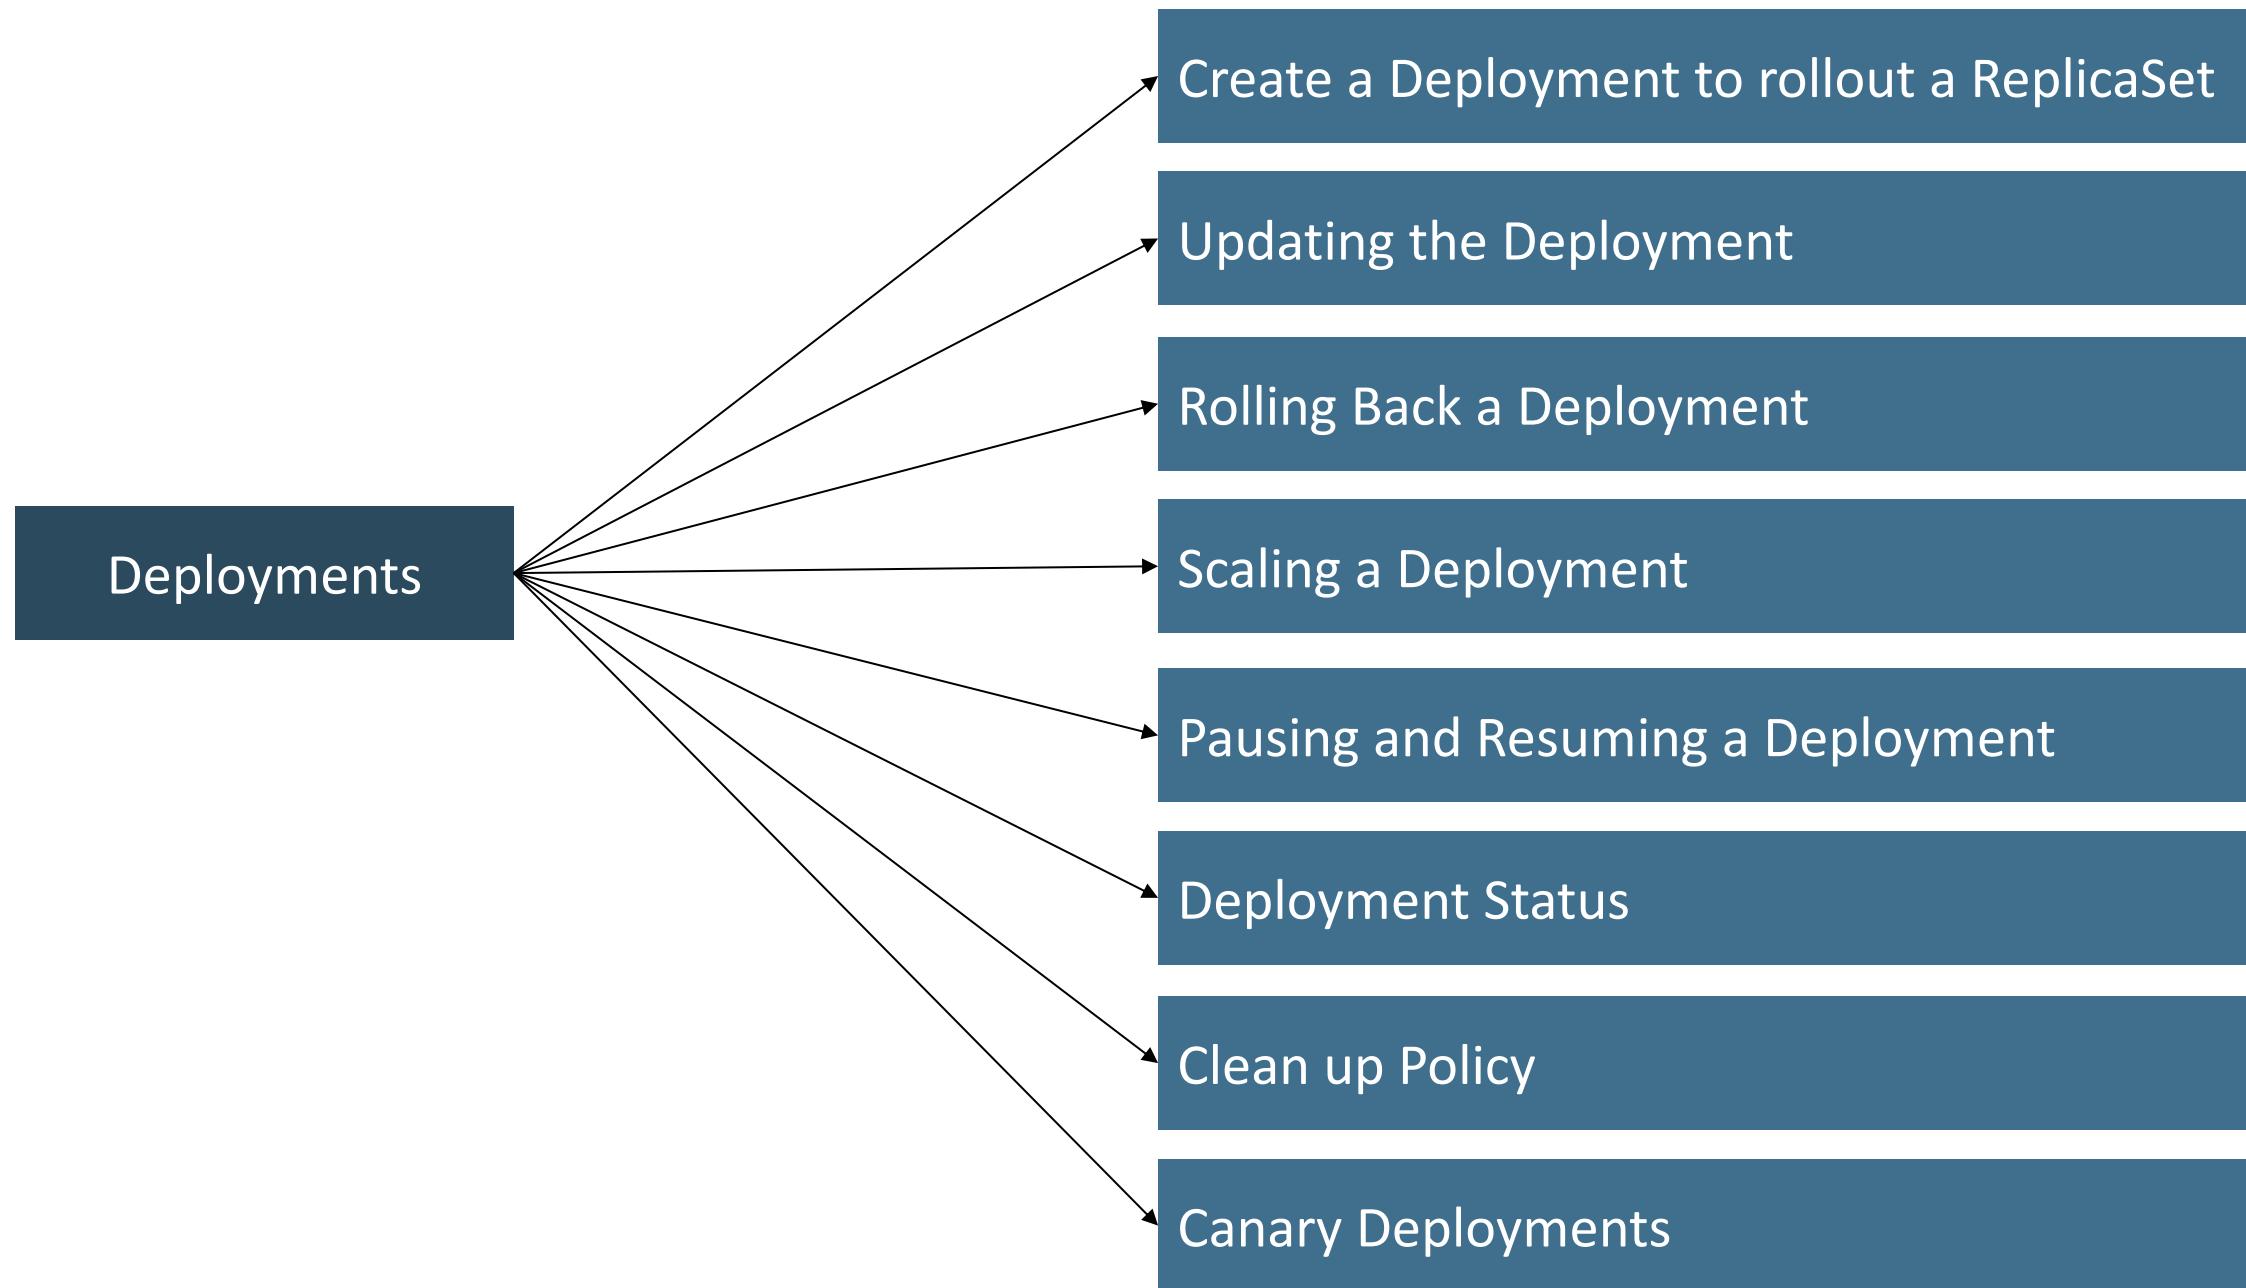
# Kubernetes Deployments



# Kubernetes – Deployments



# Kubernetes - Deployment



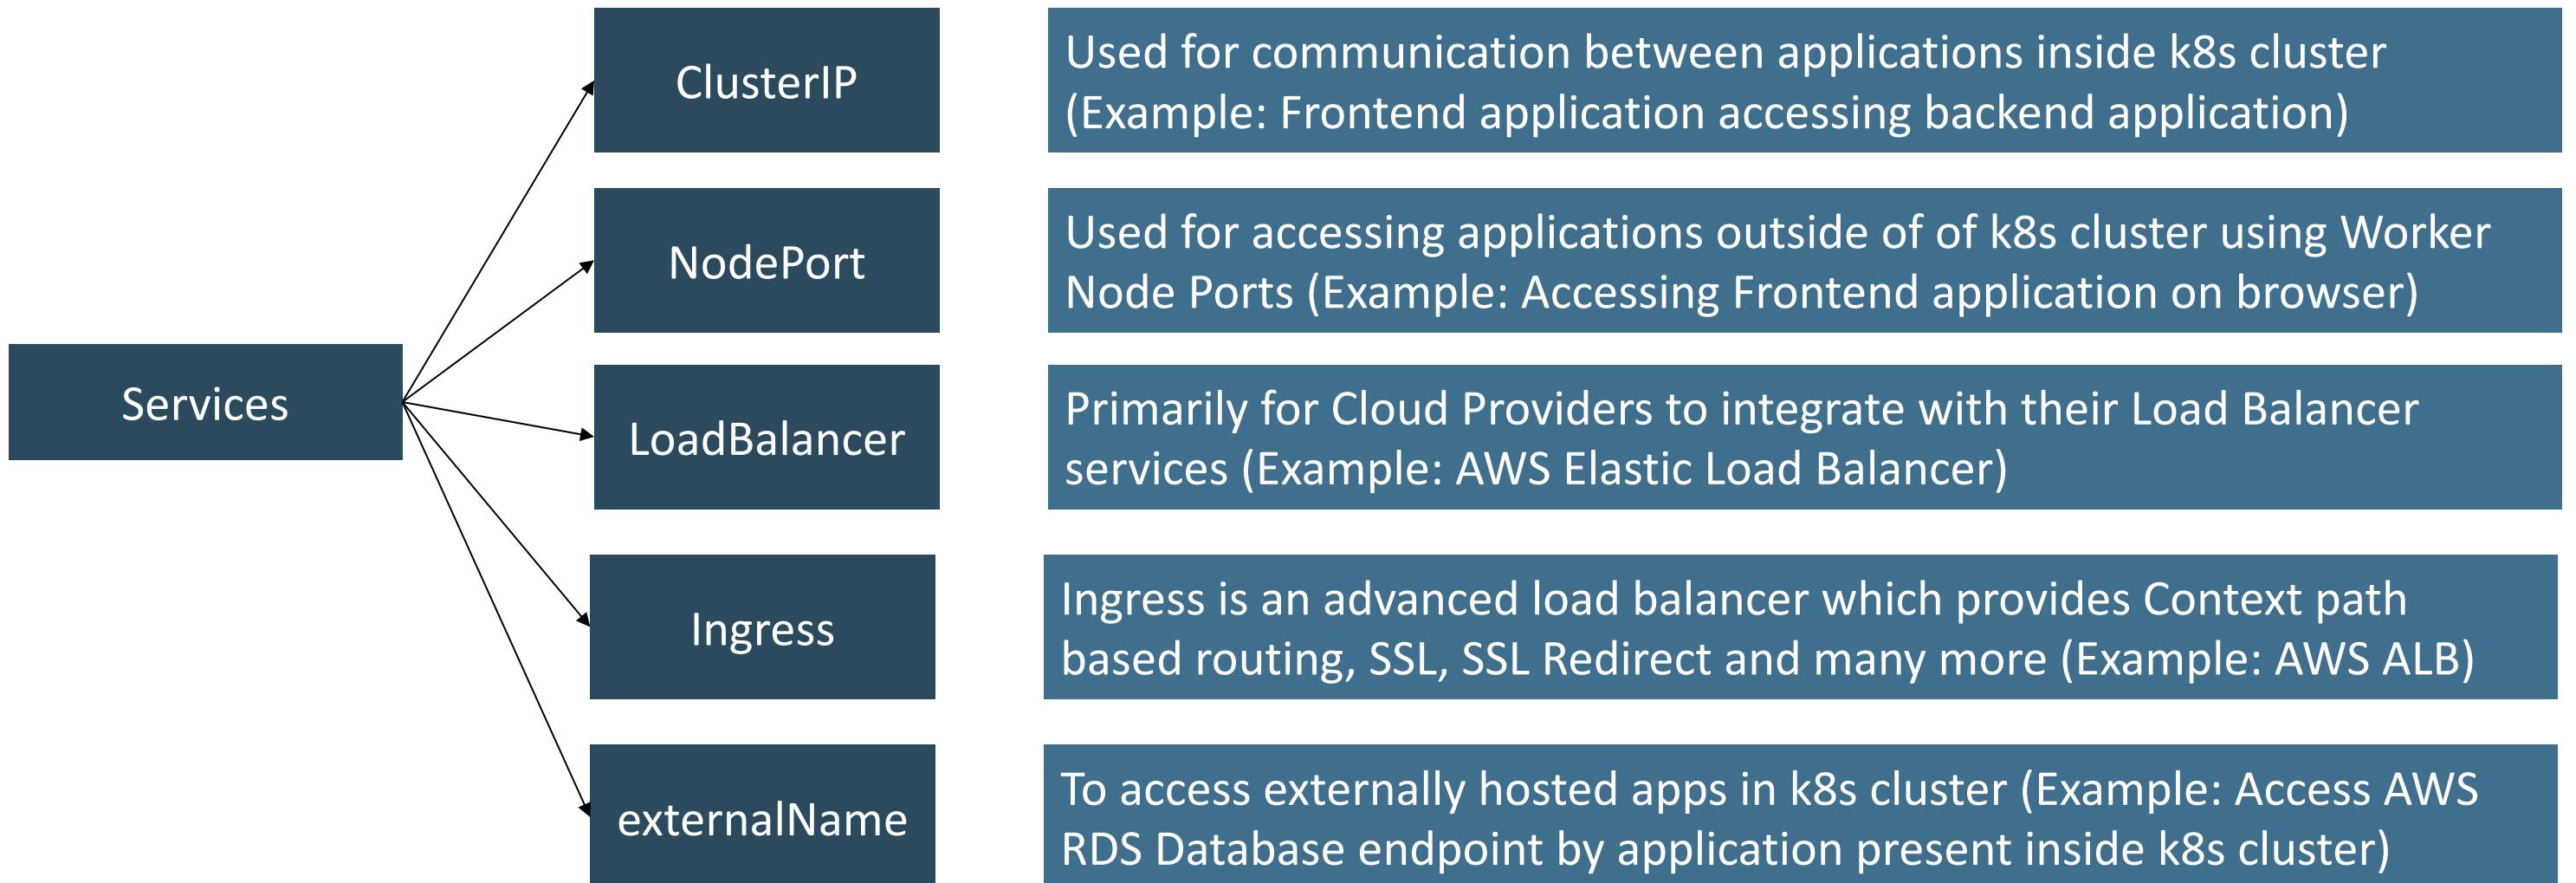
# Kubernetes Deployments Demo

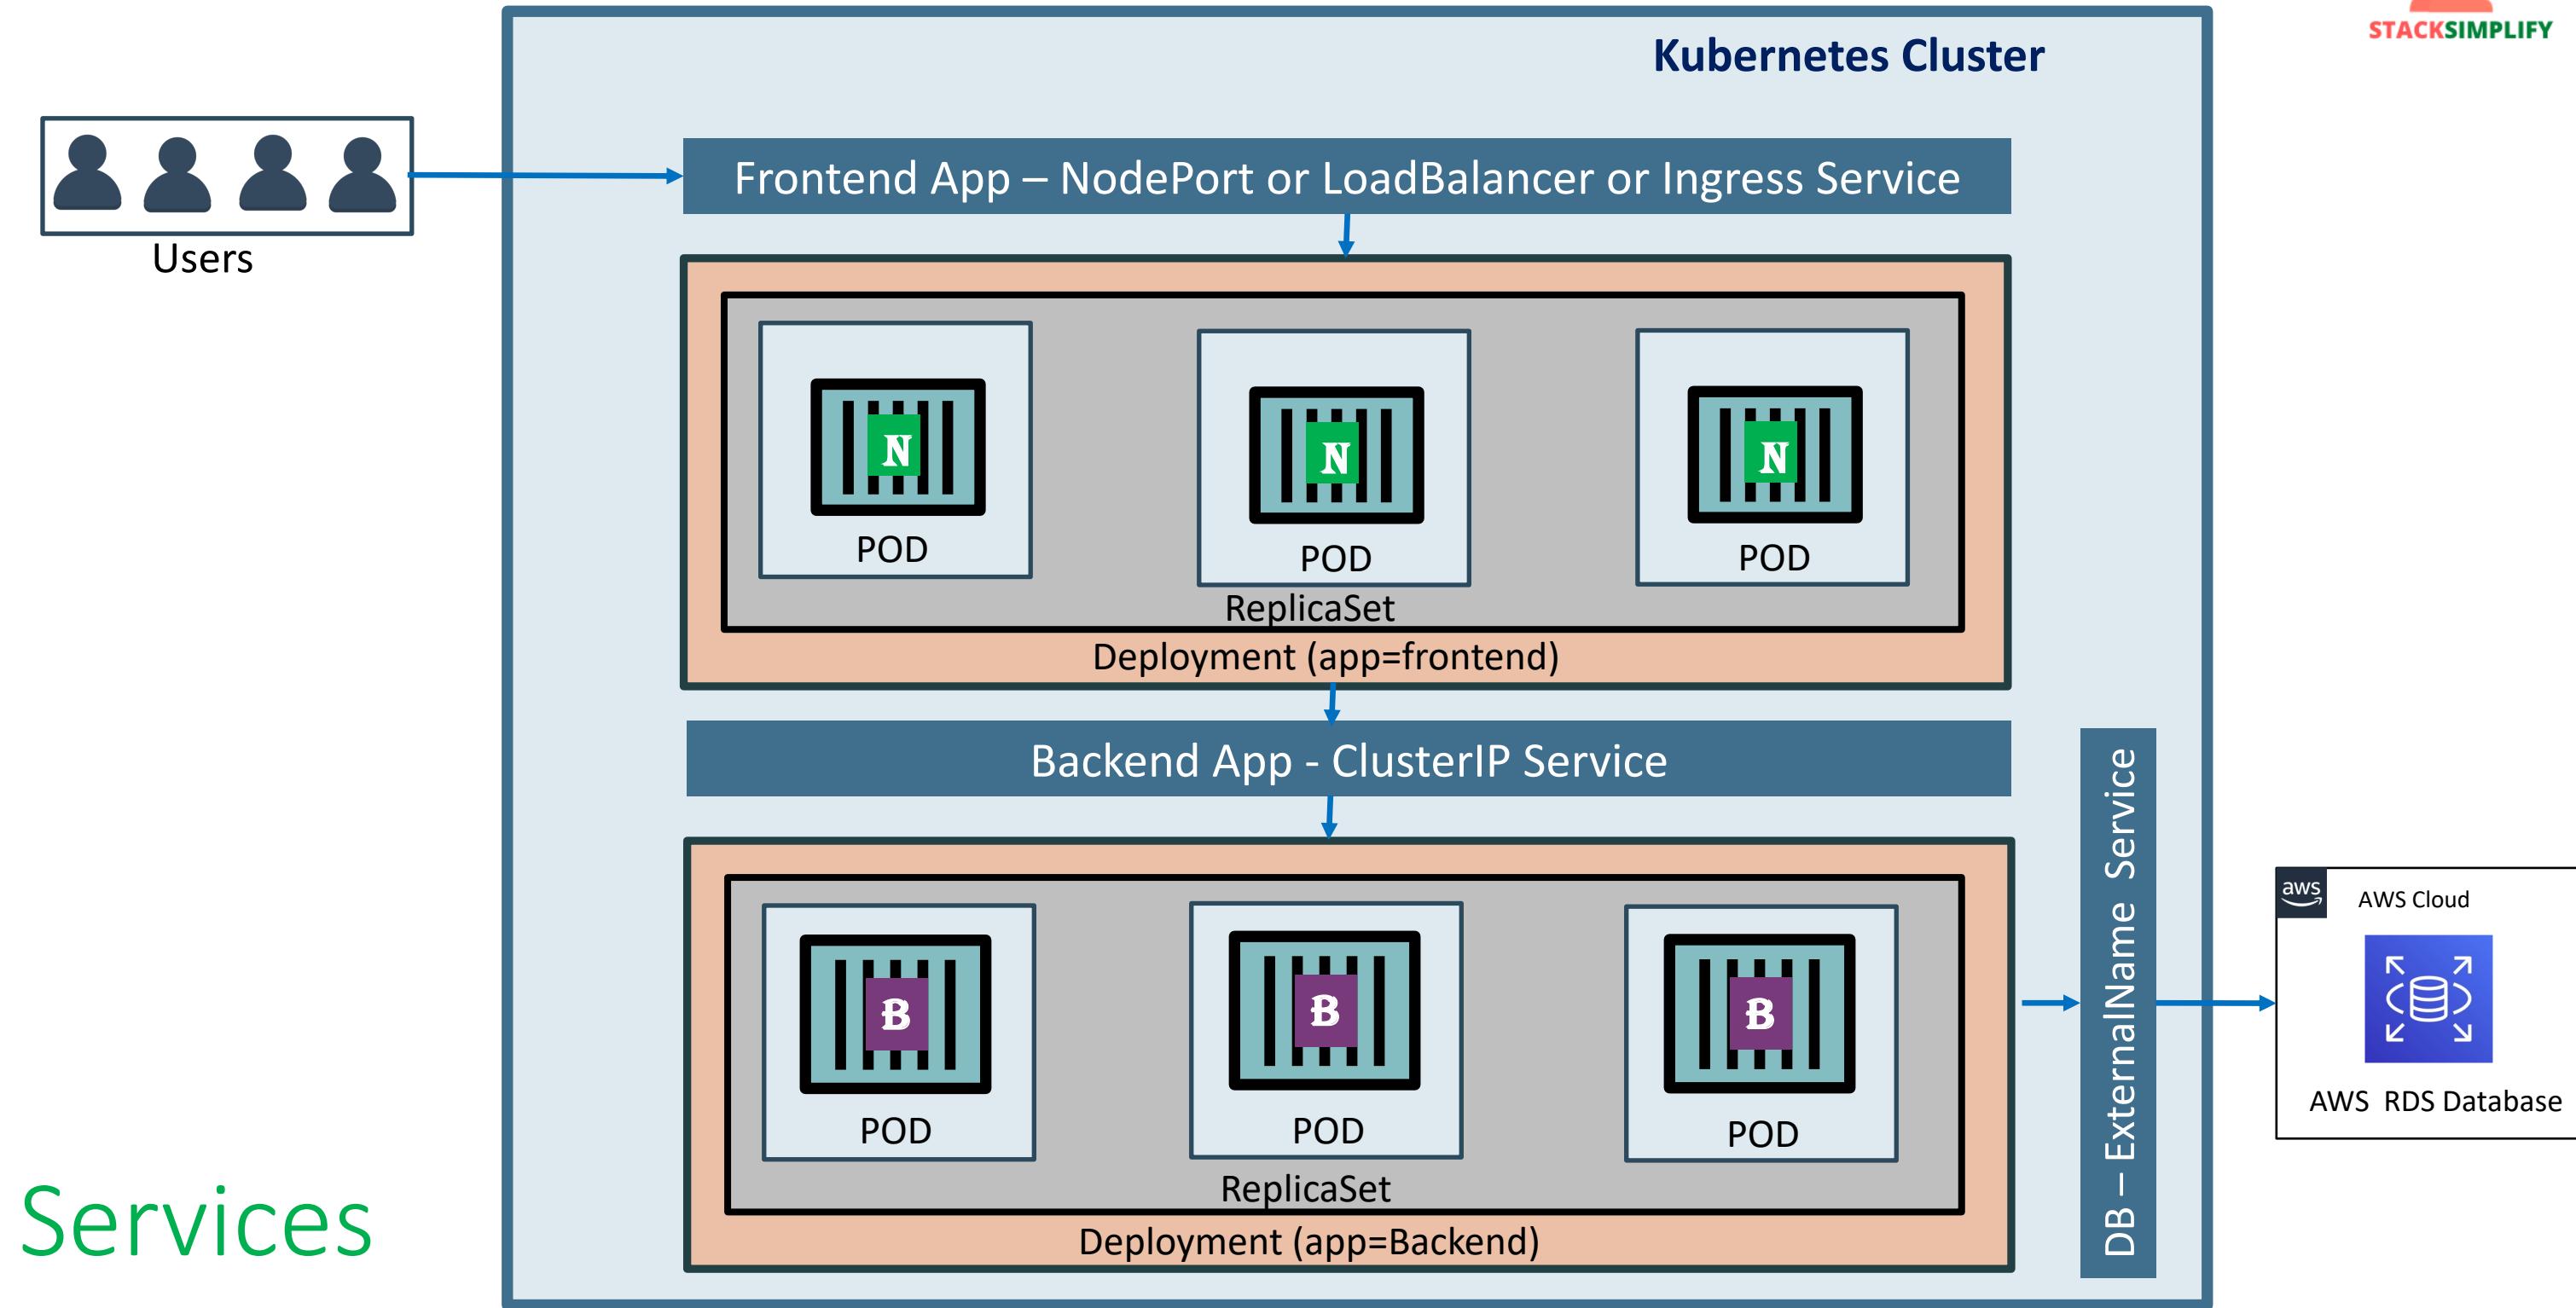


# Kubernetes Services



# Kubernetes - Services



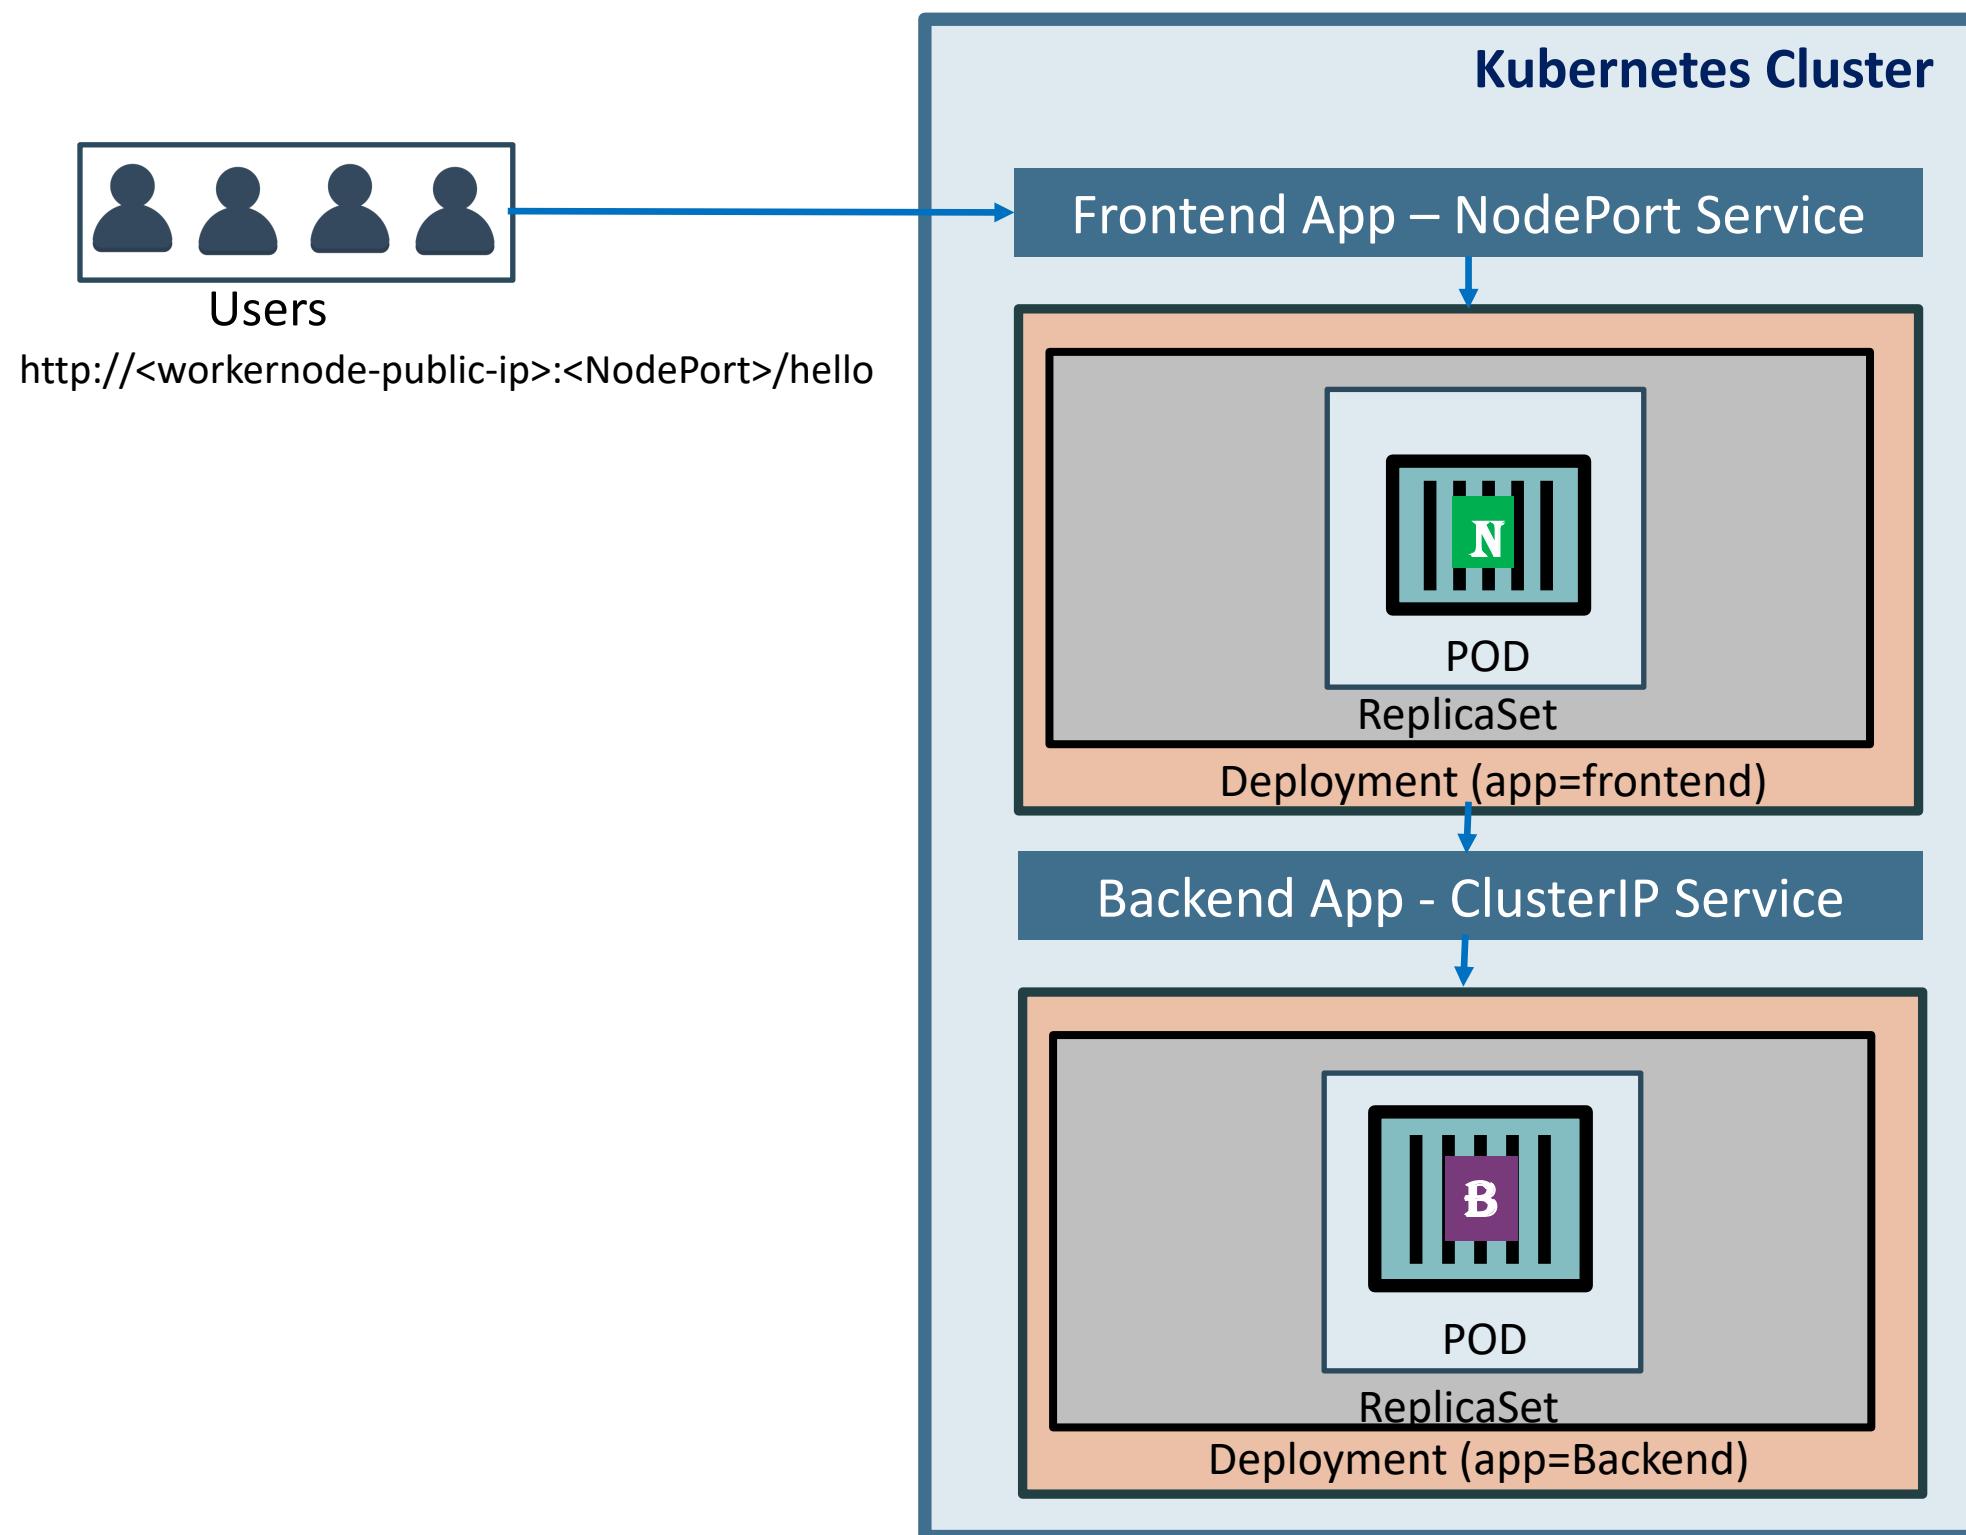


# Services

# Kubernetes Services Demo



# Services Demo



# Kubernetes YAML Basics



# YAML Basics

- YAML is **not** a Markup Language
- YAML is used to **store information** about different things
- We can use YAML to **define key, Value pairs** like variables, lists and objects
- YAML is very similar to **JSON** (Javascript Object Notation)
- YAML primarily focuses on **readability** and **user friendliness**
- YAML is designed to be **clean and easy to read**
- We can define YAML files with two different extensions
  - abc.**yml**
  - abc.**yaml**

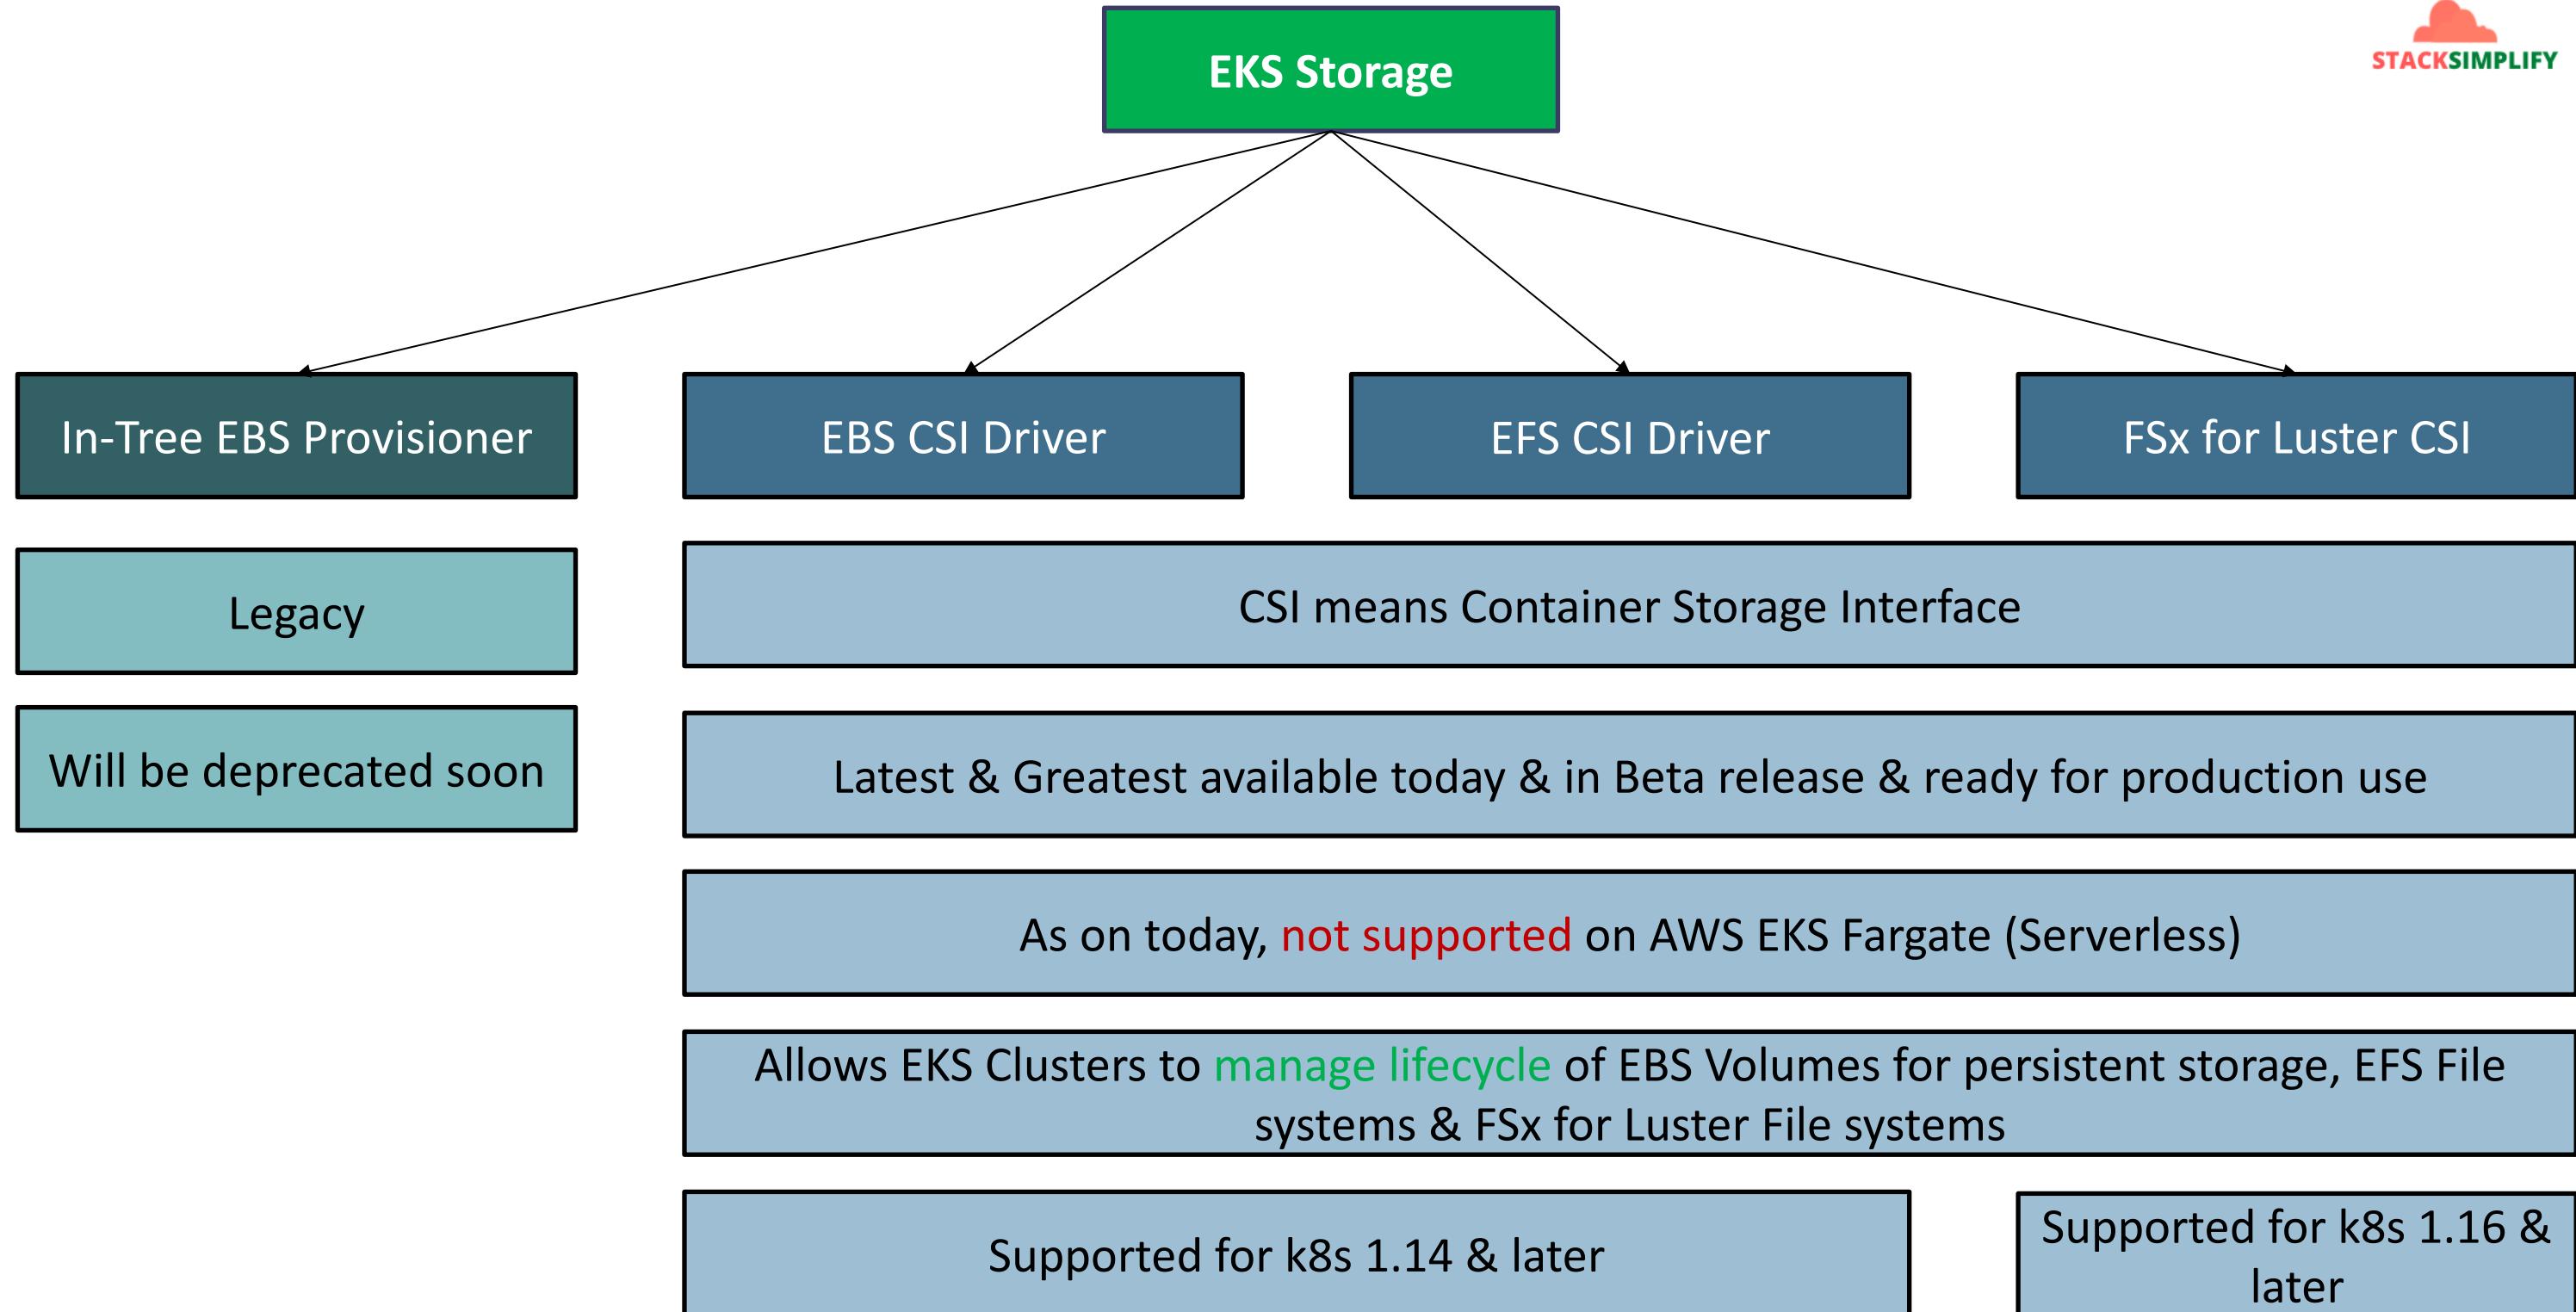
# YAML Basics

- YAML Comments
- YAML Key Value Pairs
- YAML Dictionary or Map
- YAML Array / Lists
- YAML Spaces
- YAML Document Separator



# AWS EKS Storage







# AWS EKS Storage EBS CSI Driver

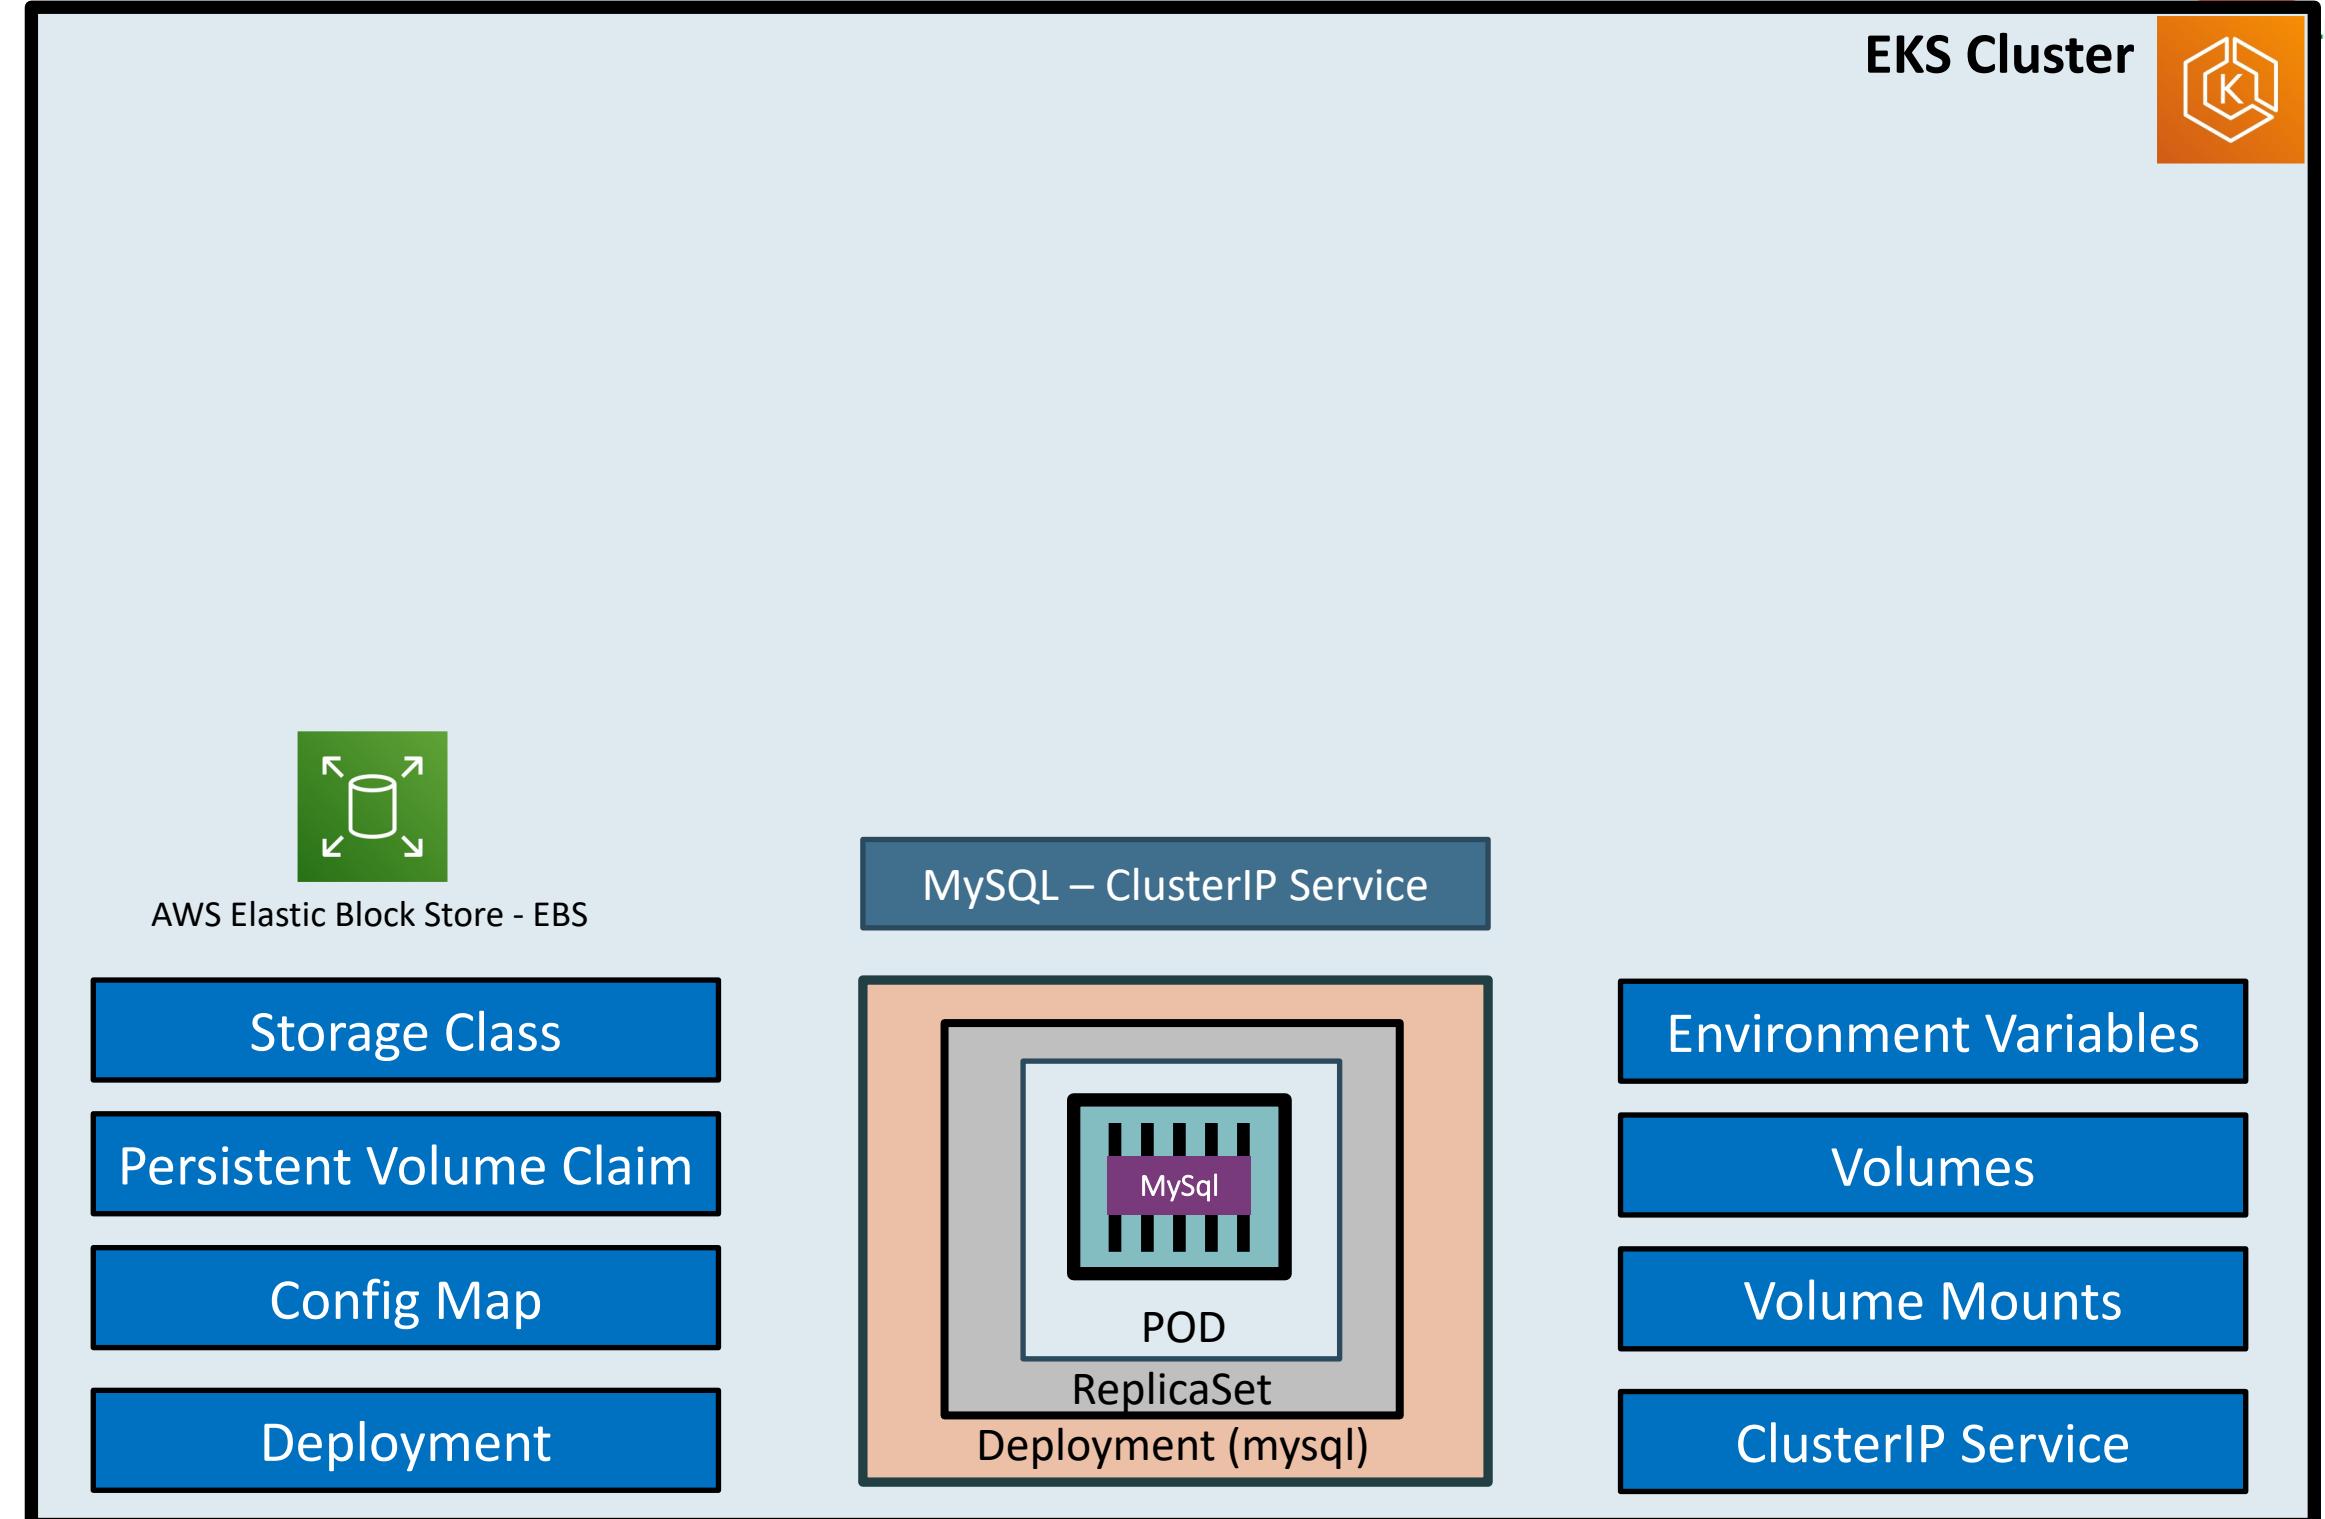


# AWS Elastic Block Store - Introduction

- EBS provides **block level storage volumes** for use with **EC2 & Container instances**.
- We can mount these **volumes as devices** on our EC2 & Container instances.
- EBS volumes that are attached to an instance are **exposed as storage volumes that persist independently** from the life of the EC2 or Container instance.
- We can **dynamically change** the configuration of a volume attached to an instance.
- AWS recommends EBS for data that must be **quickly accessible** and requires **long-term persistence**.
- EBS is well suited to both **database-style applications** that rely on random reads and writes, and to **throughput-intensive applications** that perform long, continuous reads and writes.



# EKS Storage EBS CSI Driver



## EKS Cluster



# EKS Storage EBS CSI Driver



Users

`http://<workernode-public-ip>:<NodePort>/usermgmt/<apis>`

UserMgmt – NodePort Service



POD  
ReplicaSet

Deployment (UserMgmt)

NodePort Service

Deployment

Environment Variables



AWS Elastic Block Store - EBS

Storage Class

Persistent Volume Claim

Config Map

Deployment

MySQL – ClusterIP Service



POD  
ReplicaSet

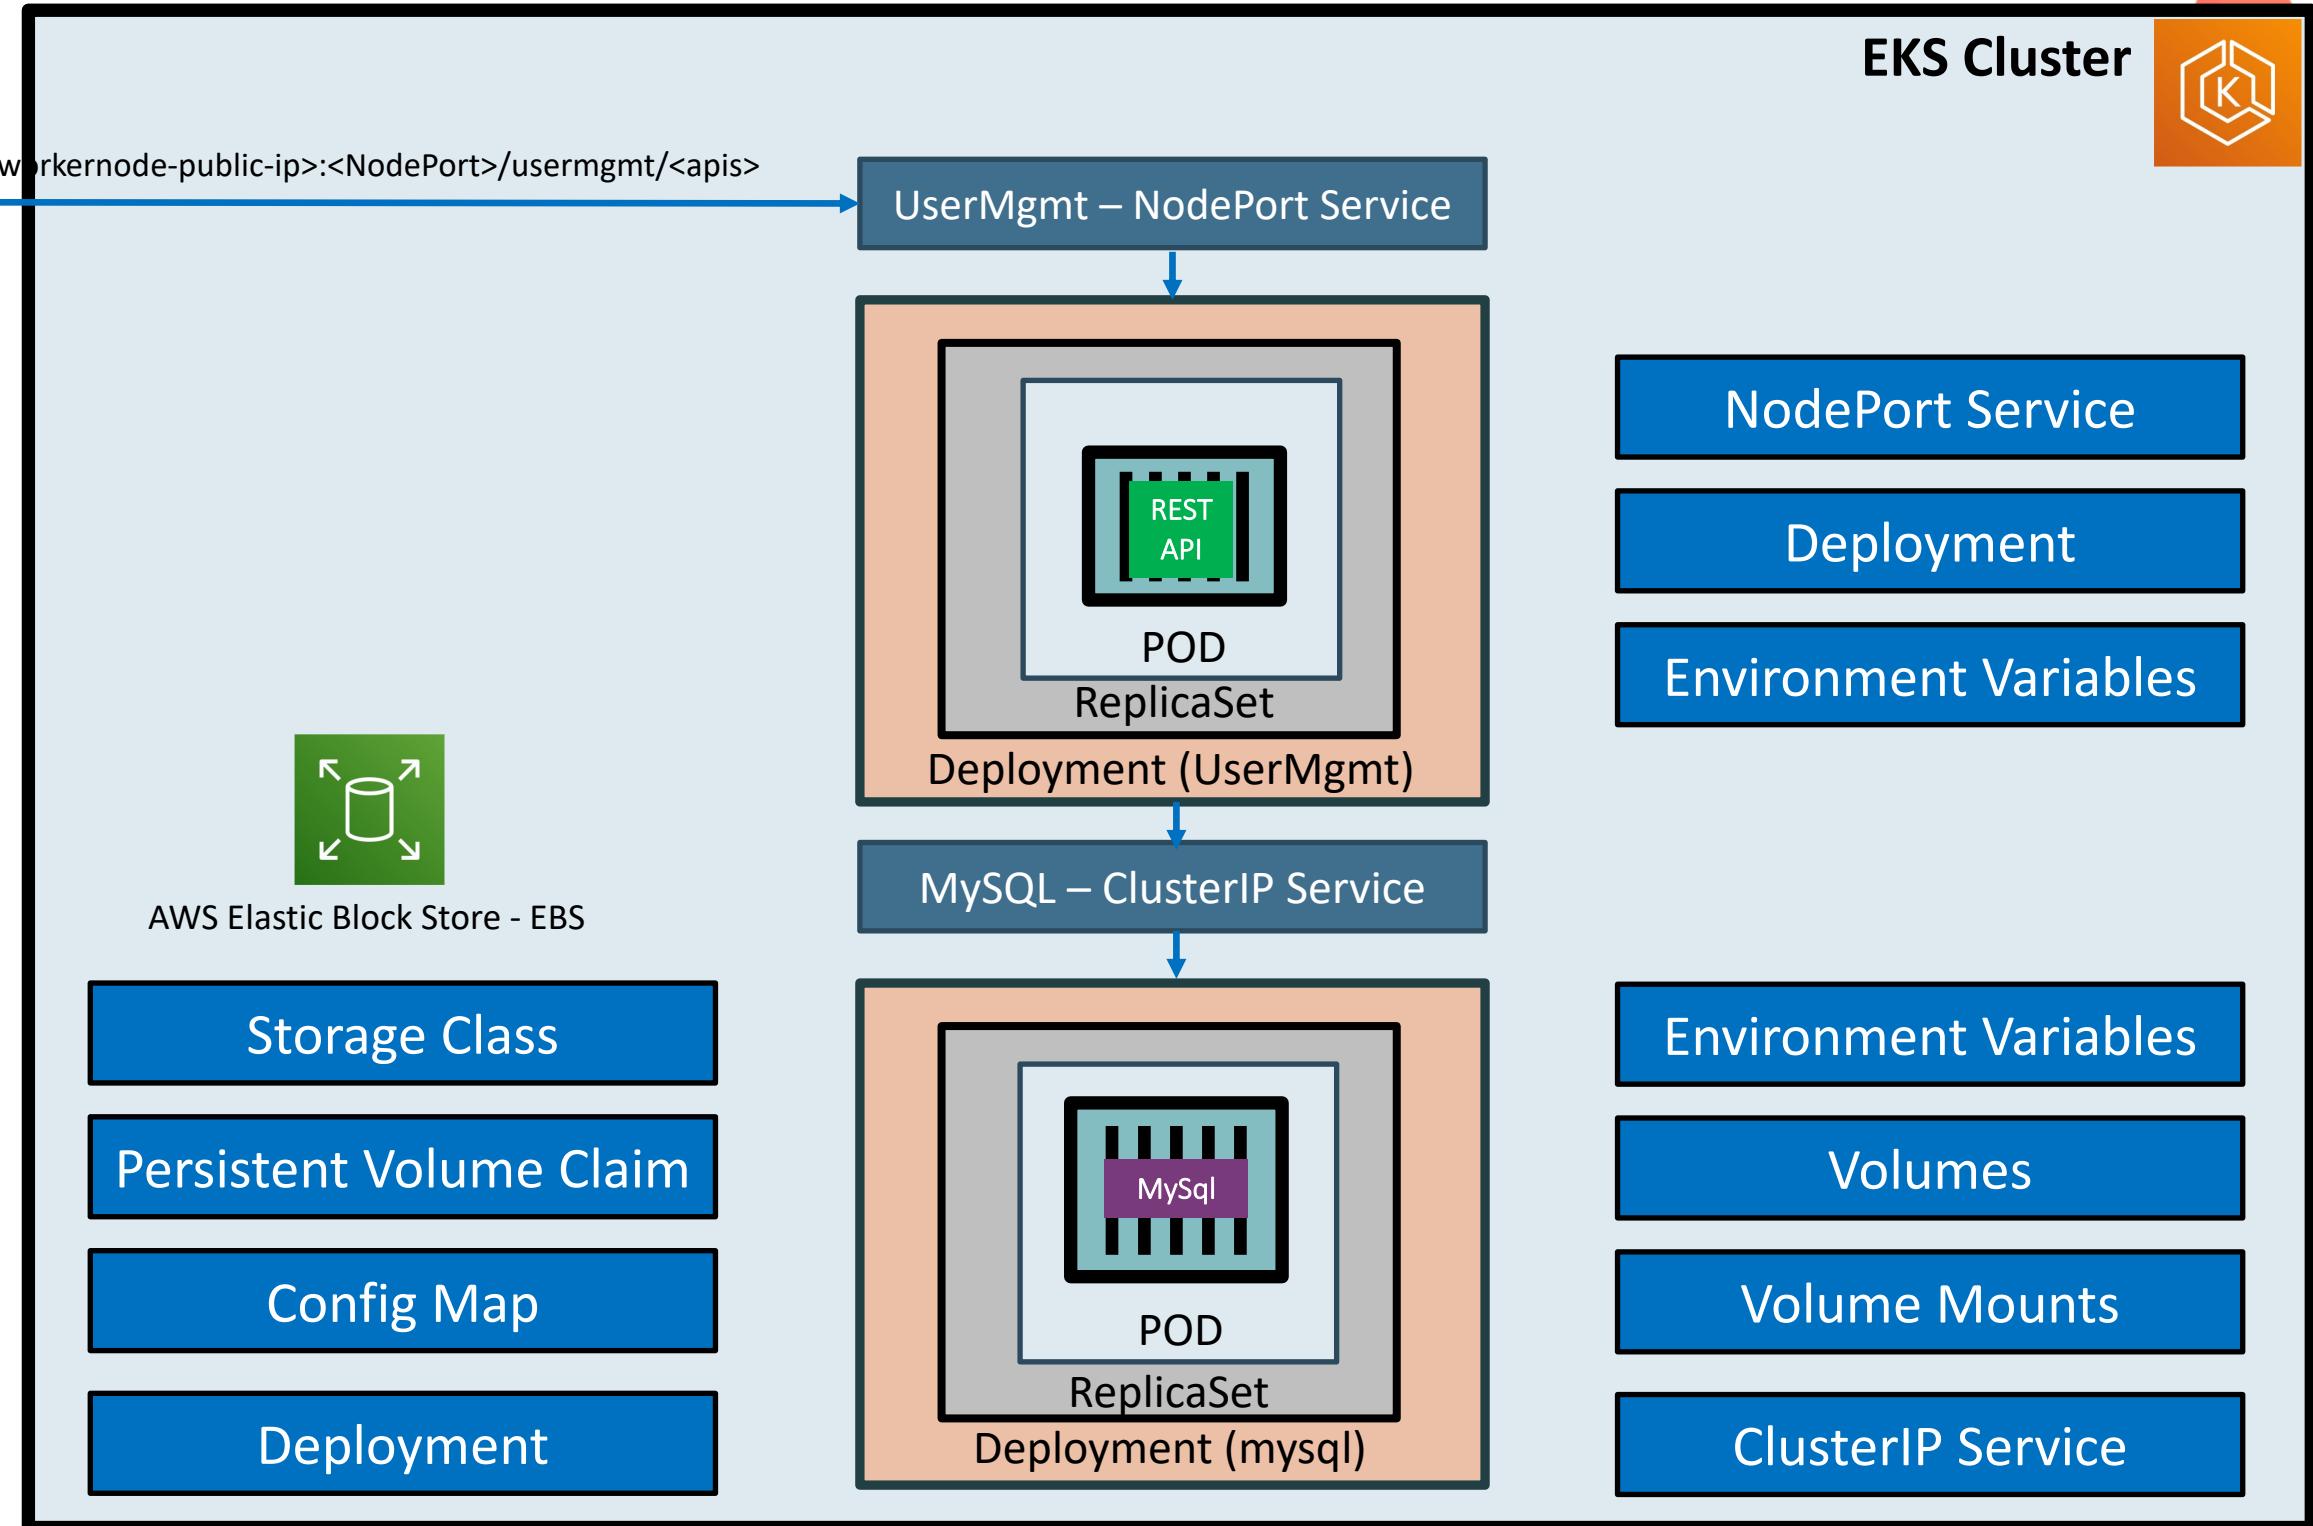
Deployment (mysql)

Environment Variables

Volumes

Volume Mounts

ClusterIP Service





Elastic Block Store



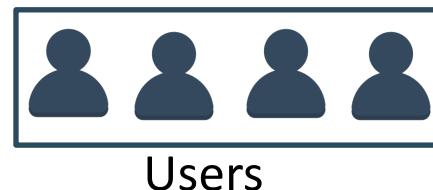
# AWS EKS Storage EBS CSI Driver Important k8s Concepts for Application Deployments



## EKS Cluster

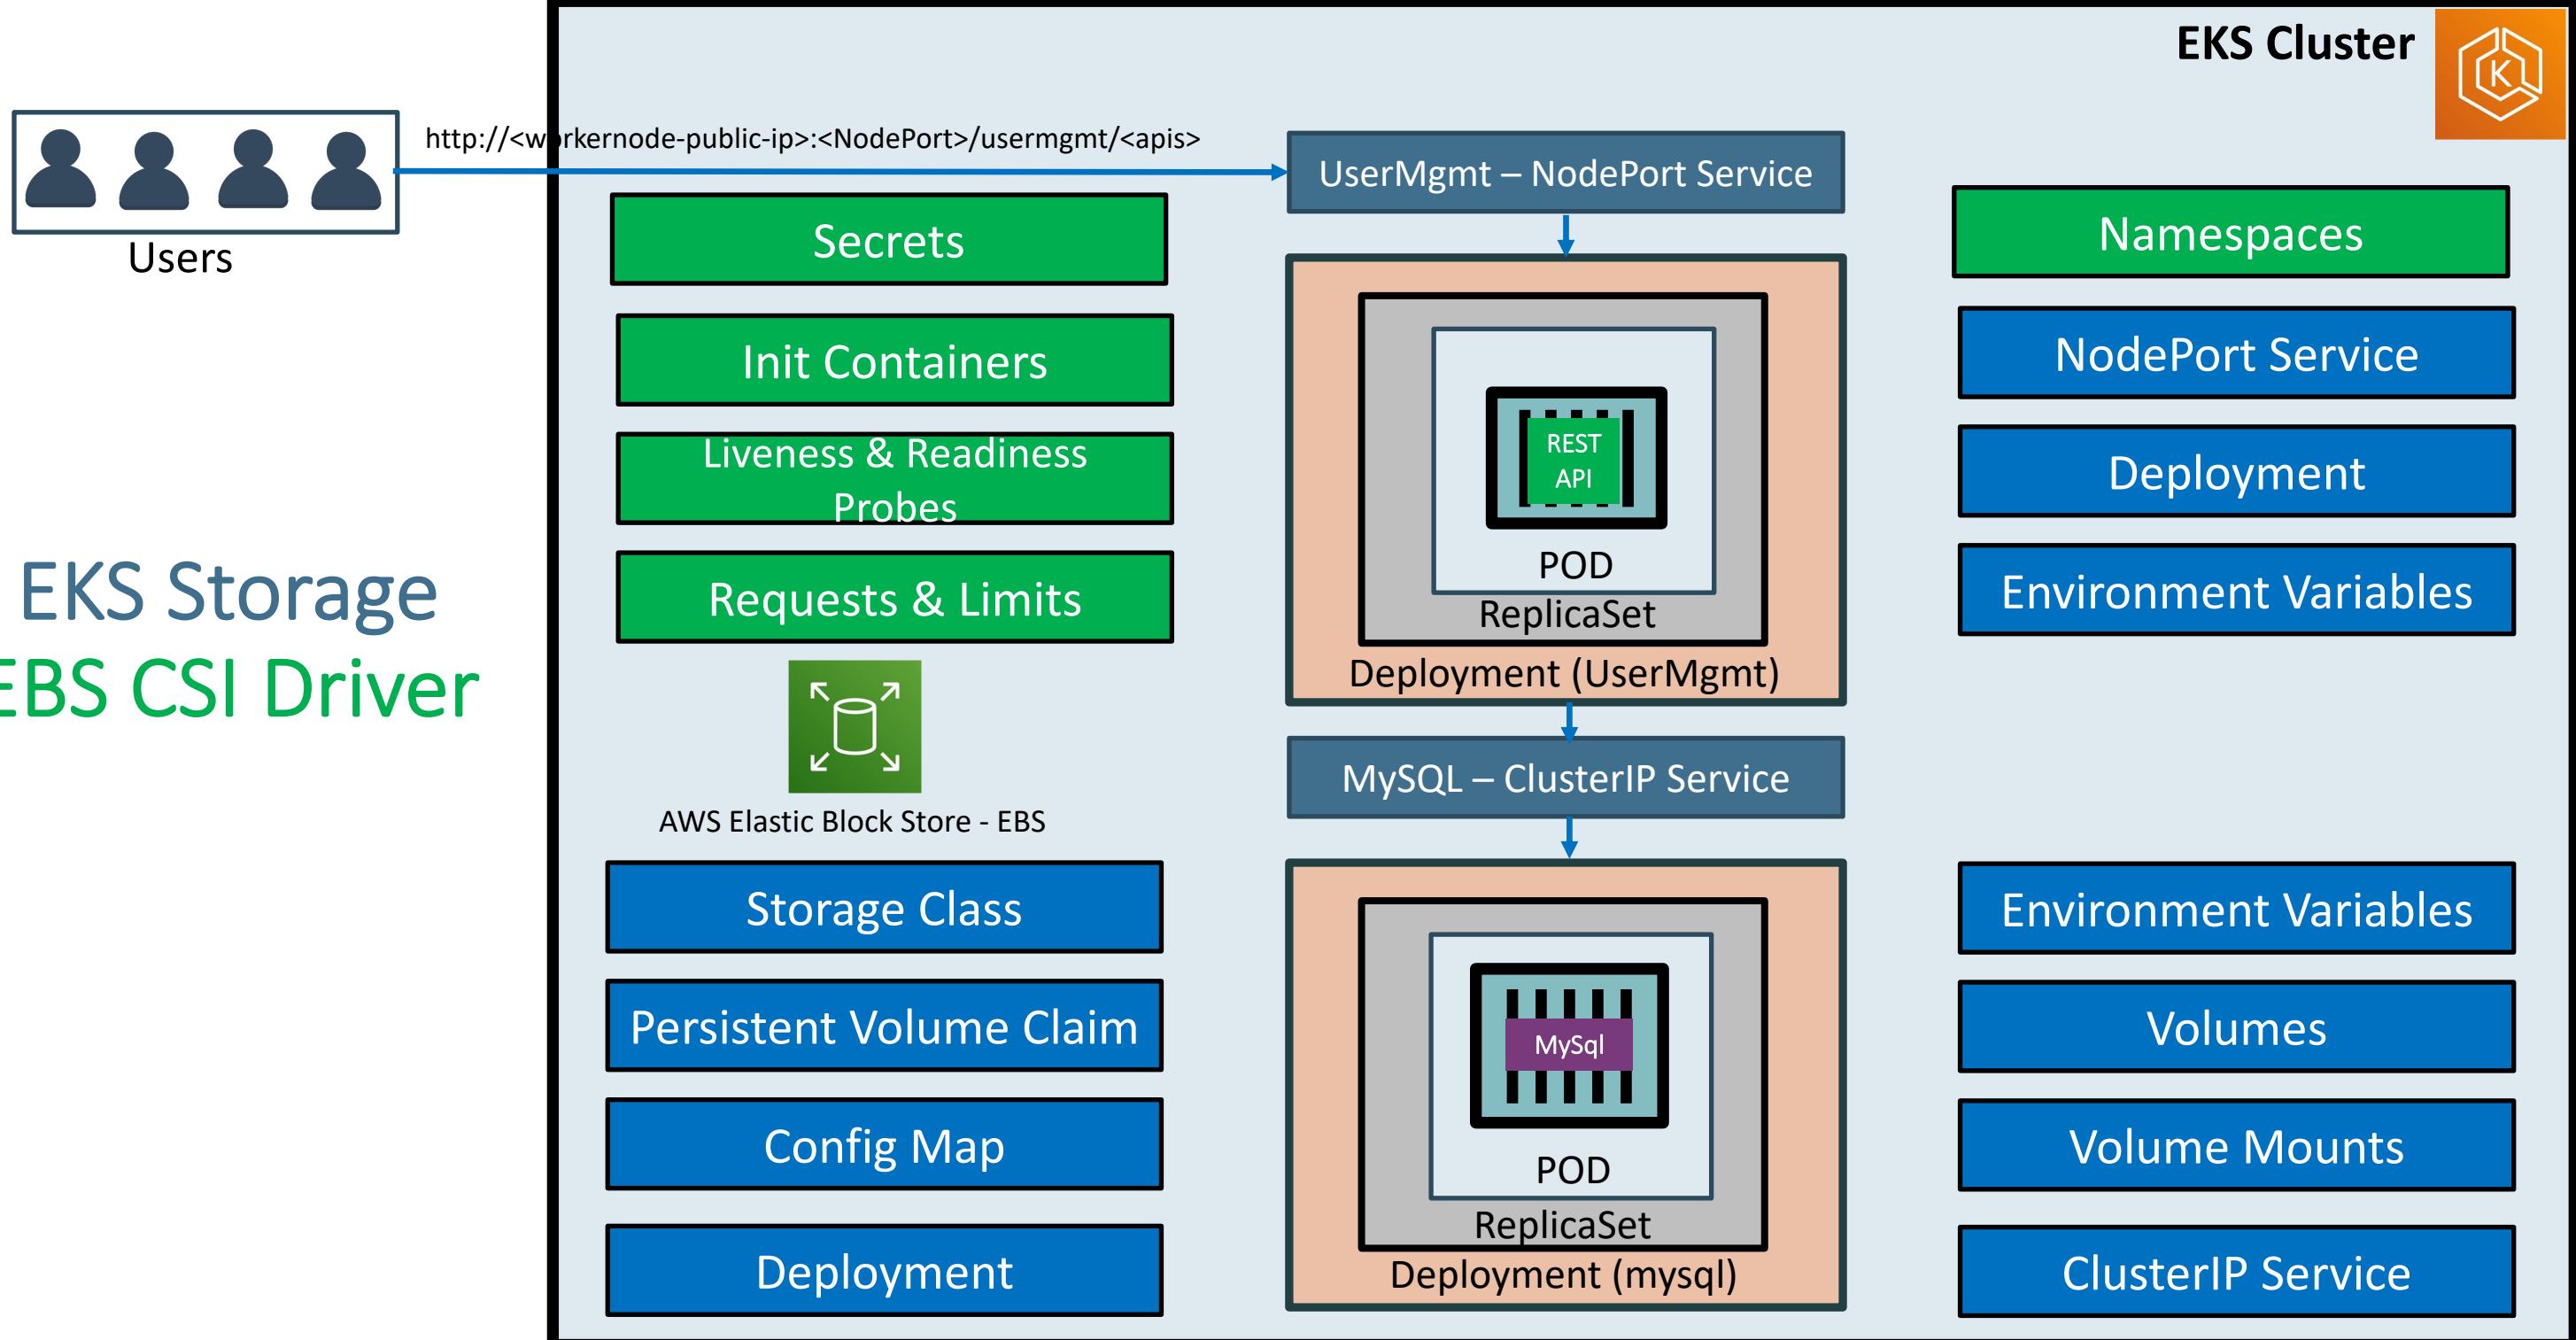


# EKS Storage EBS CSI Driver



Users

<http://<workernode-public-ip>:<NodePort>/usermgmt/<apis>>



# Probes

## Probes

### Liveness Probe

Kubelet uses liveness probes to know when to restart a container

Liveness probes could catch a **deadlock**, where an application is running, but unable to make progress and **restarting container** helps in such state

### Readiness Probe

Kubelet uses readiness probes to know when a container is ready to accept traffic

When a Pod is not ready, it is **removed** from Service load balancers based on this **readiness probe signal**.

### Startup Probe

Kubelet uses startup probes to know when a container application has started

Firstly this probe **disables** liveness & readiness checks until it **succeeds** ensuring those pods don't interfere with app startup.

## Options to define Probes

Check using Commands

```
/bin/sh –c nc -z localhost 8095
```

Check using HTTP GET Request

```
httpget path:/health-status
```

Check using TCP

```
tcpSocket Port: 8095
```

# Kubernetes Namespaces



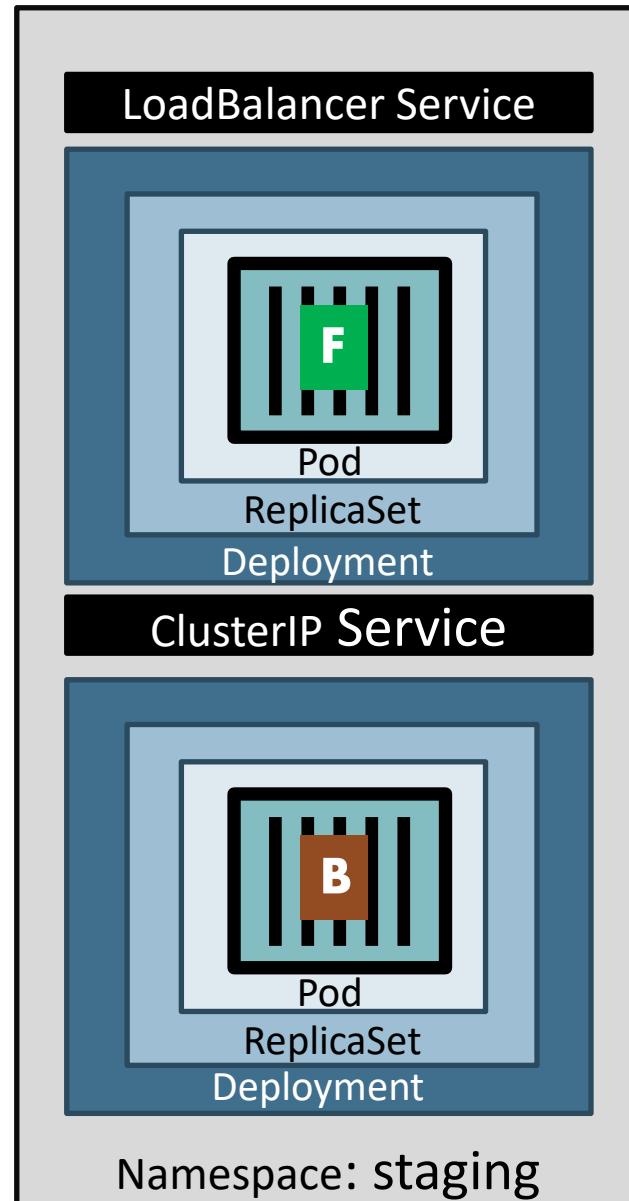
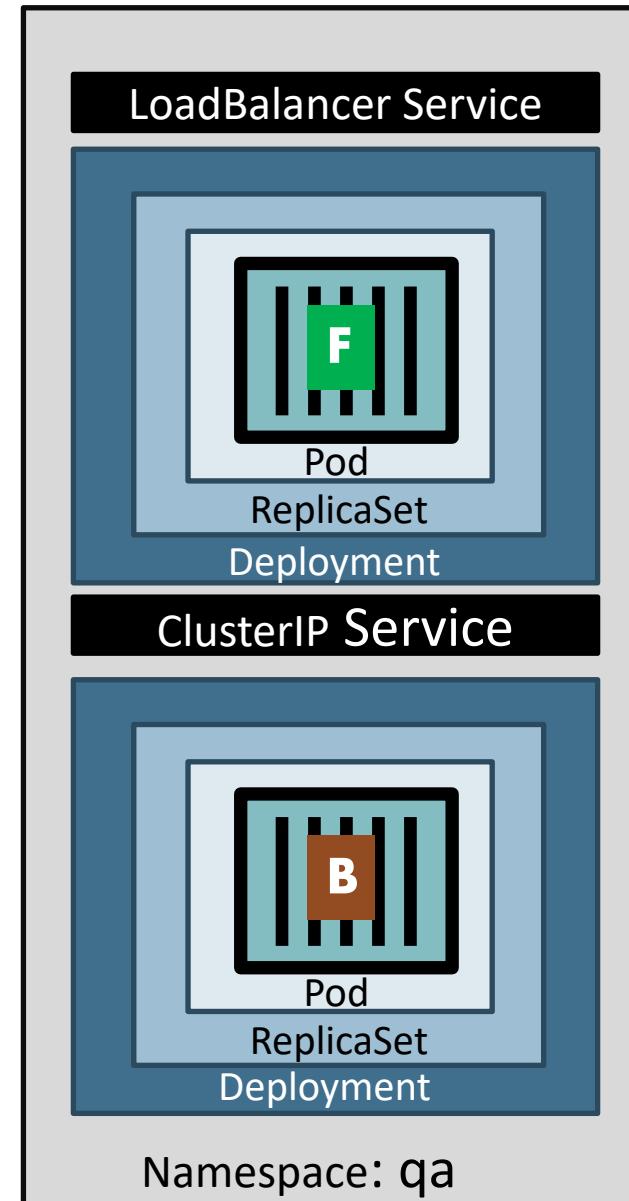
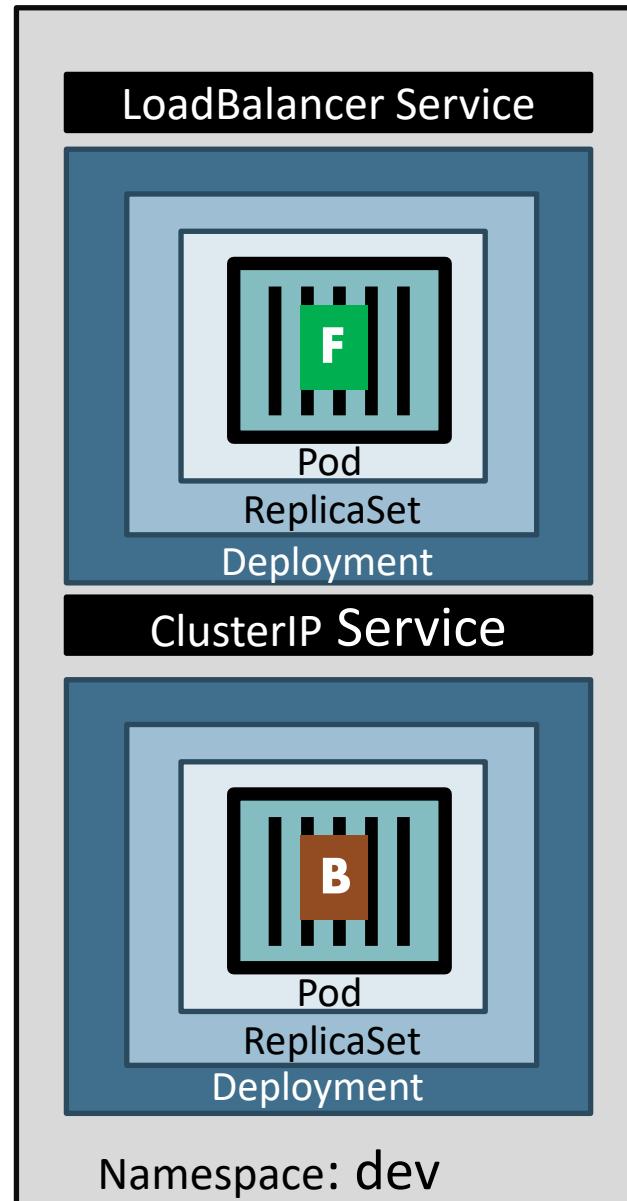
# Namespaces - Introduction

- Namespaces are also called **Virtual clusters** in our **physical** k8s cluster
- We use this in environments where we have **many users** spread across multiple teams or projects
- Clusters with **tens of users** ideally don't need to use namespaces
- Benefits**
  - Creates **isolation boundary** from other k8s objects
  - We can limit the resources like **CPU, Memory** on per namespace basis (**Resource Quota**).

```
kubectl get namespace
```

NAME	STATUS	AGE
default	Active	141m
kube-node-lease	Active	141m
kube-public	Active	141m
kube-system	Active	141m

# Namespaces



**Namespace Manifest - Declarative**

```
1 apiVersion: v1
2 kind: Namespace
3 metadata:
4   name: dev3
```

**Namespace Manifest - Imperative**

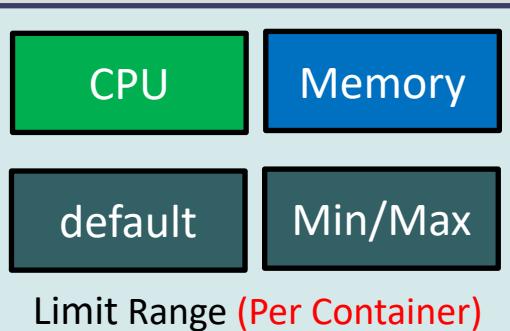
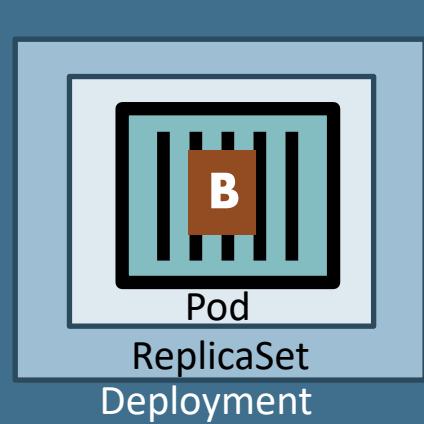
```
kubectl create namespace dev1
```

# Limit Range

LoadBalancer Service



ClusterIP Service

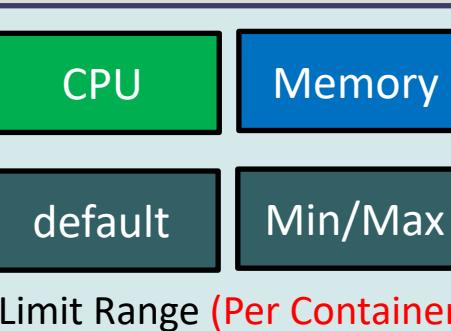
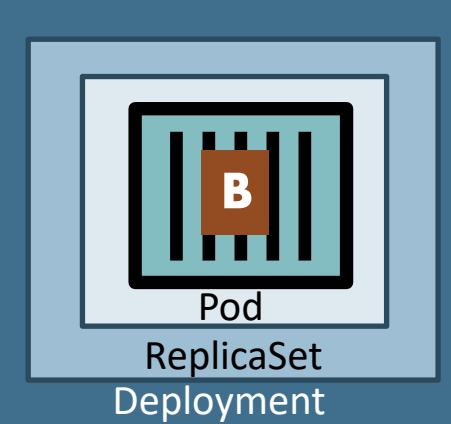


Namespace: dev

LoadBalancer Service



ClusterIP Service

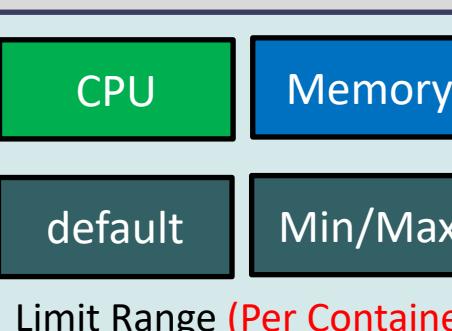
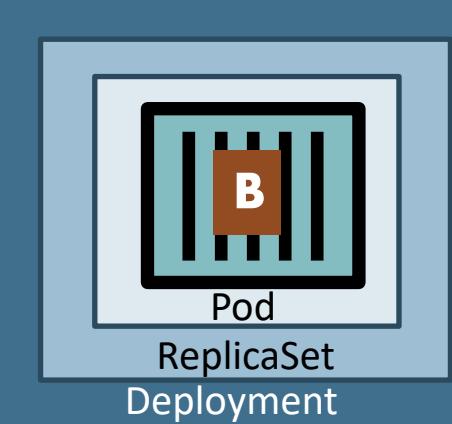


Namespace: qa

LoadBalancer Service



ClusterIP Service

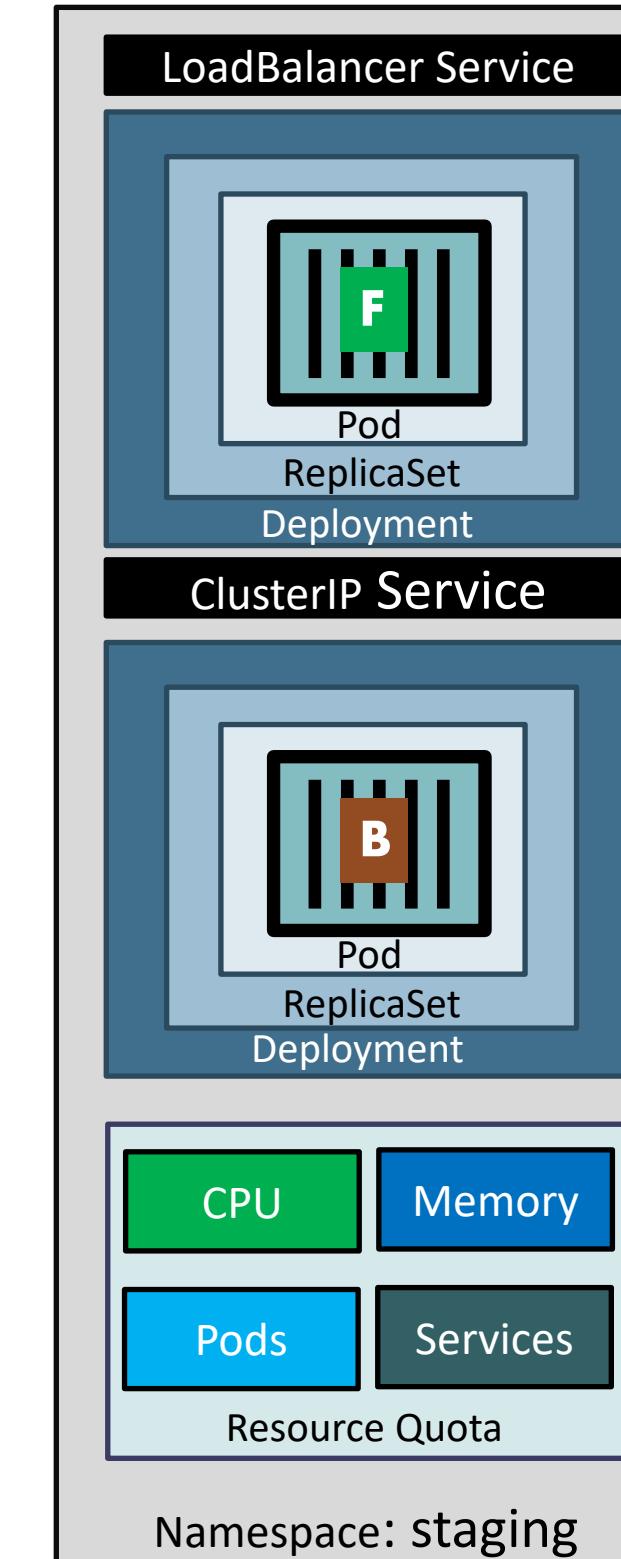
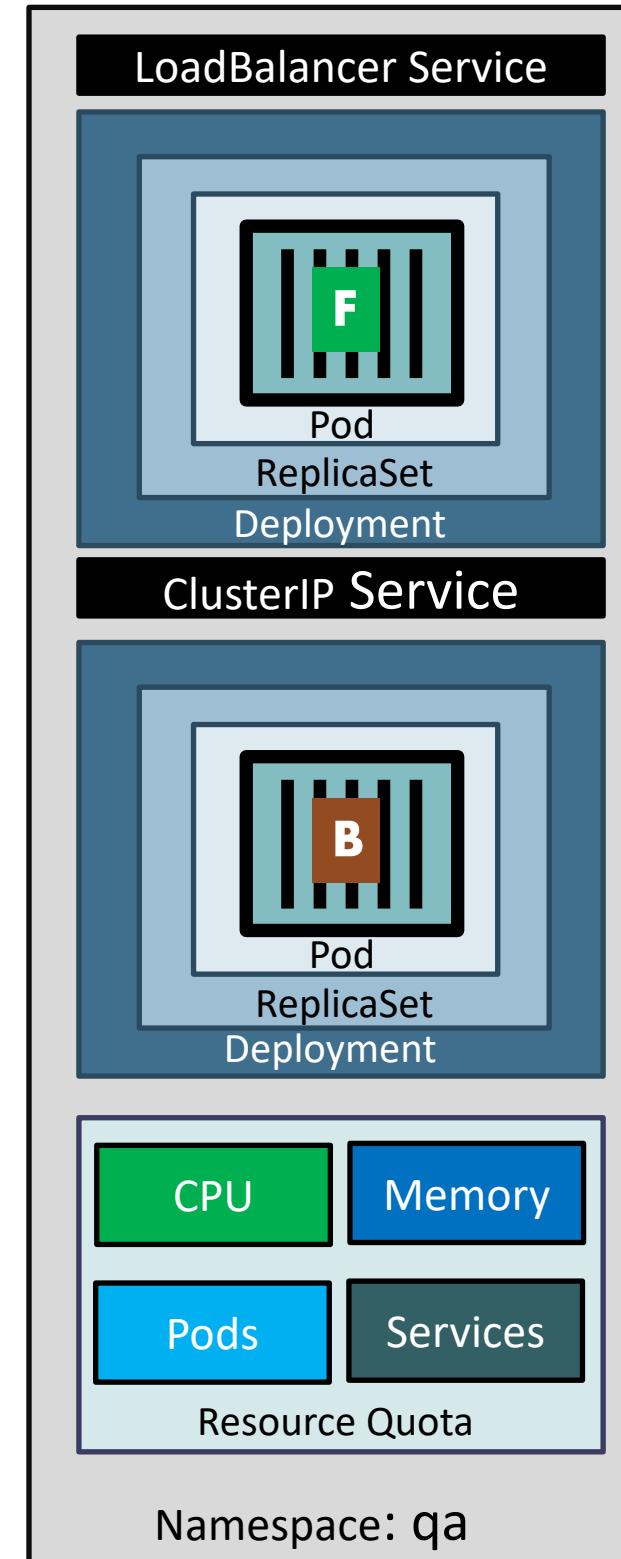
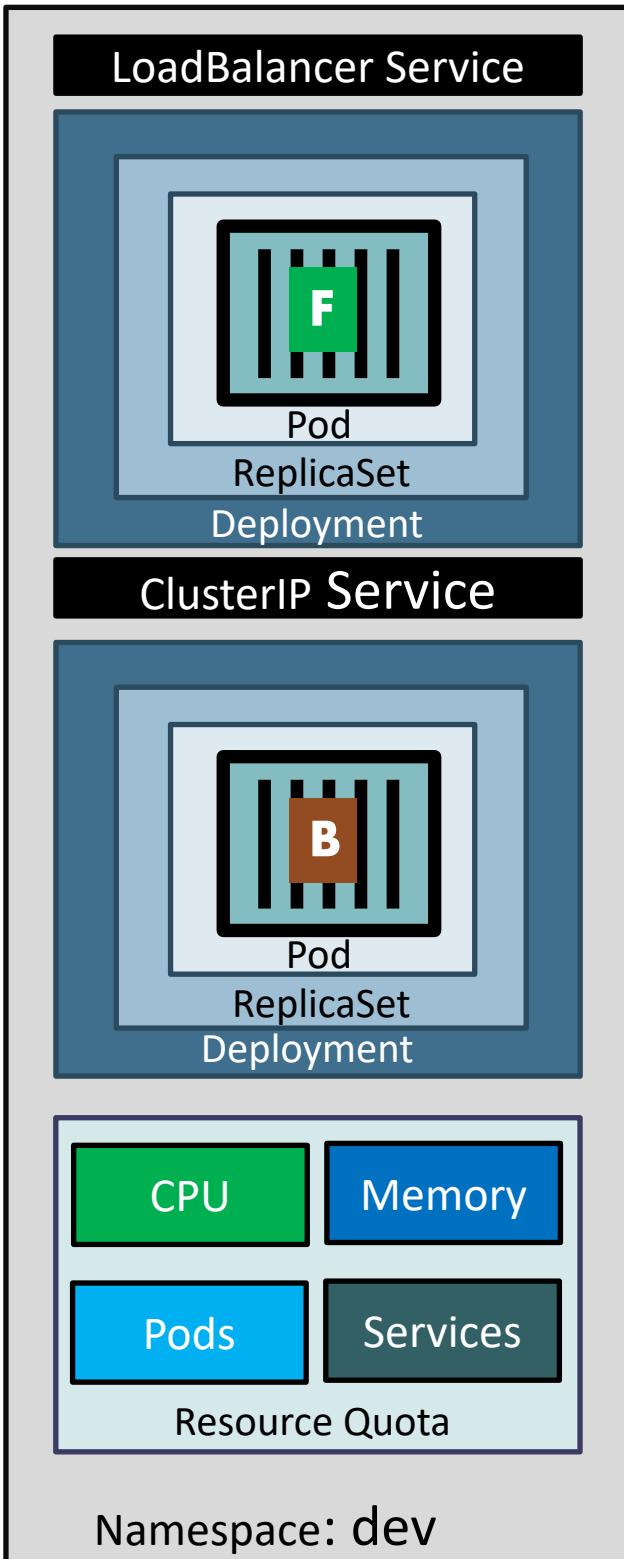


Namespace: staging

## Limit Range Manifest

```
apiVersion: v1
kind: LimitRange
metadata:
  name: default-cpu-mem-limit-range
  namespace: dev3
spec:
  limits:
  - default:
      memory: "512Mi"
      cpu: "500m"
    defaultRequest:
      memory: "256Mi"
      cpu: "300m"
  type: Container
```

# Resource Quota



## Resource Quota Manifest

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: ns-resource-quota
  namespace: dev3
spec:
  hard:
    requests.cpu: "1"
    requests.memory: 1Gi
    limits.cpu: "2"
    limits.memory: 2Gi
  pods: "5"
  configmaps: "5"
  persistentvolumeclaims: "5"
  secrets: "5"
  services: "5"
```



Amazon RDS

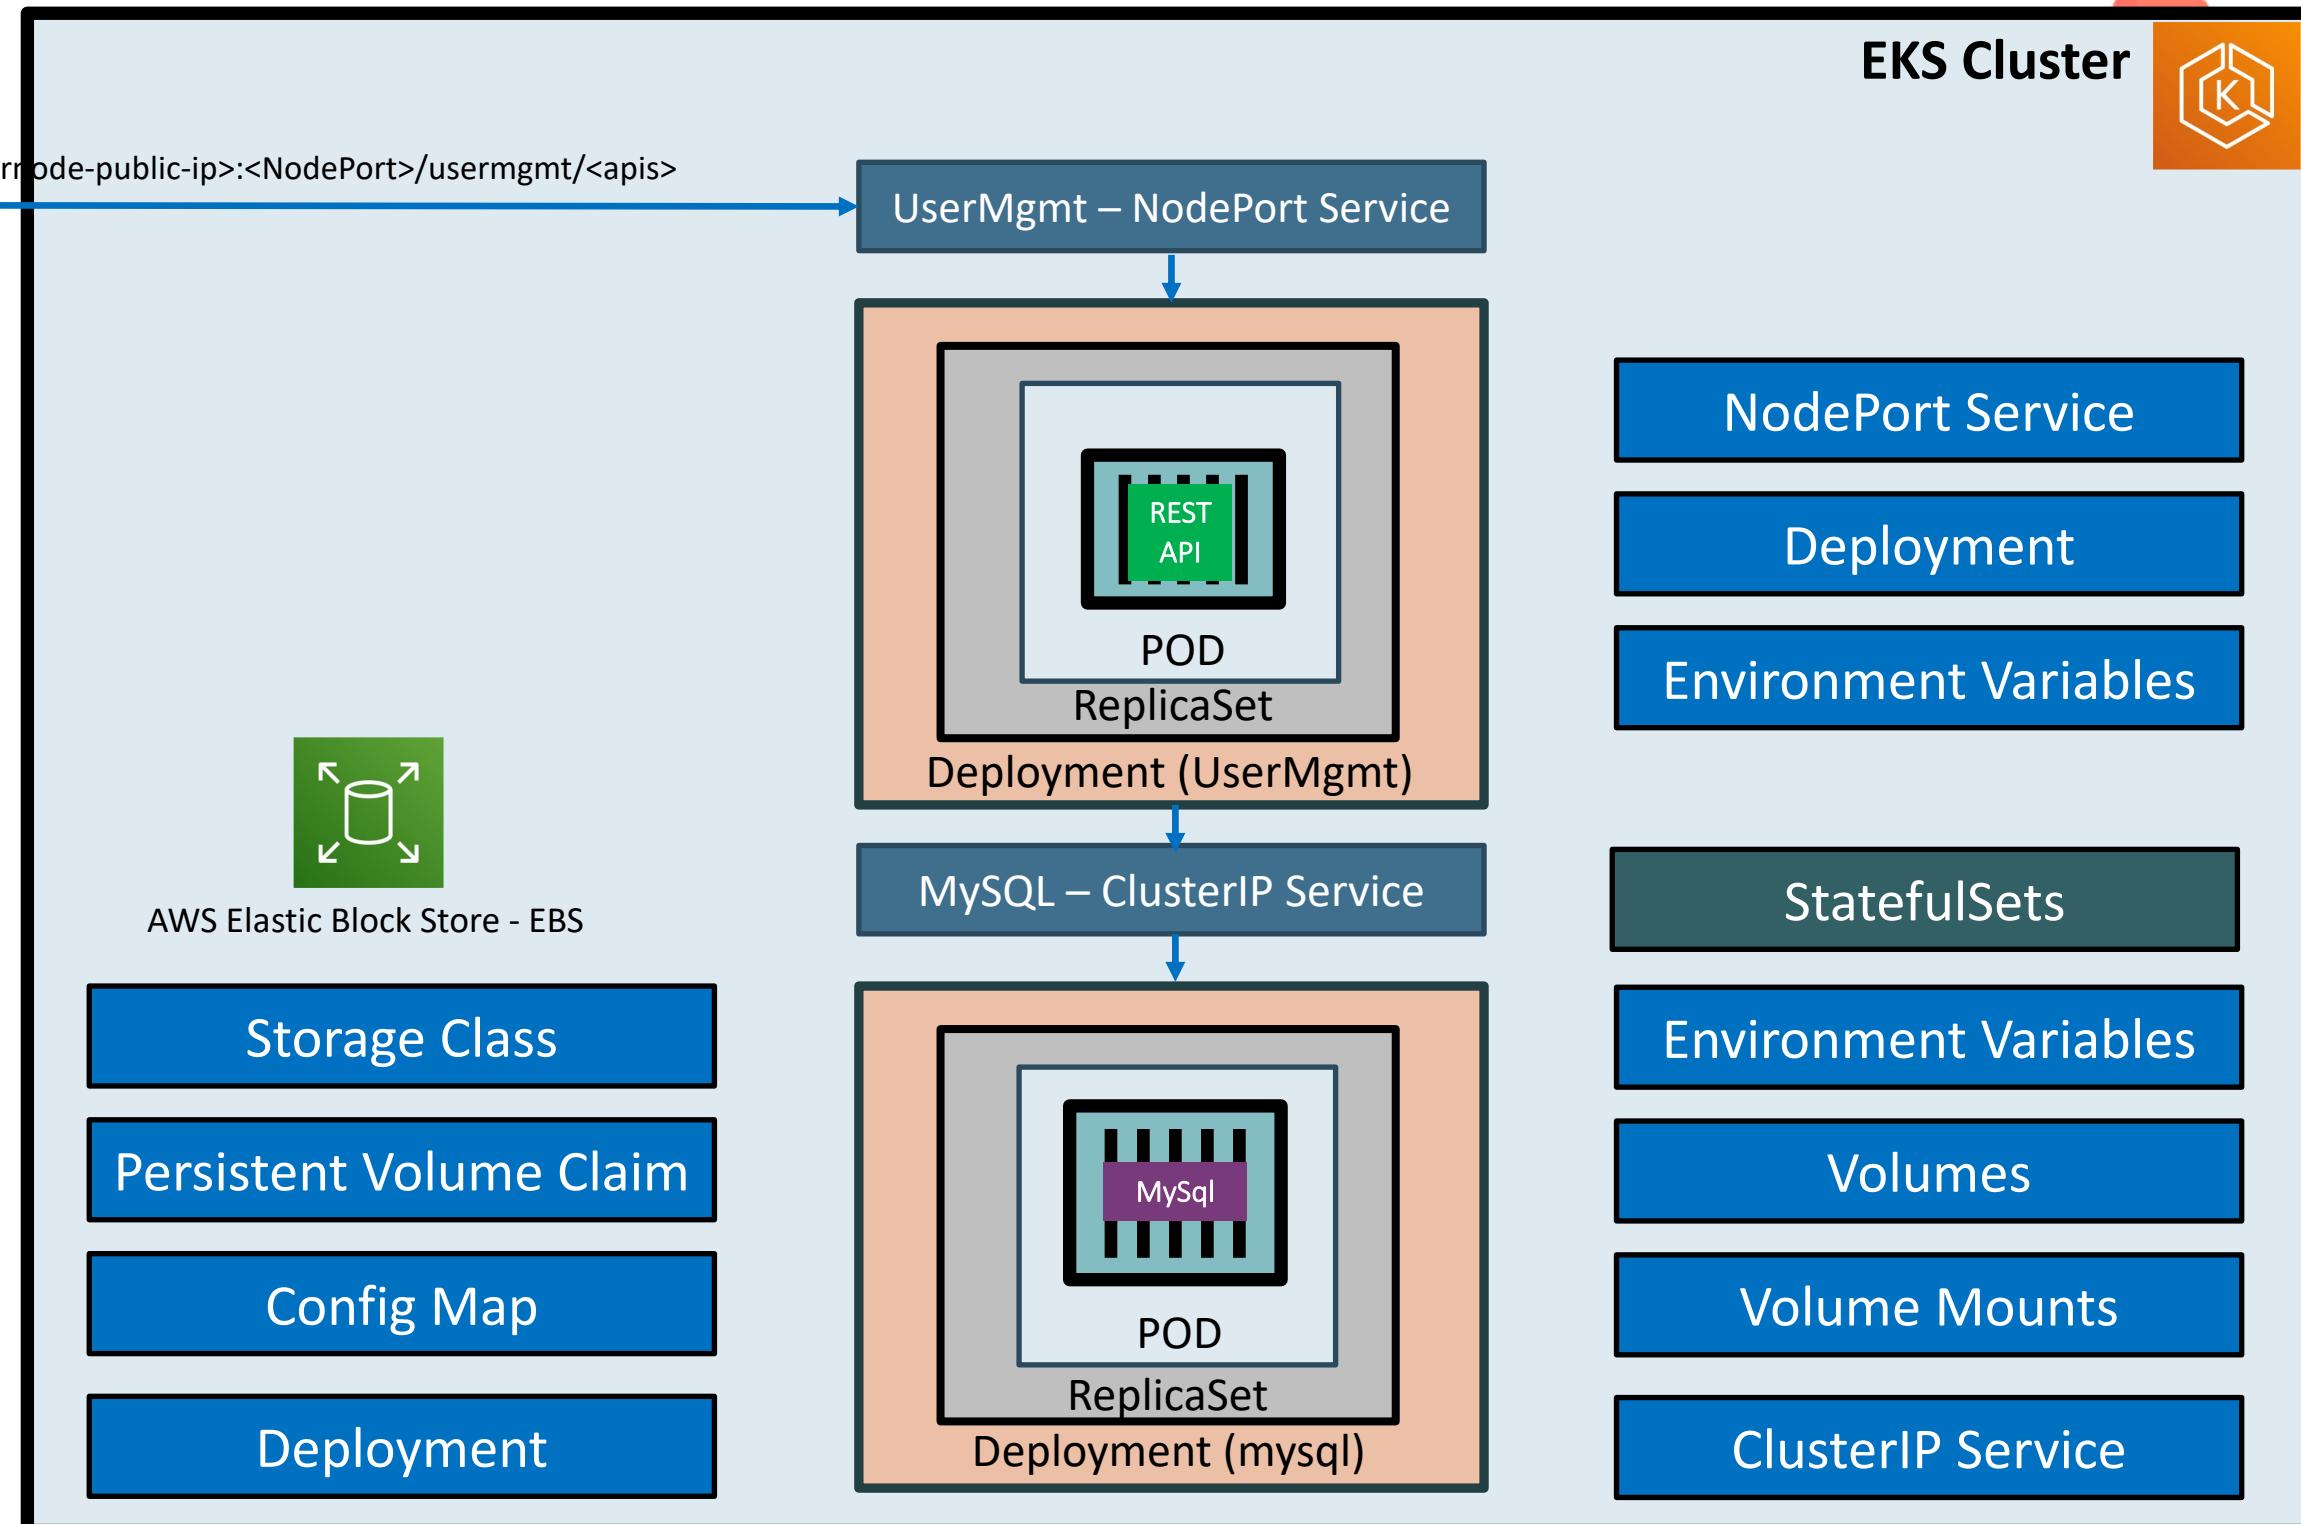
# AWS EKS & RDS Database

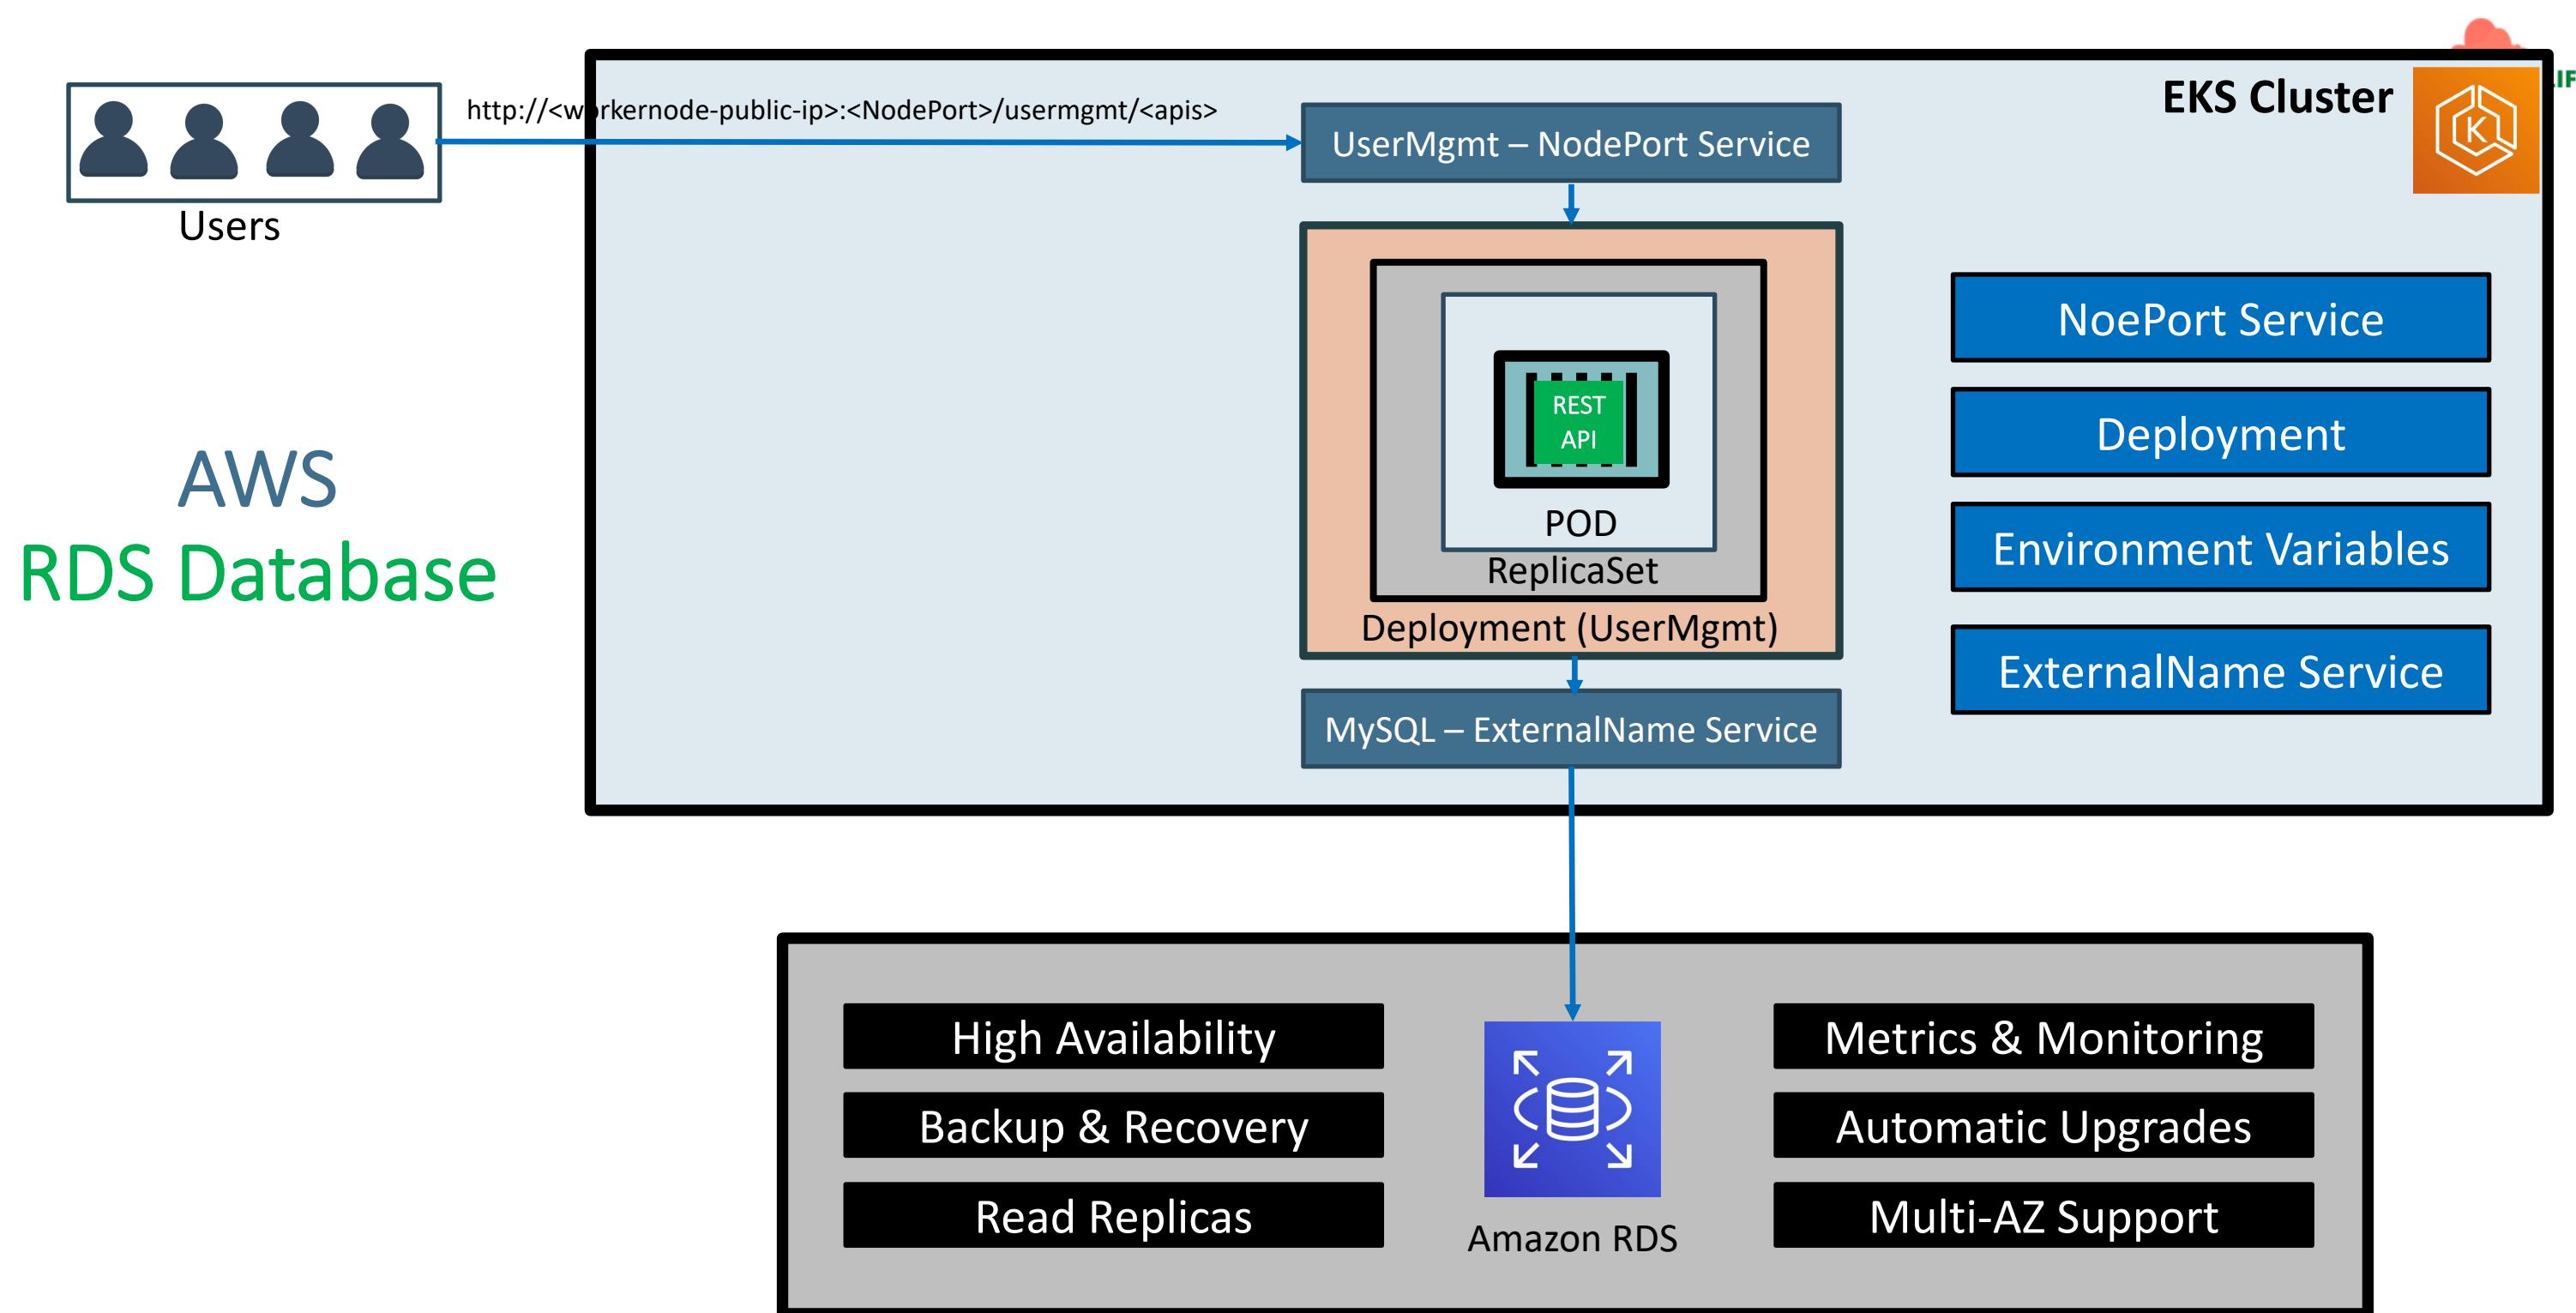


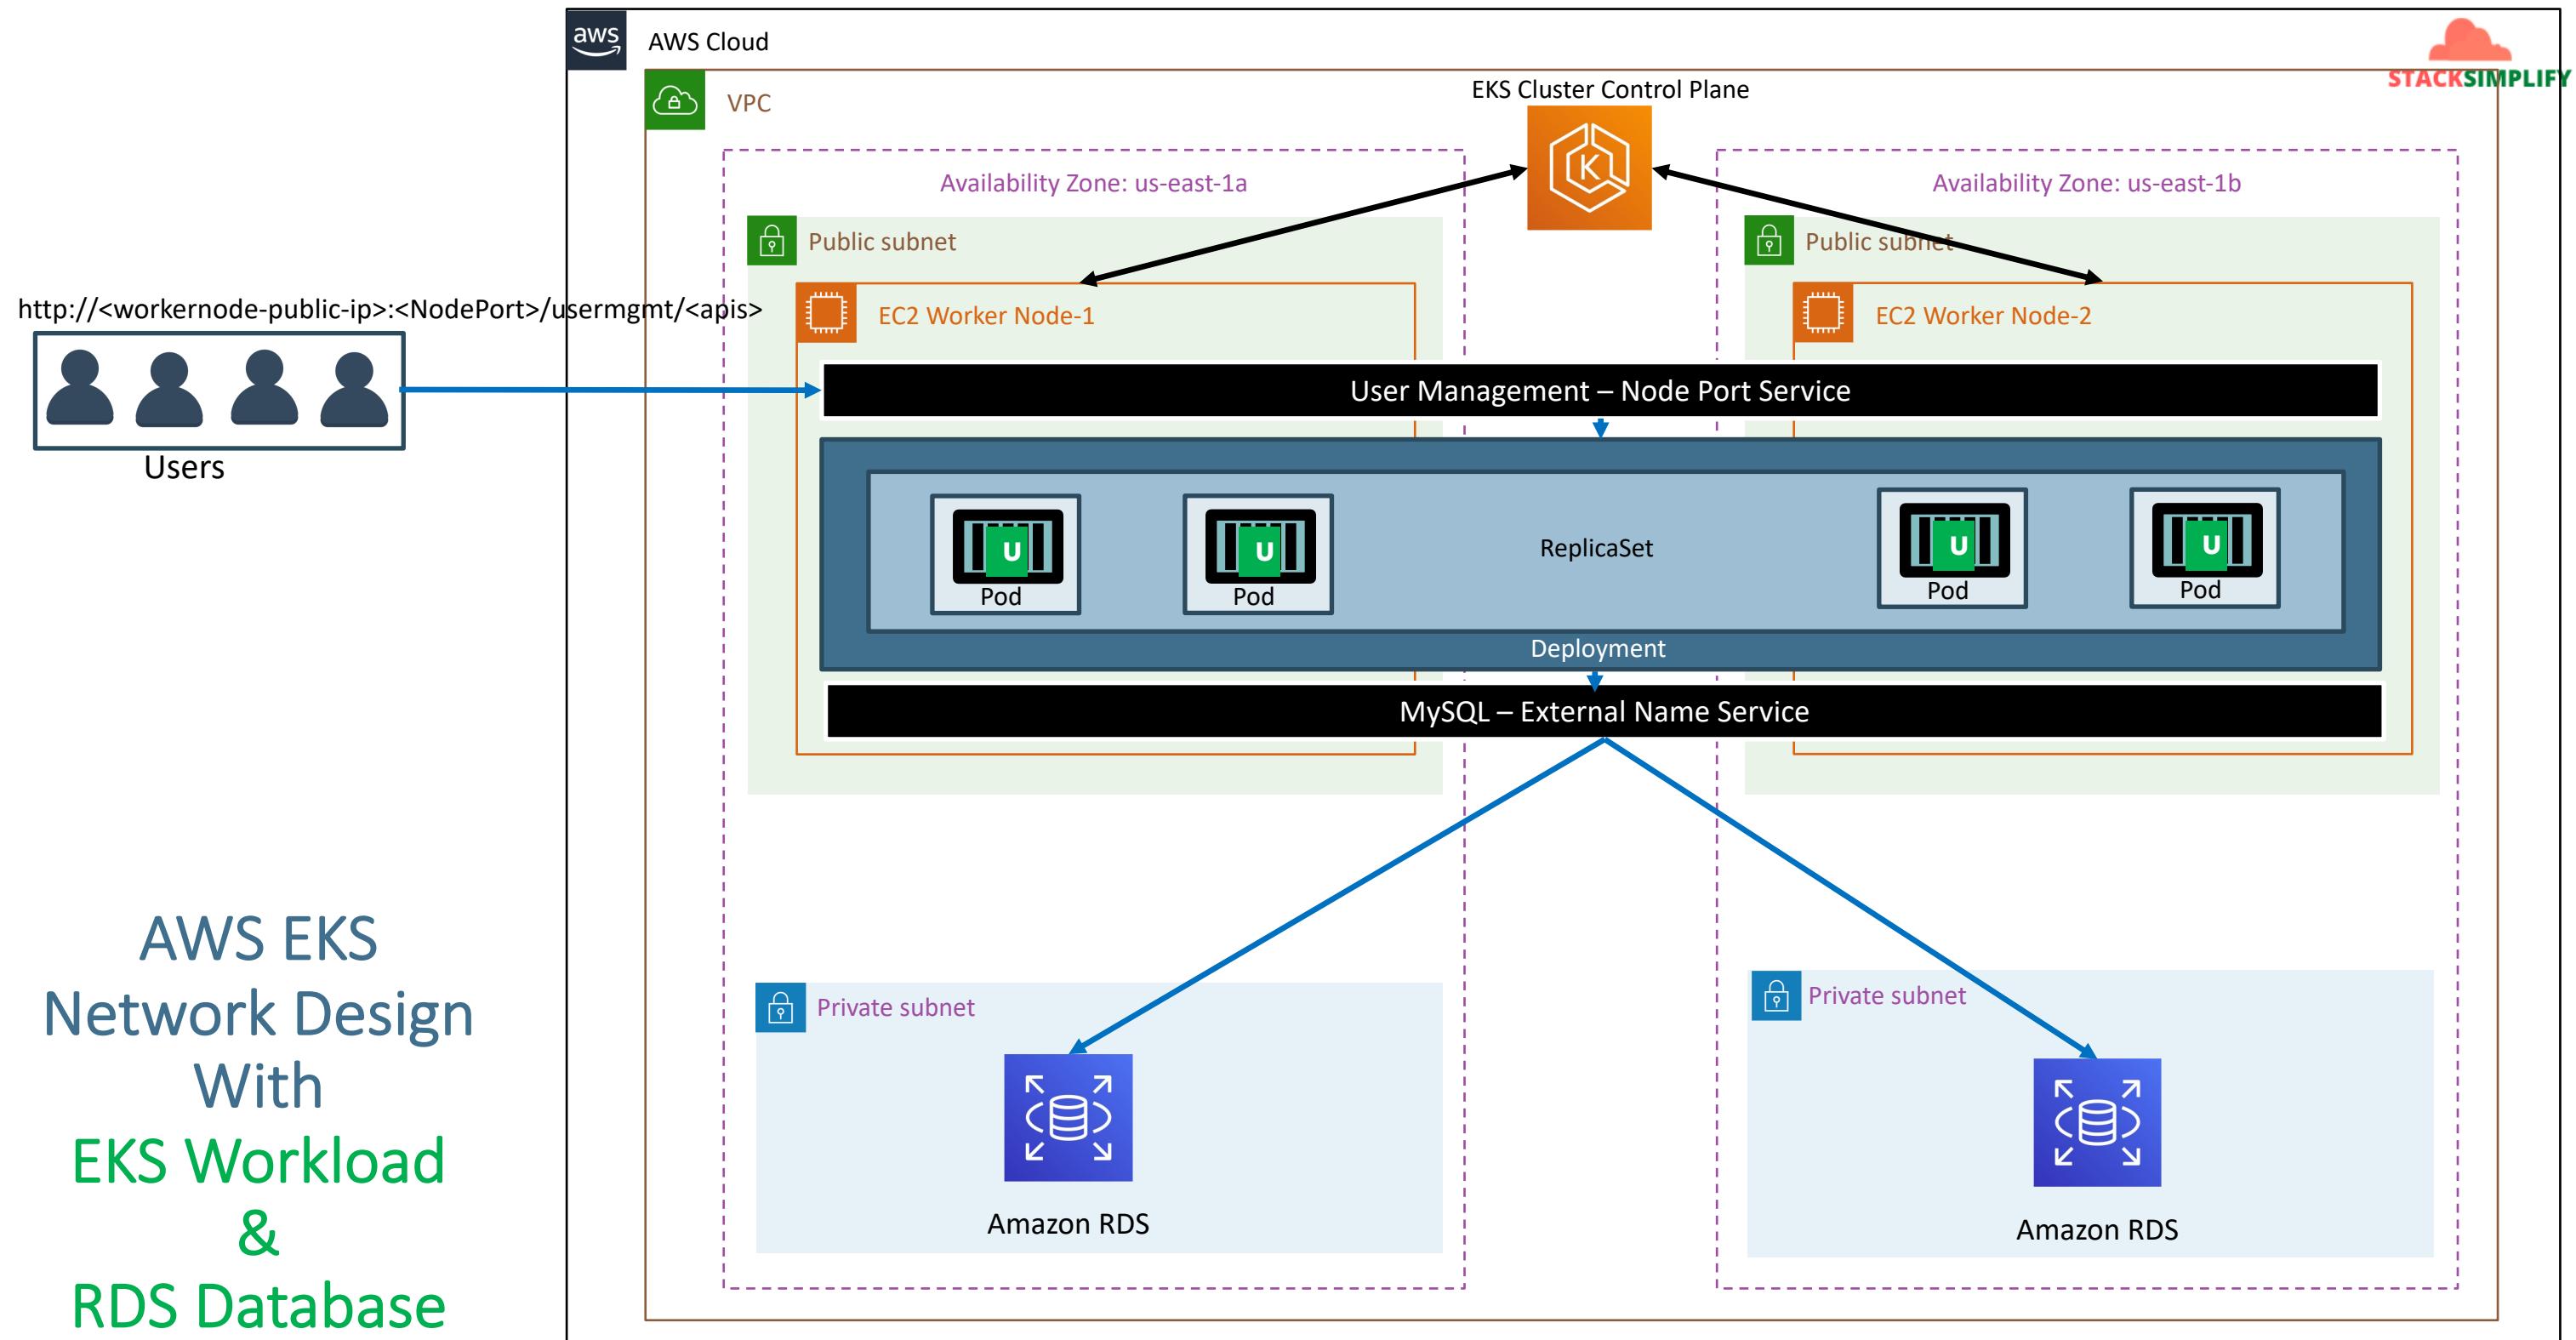


# EKS Storage EBS CSI Driver

- Drawbacks of EBS CSI for MySQL DB
- Complex setup to achieve HA
- Complex Multi-Az support for EBS
- Complex Master-Master MySQL setup
- Complex Master-Slave MySQL setup
- No Automatic Backup & Recovery
- No Auto-Upgrade MySQL

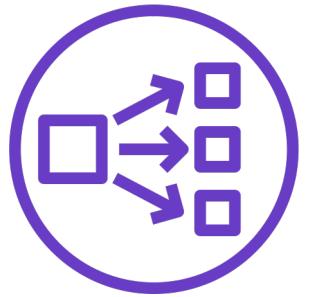




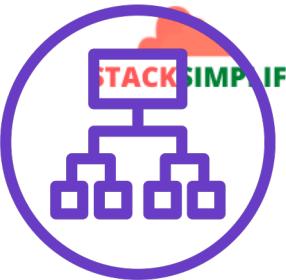




Classic  
Load Balancer



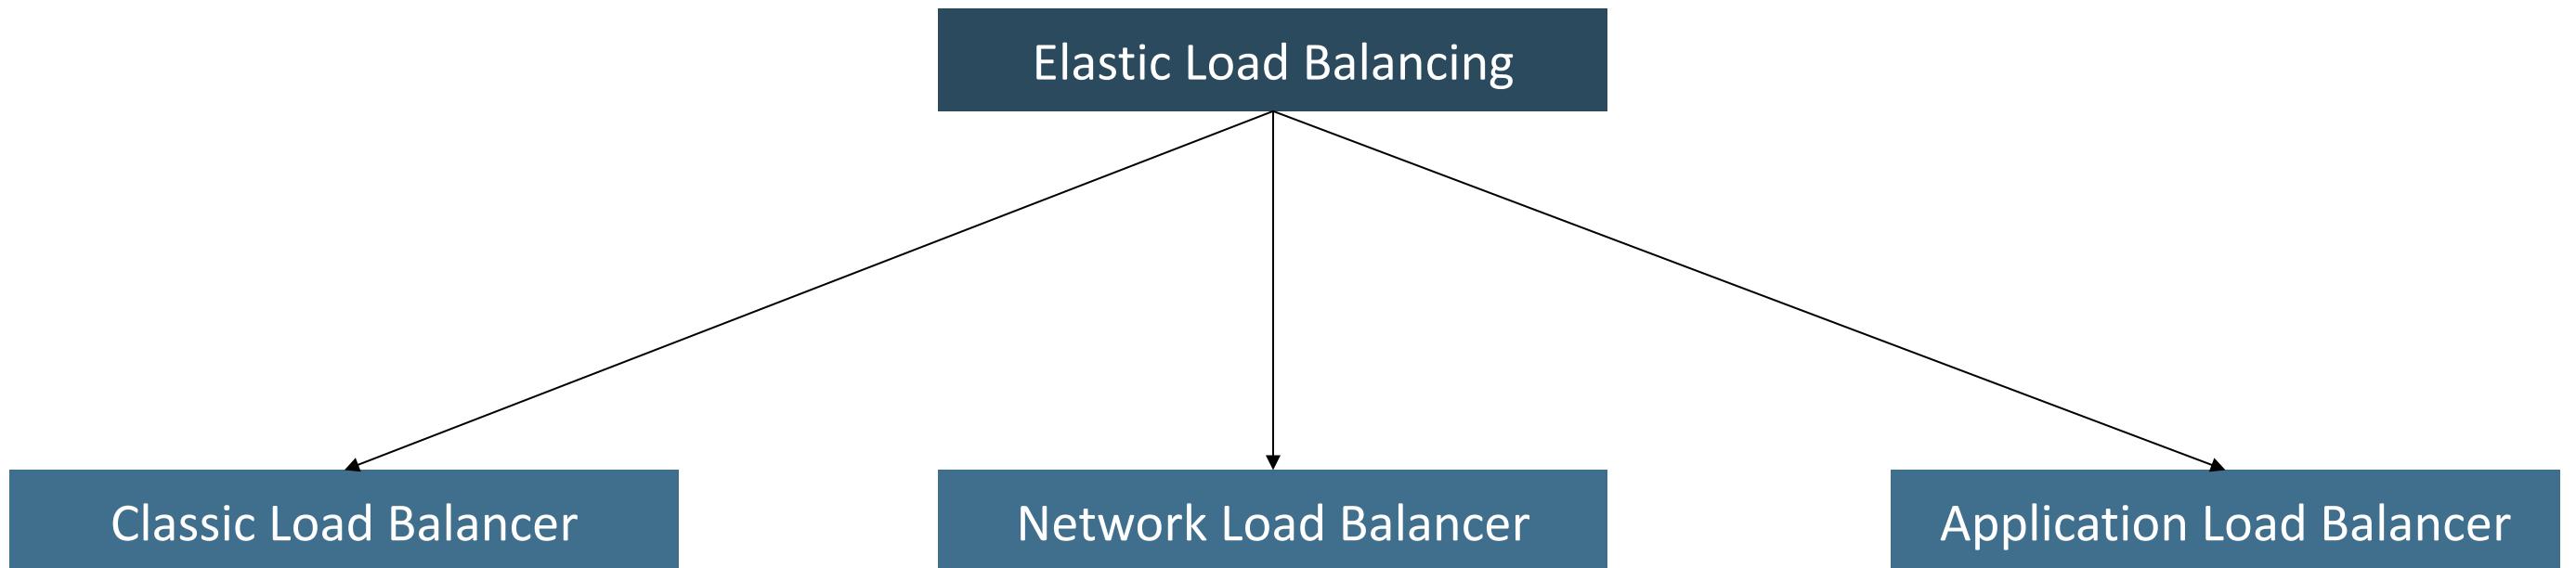
Network  
Load Balancer



Application  
Load Balancer



# AWS Elastic Load Balancing Overview



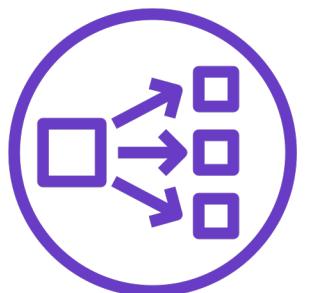
<https://aws.amazon.com/elasticloadbalancing/features/#compare>



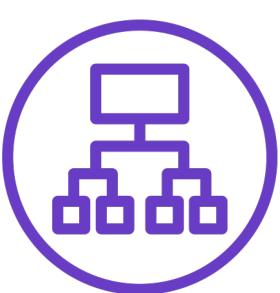
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



Application  
Load Balancer

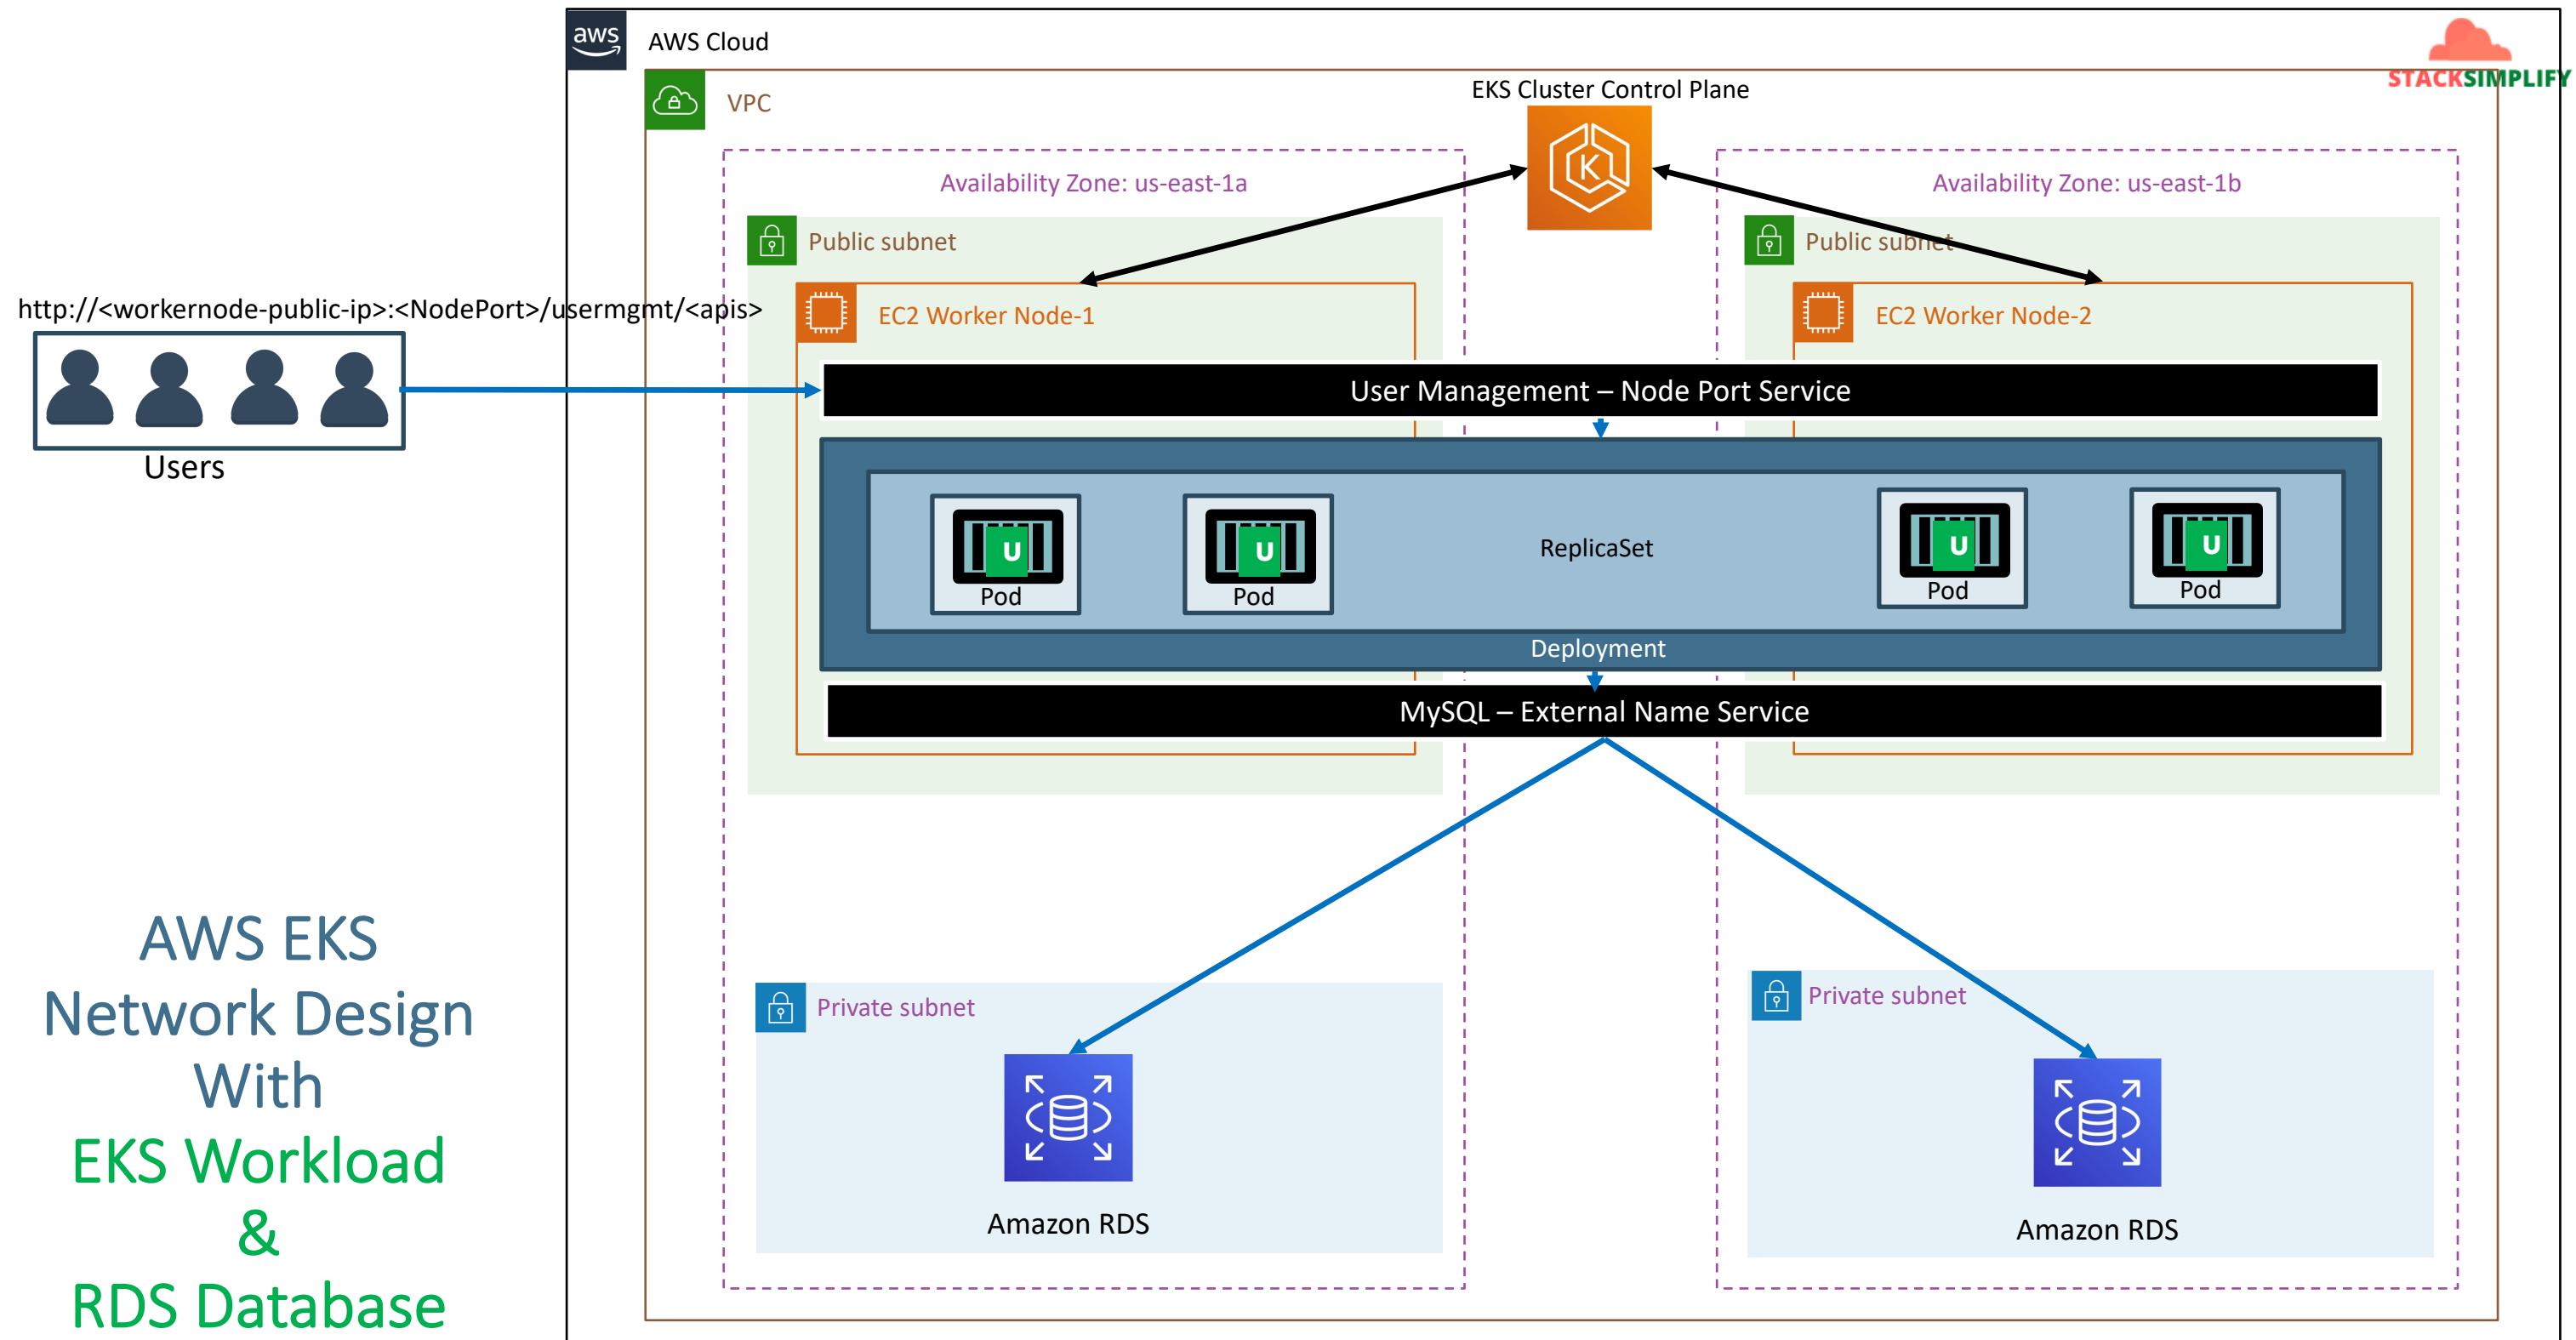


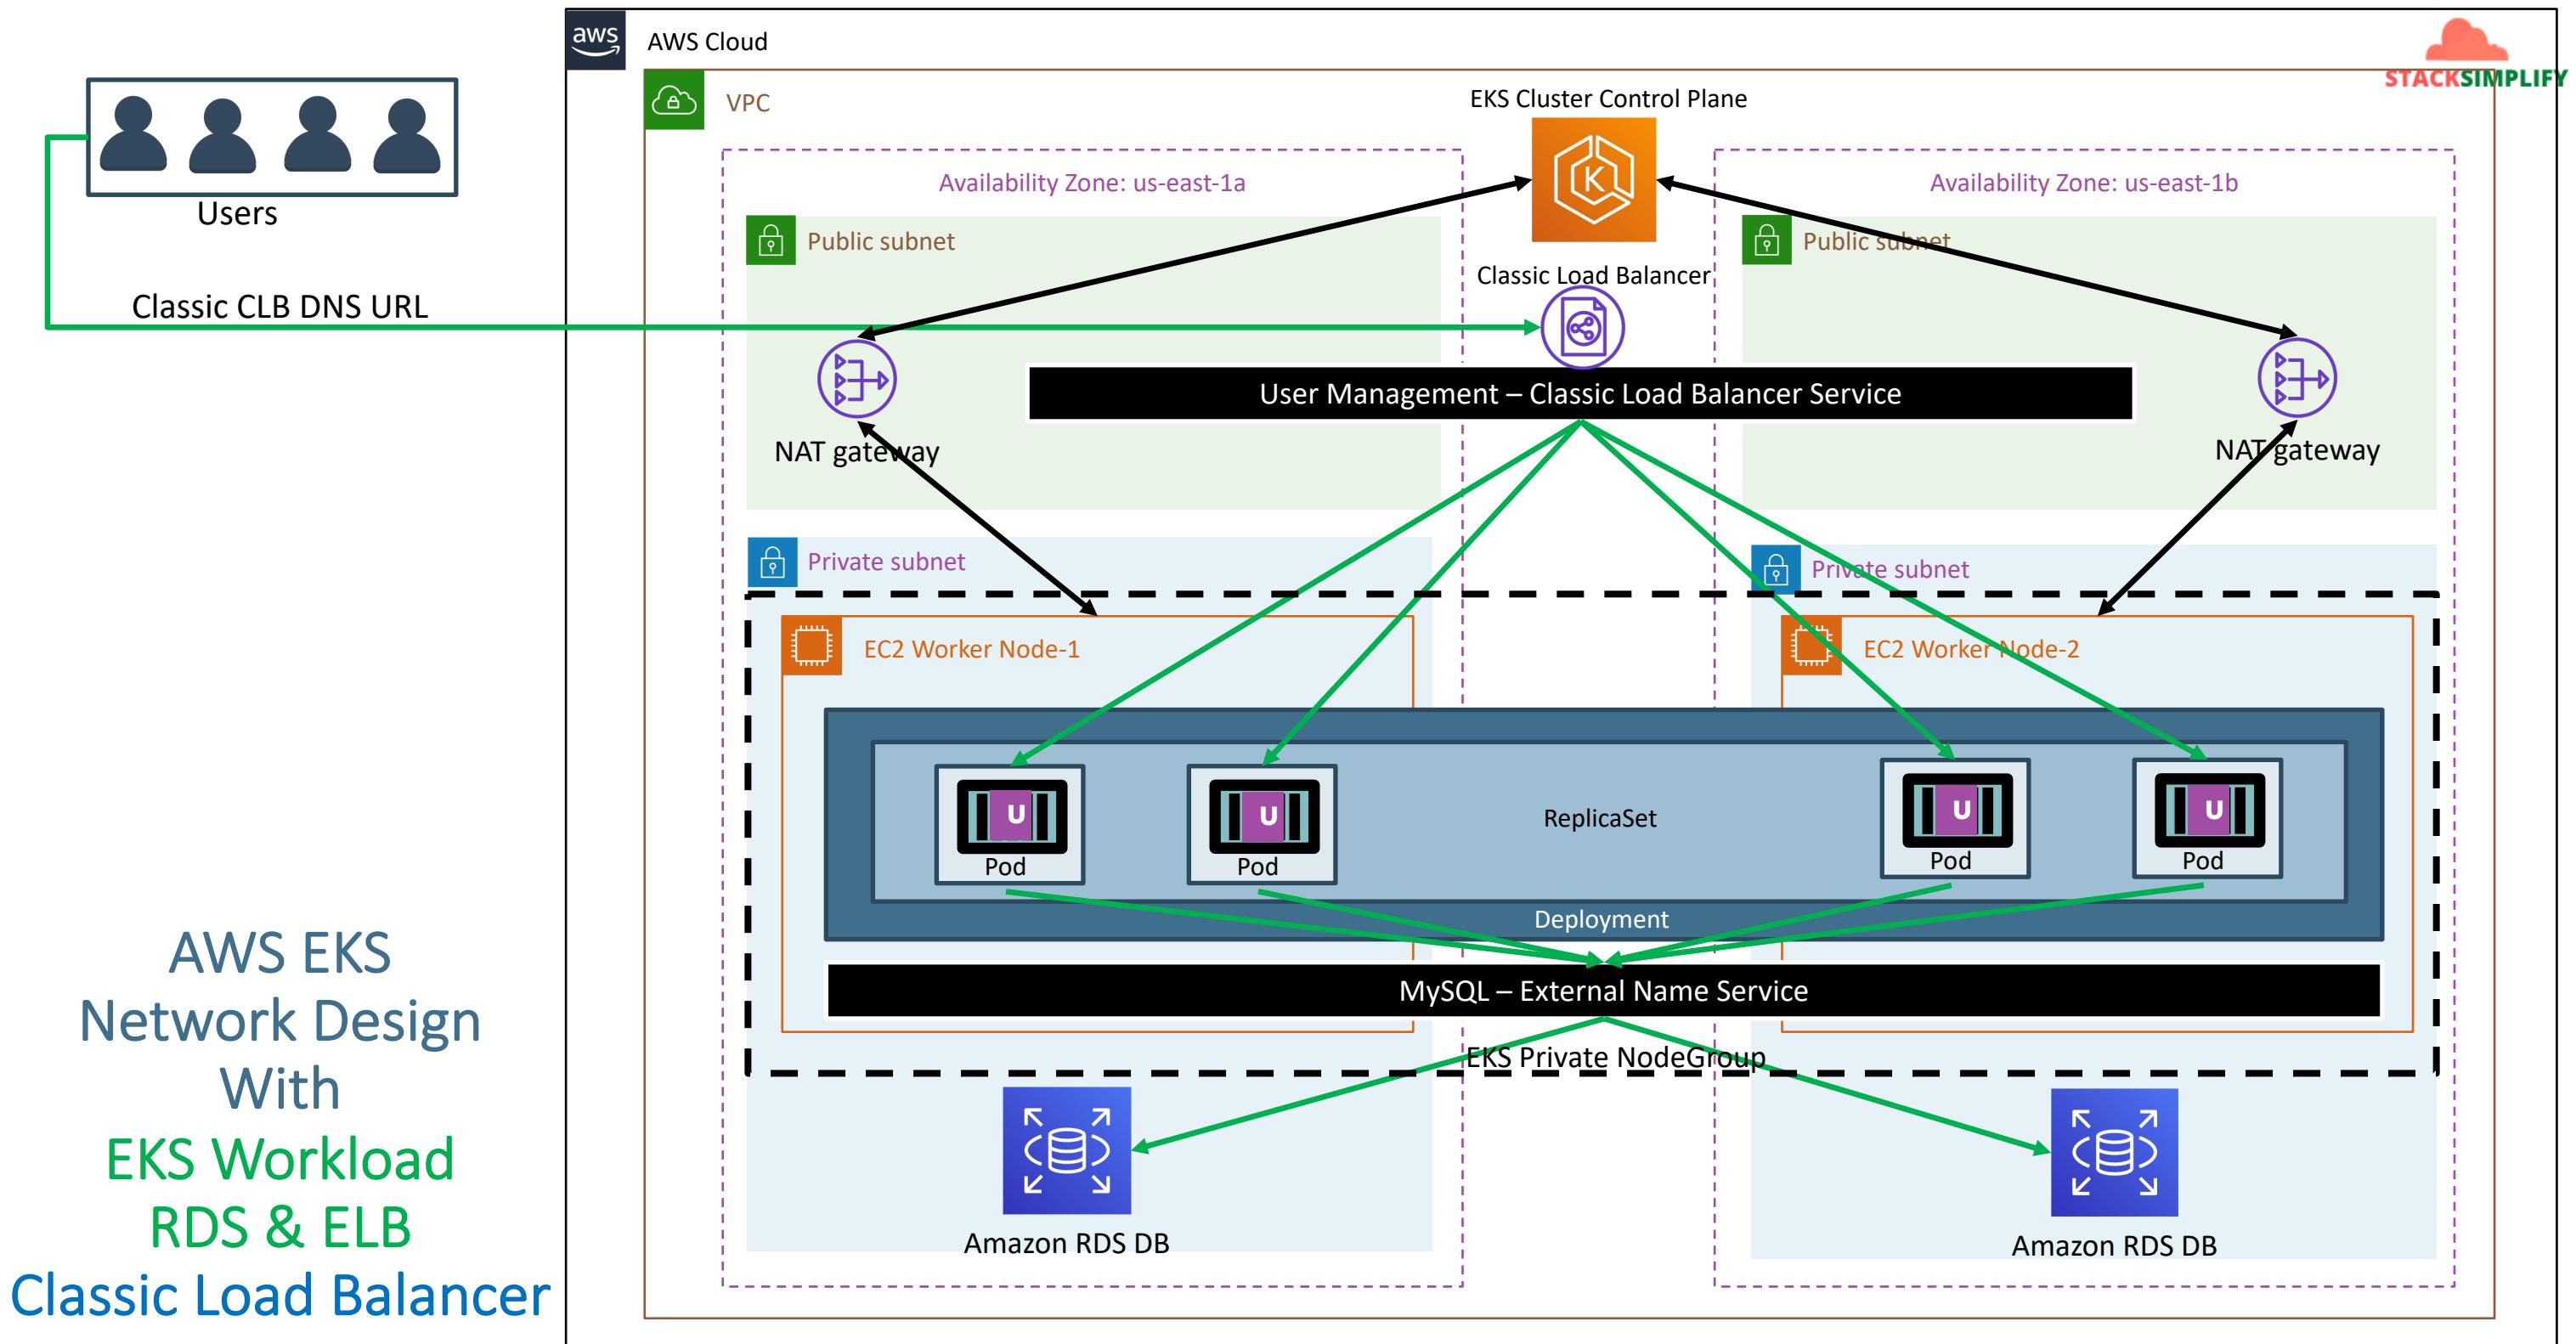
Amazon RDS



# AWS EKS & RDS & ELB **Classic Load Balancer**





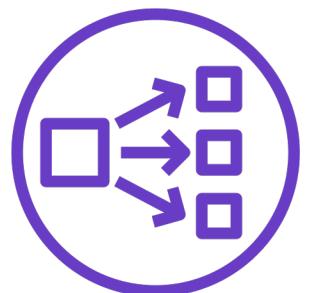




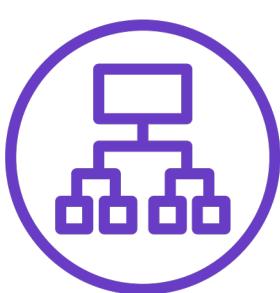
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



Application  
Load Balancer



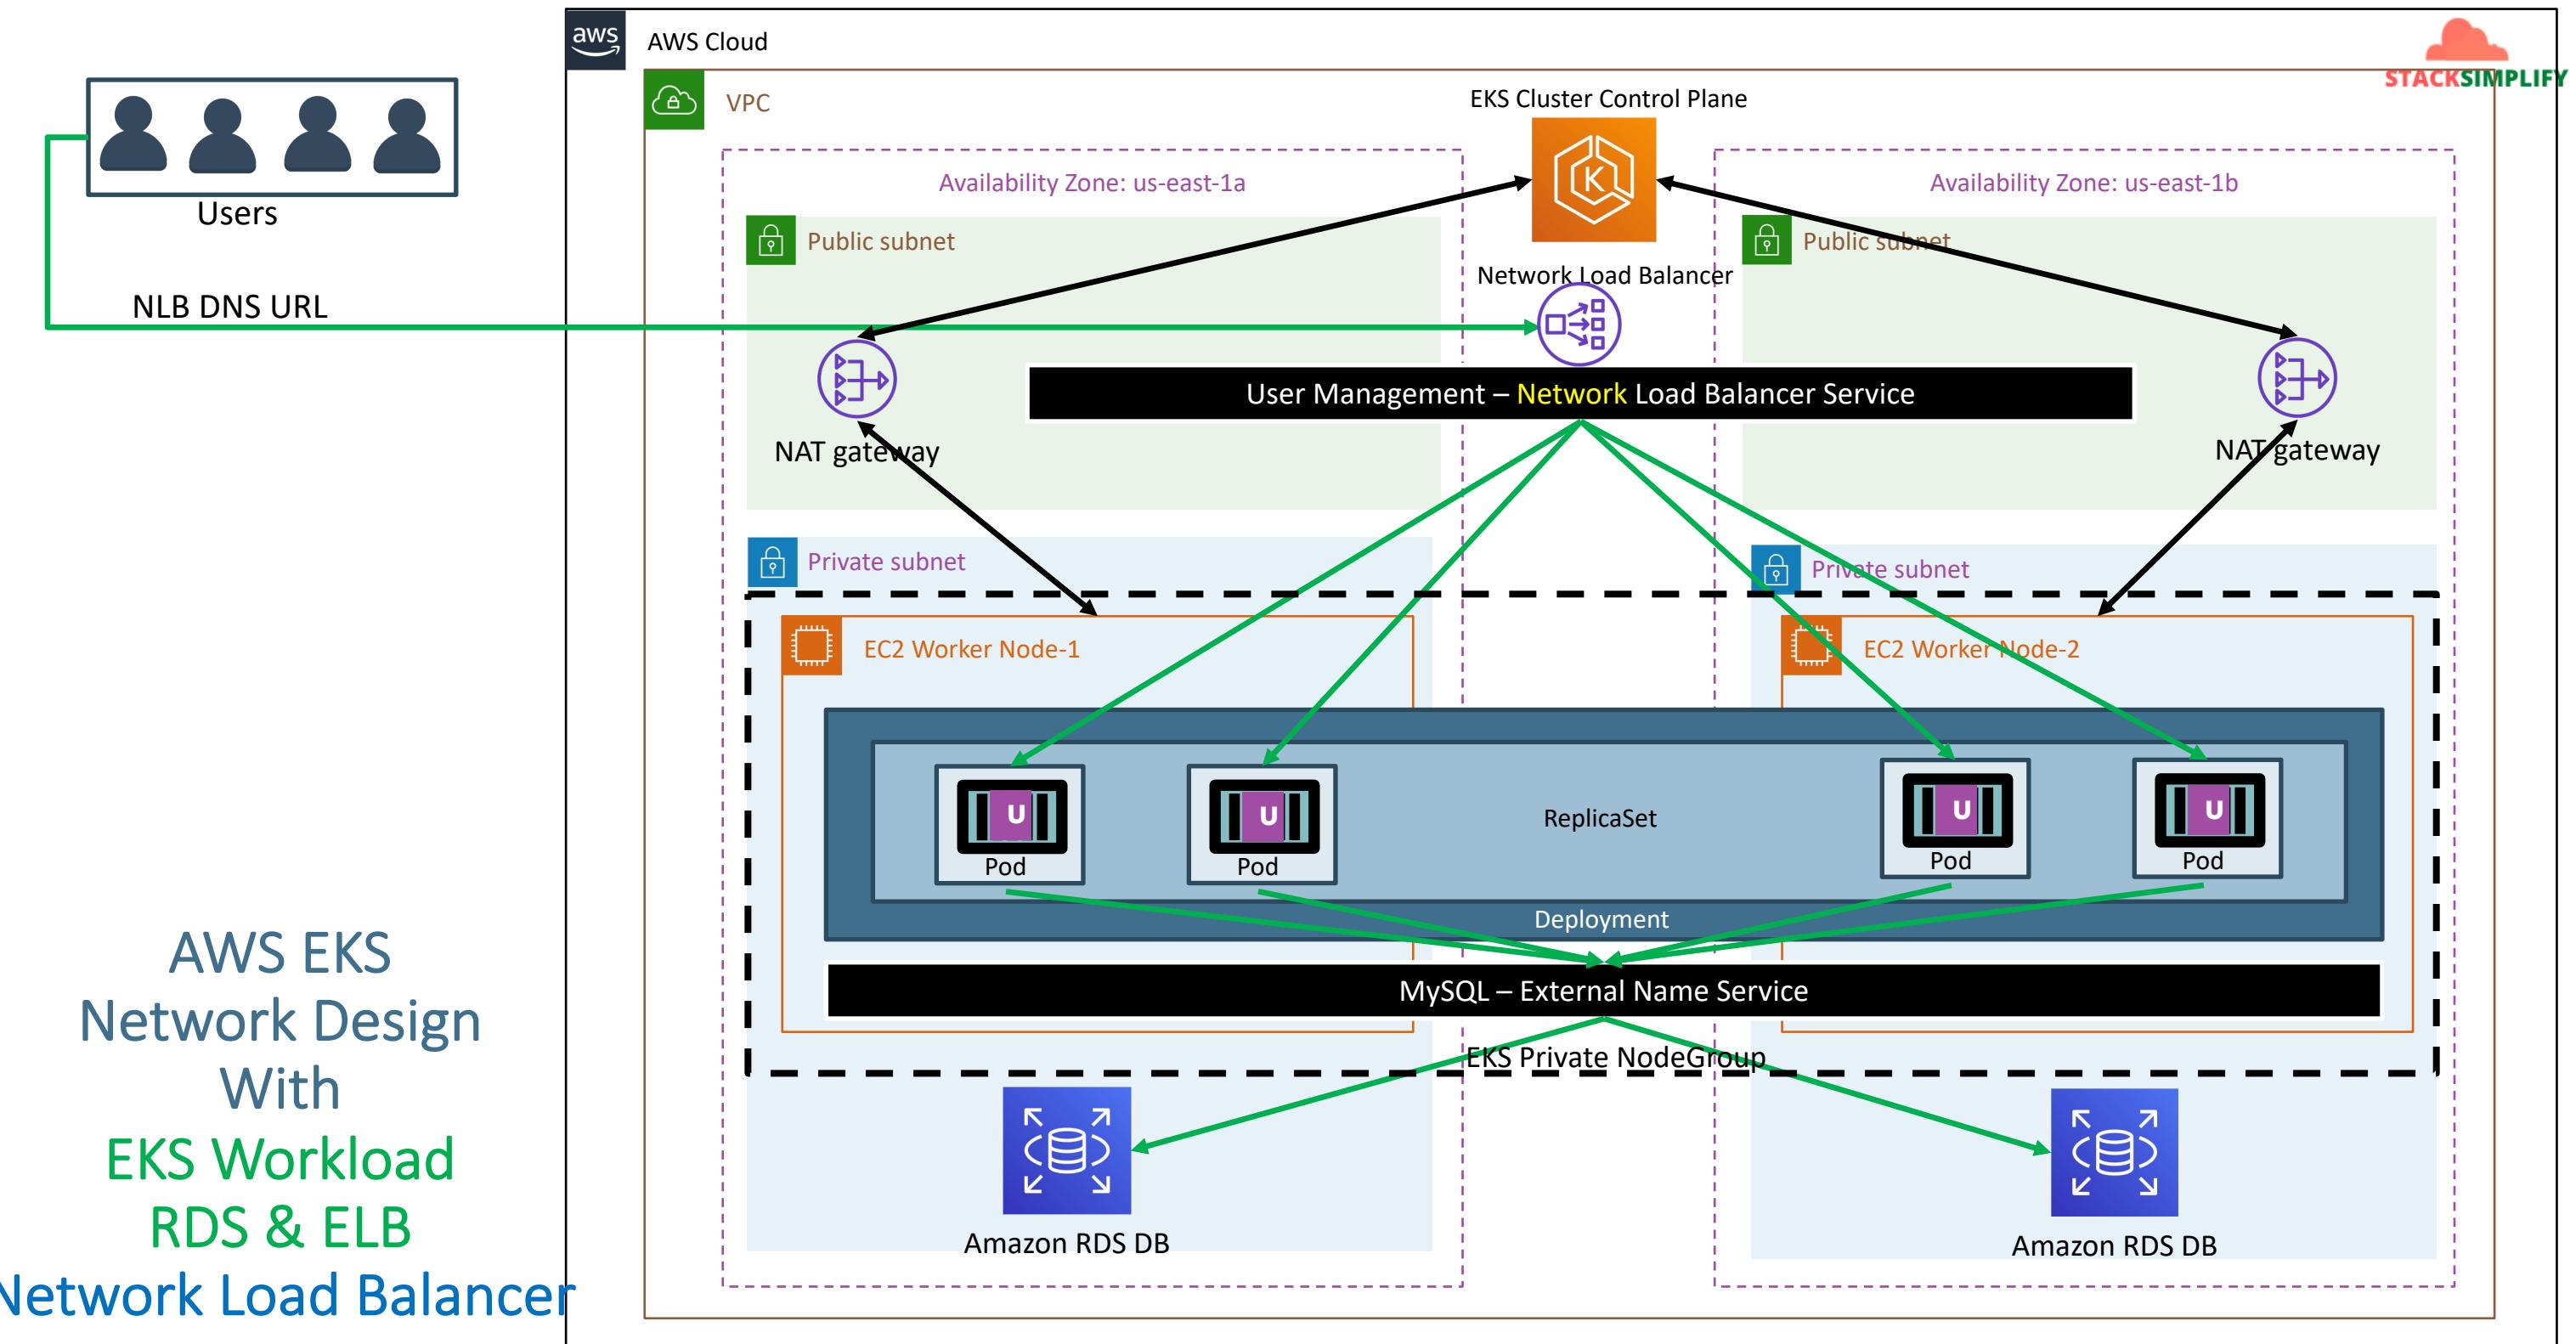
Amazon RDS



AWS EKS  
&  
RDS & ELB

Network Load Balancer

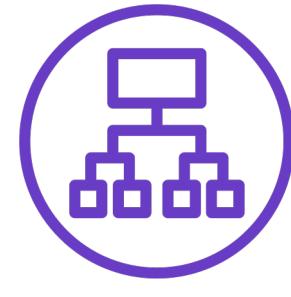




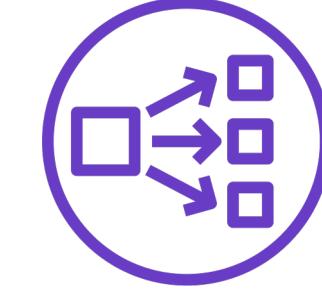
# AWS Load Balancer Controller Updates - Start



Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress



# AWS EKS

## Load Balancer Controller

## Kubernetes Ingress

## Application Load Balancer



14 Demos

AWS Load Balancer  
Controller Install

D1

Ingress Basics

D2

Ingress Context Path  
Routing

D3

Ingress SSL

D4

Ingress SSL Redirect

D5

External DNS Install

D6

Ingress + External DNS

D7

D8

k8s Service + External DNS

D9

Ingress Name based Virtual  
Host Routing

D10

SSL Discovery - Host

D11

SSL Discovery - TLS

D12

Ingress Groups

D13

Ingress Target Type - IP

D14

Ingress Internal ALB

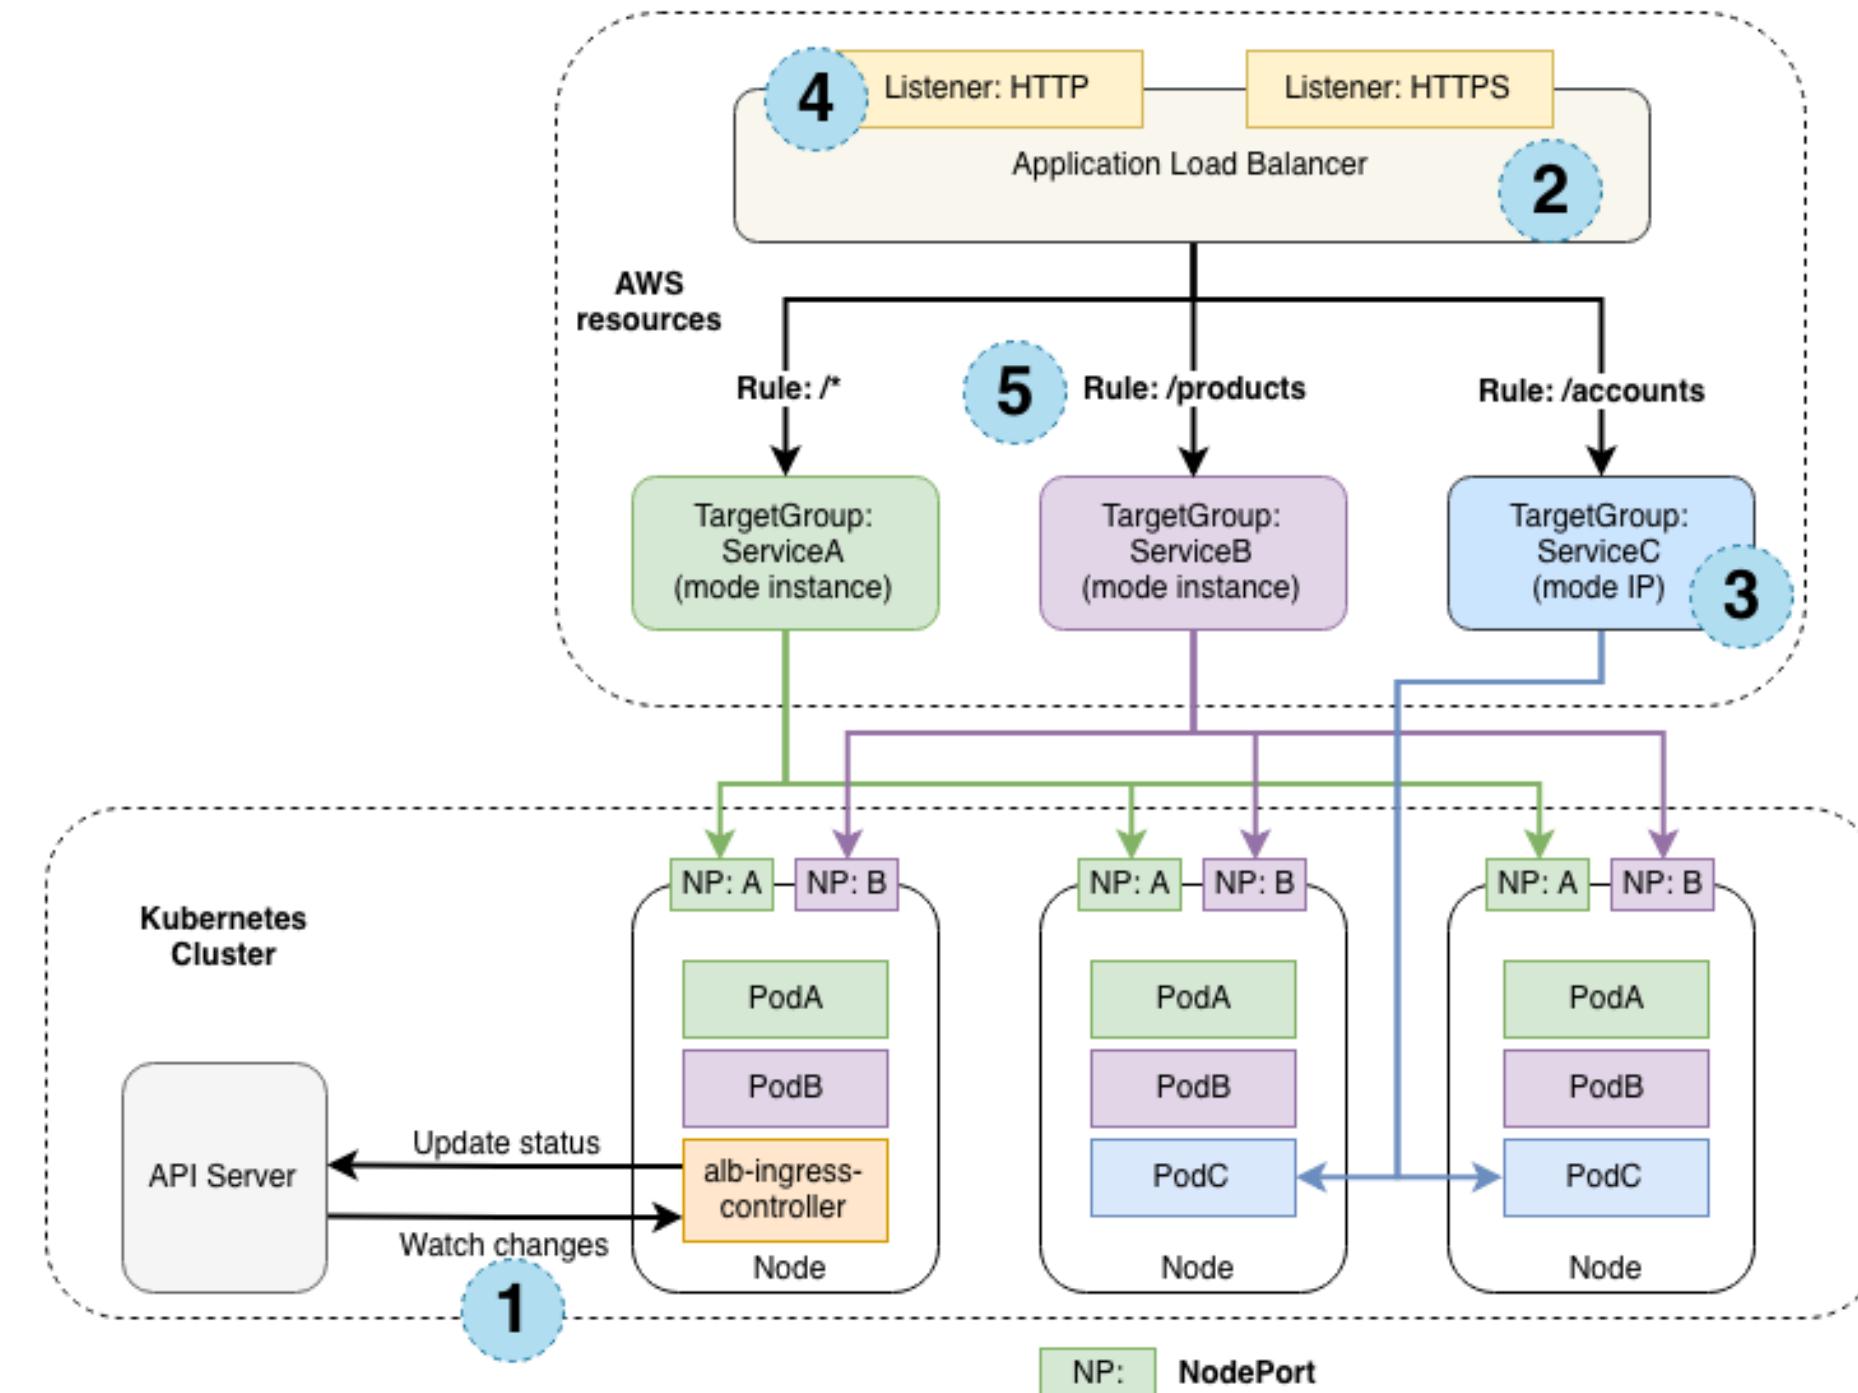
EKS  
ALB  
Ingress

Git Repo: <https://github.com/stacksimplify/aws-eks-kubernetes-masterclass/tree/master/08-NEW-ELB-Application-LoadBalancers>

-  08-01-Load-Balancer-Controller-Install
-  08-02-ALB-Ingress-Basics
-  08-03-ALB-Ingress-ContextPath-Based-R...
-  08-04-ALB-Ingress-SSL
-  08-05-ALB-Ingress-SSL-Redirect
-  08-06-Deploy-ExternalDNS-on-EKS
-  08-07-Use-ExternalDNS-for-k8s-Ingress

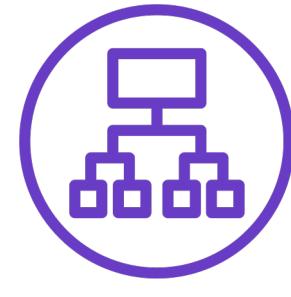
-  08-08-Use-ExternalDNS-for-k8s-Service
-  08-09-NameBasedVirtualHost-Routing
-  08-10-Ingress-SSL-Discovery-host
-  08-11-Ingress-SSL-Discovery-tls
-  08-12-IngressGroups
-  08-13-Ingress-TargetType-IP
-  08-14-Ingress-Internal-LB

# How Ingress Works?

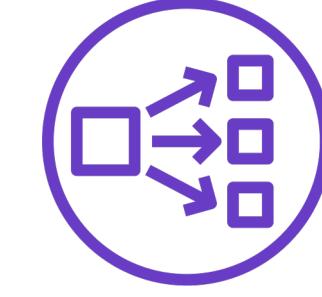




Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress



# AWS EKS

## Load Balancer Controller

## Install



# Kubernetes Ingress

AWS ALB Ingress Controller

AWS Application Load Balancer Only

Doesn't support latest Ingress API Version  
(apiVersion: networking.k8s.io/v1) from  
Kubernetes Version 1.22 onwards

No Development releases, Soon it will be  
deprecated

Renamed & Redesigned

AWS Load Balancer Controller

AWS Application and Network Load Balancers

K8s Ingress Resource

K8s Service Resource

Supports latest and greatest from Ingress perspective

AWS Application Load Balancer

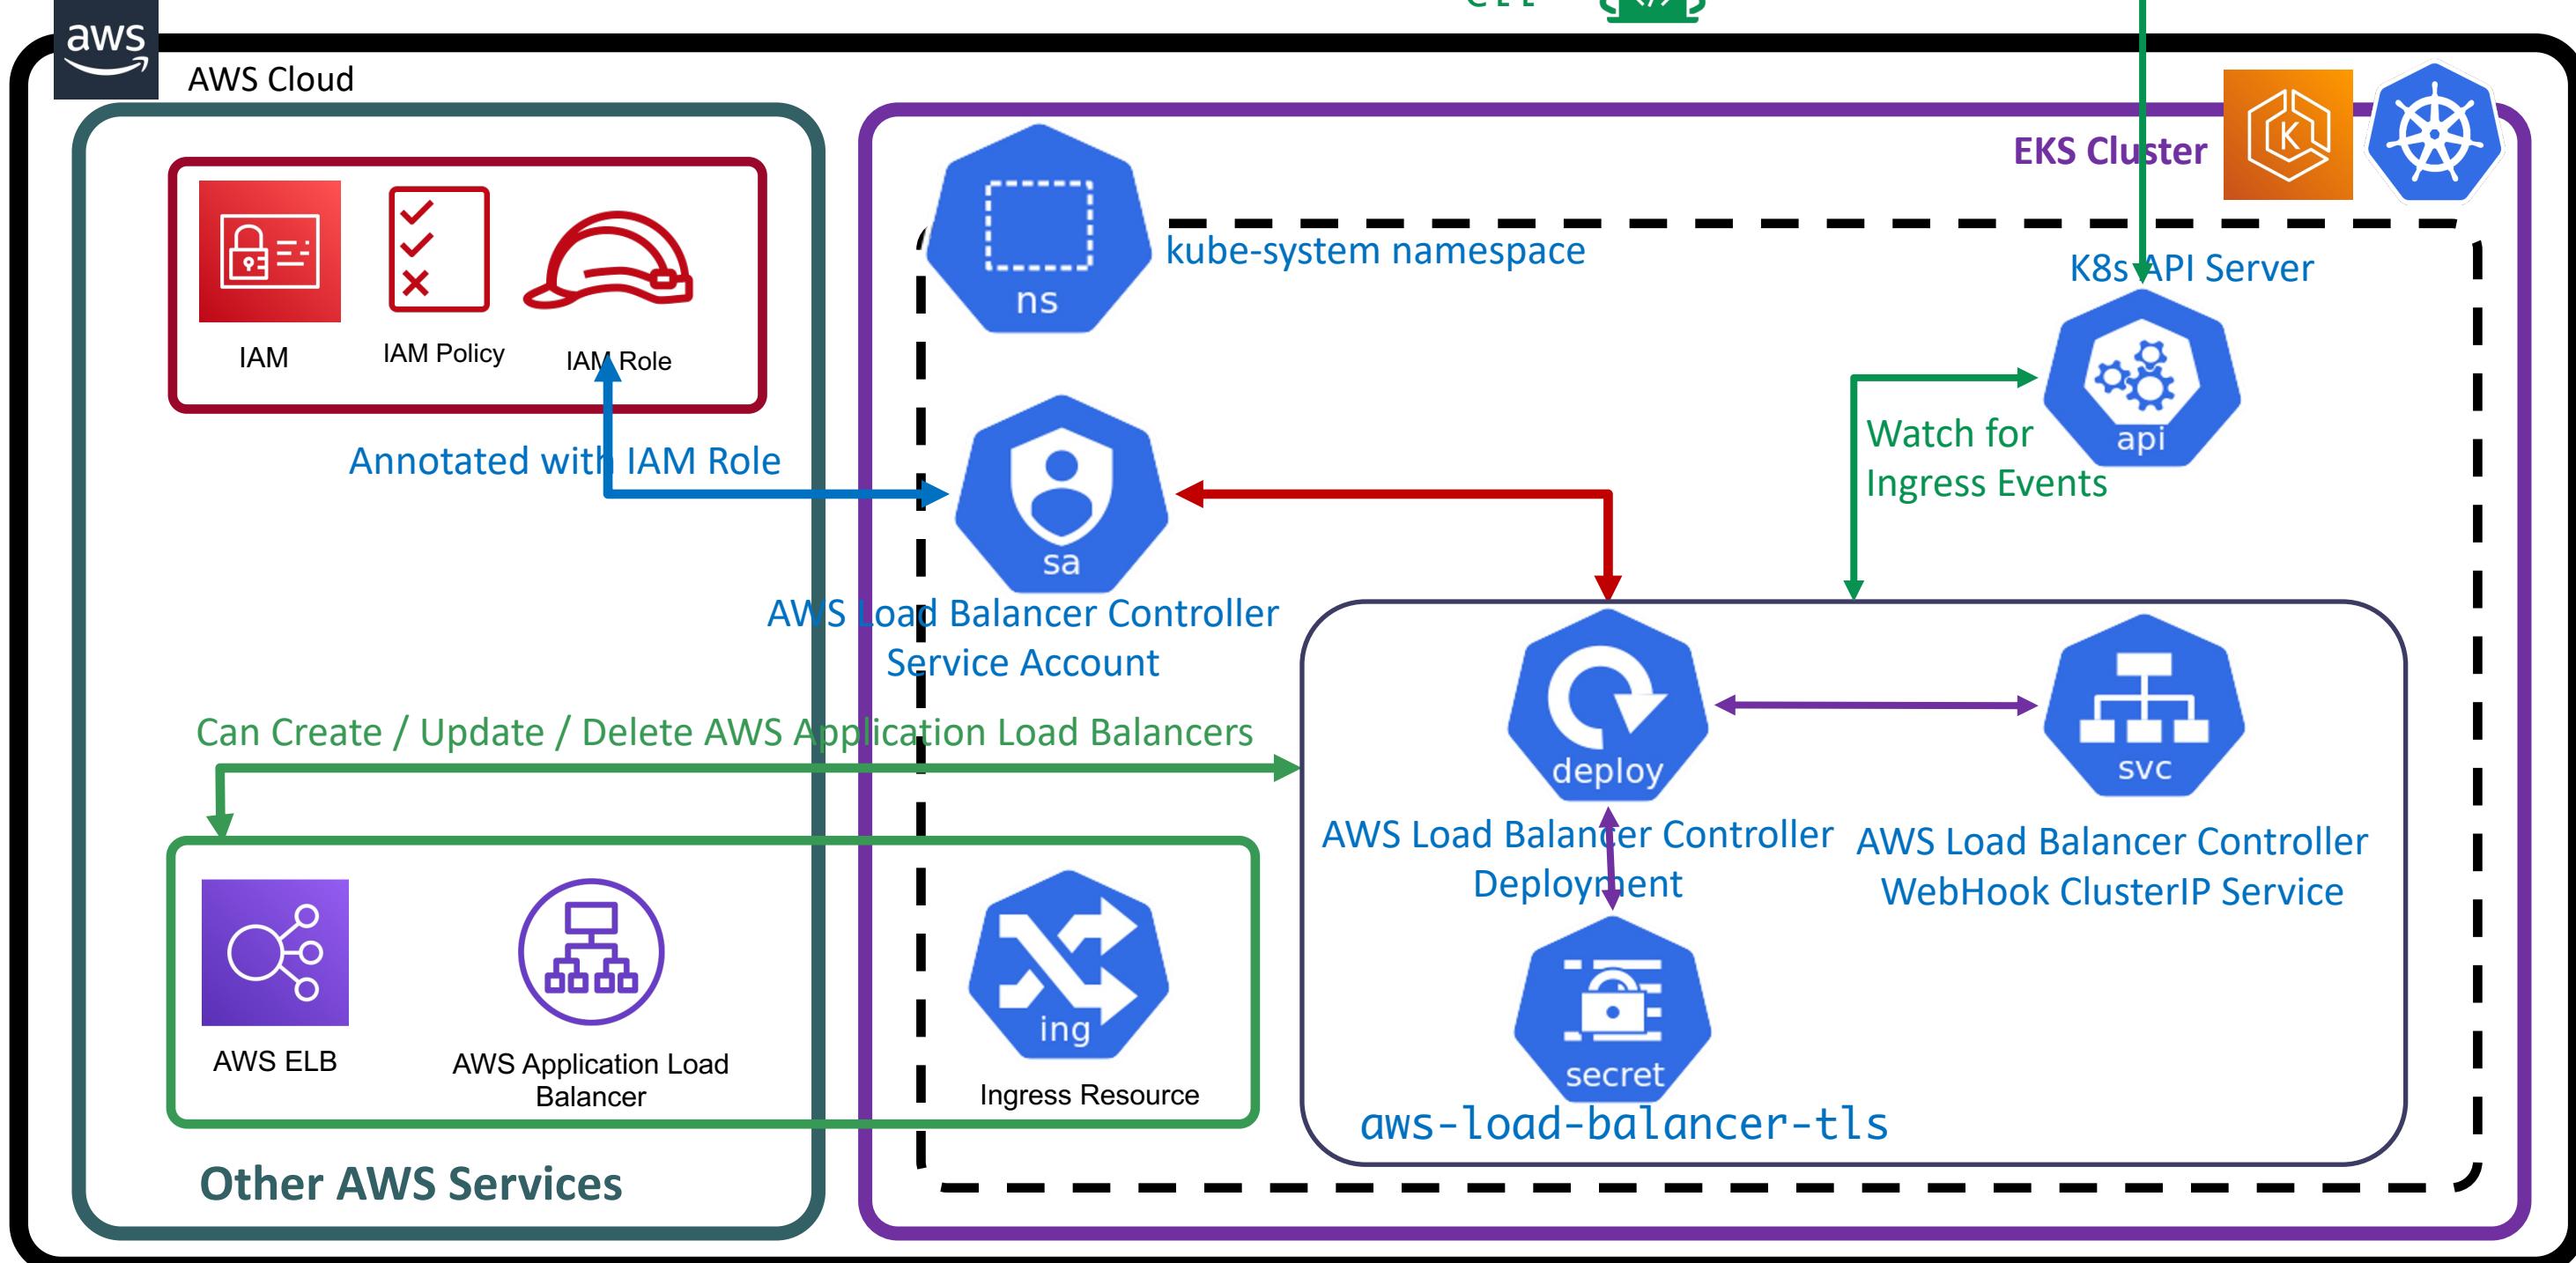
AWS Network Load Balancer

# AWS Load Balancer Controller

kubectl  
cli

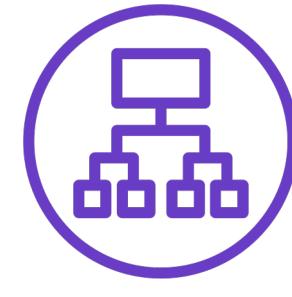
Deploy Ingress k8s  
manifest

STACKSIMPLIFY

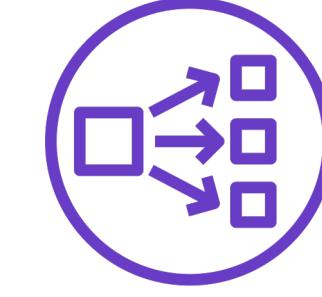




Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress



# AWS EKS

## Load Balancer Controller

### Kubernetes Ingress Class



# Kubernetes Ingress Class

AWS EKS ALB  
Ingress Controller

NGINX Ingress  
Controller

AKS Application  
Gateway Ingress  
Controller

If we have multiple Ingress Controllers running in our Kubernetes cluster, how to identify to which Ingress Controller our Ingress Resource/Service should be associated to ?

Kubernetes Object  
Kind: IngressClass

**More Ingress Controllers:** <https://kubernetes.io/docs/concepts/services-networking/ingress-controllers/>

# Kubernetes Ingress Class

```
apiVersion: networking.k8s.io/v1
kind: IngressClass
metadata:
  name: my-aws-ingress-class
  annotations:
    ingressclass.kubernetes.io/is-default-class: "true"
spec:
  controller: ingress.k8s.aws/alb
```

# Ingress Class

```
apiVersion: networking.k8s.io/v1
kind: IngressClass
metadata:
  name: my-aws-ingress-class
  #annotations:
    #ingressclass.kubernetes.io/is-default-class: "true"
spec:
  controller: ingress.k8s.aws/alb
```

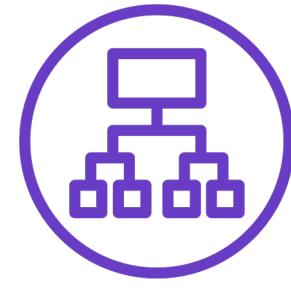
If **IngressClass** is not defined as **is-default-class:true** then in Ingress Resource we need to define the **spec.ingressClassName** in Ingress Resource

# Ingress Resource

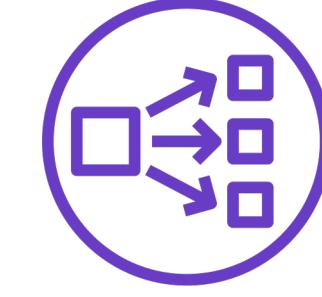
```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: ingress-nginxapp1
  labels:
    app: app1-nginx
  annotations:
    alb.ingress.kubernetes.io/load-balancer-name: app1ingress
    alb.ingress.kubernetes.io/scheme: internet-facing
spec:
  ingressClassName: my-aws-ingress-class # Ingress Class
  defaultBackend:
    service:
      name: app1-nginx-nodeport-service
      port:
        number: 80
```



Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress



# AWS EKS Load Balancer Controller



## Kubernetes Ingress Basics

# Ingress Manifest – Key Items

**Ingress Annotations**  
(Load Balancer Settings)

**Ingress Spec**  
**Ingress Class Name**  
(Which Ingress Controller to use)

**Ingress Spec**  
(Define Ingress Routing Rules, Default Backend)

**Annotations Reference:** <https://kubernetes-sigs.github.io/aws-load-balancer-controller/v2.3/guide/ingress/annotations/#annotations>

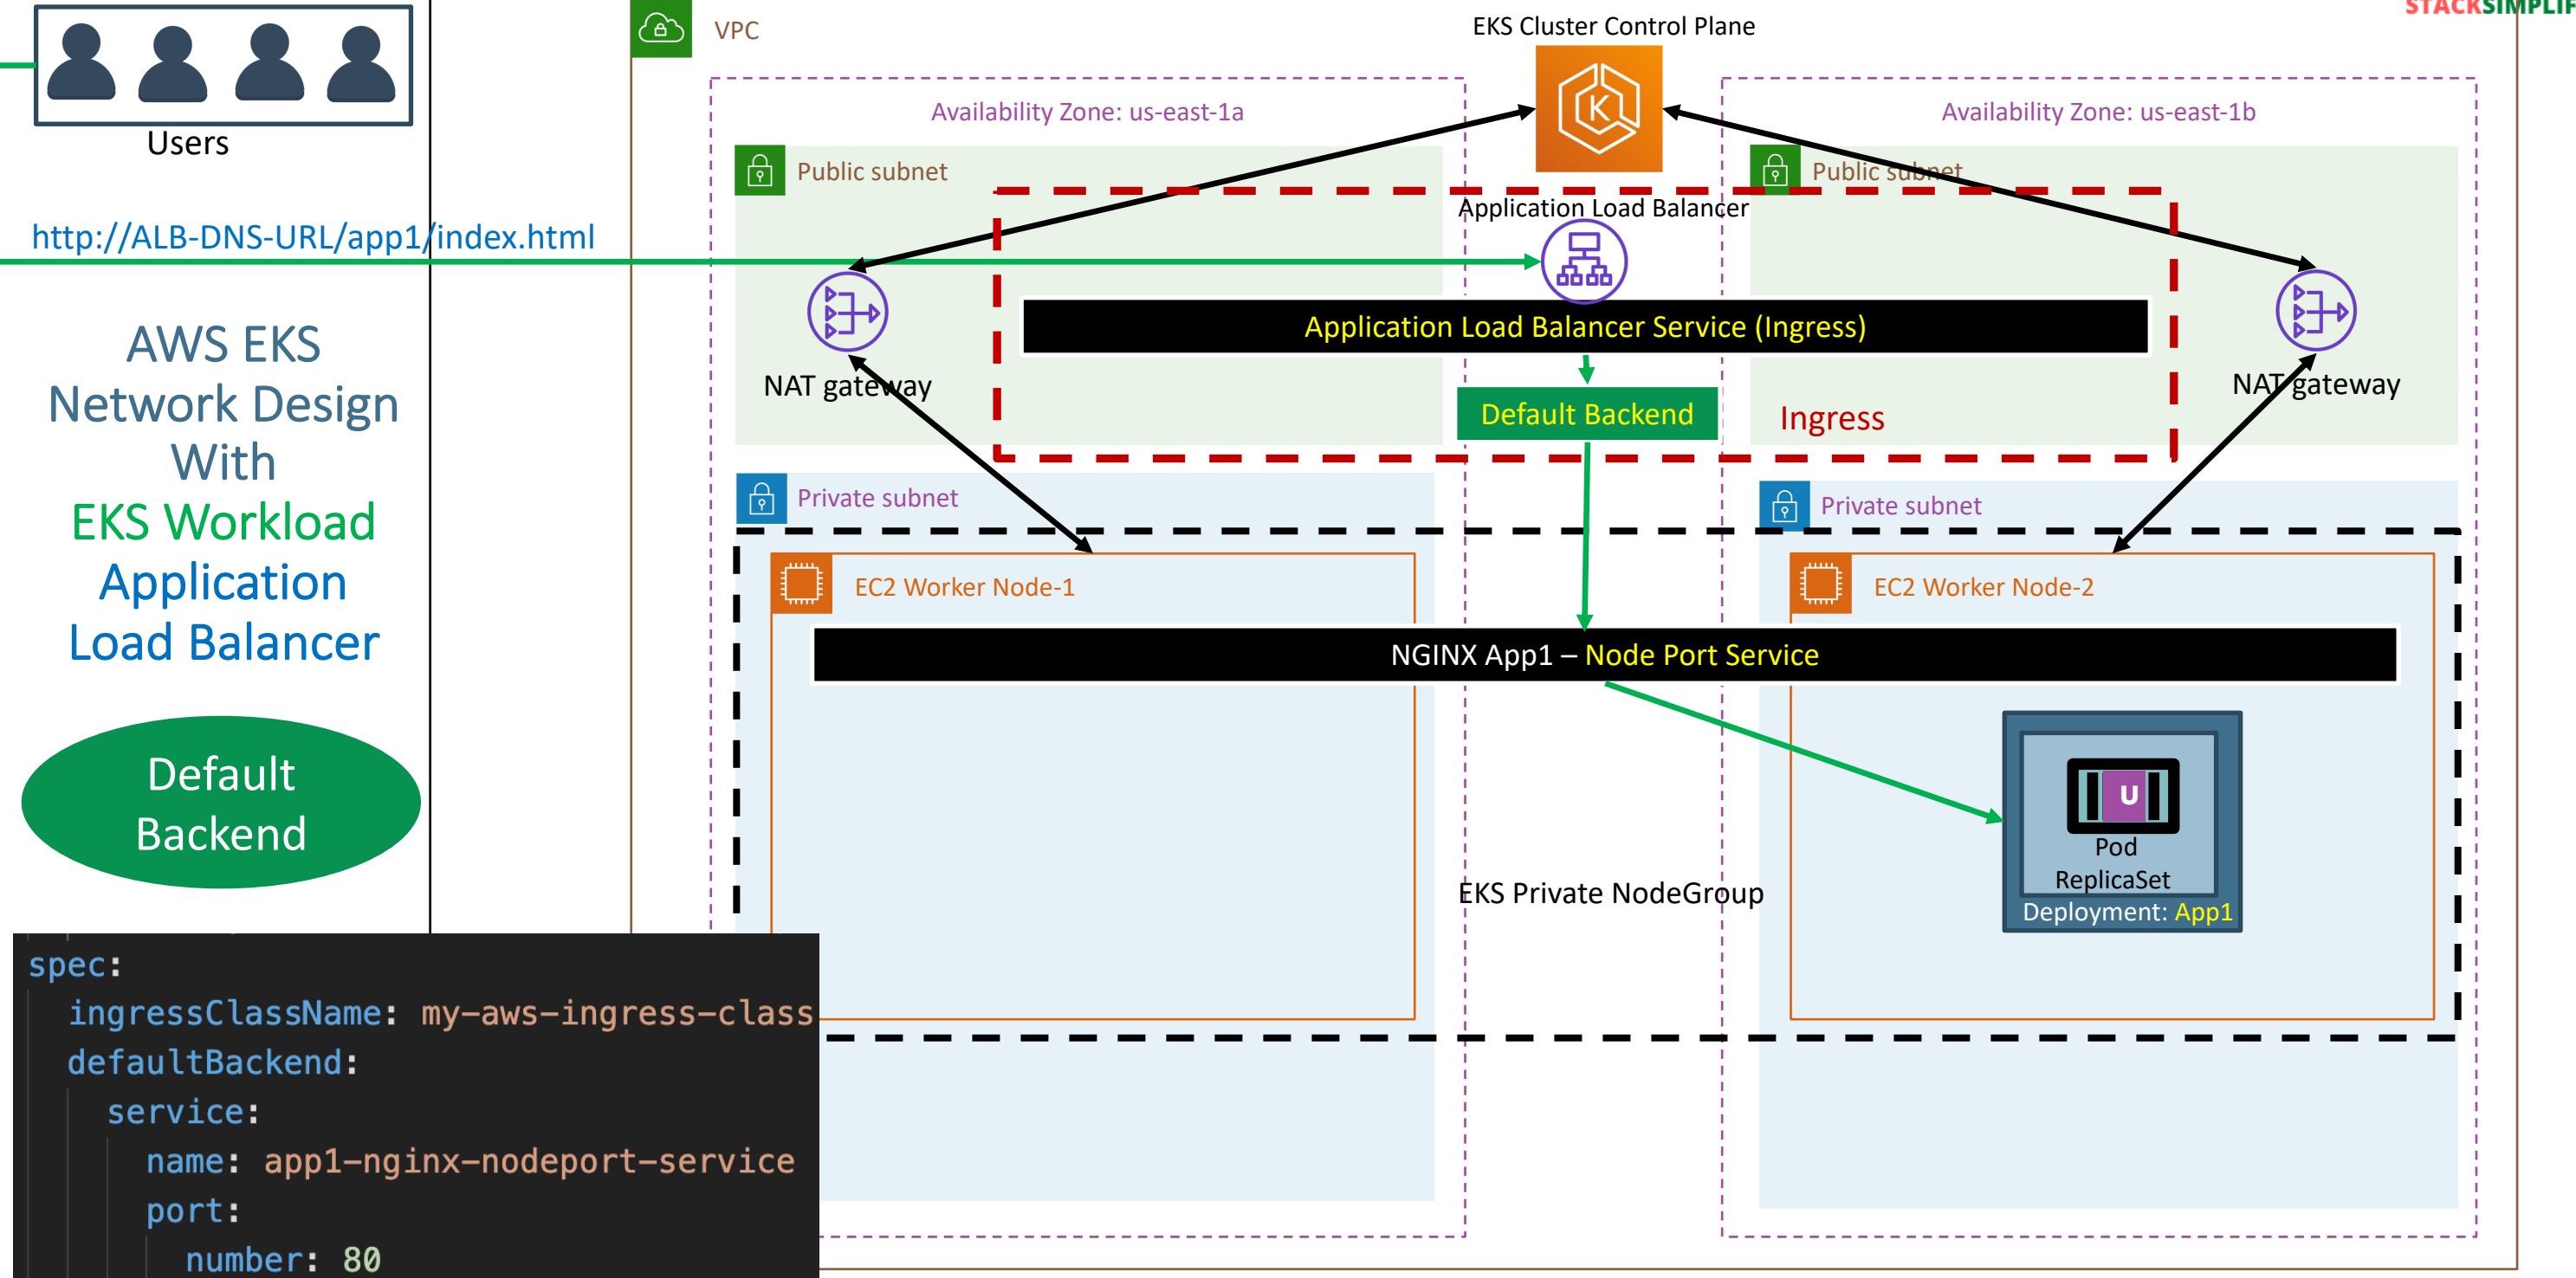
# Two Ingress Basic Demos

Demo-1

Demo-2

Ingress Service  
with  
Default Backend

Ingress Service  
with  
Ingress Rules





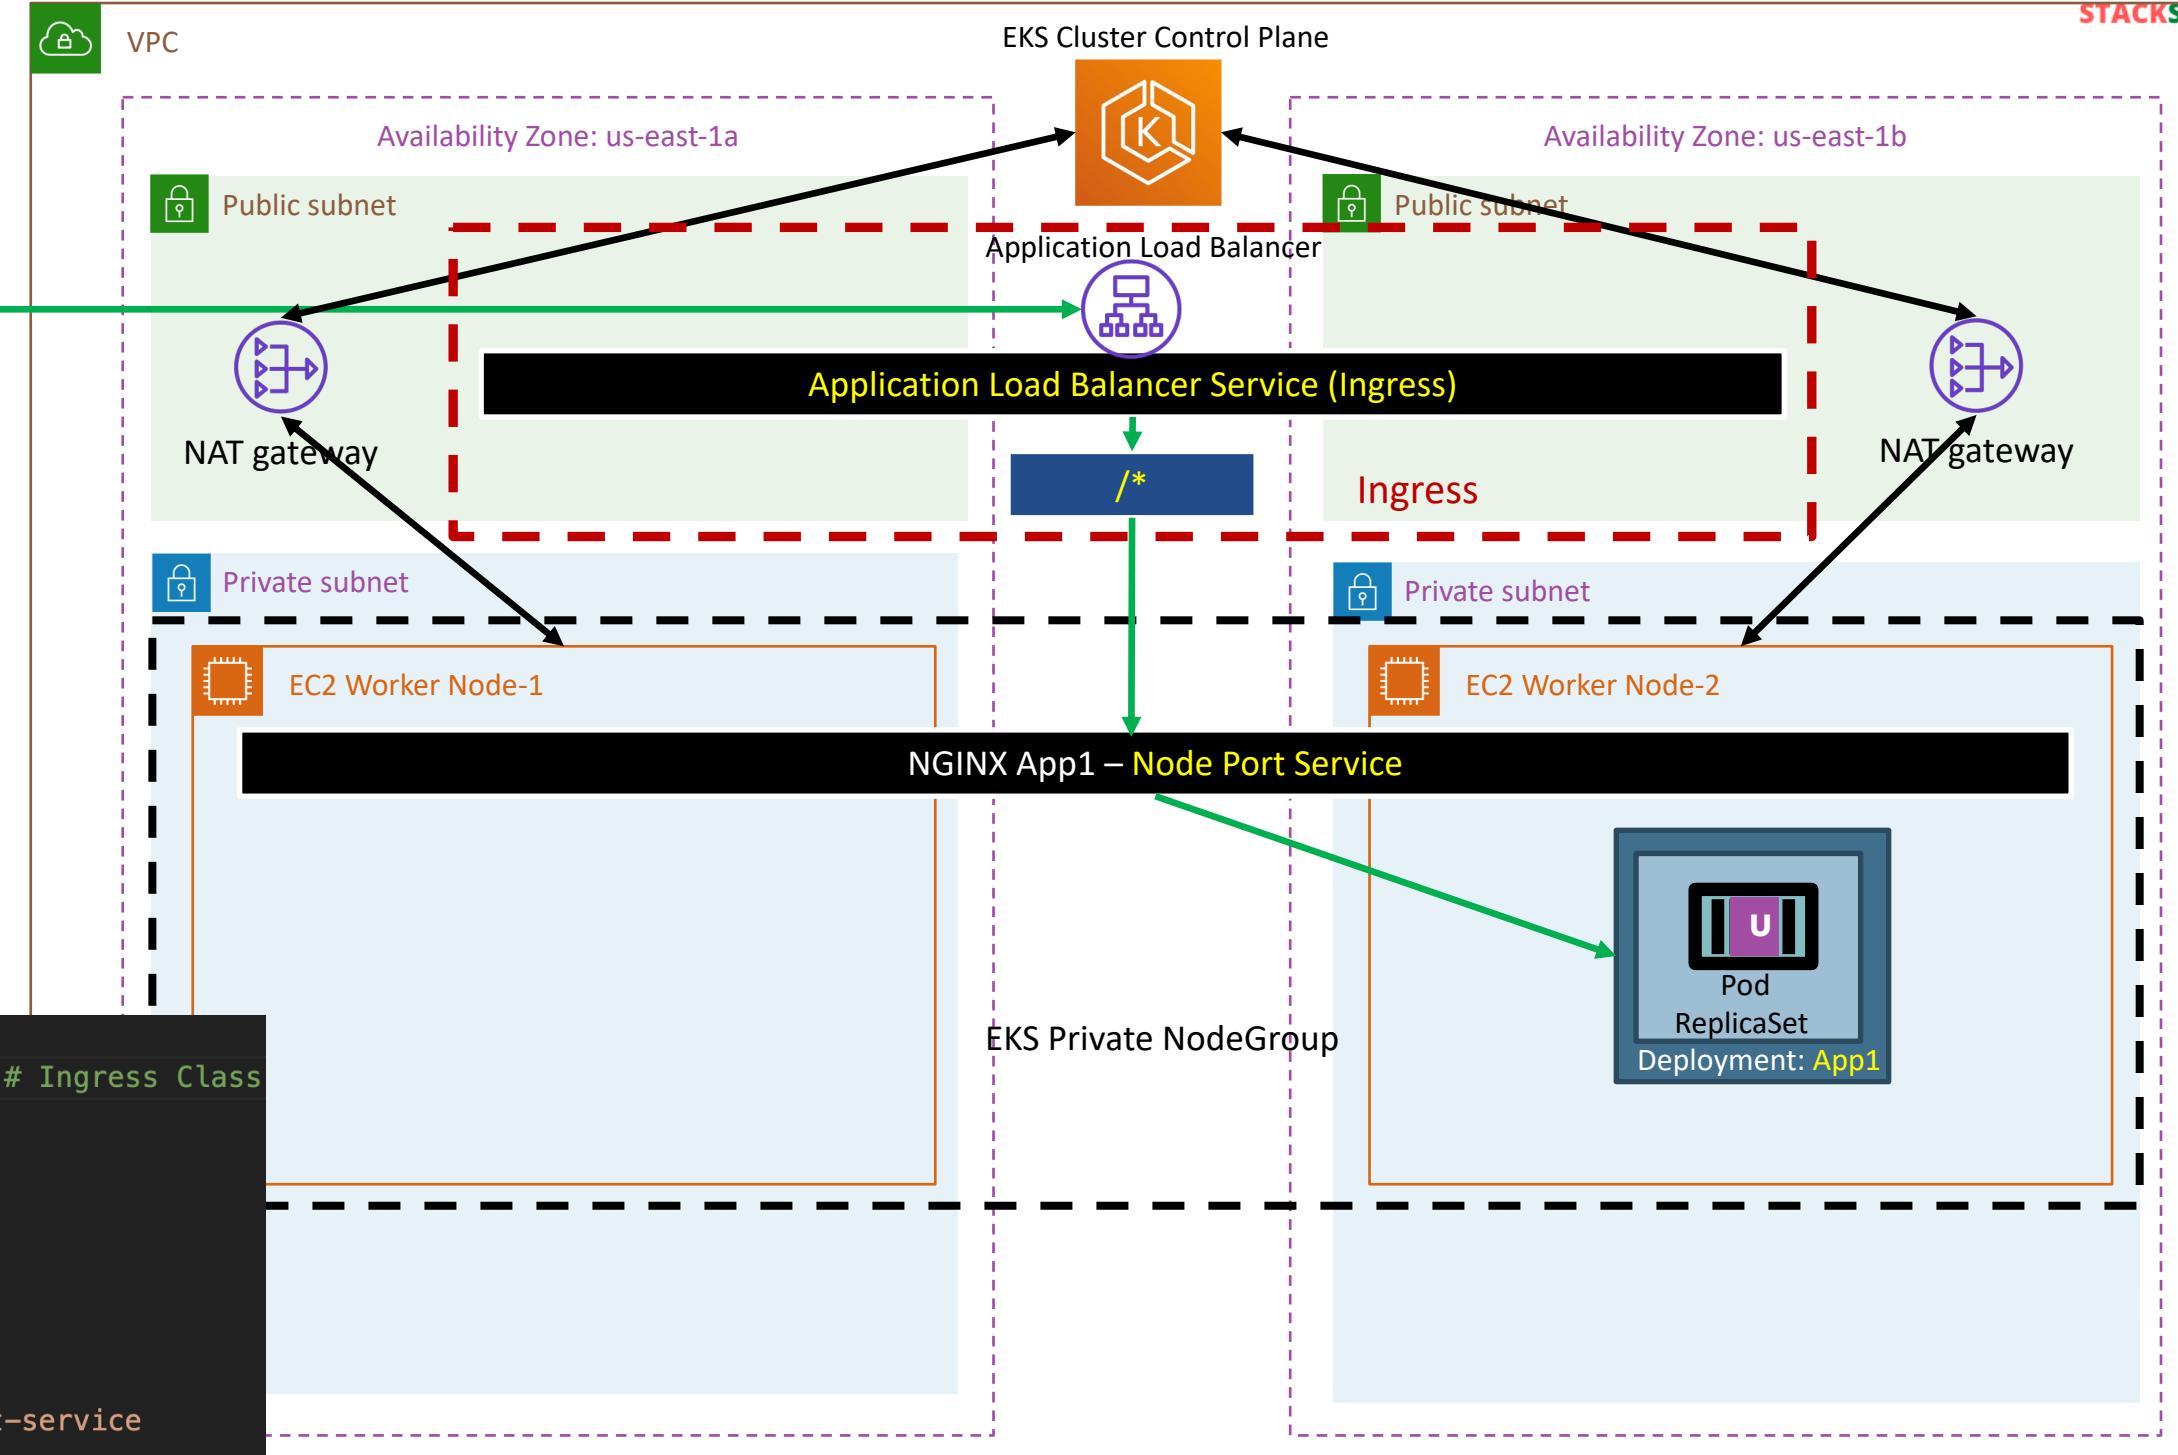
Users

http://ALB-DNS-URL/app1/index.html

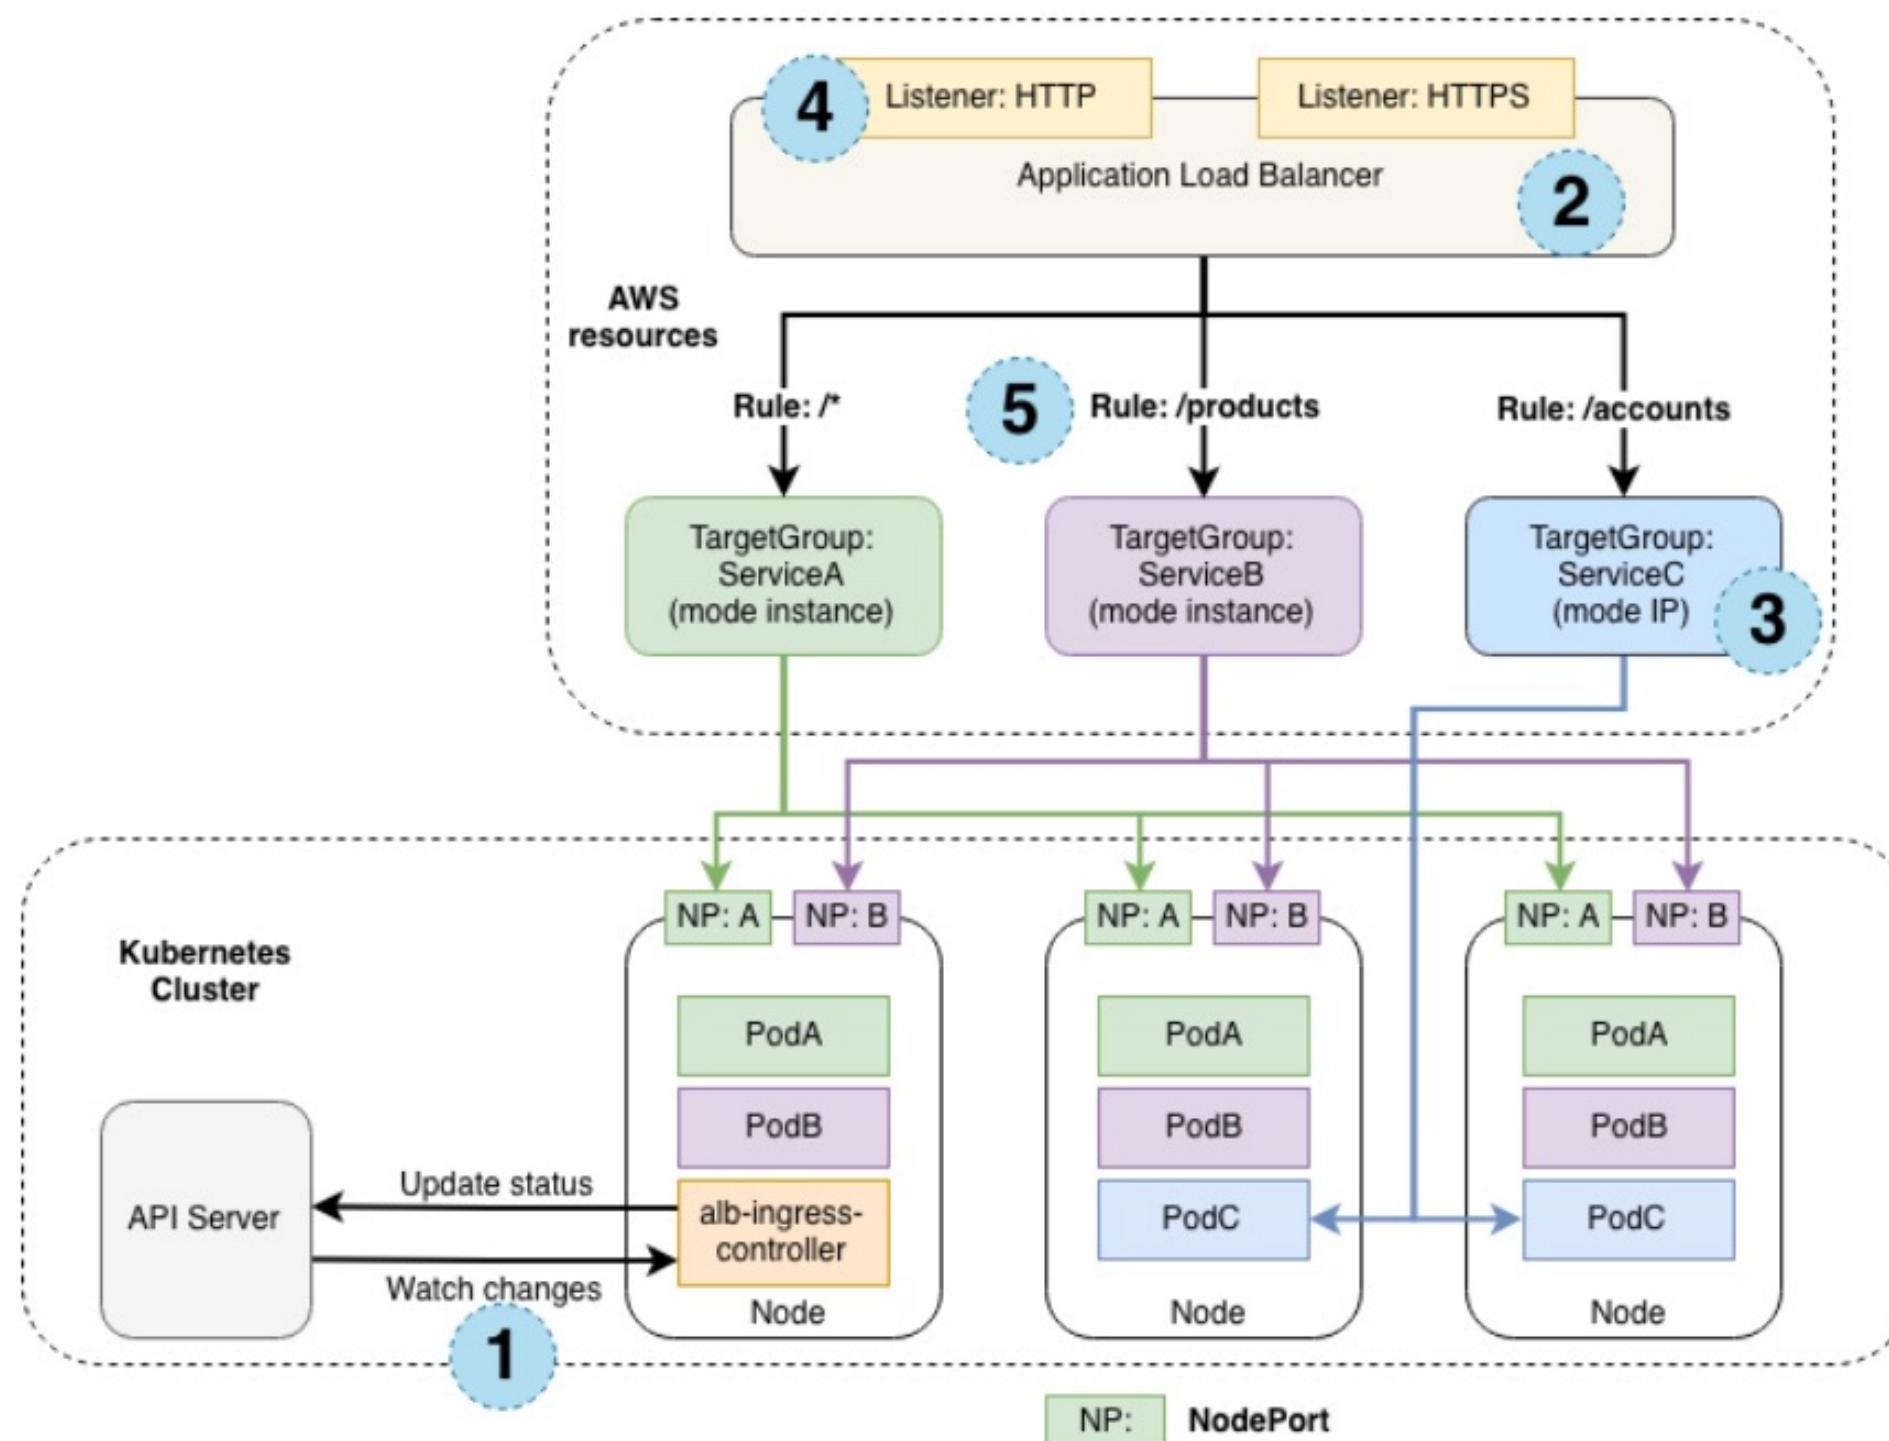
## AWS EKS Network Design With EKS Workload Application Load Balancer

### Ingress Rules

```
spec:  
  ingressClassName: my-aws-ingress-class # Ingress Class  
rules:  
  - http:  
    paths:  
      - path: /  
        pathType: Prefix  
        backend:  
          service:  
            name: app1-nginx-nodeport-service  
            port:  
              number: 80
```



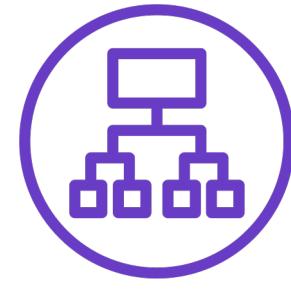
# How AWS Load Balancer controller works?



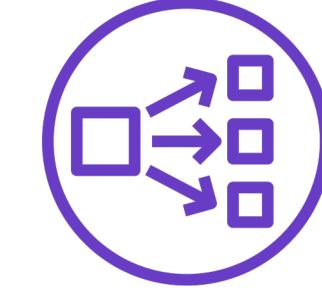
Reference: <https://kubernetes-sigs.github.io/aws-load-balancer-controller/v2.3/how-it-works/>



Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress



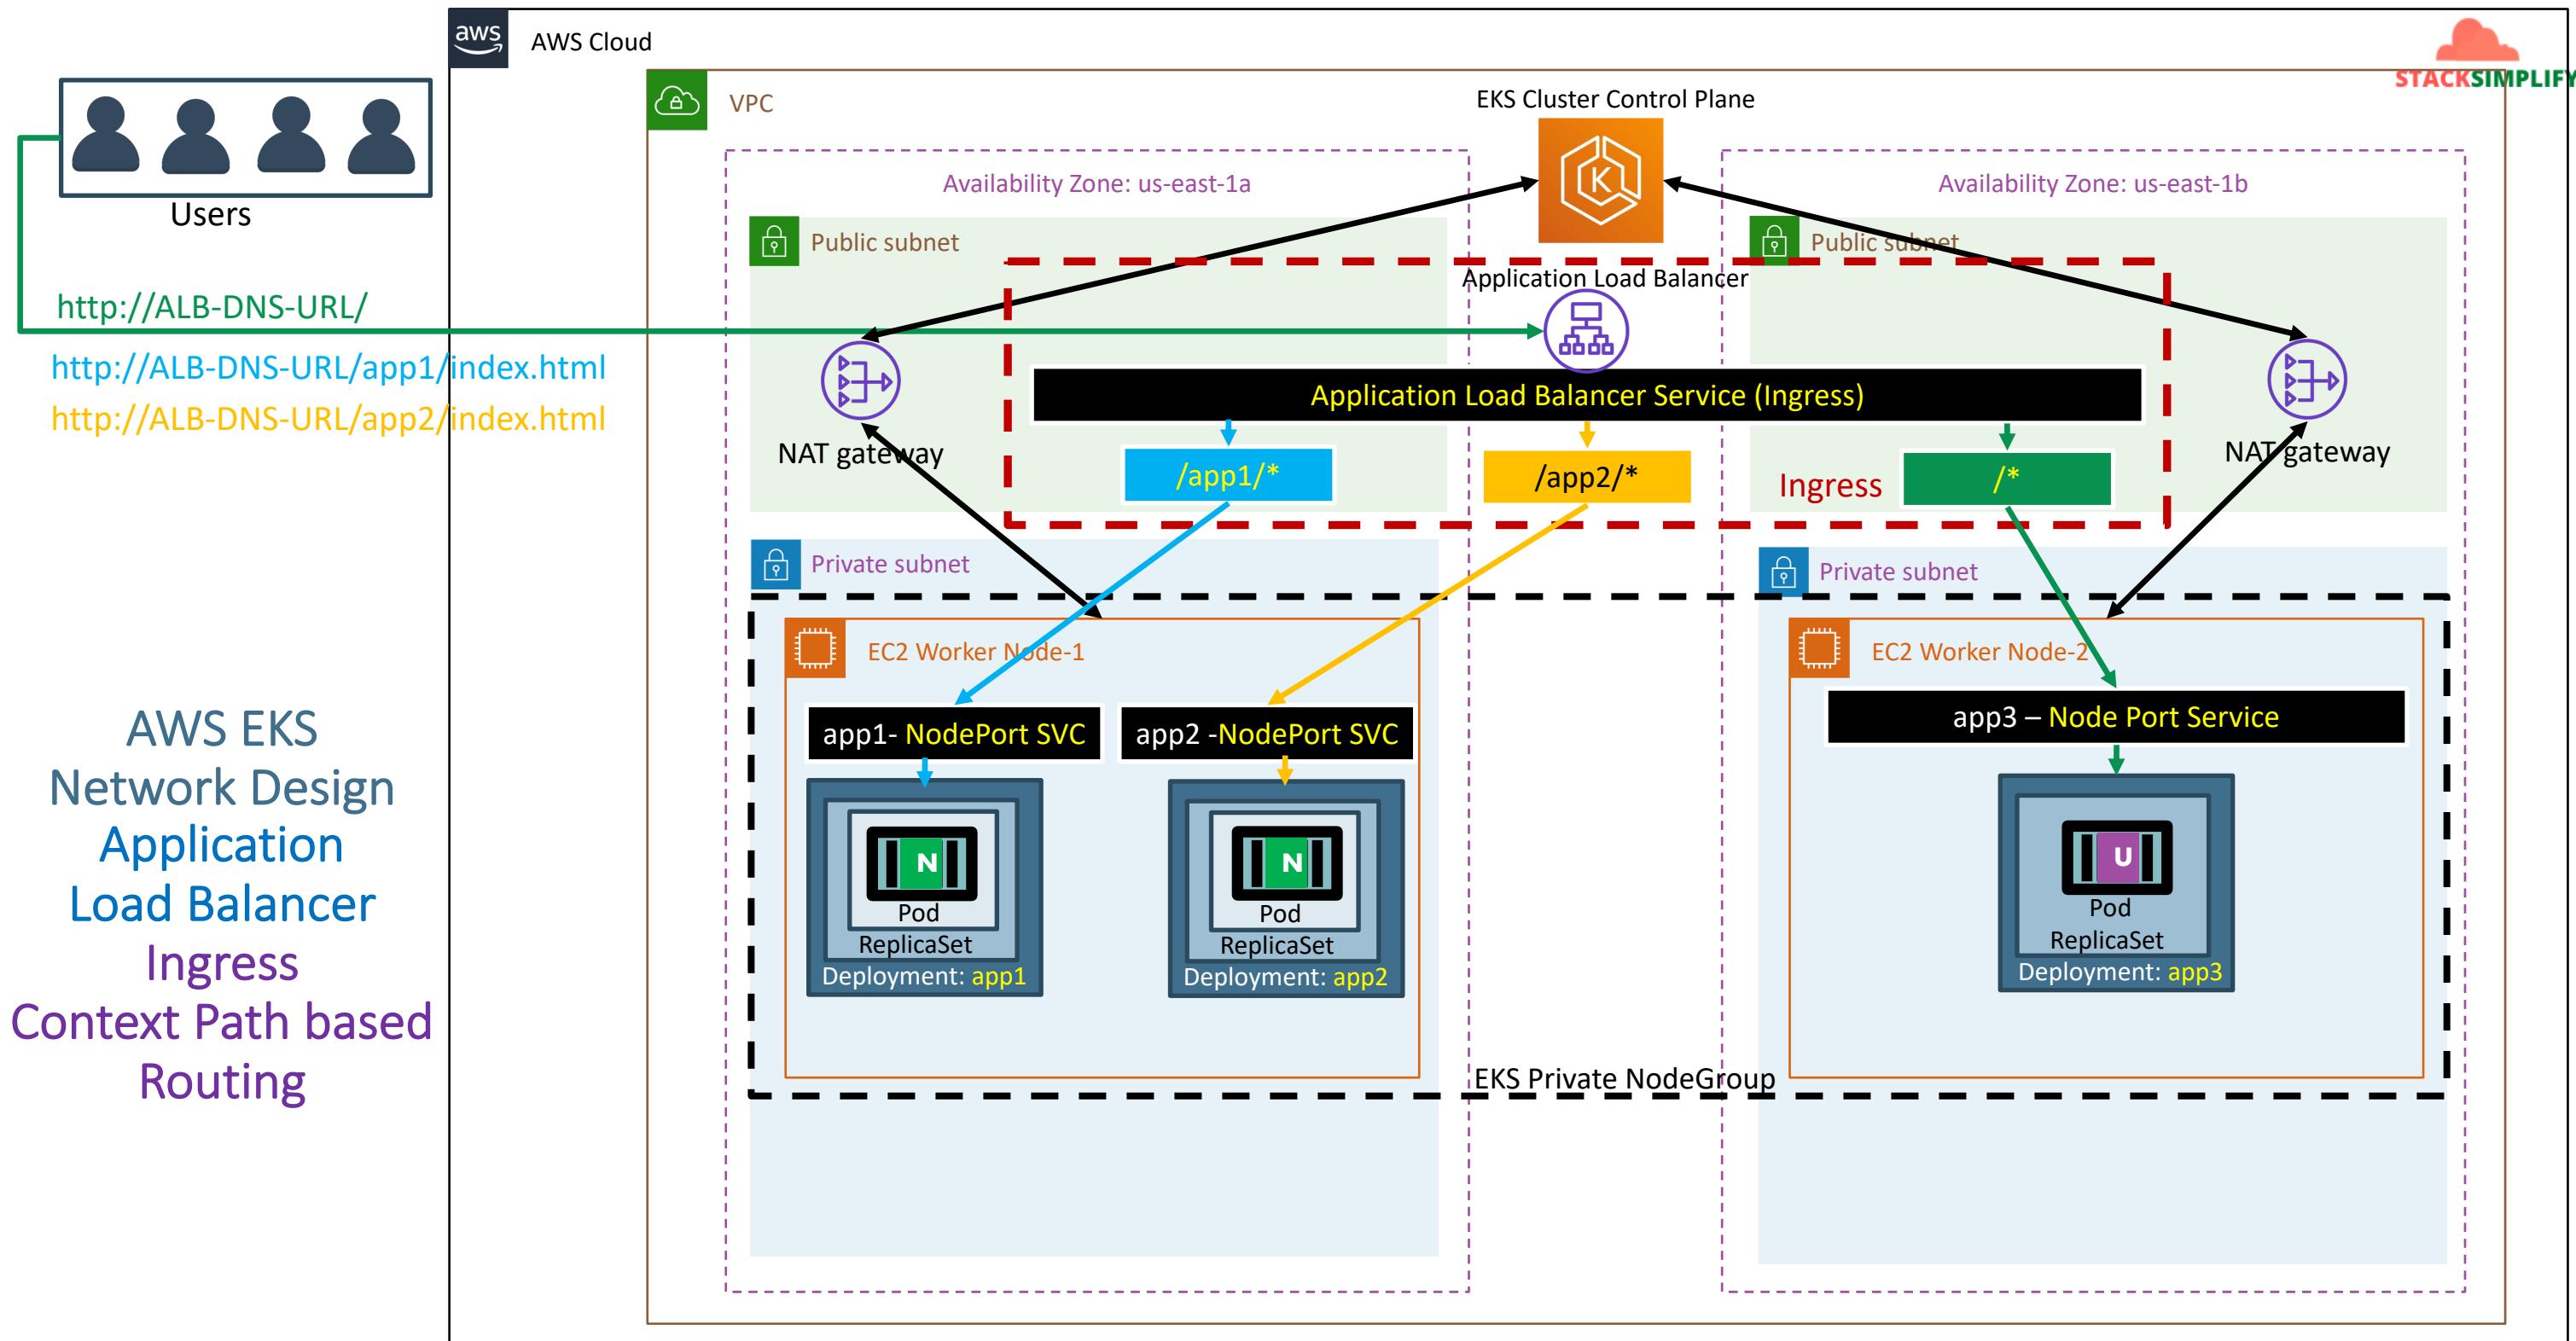
# AWS EKS

## Load Balancer Controller

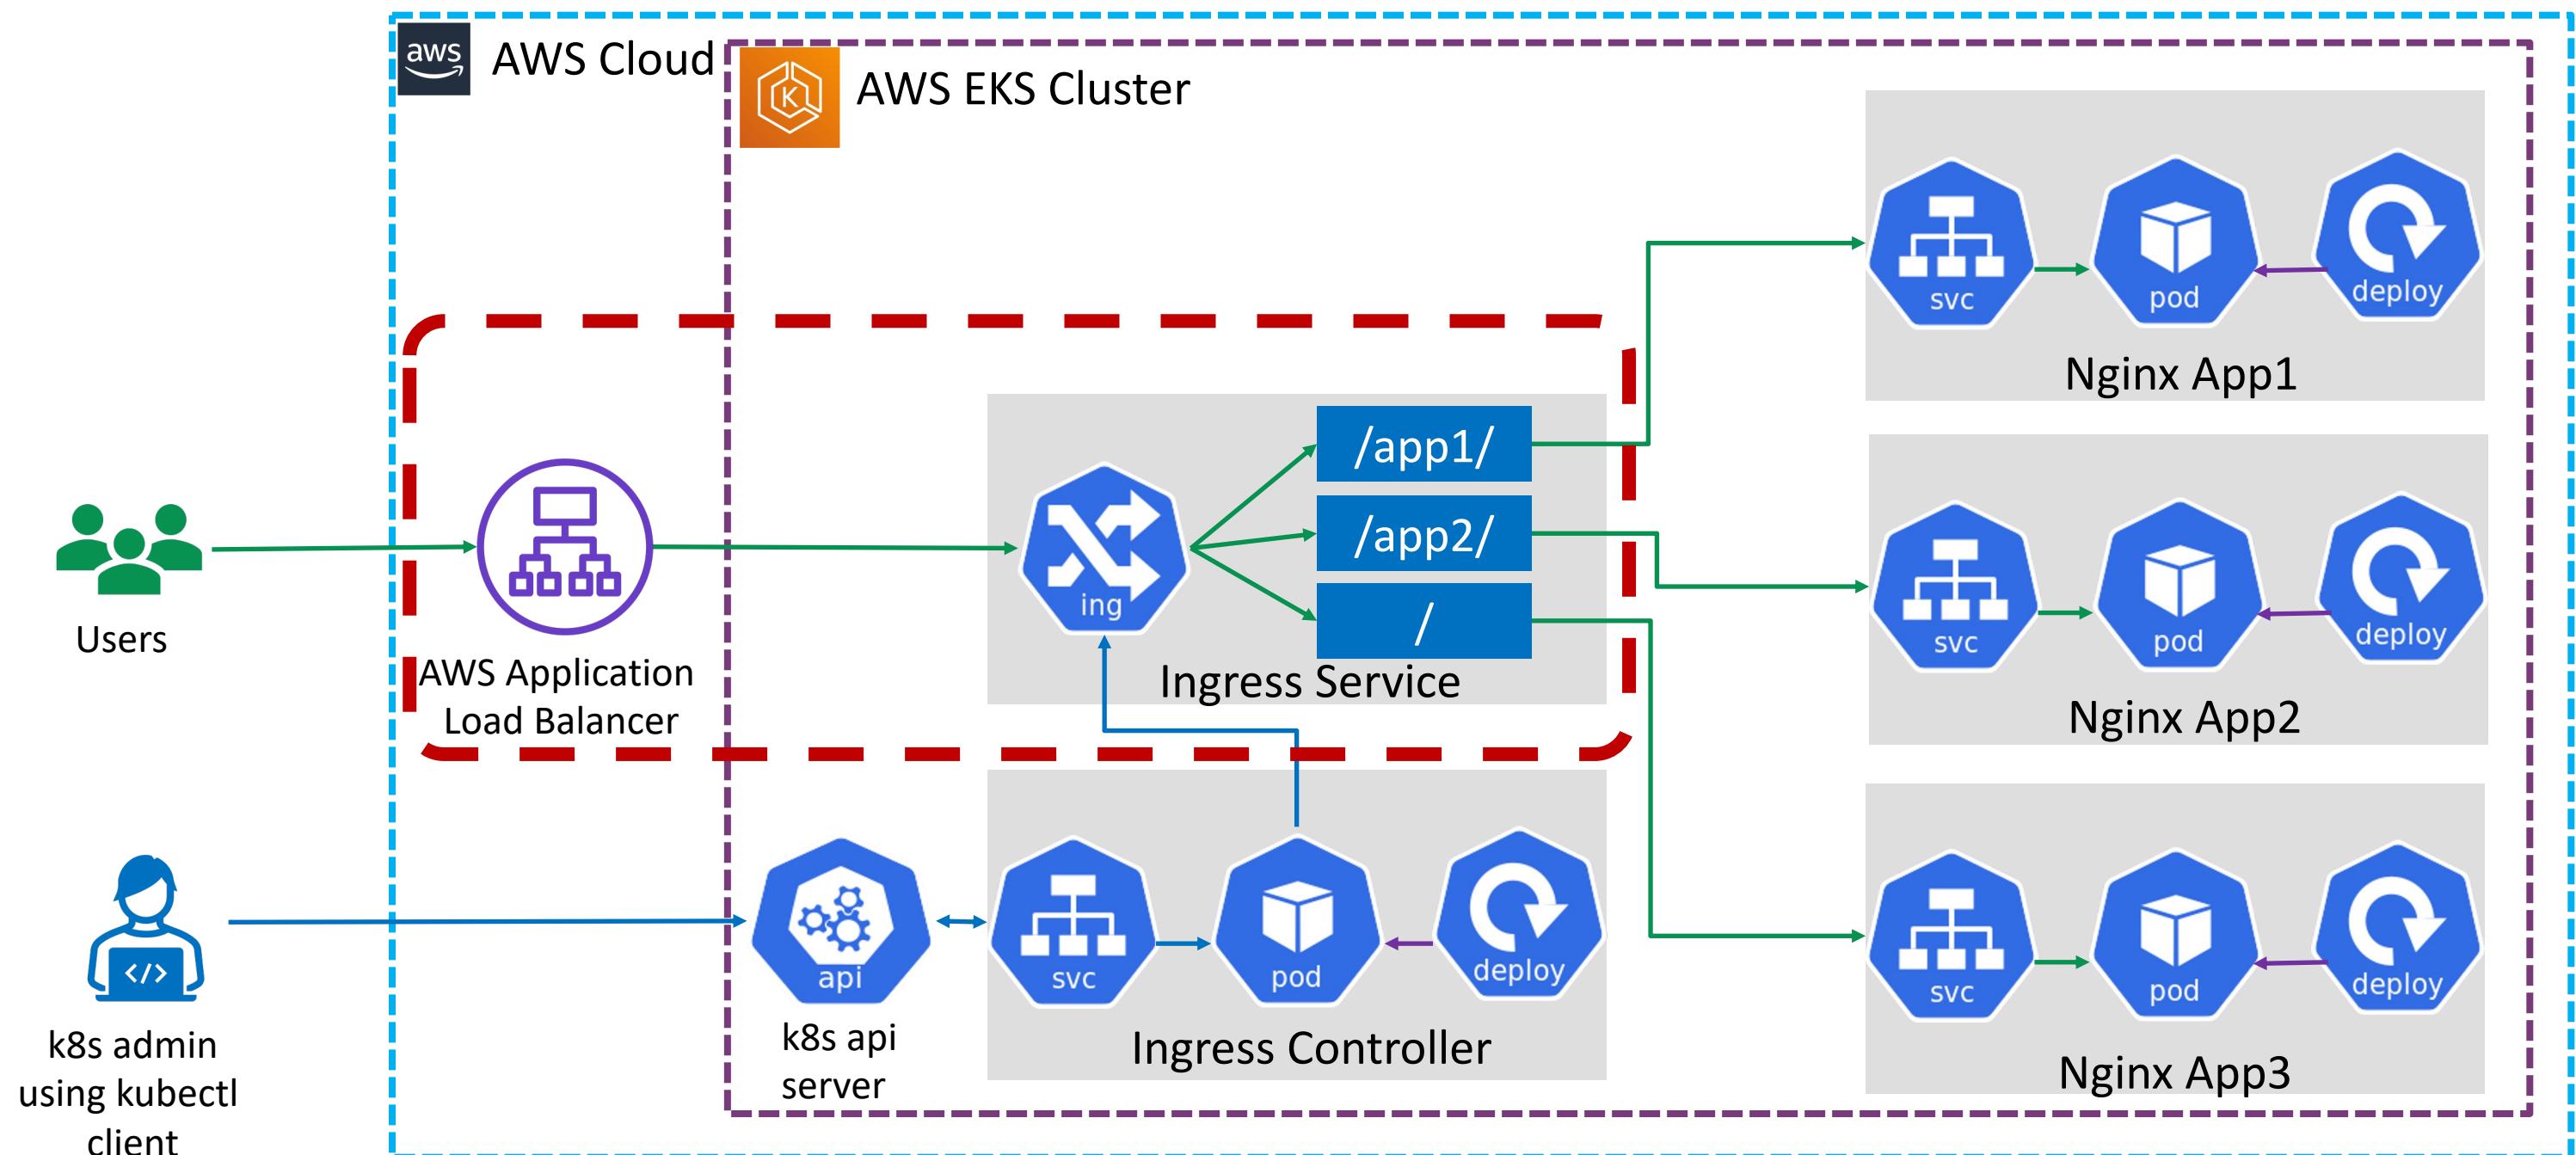
## Kubernetes Ingress

## Context Path based Routing



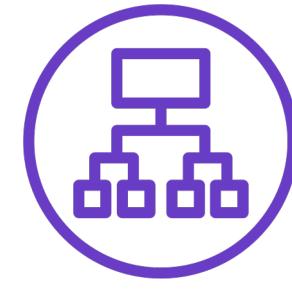


# AWS EKS ALB Ingress – Context Path Based Routing

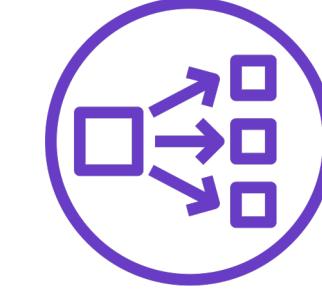




Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress



# AWS EKS

## Load Balancer Controller

### Kubernetes Ingress

### SSL



# ALB Ingress SSL

**Task-1:** Register AWS Route53 DNS Domain

**Task-2:** Create SSL Certificate in AWS Certificate Manager

**Task-3:** Update SSL Annotations in Ingress Service

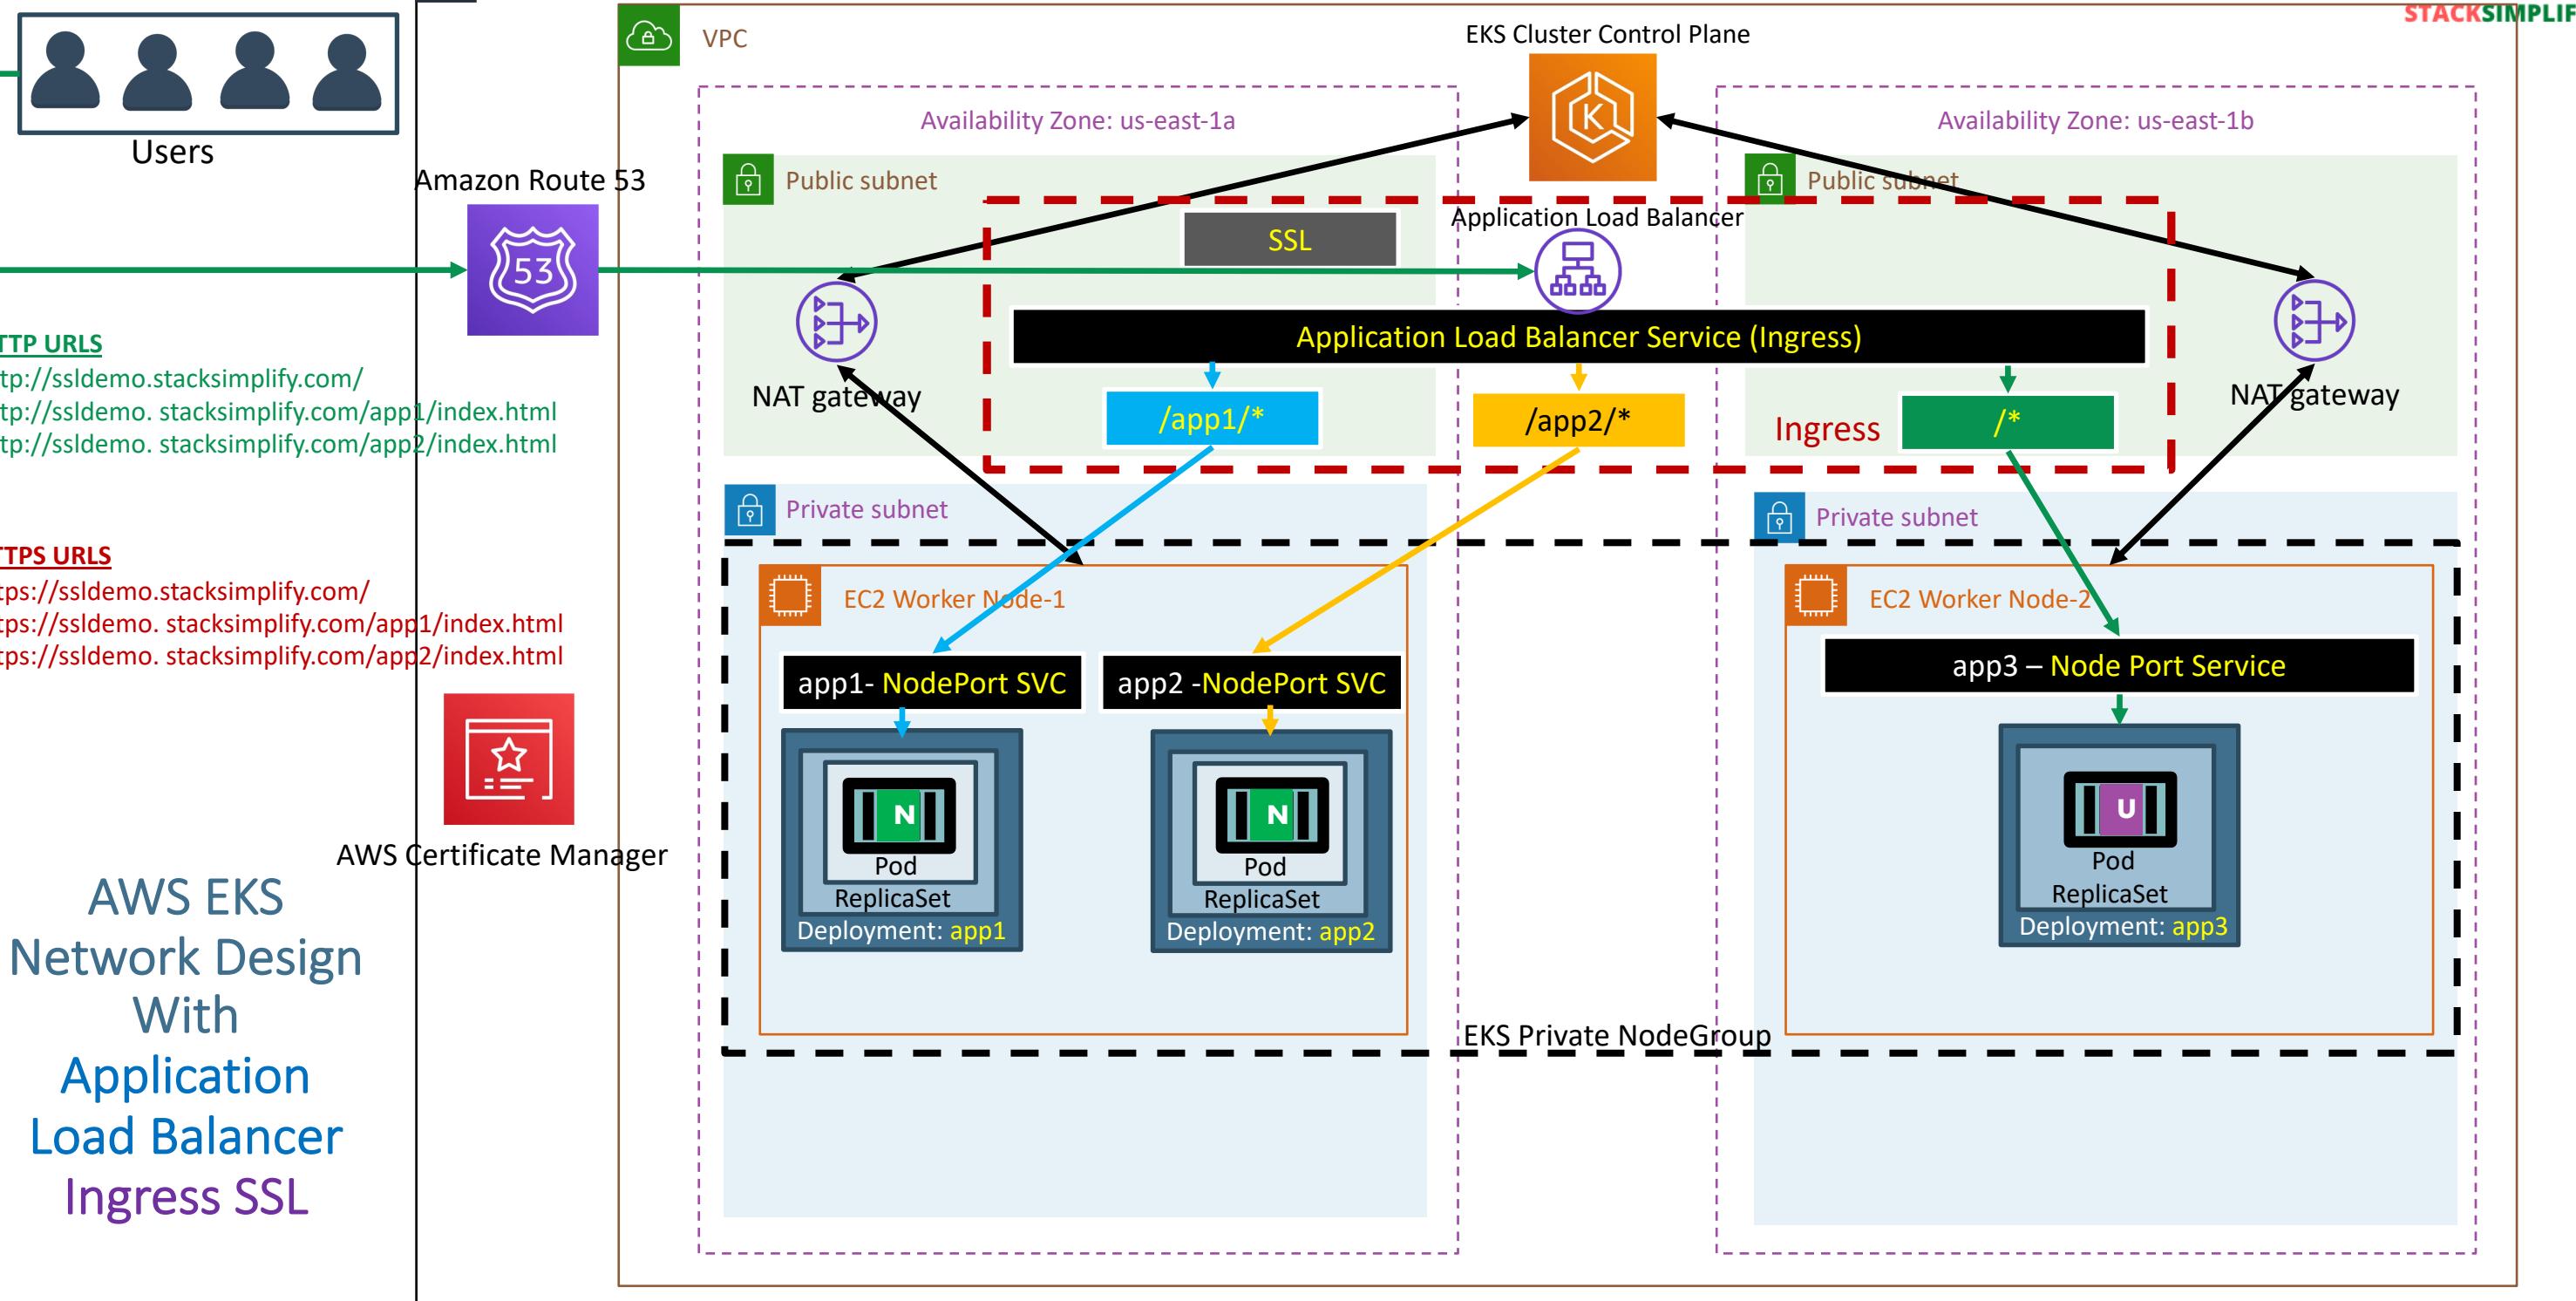
**Task-4:** Deploy k8s manifests and Test



AWS  
Route 53



AWS  
Certificate Manager



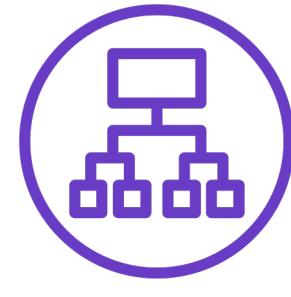
# Register a Domain using AWS Route53

**stacksimplify.com** is already a registered domain used for our websites

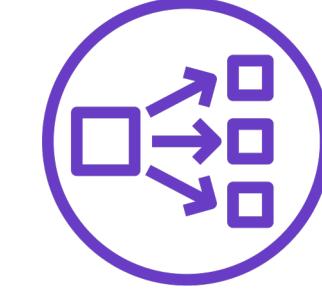
We are going to use Domain named **kubeoncloud.com** to register a new domain using AWS Route53



Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



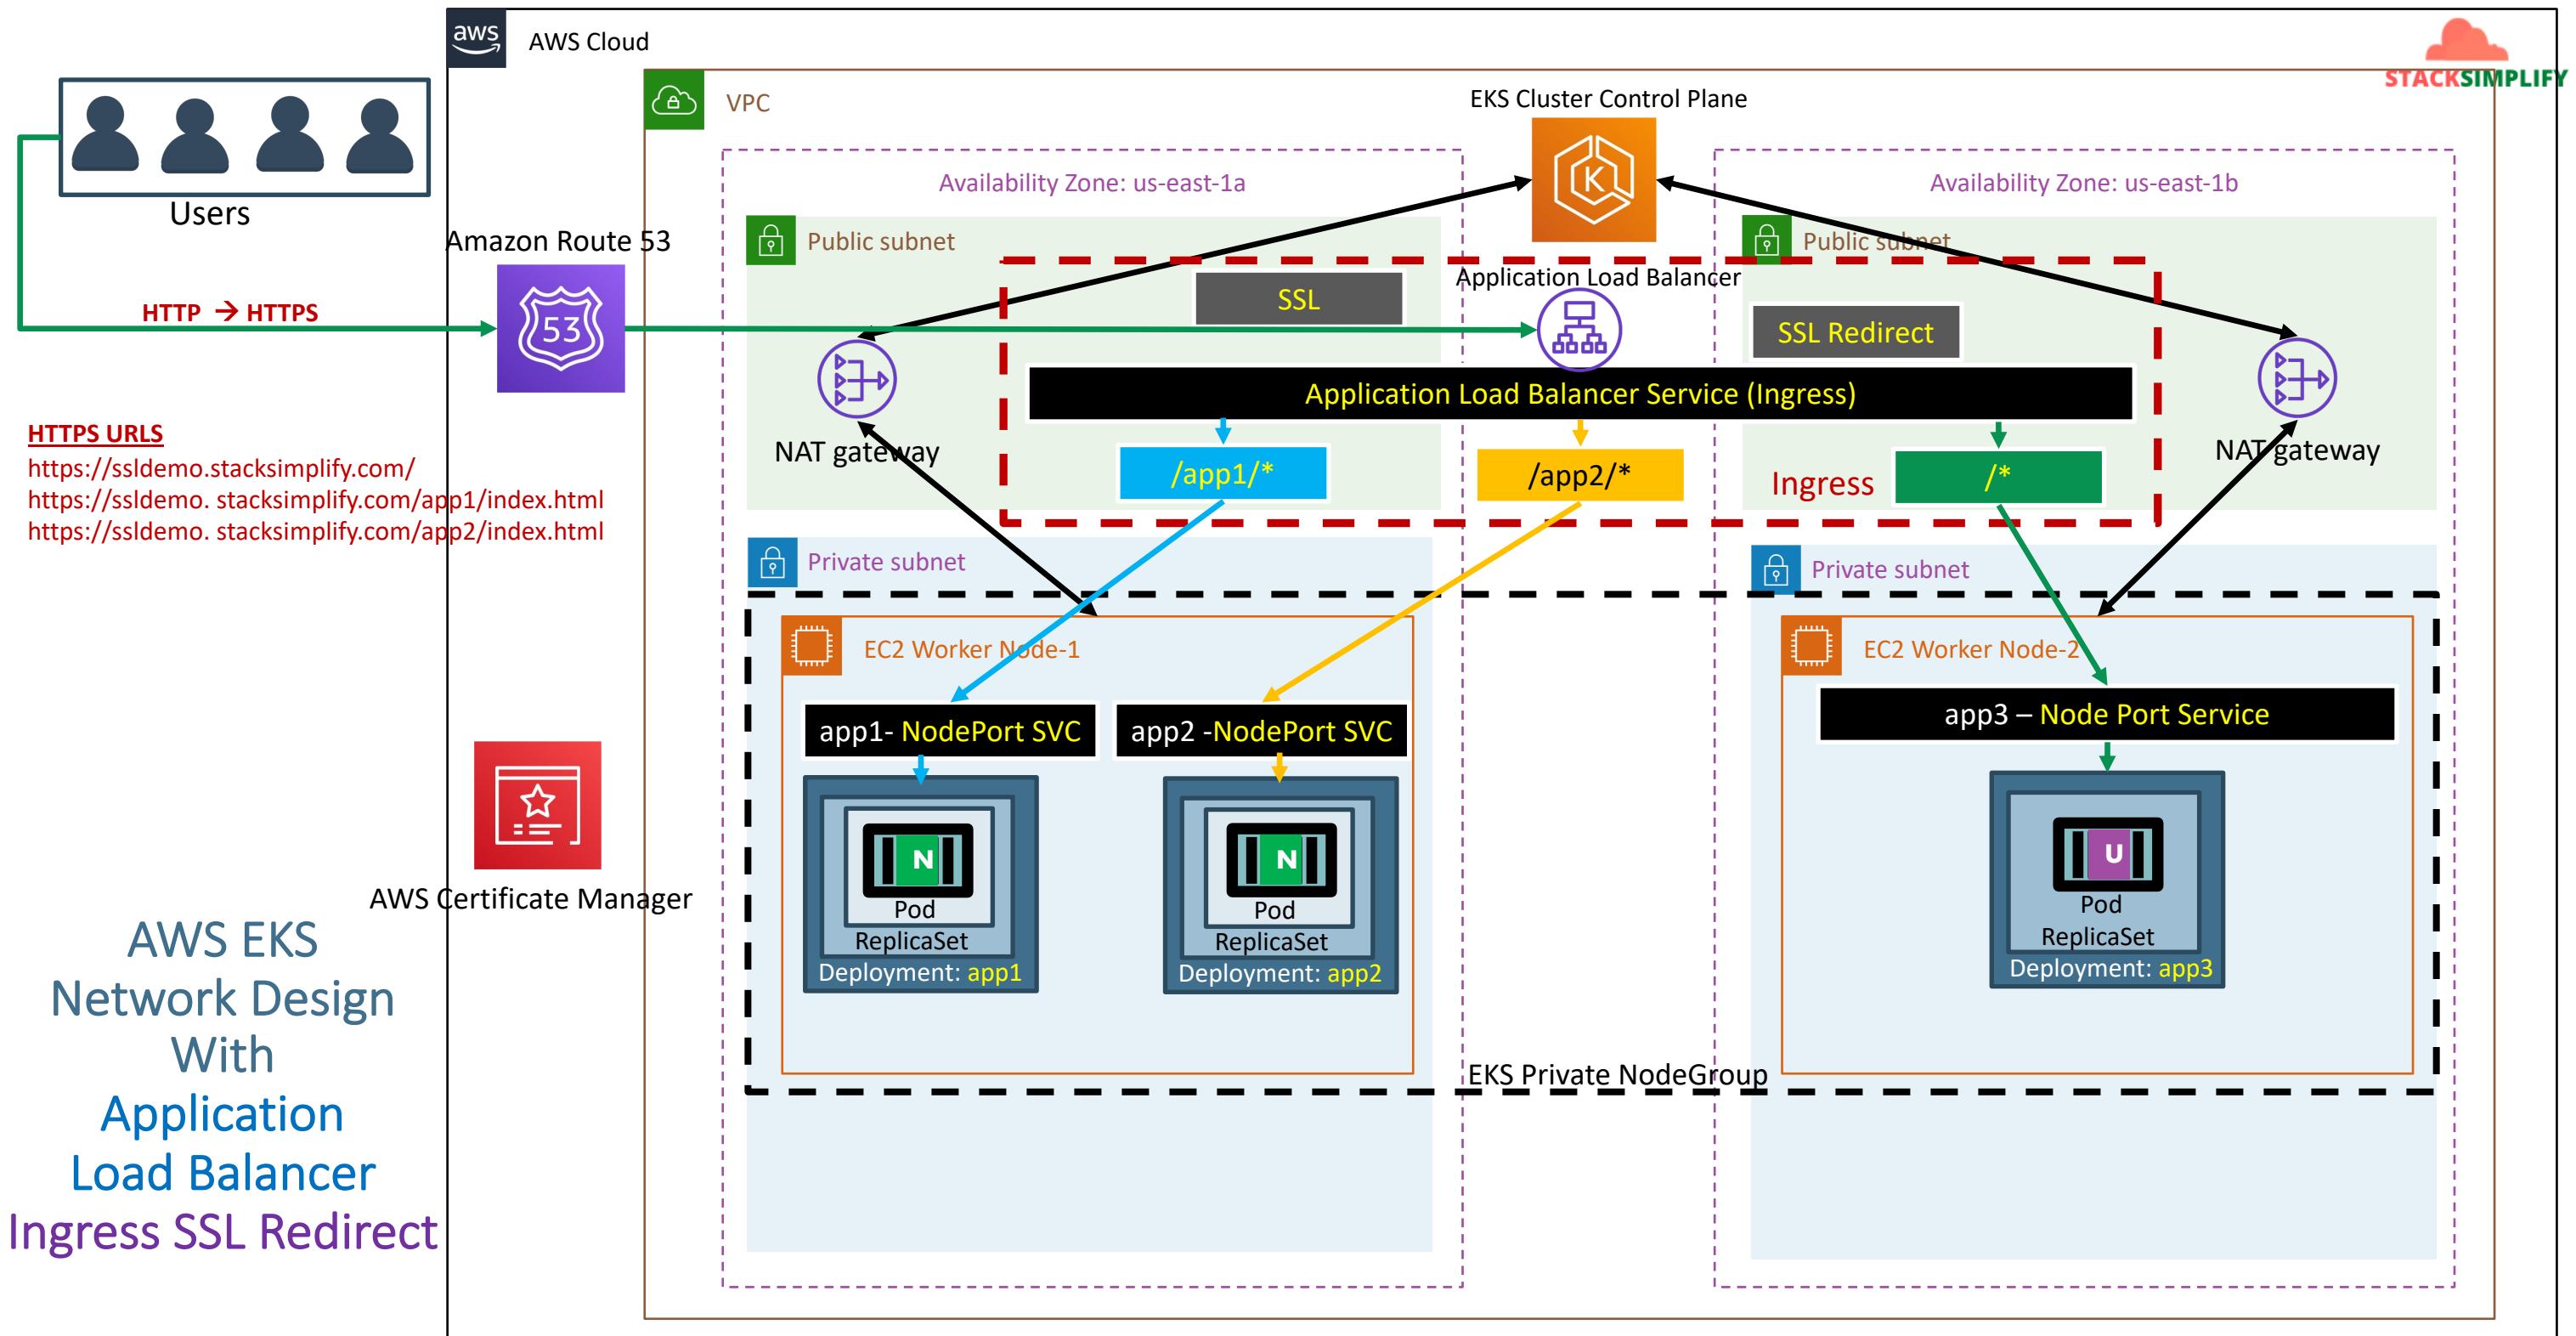
Kubernetes  
Ingress



# AWS EKS Load Balancer Controller

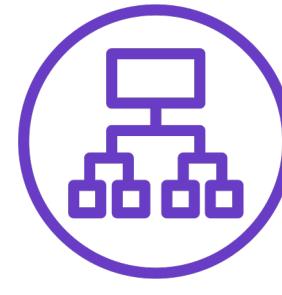
## Kubernetes Ingress SSL Redirect



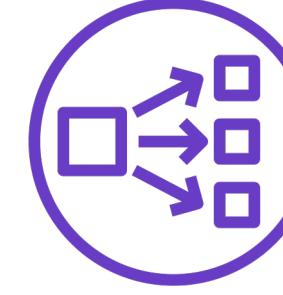




Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress



External  
DNS



# AWS EKS

## Load Balancer Controller

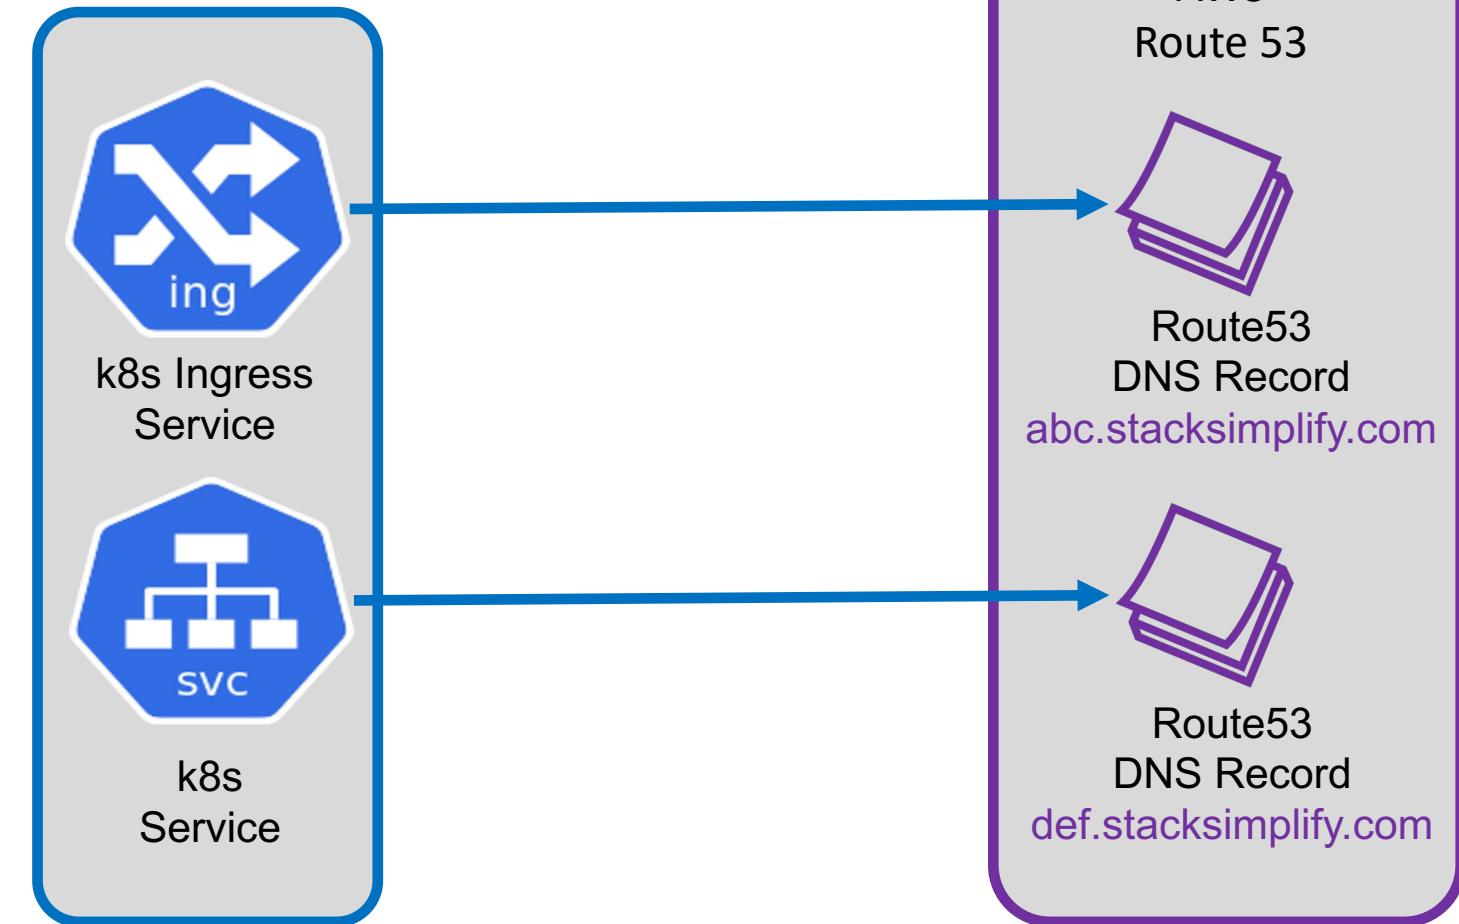
### External DNS Install



# Kubernetes External DNS

In previous SSL Demos, we added AWS Route53 DNS Record **manually** (`ssldemo101.stack simplify.com`)

With External DNS we can **automatically** add it for a Kubernetes **Ingress Service** or Kubernetes **Service** by defining it as Annotation



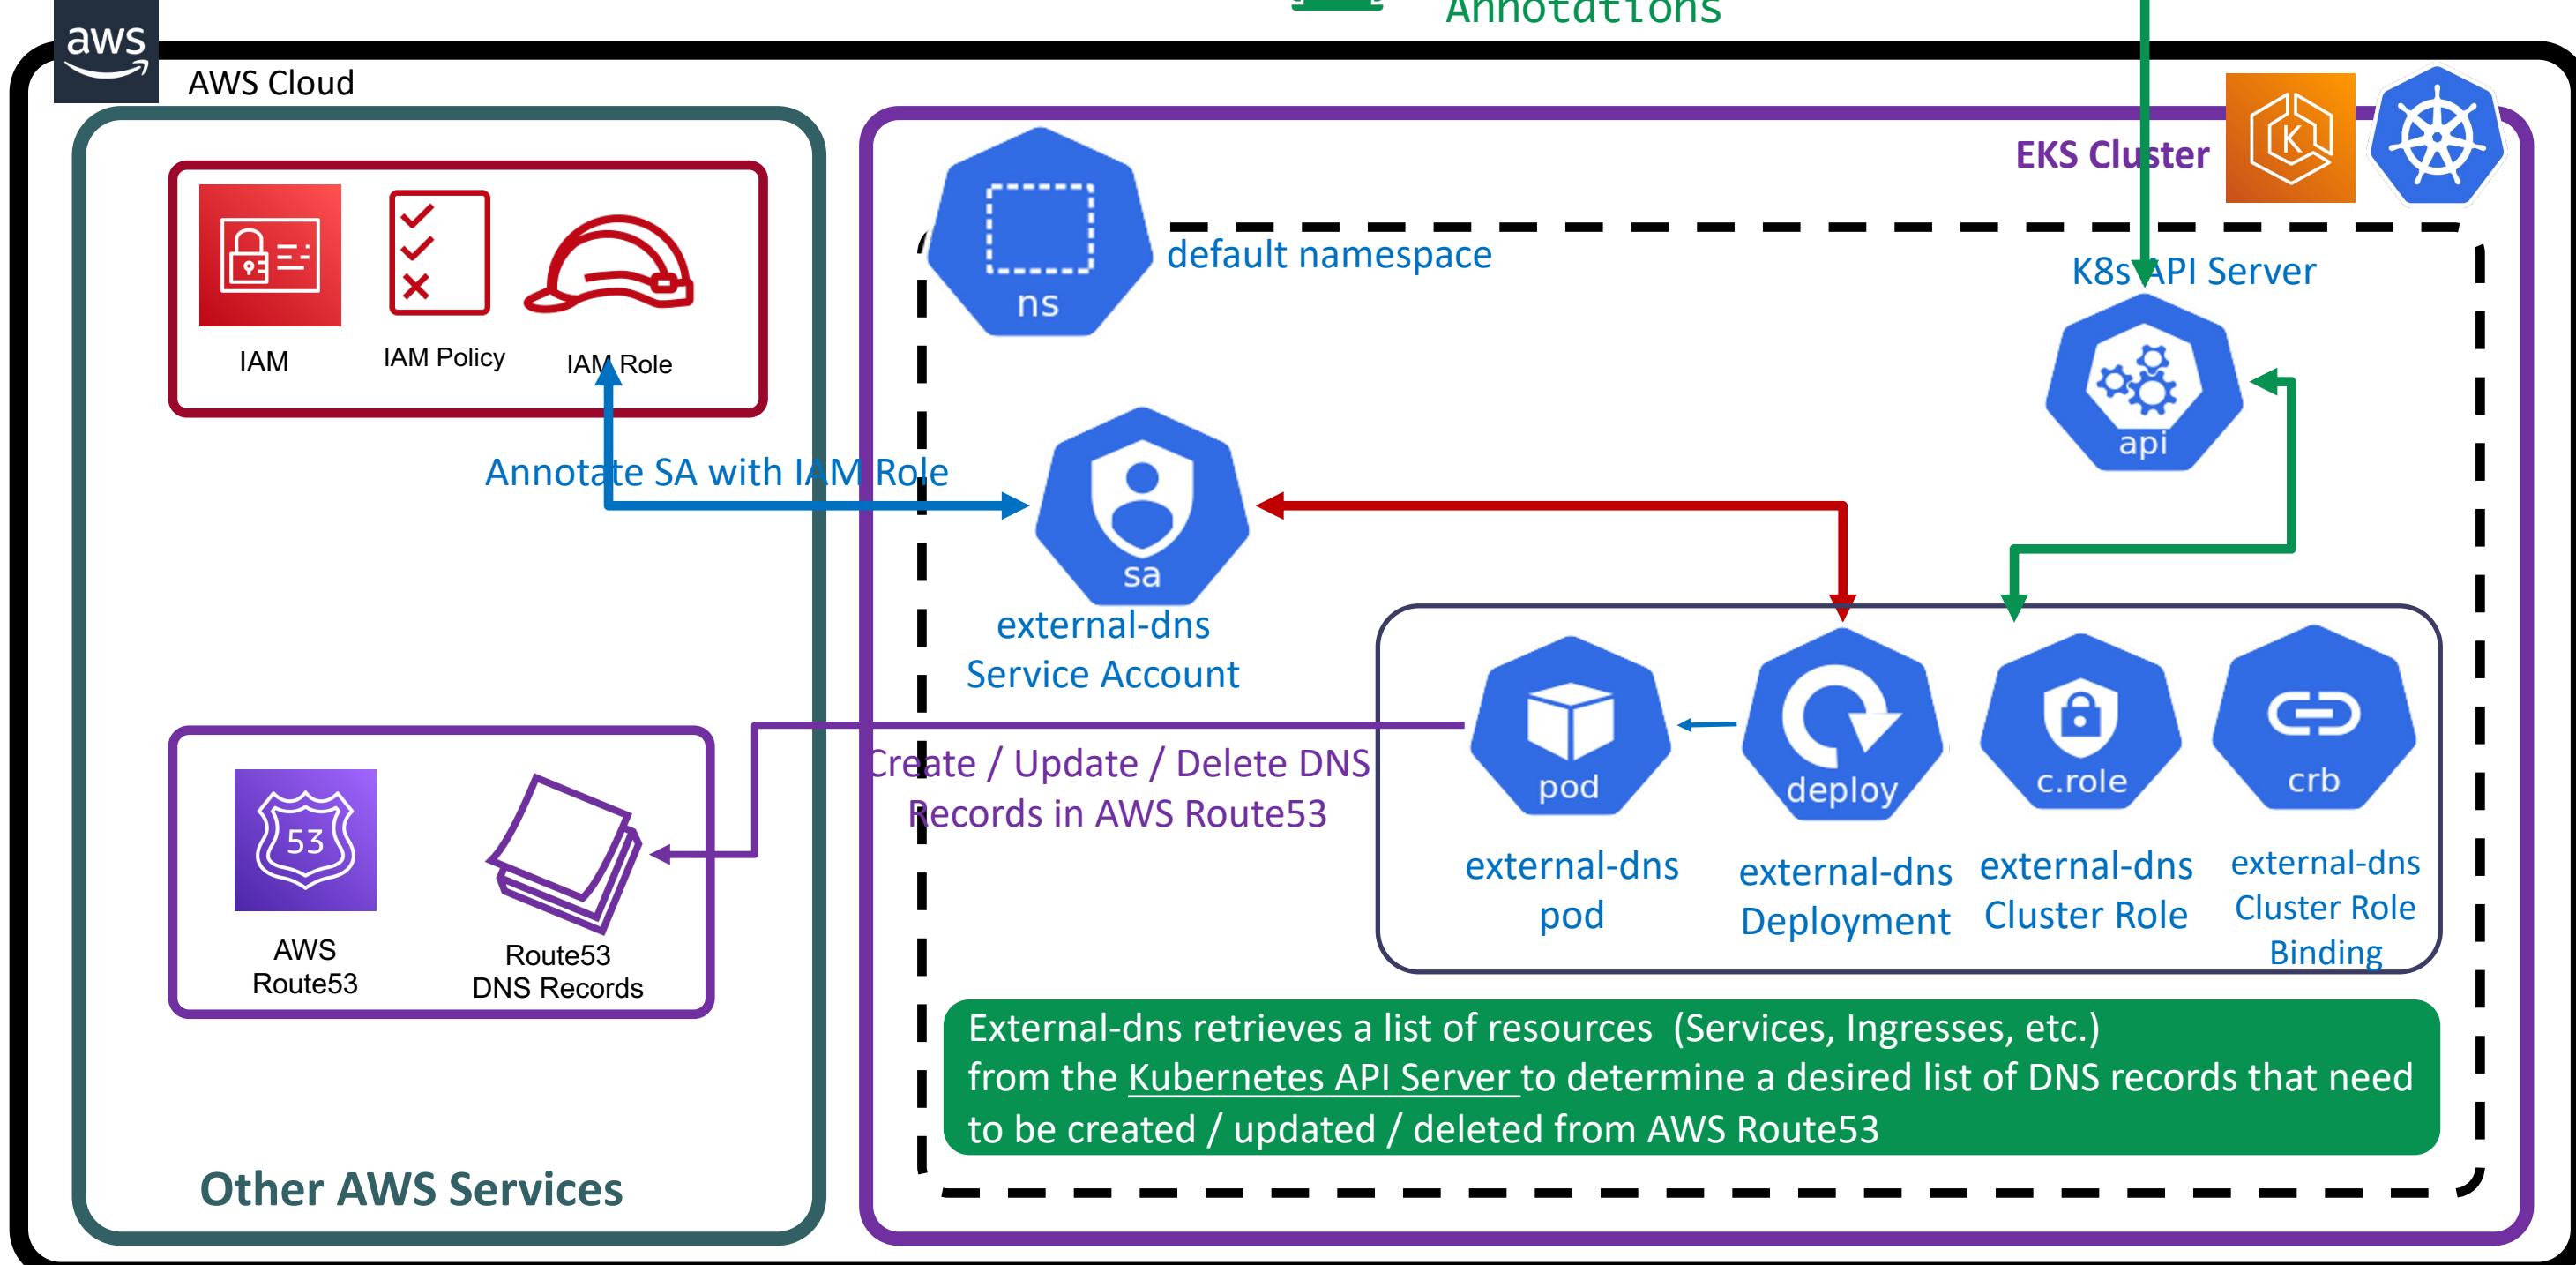
```
# External DNS - For creating a Record Set in Route53
```

```
external-dns.alpha.kubernetes.io/hostname: dnstest901.stack simplify.com
```

# AWS EKS + External DNS

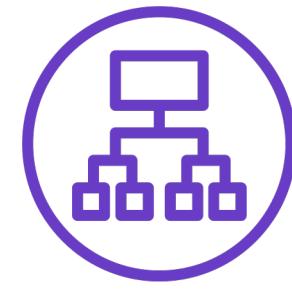


Deploy Ingress /service k8s  
Manifest with external-dns  
Annotations

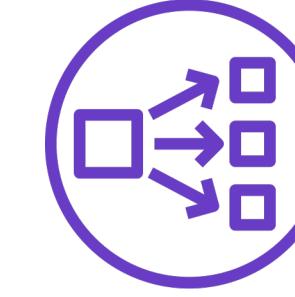




Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress



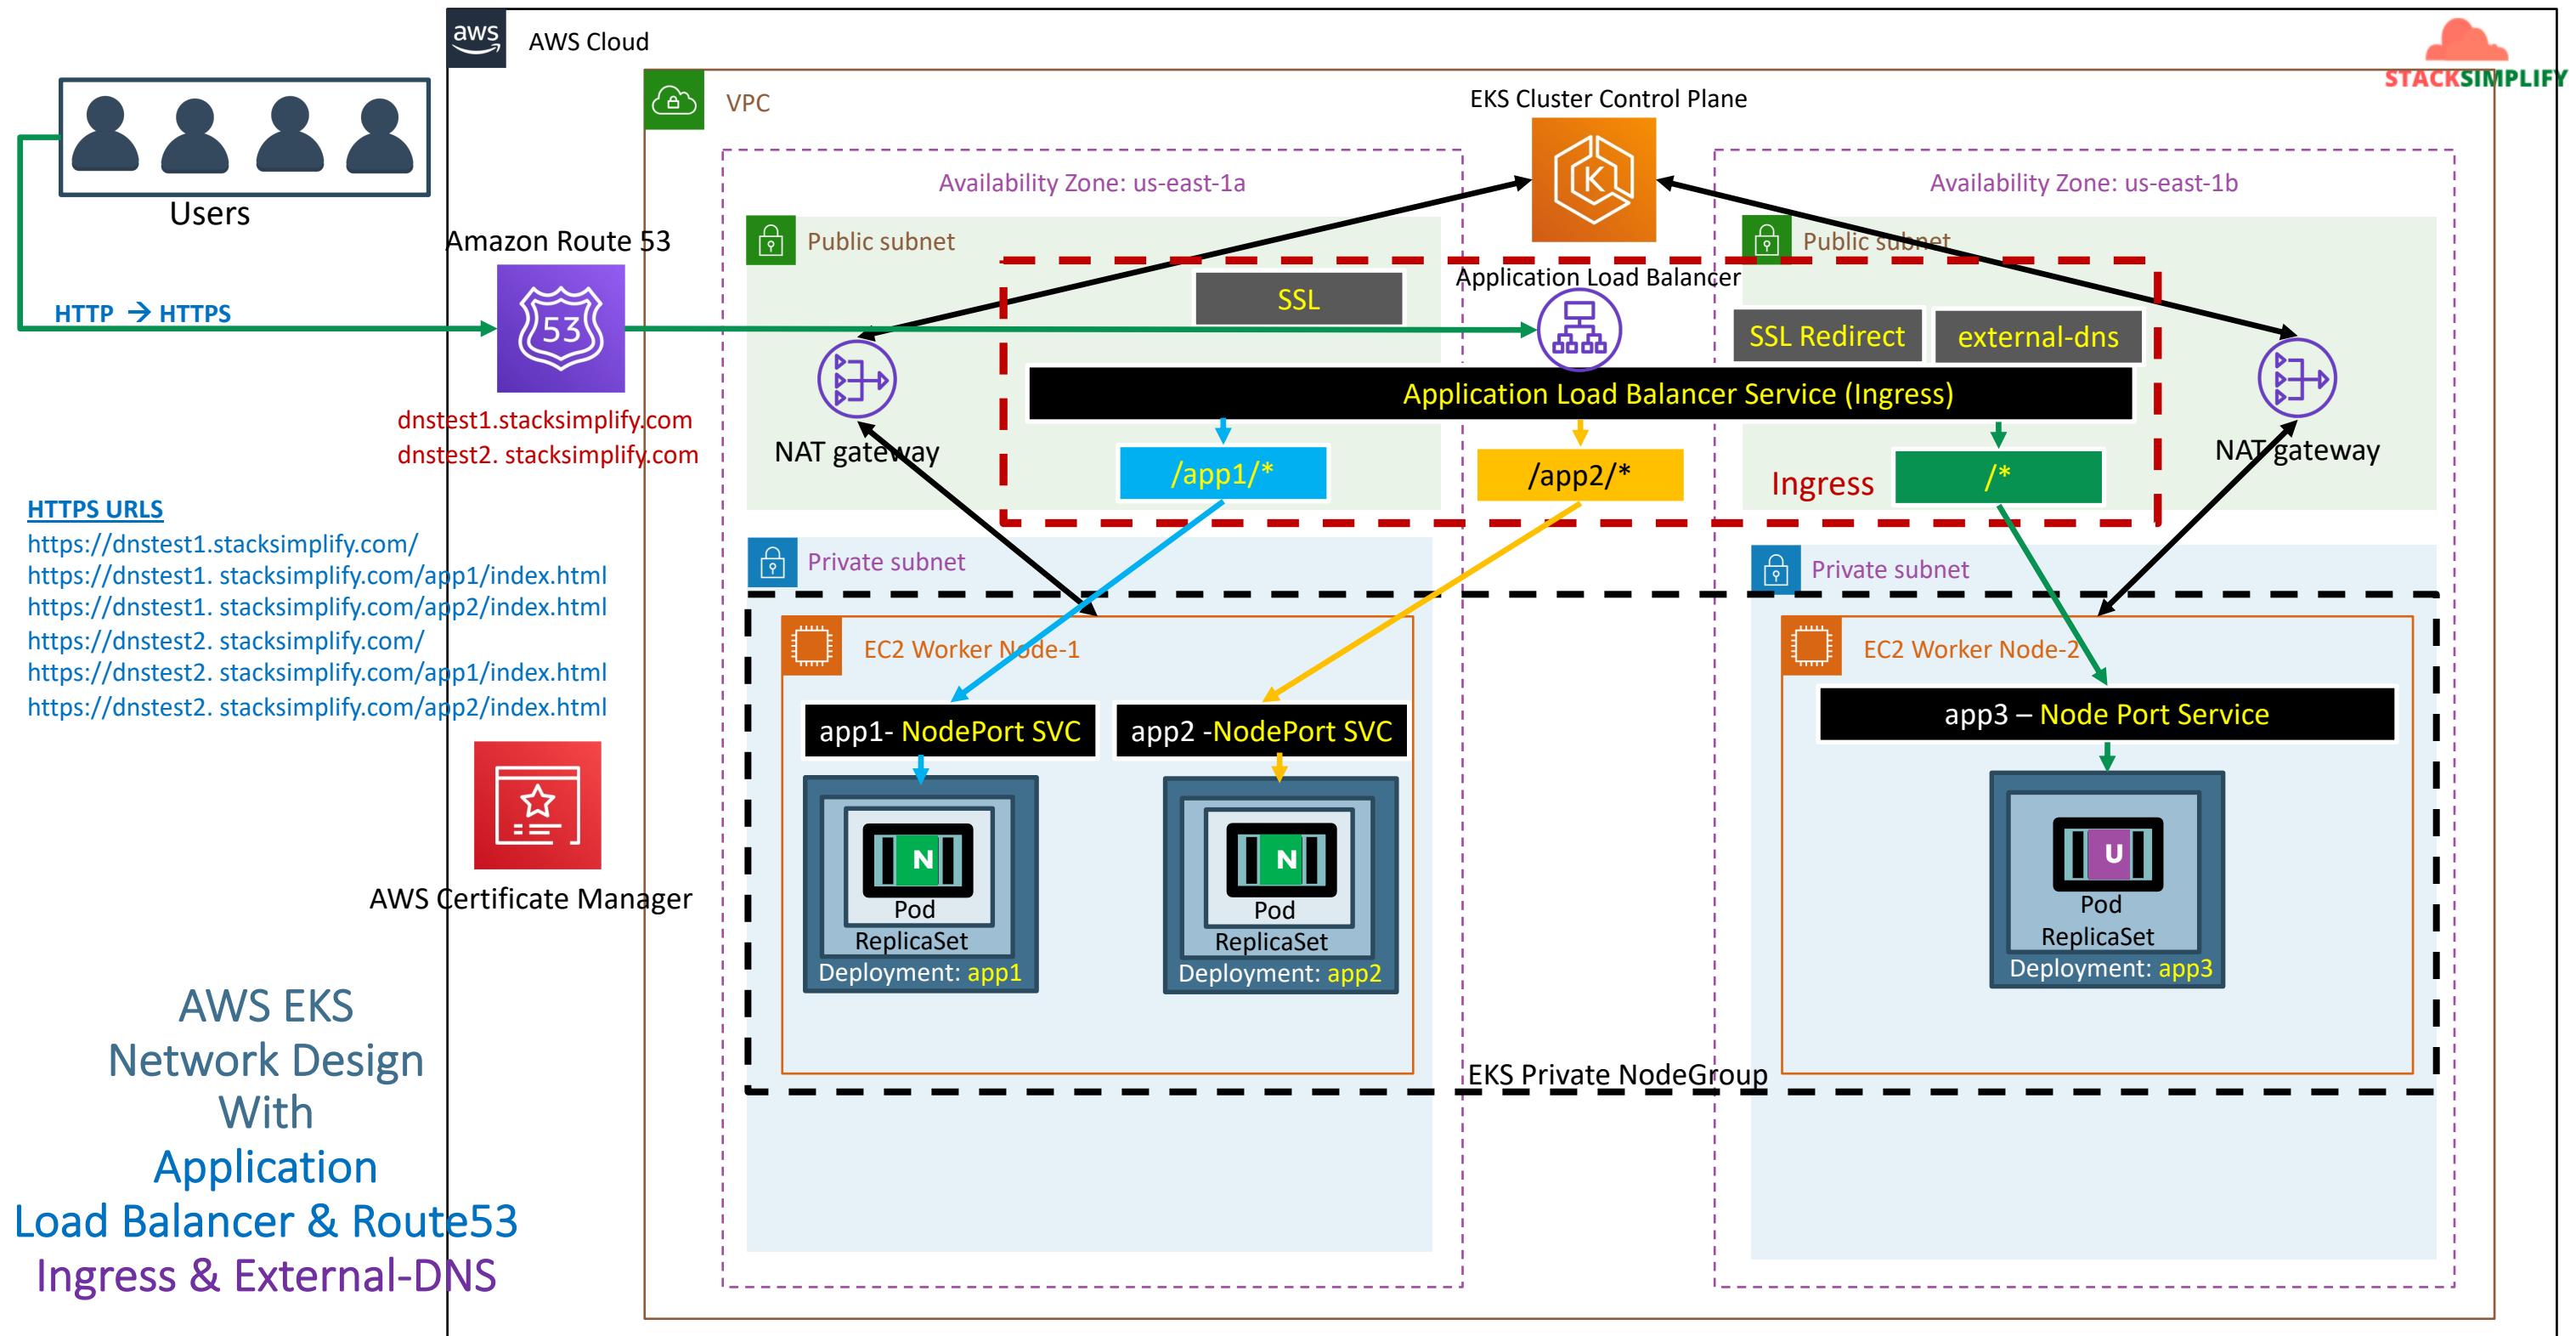
External  
DNS



# AWS EKS Load Balancer Controller

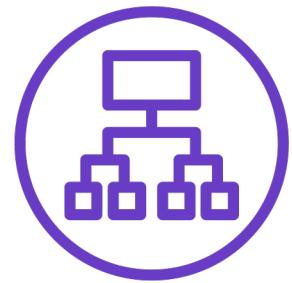
## Kubernetes Ingress & External DNS



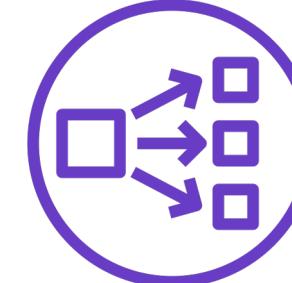




Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress



External  
DNS



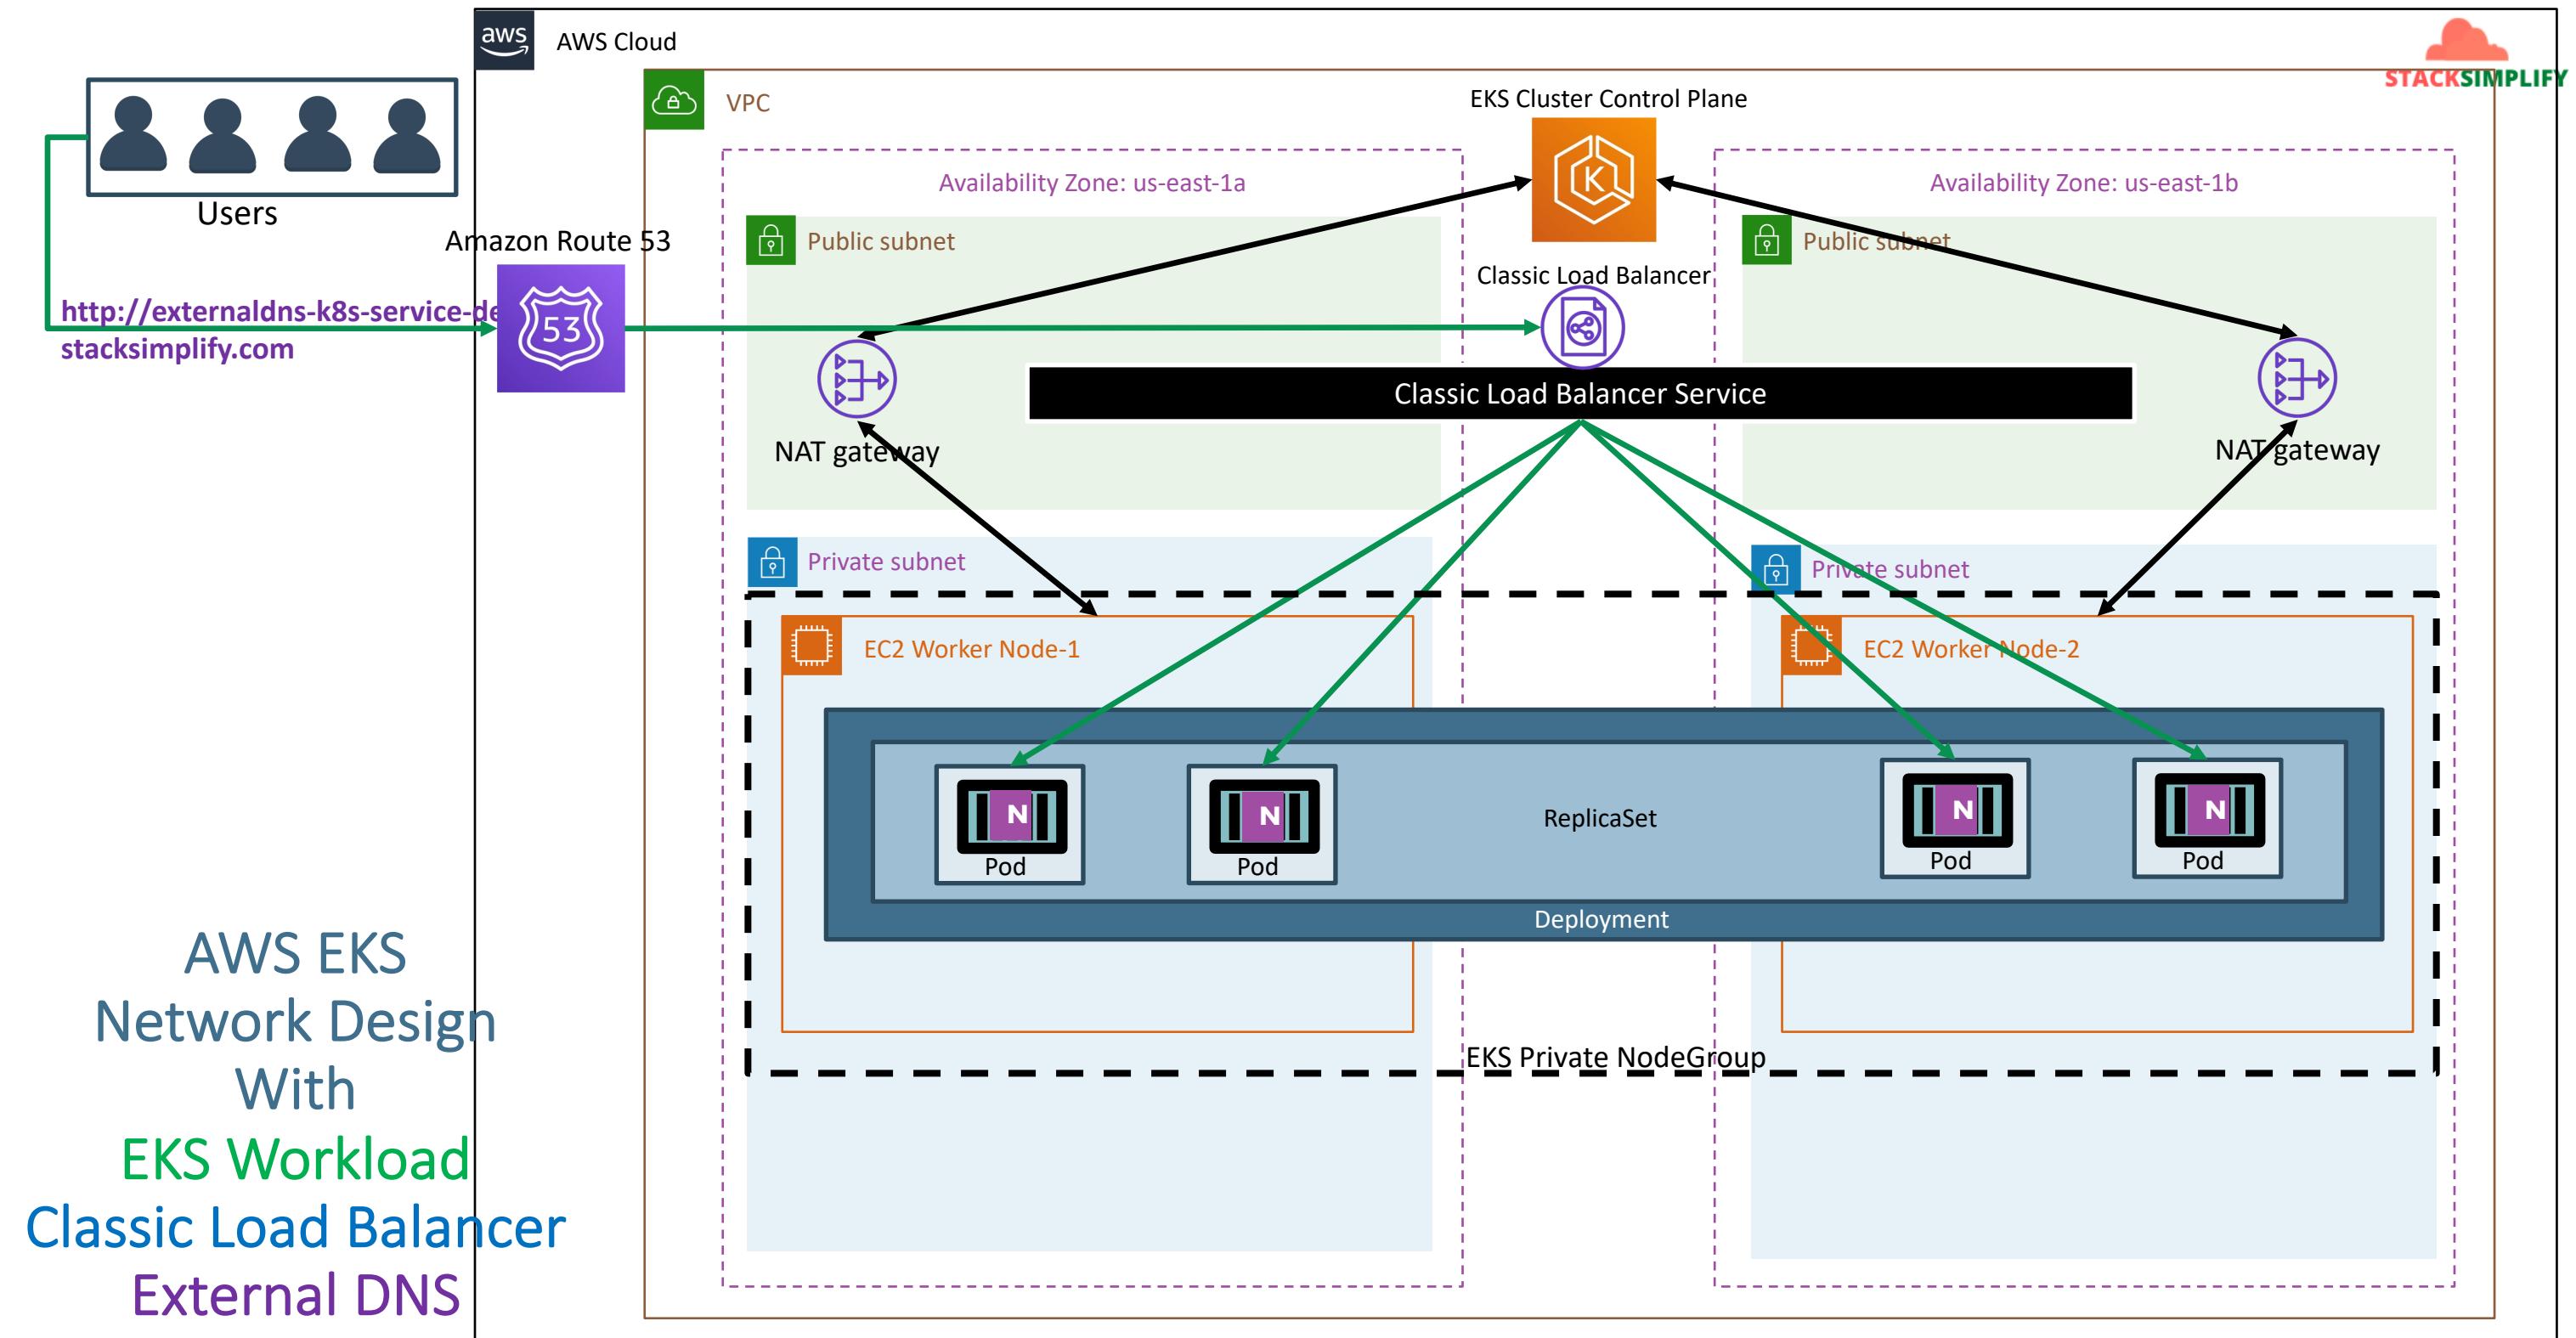
# AWS EKS Load Balancer Controller

## Kubernetes Service & External DNS



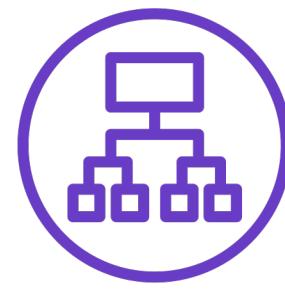
# Kubernetes Service with External DNS

```
apiVersion: v1
kind: Service
metadata:
  name: app1-nginx-loadbalancer-service
  labels:
    app: app1-nginx
  annotations:
    external-dns.alpha.kubernetes.io/hostname: externaldns-k8s-service-demo.stackimplify.
spec:
  type: LoadBalancer
  selector:
    app: app1-nginx
  ports:
    - port: 80
      targetPort: 80
```

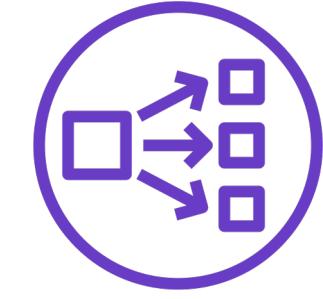




Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress



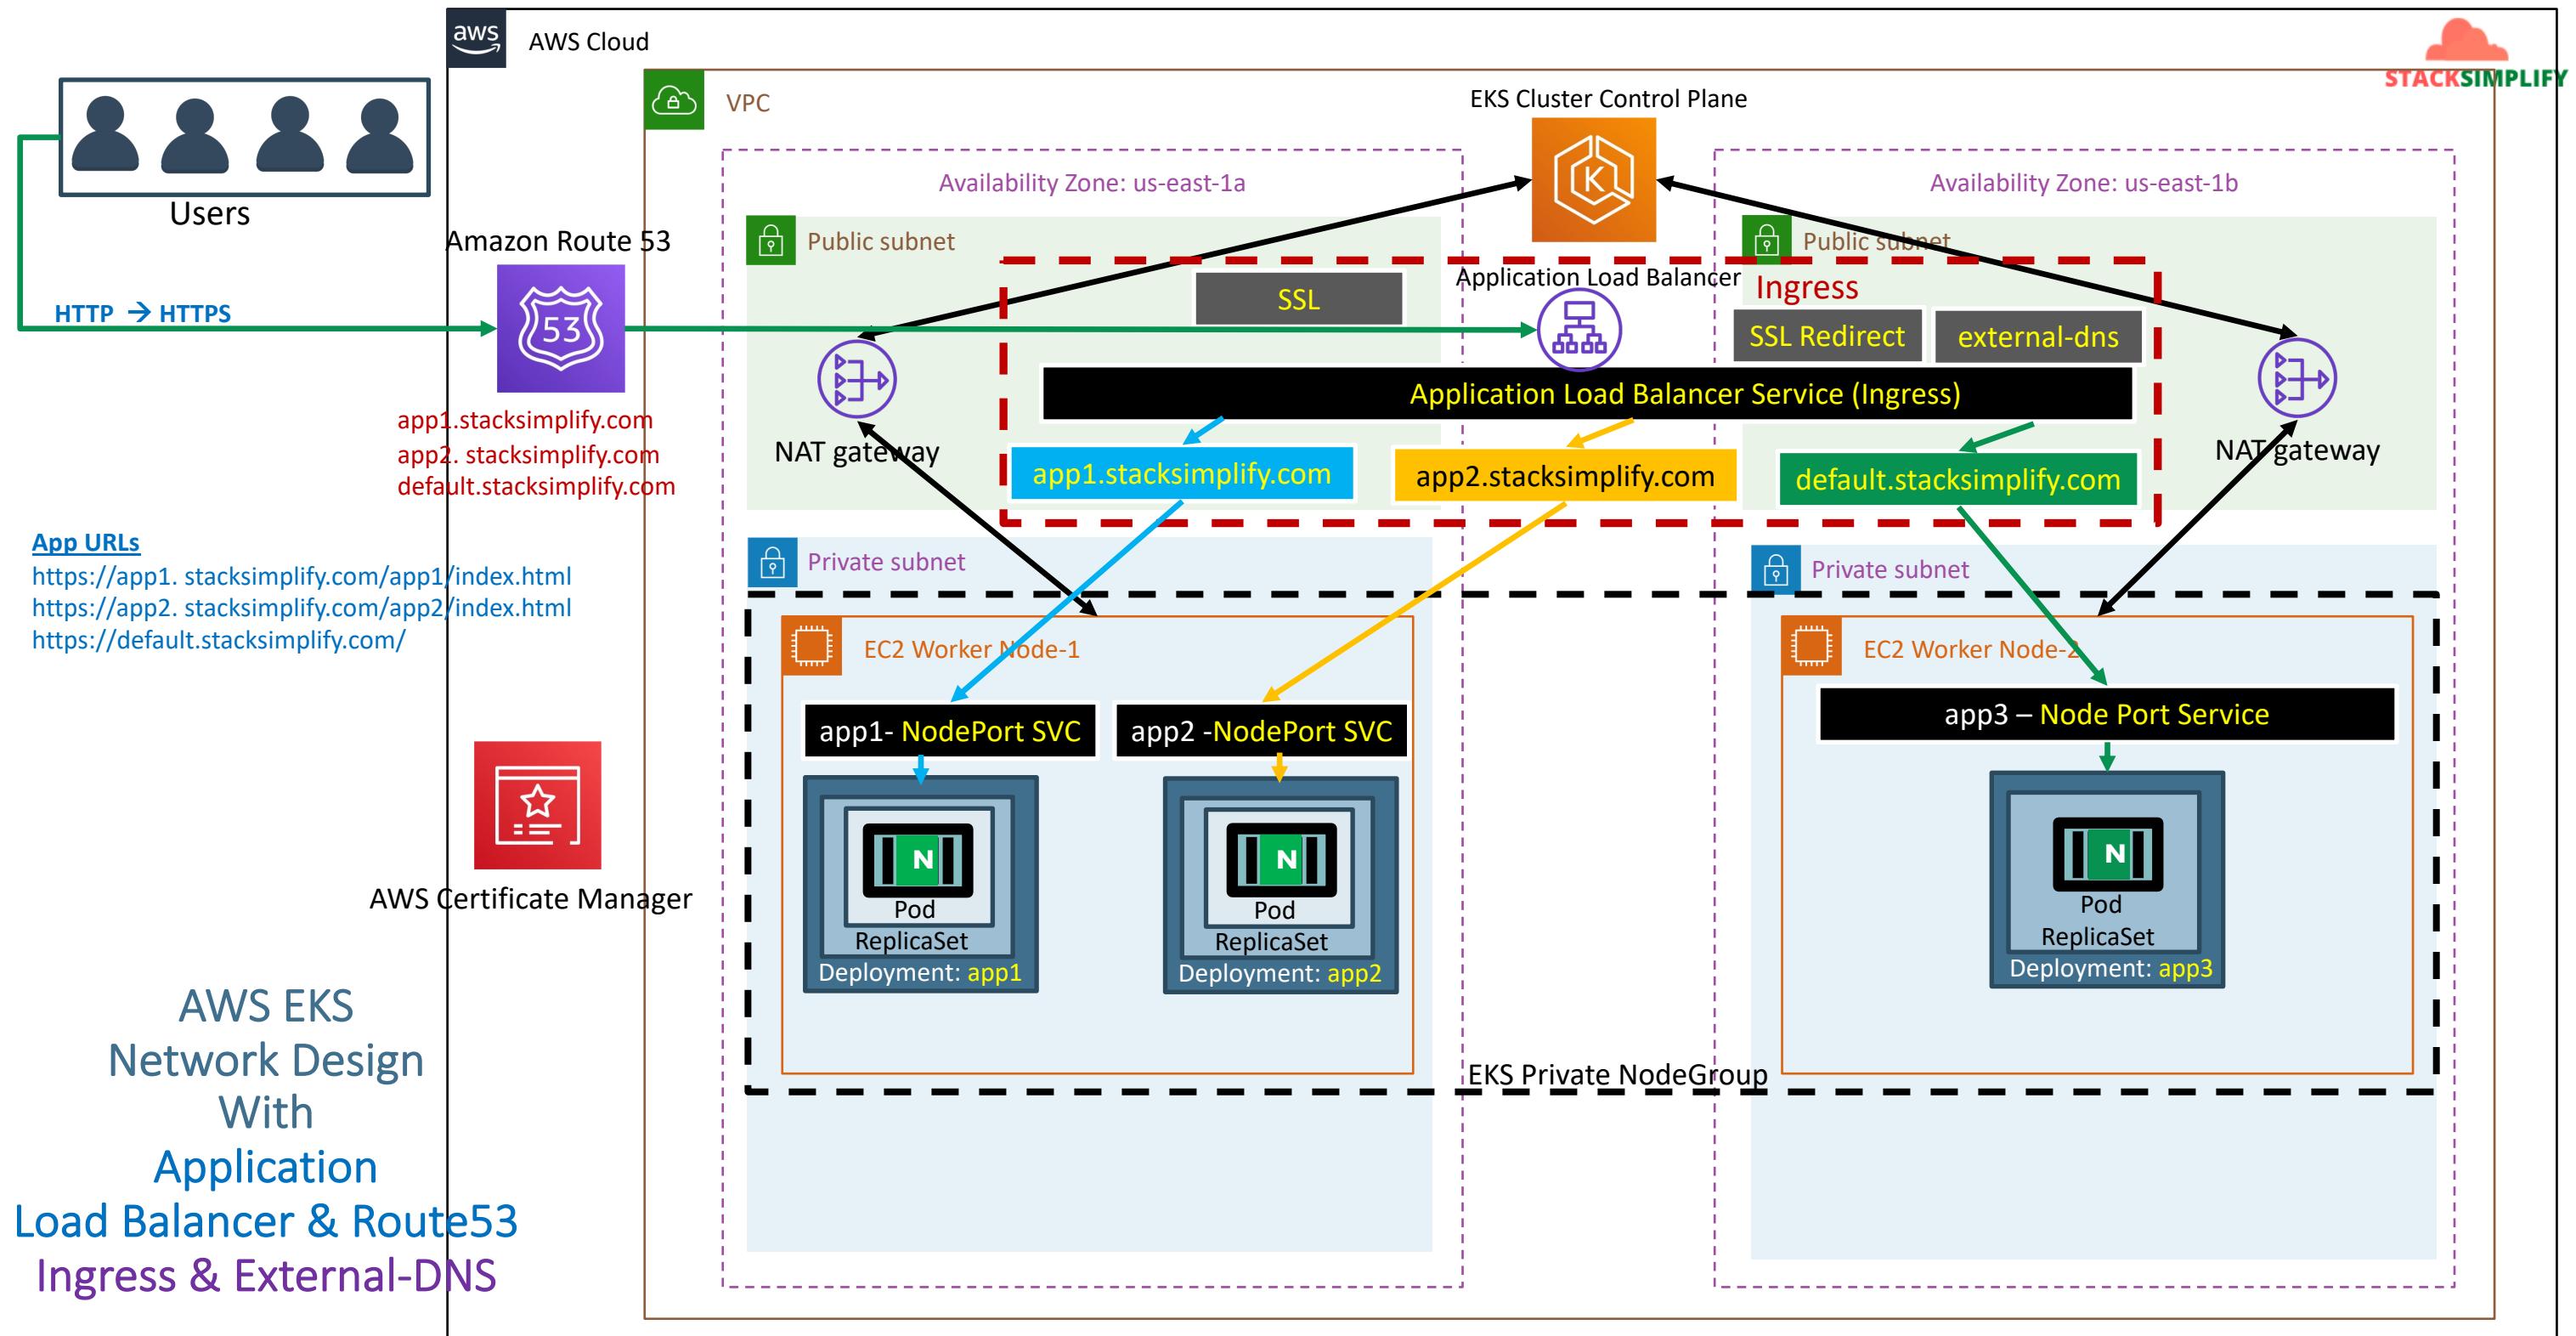
External  
DNS



# AWS EKS Load Balancer Controller



## Kubernetes Ingress Name based Virtual Host / Host Header Routing



**AWS EKS Network Design With Application Load Balancer & Route53 Ingress & External-DNS**

# Name based Virtual Host Routing

```
rules:  
- host: app101.stack simplify.com  
  http:  
    paths:  
      - path: /  
        pathType: Prefix  
    backend:  
      service:  
        name: app1-nginx-nodeport-service  
        port:  
          number: 80  
  
- host: app201.stack simplify.com  
  http:  
    paths:  
      - path: /  
        pathType: Prefix  
    backend:  
      service:  
        name: app2-nginx-nodeport-service  
        port:  
          number: 80
```

Requests with **app101.stack simplify.com** will go to **App1 Node Port Service**

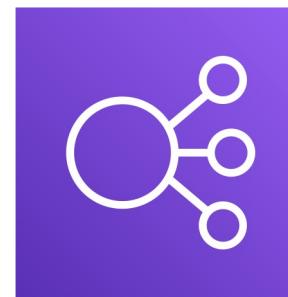
Requests with **app201.stack simplify.com** will go to **App2 Node Port Service**

# Named based Virtual Host Routing

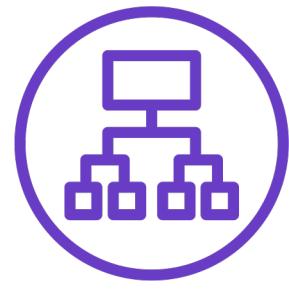
```
# External DNS - For creating a Record Set in Route53
external-dns.alpha.kubernetes.io/hostname: default101.stackimplify.com

spec:
  ingressClassName: my-aws-ingress-class    # Ingress Class
  defaultBackend:
    service:
      name: app3-nginx-nodeport-service
      port:
        number: 80
```

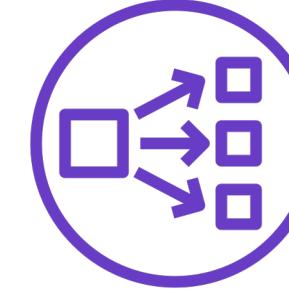
Requests with **default101.stackimplify.com** will go to App3 Node Port Service (Default Backend)



Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress



External  
DNS



# AWS EKS Load Balancer Controller



## Kubernetes Ingress SSL Discovery – Host & TLS

# Ingress SSL Certificate Discovery using Host



```
rules:  
- host: app101.stack simplify.com  
  http:  
    paths:  
      - path: /  
        pathType: Prefix  
      backend:  
        service:  
          name: app1-nginx-nodeport-service  
          port:  
            number: 80  
- host: app201.stack simplify.com  
  http:  
    paths:  
      - path: /  
        pathType: Prefix  
      backend:  
        service:  
          name: app2-nginx-nodeport-service  
          port:  
            number: 80
```

## Demo-1: SSL Discovery – Host

TLS certificates for ALB Listeners can be automatically discovered with **hostnames** from Ingress resources if the **alb.ingress.kubernetes.io/certificate-arn** annotation is not specified.

The controller will attempt to discover **TLS certificates** from the **tls** field in Ingress and **host** field in Ingress rules.

**Important Note:** You need to explicitly specify to use HTTPS listener with listen-ports annotation.  
**alb.ingress.kubernetes.io/listen-ports: '[{"HTTPS":443}]'**

# Ingress SSL Certificate Discovery using tls

```
# External DNS – For creating a Record Set in Route53
external-dns.alpha.kubernetes.io/hostname: certdiscovery-tls-901.stackimplify.com

spec:
  ingressClassName: my-aws-ingress-class    # Ingress Class
  defaultBackend:
    service:
      name: app3-nginx-nodeport-service
      port:
        number: 80
    tls:
      - hosts:
          - "*.stackimplify.com"
rules:
  - http:
      paths:
        - path: /app1
          pathType: Prefix
          backend:
            service:
              name: app1-nginx-nodeport-service
```

Demo-2: SSL  
Discovery – tls

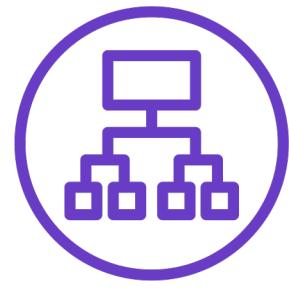
# Comment certificate-arn Annotation

Comment **alb.ingress.kubernetes.io/certificate-arn**  
Annotation when using host / tls options

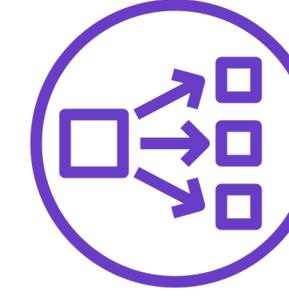
```
## SSL Settings
alb.ingress.kubernetes.io/listen-ports: '[{"HTTPS":443}, {"HTTP":80}]'
#alb.ingress.kubernetes.io/certificate-arn: arn:aws:acm:us-east-1:180789647333:certificate/632a3ff6-3f6d-464c-9121-b9d97481a76b
#alb.ingress.kubernetes.io/ssl-policy: ELBSecurityPolicy-TLS-1-1-2017-01 #Optional (Picks default if not used)
# SSL Redirect Setting
alb.ingress.kubernetes.io/ssl-redirect: '443'
```



Elastic Load  
Balancing



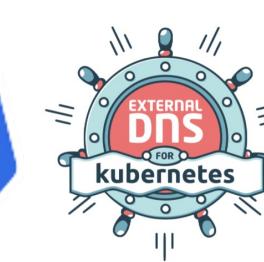
Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress

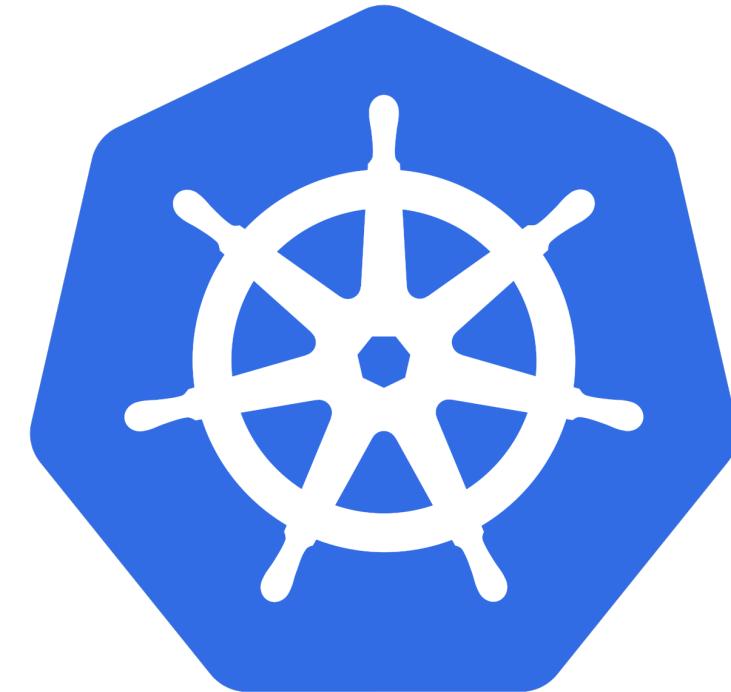


External  
DNS



# AWS EKS Load Balancer Controller

## Kubernetes Ingress Ingress Groups



# Without Ingress Groups

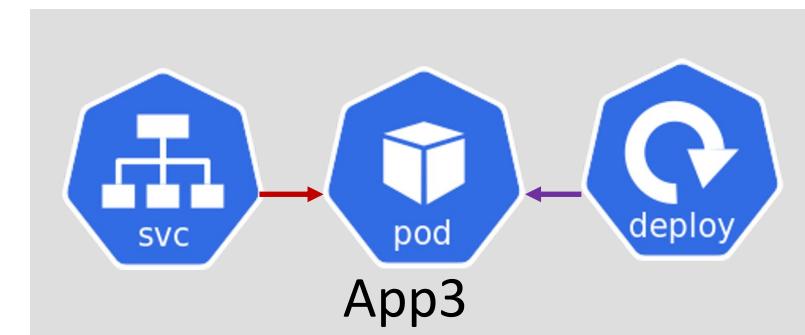
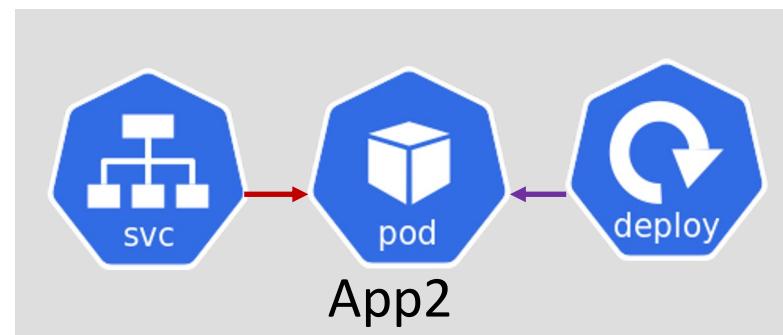
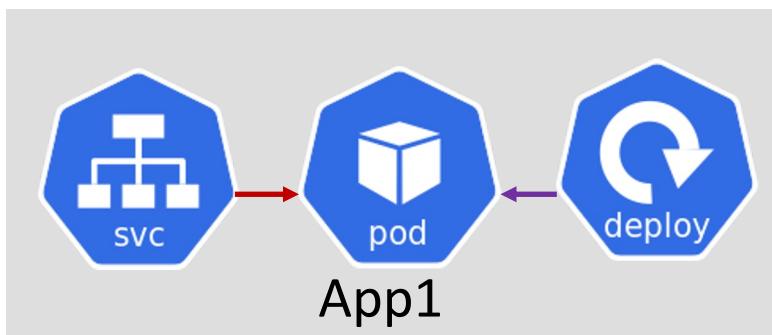
Single Ingress Resource for all three Apps

**Example:** If we have 50 apps which needs different rules in them and all should be part of single ALB, maintaining such huge k8s manifest file with configs becomes tedious and confusing.



Any changes to App3 we need to update the “my-apps” ingress resource

With **Ingress Groups** we can simplify our Kubernetes Ingress Manifest when dealing with **multiple apps** requiring single ALB



# With Ingress Groups

alb.ingress.kubernetes.io/load-balancer-name: ingress-groups-demo

Ingress Group: myapps.web

Priority: 10



Ingress Service: app1-ingress

Priority: 20

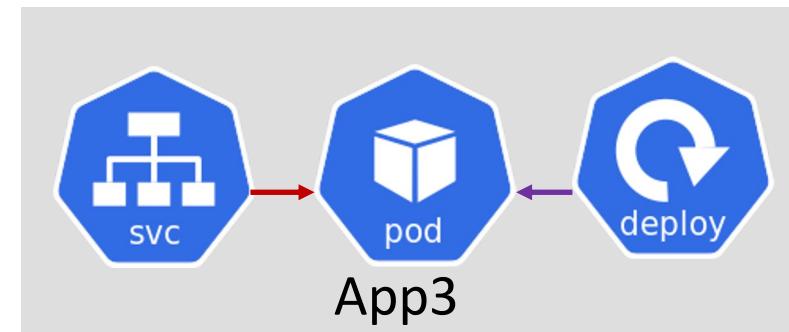
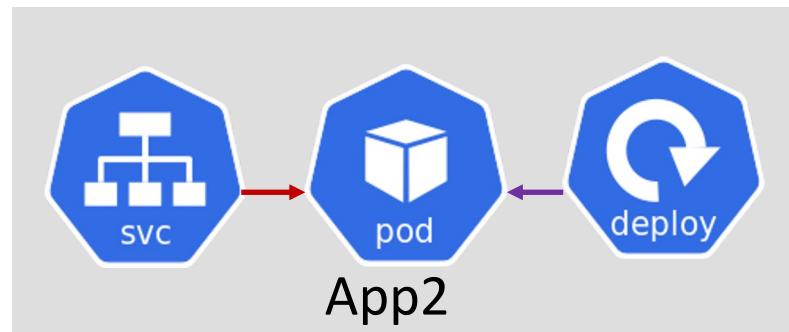
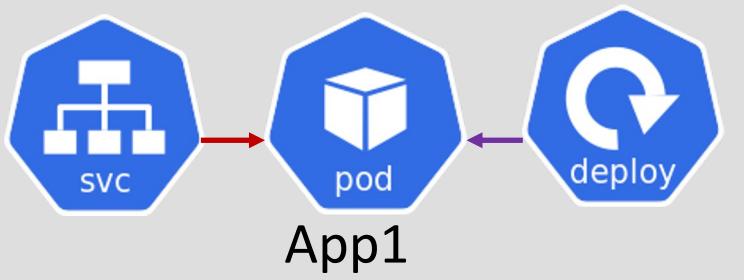


Ingress Service: app2-ingress

Priority: 30



Ingress Service: app3-ingress



# Ingress Groups

1

IngressGroup feature enables you to group multiple Ingress resources together.

2

The controller will automatically merge Ingress rules for all Ingresses within IngressGroup and support them with a single ALB.

3

**VERY VERY IMPORTANT NOTE:** In addition, most annotations defined on an Ingress only applies to the paths defined by that Ingress.

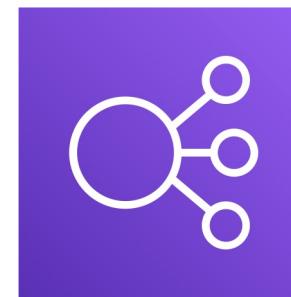
Sample

```
# Ingress Groups
alb.ingress.kubernetes.io/group.name: myapps.web
alb.ingress.kubernetes.io/group.order: '10'
```

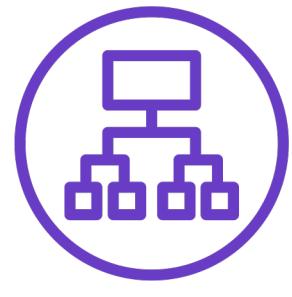
# Ingress Groups

```
✓ 08-12-IngressGroups
  ✓ kube-manifests
    ✓ app1
      ! 01-Nginx-App1-Deployment-and-NodePortService.yml
      ! 02-App1-Ingress.yml
    ✓ app2
      ! 01-Nginx-App2-Deployment-and-NodePortService.yml
      ! 02-App2-Ingress.yml
    ✓ app3
      ! 01-Nginx-App3-Deployment-and-NodePortService.yml
      ! 03-App3-Ingress-default-backend.yml
```

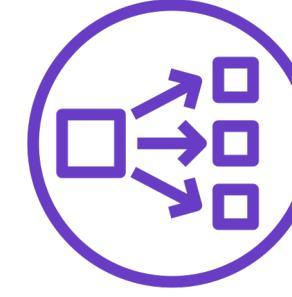
3 Ingress Kubernetes  
Manifests creates 3  
Ingress Resources and  
merge to create a Single  
AWS ALB.



Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress



External  
DNS



# AWS EKS Load Balancer Controller

## Kubernetes Ingress Target Type: IP

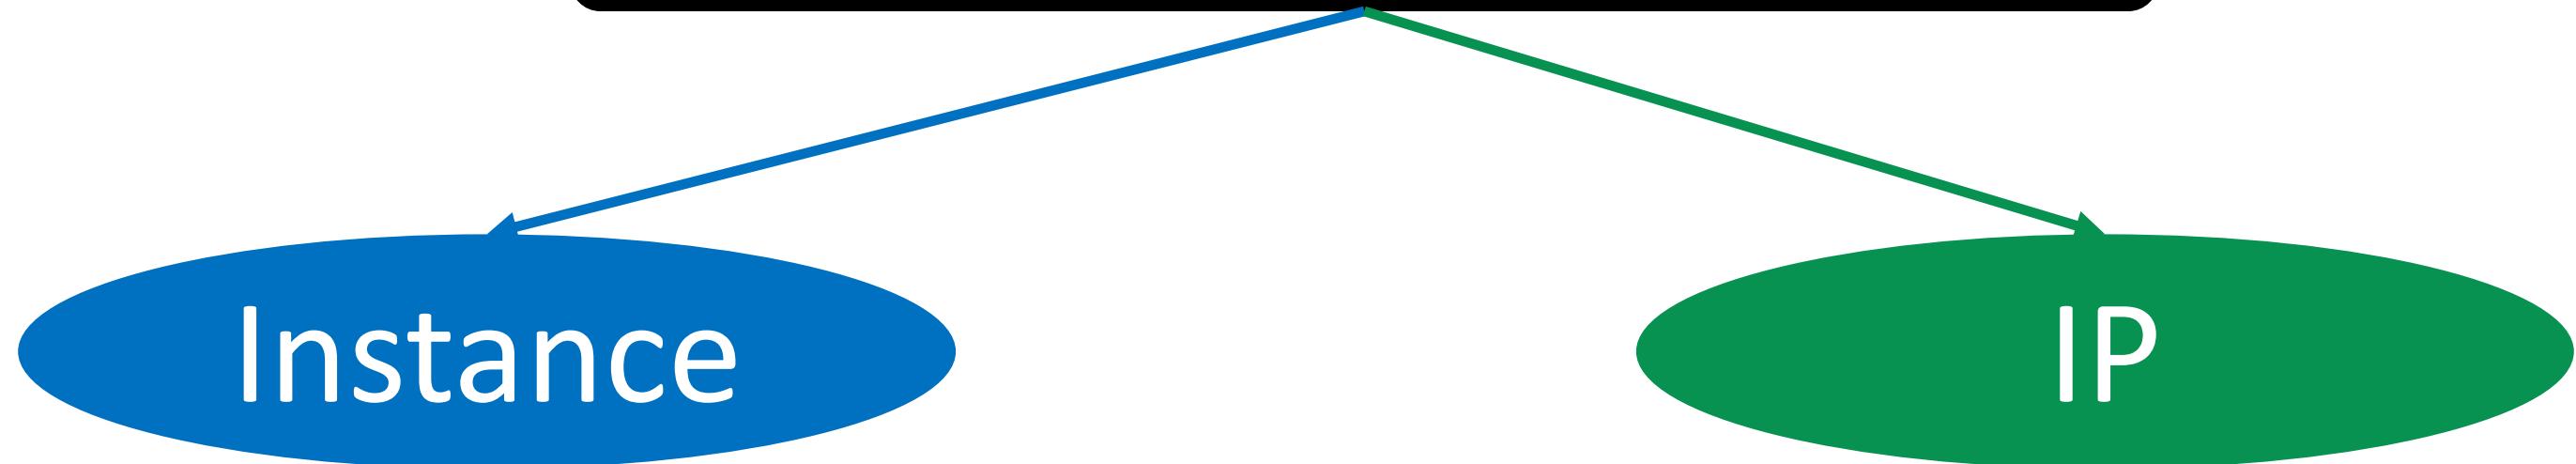


# ALB Ingress Target Types

What is  
**alb.ingress.kubernetes.io/target-type**  
Annotation ?

**target-type** specifies how to route  
traffic to pods.

## ALB Ingress Target Types



# ALB Ingress Target Types

Instance

instance mode will route traffic to all ec2 instances within cluster on NodePort opened for your service.

service must be of type "NodePort" or "LoadBalancer" to use instance mode

Default Target Type is Instance Mode which means no need to define this annotation in Ingress, by default it takes Instance Mode

```
# Target Type: Instance  
alb.ingress.kubernetes.io/target-type: instance
```

IP

ip mode will route traffic directly to the pod IP.

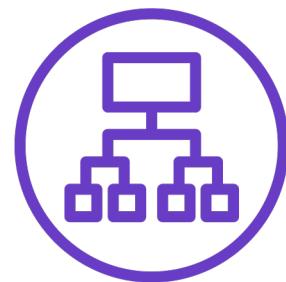
ip mode is required for sticky sessions to work with Application Load Balancers.

Need to define the annotation explicitly as target-type: ip

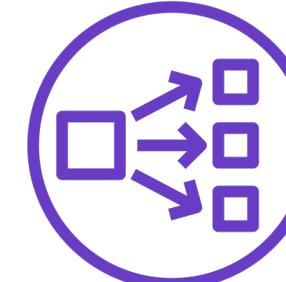
```
# Target Type: IP  
alb.ingress.kubernetes.io/target-type: ip
```



Elastic Load  
Balancing



Application  
Load Balancer



Network  
Load Balancer



Kubernetes  
Ingress



External  
DNS



# AWS EKS

## Load Balancer Controller

**Kubernetes Ingress**  
**Internal Application**  
**Load Balancer**



# Ingress ALB Internal Load Balancer

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: ingress-internal-lb-demo
annotations:
  # Load Balancer Name
  alb.ingress.kubernetes.io/load-balancer-name: ingress-internal-lb
  # Ingress Core Settings
  #kubernetes.io/ingress.class: "alb" (OLD INGRESS CLASS NOTATION -
  # Creates External Application Load Balancer
  #alb.ingress.kubernetes.io/scheme: internet-facing
  # Creates Internal Application Load Balancer
  #alb.ingress.kubernetes.io/scheme: internal
  # Health Check Settings
  alb.ingress.kubernetes.io/healthcheck-protocol: HTTP
  alb.ingress.kubernetes.io/healthcheck-port: traffic-port
```

Creates AWS  
ALB Internal  
Load Balancer

# Curl Pod

```
apiVersion: v1
kind: Pod
metadata:
  name: curl-pod
spec:
  containers:
  - name: curl
    image: curlimages/curl
    command: [ "sleep", "600" ]
```

We **cannot** connect to Internal LB directly from internet to test it.

We will deploy a **curl pod** in EKS Cluster so we can connect to curl pod in EKS Cluster and test the **Internal LB Endpoint** using **curl** command.

# Command to Test Ingress AWS Internal ALB using curl pod

```
# Deploy curl-pod
kubectl apply -f kube-manifests-curl

# Will open up a terminal session into the container
kubectl exec -it curl-pod -- sh

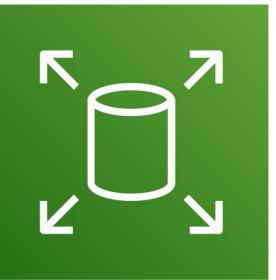
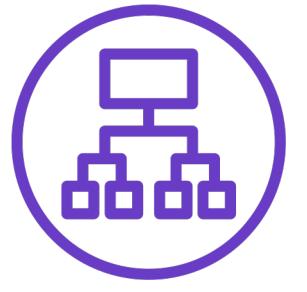
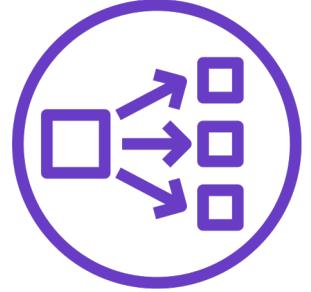
# We can now curl external addresses or internal services:
curl http://google.com/
curl <INTERNAL-INGRESS-LB-DNS>

# Default Backend Curl Test
curl internal-ingress-internal-lb-1839544354.us-east-1.elb.amazonaws.com

# App1 Curl Test
curl internal-ingress-internal-lb-1839544354.us-east-1.elb.amazonaws.com/app1/index.html

# App2 Curl Test
curl internal-ingress-internal-lb-1839544354.us-east-1.elb.amazonaws.com/app2/index.html
```

# AWS Load Balancer Controller Updates - End



Elastic Load  
Balancing

Classic  
Load Balancer

Network  
Load Balancer

Application  
Load Balancer

Certificate  
Manager

Route53

Elastic Block  
Store

Amazon RDS

# AWS EKS

## Fargate Profiles

## Serverless

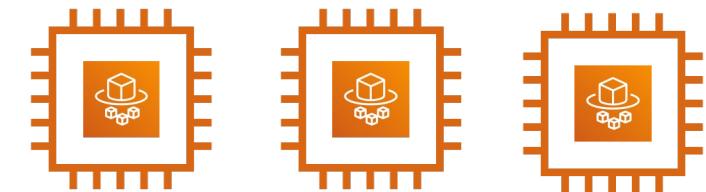
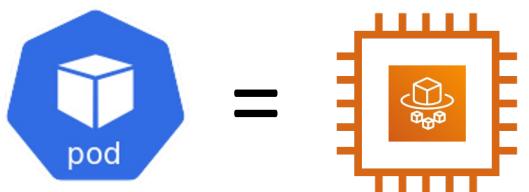


Fargate  
Profiles

# What is Fargate?

- Fargate is a **Serverless compute platform** for containers on AWS
- Fargate provides **on-demand, right-sized compute capacity** for containers
- EKS integrates Kubernetes with Fargate by using **controllers** that are built by AWS using the **upstream, extensible model** provided by Kubernetes.
- These controllers run as part of the **EKS managed Kubernetes control plane** and are responsible for **scheduling** native Kubernetes pods onto Fargate.
- The **Fargate controllers** include a **new scheduler** that runs alongside the **default Kubernetes scheduler** in addition to several mutating and validating admission controllers.
- When we start a pod that **meets the criteria for running on Fargate**, the Fargate controllers running in the cluster recognize, update, and **schedule the pod onto Fargate**.

# AWS EKS on Fargate



## Bring existing pods

- We don't need to change our existing pods
- Fargate works with existing workflows and services that run on kubernetes

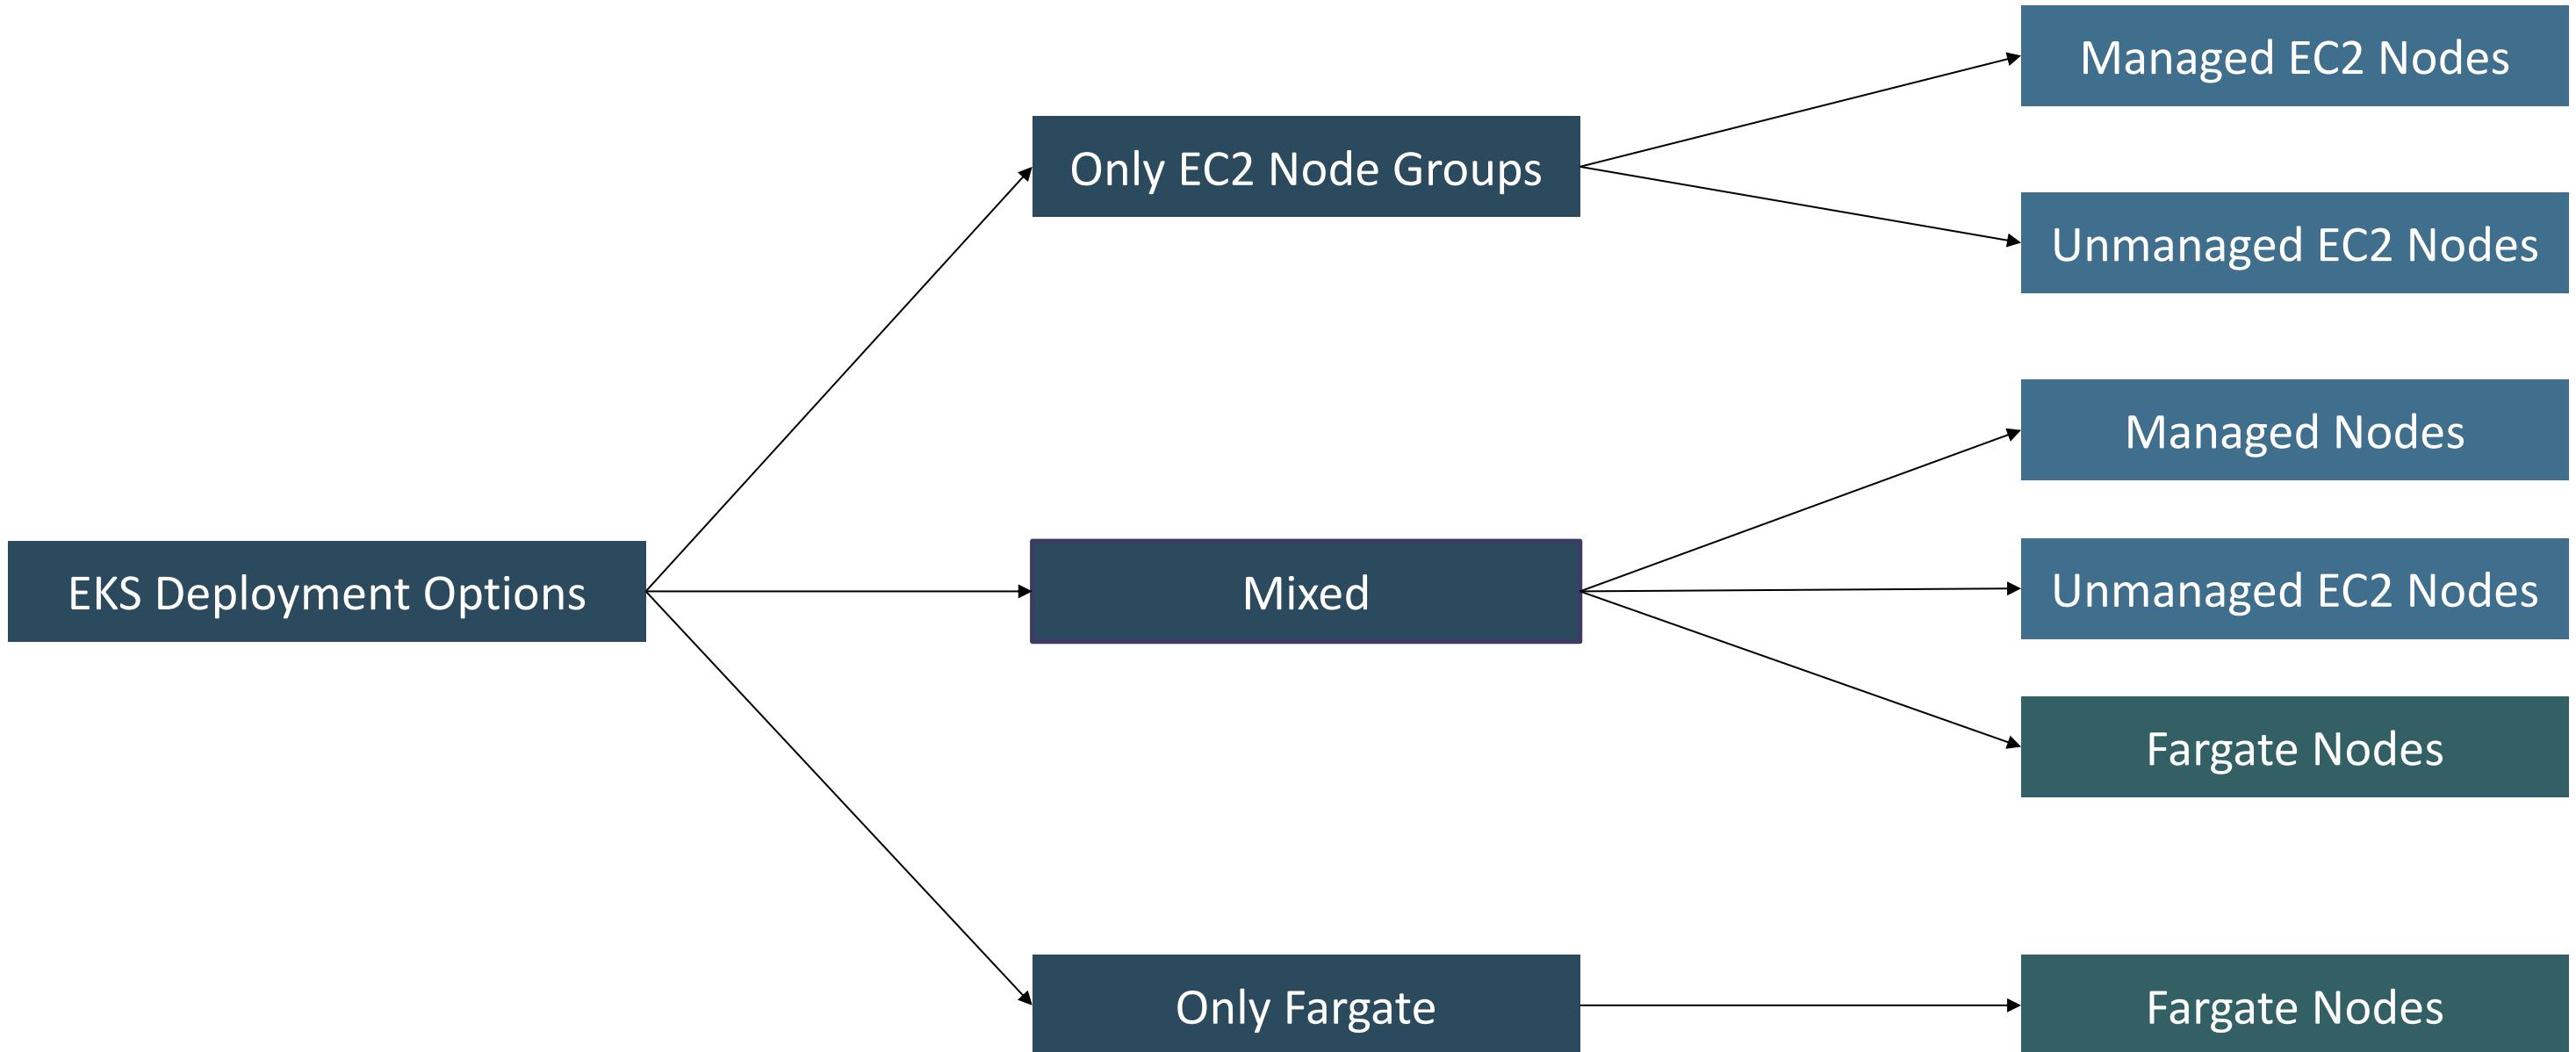
## Production Ready

- Launch pods easily.
- Easily run pods across Azs for HA
- Each pod runs in an isolated compute environment

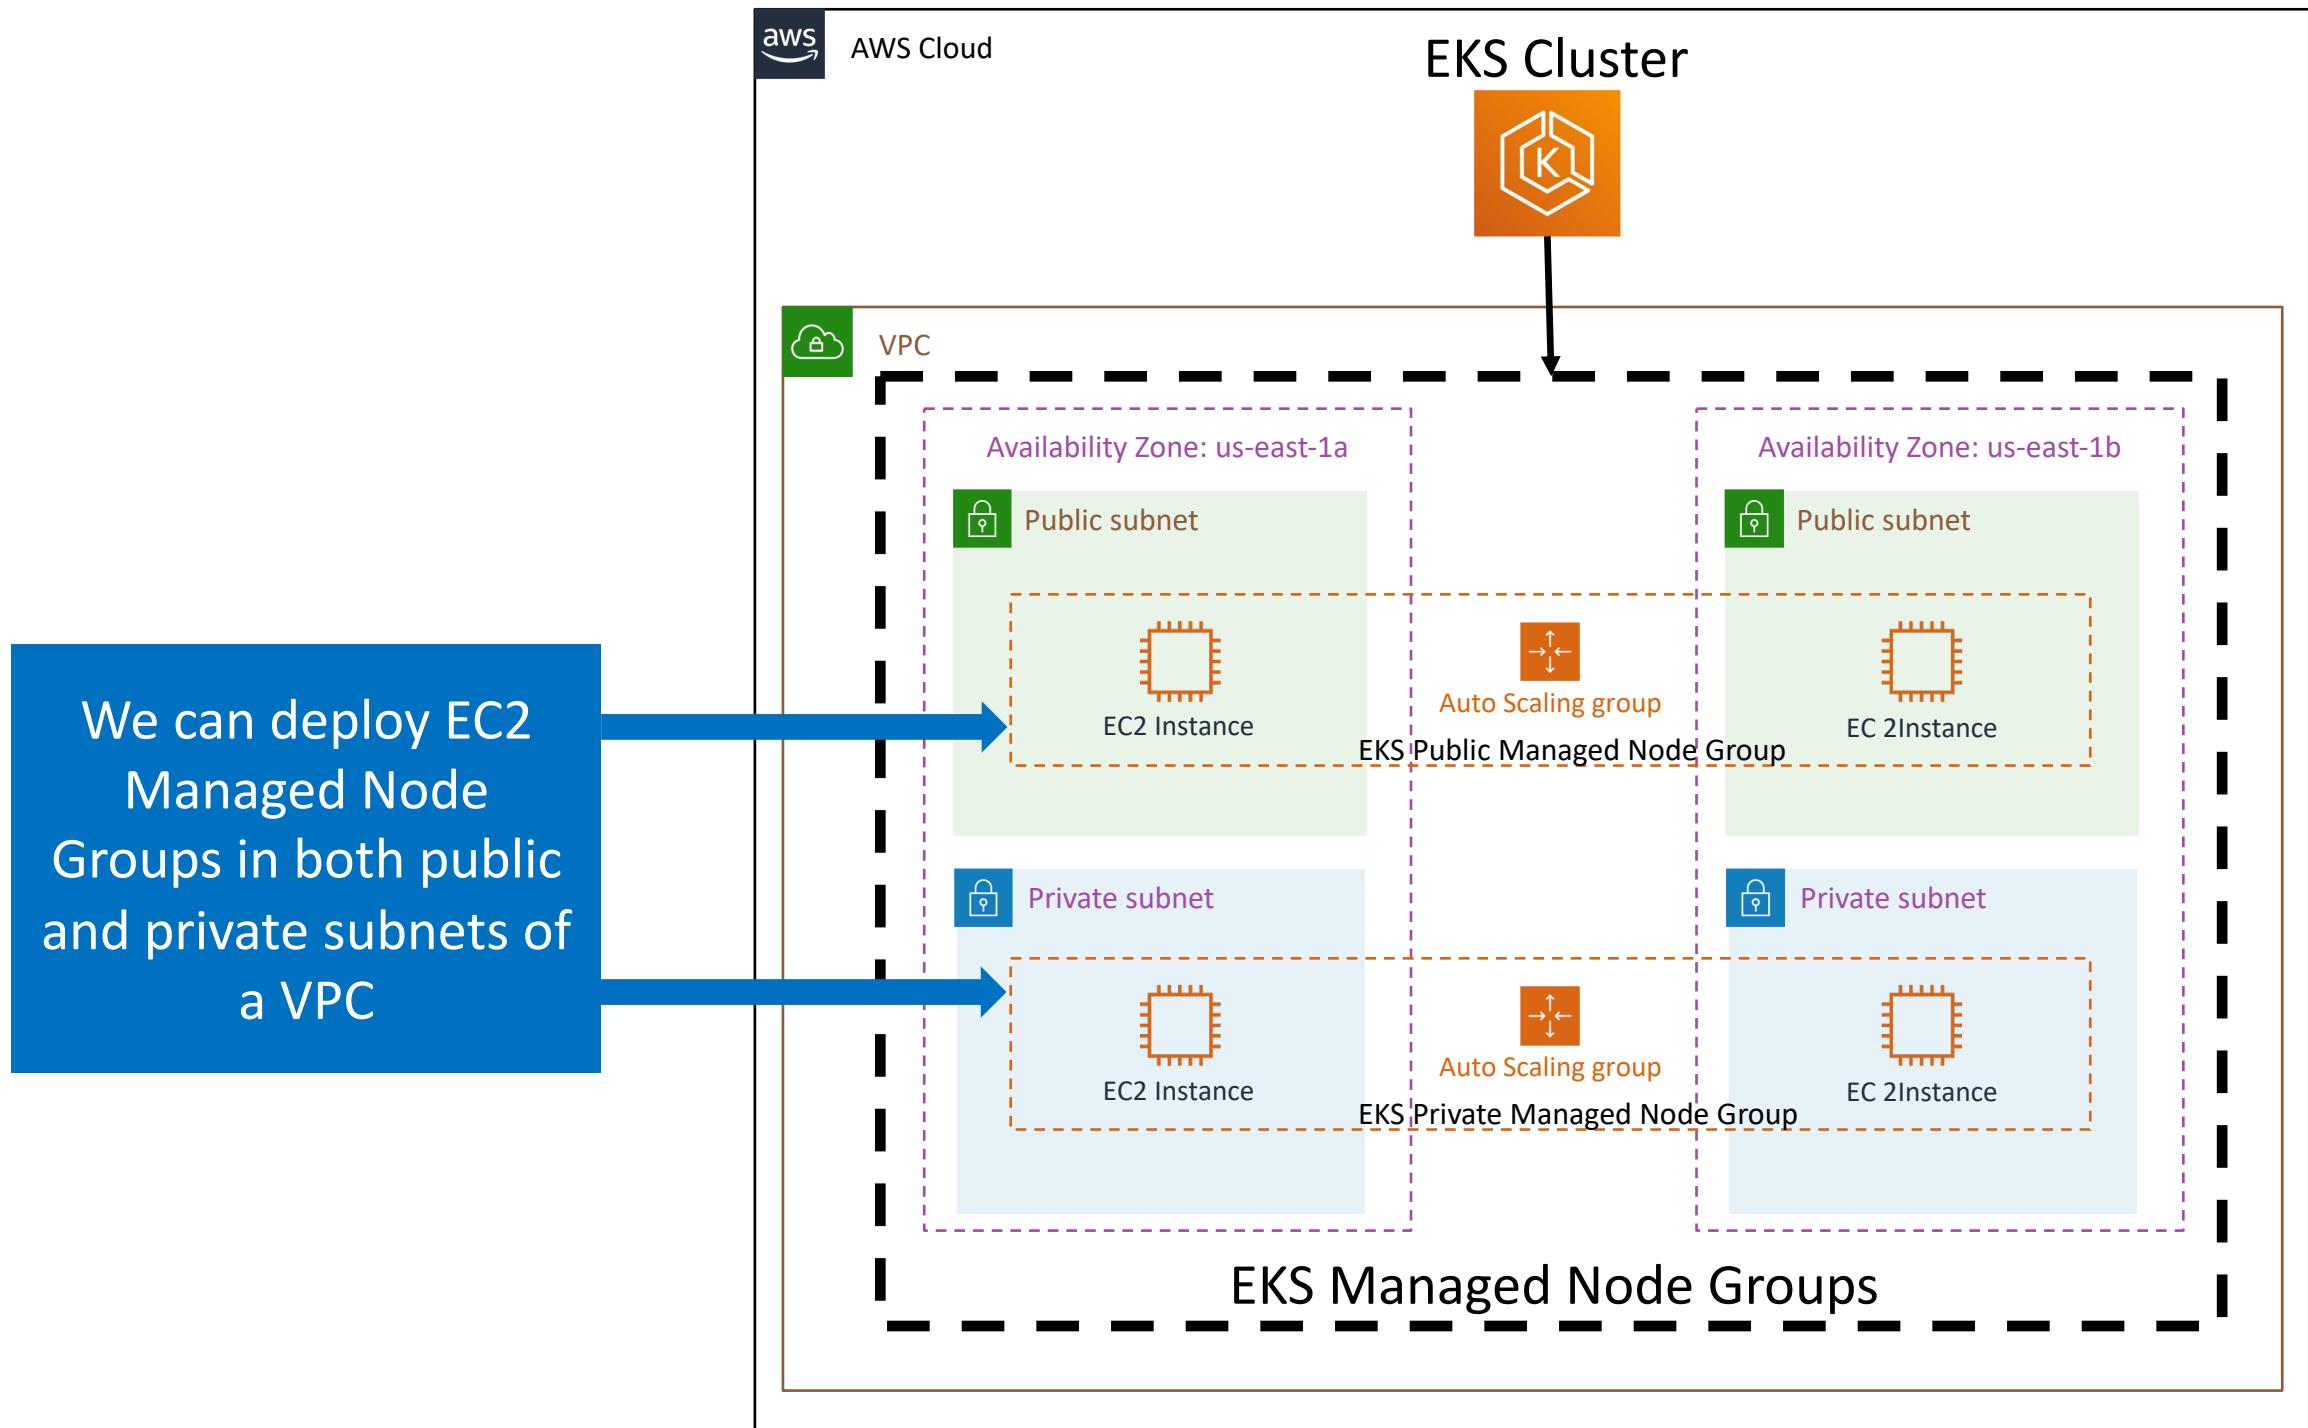
## Rightsized and Integrated

- Only pay for resources you need to run your pods
- Includes native AWS integrations for networking and security

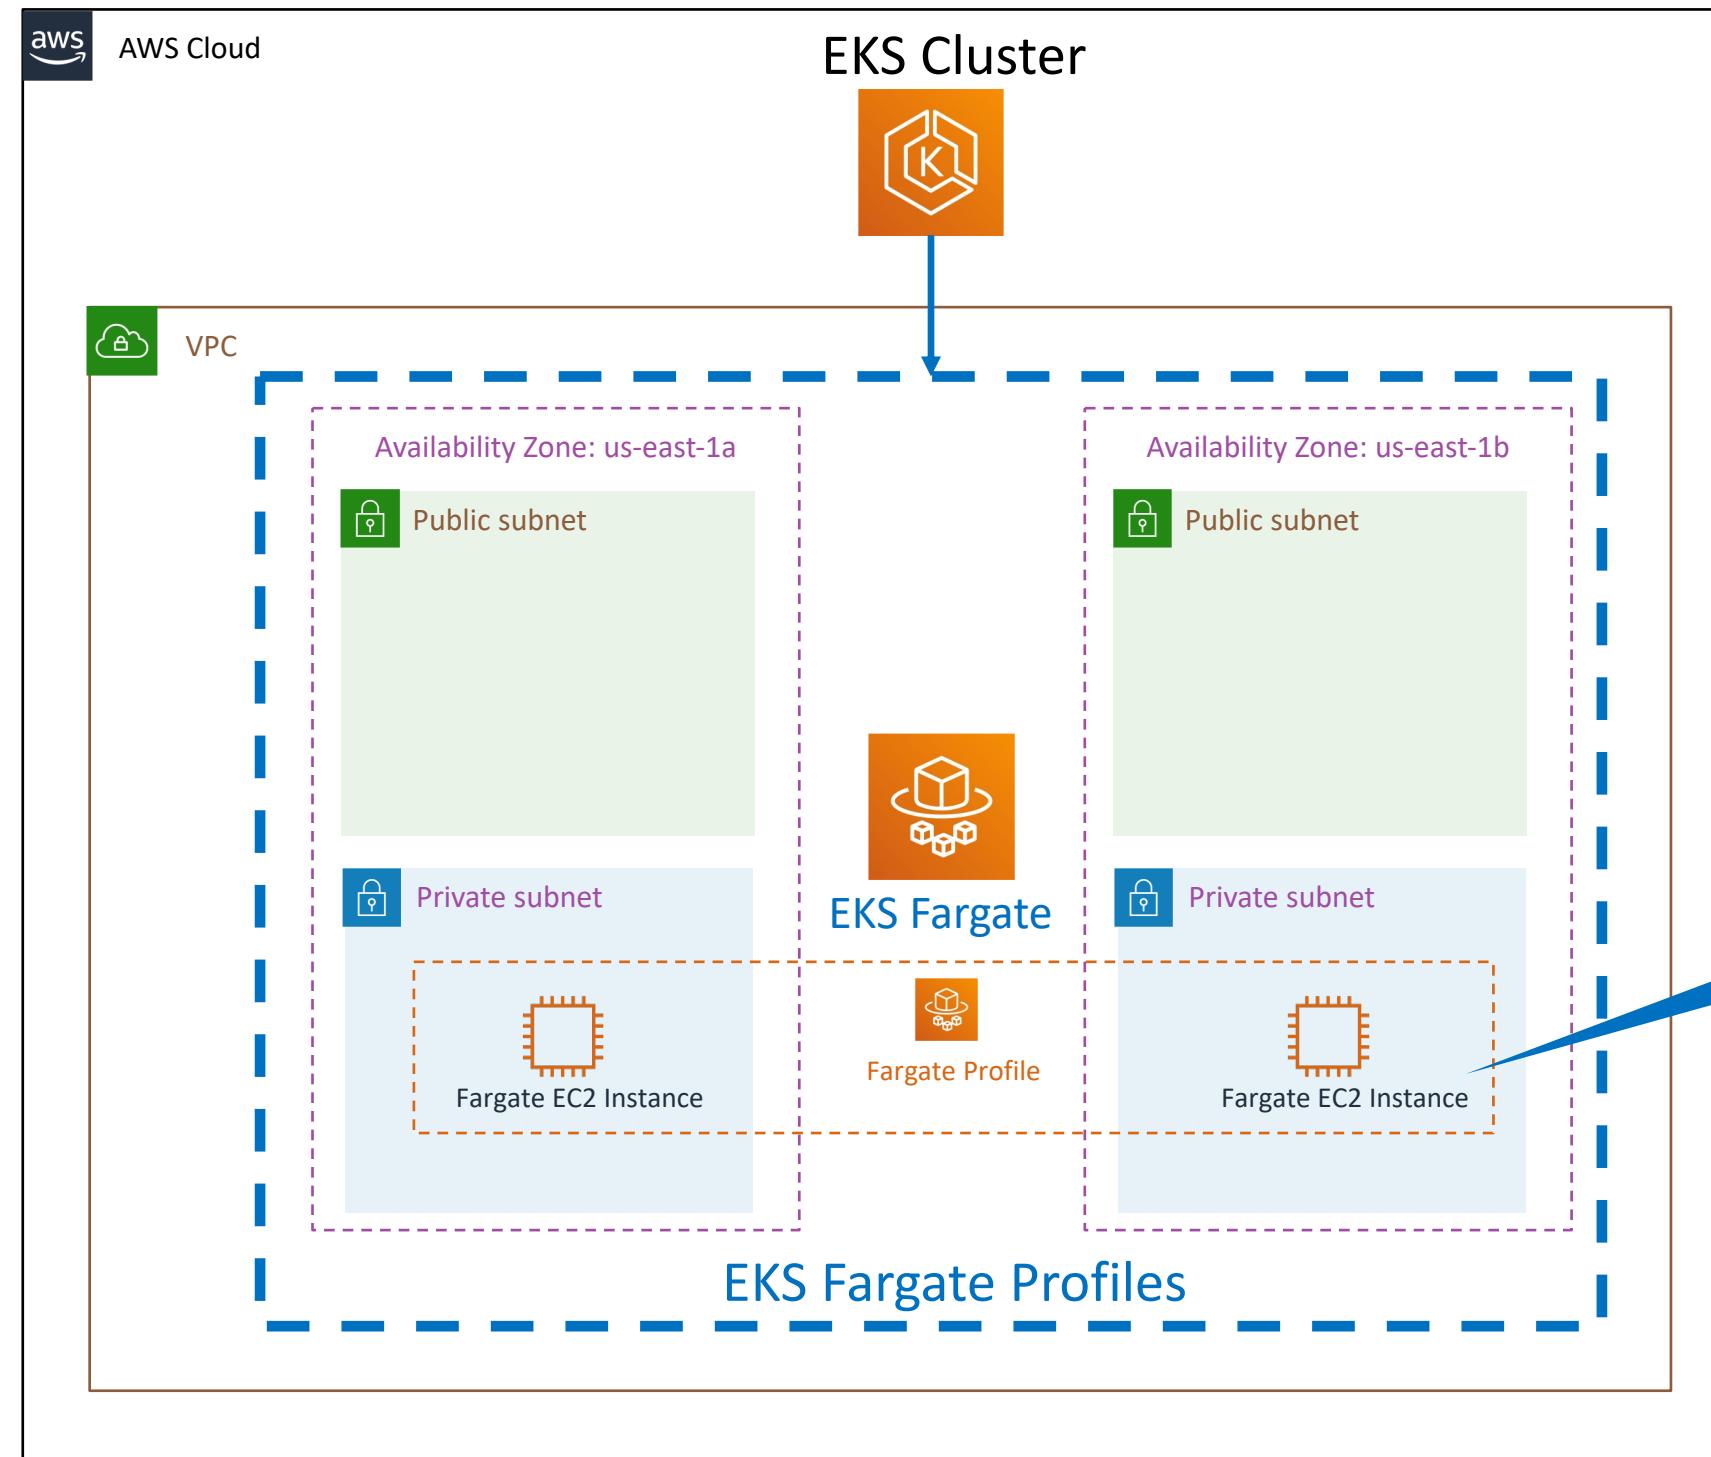
# EKS Deployment Options



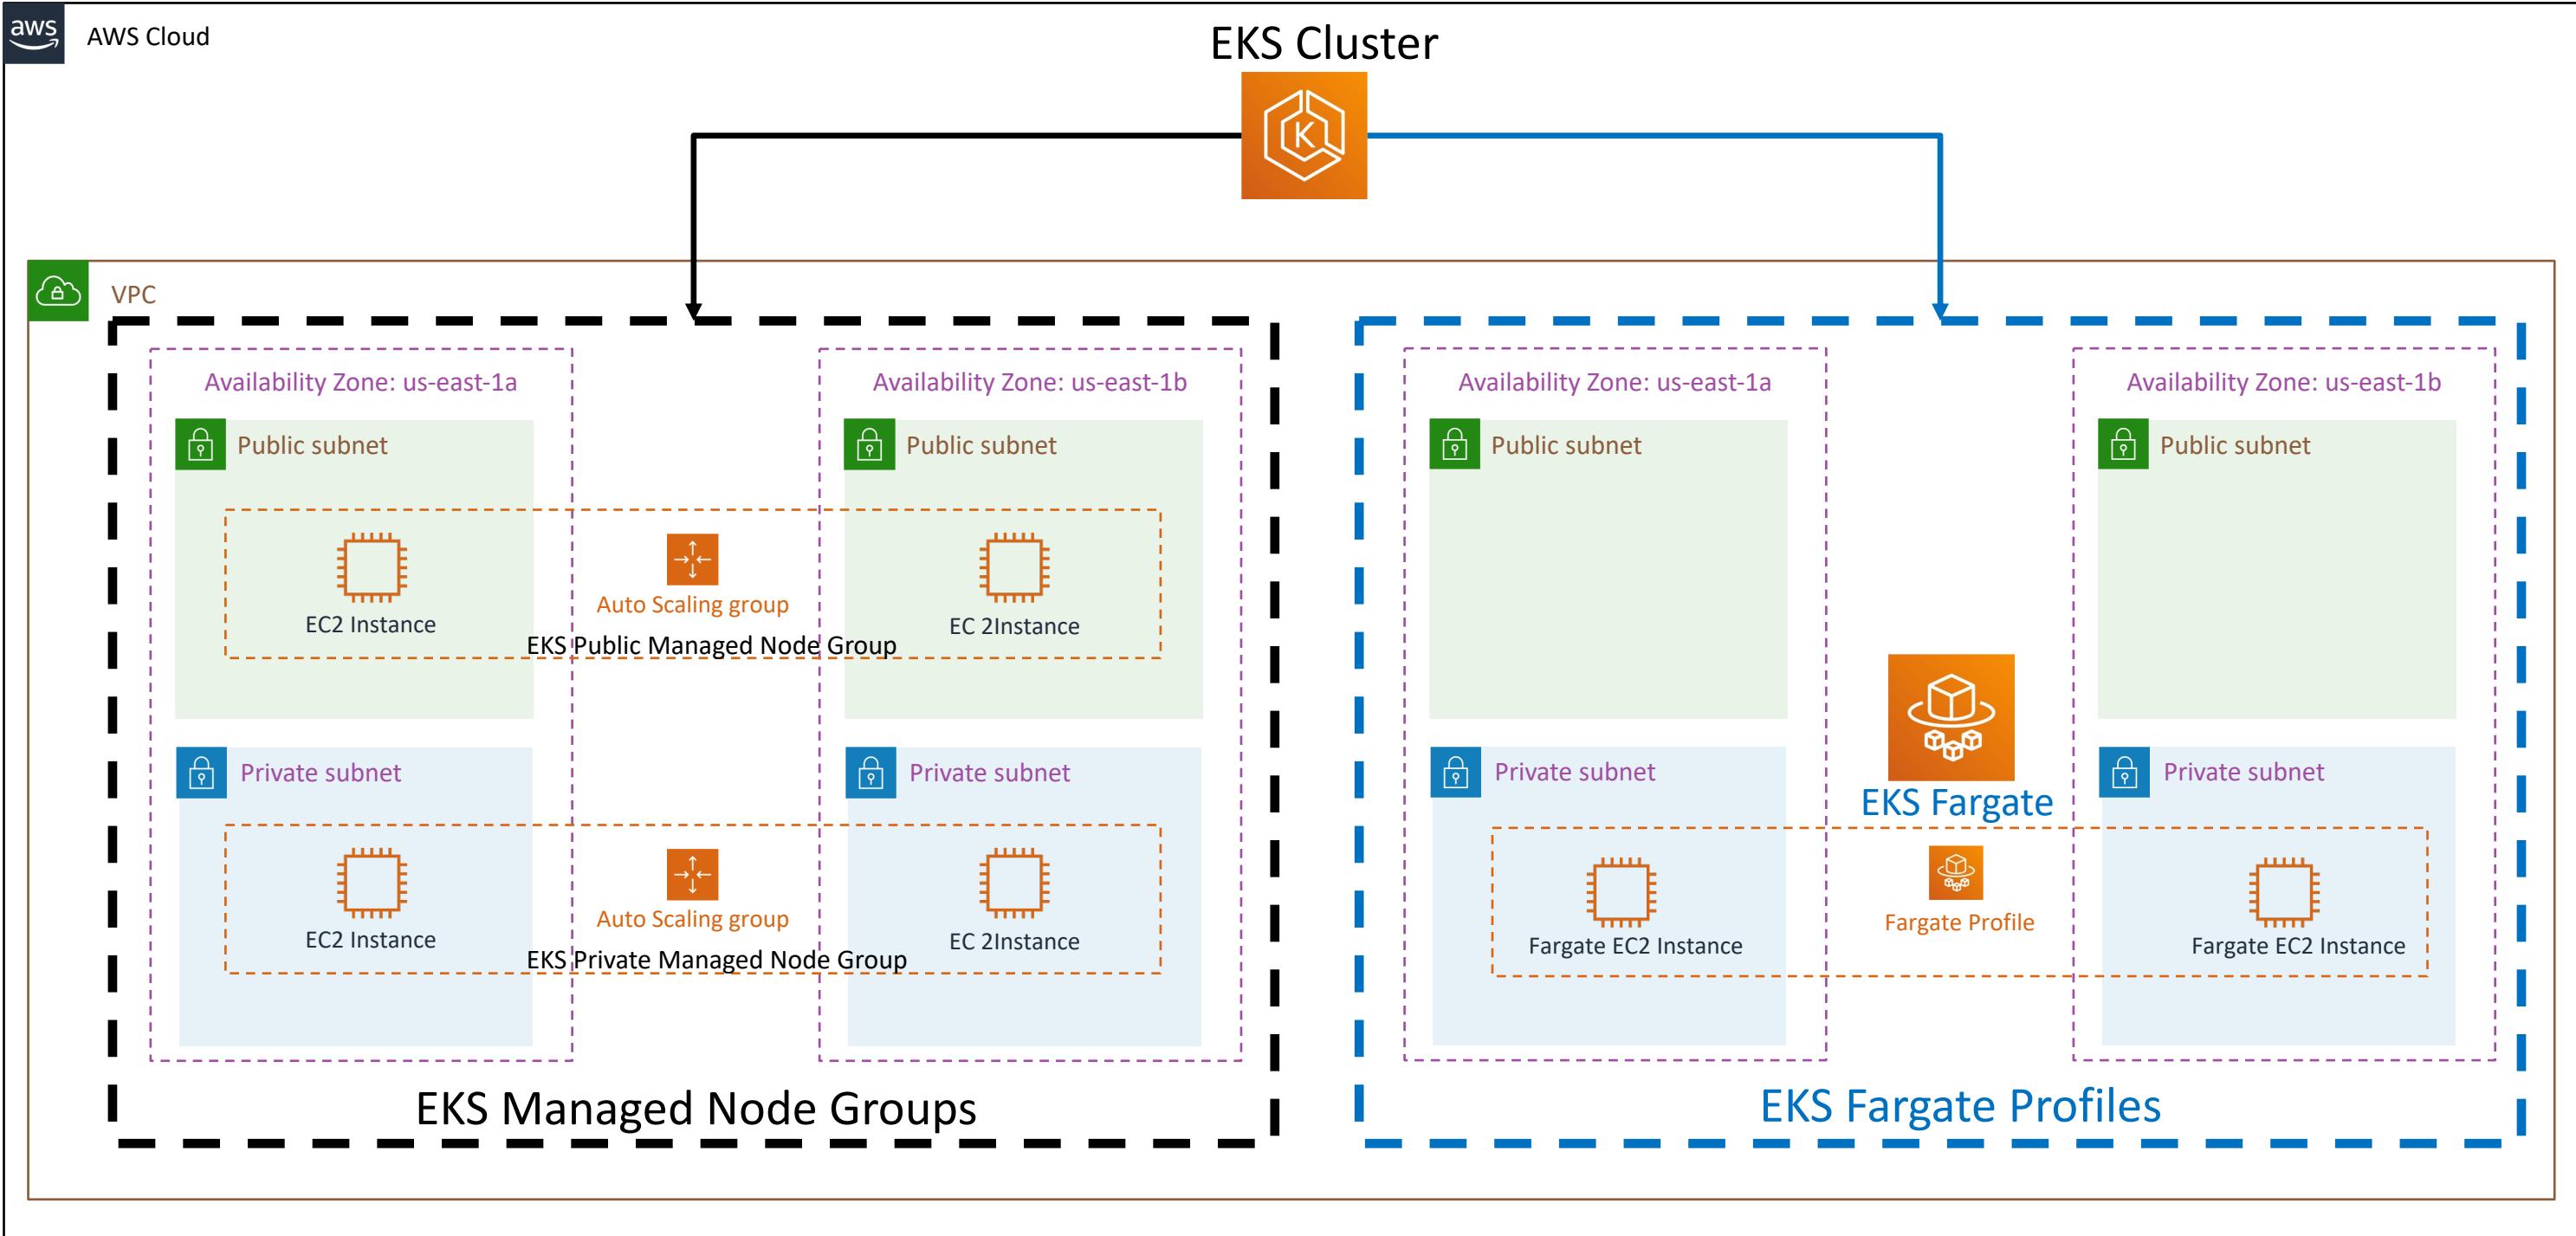
# EKS Deployment Options – EC2 Node Groups



# EKS Deployment Options – Only Fargate



# EKS Deployment Options - Mixed



# EKS Fargate vs Managed vs Unmanaged Nodes

	Fargate	Managed nodes	Unmanaged nodes
Units of work	Pod	Pod and EC2	Pod and EC2
Unit of charge	Pod	EC2	EC2
Host lifecycle	There is no visible host	AWS (SSH is allowed)	Customer
Host AMI	There is no visible host	AWS vetted AMIs	Customer BYO
Host : Pods	1 : 1	1 : many	1 : many

# EKS Fargate Considerations

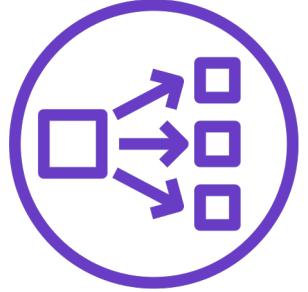
- There are many considerations we need to be **aware of** before we decide our Kubernetes workloads to run on Fargate.
- Documentation Link
- <https://docs.aws.amazon.com/eks/latest/userguide/fargate.html>



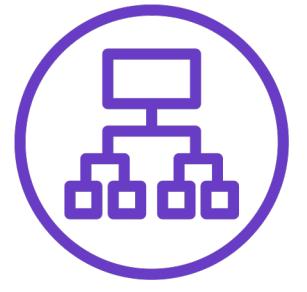
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



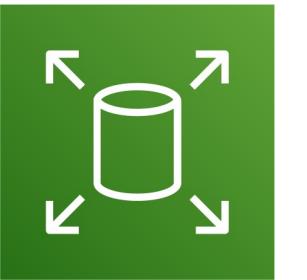
Application  
Load Balancer



Certificate  
Manager



Route53



Elastic Block  
Store



Amazon RDS



# AWS EKS

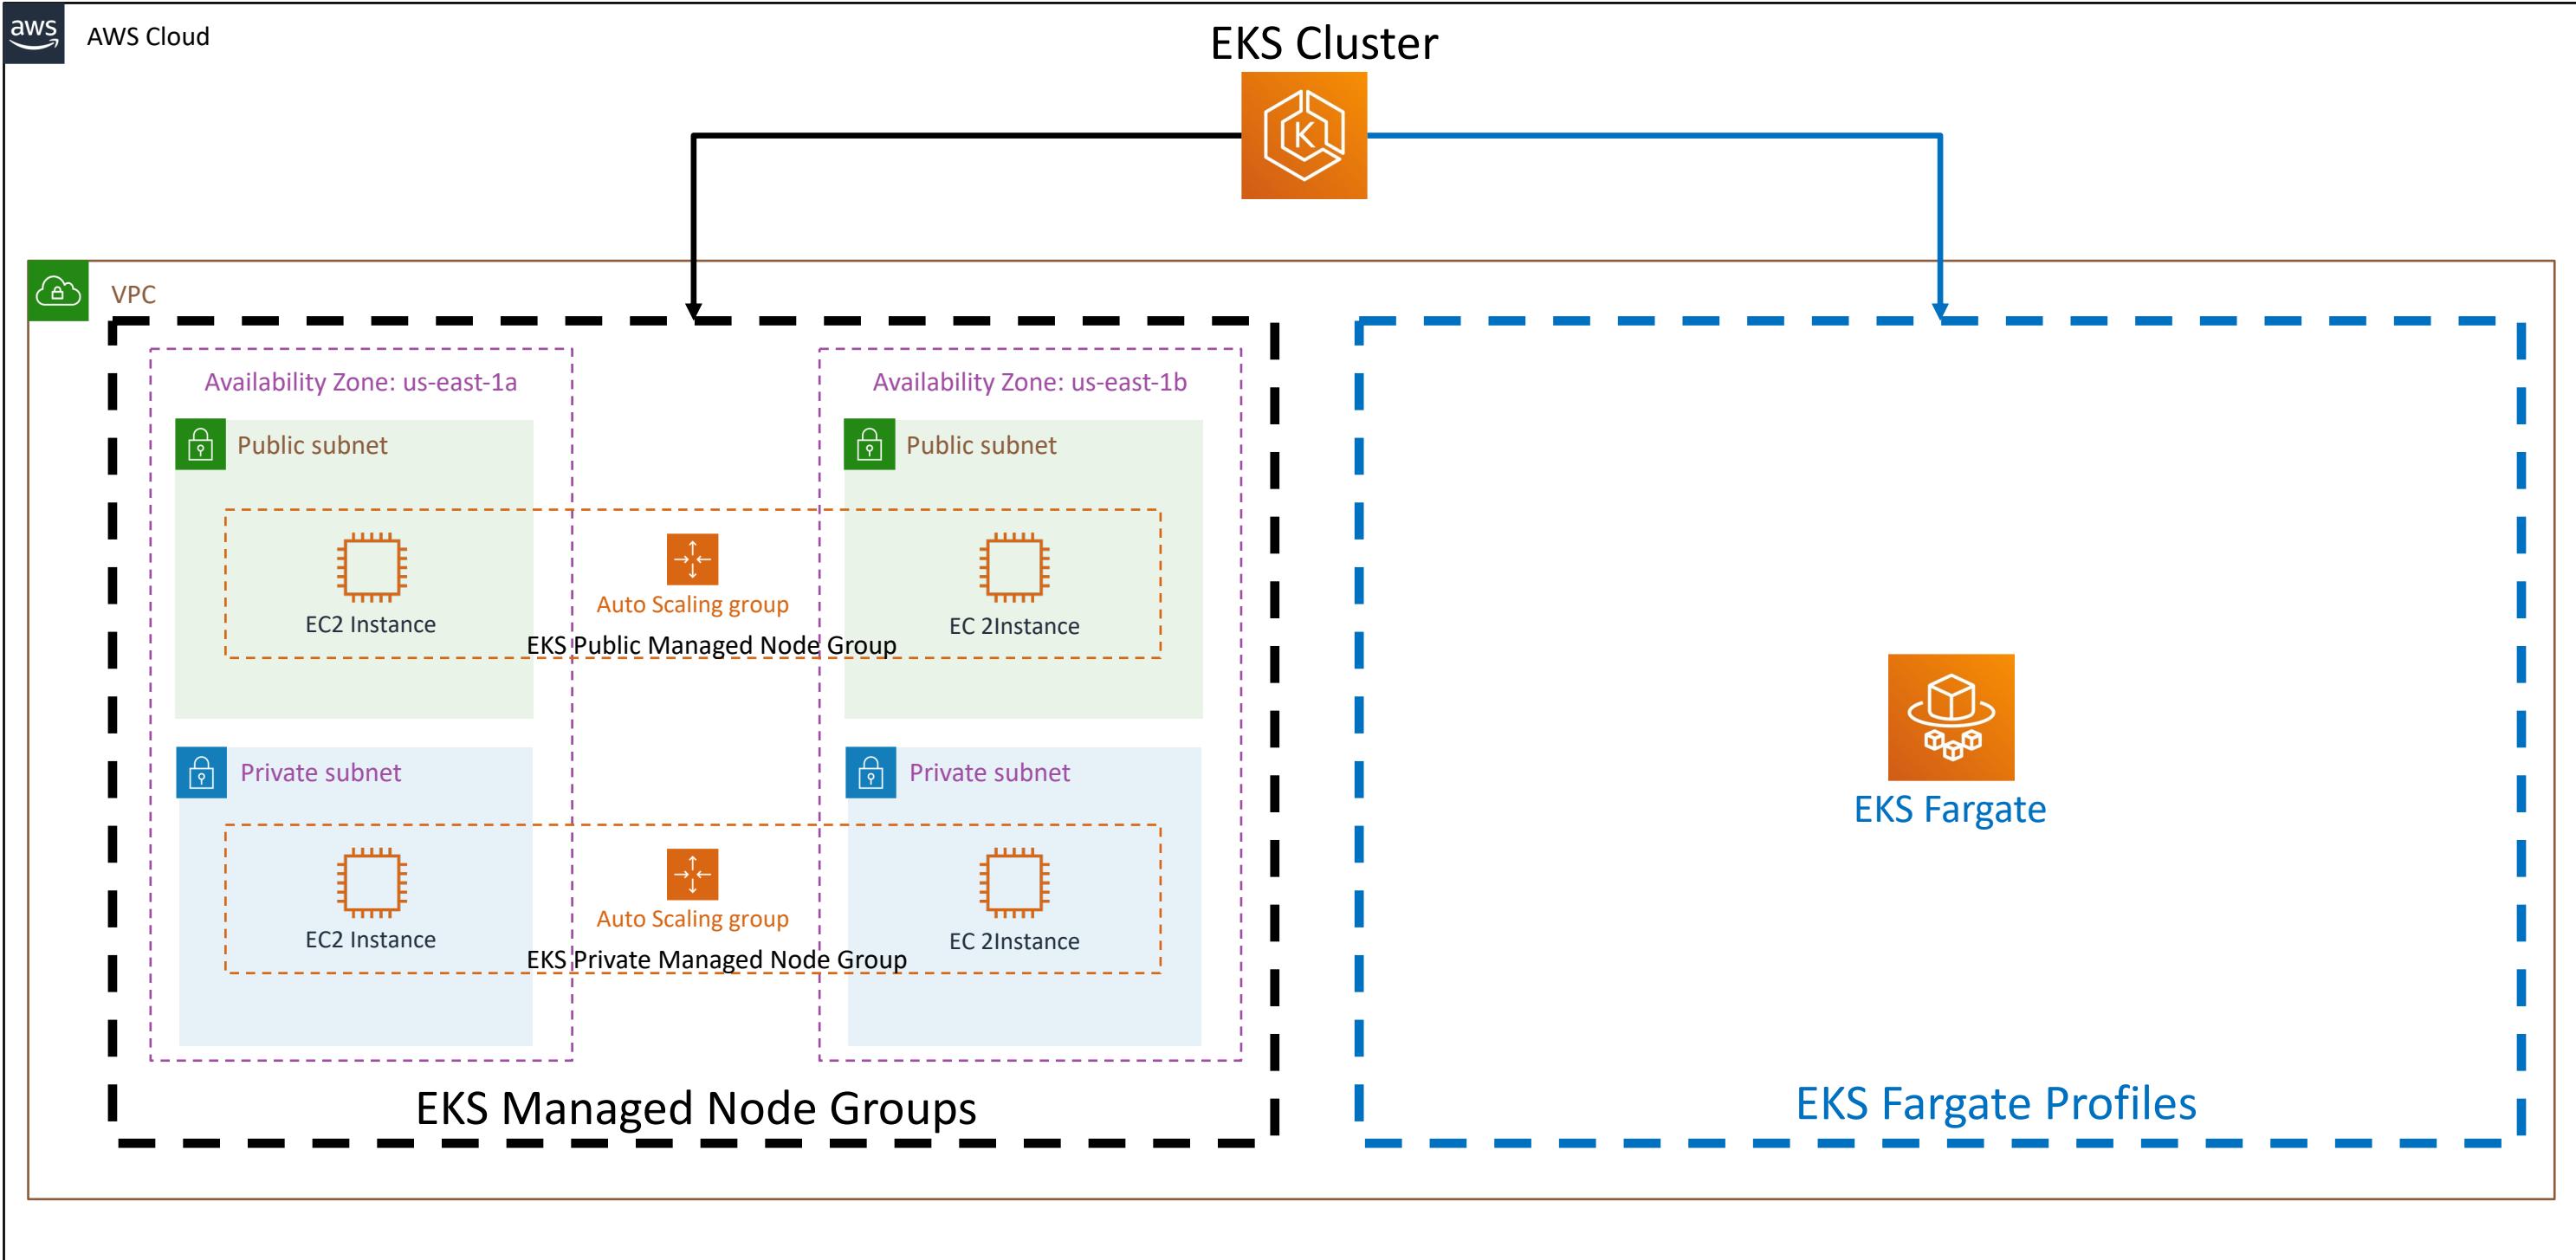
## Fargate Profiles

### Basics



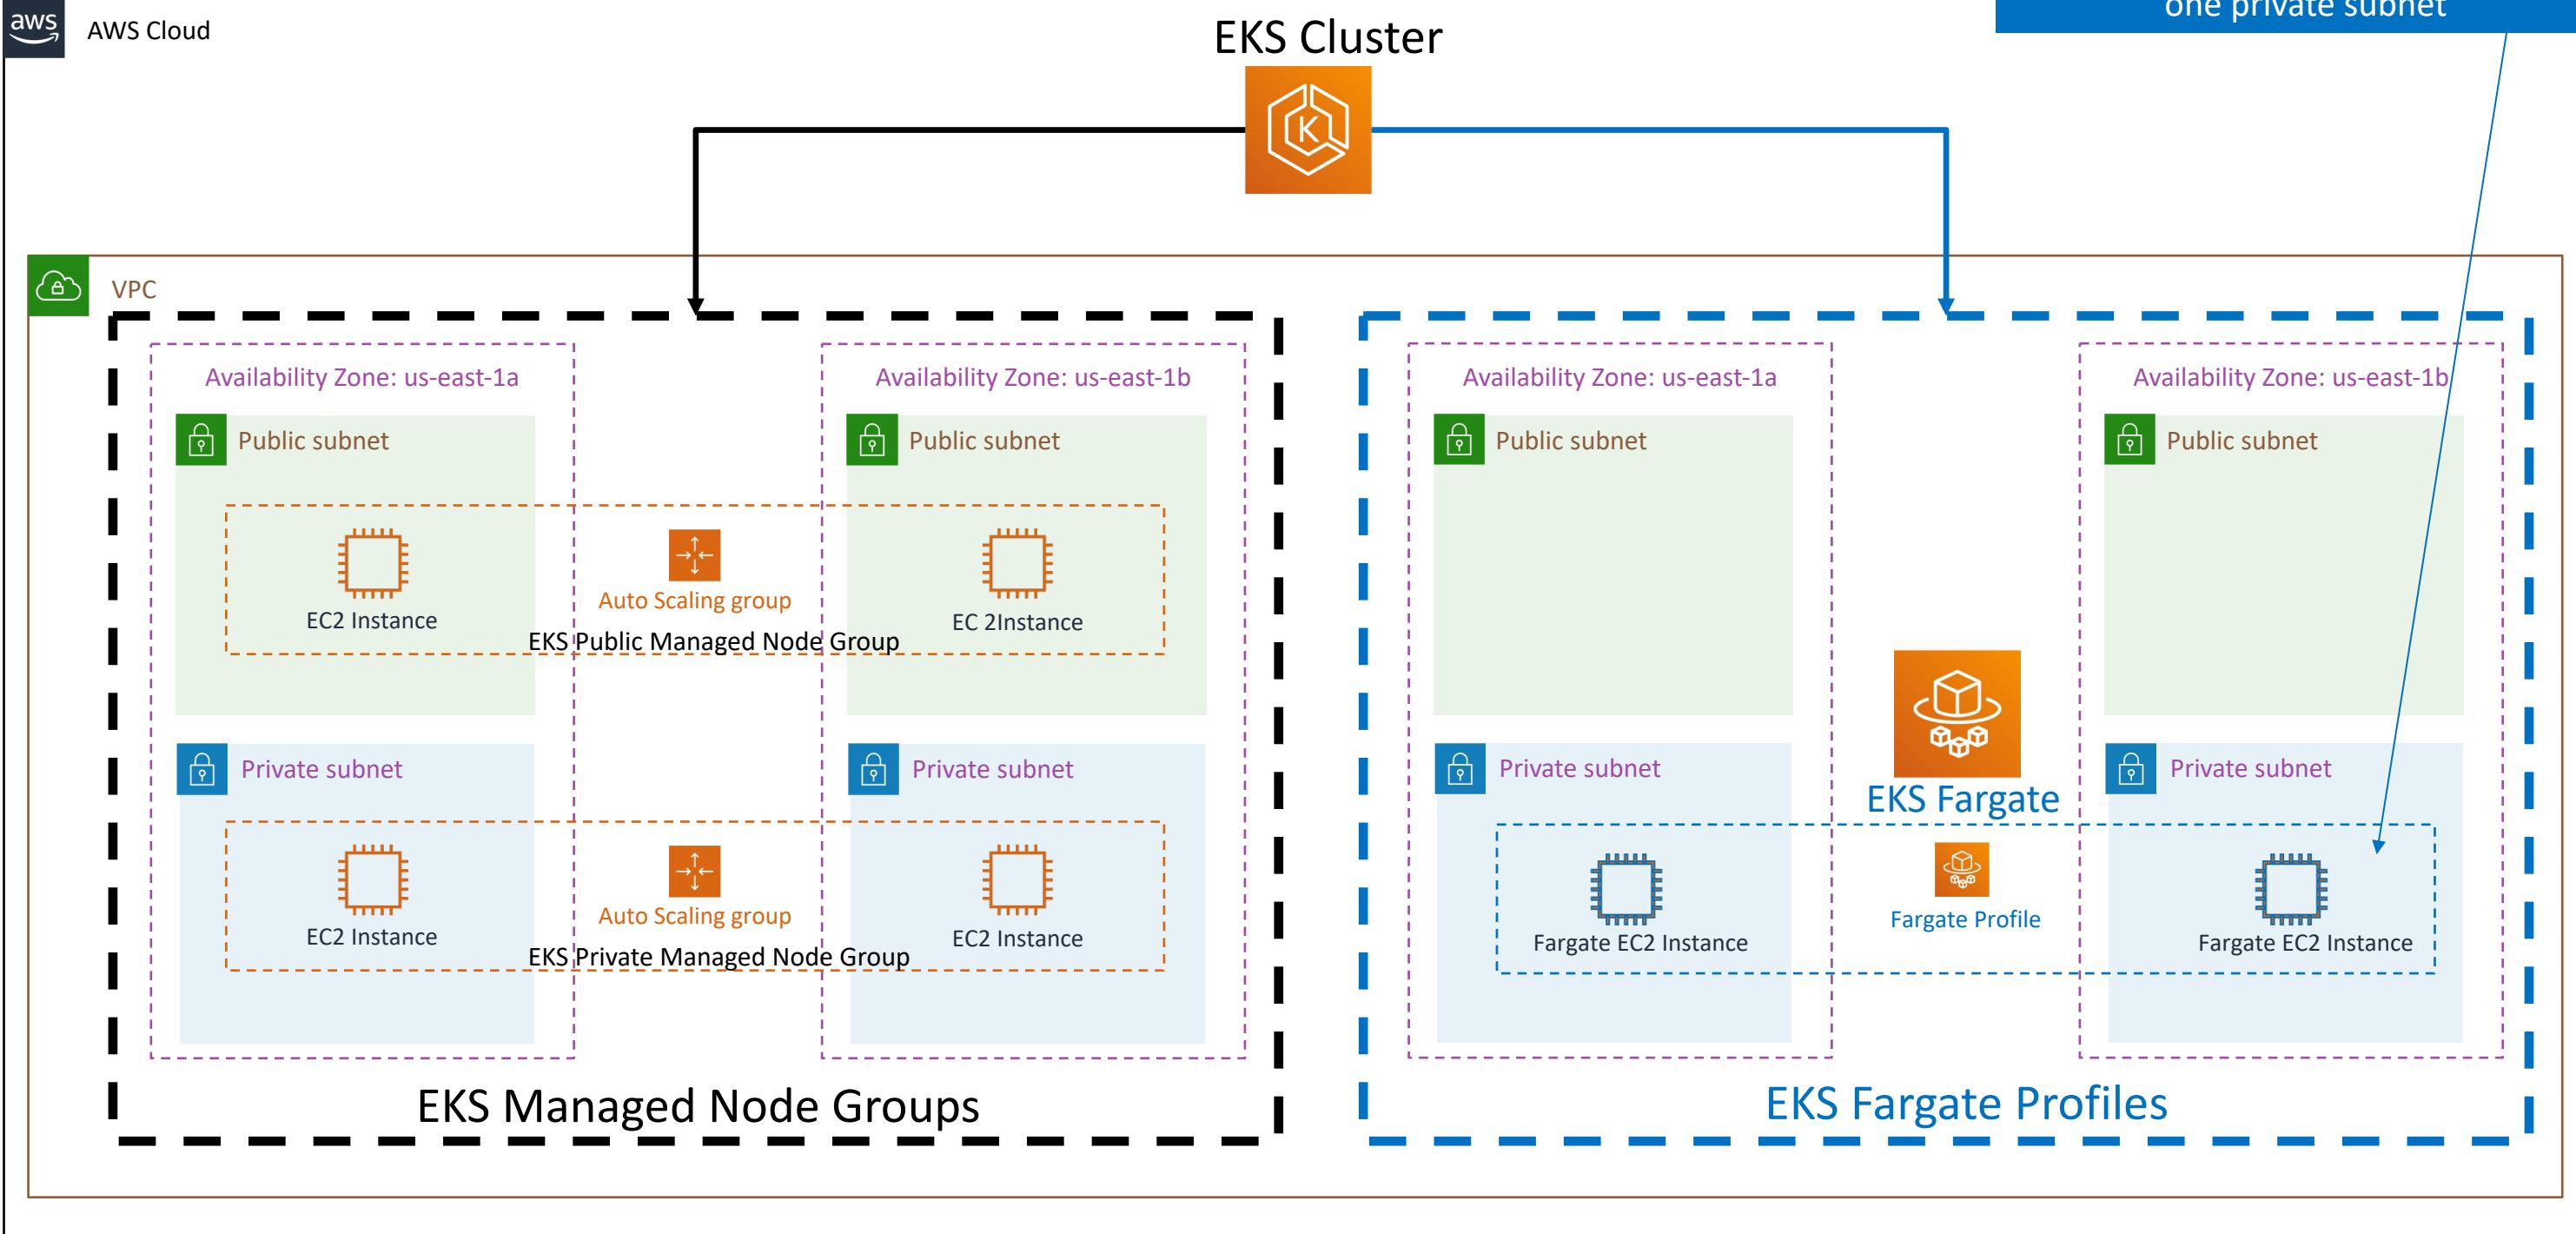
Fargate  
Profiles

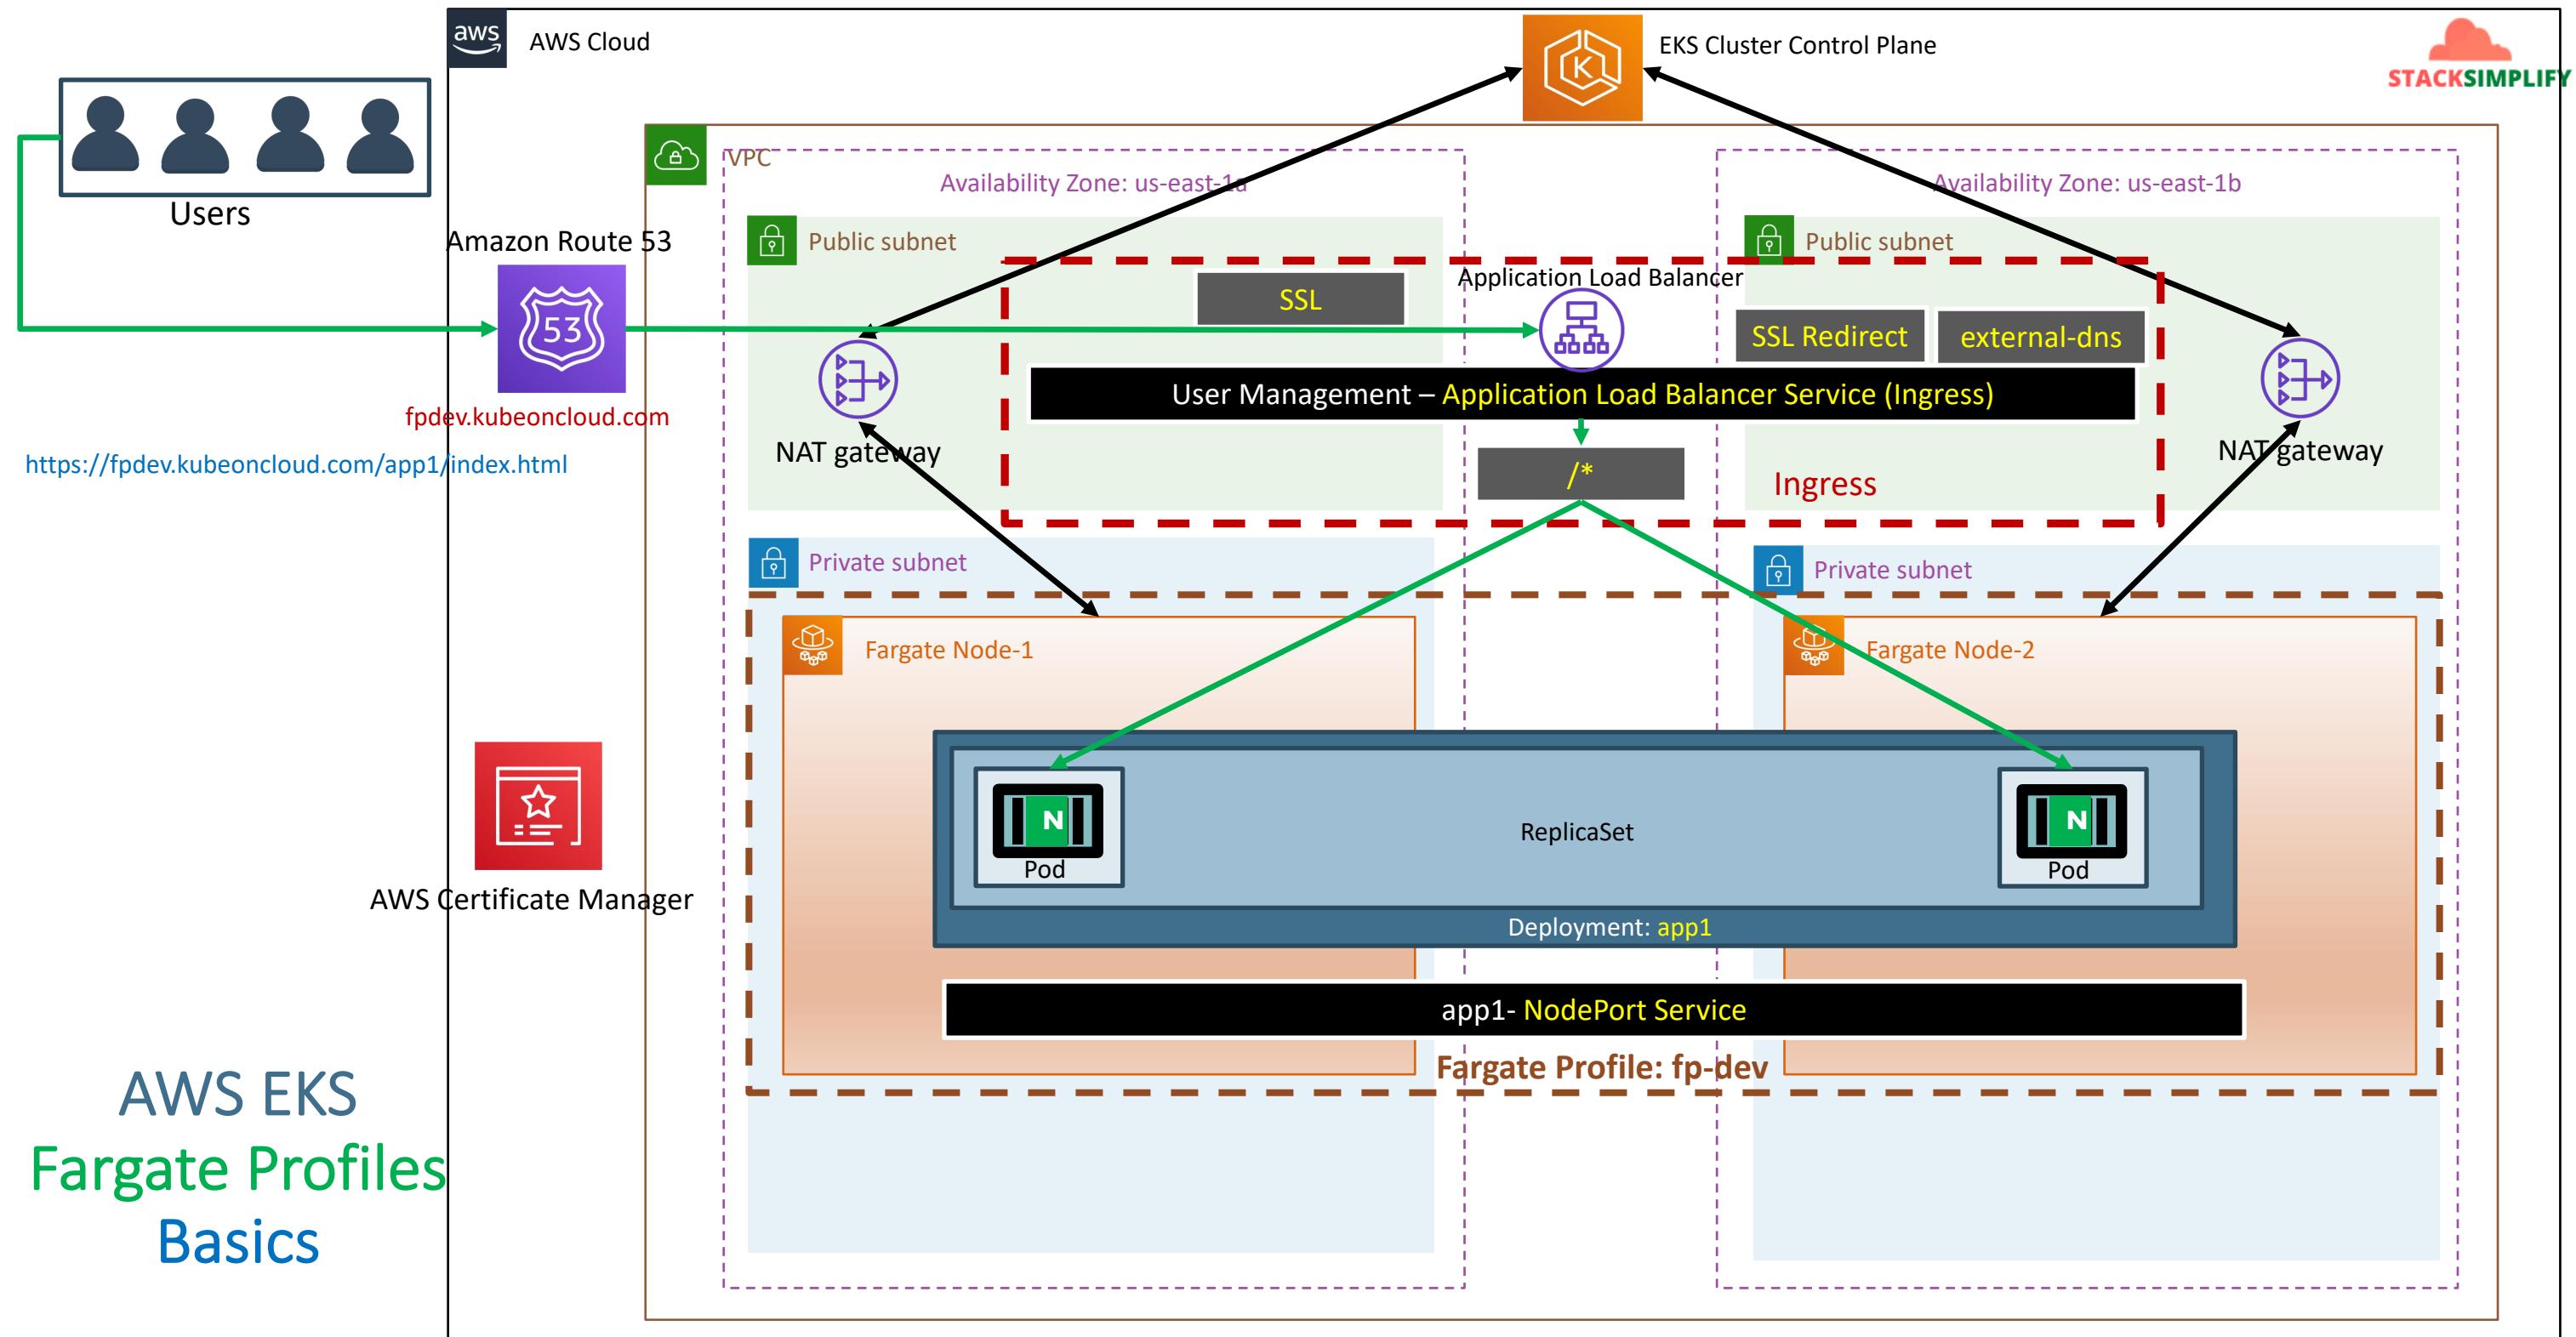
# EKS Deployment Options - Mixed

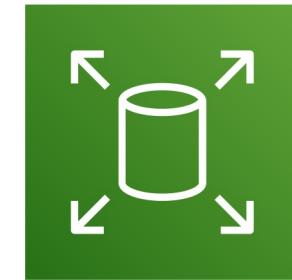
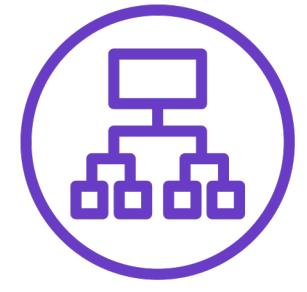
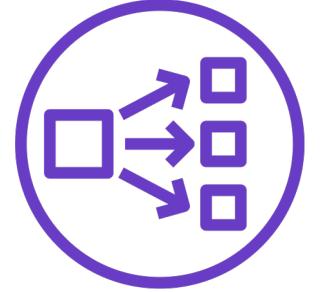


# EKS Deployment Options - Mixed

Fargate Profiles can be deployed to EKS Cluster only when we have at least one private subnet







Elastic Load  
Balancing

Classic  
Load Balancer

Network  
Load Balancer

Application  
Load Balancer

Certificate  
Manager

Route53

Elastic Block  
Store

Amazon RDS



# AWS EKS

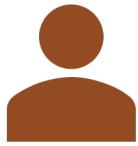
## Fargate Profiles

### Advanced with YAML



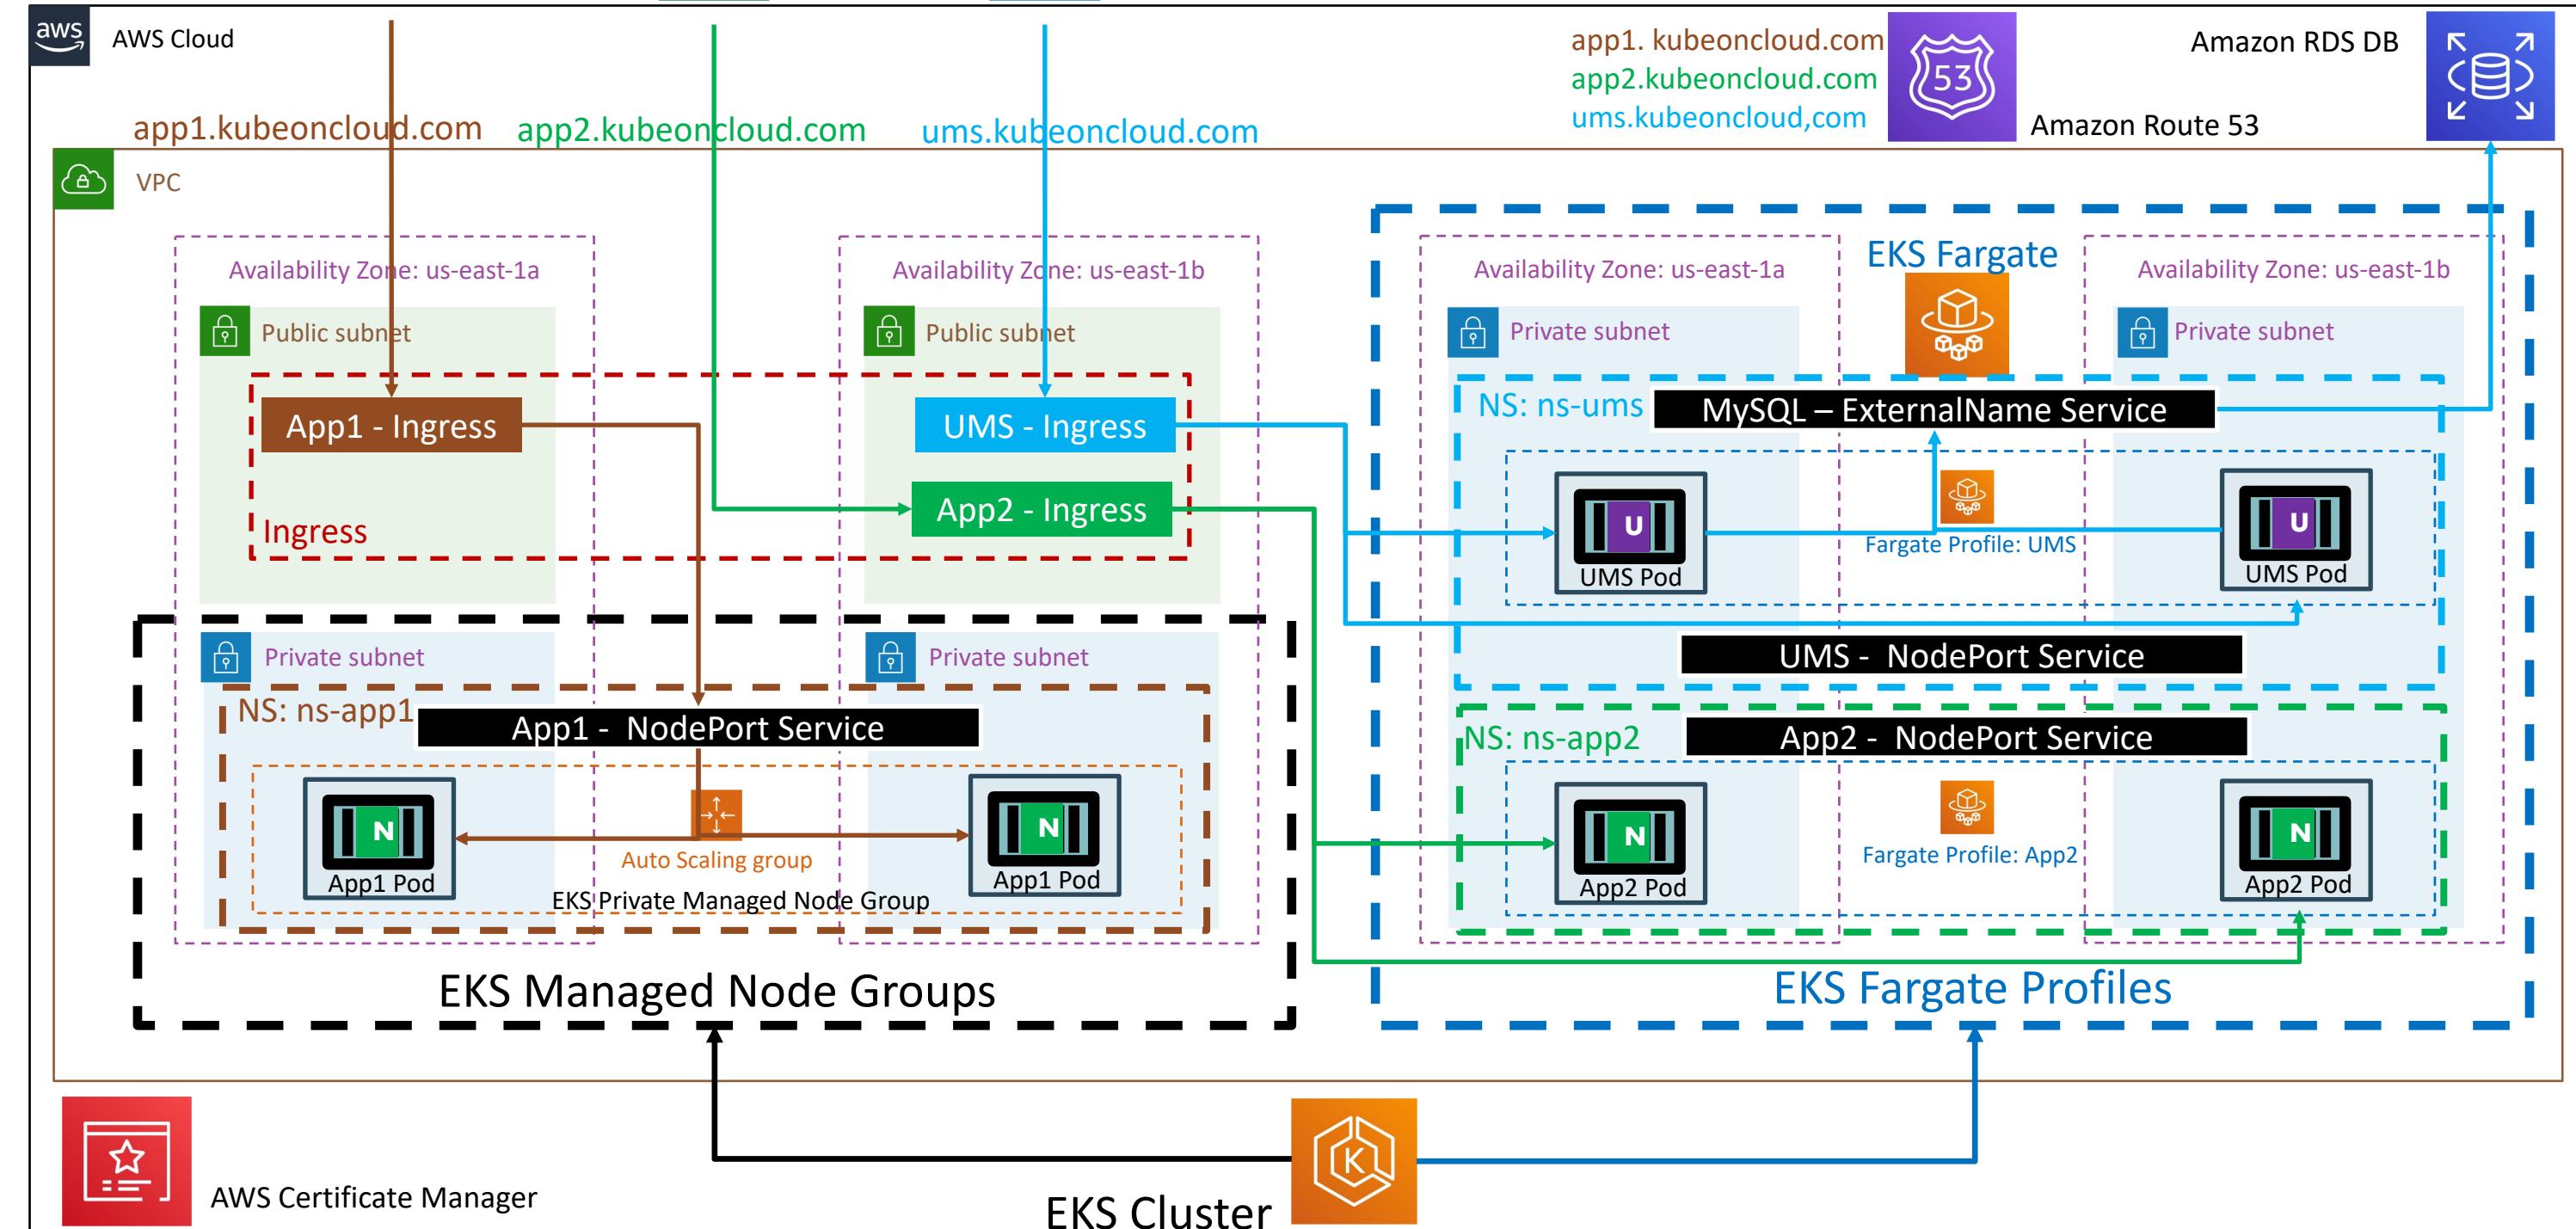
Fargate  
Profiles

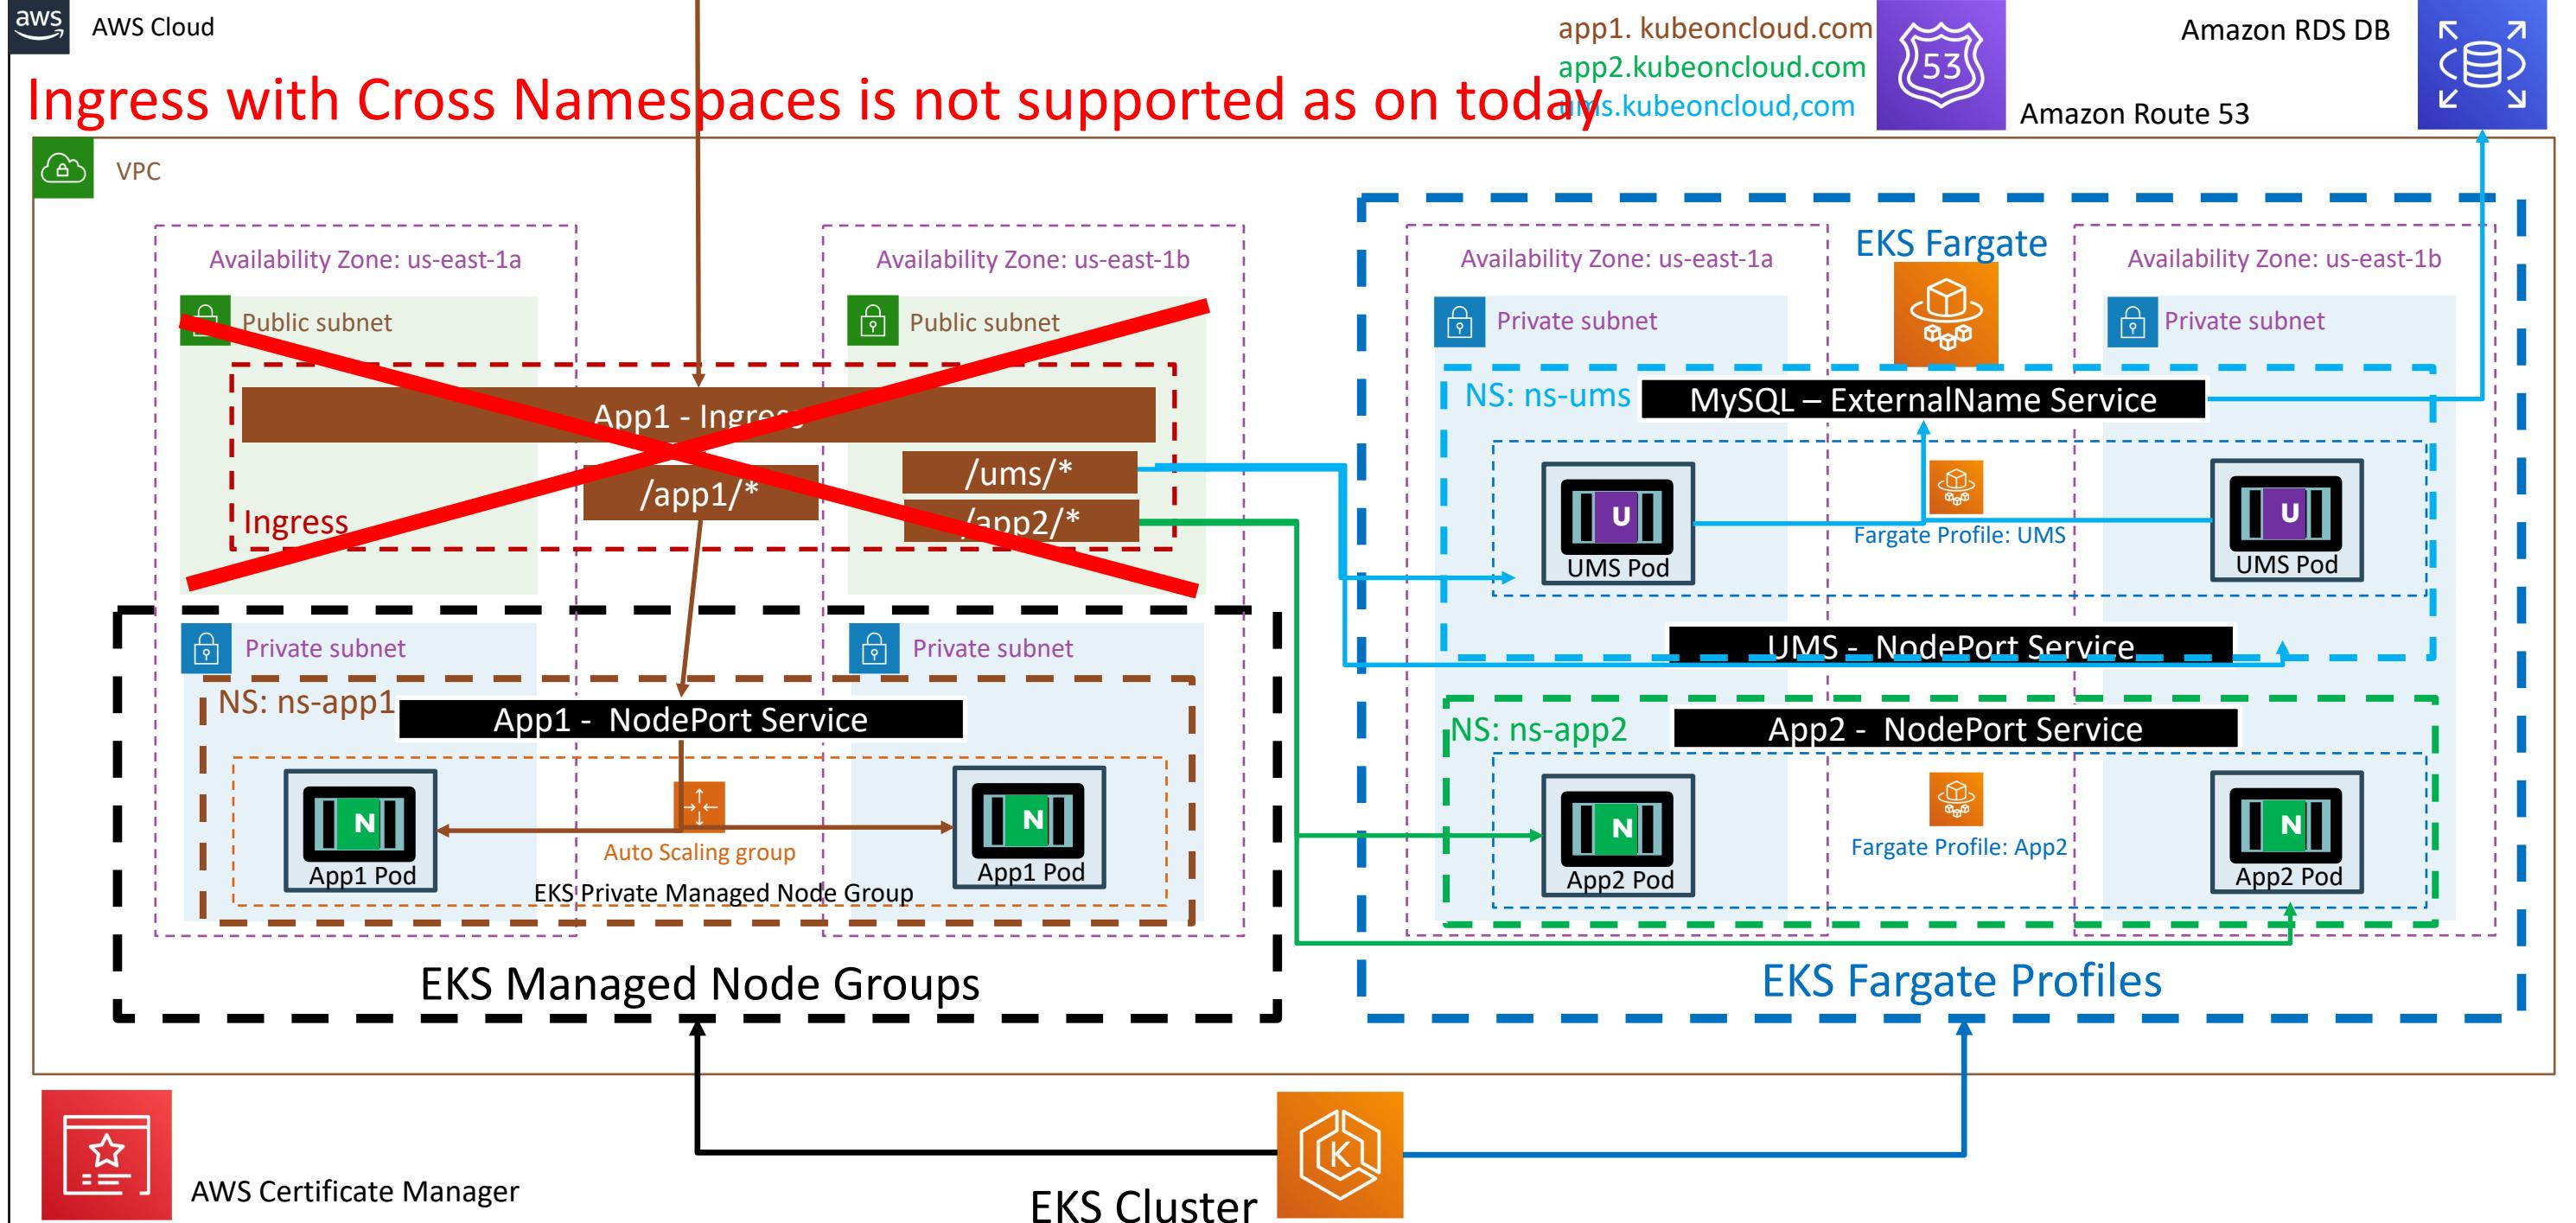
Users



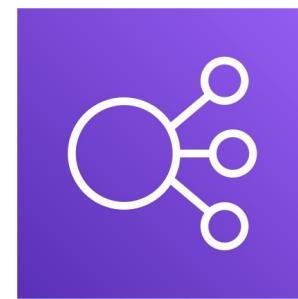
# EKS Deployment Options – Mixed Mode - 3 Apps

StackSimplify

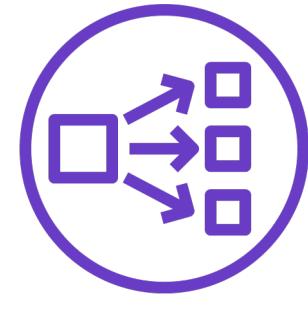




# AWS Load Balancer Controller & AWS Network Load Balancer – NEW TOPIC ADDITION - START



Elastic Load  
Balancing



Network  
Load Balancer



Kubernetes  
Service



External  
DNS



# AWS EKS

## Load Balancer Controller

## Kubernetes Service

## Network Load Balancer



Demo-1: NLB Basics

Demo-2: NLB TLS  
SSL

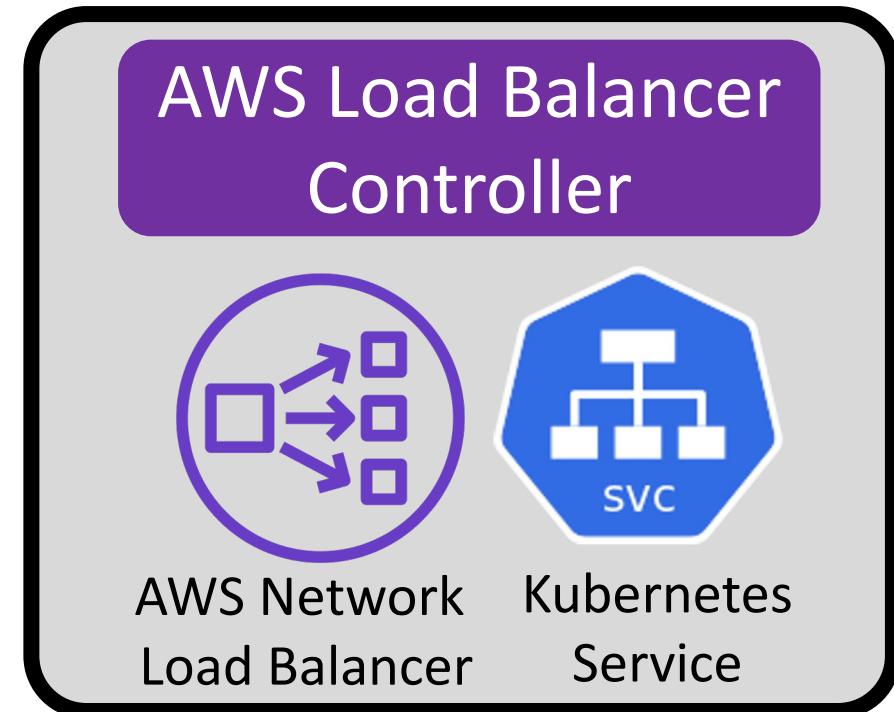
Demo-3: NLB &  
External DNS

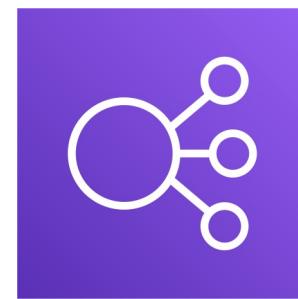
6 Demos

Demo-4: NLB &  
Elastic IP

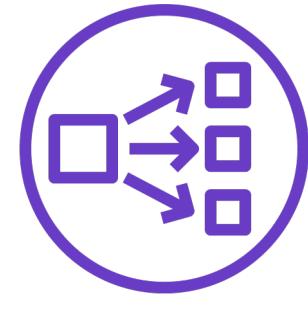
Demo-5: Internal  
NLB

Demo-6: NLB &  
Fargate





Elastic Load  
Balancing



Network  
Load Balancer



Kubernetes  
Service



External  
DNS



# AWS EKS

## Load Balancer Controller

### Kubernetes Service

### Network Load Balancer - Basics



# Network Load Balancer

## What is Network Load Balancer ?

Network traffic is load balanced at L4 of the OSI model.

Supports TCP, UDP and TLS Protocols

**ELB Features Comparison for Additional Understanding:**

<https://aws.amazon.com/elasticloadbalancing/features/>

# Options for creating Network Load Balancer using Kubernetes on EKS Cluster

## AWS cloud provider load balancer controller

Legacy Controller which can create AWS Classic LB and Network LB

Creates very basic Network Load Balancer

This controller will receive only critical bug fixes and in long run it will be deprecated

This controller will not have any new features added to it

- ✓ 07-ELB-Classic-and-Network-LoadBalancers
  - > 07-01-Create-EKS-Private-NodeGroup
  - > 07-02-Classic-LoadBalancer-CLB
  - > 07-03-Network-LoadBalancer-NLB

## AWS Load Balancer Controller

Latest and greatest which can create AWS ALB and NLB

Supports Protocols TCP, UDP & TLS (SSL) and for Health Check it supports TCP, HTTP & HTTPS

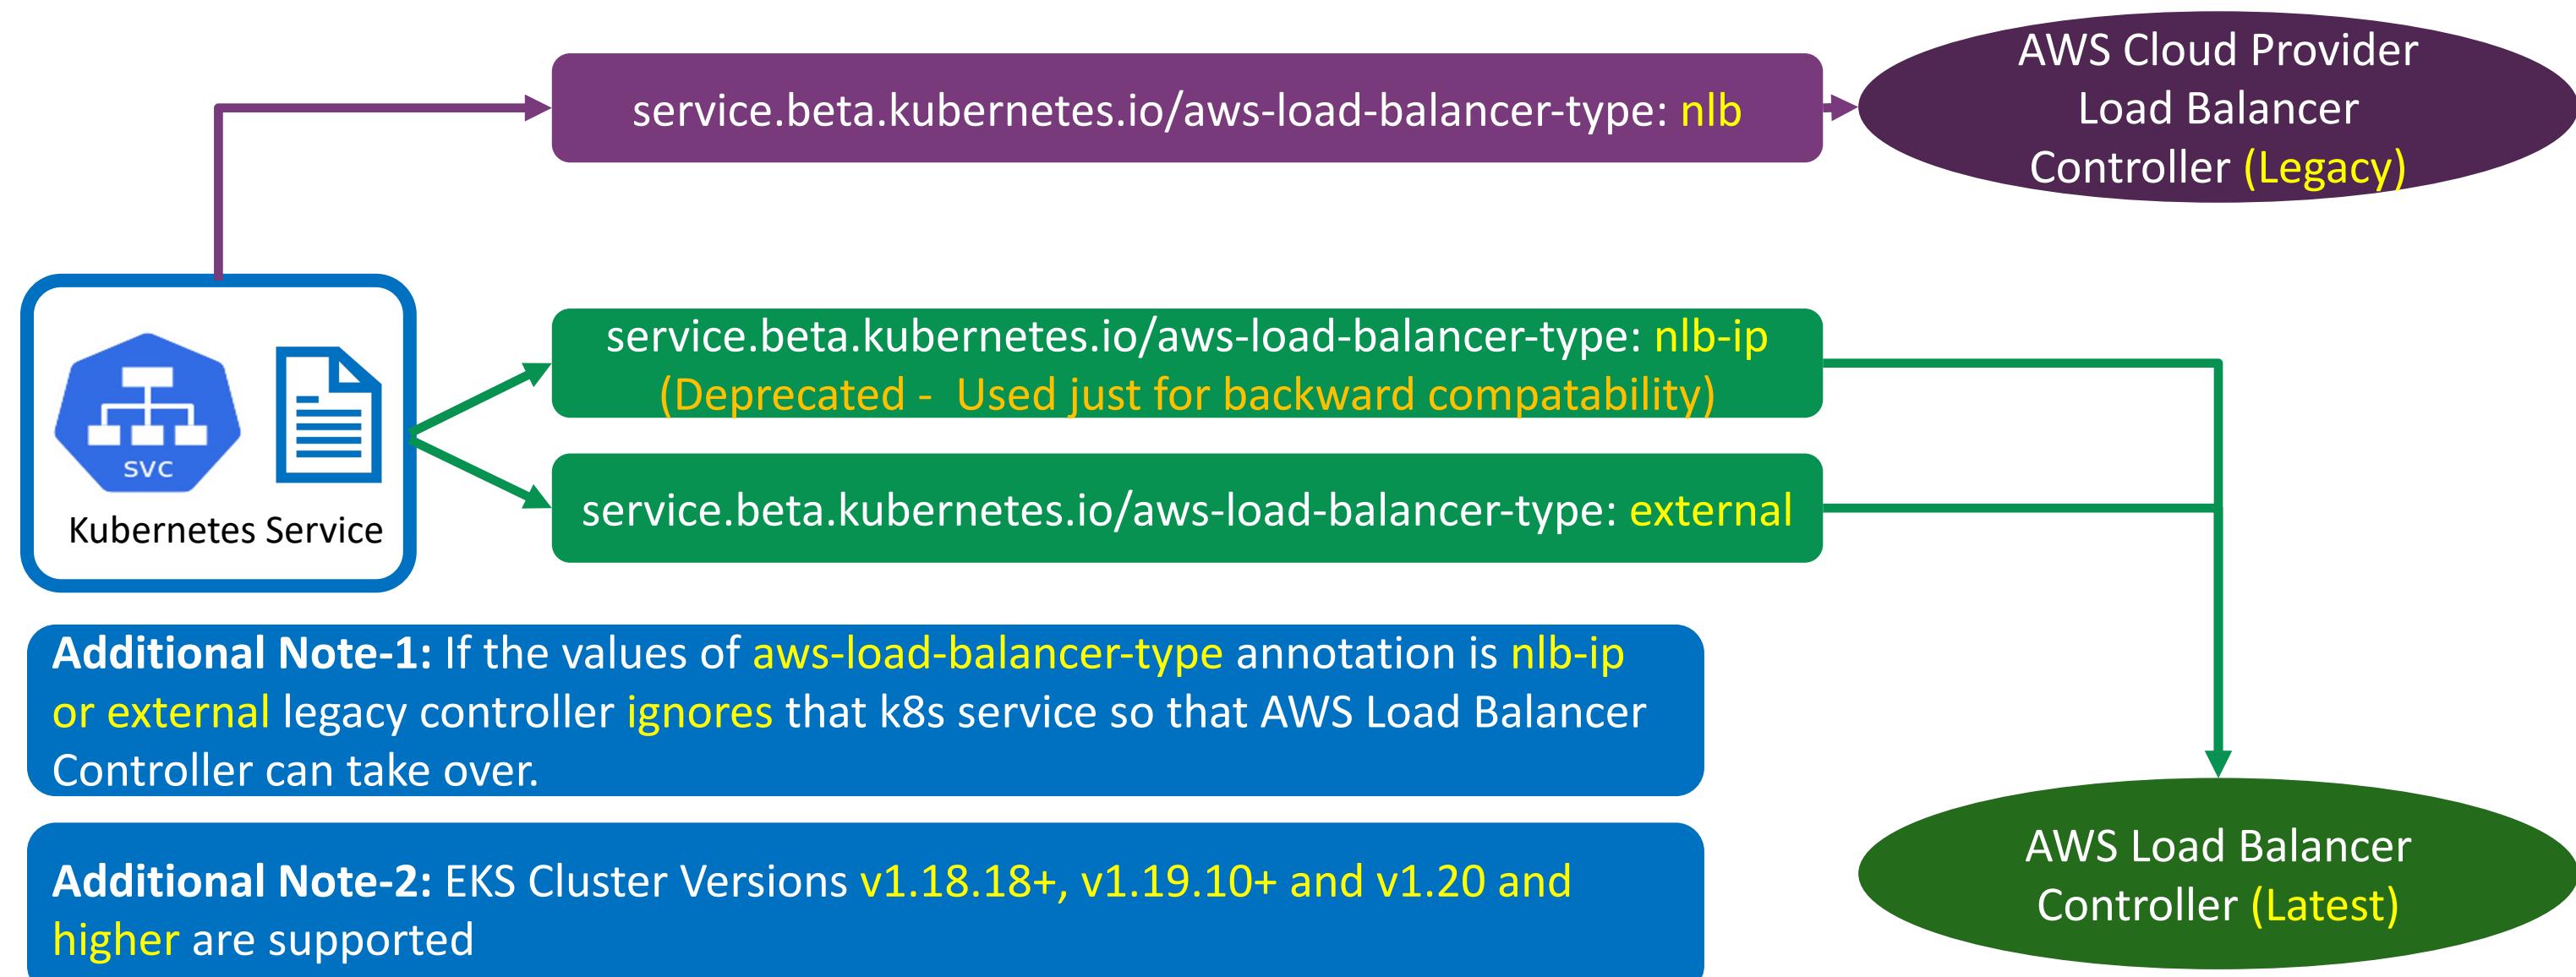
Supports Instance and IP Mode(For Fargate)

Supports Internal, External NLB and Custom Subnet Discovery

Supports Static Elastic IP and Access Control

Supports 25+ new Annotations  
<https://kubernetes-sigs.github.io/aws-load-balancer-controller/v2.4/guide/service/annotations/>

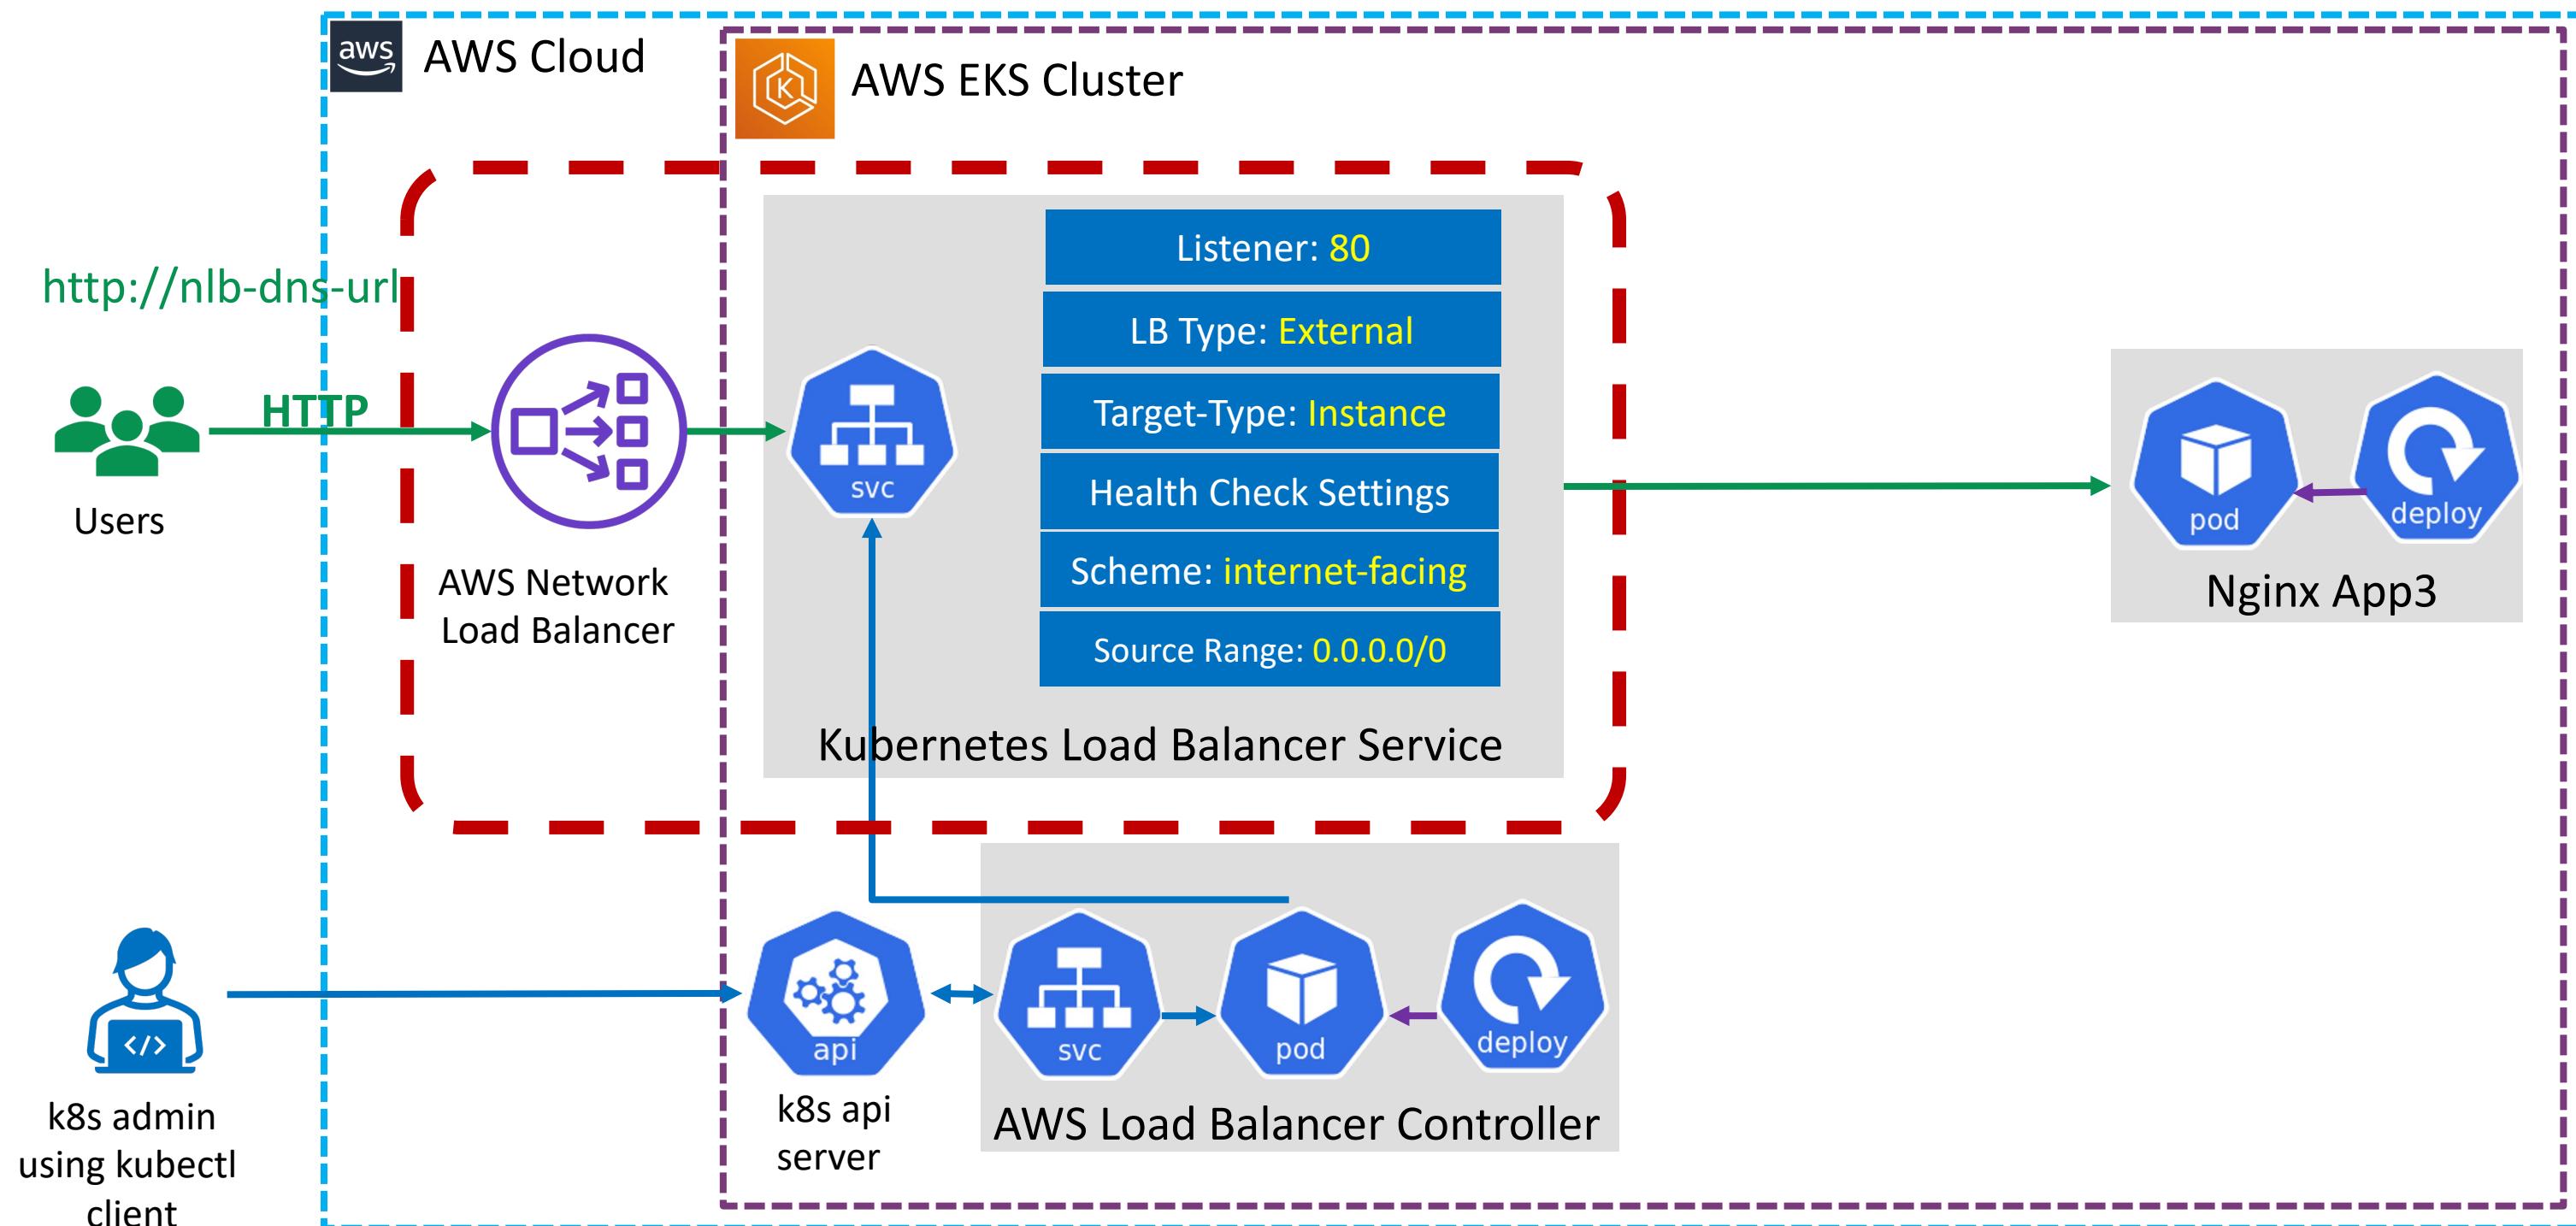
# How to ensure our Network LB created using Kubernetes Service will be associated with only latest AWS Load Balancer Controller ?



```
apiVersion: v1
kind: Service
metadata:
  name: lbc-network-lb-demo
  annotations:
    # Traffic Routing
    service.beta.kubernetes.io/aws-load-balancer-name: lbc-network-lb-demo
    service.beta.kubernetes.io/aws-load-balancer-type: external
    service.beta.kubernetes.io/aws-load-balancer-nlb-target-type: instance
    #service.beta.kubernetes.io/aws-load-balancer-subnets: subnet-xxxx, mySubnet ## Subnets are auto-detected
    # Health Check Settings
    service.beta.kubernetes.io/aws-load-balancer-healthcheck-protocol: http
    service.beta.kubernetes.io/aws-load-balancer-healthcheck-port: traffic-port
    service.beta.kubernetes.io/aws-load-balancer-healthcheck-path: /index.html
    service.beta.kubernetes.io/aws-load-balancer-healthcheck-healthy-threshold: "3"
    service.beta.kubernetes.io/aws-load-balancer-healthcheck-unhealthy-threshold: "3"
    service.beta.kubernetes.io/aws-load-balancer-healthcheck-interval: "10" # The controller currently ignores this value
    # Access Control
    service.beta.kubernetes.io/load-balancer-source-ranges: 0.0.0.0/0
    service.beta.kubernetes.io/aws-load-balancer-scheme: "internet-facing"
    # AWS Resource Tags
    service.beta.kubernetes.io/aws-load-balancer-additional-resource-tags: Environment=dev,Team=test
spec:
  type: LoadBalancer
  selector:
    app: app3-nginx
  ports:
    - port: 80
      targetPort: 80
```

This Kubernetes Service creates Network Load Balancer which will be associated to latest AWS Load Balancer Controller

# AWS EKS NLB – Basics

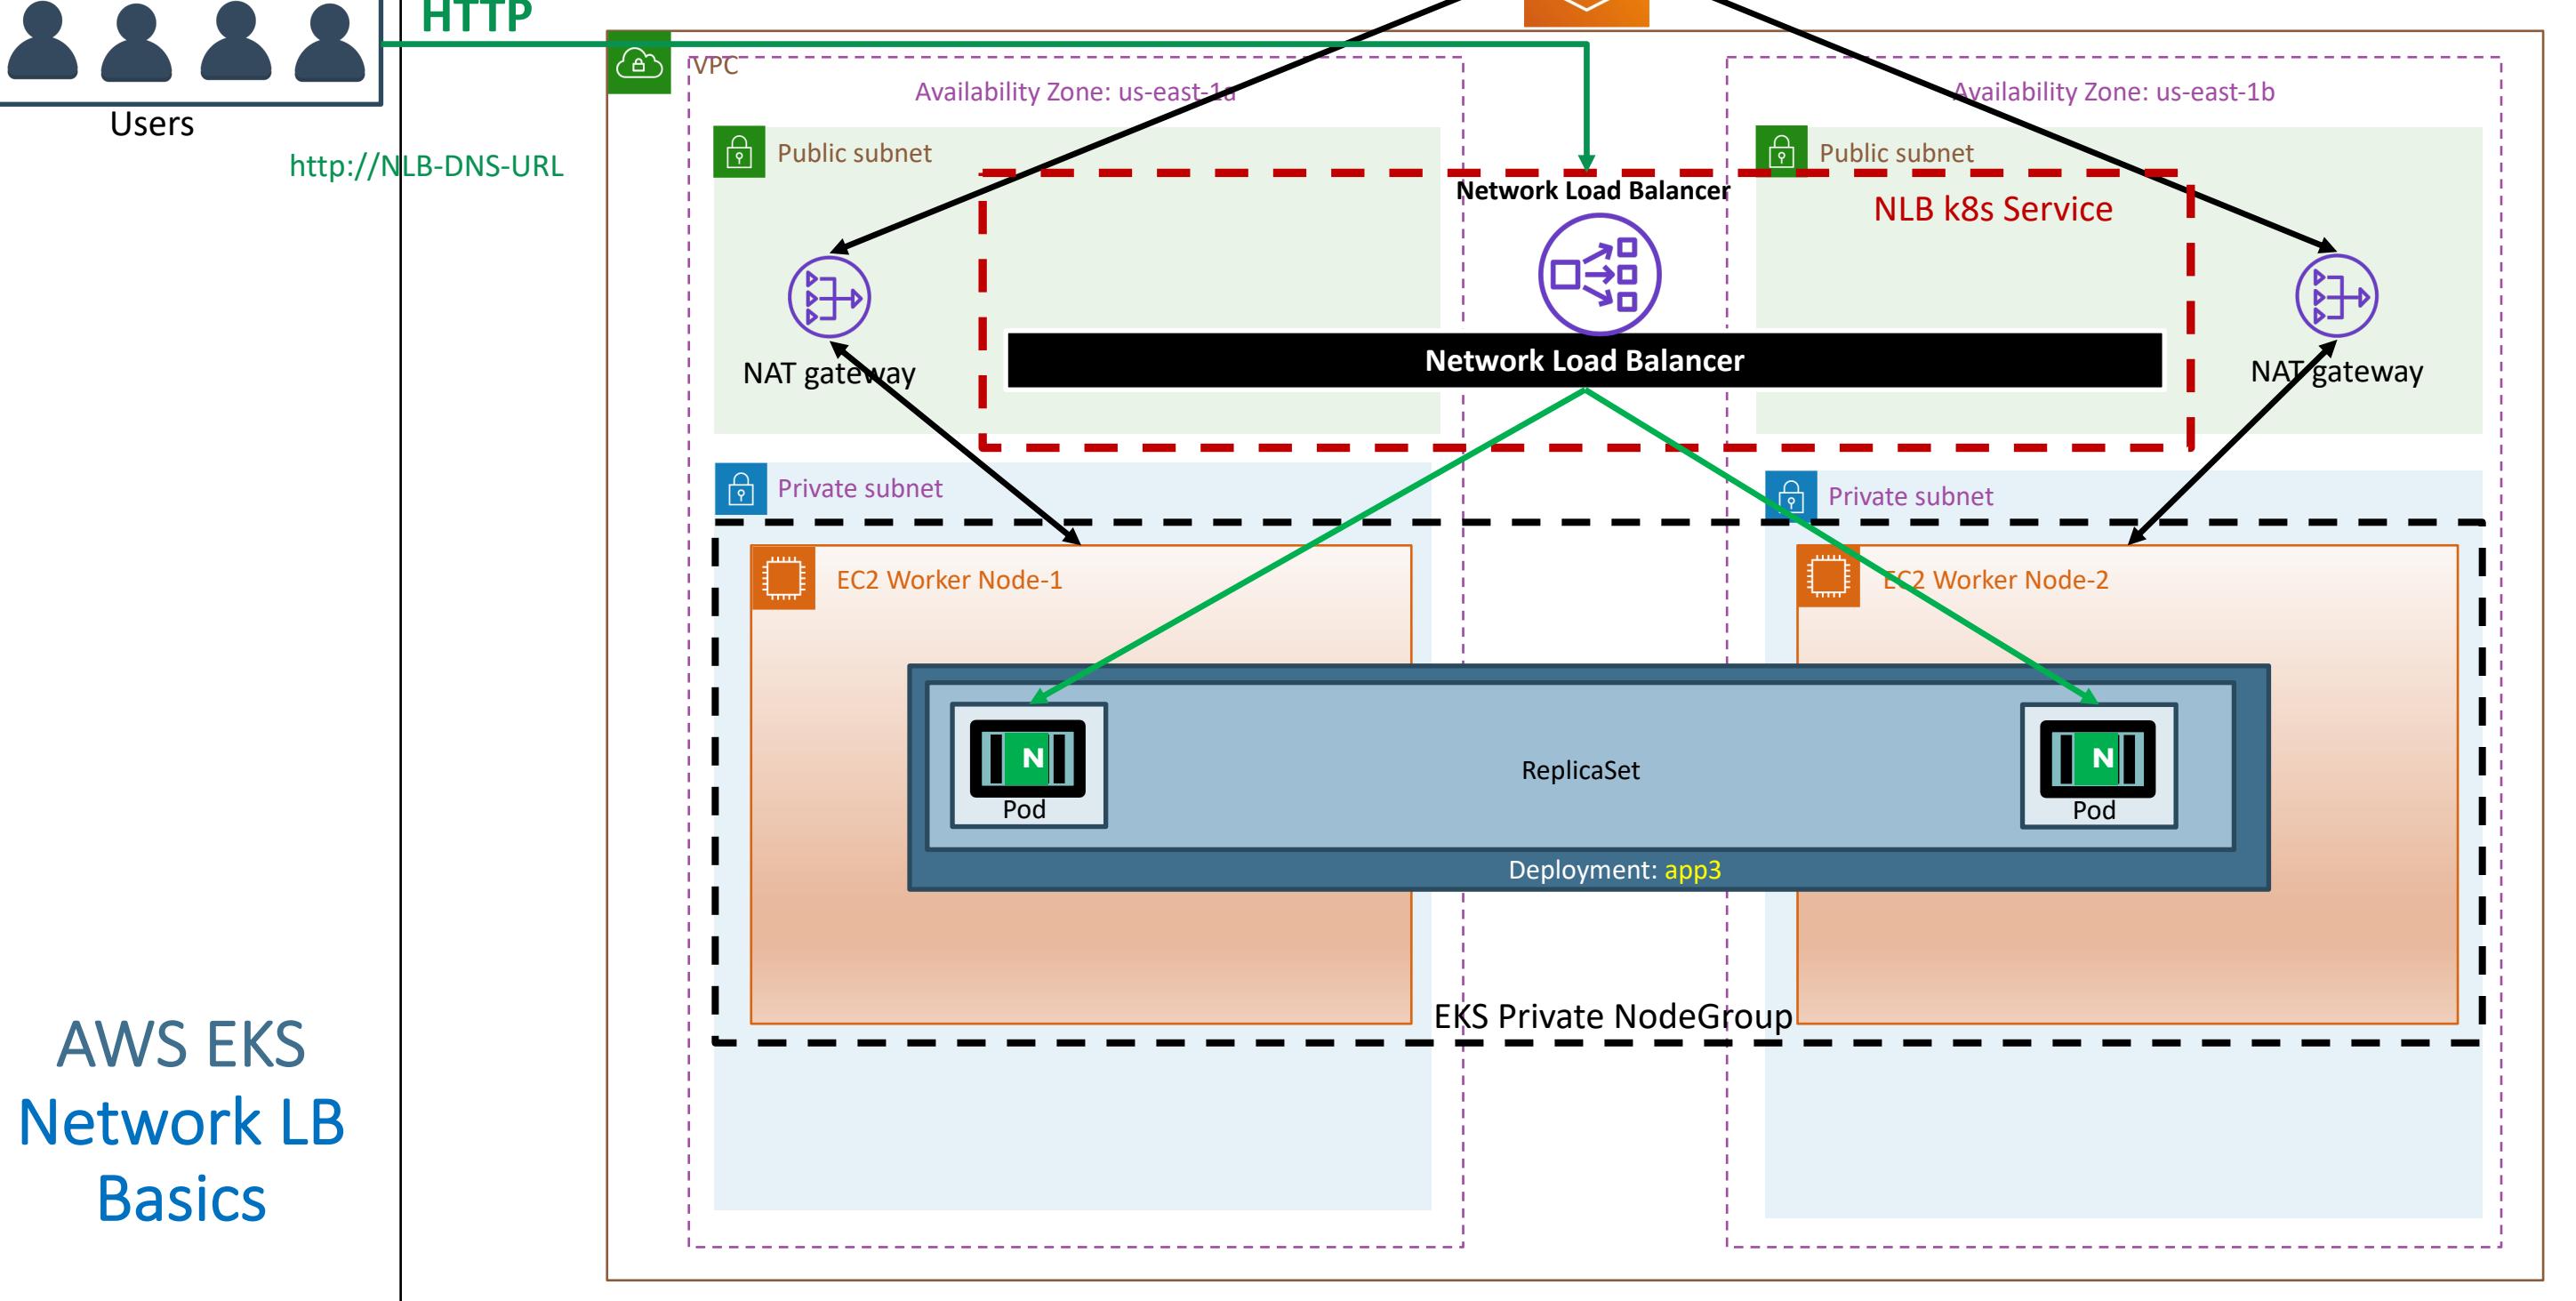




Users

HTTP

http://NLB-DNS-URL



**annotations:**

```
# Traffic Routing
service.beta.kubernetes.io/aws-load-balancer-name: lbc-network-lb-demo
service.beta.kubernetes.io/aws-load-balancer-type: external
service.beta.kubernetes.io/aws-load-balancer-nlb-target-type: instance
#service.beta.kubernetes.io/aws-load-balancer-subnets: subnet-xxxx, mySubnet ## Subnets

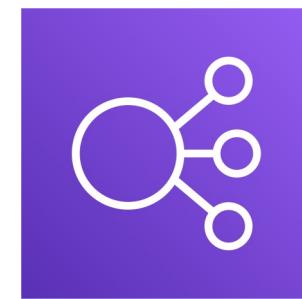
# Health Check Settings
service.beta.kubernetes.io/aws-load-balancer-healthcheck-protocol: http
service.beta.kubernetes.io/aws-load-balancer-healthcheck-port: traffic-port
service.beta.kubernetes.io/aws-load-balancer-healthcheck-path: /index.html
service.beta.kubernetes.io/aws-load-balancer-healthcheck-healthy-threshold: "3"
service.beta.kubernetes.io/aws-load-balancer-healthcheck-unhealthy-threshold: "3"
service.beta.kubernetes.io/aws-load-balancer-healthcheck-interval: "10" # The controller

# Access Control
service.beta.kubernetes.io/load-balancer-source-ranges: 0.0.0.0/0
service.beta.kubernetes.io/aws-load-balancer-scheme: "internet-facing"

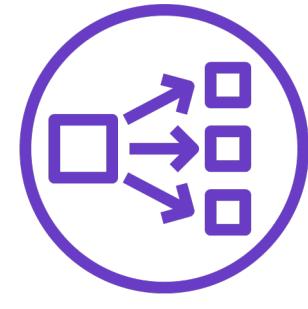
# AWS Resource Tags
service.beta.kubernetes.io/aws-load-balancer-additional-resource-tags: Environment=dev,
```

Kubernetes  
Service of  
Type NLB

Annotations



Elastic Load  
Balancing



Network  
Load Balancer



Kubernetes  
Service



External  
DNS



# AWS EKS

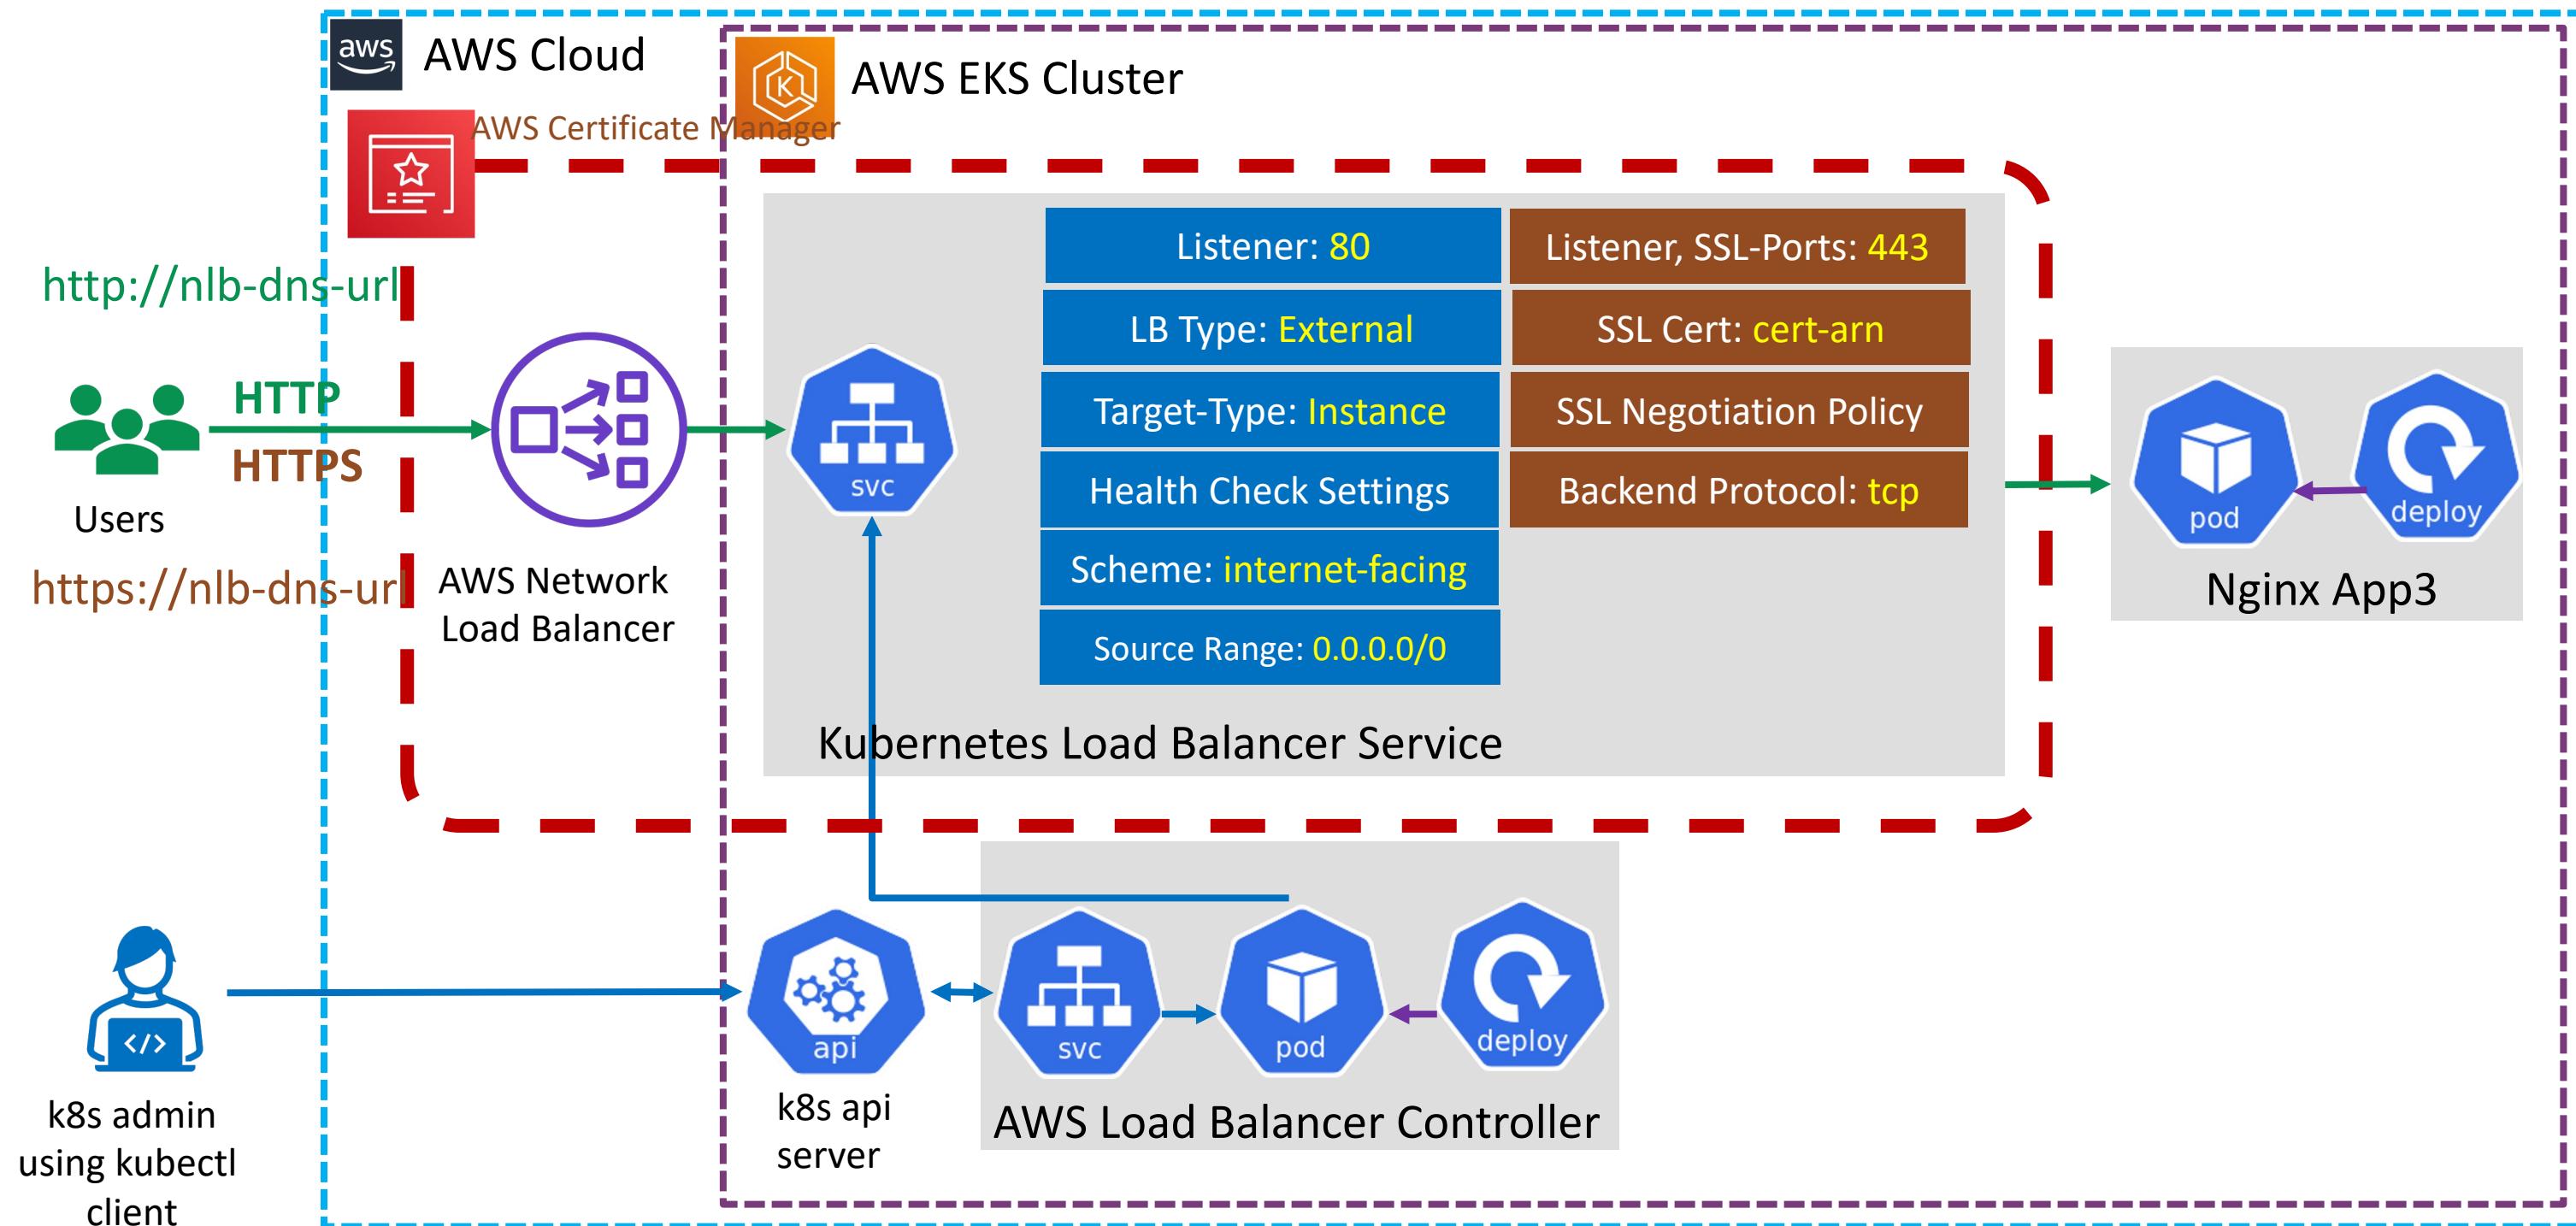
## Load Balancer Controller

### Kubernetes Service

### Network Load Balancer - TLS



# AWS EKS NLB – TLS (HTTP + HTTPS)





AWS Cloud

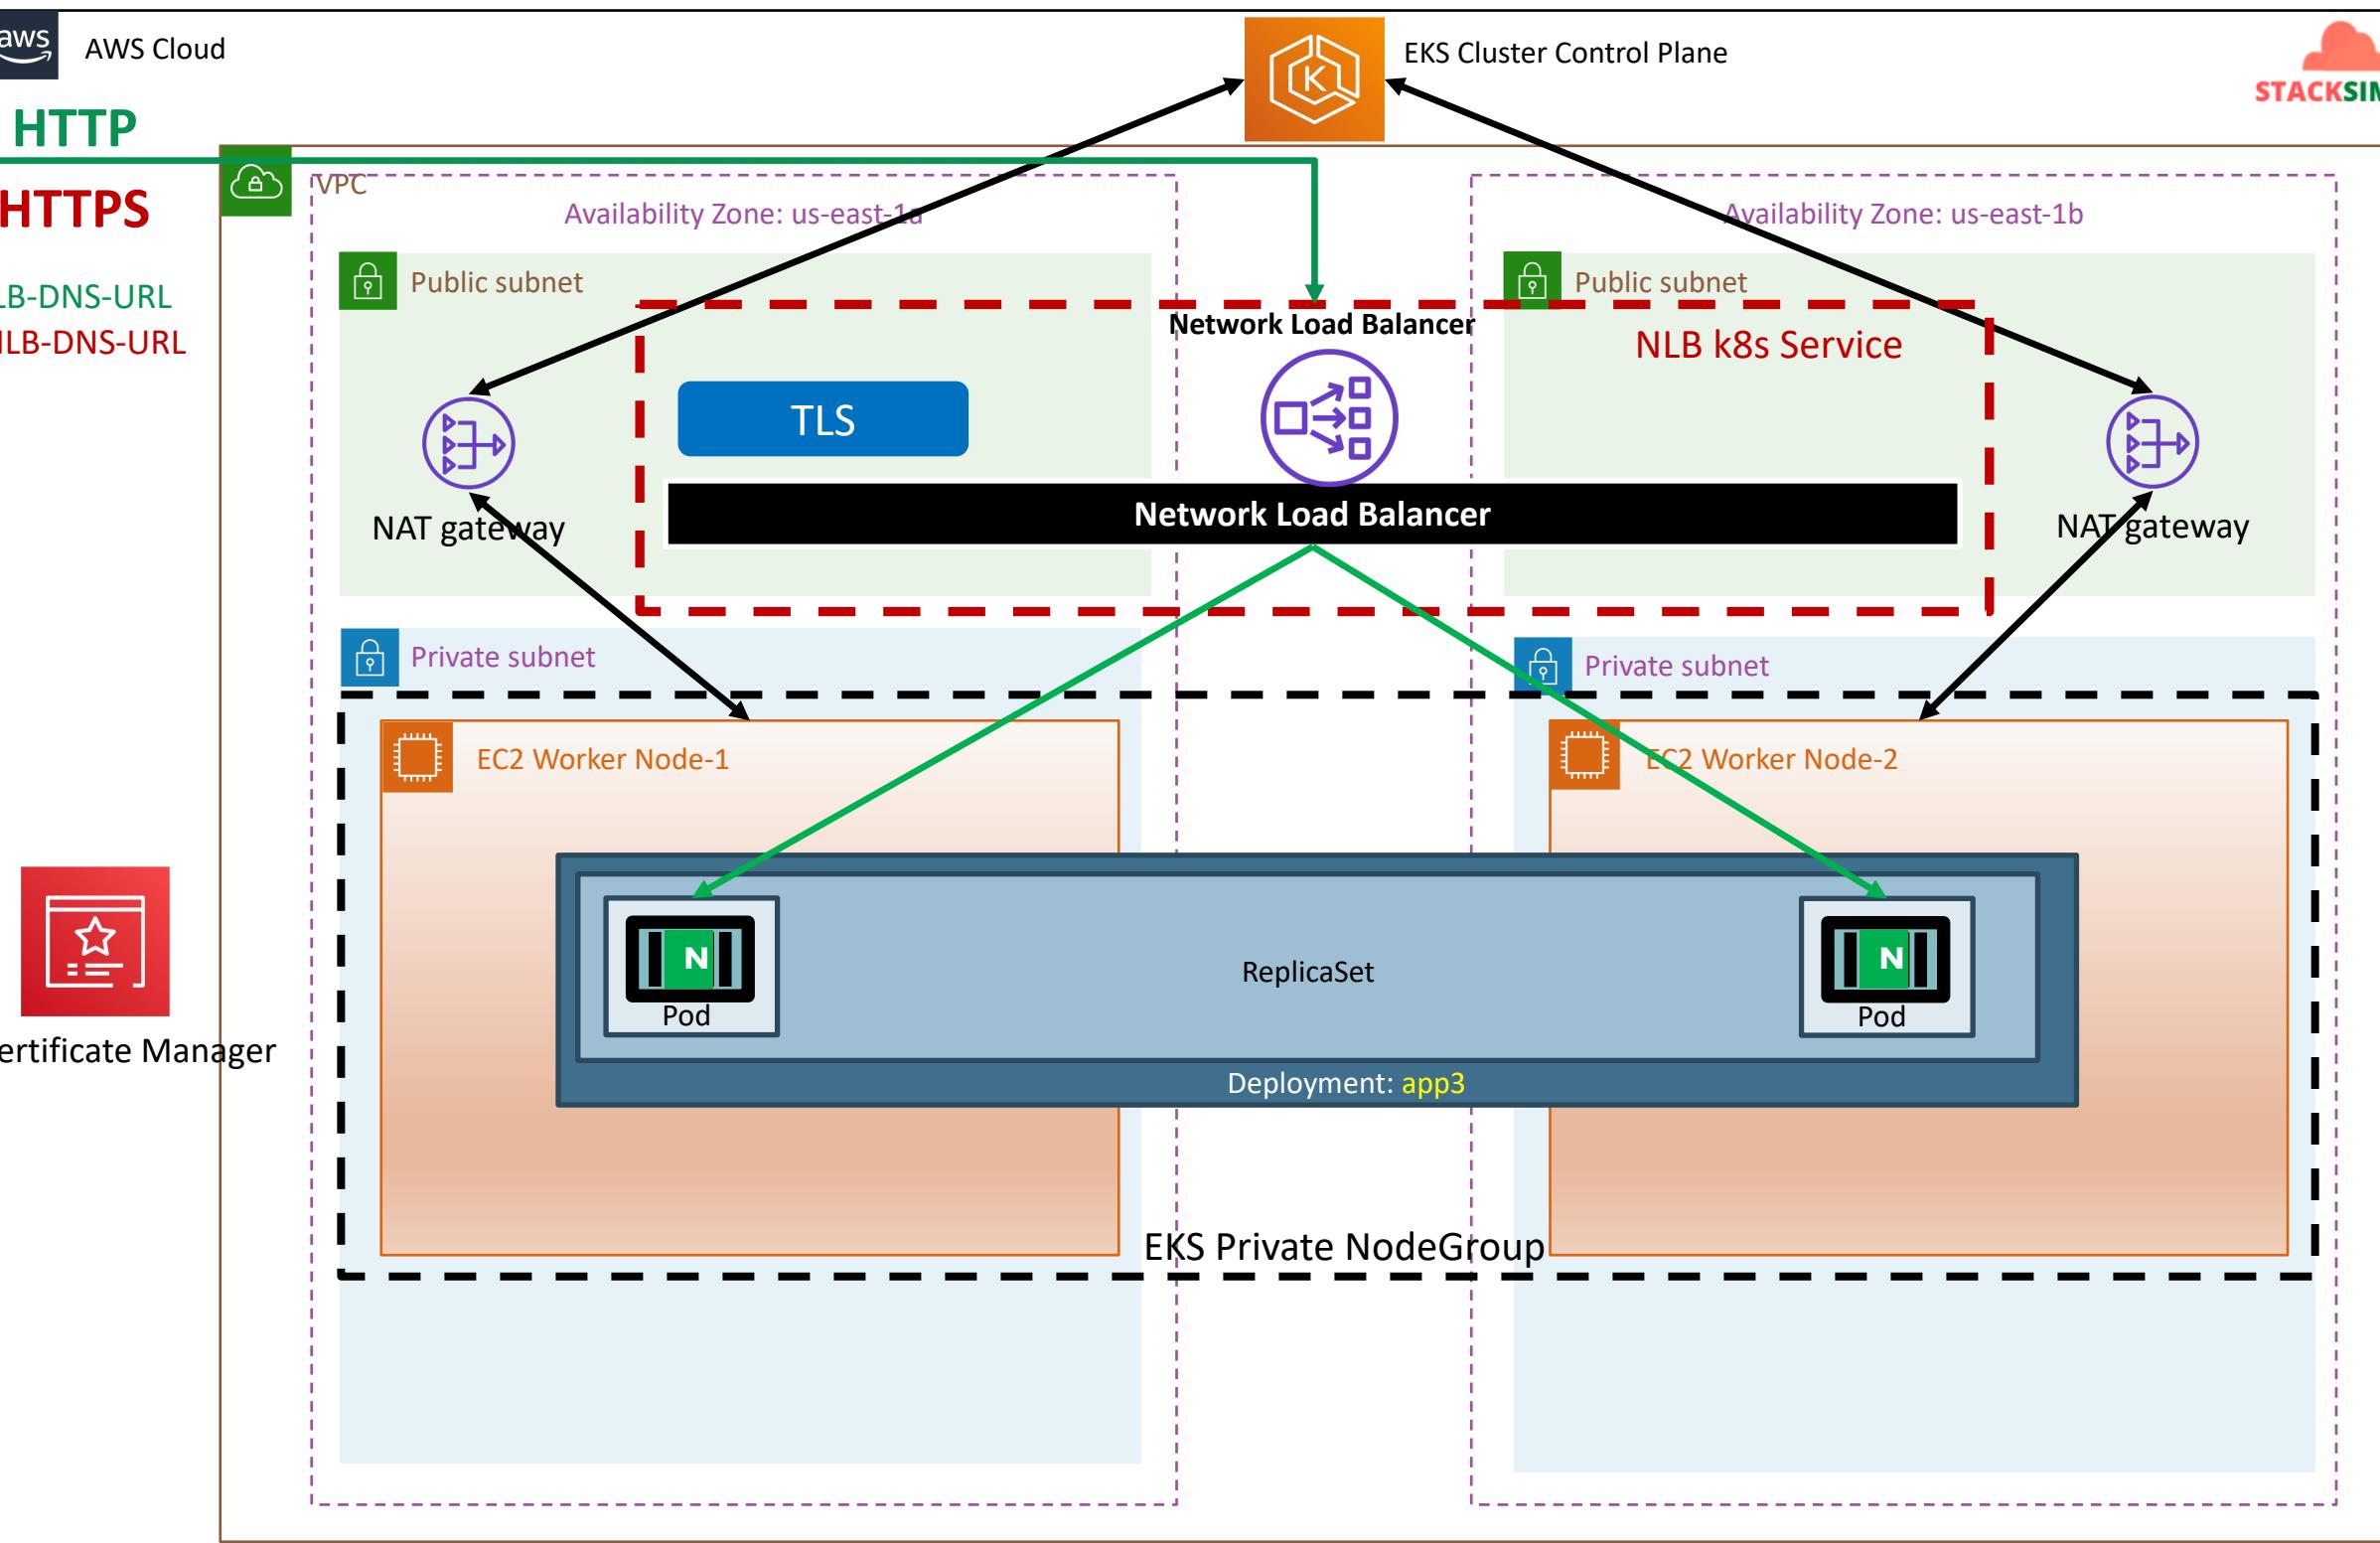
**HTTP**  
**HTTPS**

Users

<http://NLB-DNS-URL>  
<https://NLB-DNS-URL>

AWS Certificate Manager

# AWS EKS Network LB With TLS



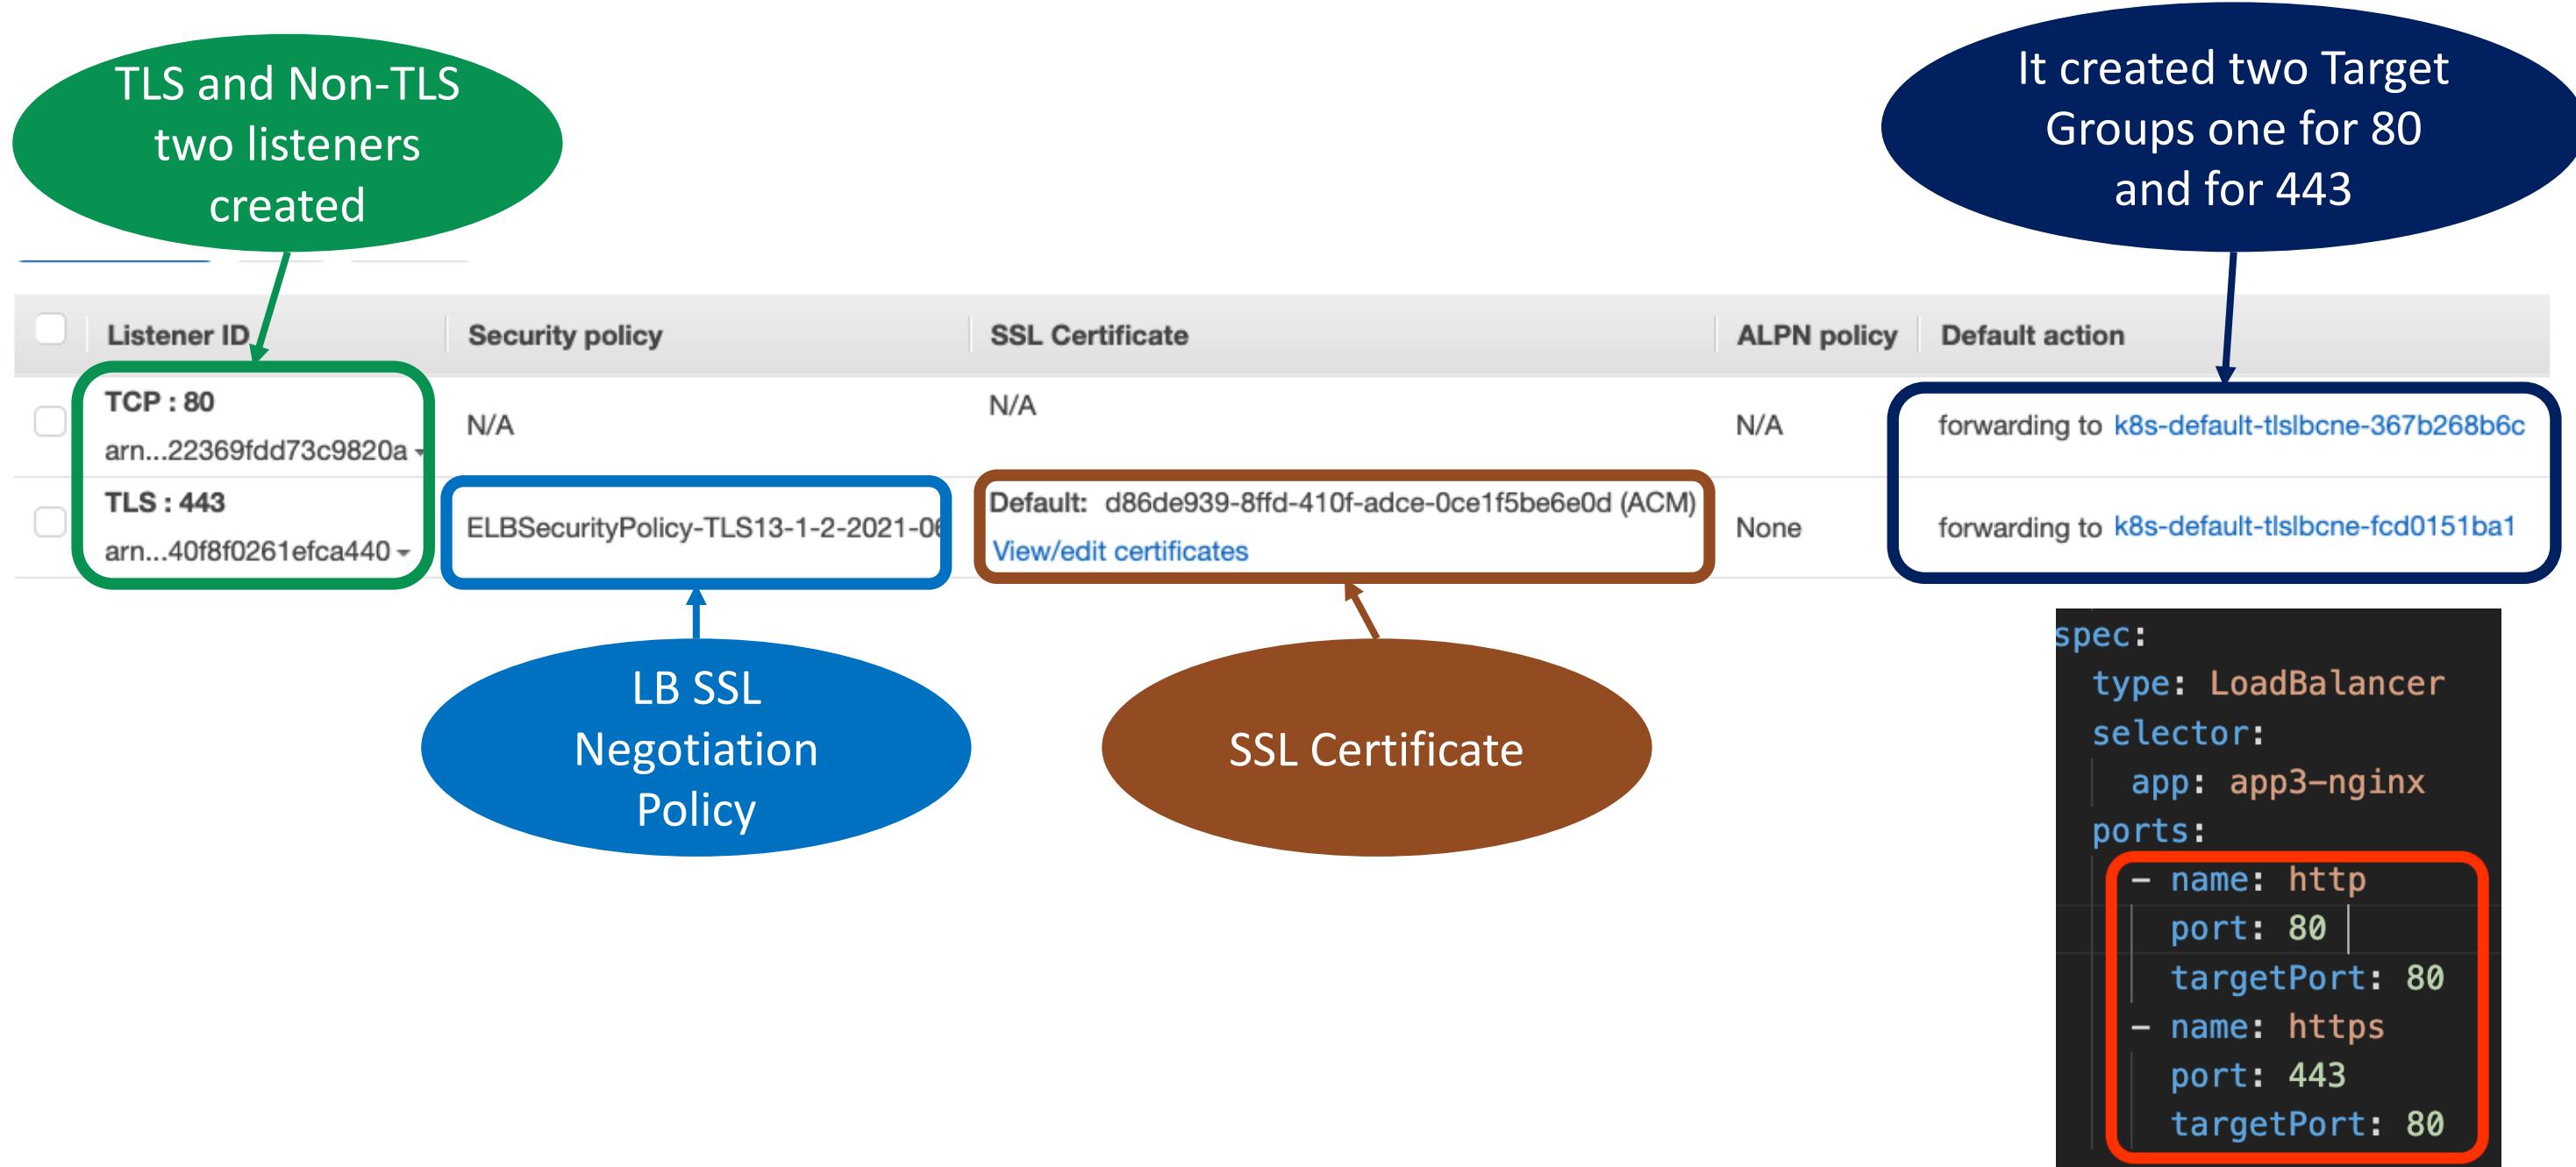
# AWS EKS NLB -Annotations

```
# TLS
service.beta.kubernetes.io/aws-load-balancer-ssl-cert: arn:aws:acm:us-east-1:180789647333:certificate/d8
service.beta.kubernetes.io/aws-load-balancer-ssl-ports: 443, # Specify this annotation if you need both
service.beta.kubernetes.io/aws-load-balancer-ssl-negotiation-policy: ELBSecurityPolicy-TLS13-1-2-2021-06
service.beta.kubernetes.io/aws-load-balancer-backend-protocol: tcp # specifies whether to use TLS or TCP
spec:
  type: LoadBalancer
  selector:
    app: app3-nginx
  ports:
    - name: http
      port: 80
      targetPort: 80
    - name: https
      port: 443
      targetPort: 80
```

# AWS EKS NLB -Annotations

```
# TLS
service.beta.kubernetes.io/aws-load-balancer-ssl-cert: arn:aws:acm:us-east-1:123456789012:certificate/12345678-1234-1234-1234-123456789012
service.beta.kubernetes.io/aws-load-balancer-ssl-ports: 443, # Specify this port if your application uses https
service.beta.kubernetes.io/aws-load-balancer-ssl-negotiation-policy: ELBSecurityPolicy-TLS-1-2-2017-1
service.beta.kubernetes.io/aws-load-balancer-backend-protocol: tcp # specifies the protocol for the target group
spec:
  type: LoadBalancer
  selector:
    app: app3-nginx
  ports:
    - name: http
      port: 80          # Creates NLB Port 80 Listener
      targetPort: 80   # Creates NLB Port 80 Target Group-1
    - name: https
      port: 443         # Creates NLB Port 443 Listener
      targetPort: 80   # Creates NLB Port 80 Target Group-2
    - name: http81
      port: 81          # Creates NLB Port 81 Listener
      targetPort: 80   # Creates NLB Port 80 Target Group-3
    - name: http82
      port: 82          # Creates NLB Port 82 Listener
      targetPort: 80   # Creates NLB Port 80 Target Group-4
```

# AWS EKS NLB –Load Balancer



# AWS EKS NLB -Target Groups

EC2 > Target groups

**Target groups (2) [Info](#)**

[Actions ▾](#) [Create target group](#)

[C](#)

[Search or filter target groups](#)

< 1 > [⚙️](#)

Name	ARN	Port	Protocol	Target type	Load balancer
<a href="#">k8s-default-tlslbcne-367b268b6c</a>	<input type="button"/> <a href="#">arn:aws:elasticloadbalancin...</a>	30810	TCP	Instance	<a href="#">tls-lbc-network-lb</a>
<a href="#">k8s-default-tlslbcne-fcd0151ba1</a>	<input type="button"/> <a href="#">arn:aws:elasticloadbalancin...</a>	32162	TCP	Instance	<a href="#">tls-lbc-network-lb</a>

# AWS EKS NLB -Target Groups

arn:aws:elasticloadbalancing:us-east-1:180789647333:targetgroup/k8s-default-tlsLbcne-367b268b6c/e564572efca1f1aa

Details					
Target type Instance	Protocol : Port TCP: 30810	VPC <a href="#">vpc-0165a396e41e292a3</a>	IP address type IPv4		
Load balancer <a href="#">tls-lbc-network-lb</a>					
Total targets 2	<b>Healthy</b> <input checked="" type="checkbox"/> 2	Unhealthy <input type="checkbox"/> 0	Unused <input type="checkbox"/> 0	Initial <input type="checkbox"/> 0	Draining <input type="checkbox"/> 0

**Targets**    [Monitoring](#)    [Health checks](#)    [Attributes](#)    [Tags](#)

Registered targets (2)					
<input type="text"/> Filter resources by property or value				<input type="button"/>	<a href="#">Deregister</a>
<input type="checkbox"/>	<a href="#">Instance ID</a>	Name	Port	Zone	<a href="#">Health status</a>
<input type="checkbox"/>	i-0e06f6d06e5dfa0df	eksdemo1-eksdemo1-nginx-private1-Node	30810	us-east-1a	<input checked="" type="checkbox"/> healthy
<input type="checkbox"/>	i-016128c9dd4835826	eksdemo1-eksdemo1-nginx-private1-Node	30810	us-east-1b	<input checked="" type="checkbox"/> healthy

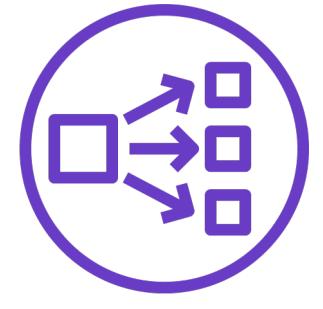
# AWS EKS NLB -Target Groups

arn:aws:elasticloadbalancing:us-east-1:180789647333:targetgroup/k8s-default-tlslbcne-367b268b6c/e564572efca1f1aa

Details													
Target type Instance	Protocol : Port TCP: 30810	VPC <a href="#">vpc-0165a396e41e292a3</a>	IP address type IPv4										
Load balancer <a href="#">tls-lbc-network-lb</a>													
Total targets 2	Healthy <span style="color: green;">✓ 2</span>	Unhealthy <span style="color: red;">✗ 0</span>	Unused <span style="color: grey;">○ 0</span>	Initial <span style="color: grey;">⌚ 0</span>	Draining <span style="color: grey;">⊖ 0</span>								
Targets	Monitoring	<b>Health checks</b>	Attributes	Tags									
<b>Health check settings</b> <div style="float: right;"><a href="#">Edit</a></div> <table border="1"> <tr> <td>Protocol HTTP</td> <td>Path <a href="#">/index.html</a></td> <td>Port Traffic port</td> <td>Healthy threshold 3 consecutive health check successes</td> </tr> <tr> <td>Unhealthy threshold 3 consecutive health check failures</td> <td>Timeout 6 seconds</td> <td>Interval 10 seconds</td> <td>Success codes 200-399</td> </tr> </table>						Protocol HTTP	Path <a href="#">/index.html</a>	Port Traffic port	Healthy threshold 3 consecutive health check successes	Unhealthy threshold 3 consecutive health check failures	Timeout 6 seconds	Interval 10 seconds	Success codes 200-399
Protocol HTTP	Path <a href="#">/index.html</a>	Port Traffic port	Healthy threshold 3 consecutive health check successes										
Unhealthy threshold 3 consecutive health check failures	Timeout 6 seconds	Interval 10 seconds	Success codes 200-399										



Elastic Load  
Balancing



Network  
Load Balancer



Kubernetes  
Service



External  
DNS



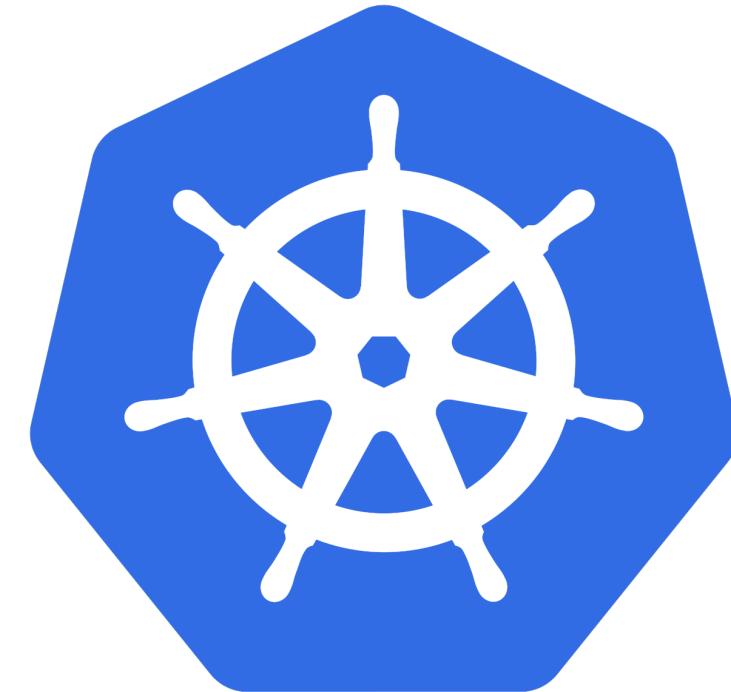
# AWS EKS

## Load Balancer Controller

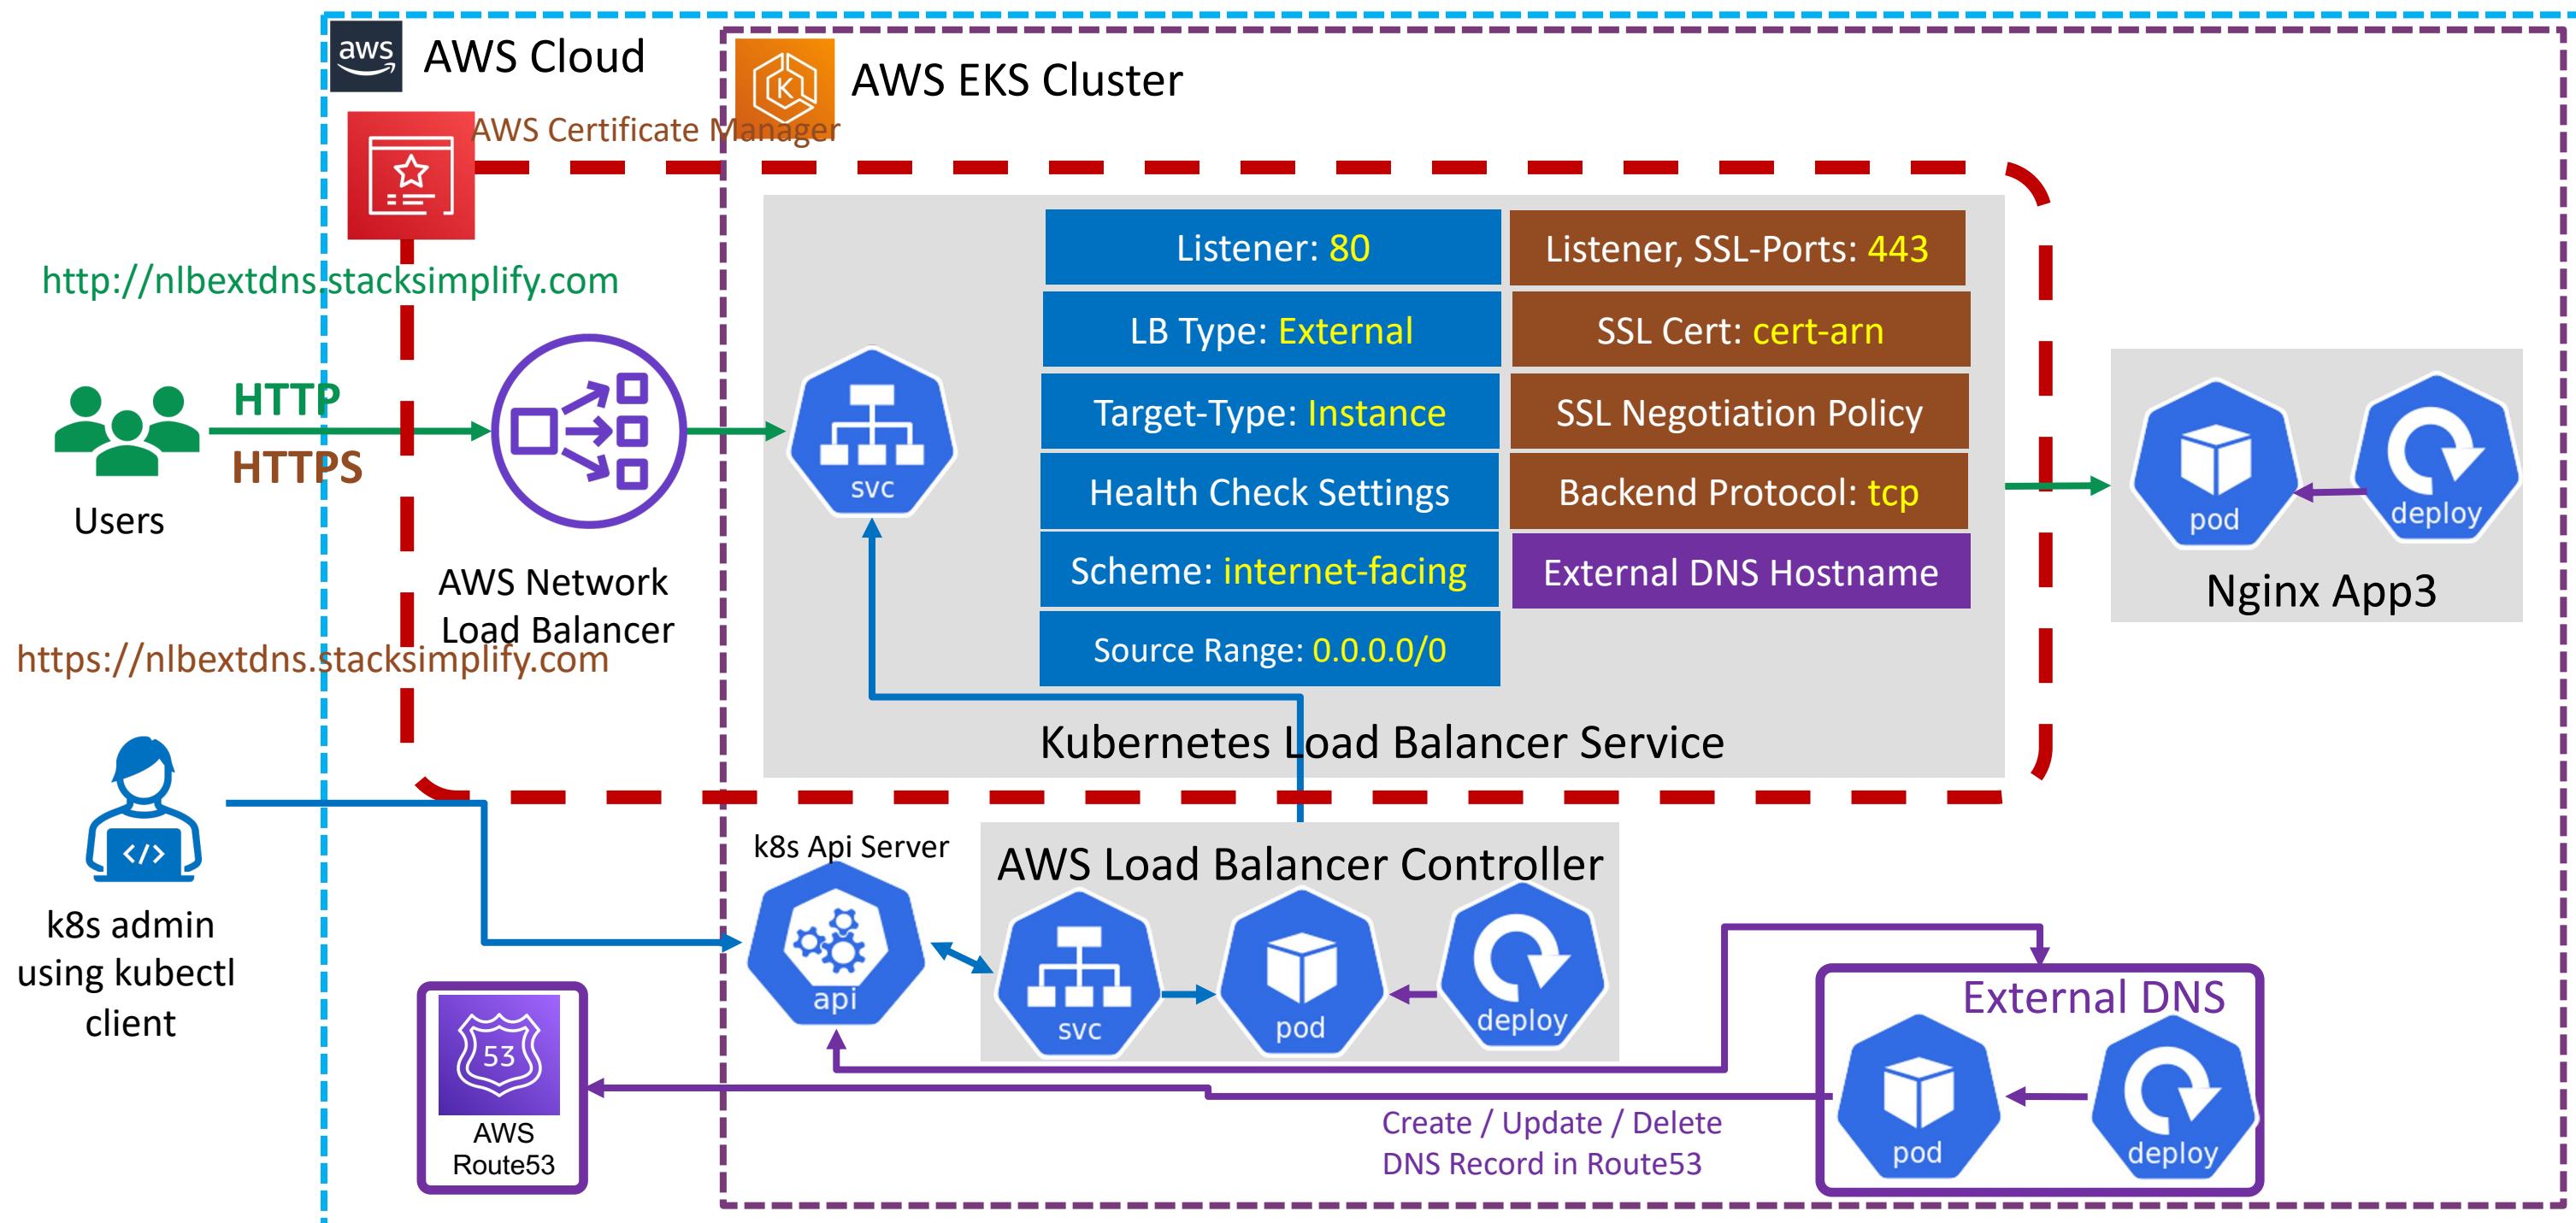
### Kubernetes Service

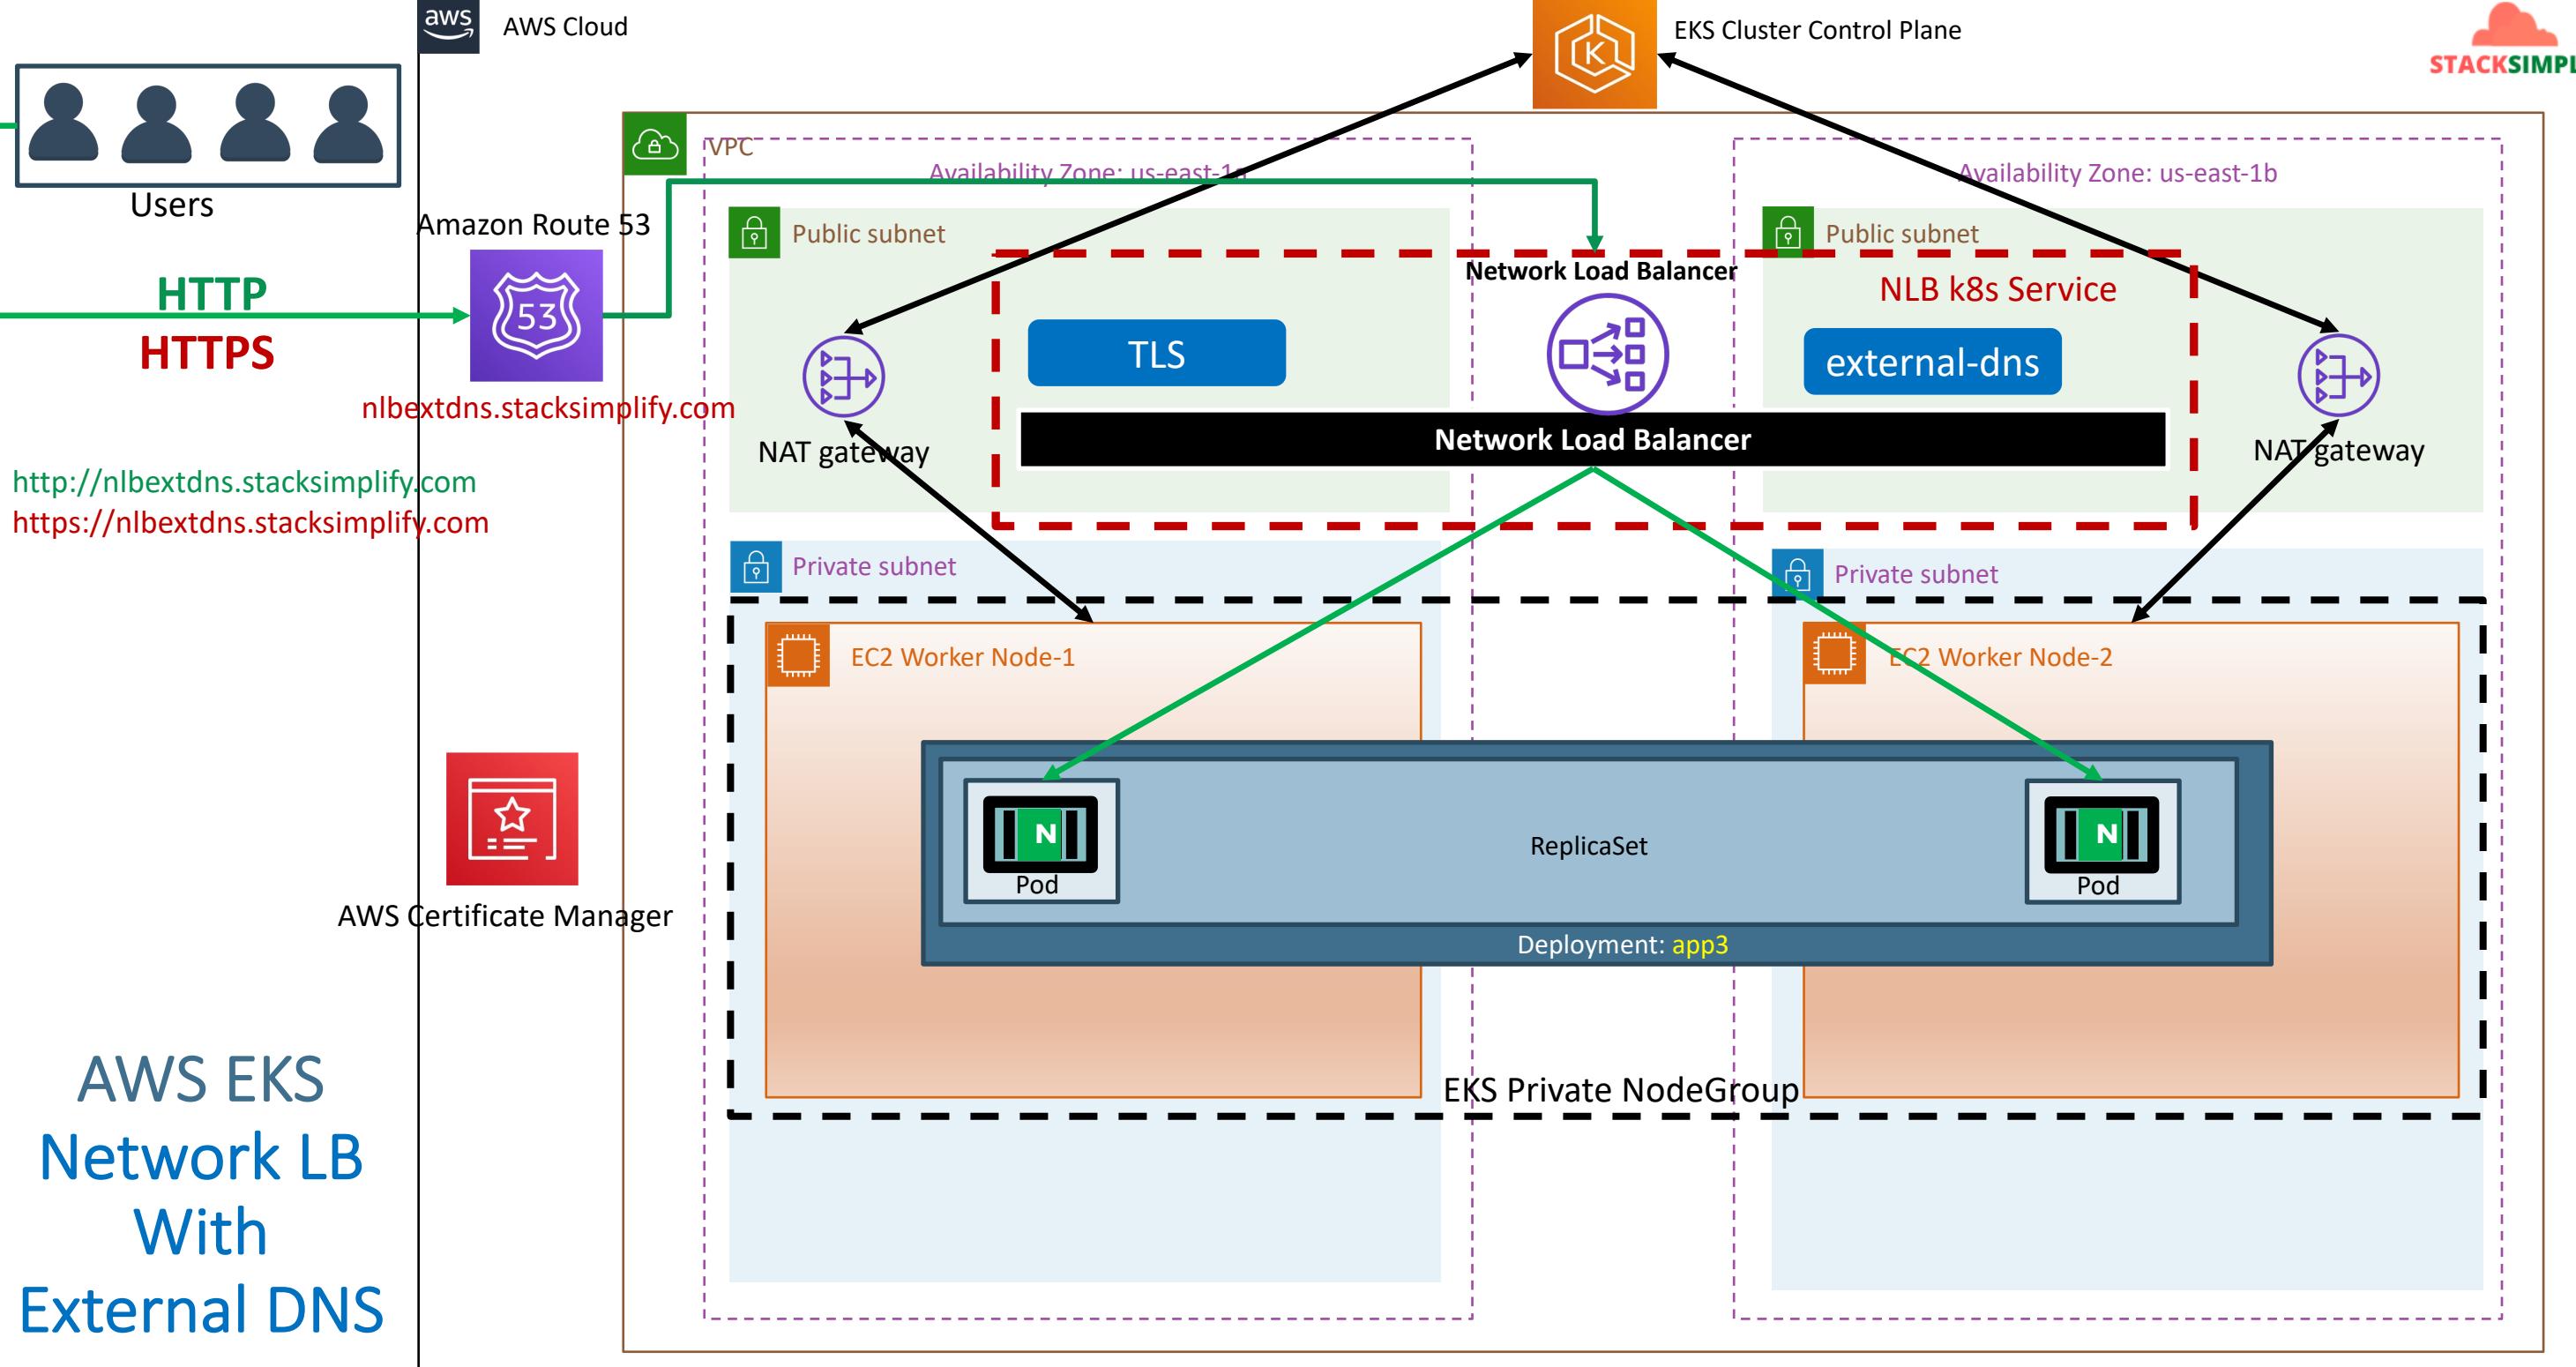
### Network Load Balancer

### External DNS



# AWS EKS NLB – External DNS



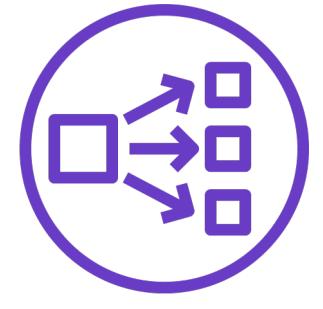


# AWS EKS NLB - External DNS Annotation

```
# External DNS – For creating a Record Set in Route53  
external-dns.alpha.kubernetes.io/hostname: nlbdns101.stacksimplify.com
```



Elastic Load  
Balancing



Network  
Load Balancer



Kubernetes  
Service



External  
DNS



# AWS EKS

## Load Balancer Controller

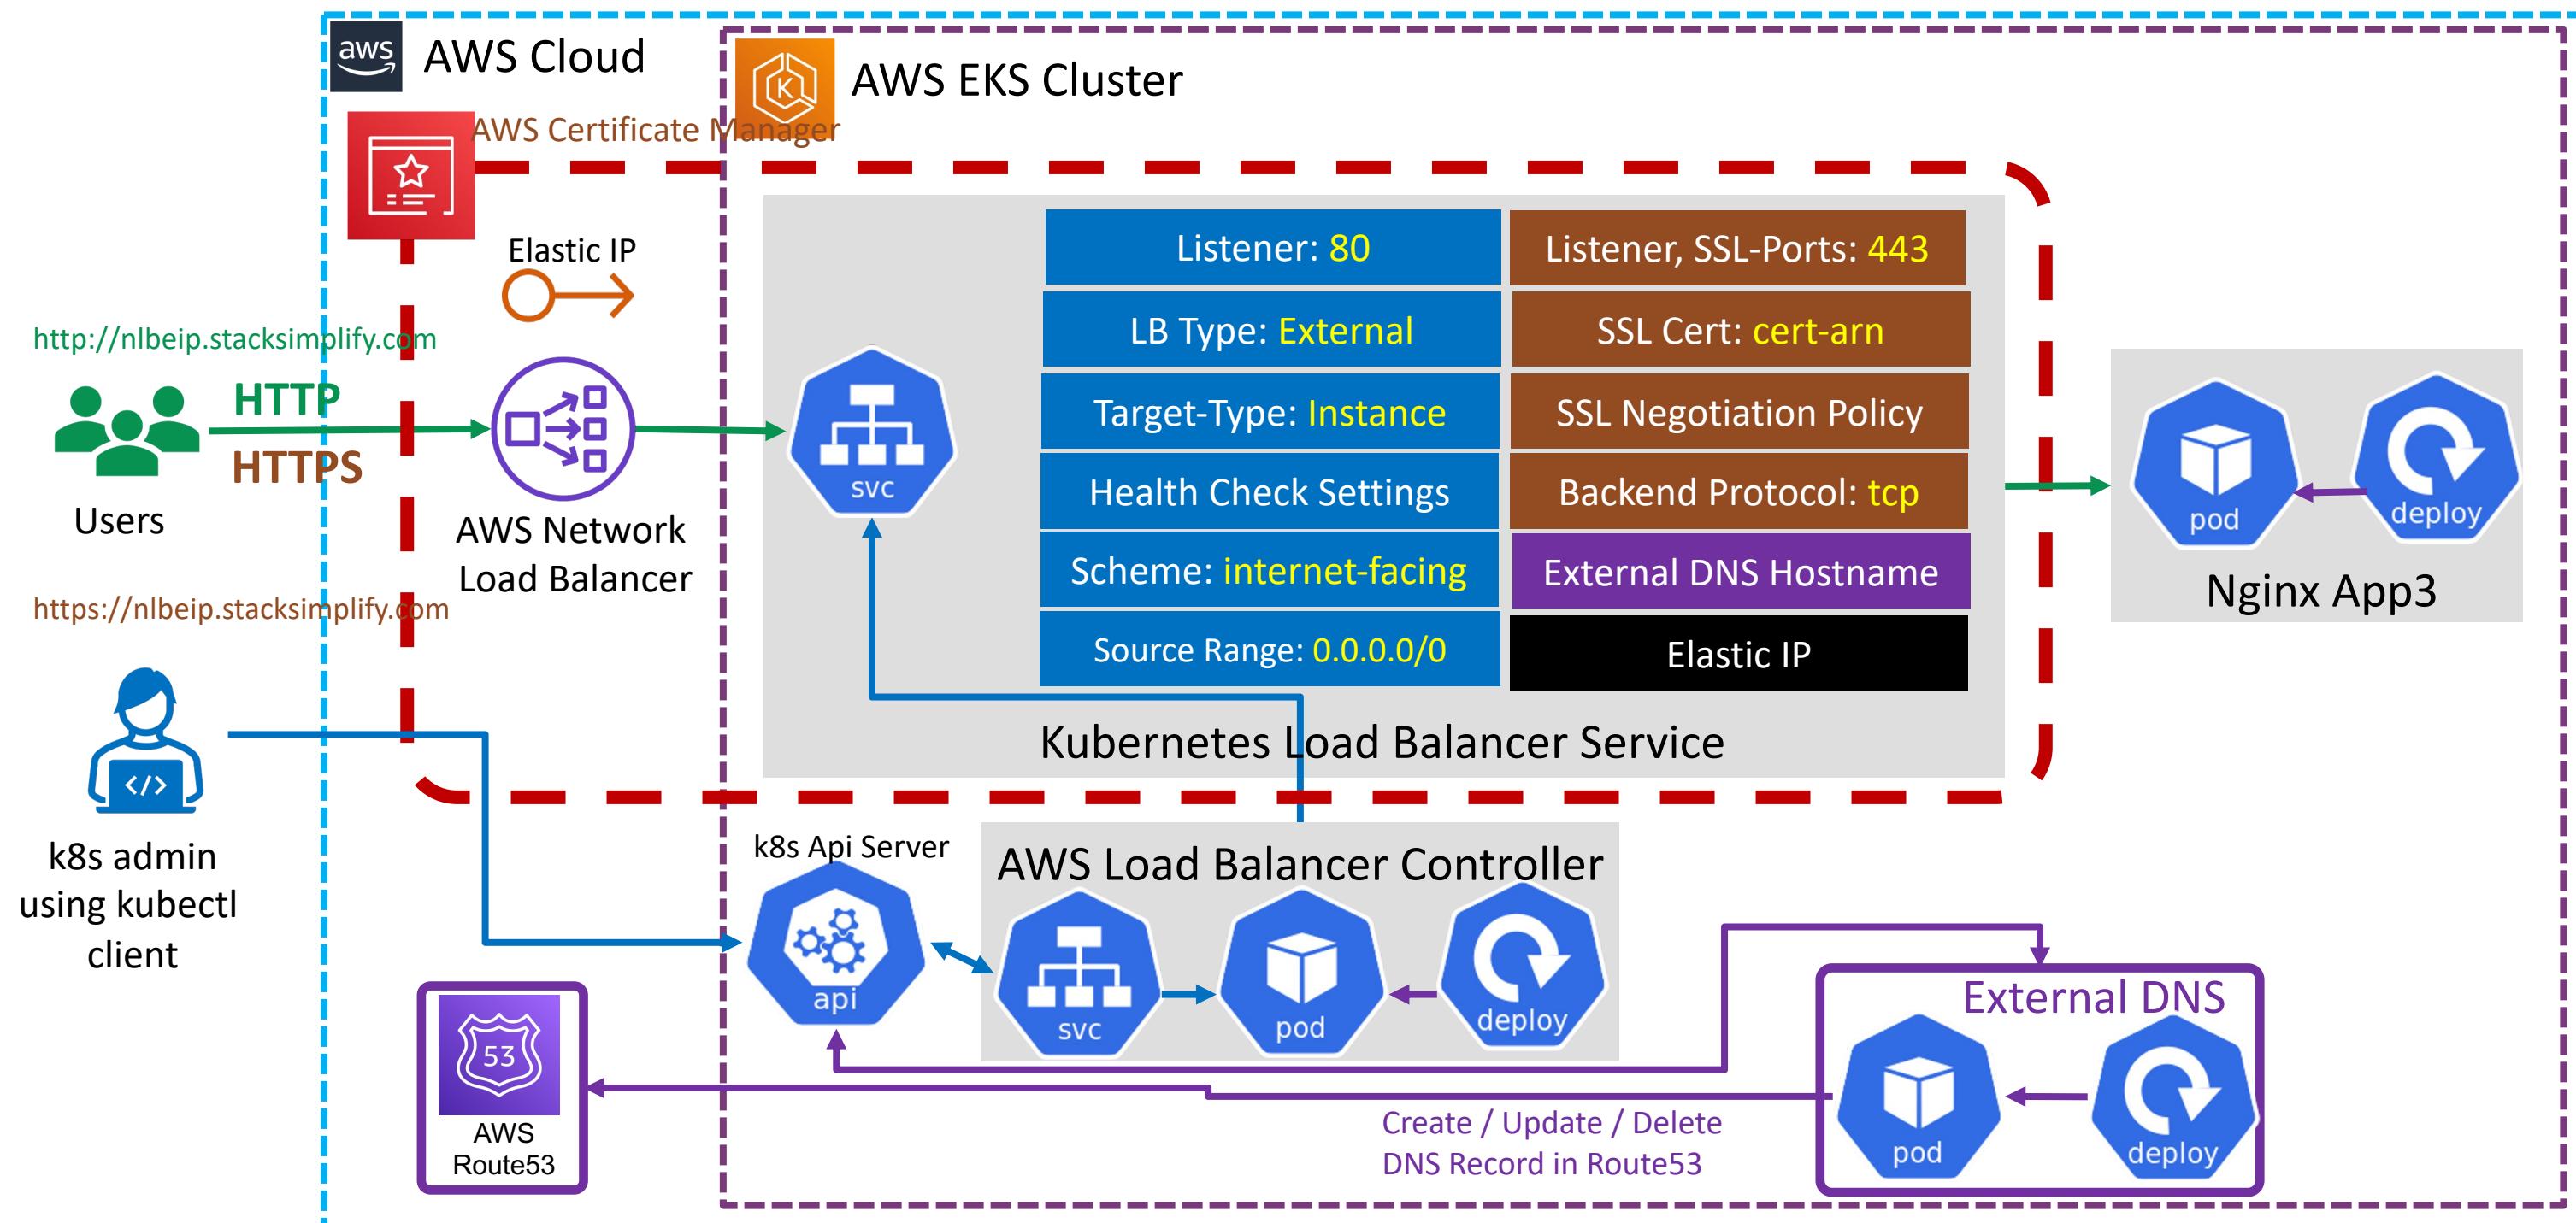
**Kubernetes Service**

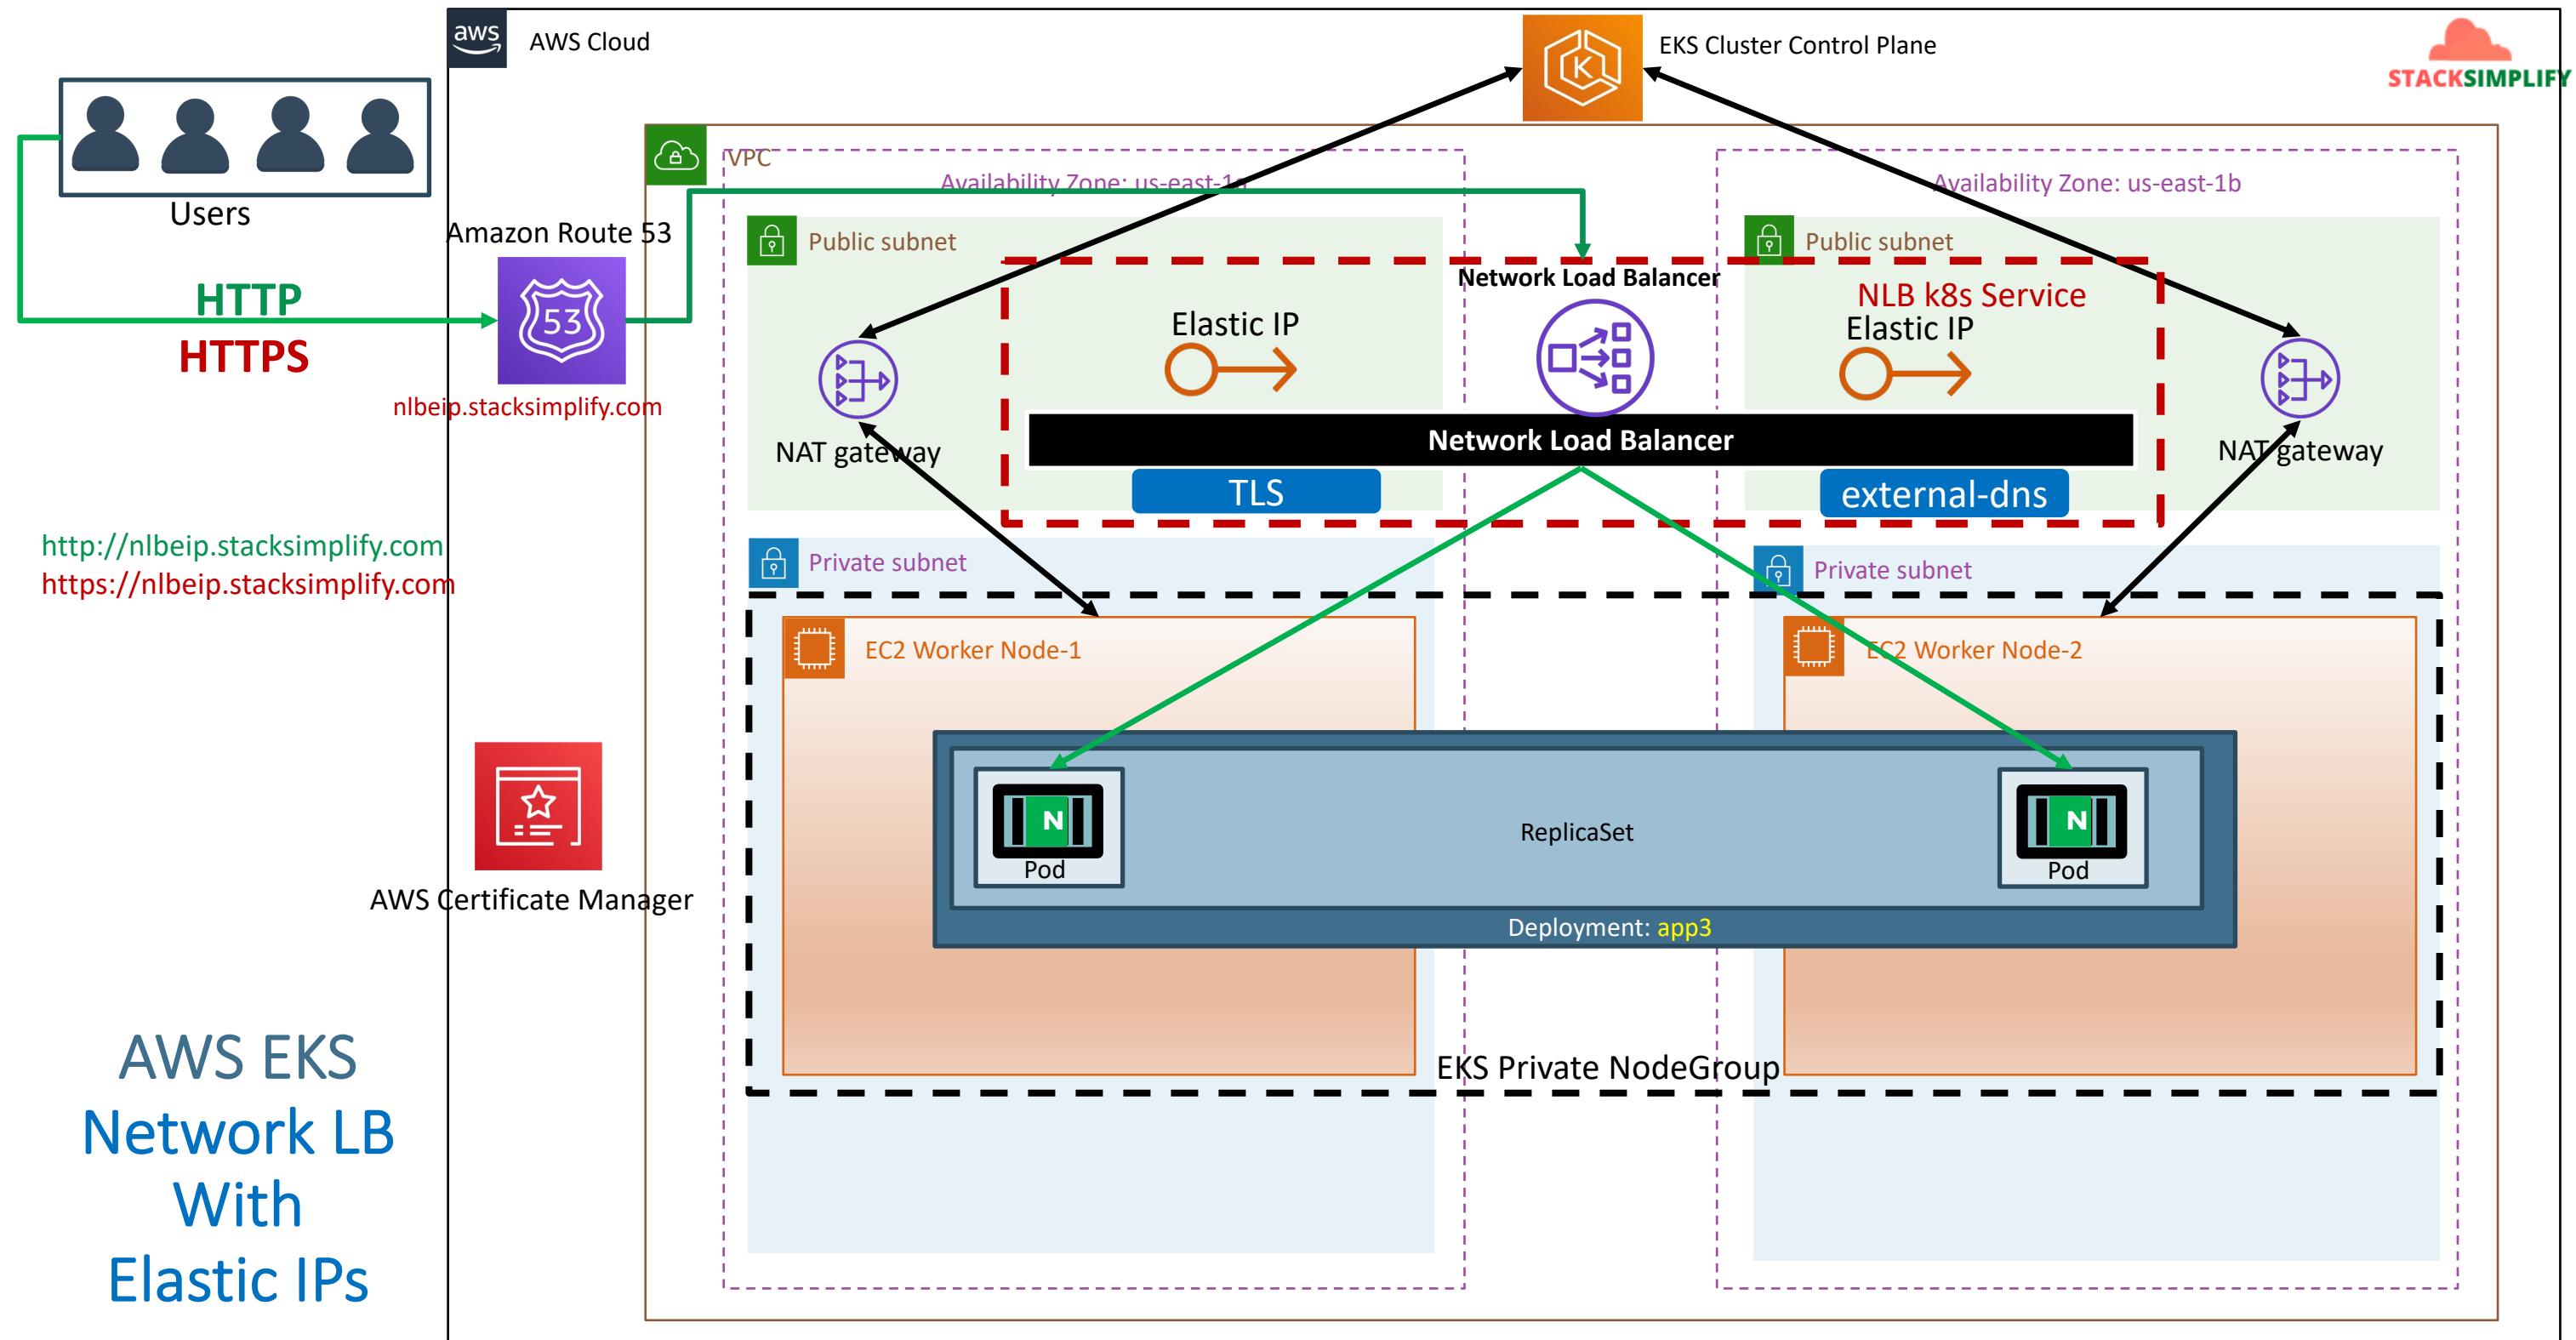
**Network Load Balancer**

**Elastic IP**



# AWS EKS NLB – External DNS





# AWS EKS Network LB With Elastic IPs

# AWS EKS NLB - Elastic IP Annotation

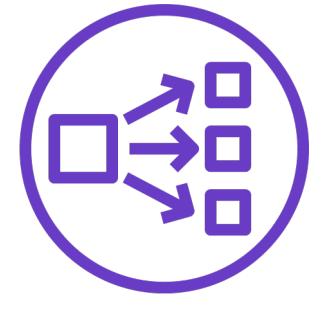
```
# Elastic IPs  
service.beta.kubernetes.io/aws-load-balancer-eip-allocations: eipalloc-07daf60991cf21f0,
```

**Important Note-1:**  
You must specify the same number of **eip allocations** as load balancer subnets annotation

**Important Note-2:**  
NLB must be **internet-facing**



Elastic Load  
Balancing



Network  
Load Balancer



Kubernetes  
Service



External  
DNS



# AWS EKS

## Load Balancer Controller

**Kubernetes Service**

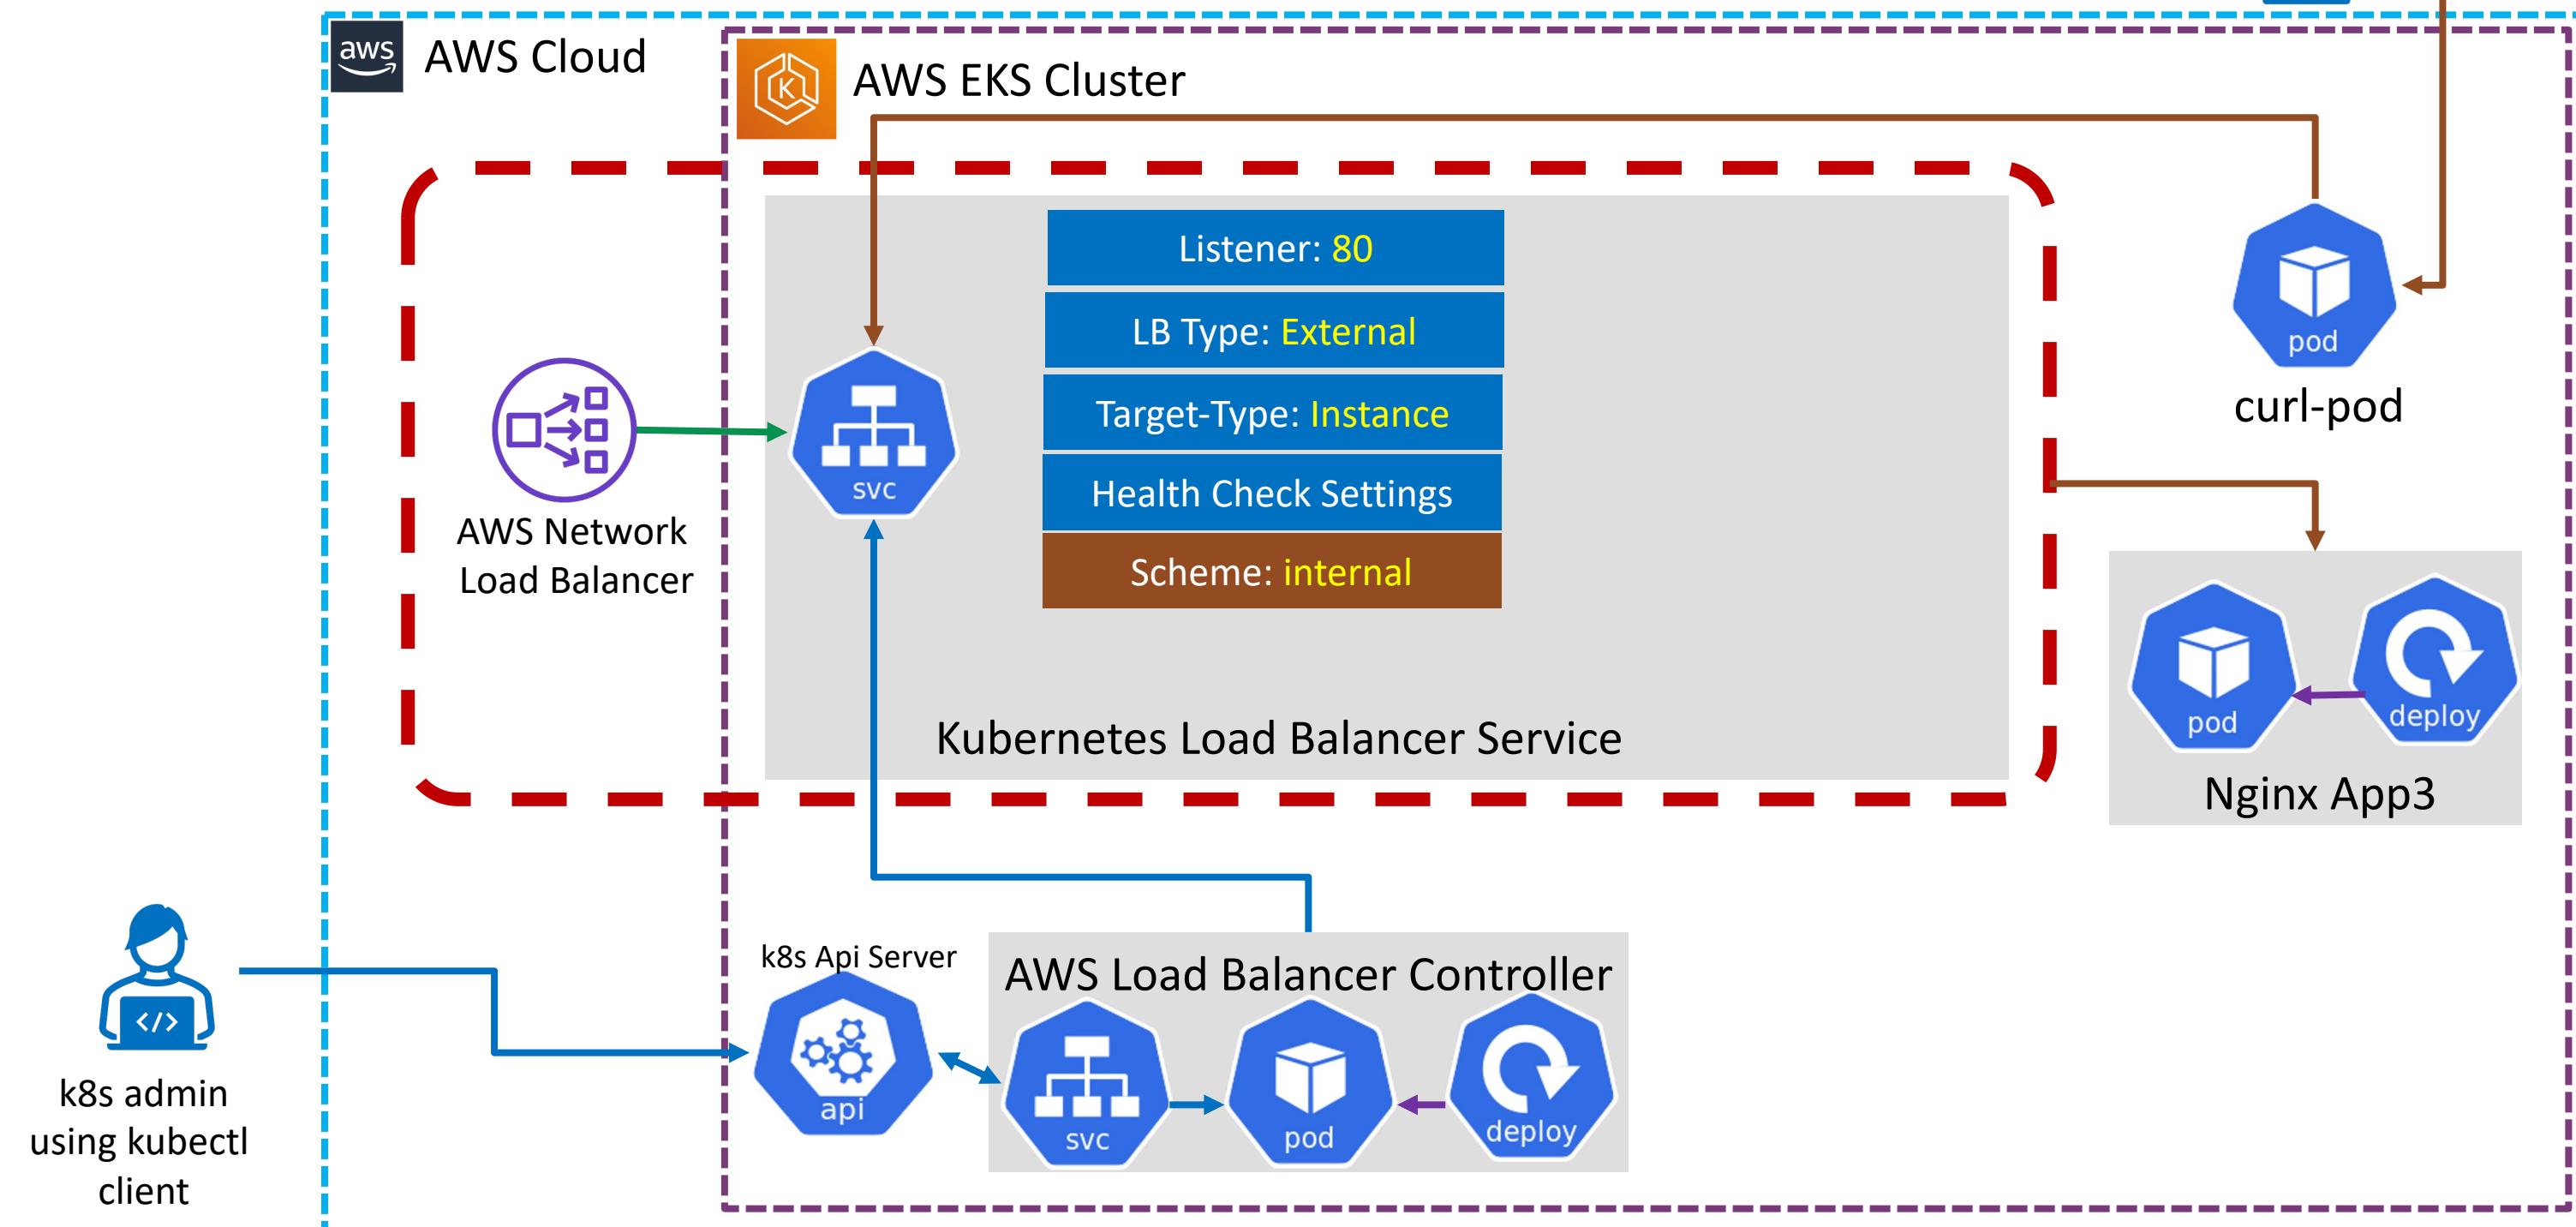
**Network Load Balancer**

**Internal LB**

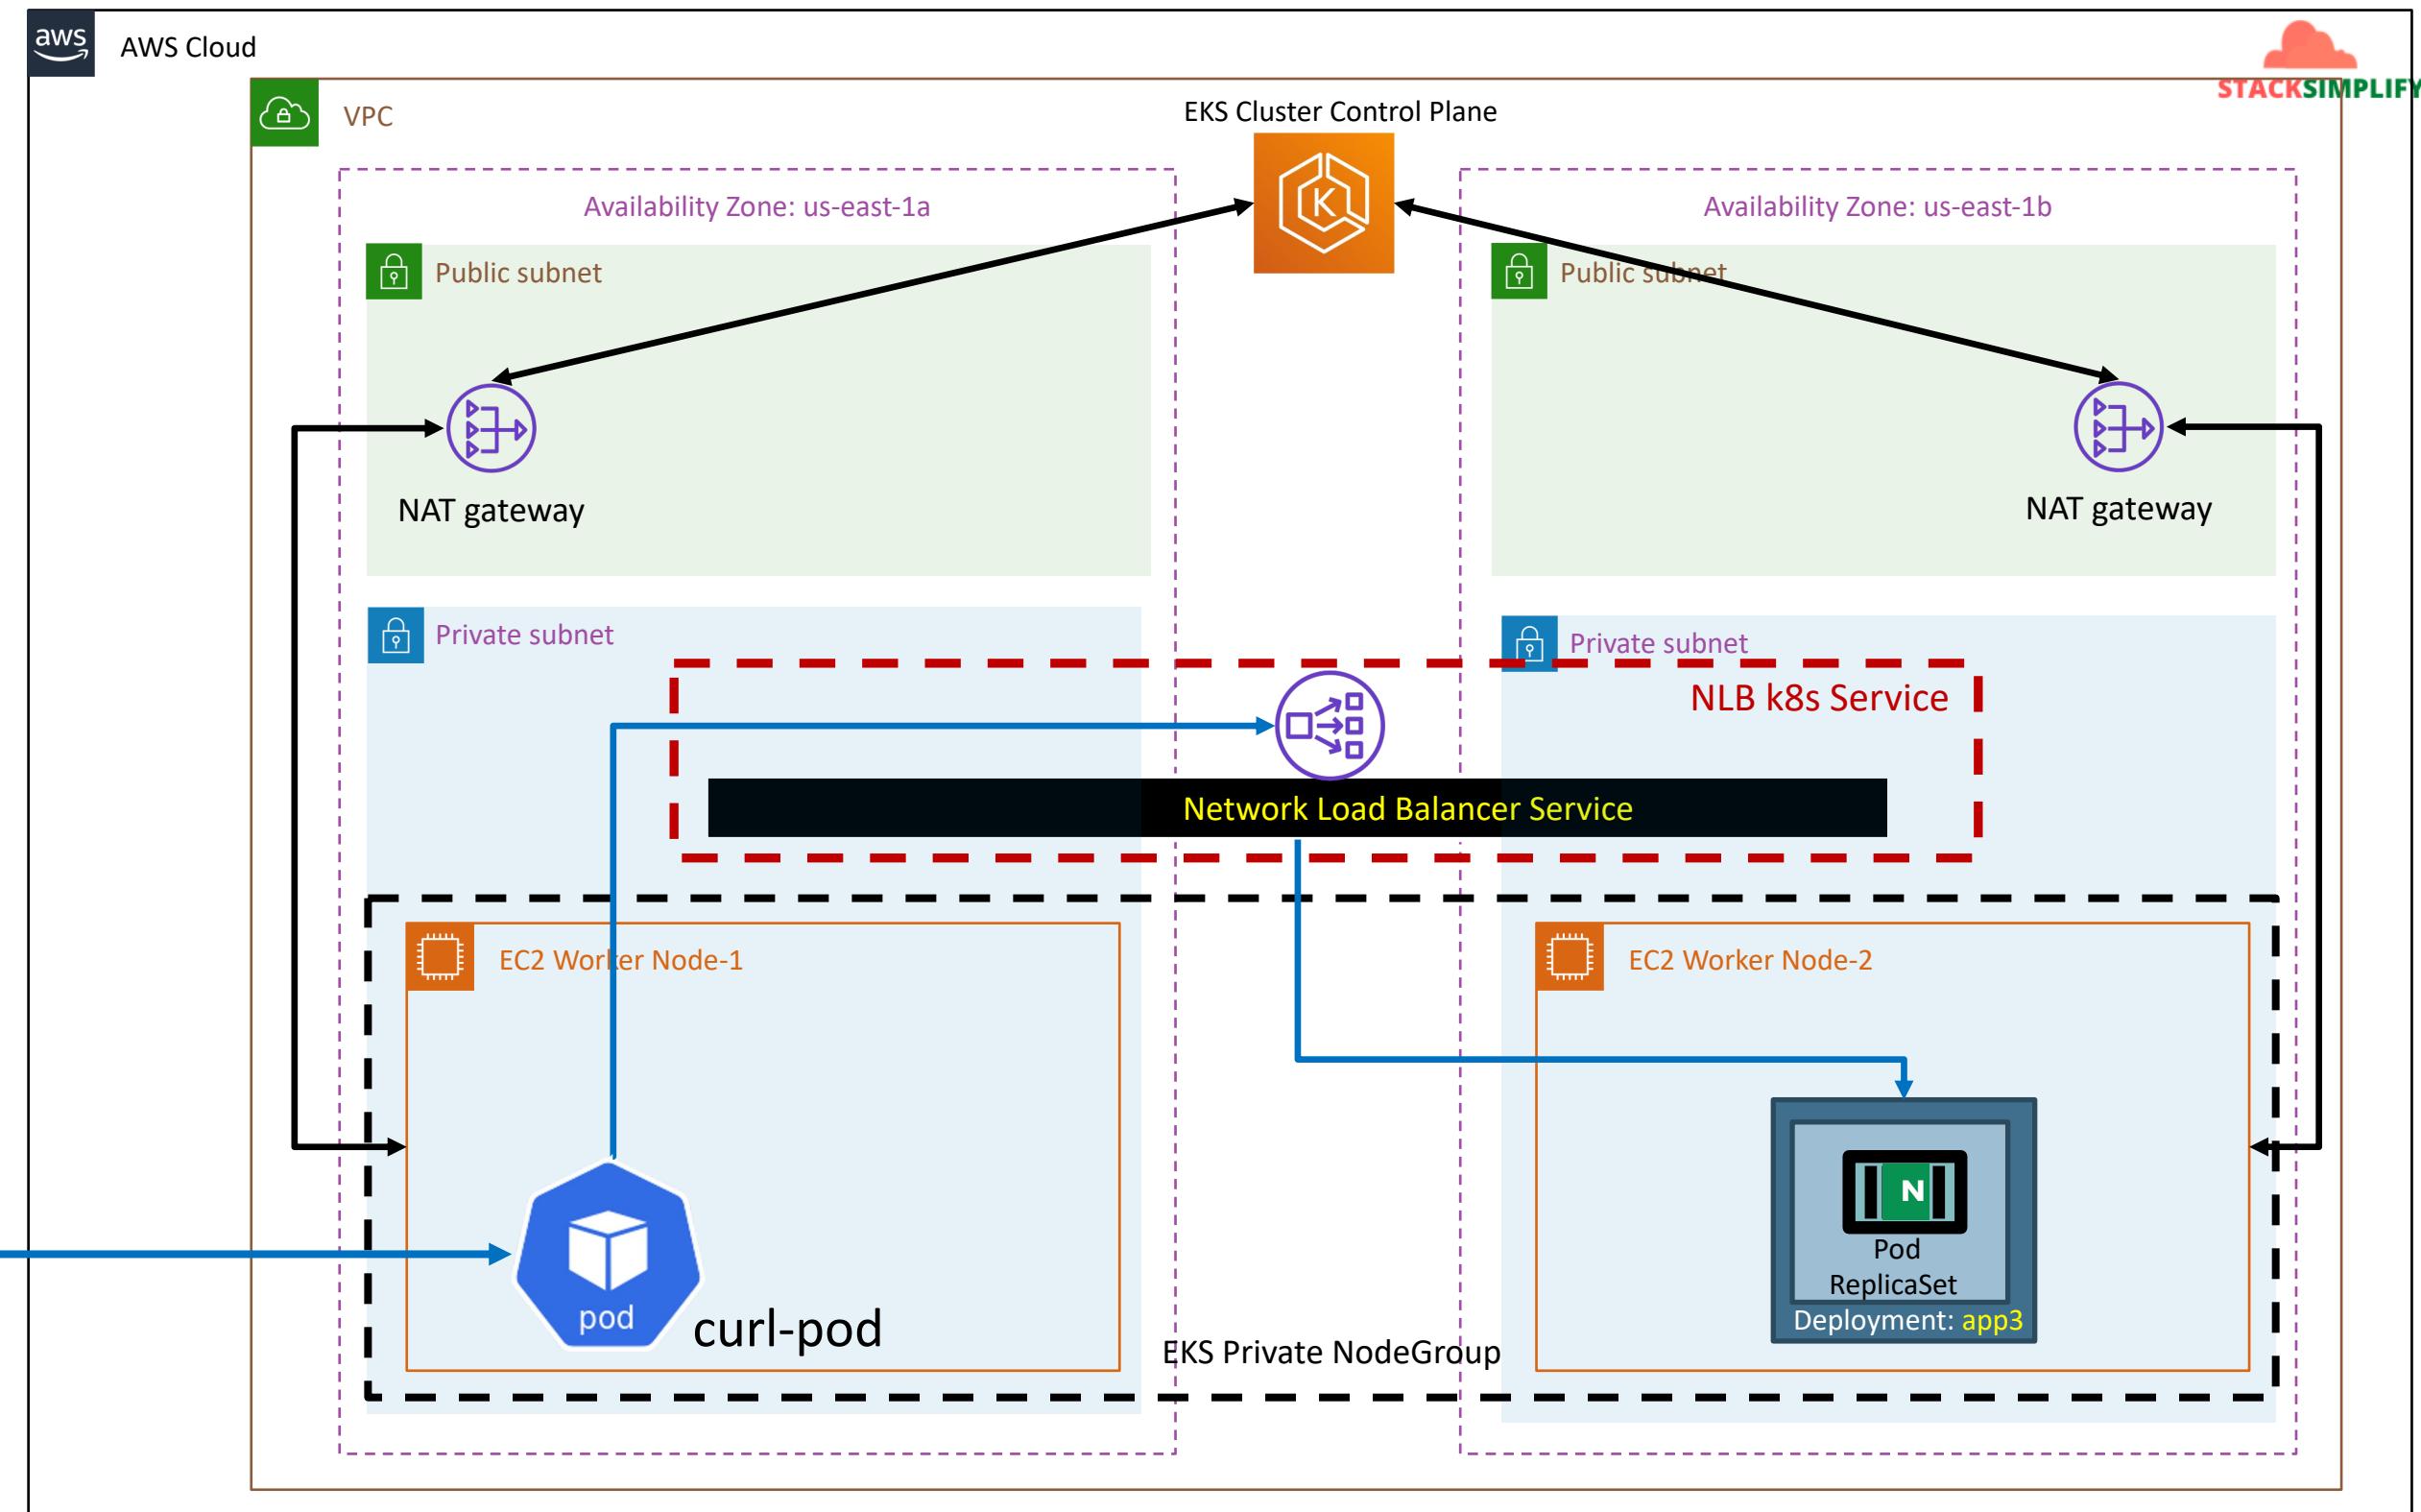


# AWS EKS NLB – Internal LB

k8s admin using  
kubectl client



# AWS EKS Network Design With Network Load Balancer Internal



# AWS EKS NLB - Internal

```
# Access Control  
service.beta.kubernetes.io/aws-load-balancer-scheme: "internal"
```

# Curl Pod

```
apiVersion: v1
kind: Pod
metadata:
  name: curl-pod
spec:
  containers:
  - name: curl
    image: curlimages/curl
    command: [ "sleep", "600" ]
```

We **cannot** connect to Internal LB directly from internet to test it.

We will deploy a **curl pod** in EKS Cluster so we can connect to curl pod in EKS Cluster and test the **Internal LB Endpoint** using **curl** command.

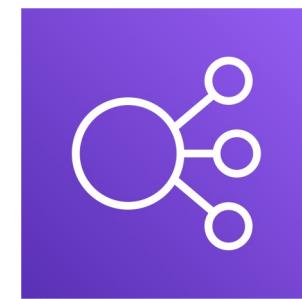
# Command to Test AWS Internal NLB using curl pod

```
# Deploy curl-pod
kubectl apply -f kube-manifests-curl

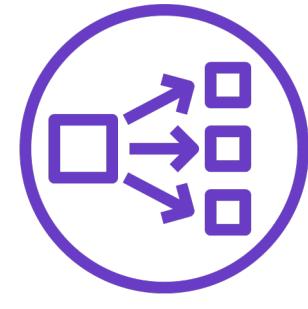
# Will open up a terminal session into the container
kubectl exec -it curl-pod -- sh

# We can now curl external addresses or internal services:
curl http://google.com/
curl <INTERNAL-NETWORK-LB-DNS>

# Internal Network LB Curl Test
curl lbc-network-lb-internal-demo-7031ade4ca457080.elb.us-east-1.amazonaws.com
```



Elastic Load  
Balancing



Network  
Load Balancer



Kubernetes  
Service



External  
DNS



# AWS EKS

## Load Balancer Controller

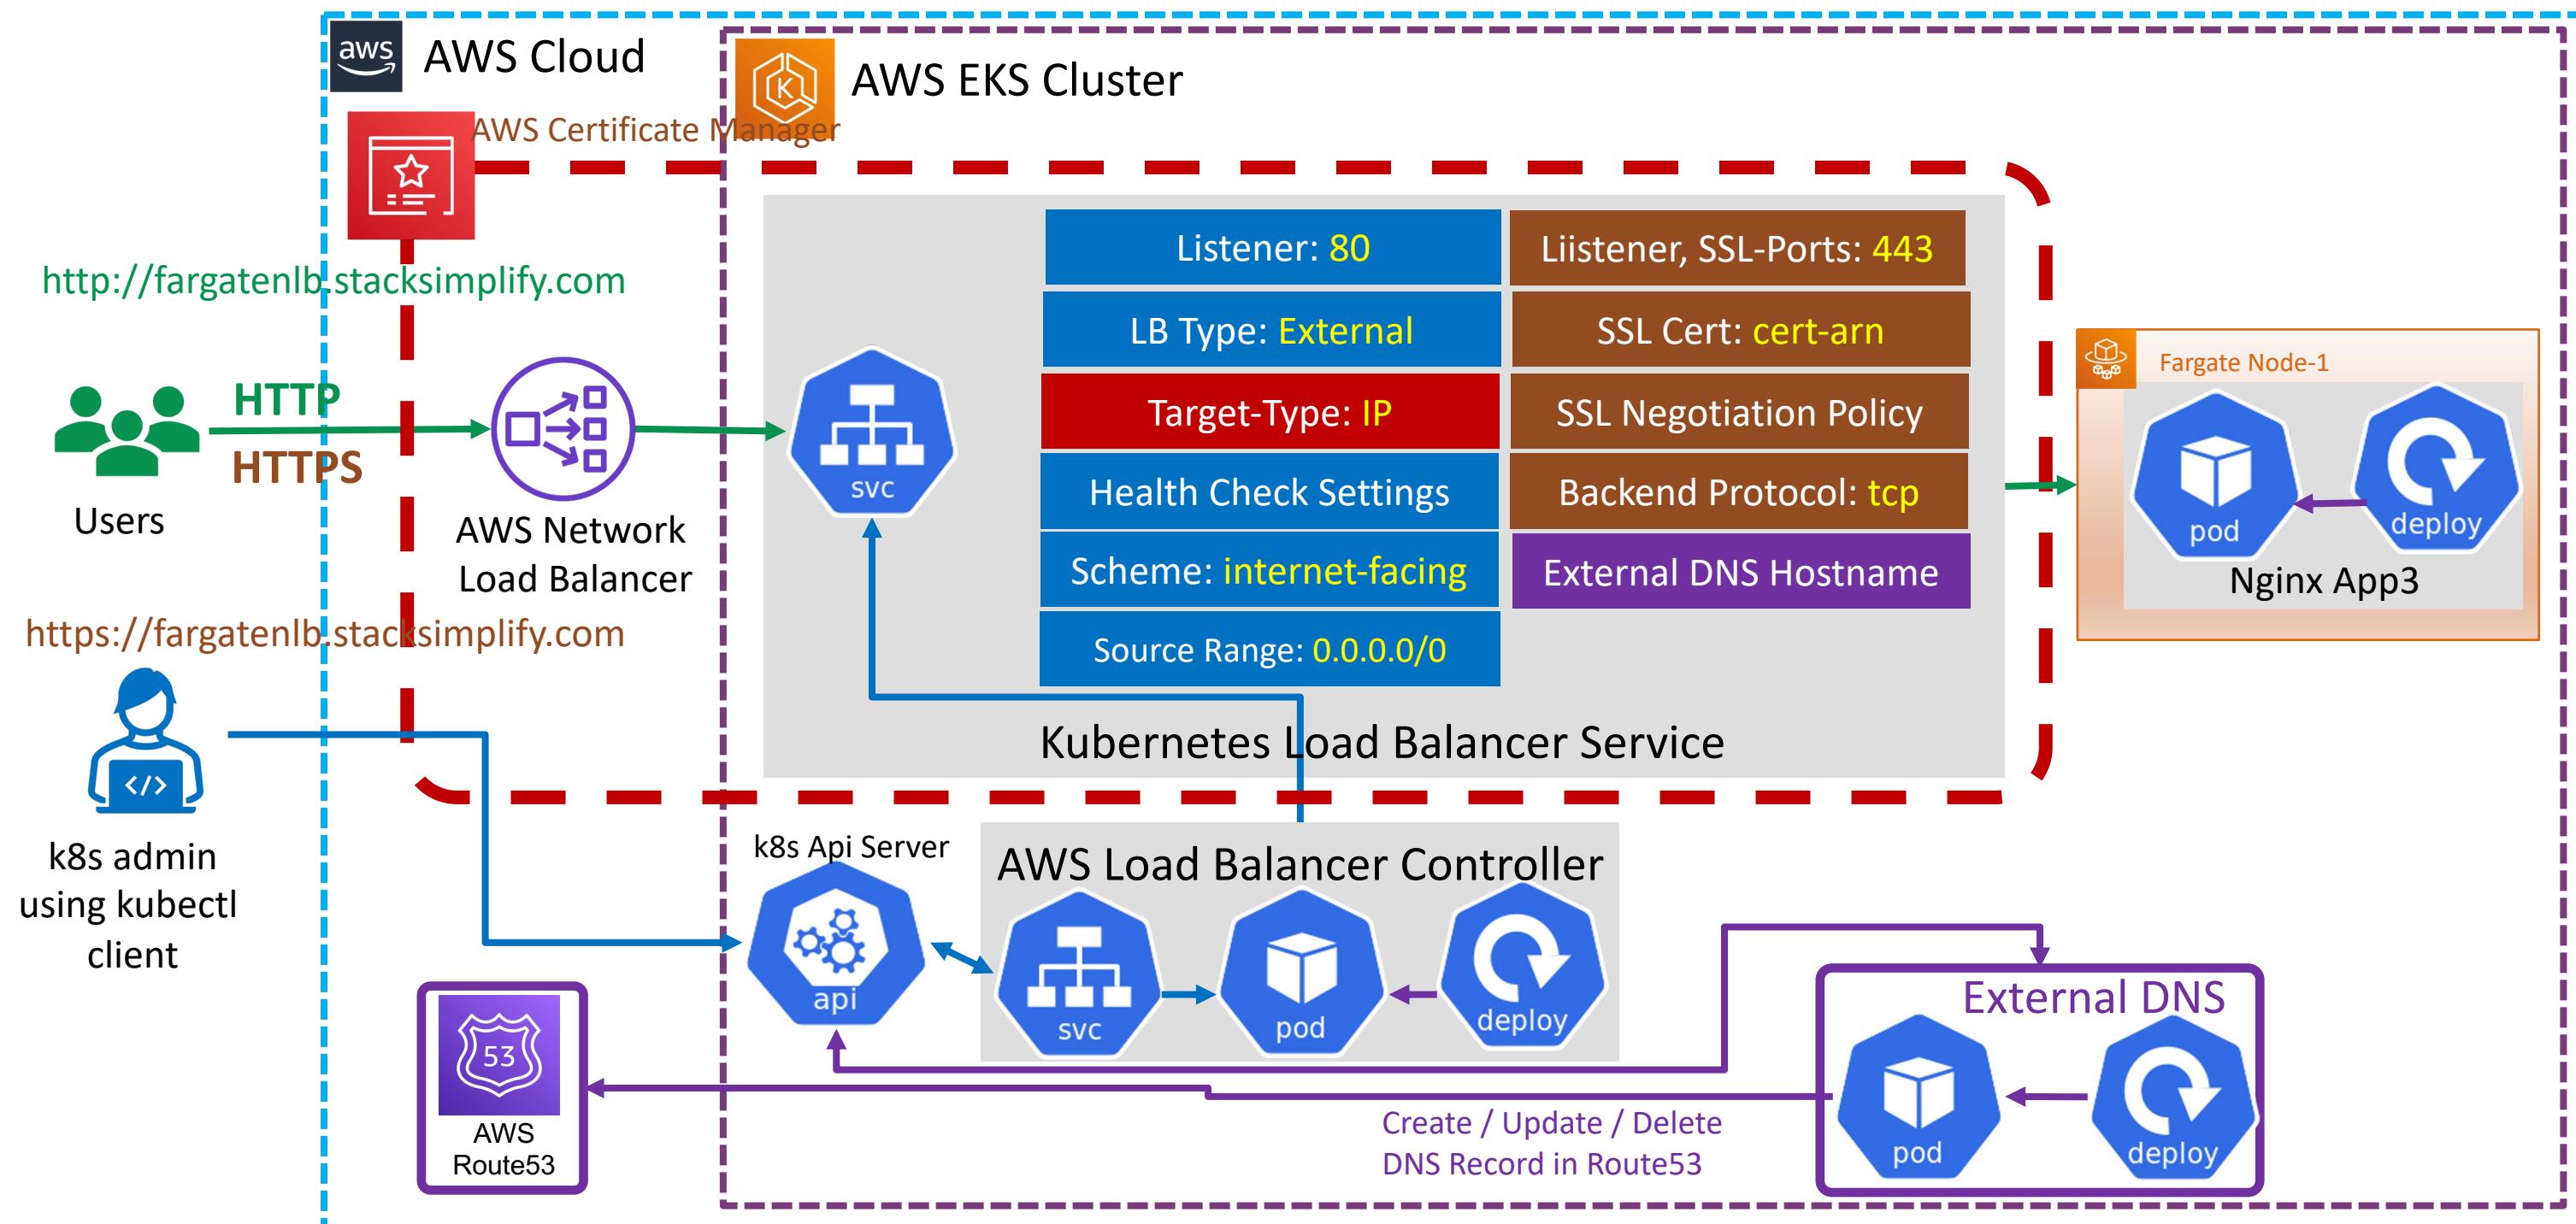
### Kubernetes Service

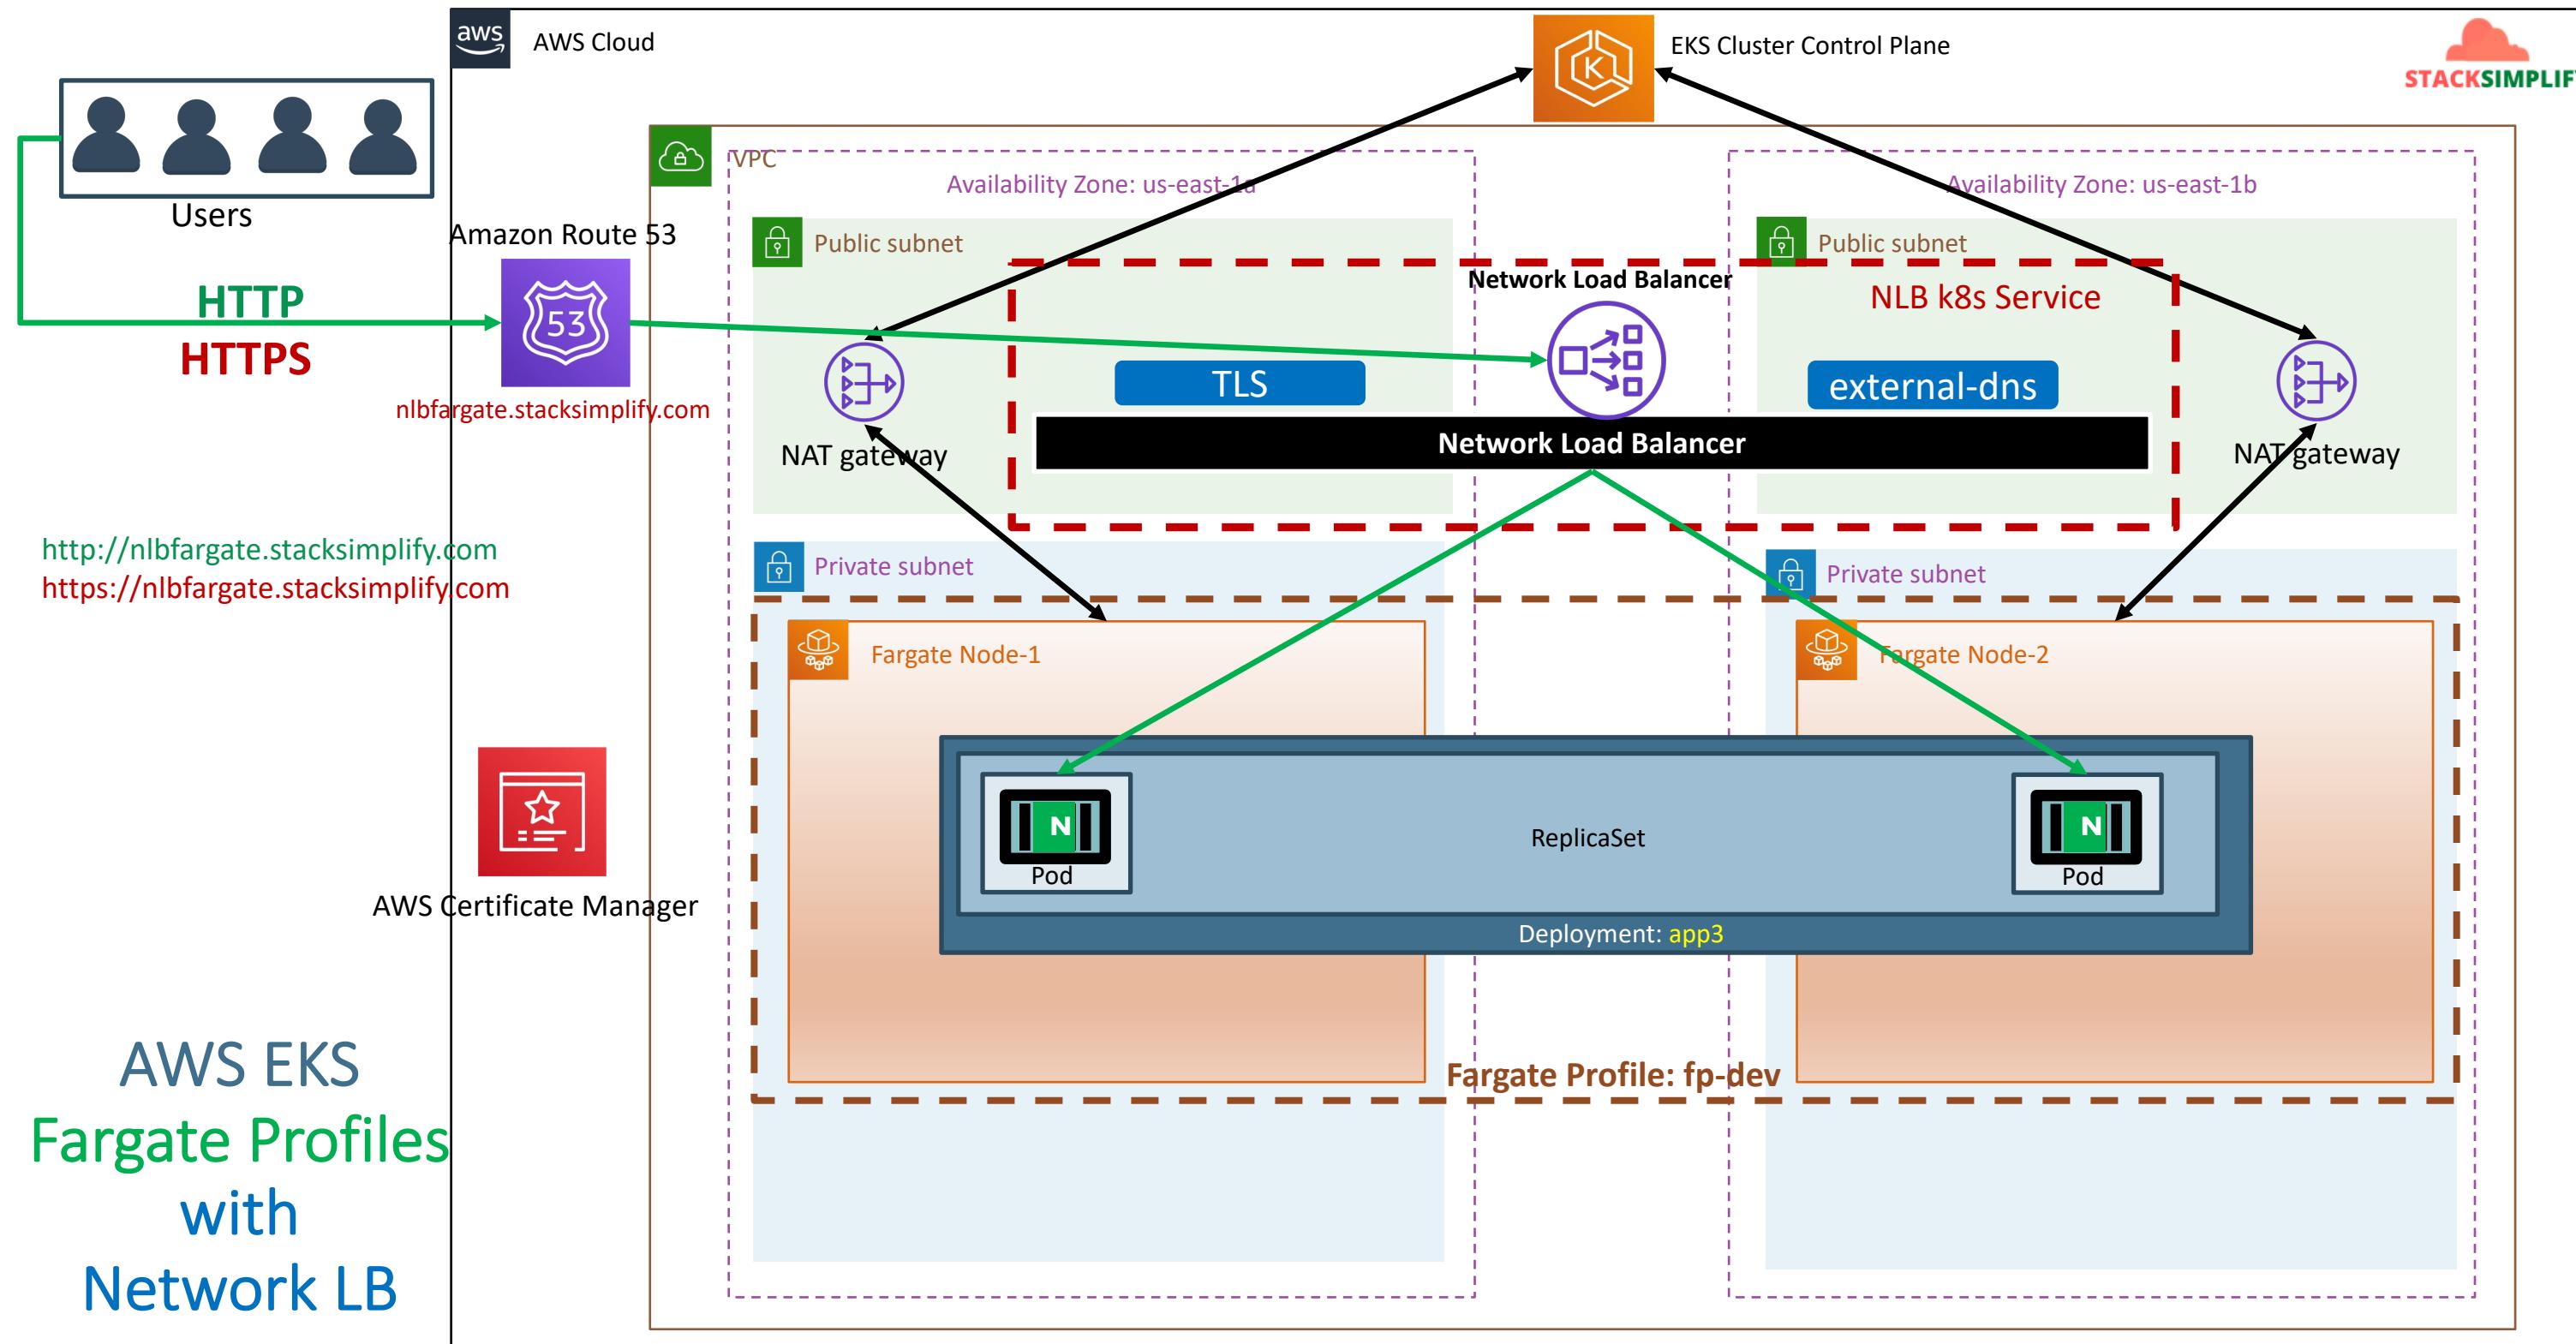
### Network Load Balancer

### with Workload running on Fargate



# AWS EKS NLB – Fargate Workloads





# AWS EKS NLB – Fargate Workloads

```
service.beta.kubernetes.io/aws-load-balancer-nlb-target-type: ip
```

Pod IP Address will be directly registered as targets in NLB Load Balancer Target Groups when we use Target Type as **IP Mode**

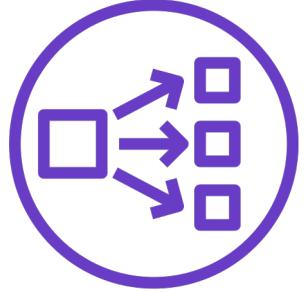
# AWS Load Balancer Controller & AWS Network Load Balancer – NEW TOPIC ADDITION - END



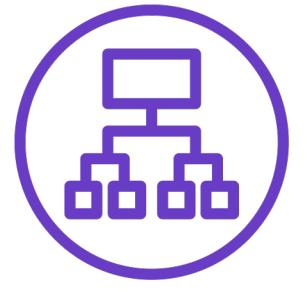
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



Application  
Load Balancer



Certificate  
Manager



Route53



Elastic Block  
Store



Amazon RDS

# AWS EKS

## ECR



# Elastic Container Registry



Fargate  
Profiles



Elastic Container  
Registry

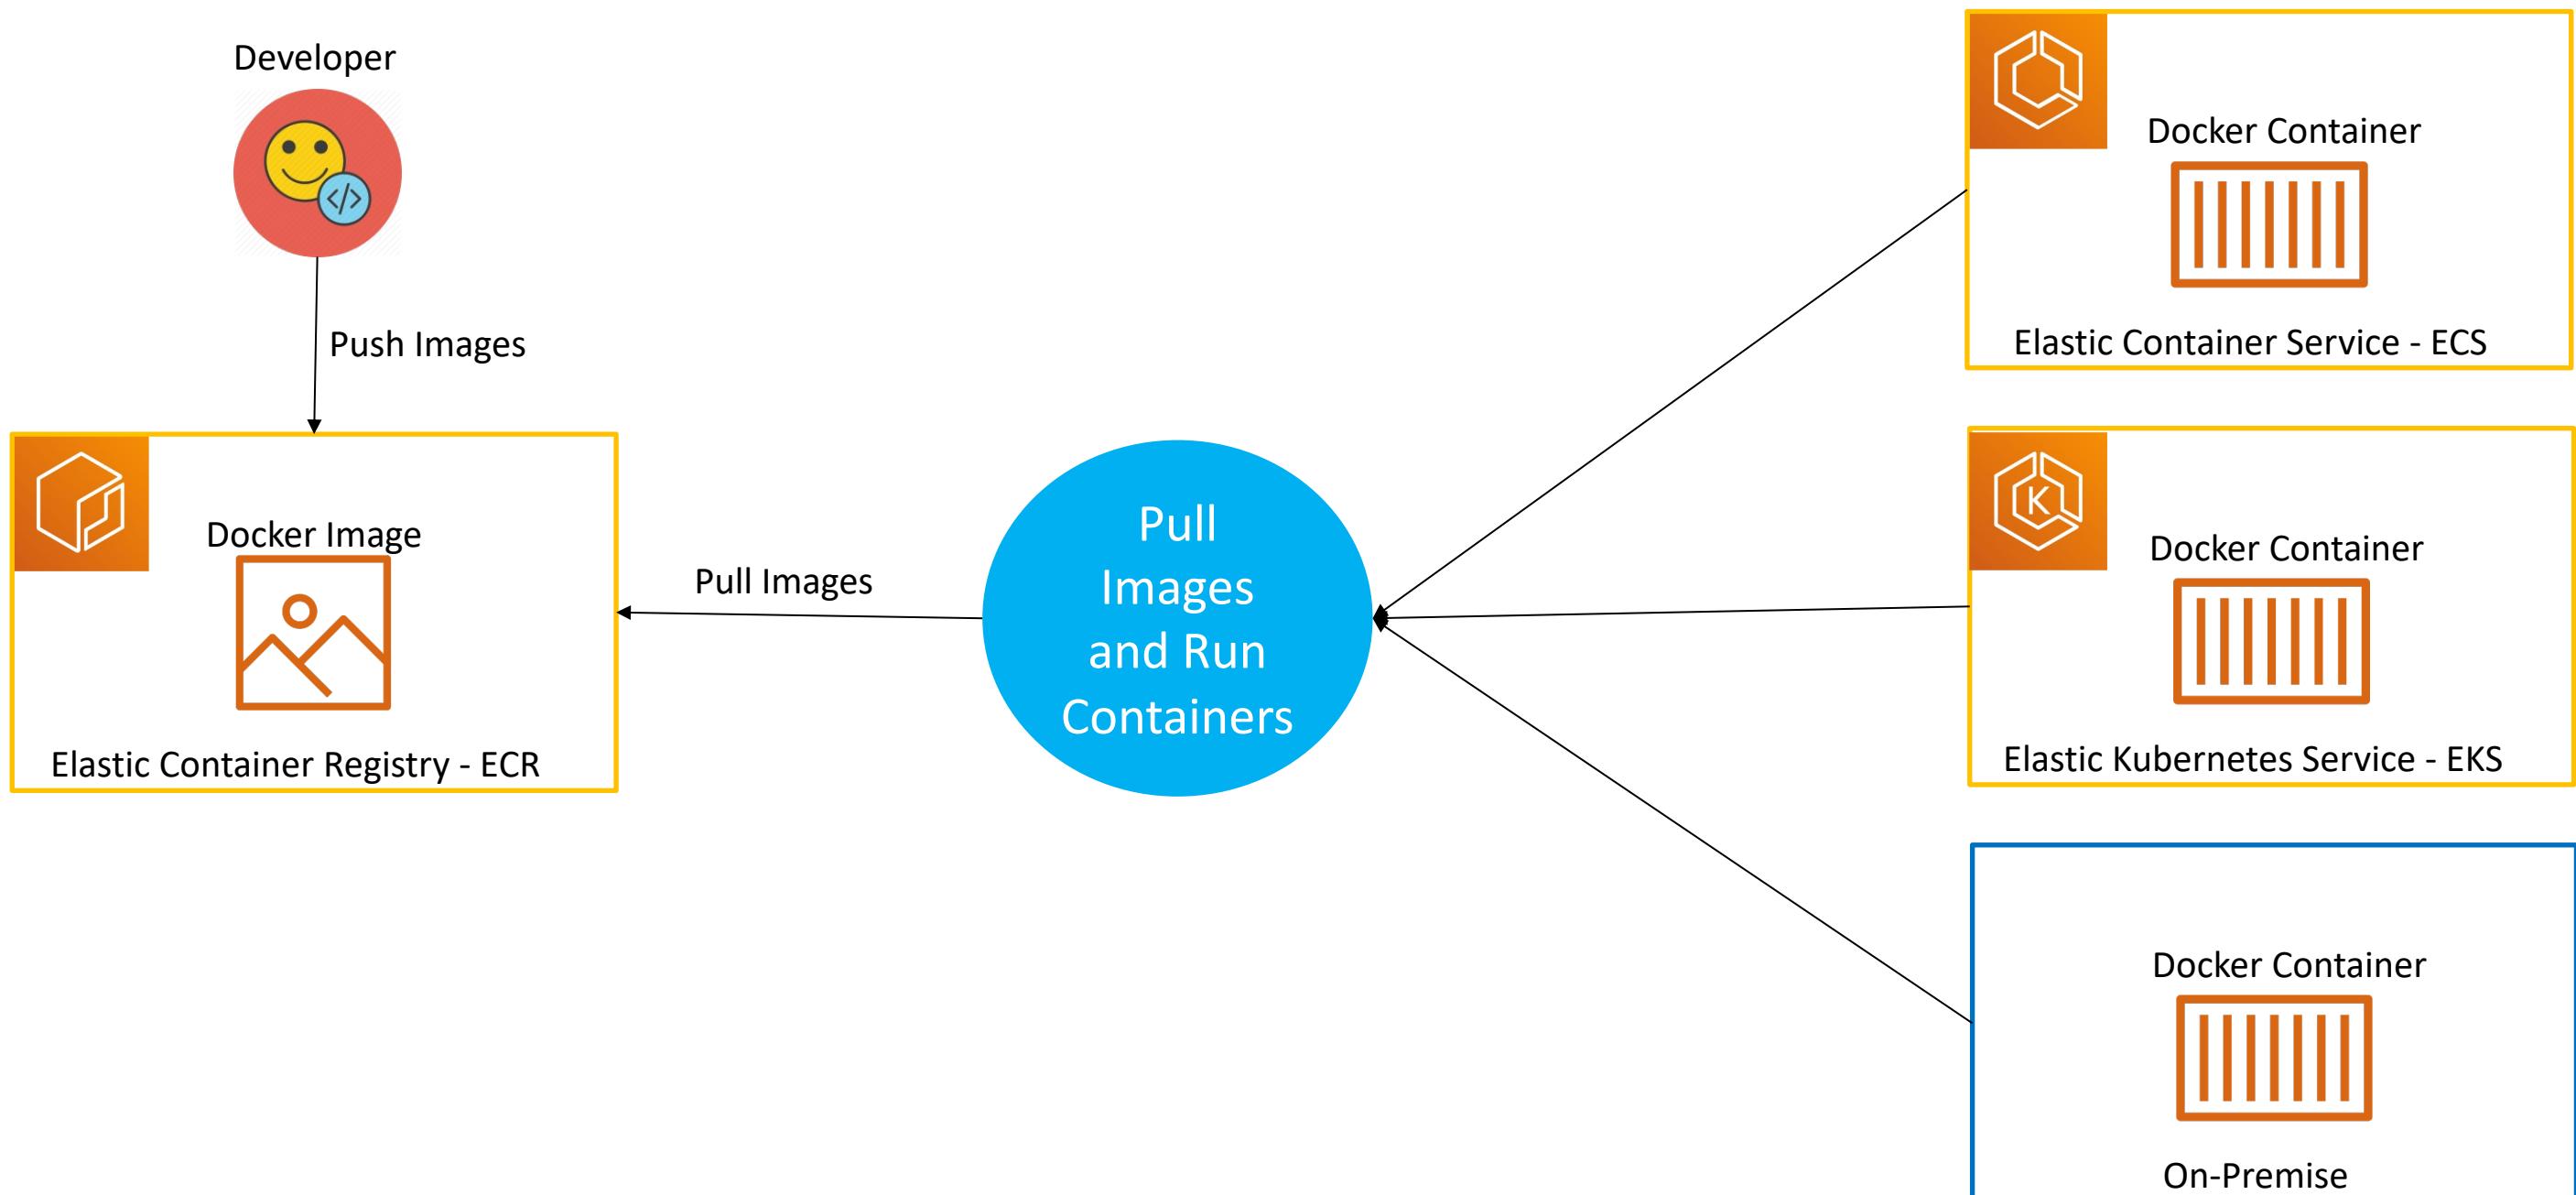
# Elastic Container Registry - ECR

- Elastic Container Registry (ECR) is a **fully-managed** Docker container registry that makes it easy for developers to store, manage, and deploy Docker container images.
- ECR is integrated with **Elastic Kubernetes Service (EKS)**, simplifying our development to production workflow.
- ECR **eliminates** the need to operate our own container repositories or worry about scaling the underlying infrastructure.
- ECR hosts our images in a **highly available** and scalable architecture, allowing us to reliably deploy containers for our applications.
- Integration with **AWS Identity and Access Management (IAM)** provides resource-level control of each repository.
- With Amazon ECR, there are **no upfront fees** or commitments. We pay only for the amount of data you store in your repositories and data transferred to the Internet.

# Elastic Container Registry - ECR

- Benefits
  - Full managed
  - Secure
  - Highly Available
  - Simplified Workflow

# How ECR Works?

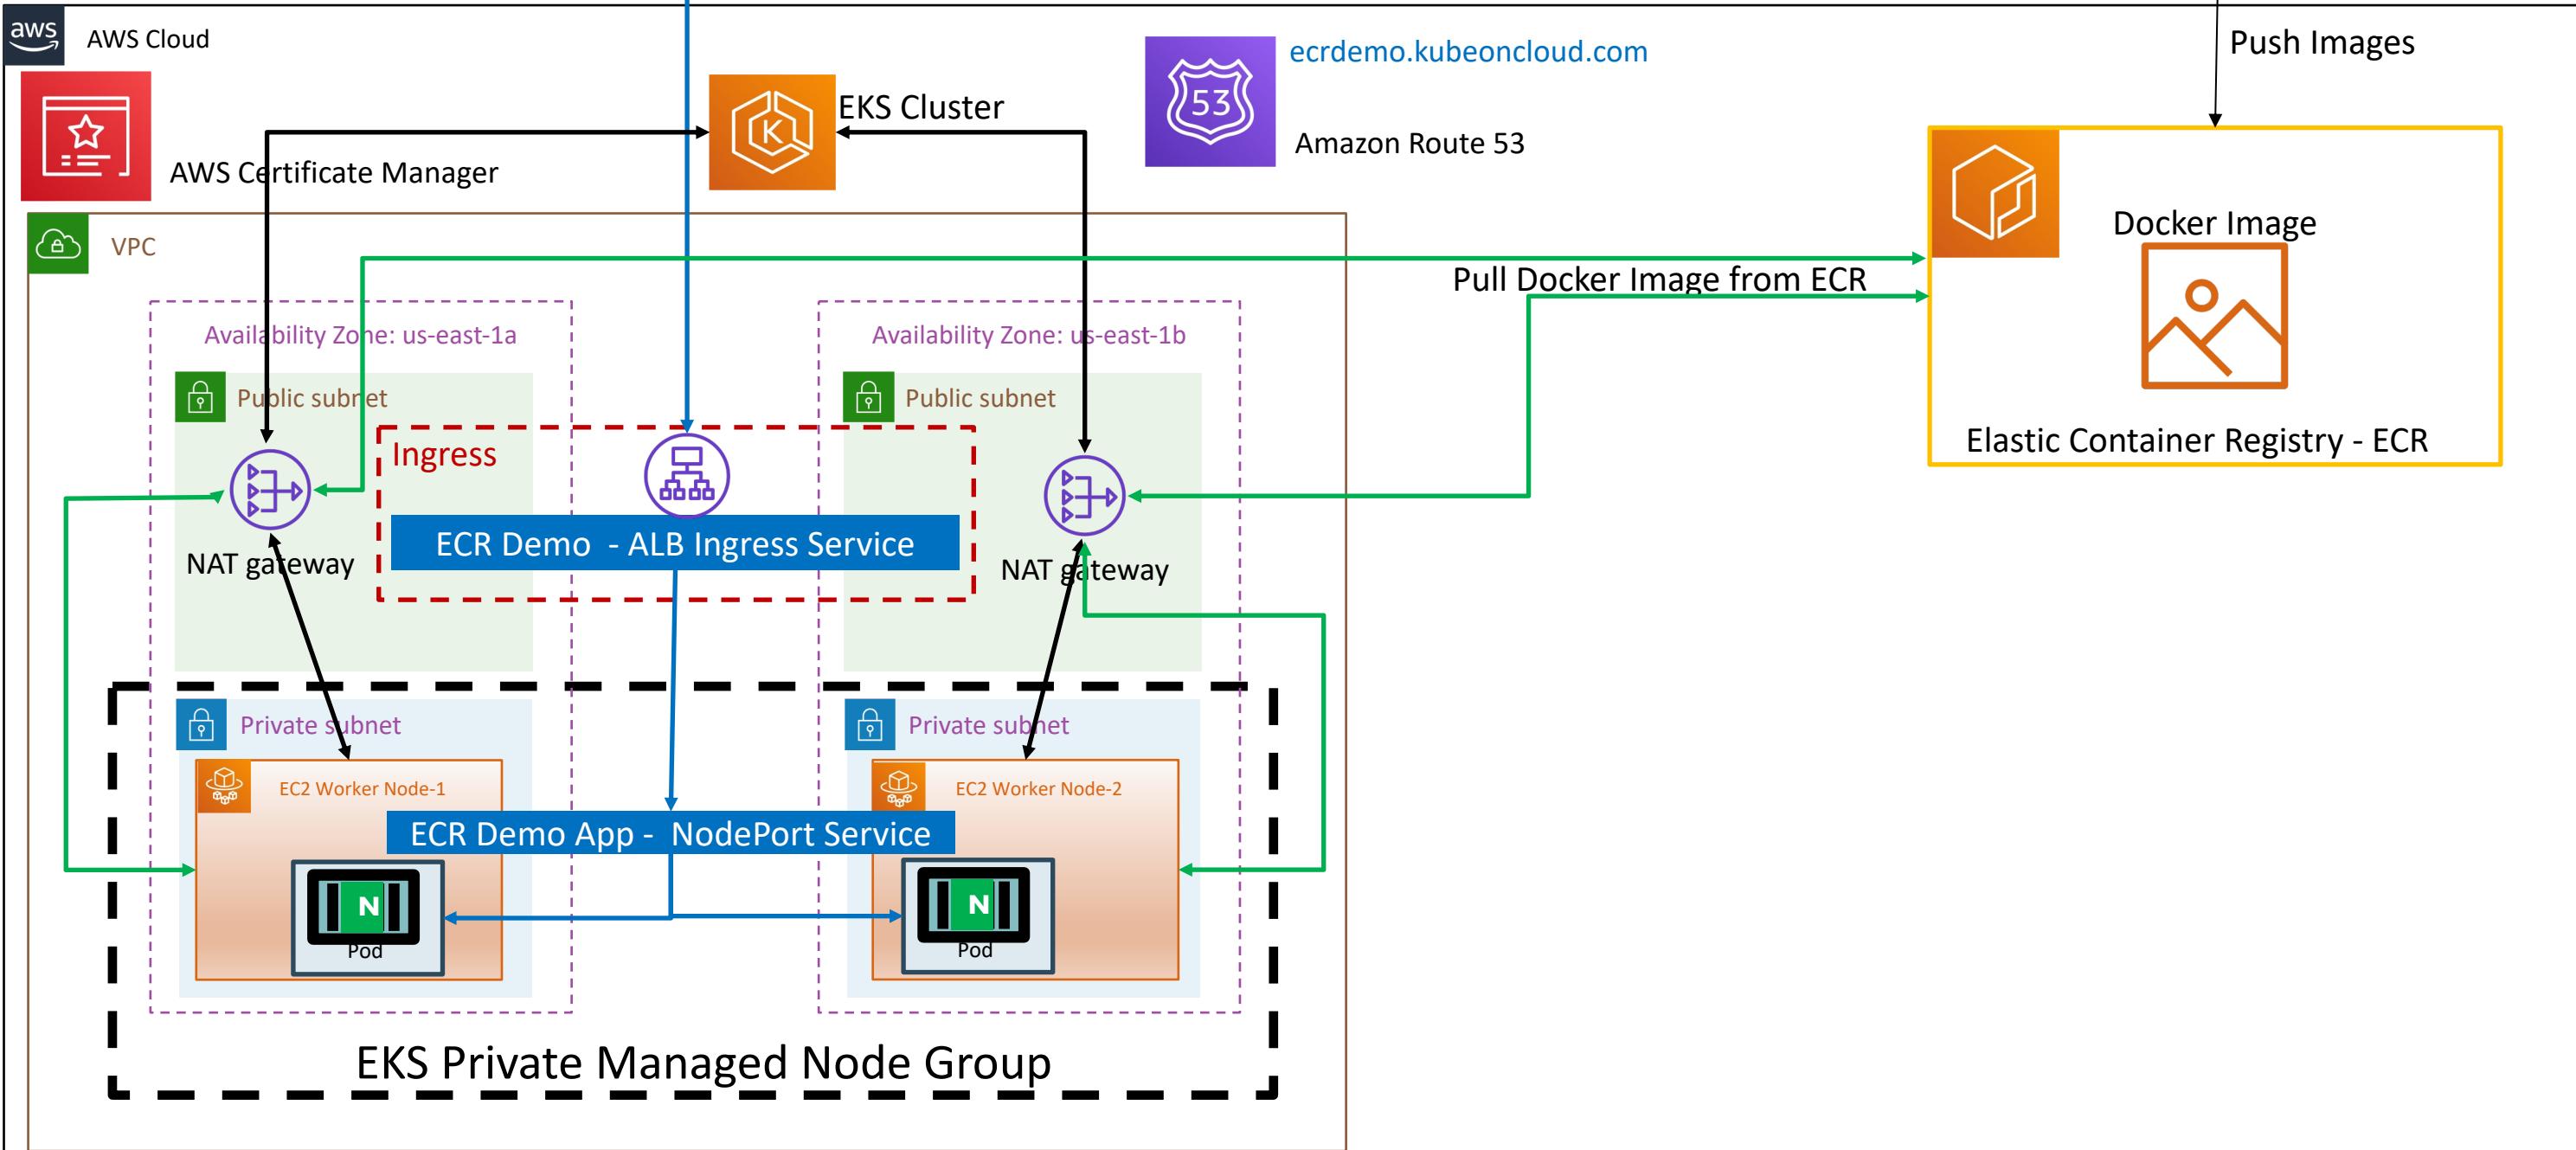


# EKS & ECR

Developer  
StackSimplify



http://ecrdemo.kubeoncloud.com

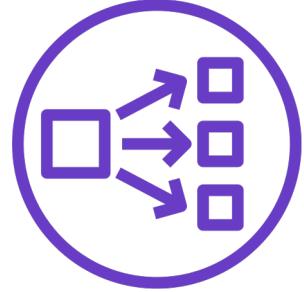




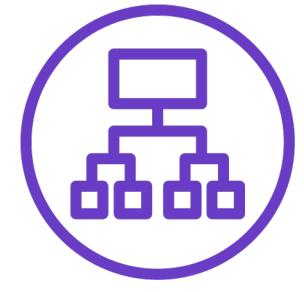
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



Application  
Load Balancer



Certificate  
Manager



Route53



Elastic Block  
Store



Amazon RDS

# AWS EKS

&

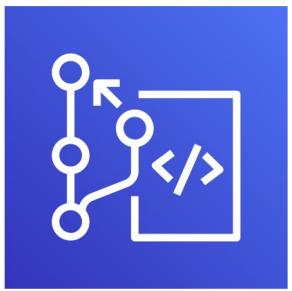
# AWS Developer Tools



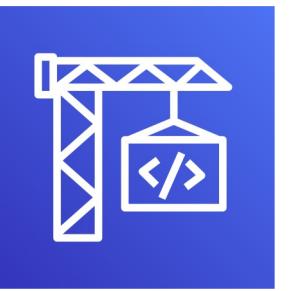
Fargate  
Profiles



Elastic Container  
Registry



Code  
Commit



Code  
Build



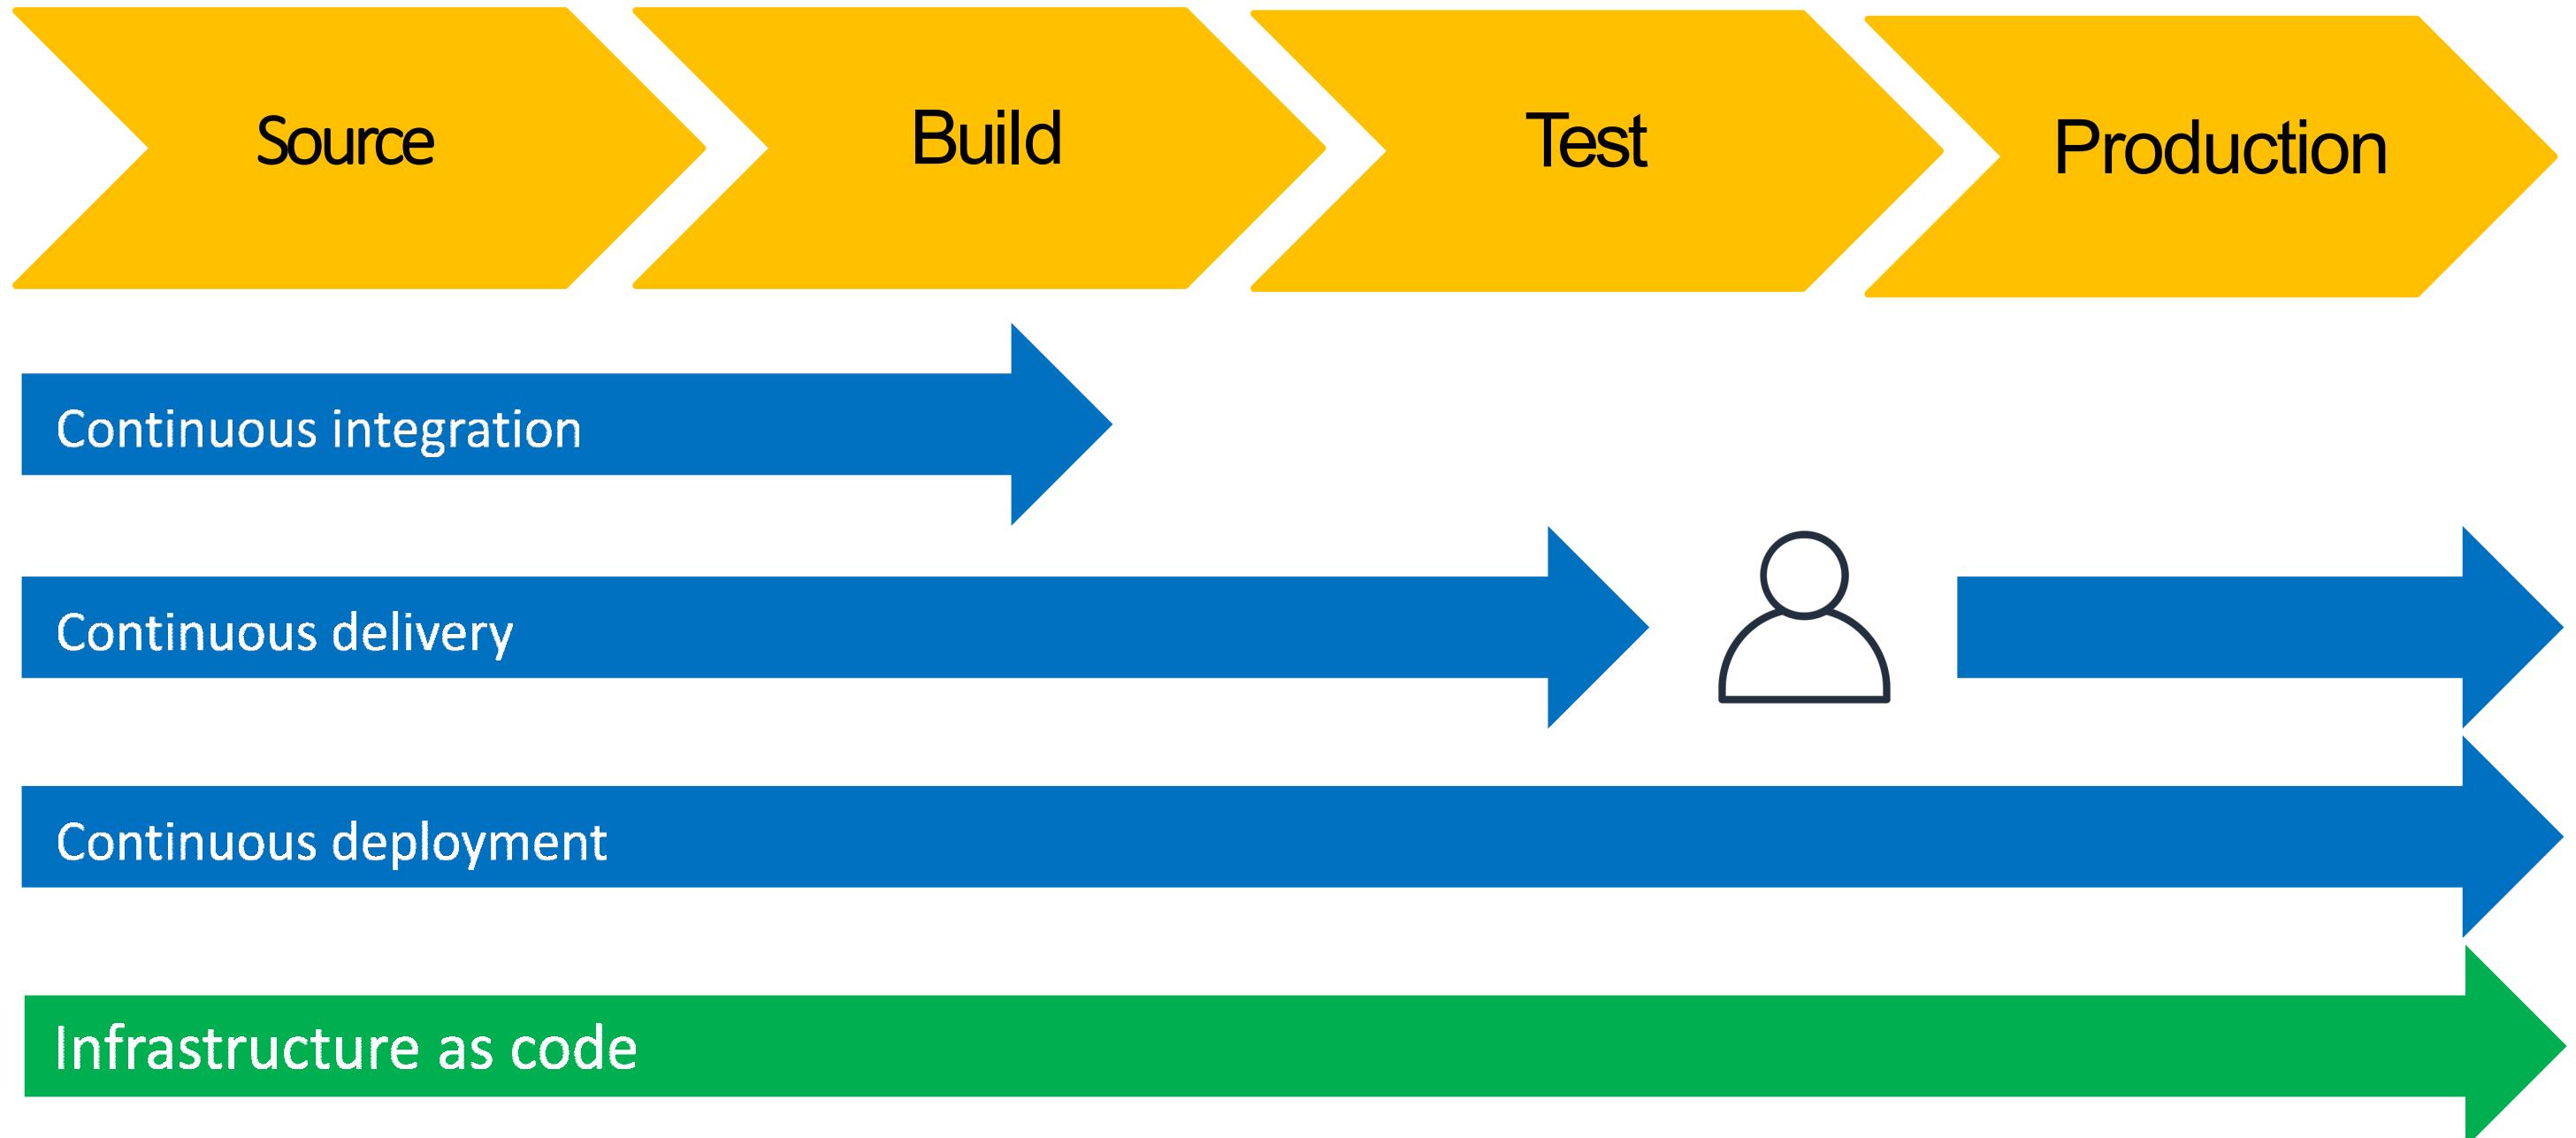
Code  
Pipeline

# Stages in Release Process



- |  |   |  |   |
|--|---|--|---|
| <ul style="list-style-type: none"><li>• Check-in source code</li><li>• Peer review new code</li><li>• Pull Request process</li></ul> | <ul style="list-style-type: none"><li>• Compile Code &amp; build artifacts (war ,jar, container images, Kubernetes manifest files)</li><li>• Unit Tests</li></ul> | <ul style="list-style-type: none"><li>• Integration tests with other systems.</li><li>• Load Testing</li><li>• UI Tests</li><li>• Security Tests</li><li>• Test Environments (Dev, QA and Staging)</li></ul> | <ul style="list-style-type: none"><li>• Deployment to production environments</li><li>• Monitor code in production to quickly detect errors</li></ul> |
|--|---|--|---|

# Stages in Release Process



# Continuous Integration



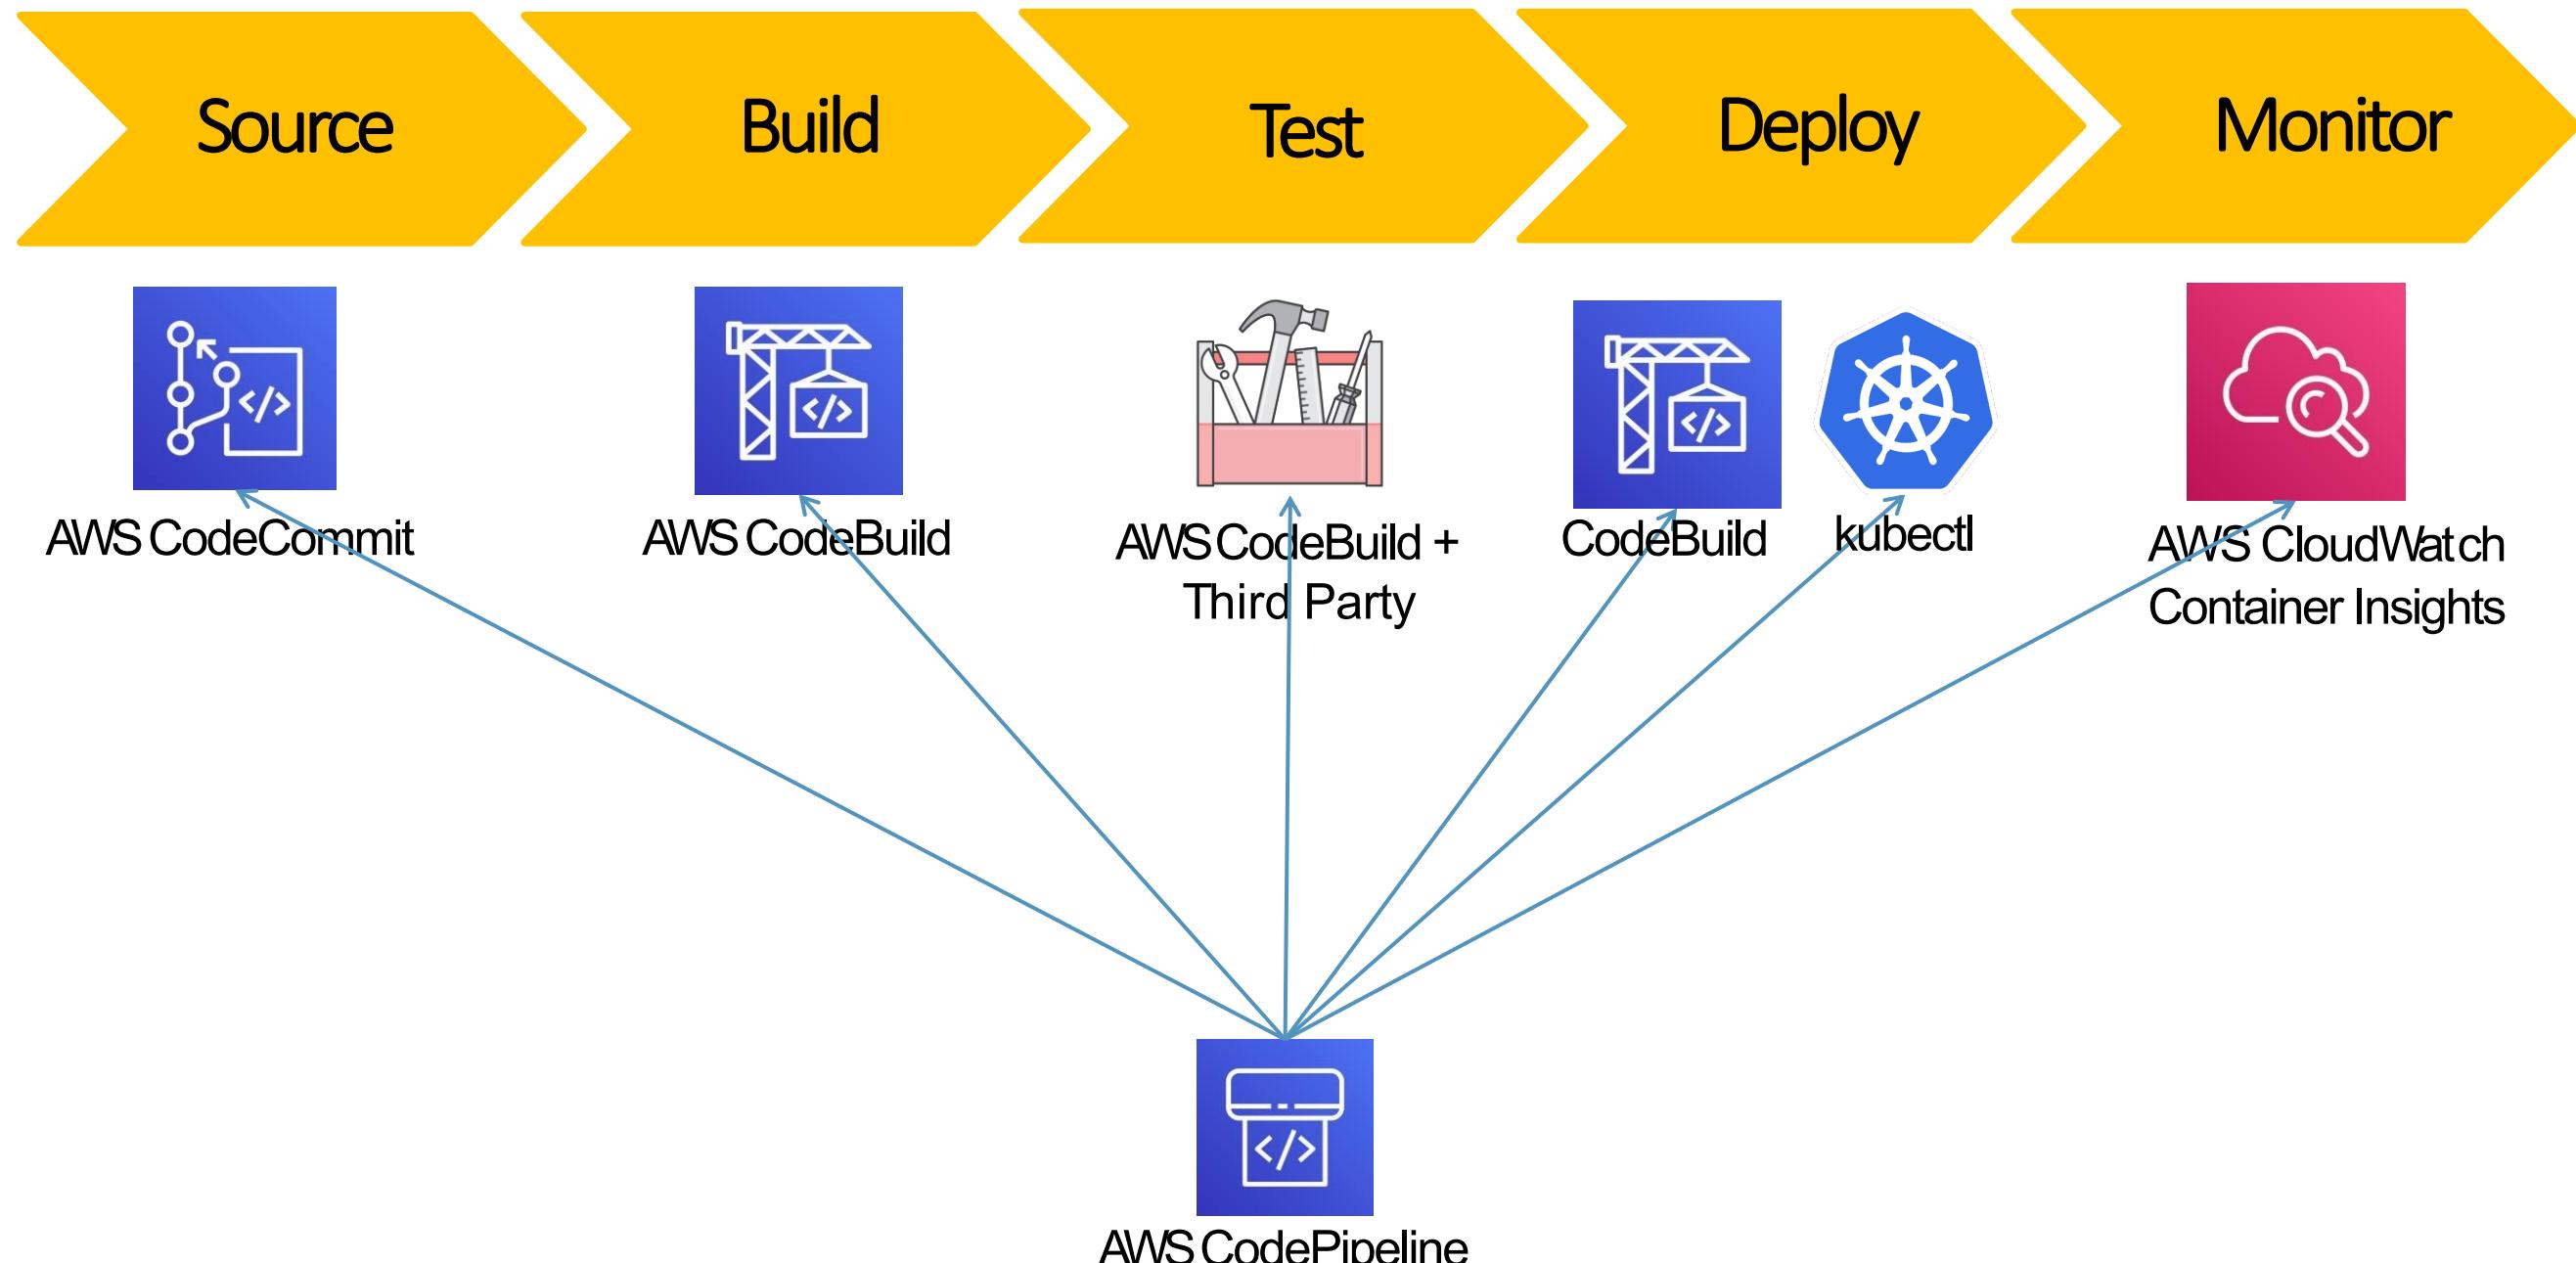
- Automatically kick off a new release when new code is checked-in
- Build and test code in a consistent, repeatable environment
- Continually have an artifact ready for deployment

# Continuous Delivery

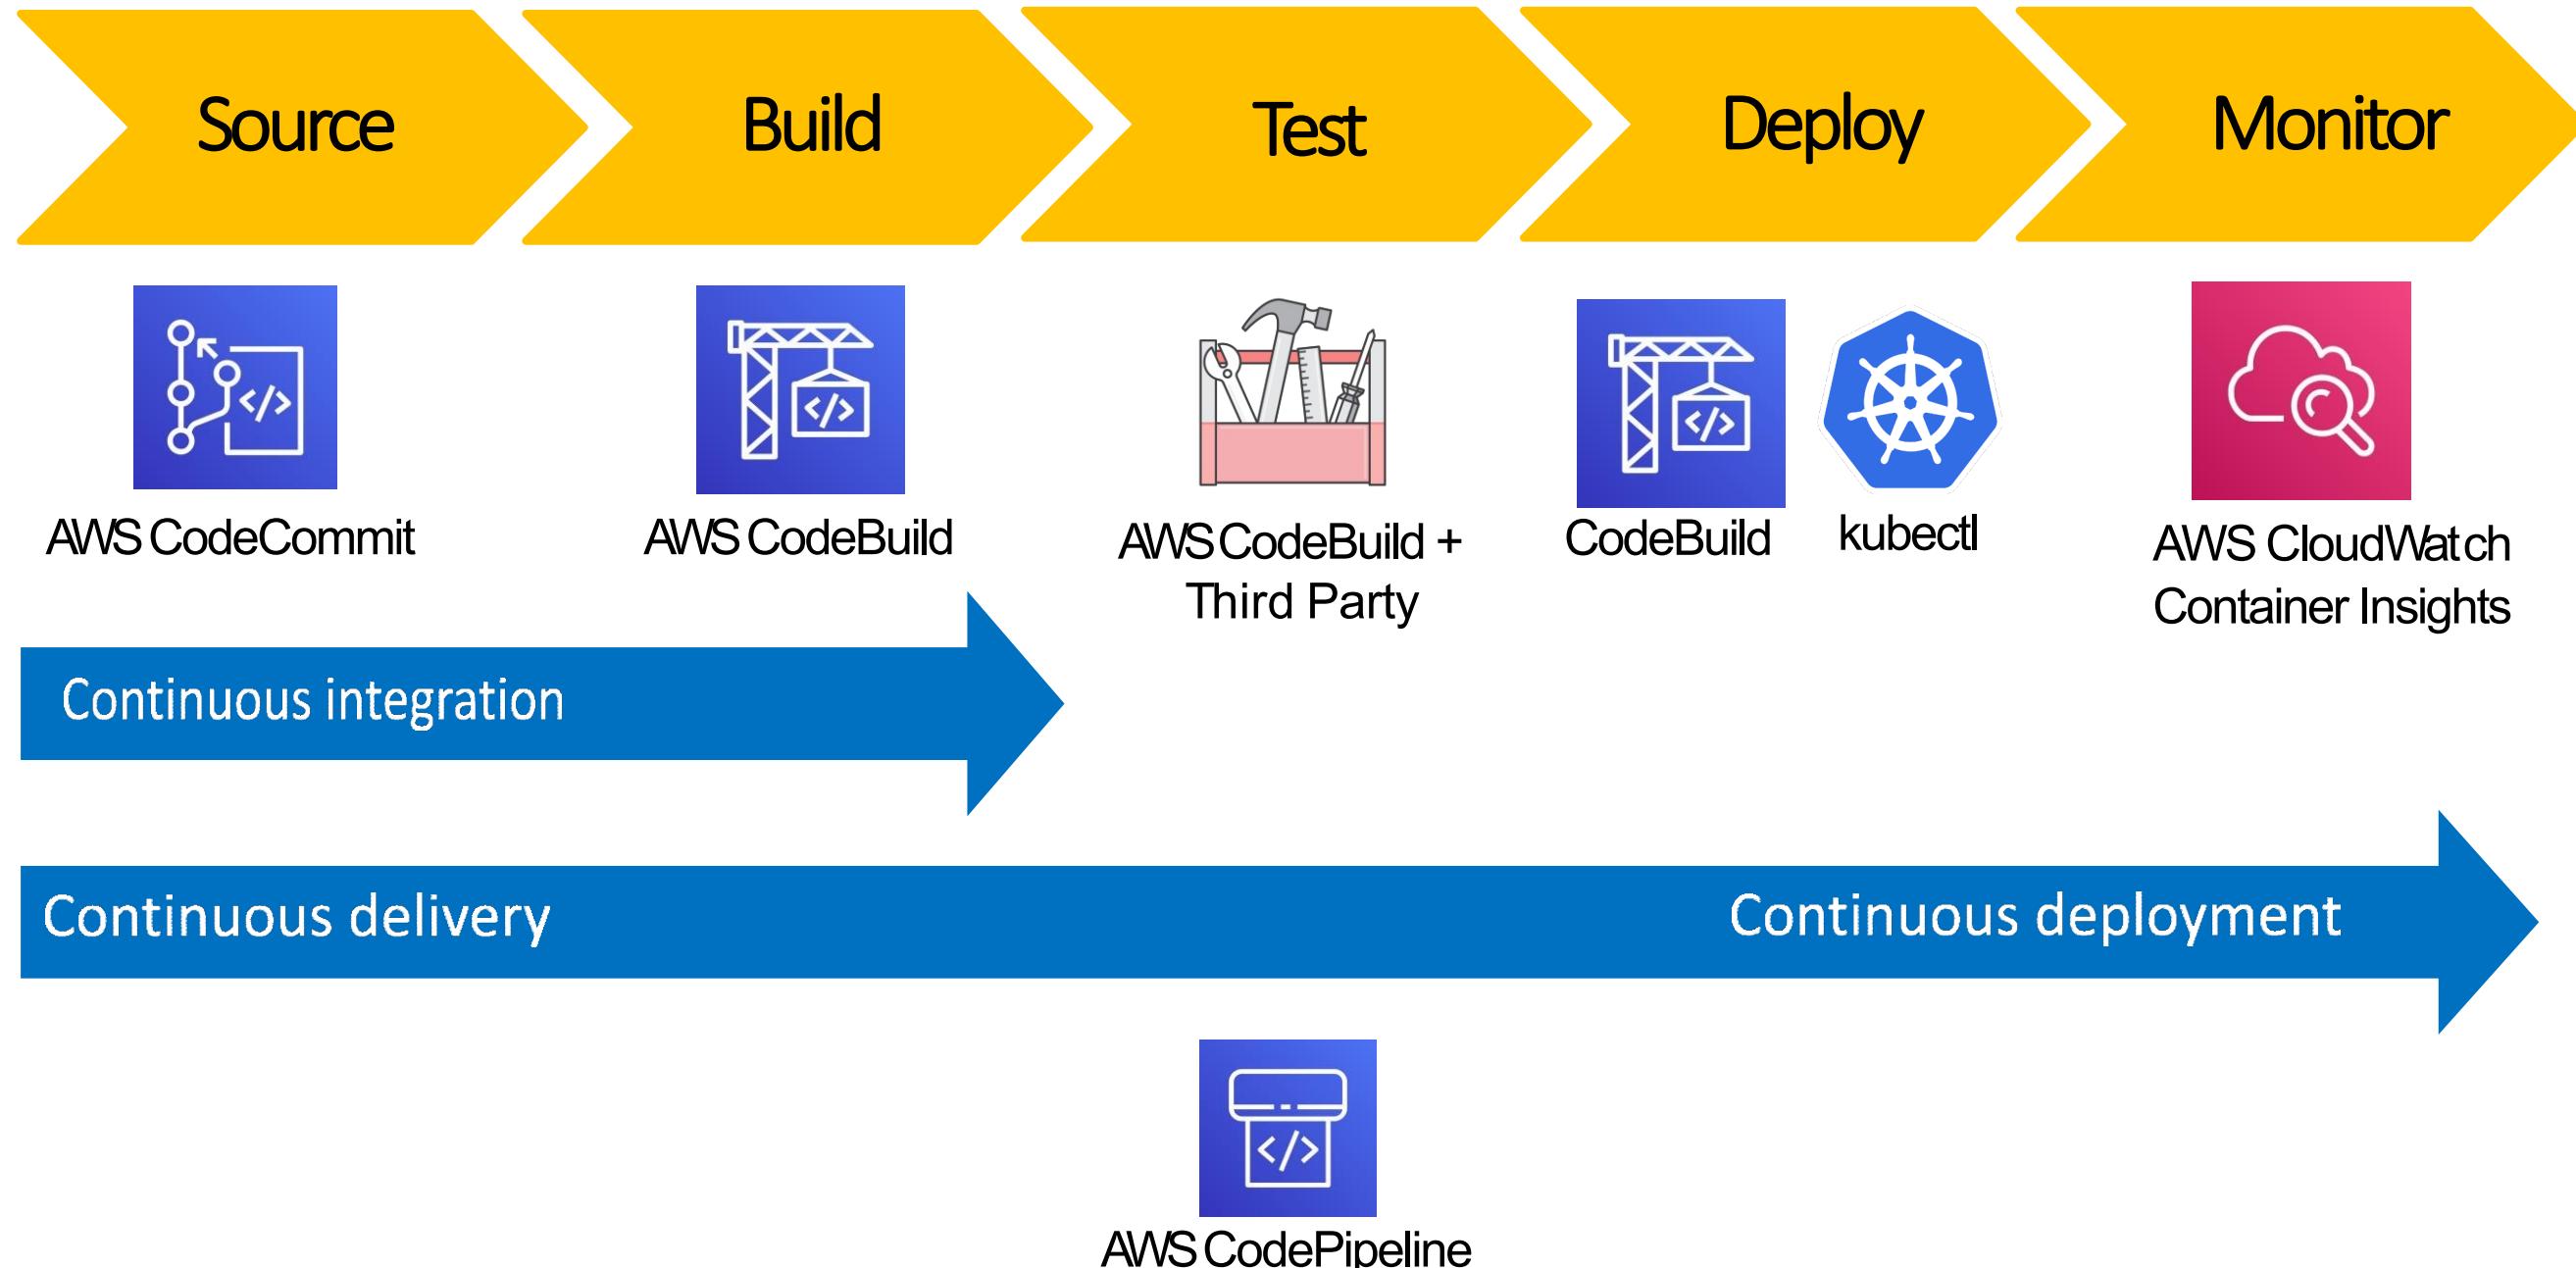


- Automatically deploy new changes to staging environments for testing
- Deploy to production safely without affecting customers
- Deliver to customers faster
- Increase deployment frequency, and reduce change lead time and change failure rate

# AWS Developer Tools or Code Services



# AWS Developer Tools or Code Services



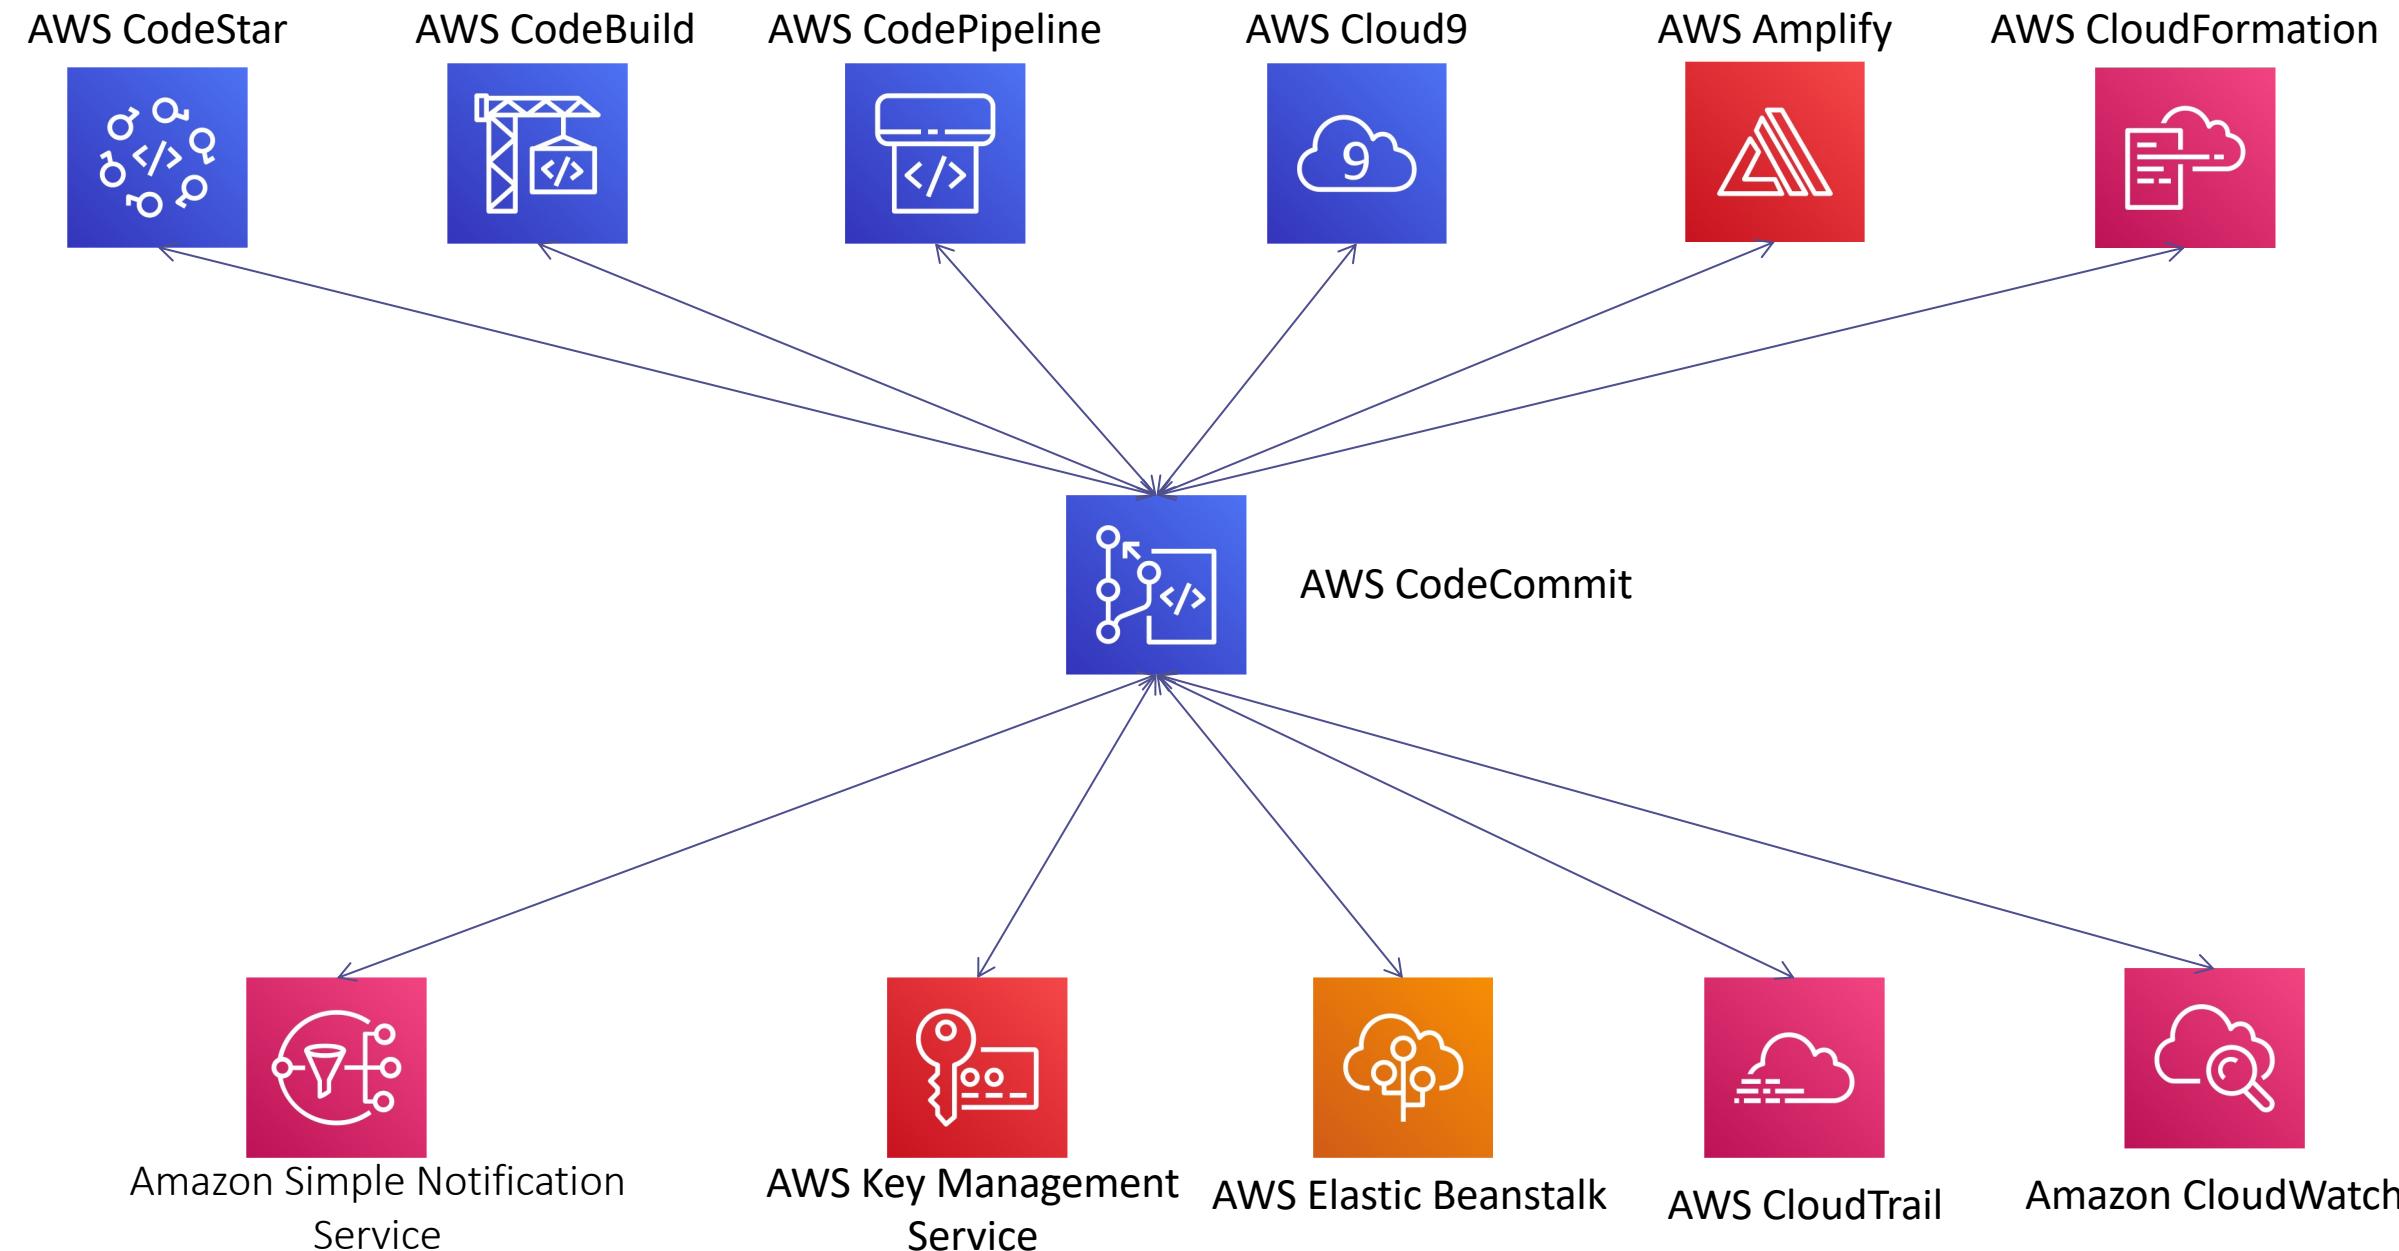
# AWS CodeCommit



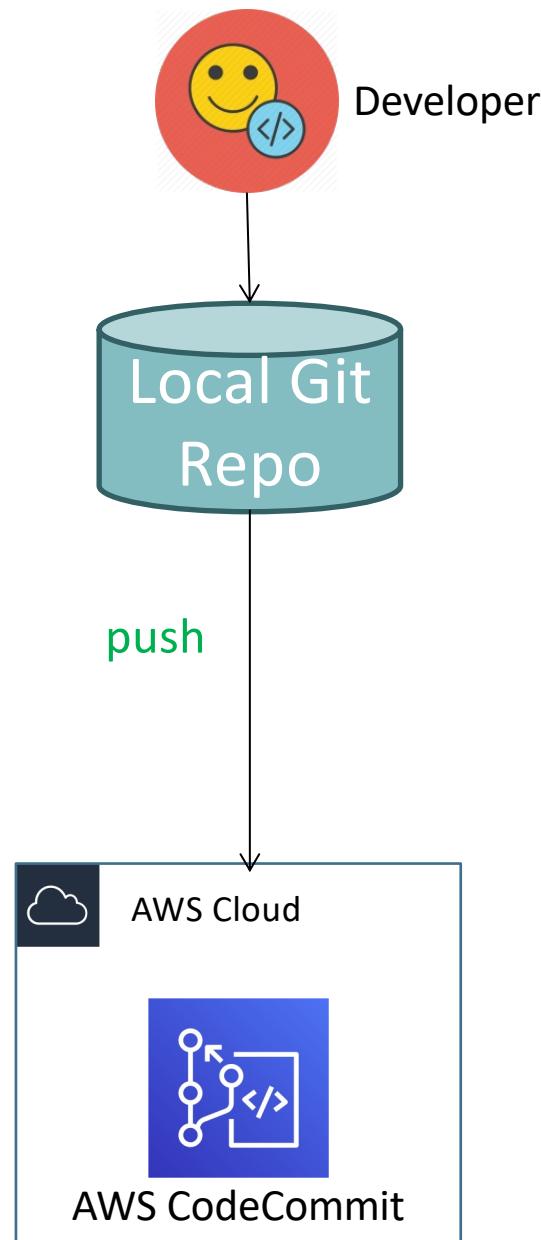
# AWS CodeCommit - Introduction

- Version Control Service hosted by AWS
- We can privately store and manage documents, source code, and binary files
- Secure & highly scalable
- Supports standard functionality of Git (CodeCommit supports Git versions 1.7.9 and later.)
- Uses a static user name and password in addition to standard SSH..

# CodeCommit – Integration with AWS Services



# CodeCommit - Steps



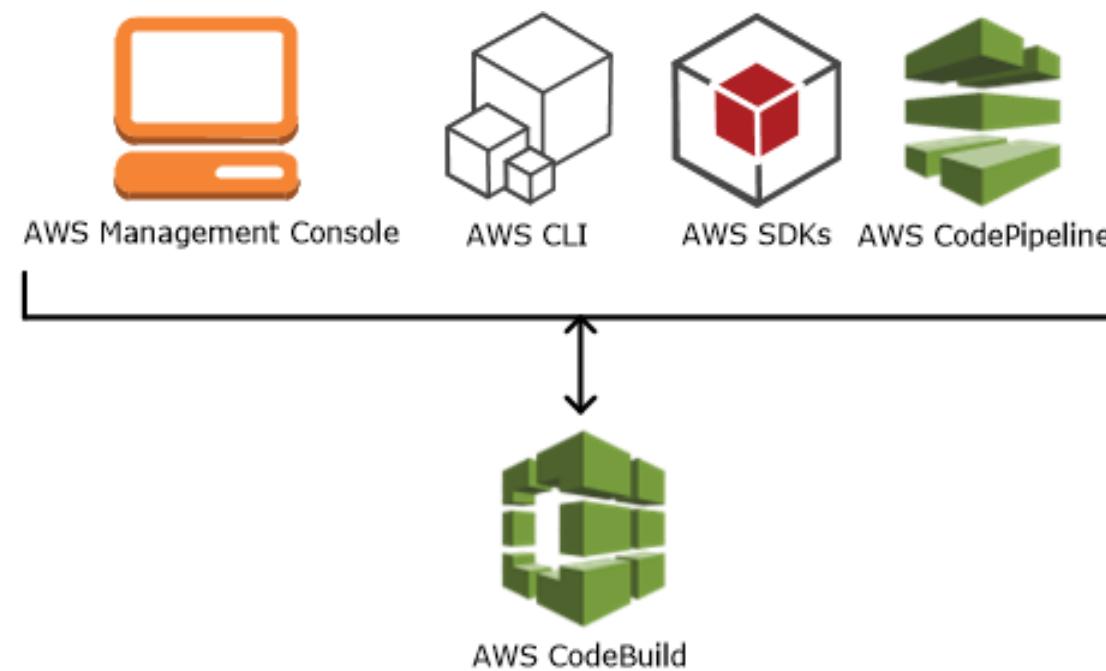
# AWS CodeBuild



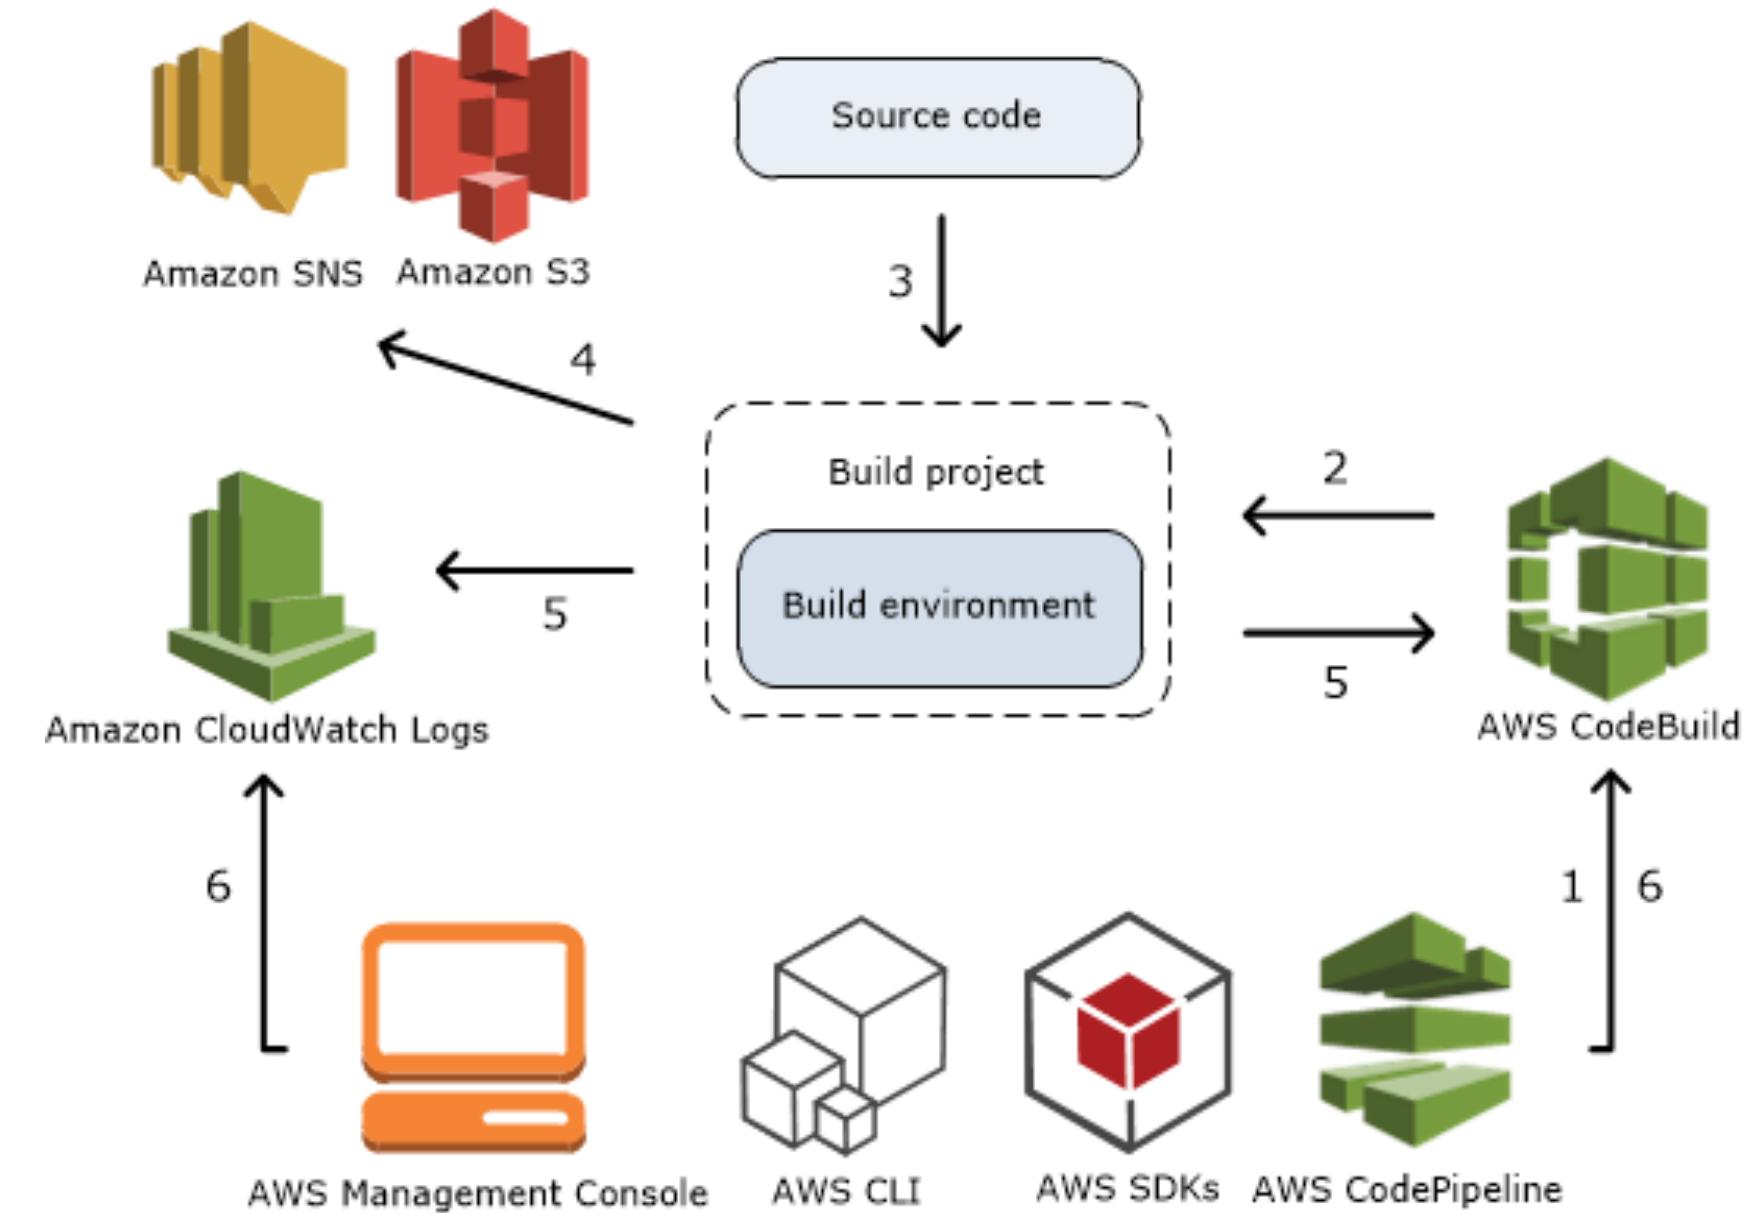
# CodeBuild - Introduction

- CodeBuild is a **fully managed** build service in the cloud.
- Compiles our **source code**, runs **unit tests**, and produces **artifacts** that are ready to deploy.
- Eliminates the need to provision, manage, and scale **our own build servers**.
- It provides **prepackaged build environments** for the most popular programming languages and build tools such as Apache Maven, Gradle, and many more.
- We can also customize build environments in CodeBuild to use **our own build tools**.
- **Scales automatically** to meet peak build requests.

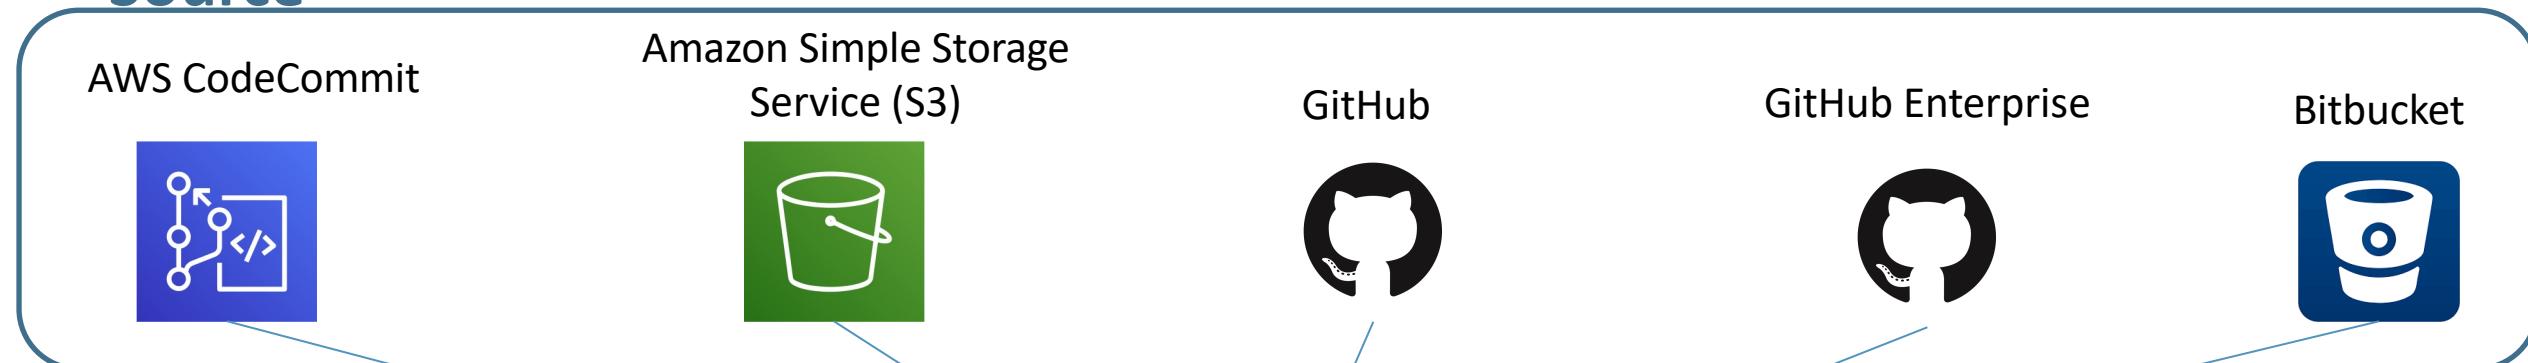
## How to run CodeBuild?



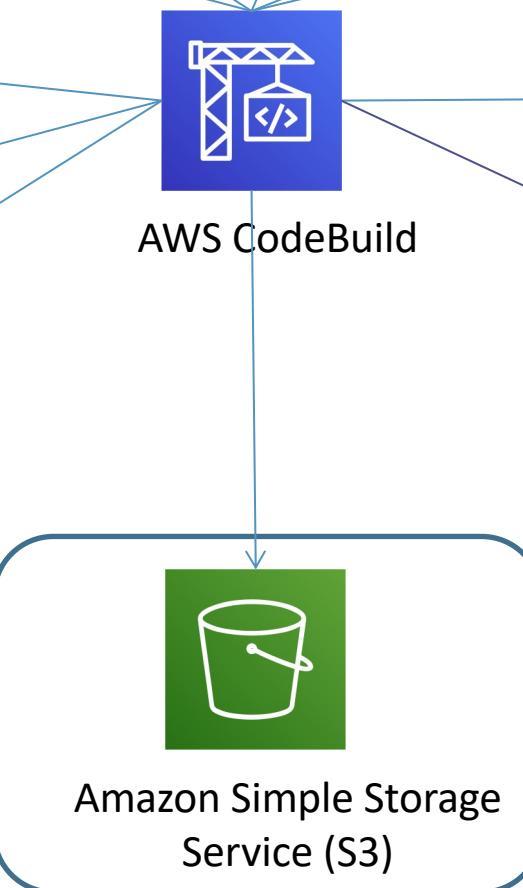
## How CodeBuild works?



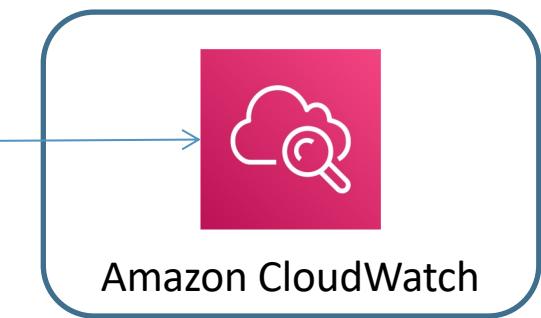
# Source



# Build Environment



# Build Artifacts



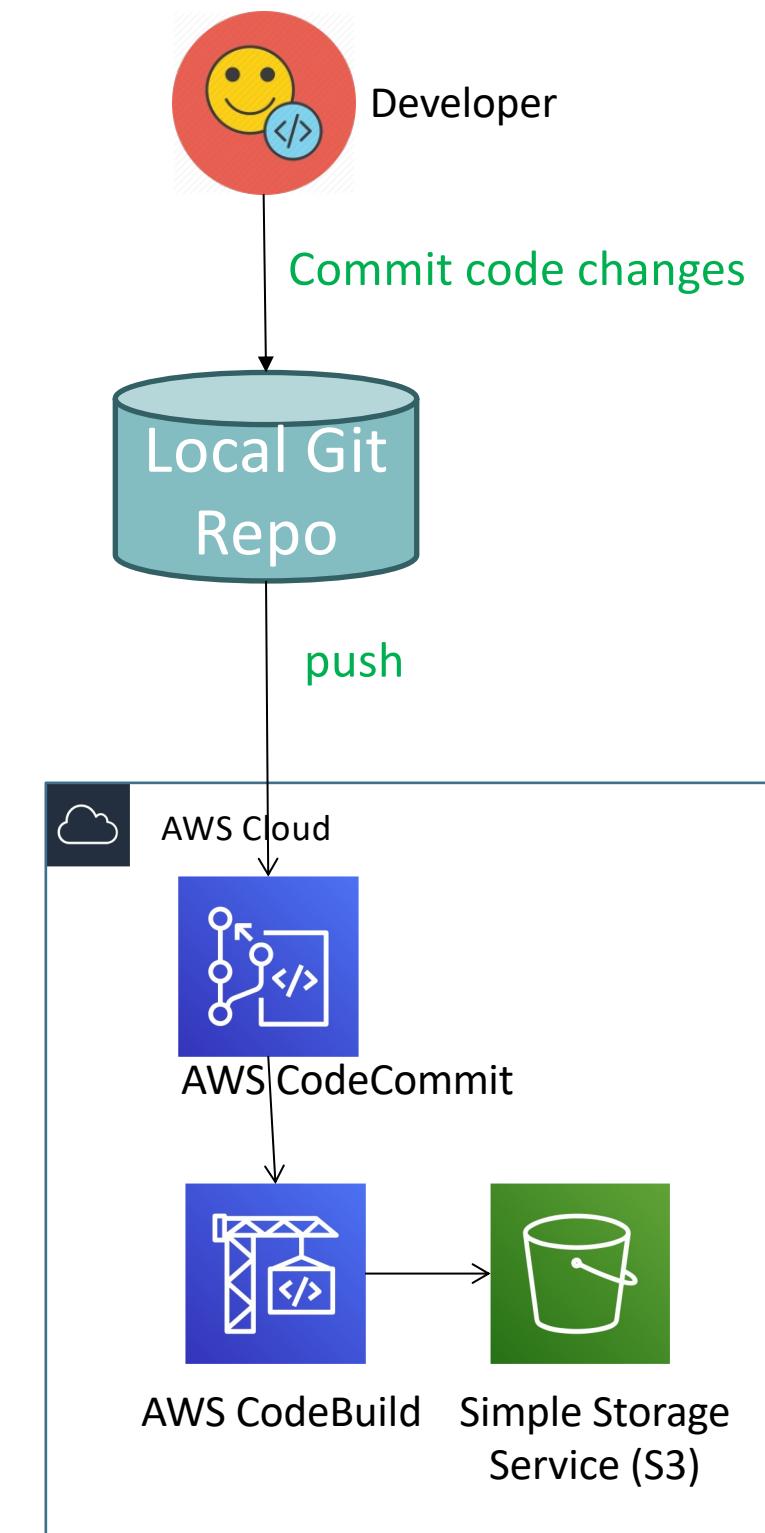
# Build Logs



# Build Notifications

# AWS CodeBuild Architecture

# CodeBuild - Steps



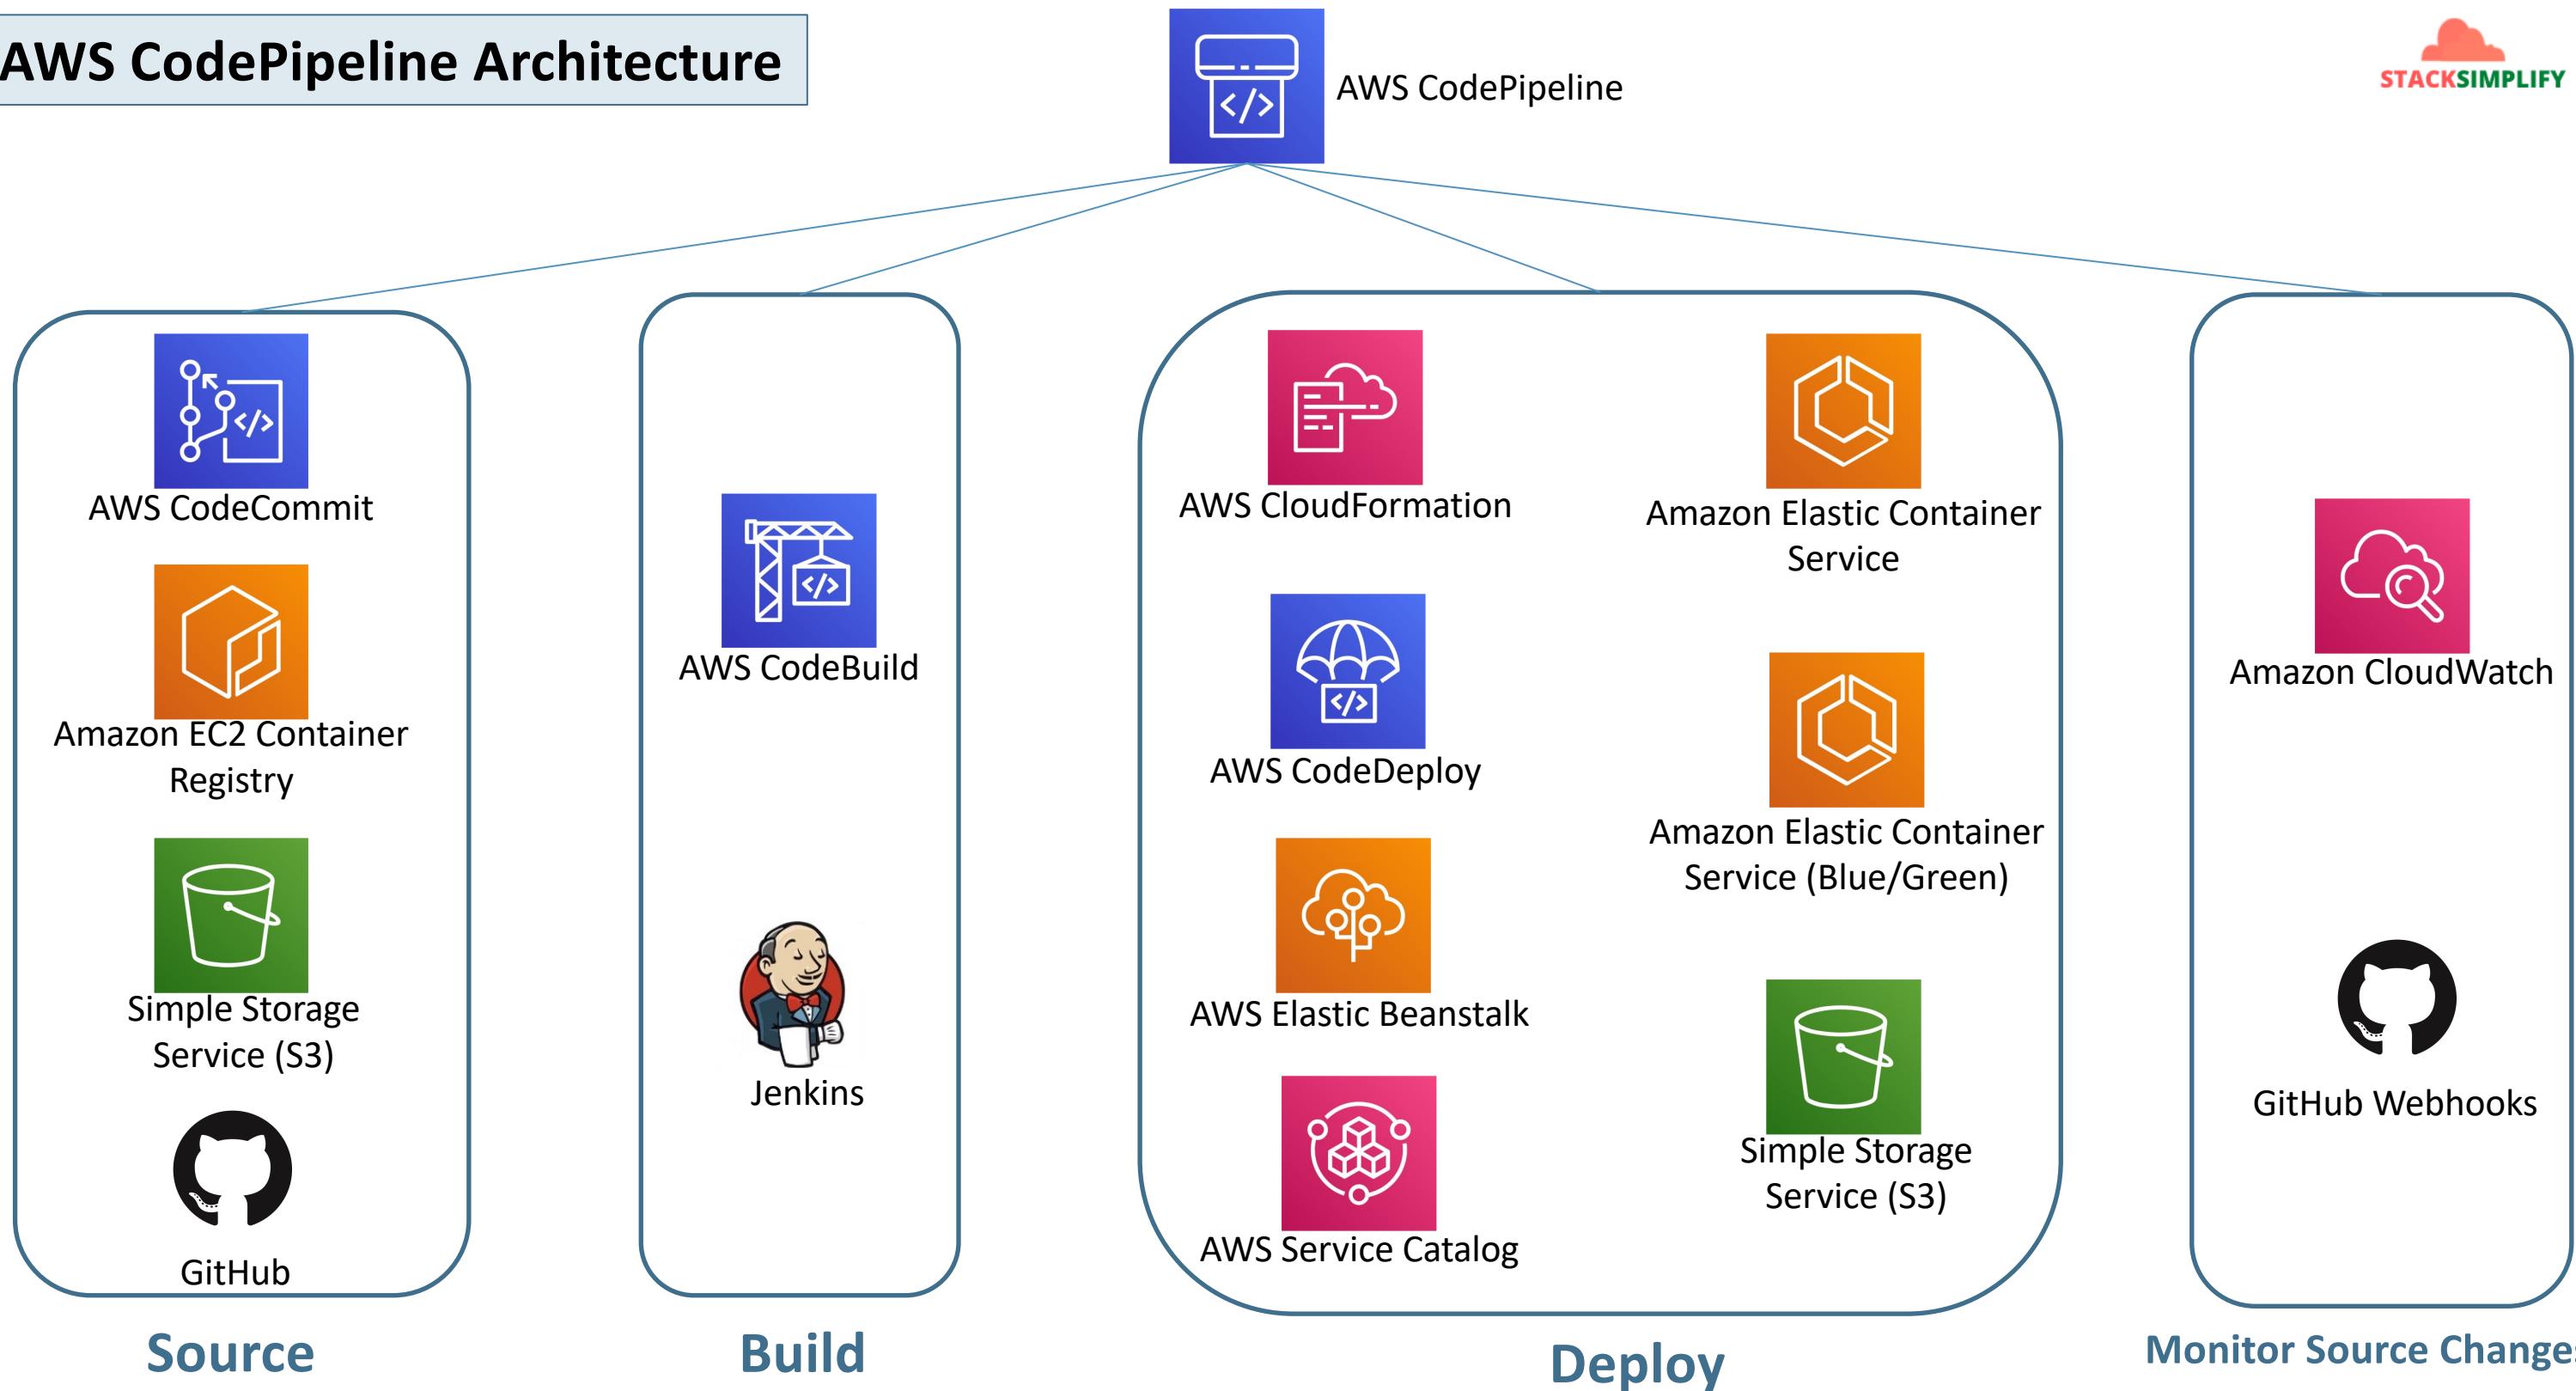
# AWS CodePipeline



# CodePipeline - Introduction

- AWS CodePipeline is a **continuous delivery service to model, visualize, and automate** the steps required to release your software.
- Benefits
  - We can **automate** our release processes.
  - We can establish a **consistent** release process.
  - We can **speed** up delivery while improving quality.
  - Supports **external tools** integration for source, build and deploy.
  - View **progress** at a glance
  - View pipeline **history details**.

# AWS CodePipeline Architecture



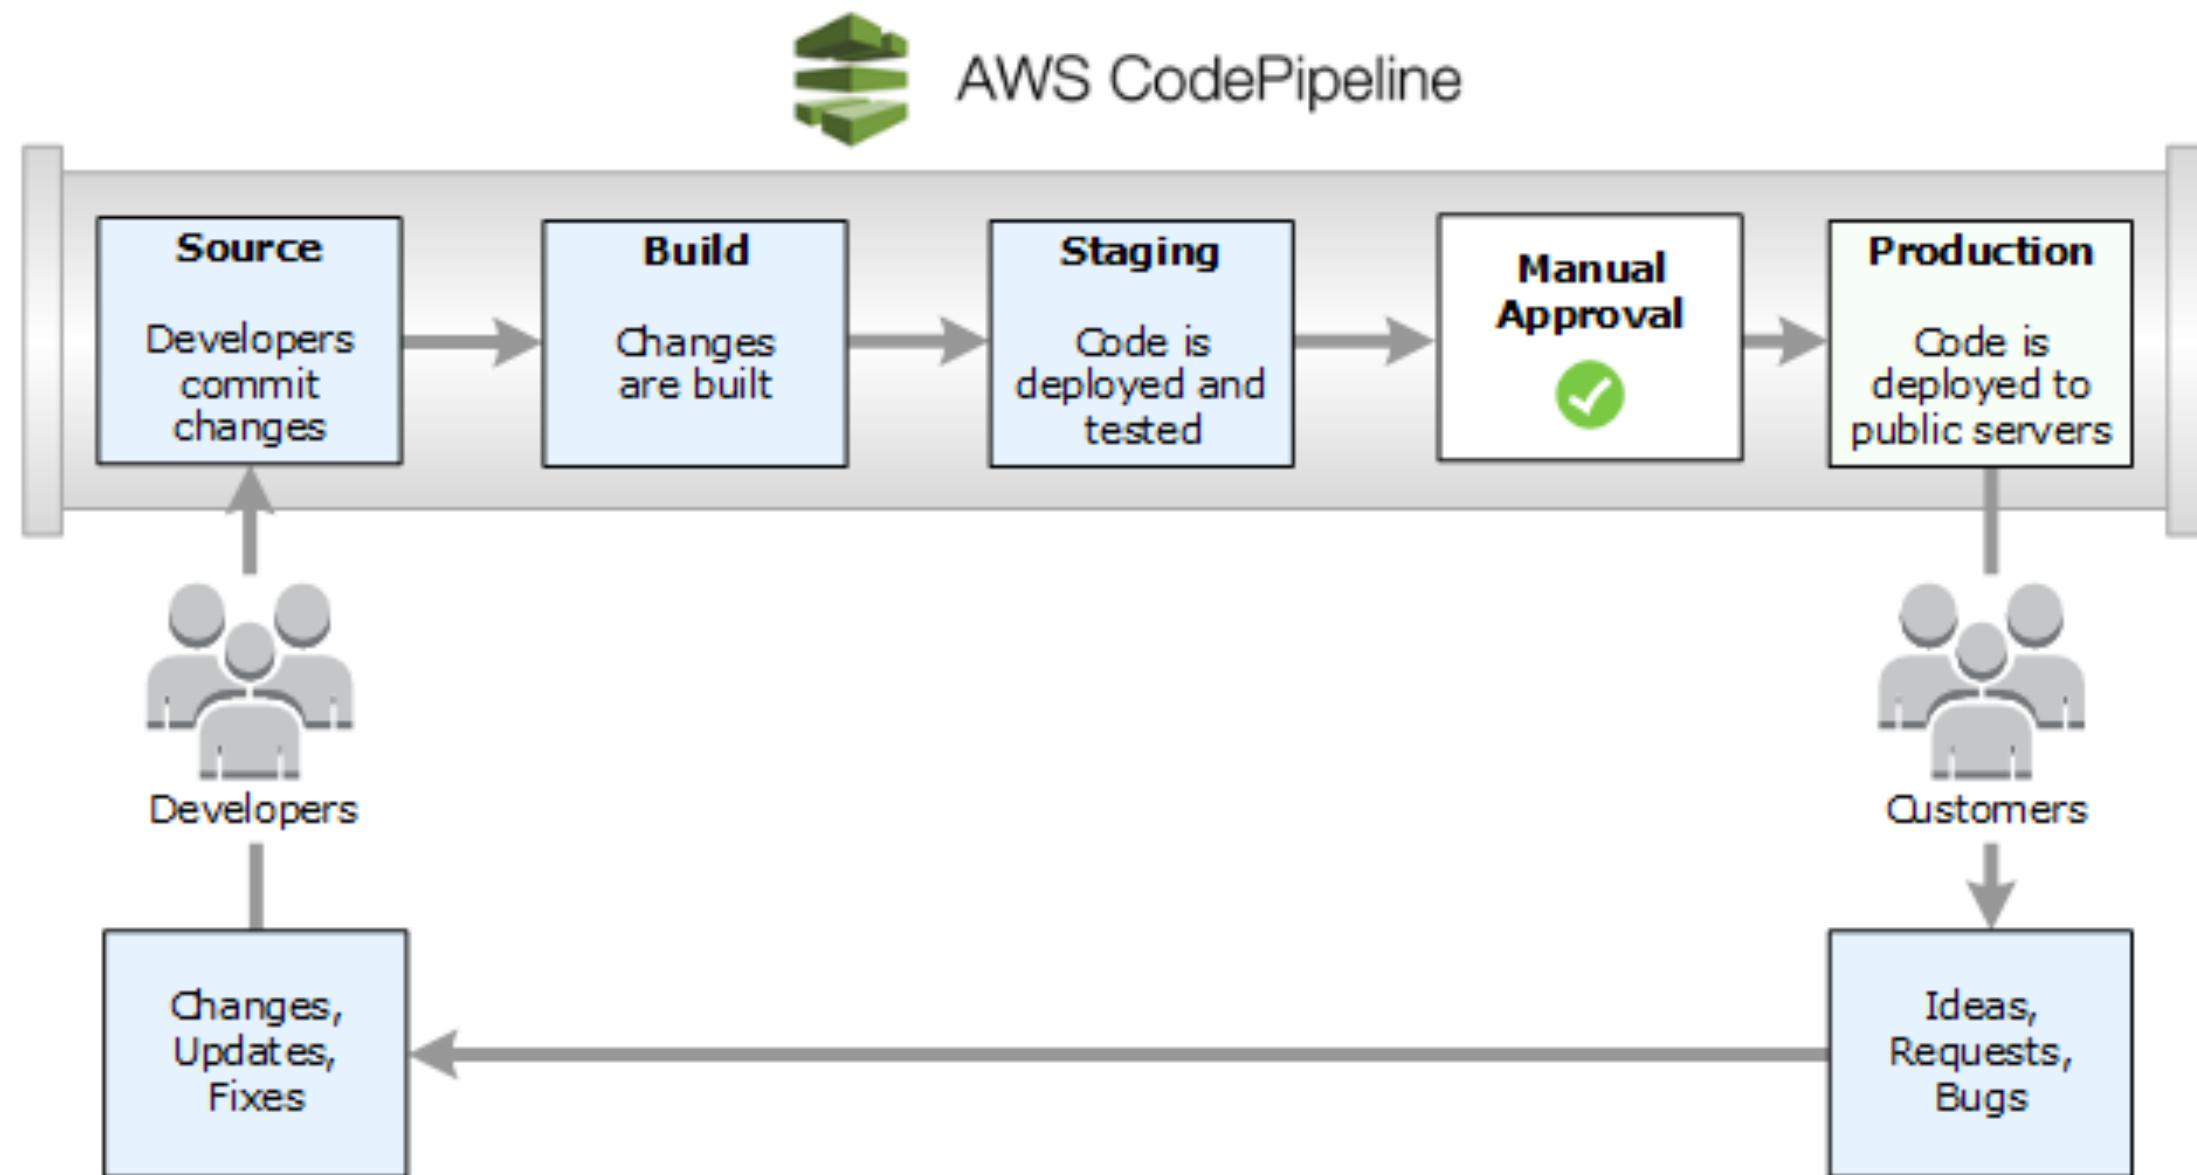
Source

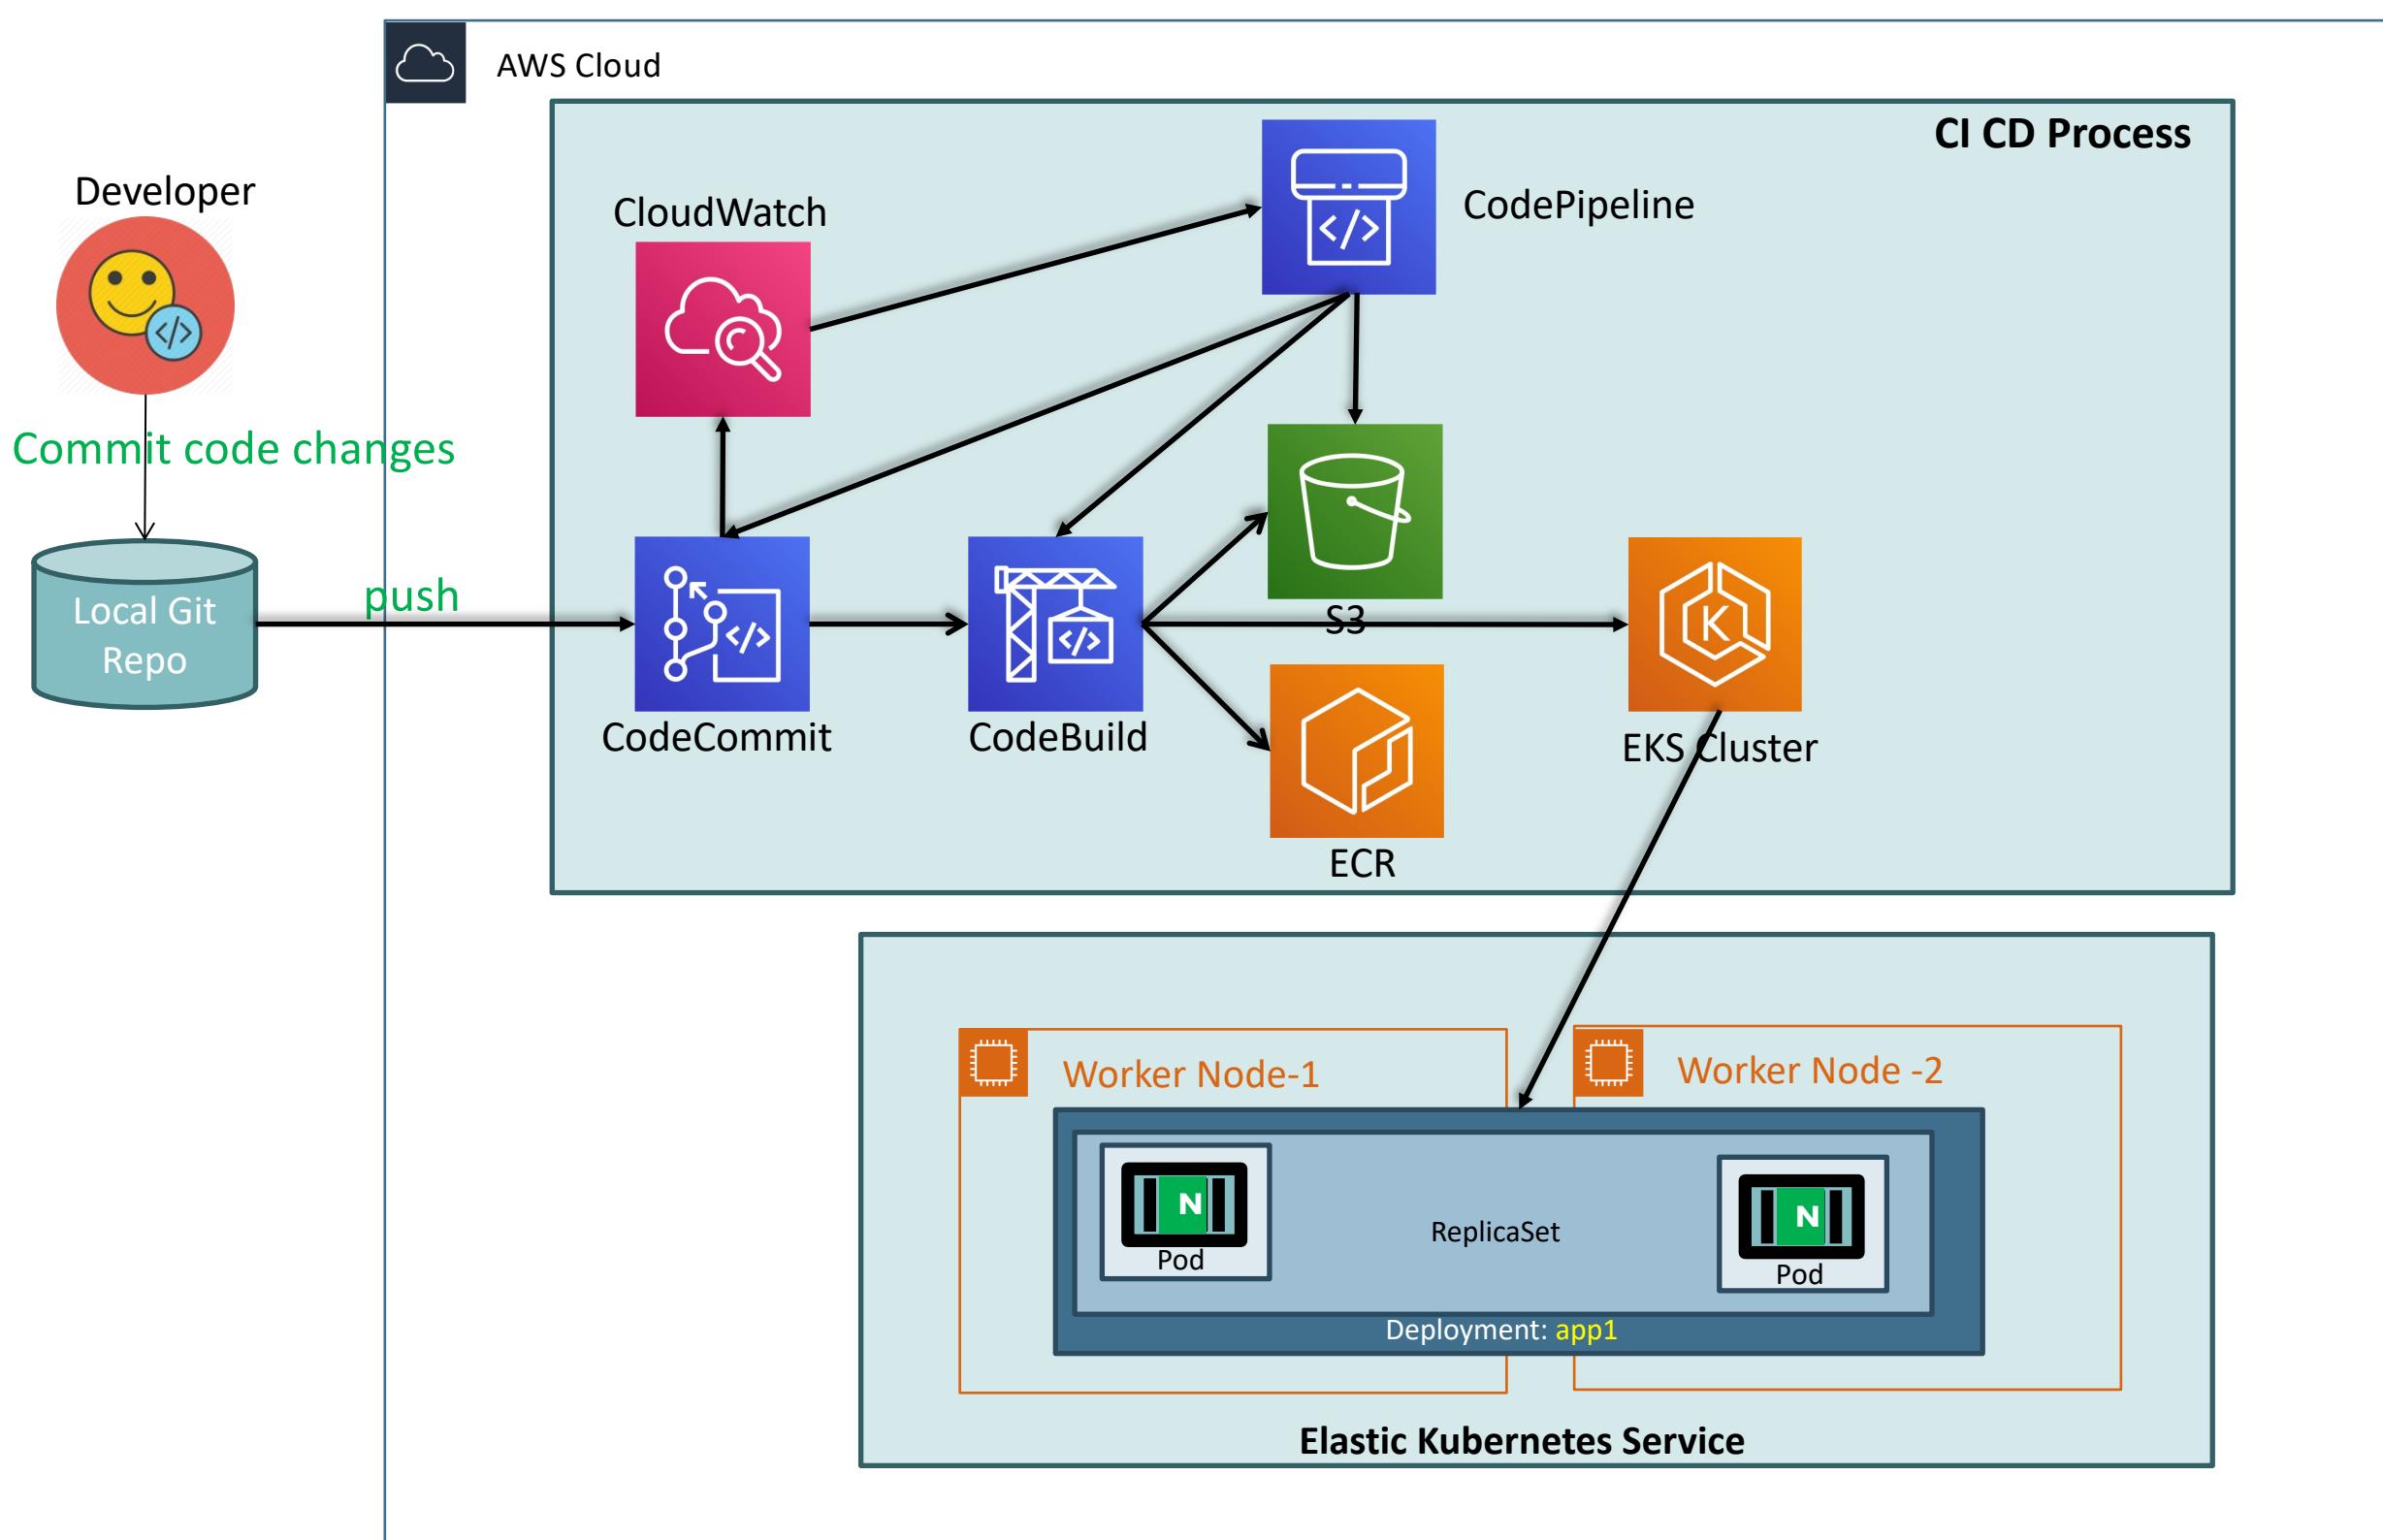
Build

Deploy

Monitor Source Changes

# Continuous Delivery



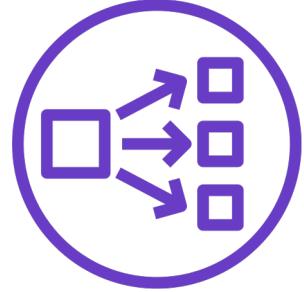




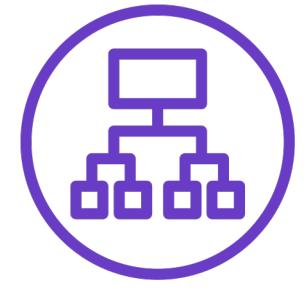
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



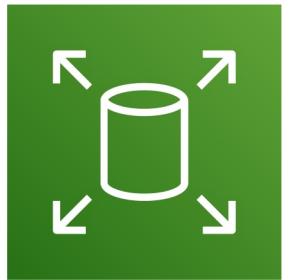
Application  
Load Balancer



Certificate  
Manager



Route53



Elastic Block  
Store



Amazon RDS



# AWS EKS

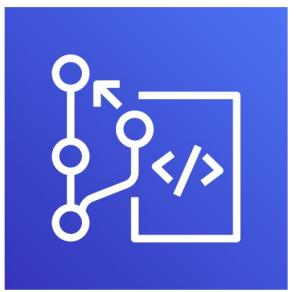
## What are Microservices?



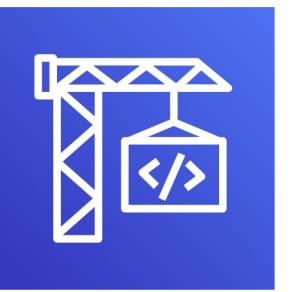
Fargate  
Profiles



Elastic Container  
Registry



Code  
Commit



Code  
Build



Code  
Pipeline



Simple Email  
Service

# What are Microservices?

- **Microservices** - also known as the **microservice architecture** - is an architectural style that structures an application as a **collection of services** that are
  - Highly maintainable and testable
  - Loosely coupled
  - Independently deployable
  - Organized around business capabilities
  - Owned by a small team

# Microservices - Benefits

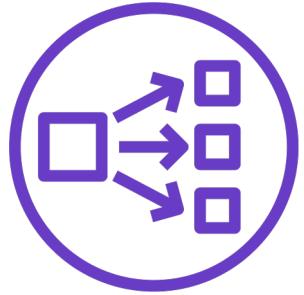
- **Developer independence:** Small teams work in parallel and can iterate faster than large teams.
- **Isolation and resilience:** If a component dies, you spin up another while the rest of the application continues to function.
- **Scalability:** Smaller components take up fewer resources and can be scaled to meet increasing demand of that component only.
- **Lifecycle automation:** Individual components are easier to fit into continuous delivery pipelines and complex deployment scenarios not possible with monoliths.
- **Relationship to the business:** Microservice architectures are split along business domain boundaries, increasing independence and understanding across the organization.



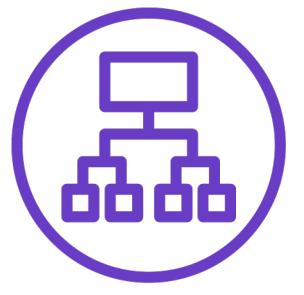
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



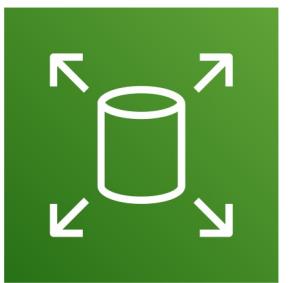
Application  
Load Balancer



Certificate  
Manager



Route53



Elastic Block  
Store



Amazon RDS



# AWS EKS

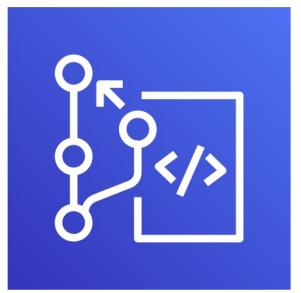
## Microservices Deployment



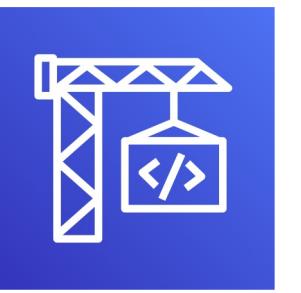
Fargate



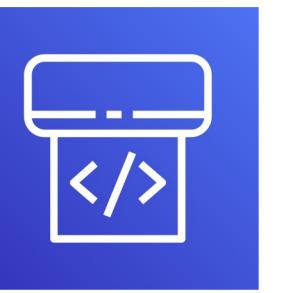
Elastic Container  
Registry



Code  
Commit



Code  
Build

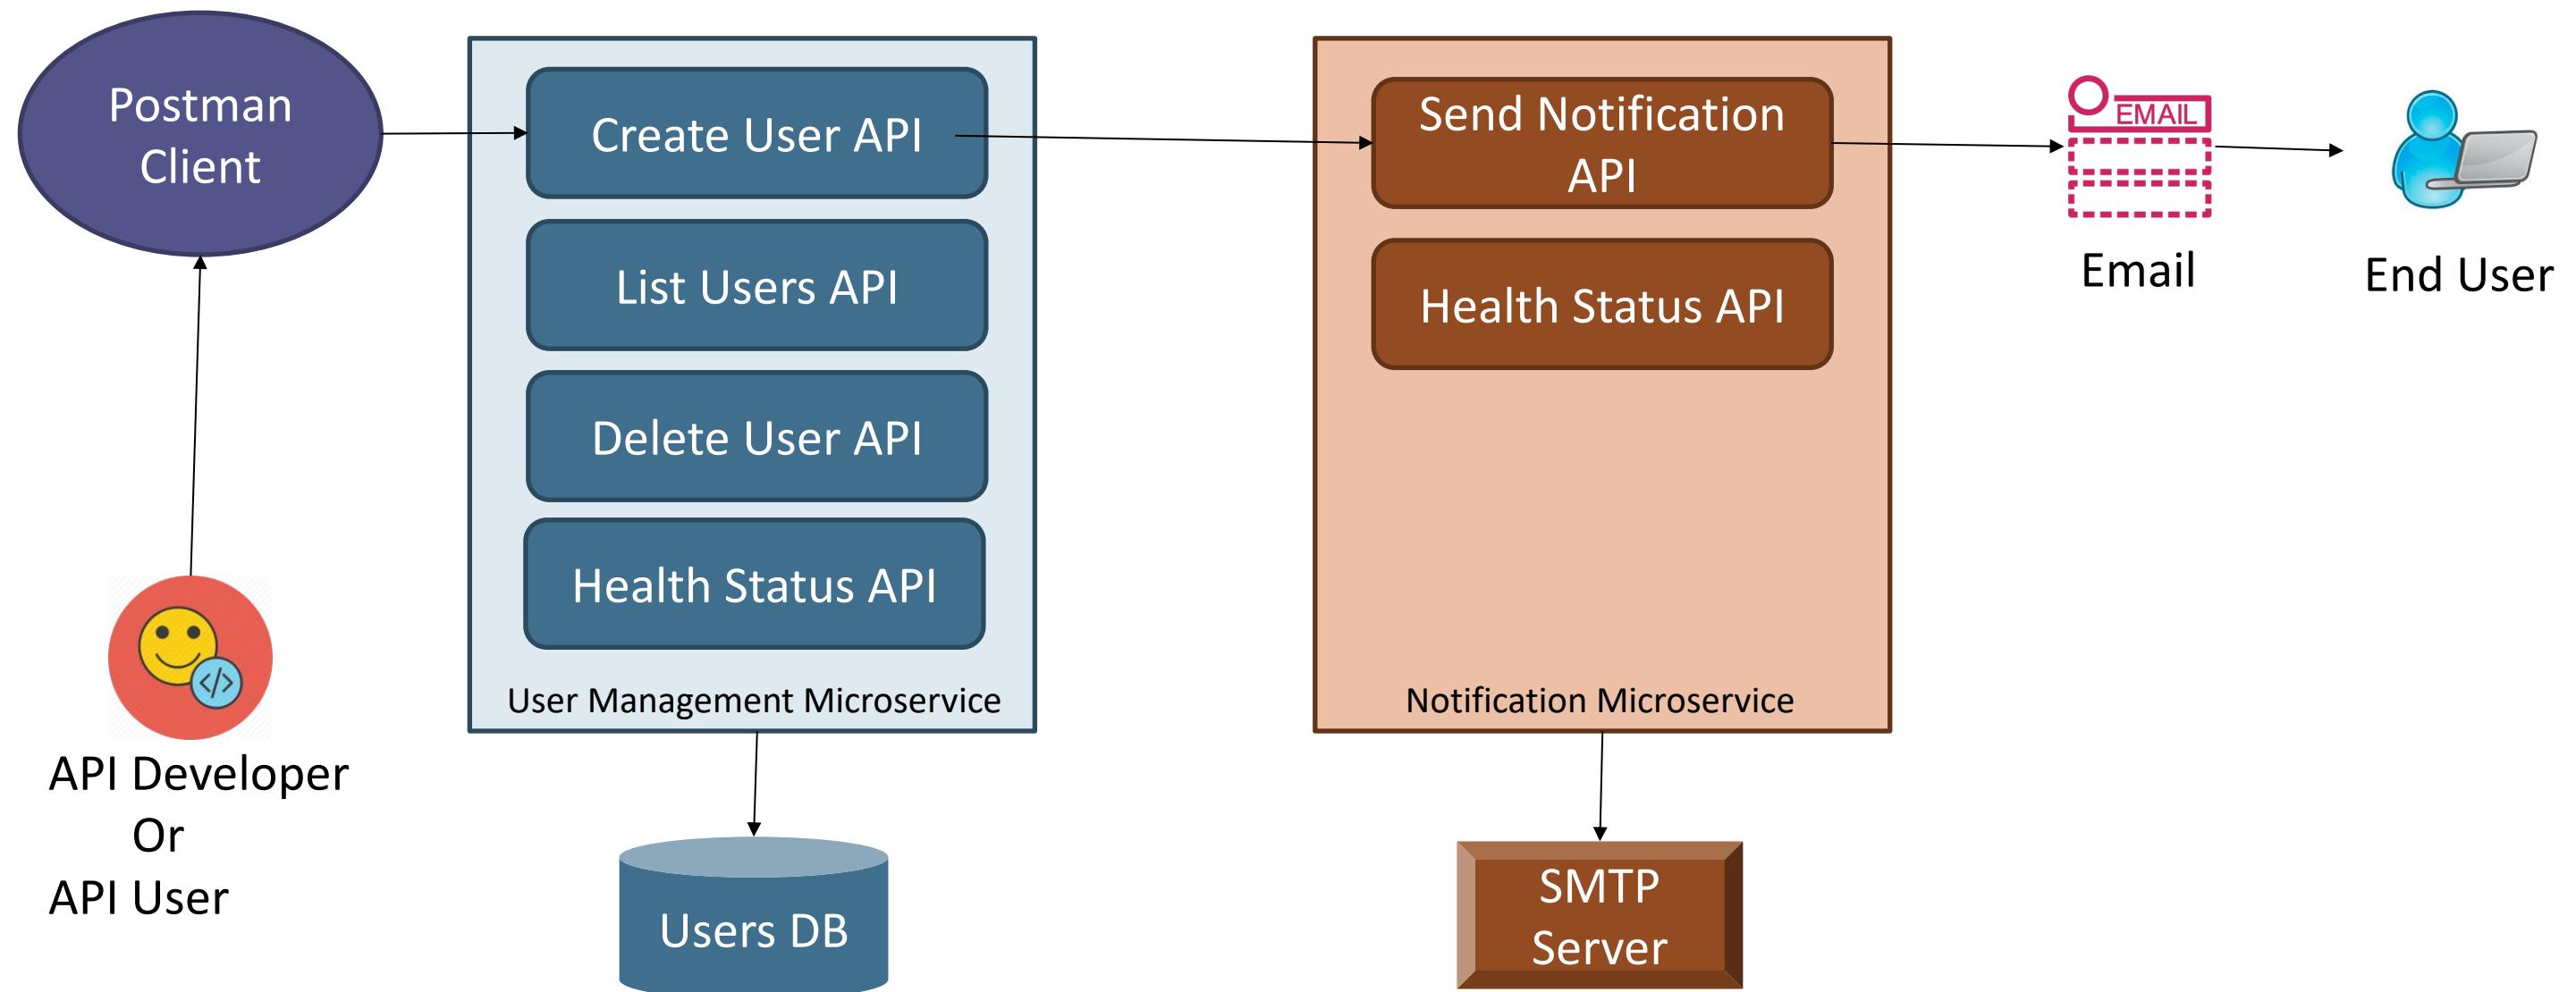


Code  
Pipeline

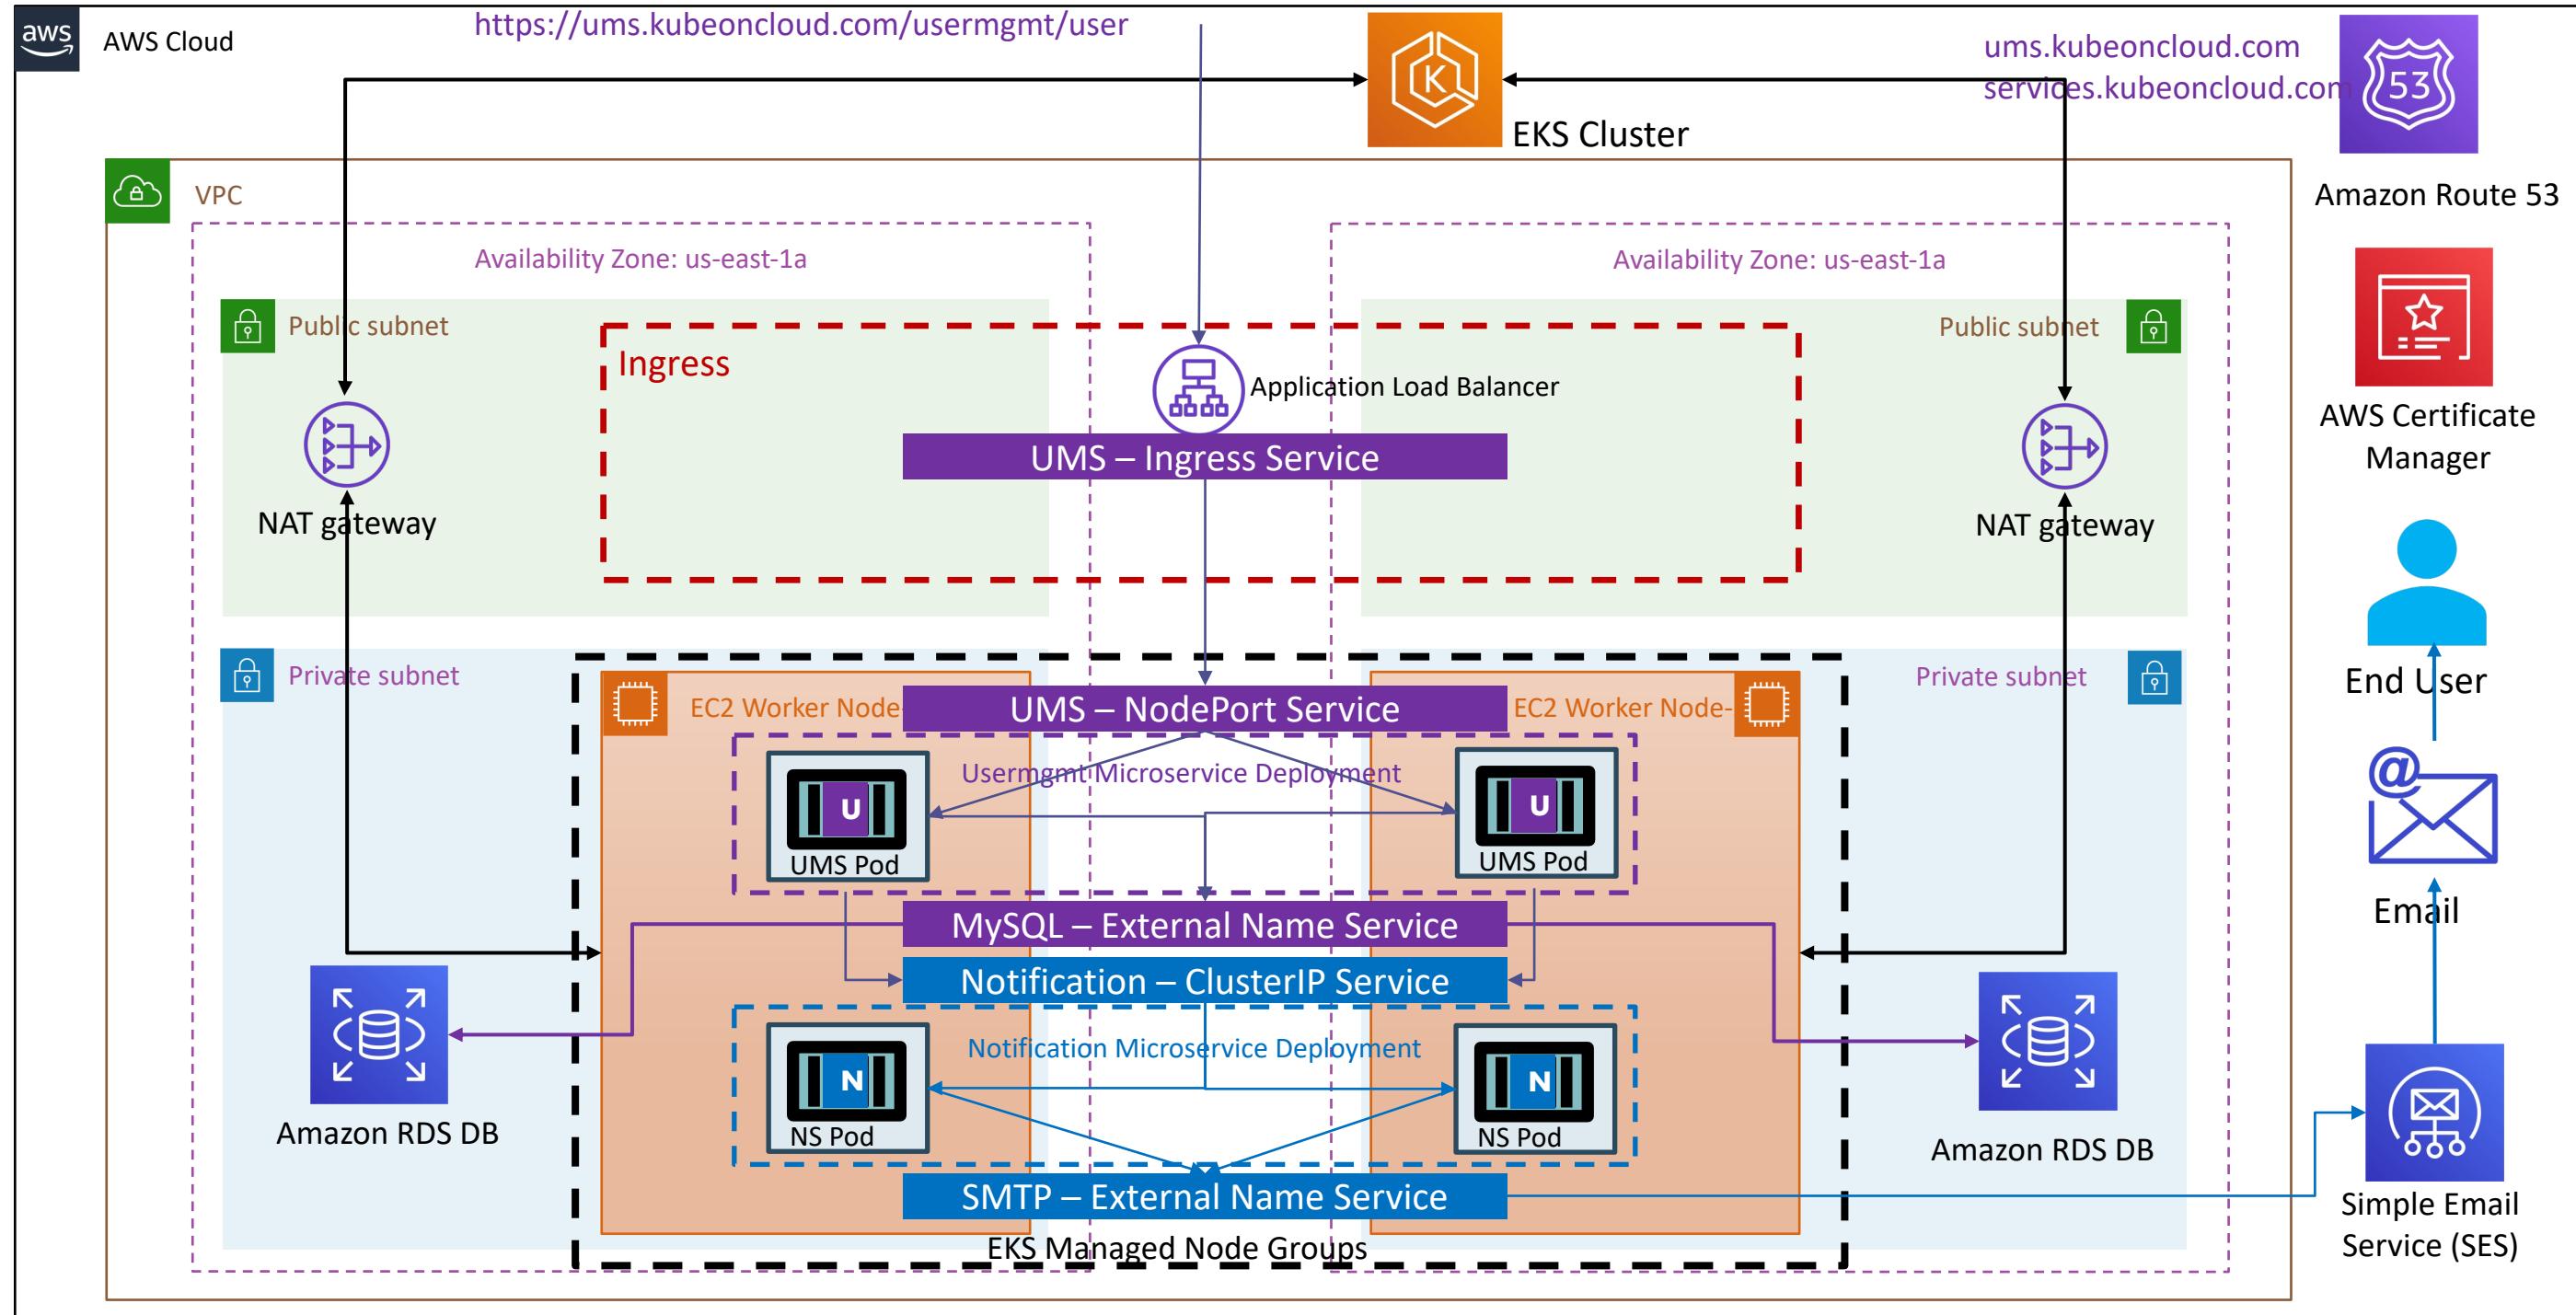
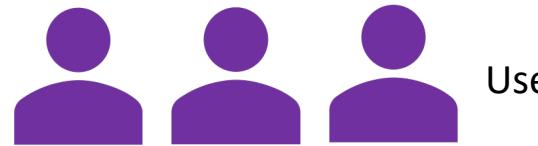


Simple Email  
Service

# Microservices



# Microservices Deployment on AWS EKS

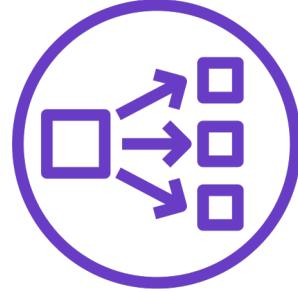




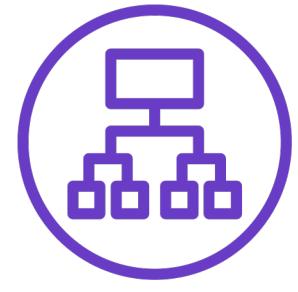
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



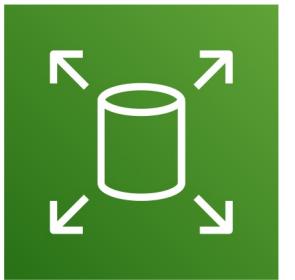
Application  
Load Balancer



Certificate  
Manager



Route53



Elastic Block  
Store



Amazon RDS



# AWS EKS

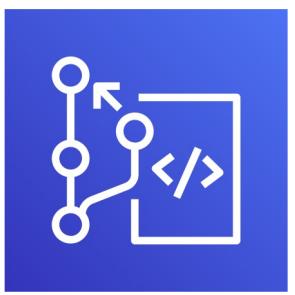
## Microservices Distributed Tracing



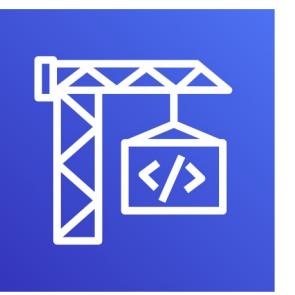
Fargate  
Profiles



Elastic Container  
Registry



Code  
Commit



Code  
Build



Code  
Pipeline



Simple Email  
Service



X-Ray

# AWS X-Ray Introduction

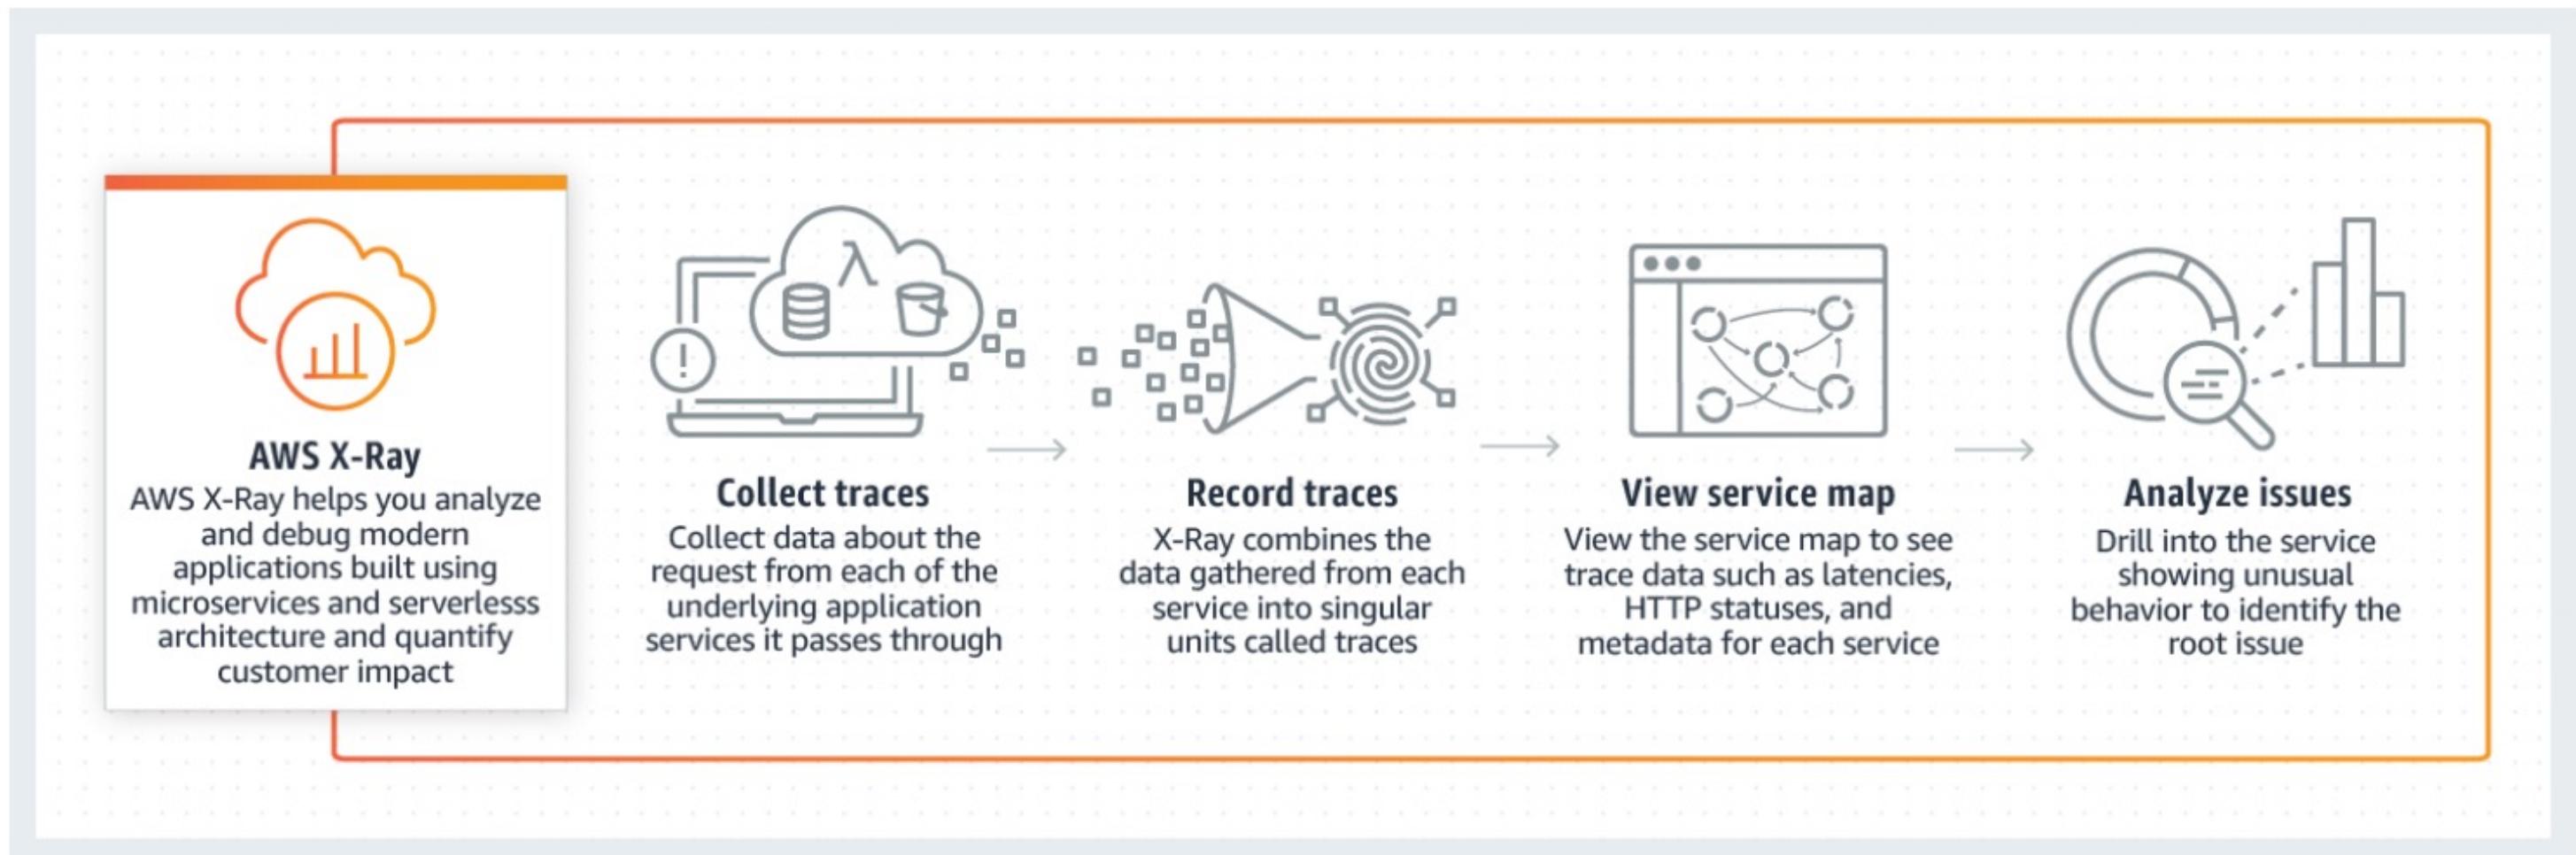
- AWS X-Ray helps analyse and debug distributed applications built using microservices architecture.
- With X-Ray, we can understand how our application and its underlying services are performing to identify and troubleshoot the root cause of performance issues and errors.
- X-Ray provides an end-to-end view of requests as they travel through our application and shows a map of our application's underlying components.
- We can also use X-Ray to analyse applications in development and in production, from simple three-tier applications to complex microservices applications consisting of thousands of services.

# AWS X-Ray - Benefits

- Review request behavior
- Discover application issues
- Improve application performance
- Ready to use with AWS
- Designed for a variety of applications

# AWS X-Ray – How it works?

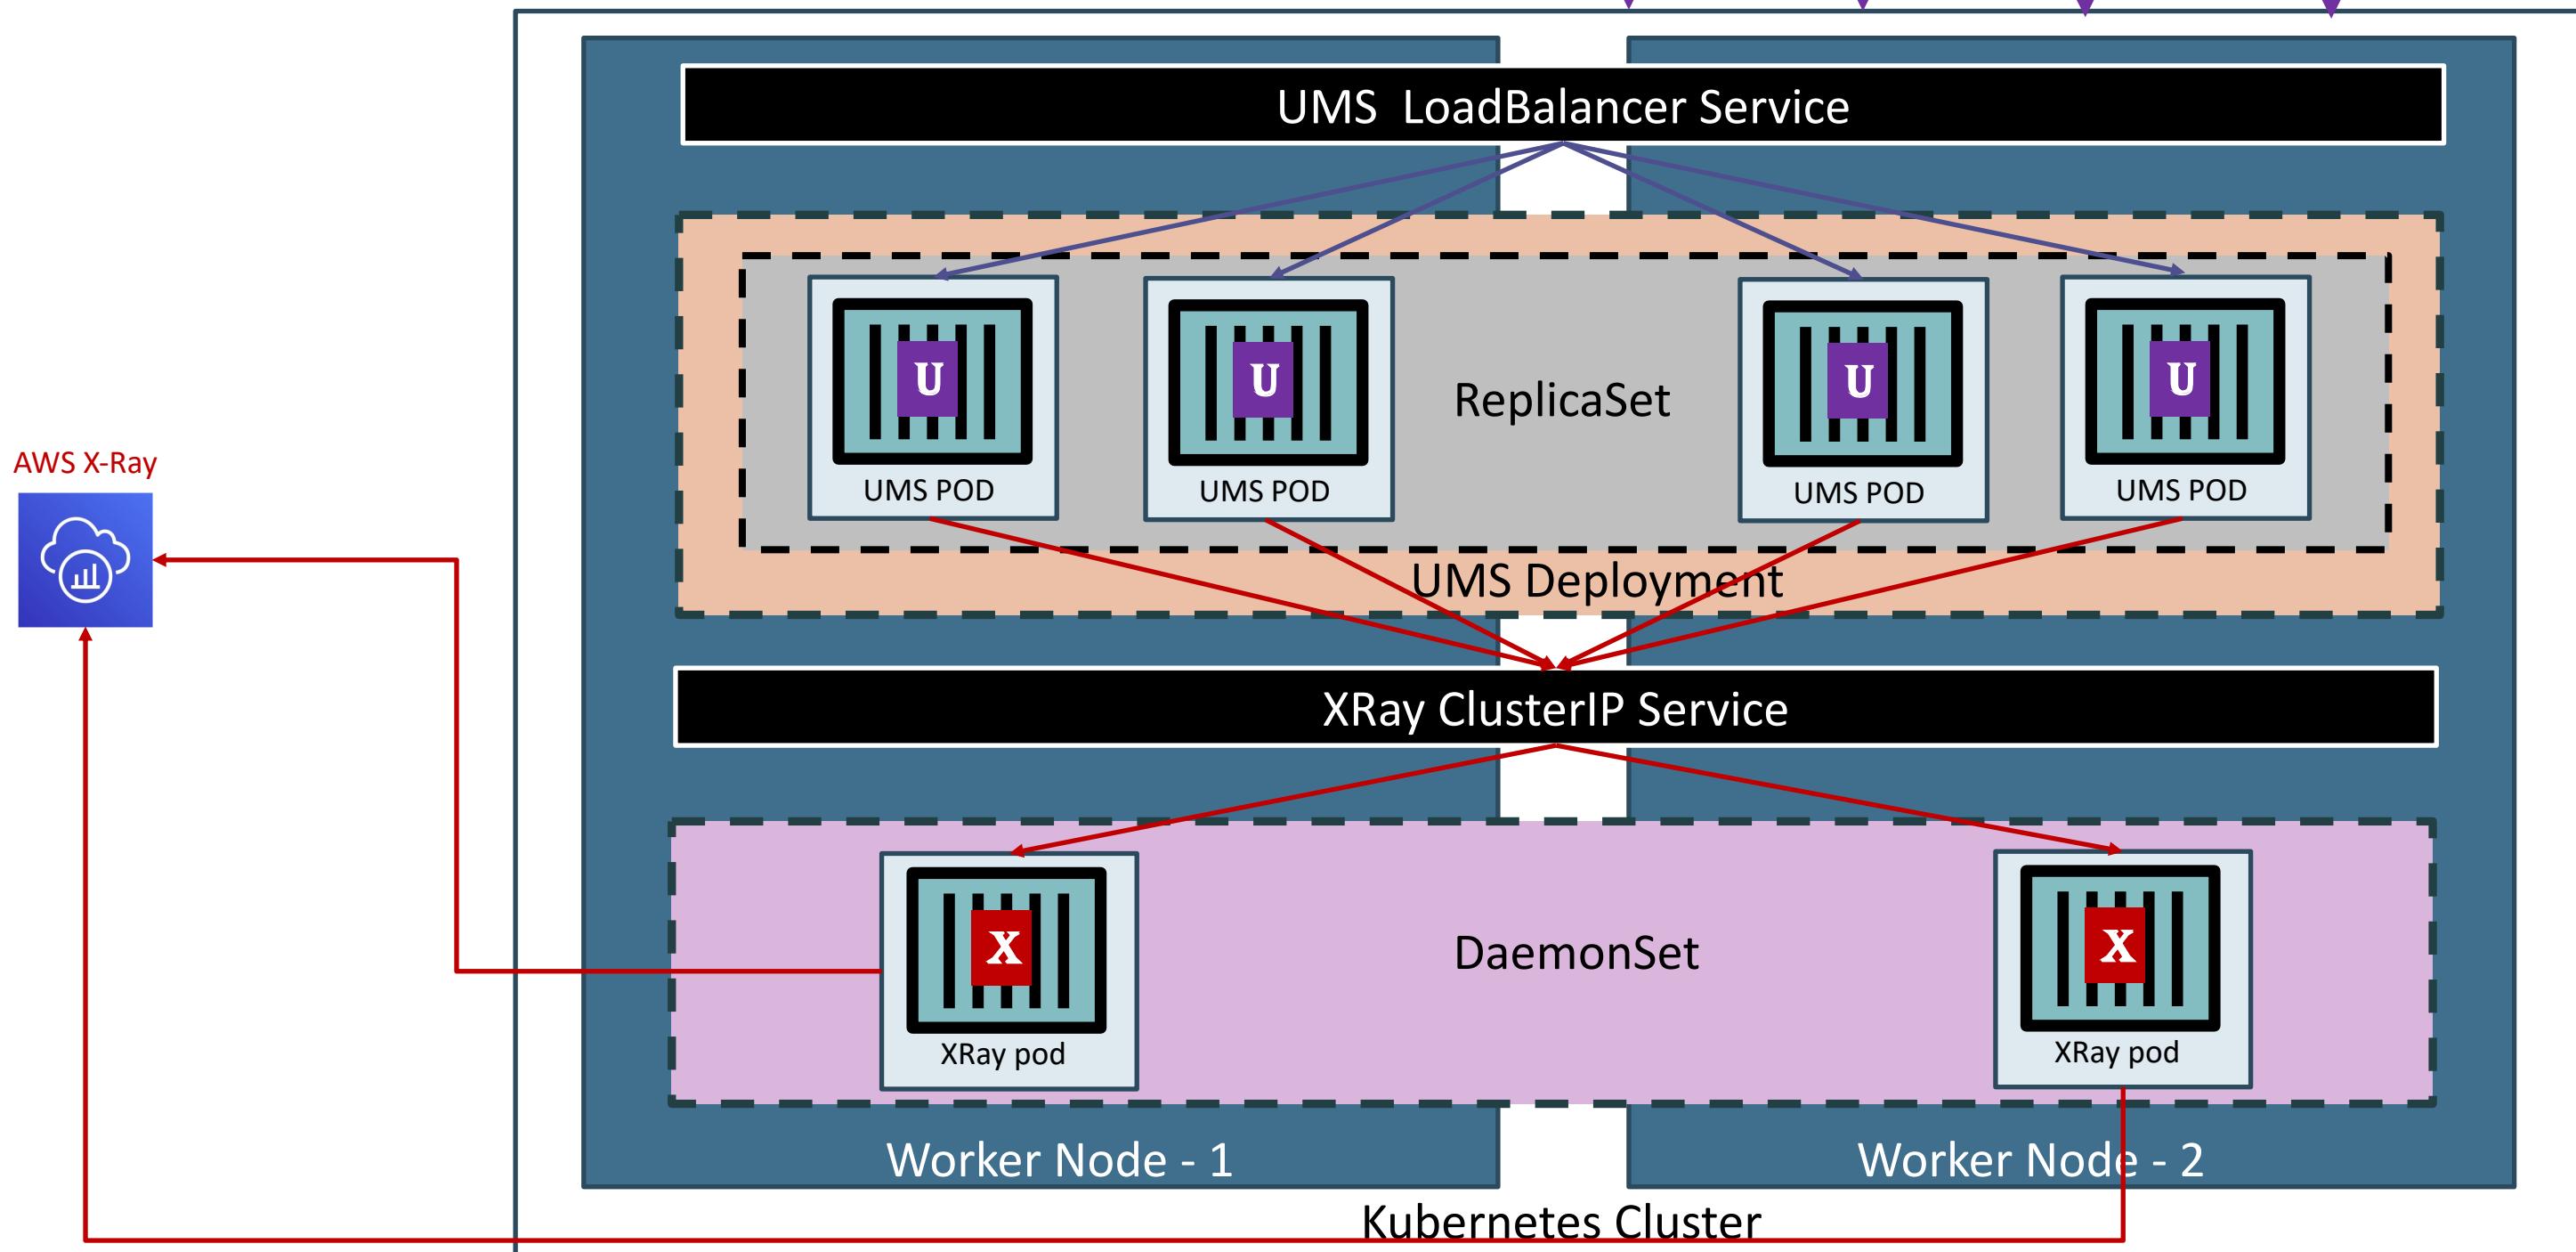
## How it works



# Kubernetes DaemonSets - Introduction

- A *DaemonSet* ensures that all (or some) Nodes run a copy of a Pod.
  - As nodes are **added** to the cluster, Pods are added to them.
  - As nodes are **removed** from the cluster, those Pods are garbage collected.
  - **Deleting** a DaemonSet will clean up the Pods it created.
- Some typical uses of a DaemonSet are:
  - running a **logs collection daemon** on every node (Example: [fluentd](#))
  - running a **node monitoring daemon** on every node (Example: [cloudwatchagent](#))
  - running an **application trace collection daemon** on every node (Example: [AWS X-Ray](#))
- In a **simple case**, one DaemonSet, covering all nodes, would be used for each type of daemon.
- A **more complex setup** might use **multiple DaemonSets for a single type of daemon**, but with different flags and/or different memory and cpu requests for different hardware types

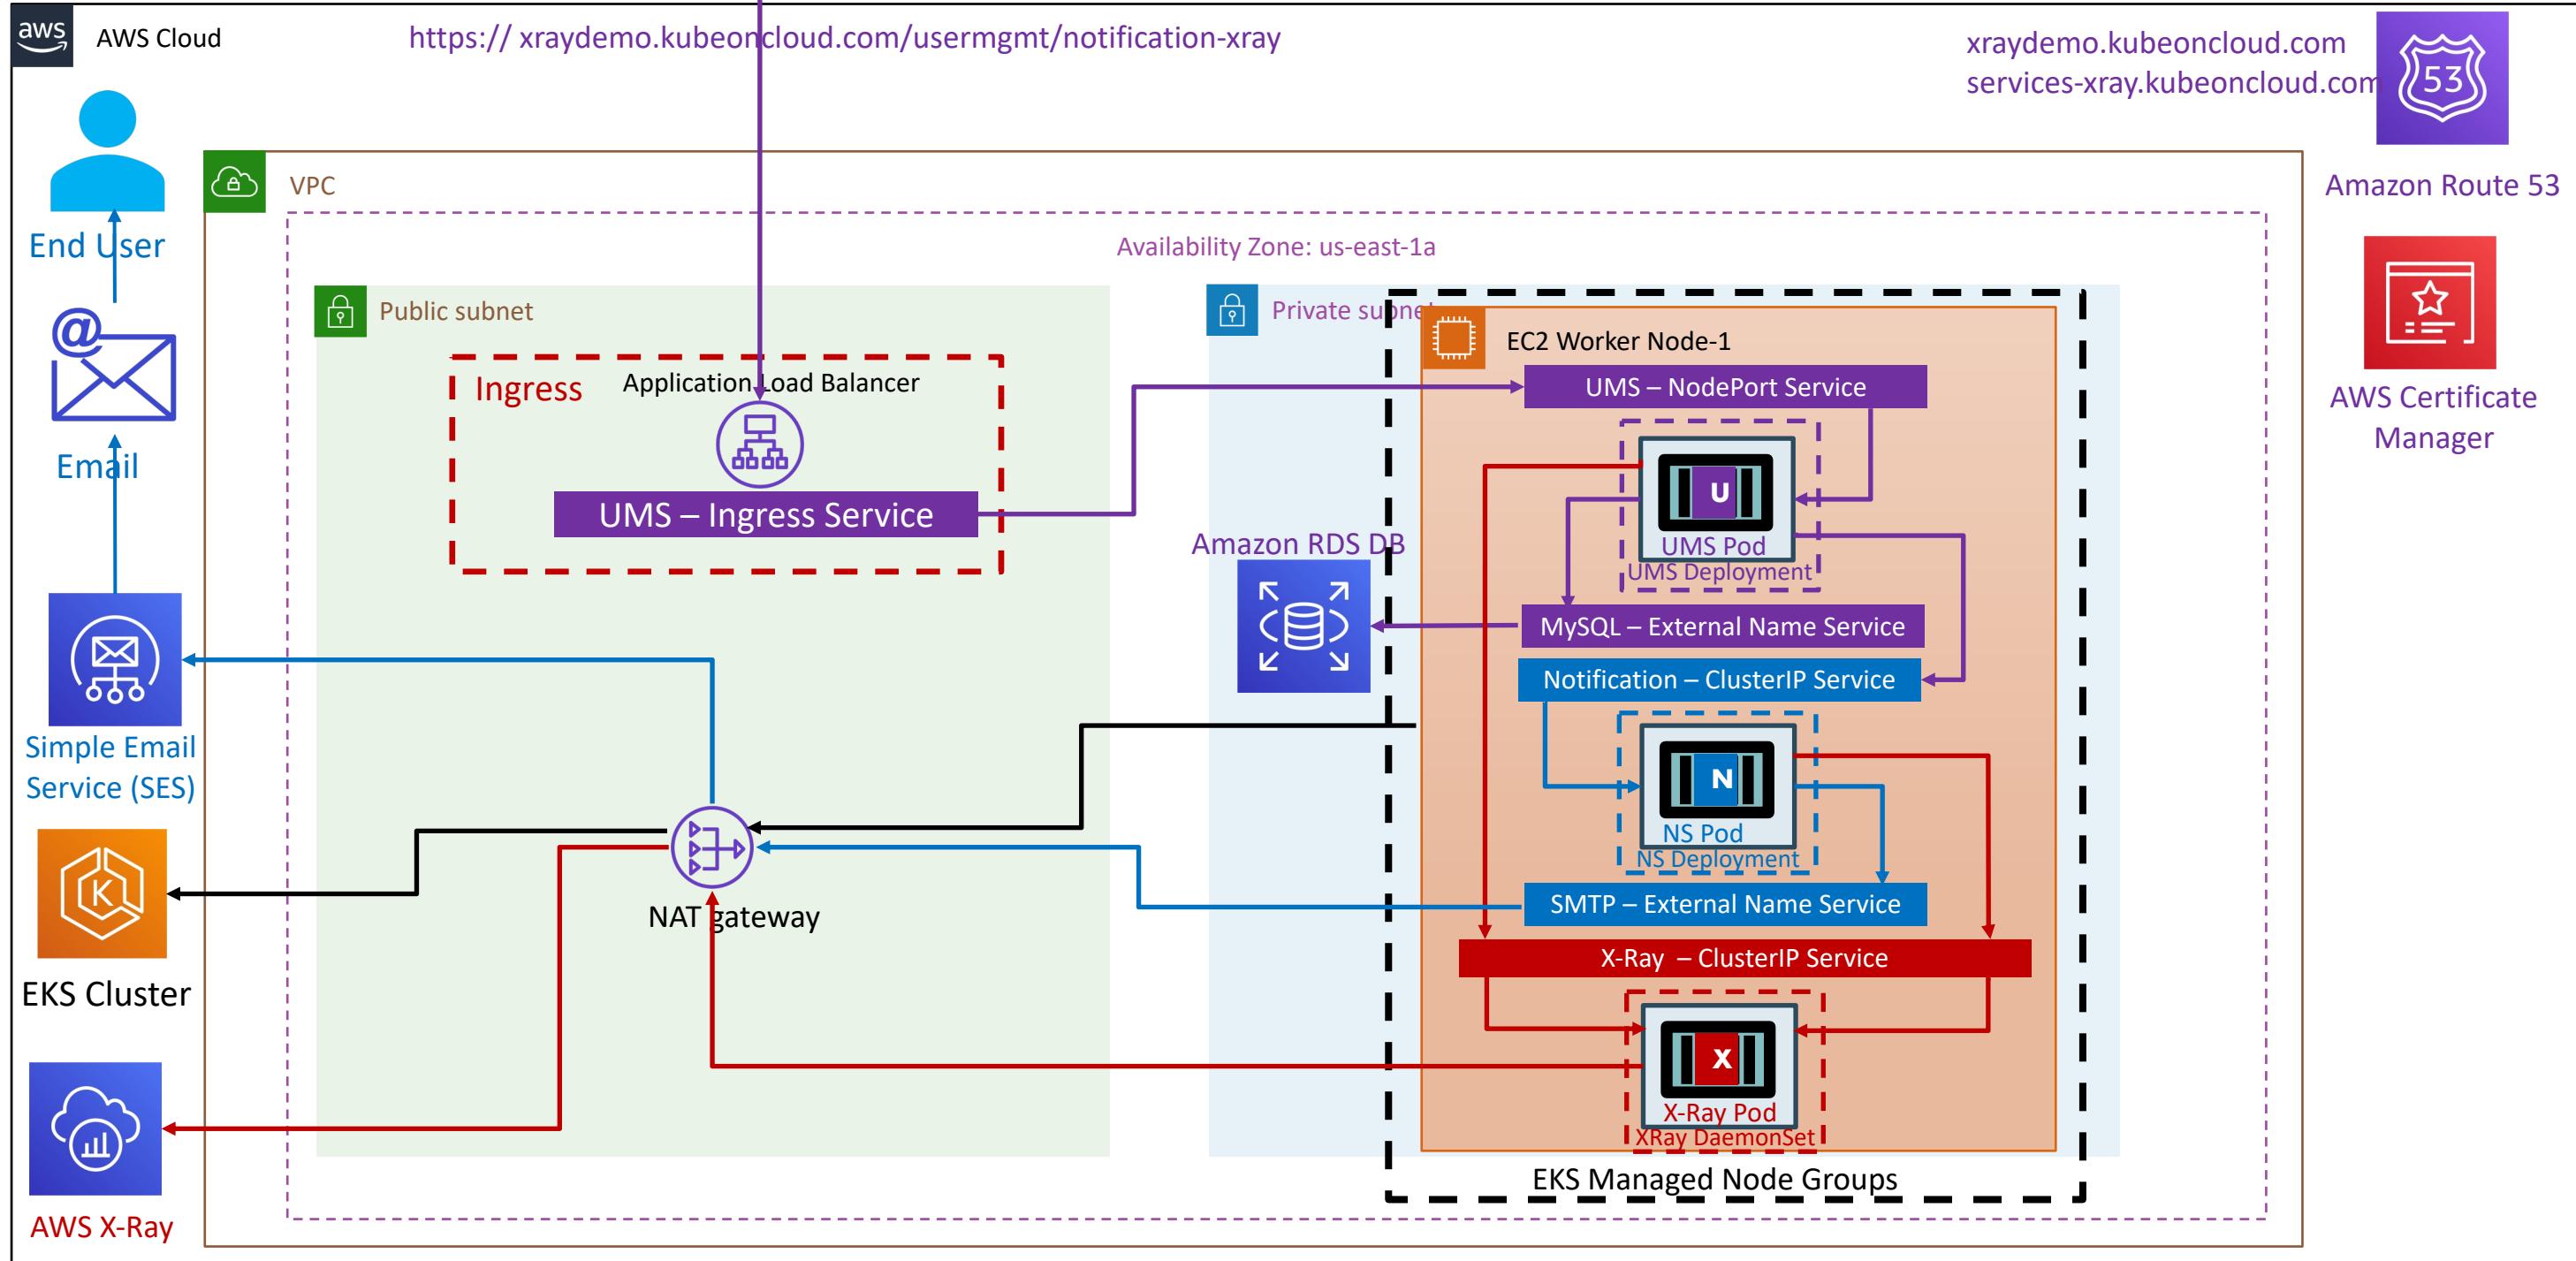
# Kubernetes – DaemonSets





# Microservices Distributed Tracing with AWS X-Ray

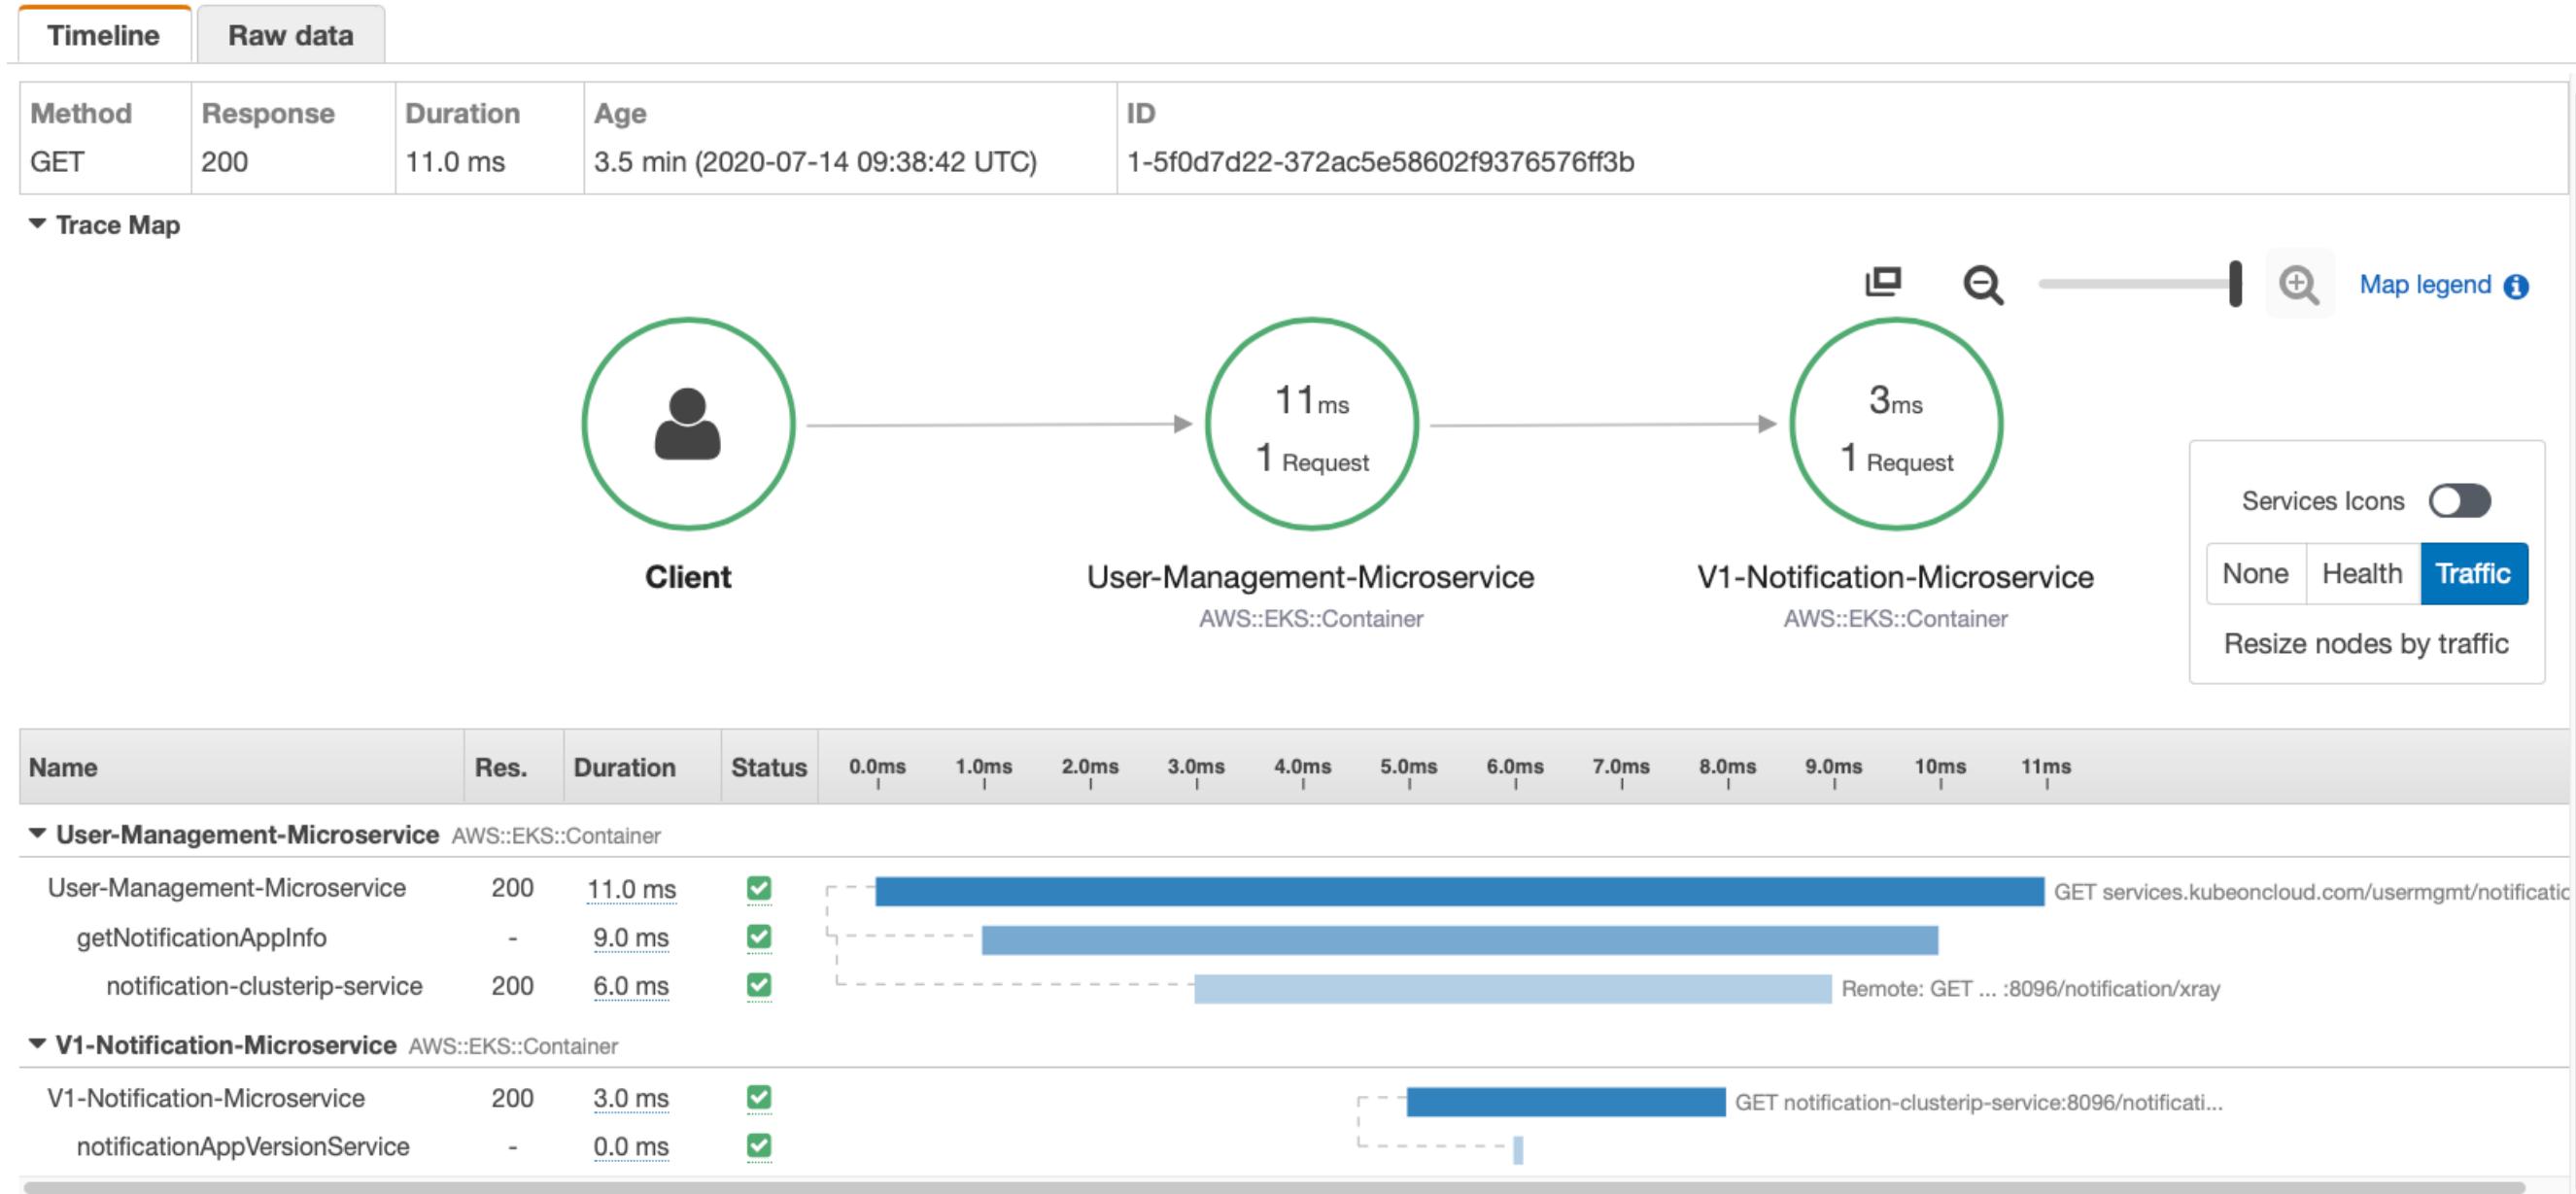
STACKSIMPLIFY



# AWS X-Ray – Service Map



# AWS X-Ray - Traces



# AWS X-Ray - Traces

## Segment - User-Management-Microservice

[Overview](#)
[Resources](#)
[Annotations](#)
[Metadata](#)
[Exceptions](#)

Segment ID

722d02f36eb066c4

Parent ID

Name

User-Management-Microservice

Origin

AWS::EKS::Container

### Time

Start time

2020-07-14 09:38:42.647 (UTC)

End time

2020-07-14 09:38:42.658 (UTC)

Duration

11.0 ms

In progress

false

### Errors & Faults

Error

false

Fault

false

### Request & Response

Request url

http://services.kubeoncloud.com/usermgmt/notification-xray

Request method

GET

Request user\_agent

PostmanRuntime/7.25.0

Request client\_ip

49.206.222.64

Request x\_forwarded\_for

true

## Segment - User-Management-Microservice

[Overview](#)
[Resources](#)
[Annotations](#)
[Metadata](#)
[Exceptions](#)

### CloudWatch Logs

```
[
  {
    "log_group": "/aws/containerinsights//application"
  }
]
```

### EC2

Availability zone

us-east-1b

Instance ID

i-0247e0ae1264716b3

Instance size

t3.medium

Ami ID

ami-01b33458af4b31b8f

### X-Ray

SDK version

2.6.1

SDK

X-Ray for Java

### Eks

Pod

usermgmt-microservice-7598cb7fcf-txwck

Containerid

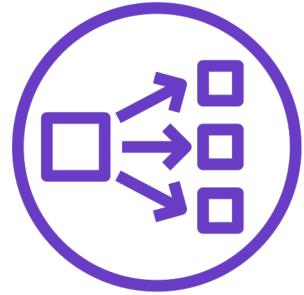
1aeabf8732ab8303eb24a138b31f963566d5d113dd571ea965f5b7ad2dd9e34a



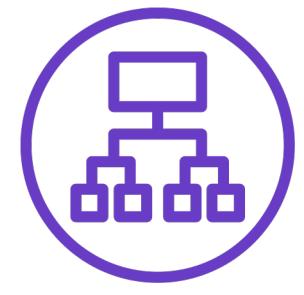
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



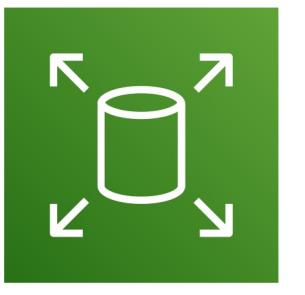
Application  
Load Balancer



Certificate  
Manager



Route53



Elastic Block  
Store



Amazon RDS



# AWS EKS

## Microservices

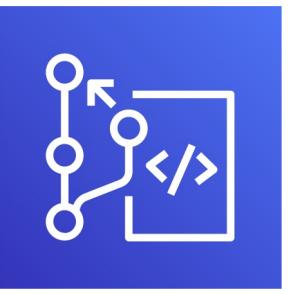
## Canary Deployments



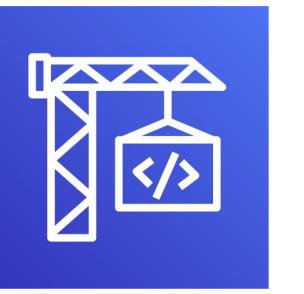
Fargate  
Profiles



Elastic Container  
Registry



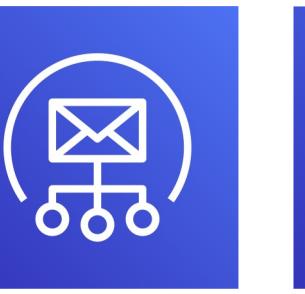
Code  
Commit



Code  
Build



Code  
Pipeline



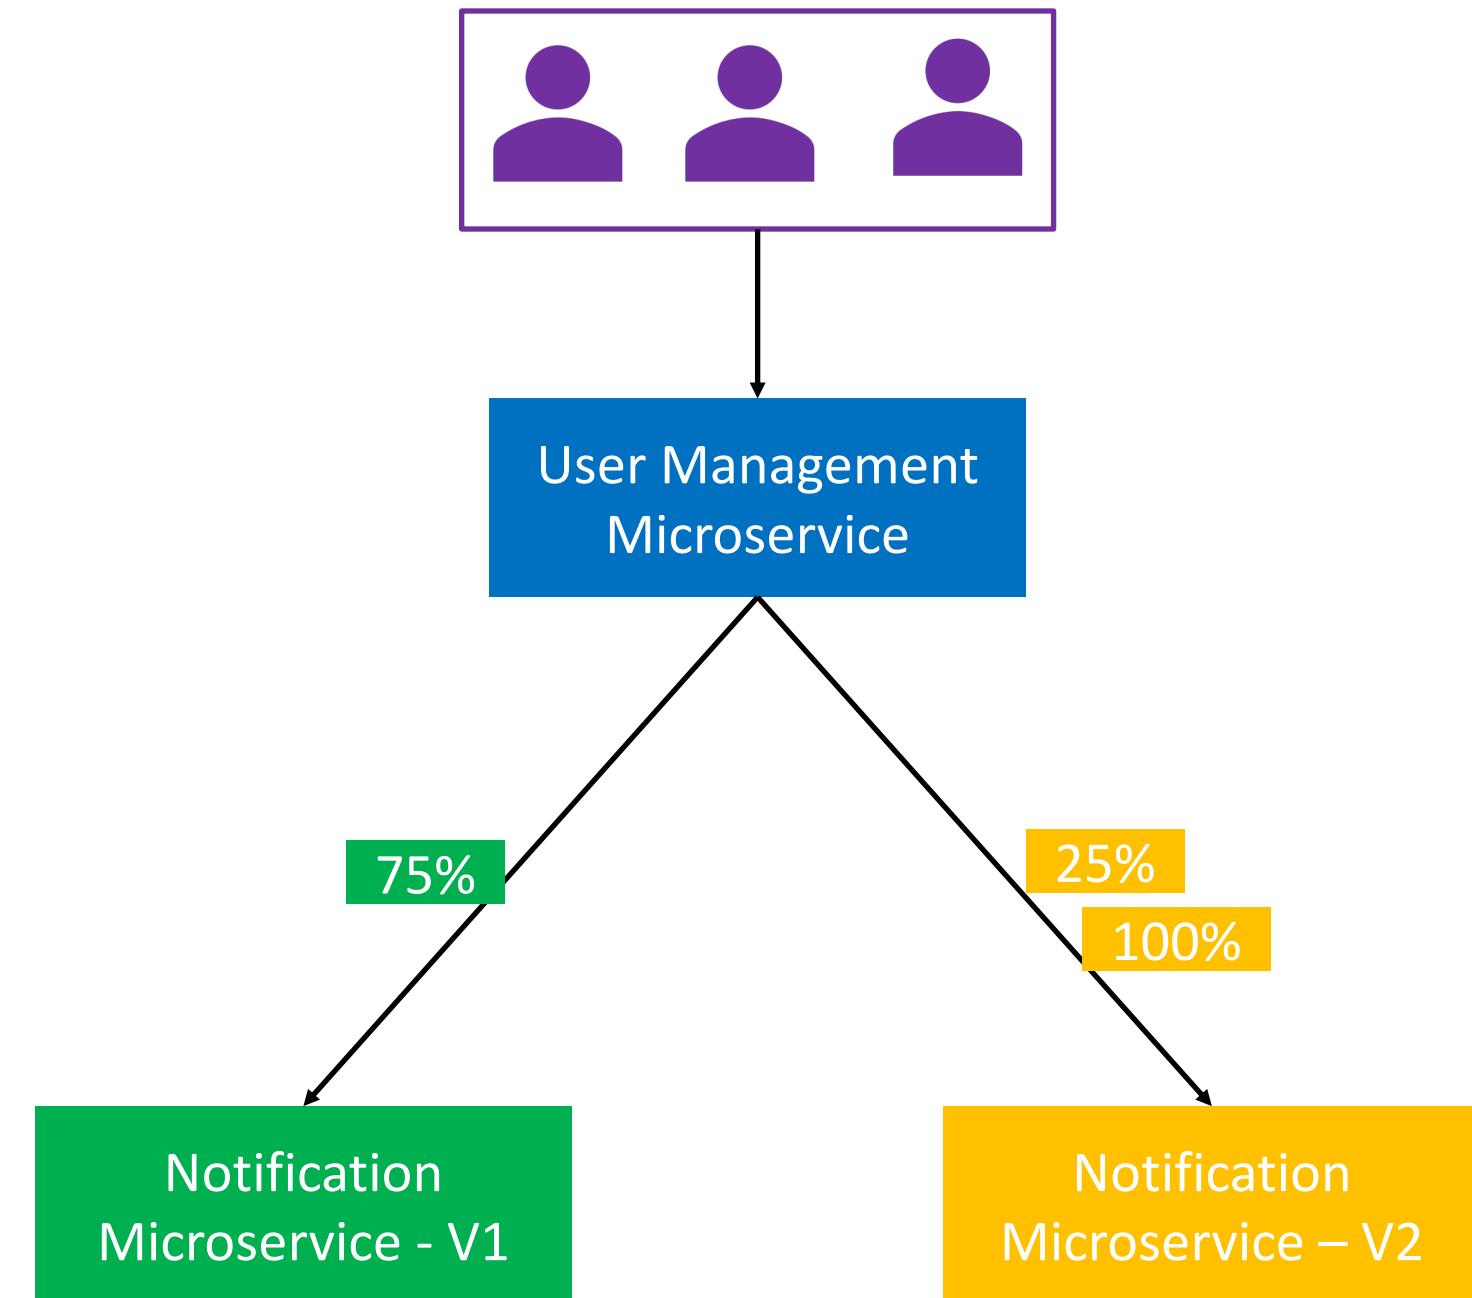
Simple Email  
Service



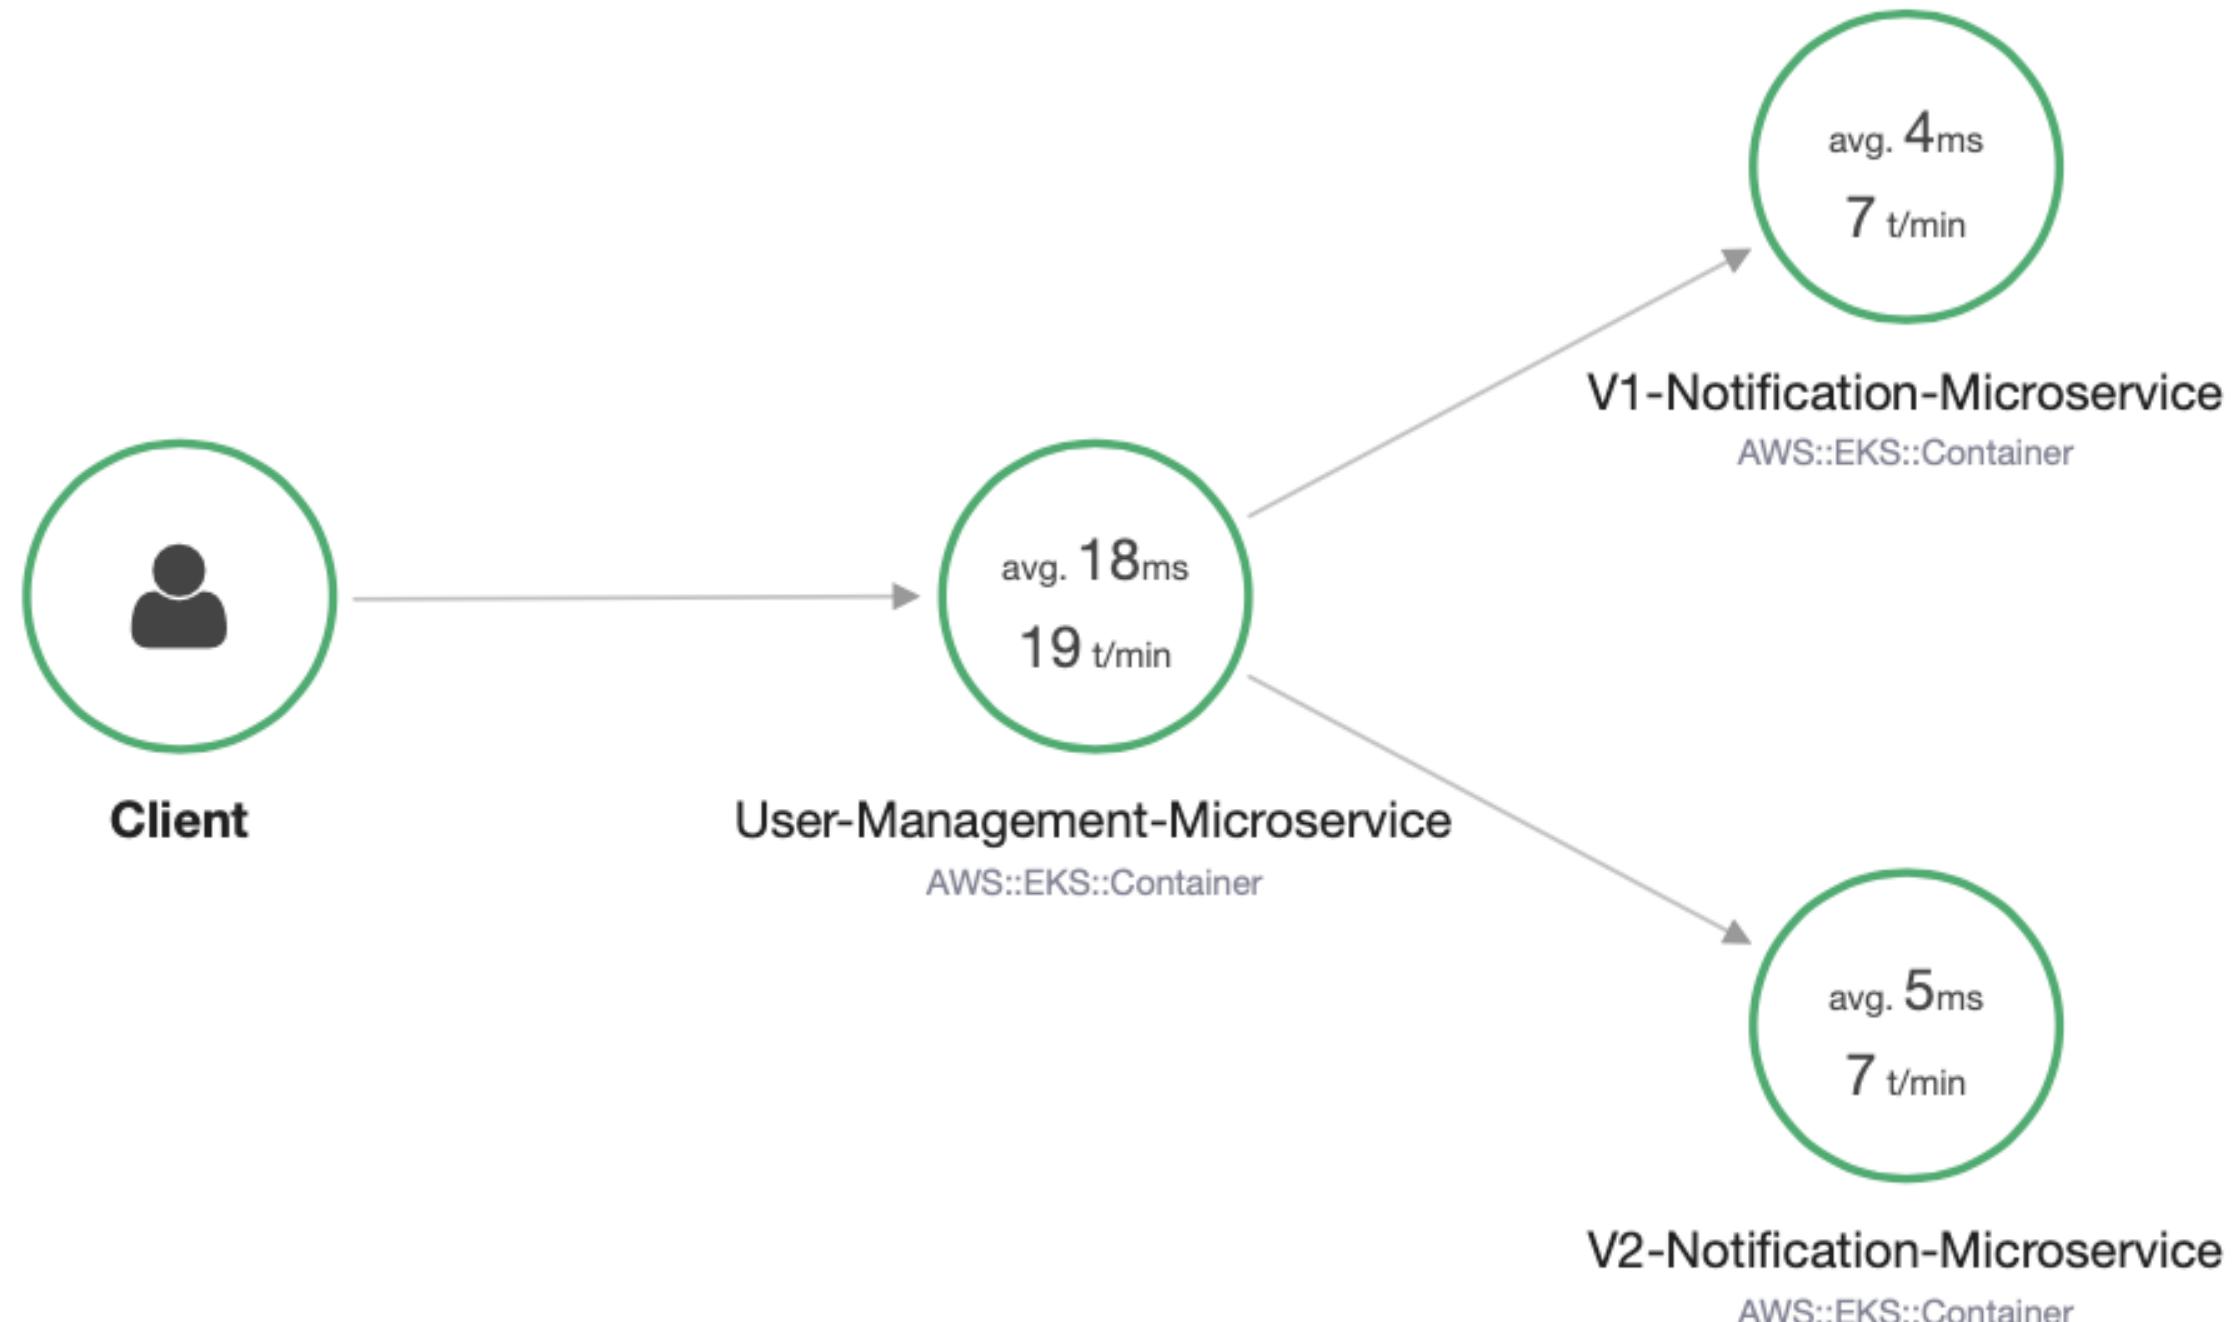
X-Ray

# What are Canary Deployments?

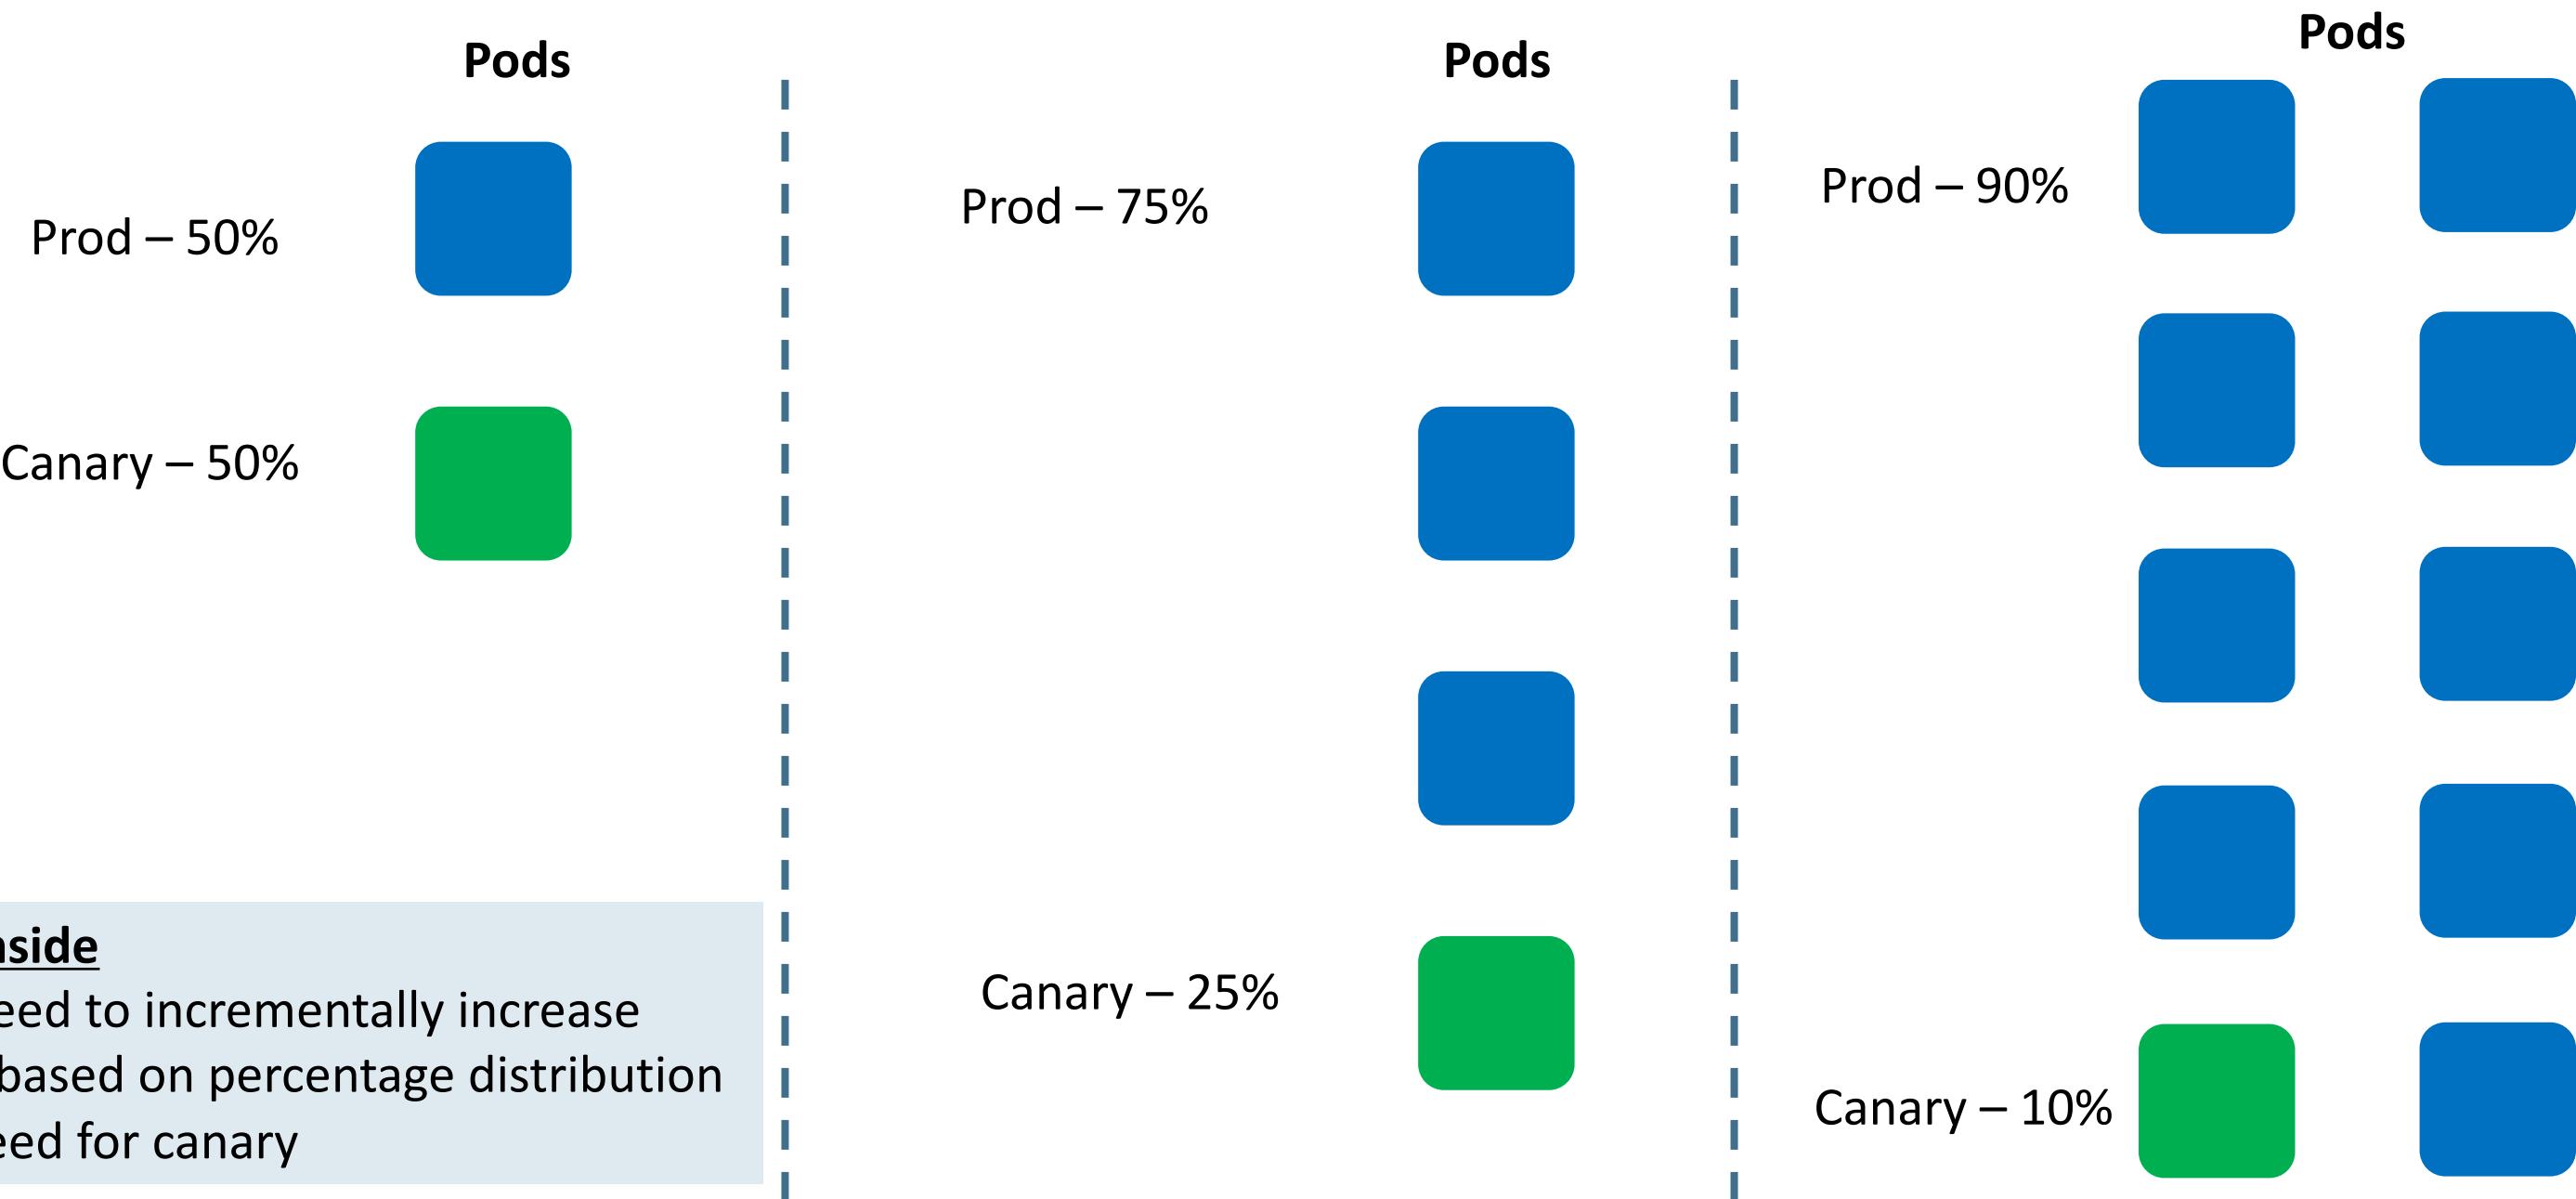
- Canaries means **incremental** rollouts
- With canaries, the **new version** of the application is slowly deployed to the Kubernetes cluster while getting a very small amount of **live traffic**
- In short, **a subset of live users** are connecting to the **new version** while the rest are still using the **previous version**
- Using canaries, we can detect **deployment issues very early** while they effect only a small subset of users
- If we **encounter any issues with a canary**, the production version is still present, and **all traffic can simply be reverted to it.**



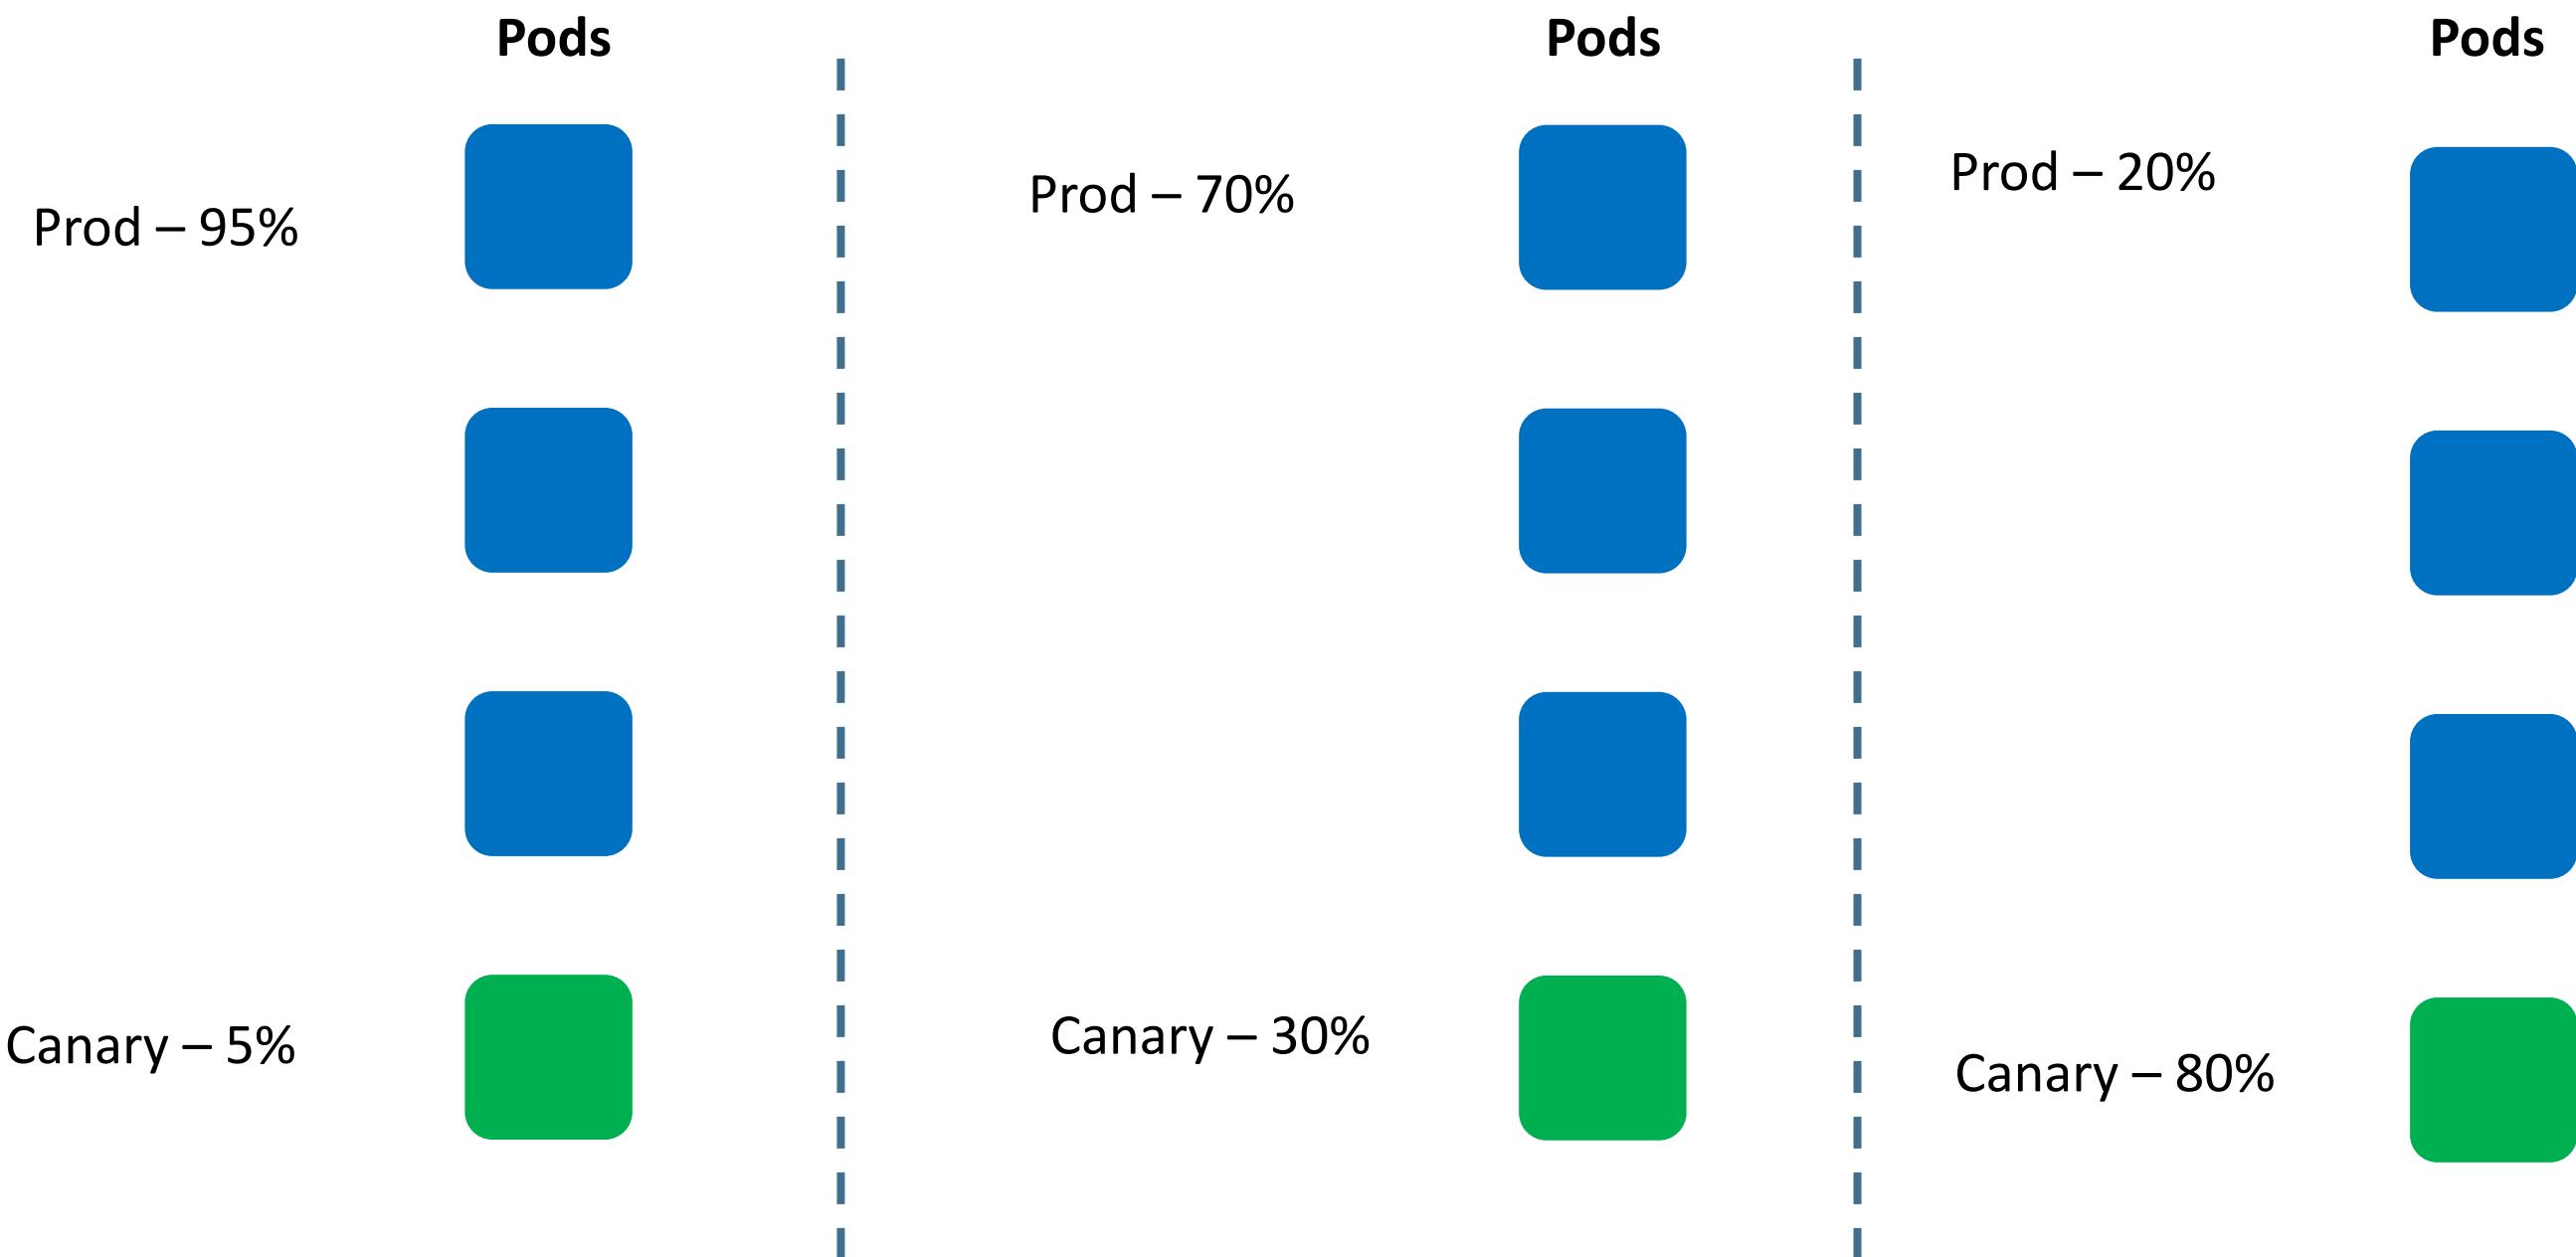
# Microservices – Canary Deployments



# Canary Deployments out of box on Kubernetes



# Canary Deployments on Kubernetes with AWS App Mesh





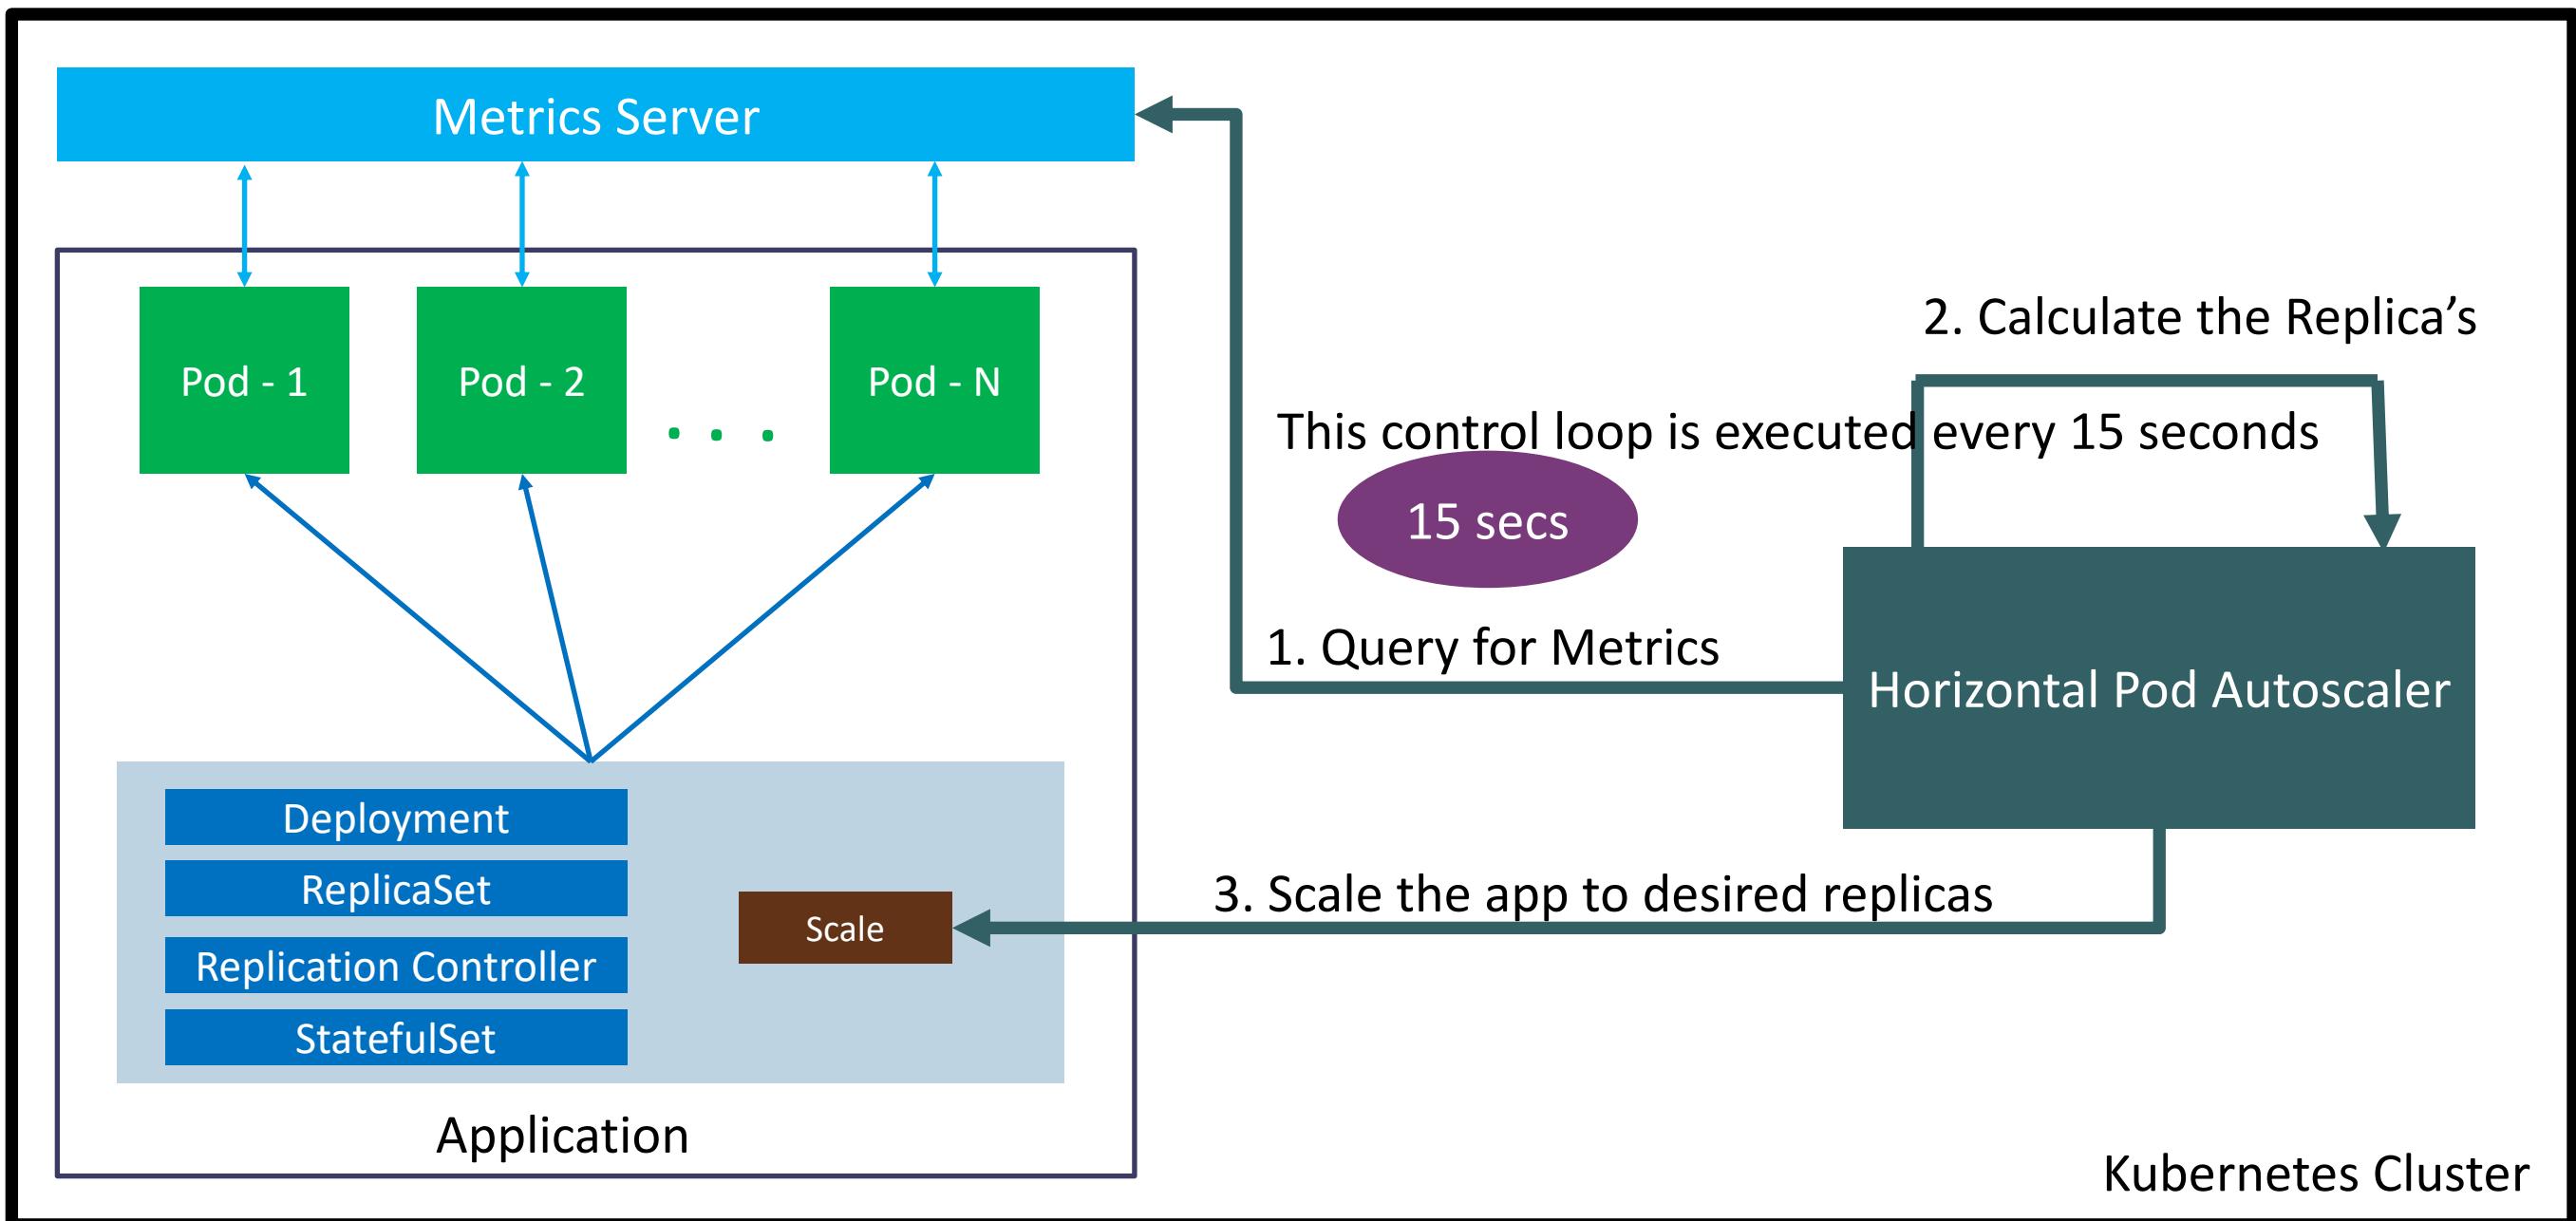
# AWS EKS Autoscaling Horizontal Pod Autoscaler



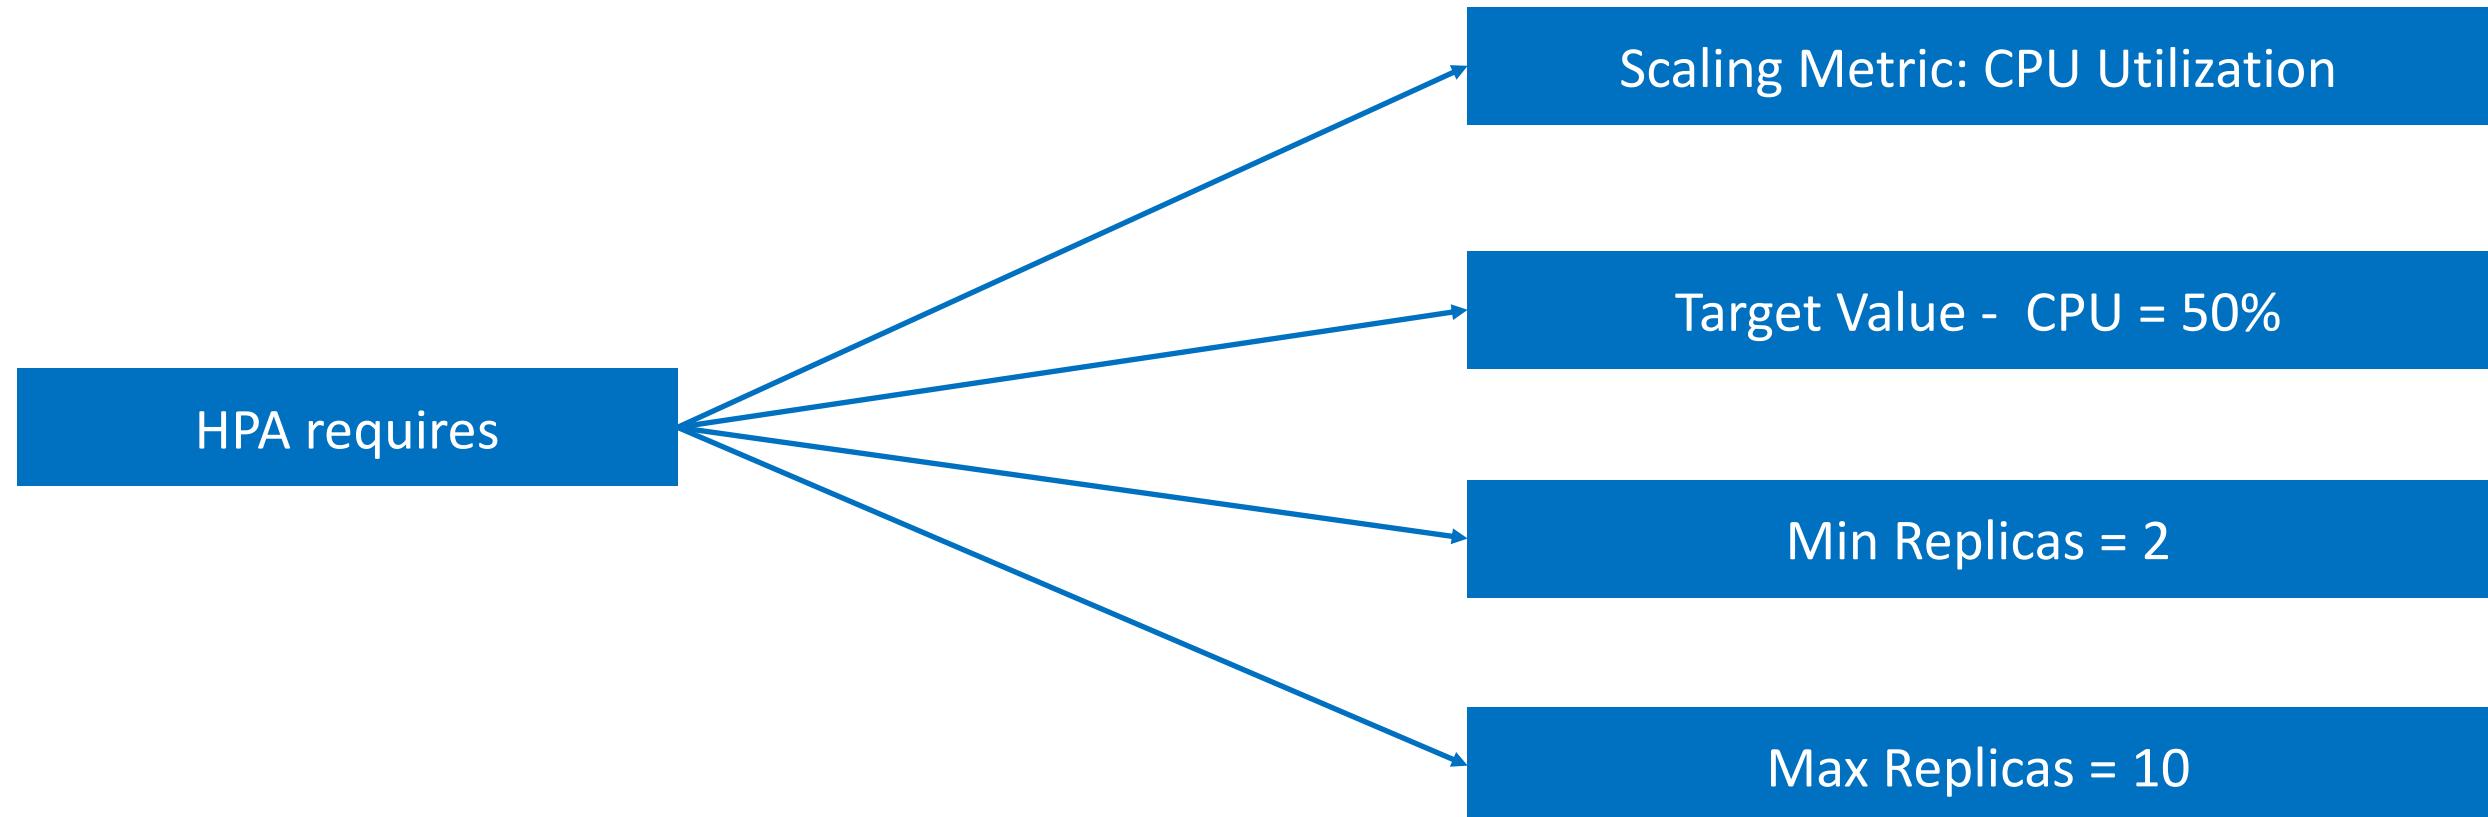
# Horizontal Pod Autoscaler – HPA - Introduction

- In a very simple note Horizontal Scaling means **increasing and decreasing** the number of **Replicas (Pods)**
- HPA **automatically scales** the number of pods in a deployment, replication controller, or replica set, stateful set based on that resource's **CPU utilization**.
- This can help our applications **scale out to meet increased demand** or **scale in when resources are not needed**, thus freeing up your worker nodes for other applications.
- When we set a **target CPU utilization percentage**, the HPA scales our application in or out to try to meet that **target**.
- HPA needs **Kubernetes metrics server** to verify CPU metrics of a pod.
- We **do not need** to **deploy or install the HPA** on our cluster to begin scaling our applications, its out of the box available as a **default** Kubernetes API resource.

# How HPA works?



# How is HPA configured?



```
kubectl autoscale deployment demo-deployment --cpu-percent=50 --min=1 --max=10
```



# AWS EKS Autoscaling Vertical Pod Autoscaler



# Vertical Pod Autoscaler – VPA - Introduction

- VPA automatically adjusts the **CPU and memory reservations** for our pods to help "**right size**" our applications.
- This adjustment can **improve cluster resource utilization** and **free up** CPU and memory for other pods.
- Benefits
  - Cluster nodes are used **efficiently**, because Pods use exactly what they need.
  - Pods are **scheduled onto nodes** that have the appropriate resources available.
  - We **don't have** to run **time-consuming benchmarking tasks** to determine the correct values for CPU and memory requests.
  - **Maintenance time is reduced**, because the autoscaler can adjust CPU and memory requests over time without any action on your part.

## VPA Components

### VPA Admission Hook

Every pod submitted to the k8s cluster goes through this webhook automatically which checks whether a **VerticalPodAutoscaler** object is referencing this pod or one of its parents (a ReplicaSet, a Deployment, etc).

### VPA Recommender

Connects to the **metrics-server** in the cluster, fetches historical and current usage data (CPU and memory) for each VPA-enabled pod and generates recommendations for **scaling up or down** the requests and limits of these pods.

### VPA Updater

**Runs every 1 minute.** If a pod is not running in the calculated recommendation range, **it evicts the currently running version of this pod**, so it can restart and go through the **VPA admission webhook** which will change the CPU and memory settings for it, before it can start.



# AWS EKS Autoscaling Cluster Autoscaler



# Cluster Autoscaler - Introduction

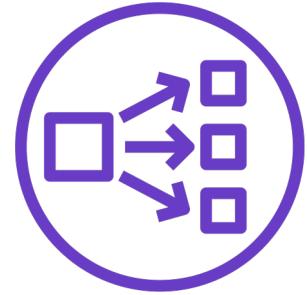
- Cluster Autoscaler is a tool that automatically adjusts the size of a Kubernetes cluster when one of the following conditions is true:
- There are pods that failed to run in the cluster due to insufficient resources.
- There are nodes in the cluster that have been underutilized for an extended period of time and their pods can be placed on other existing nodes.
- The Cluster Autoscaler modifies our worker node groups so that they scale out when we need more resources and scale in when we have underutilized resources.



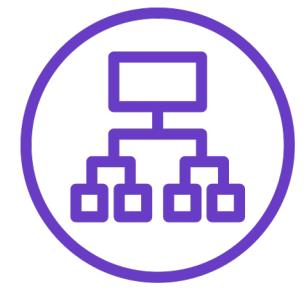
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



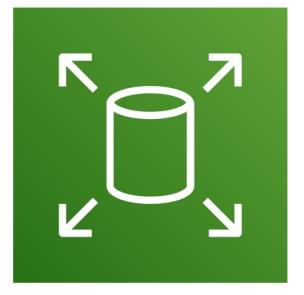
Application  
Load Balancer



Certificate  
Manager



Route53



Elastic Block  
Store



Amazon RDS



# AWS EKS

## CloudWatch

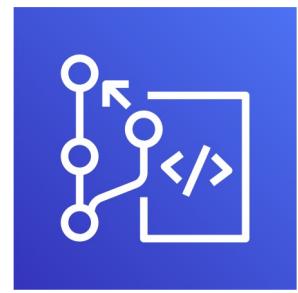
## Container Insights



Fargate  
Profiles



Elastic Container  
Registry



Code  
Commit



Code  
Build



Code  
Pipeline



Simple Email  
Service



X-Ray



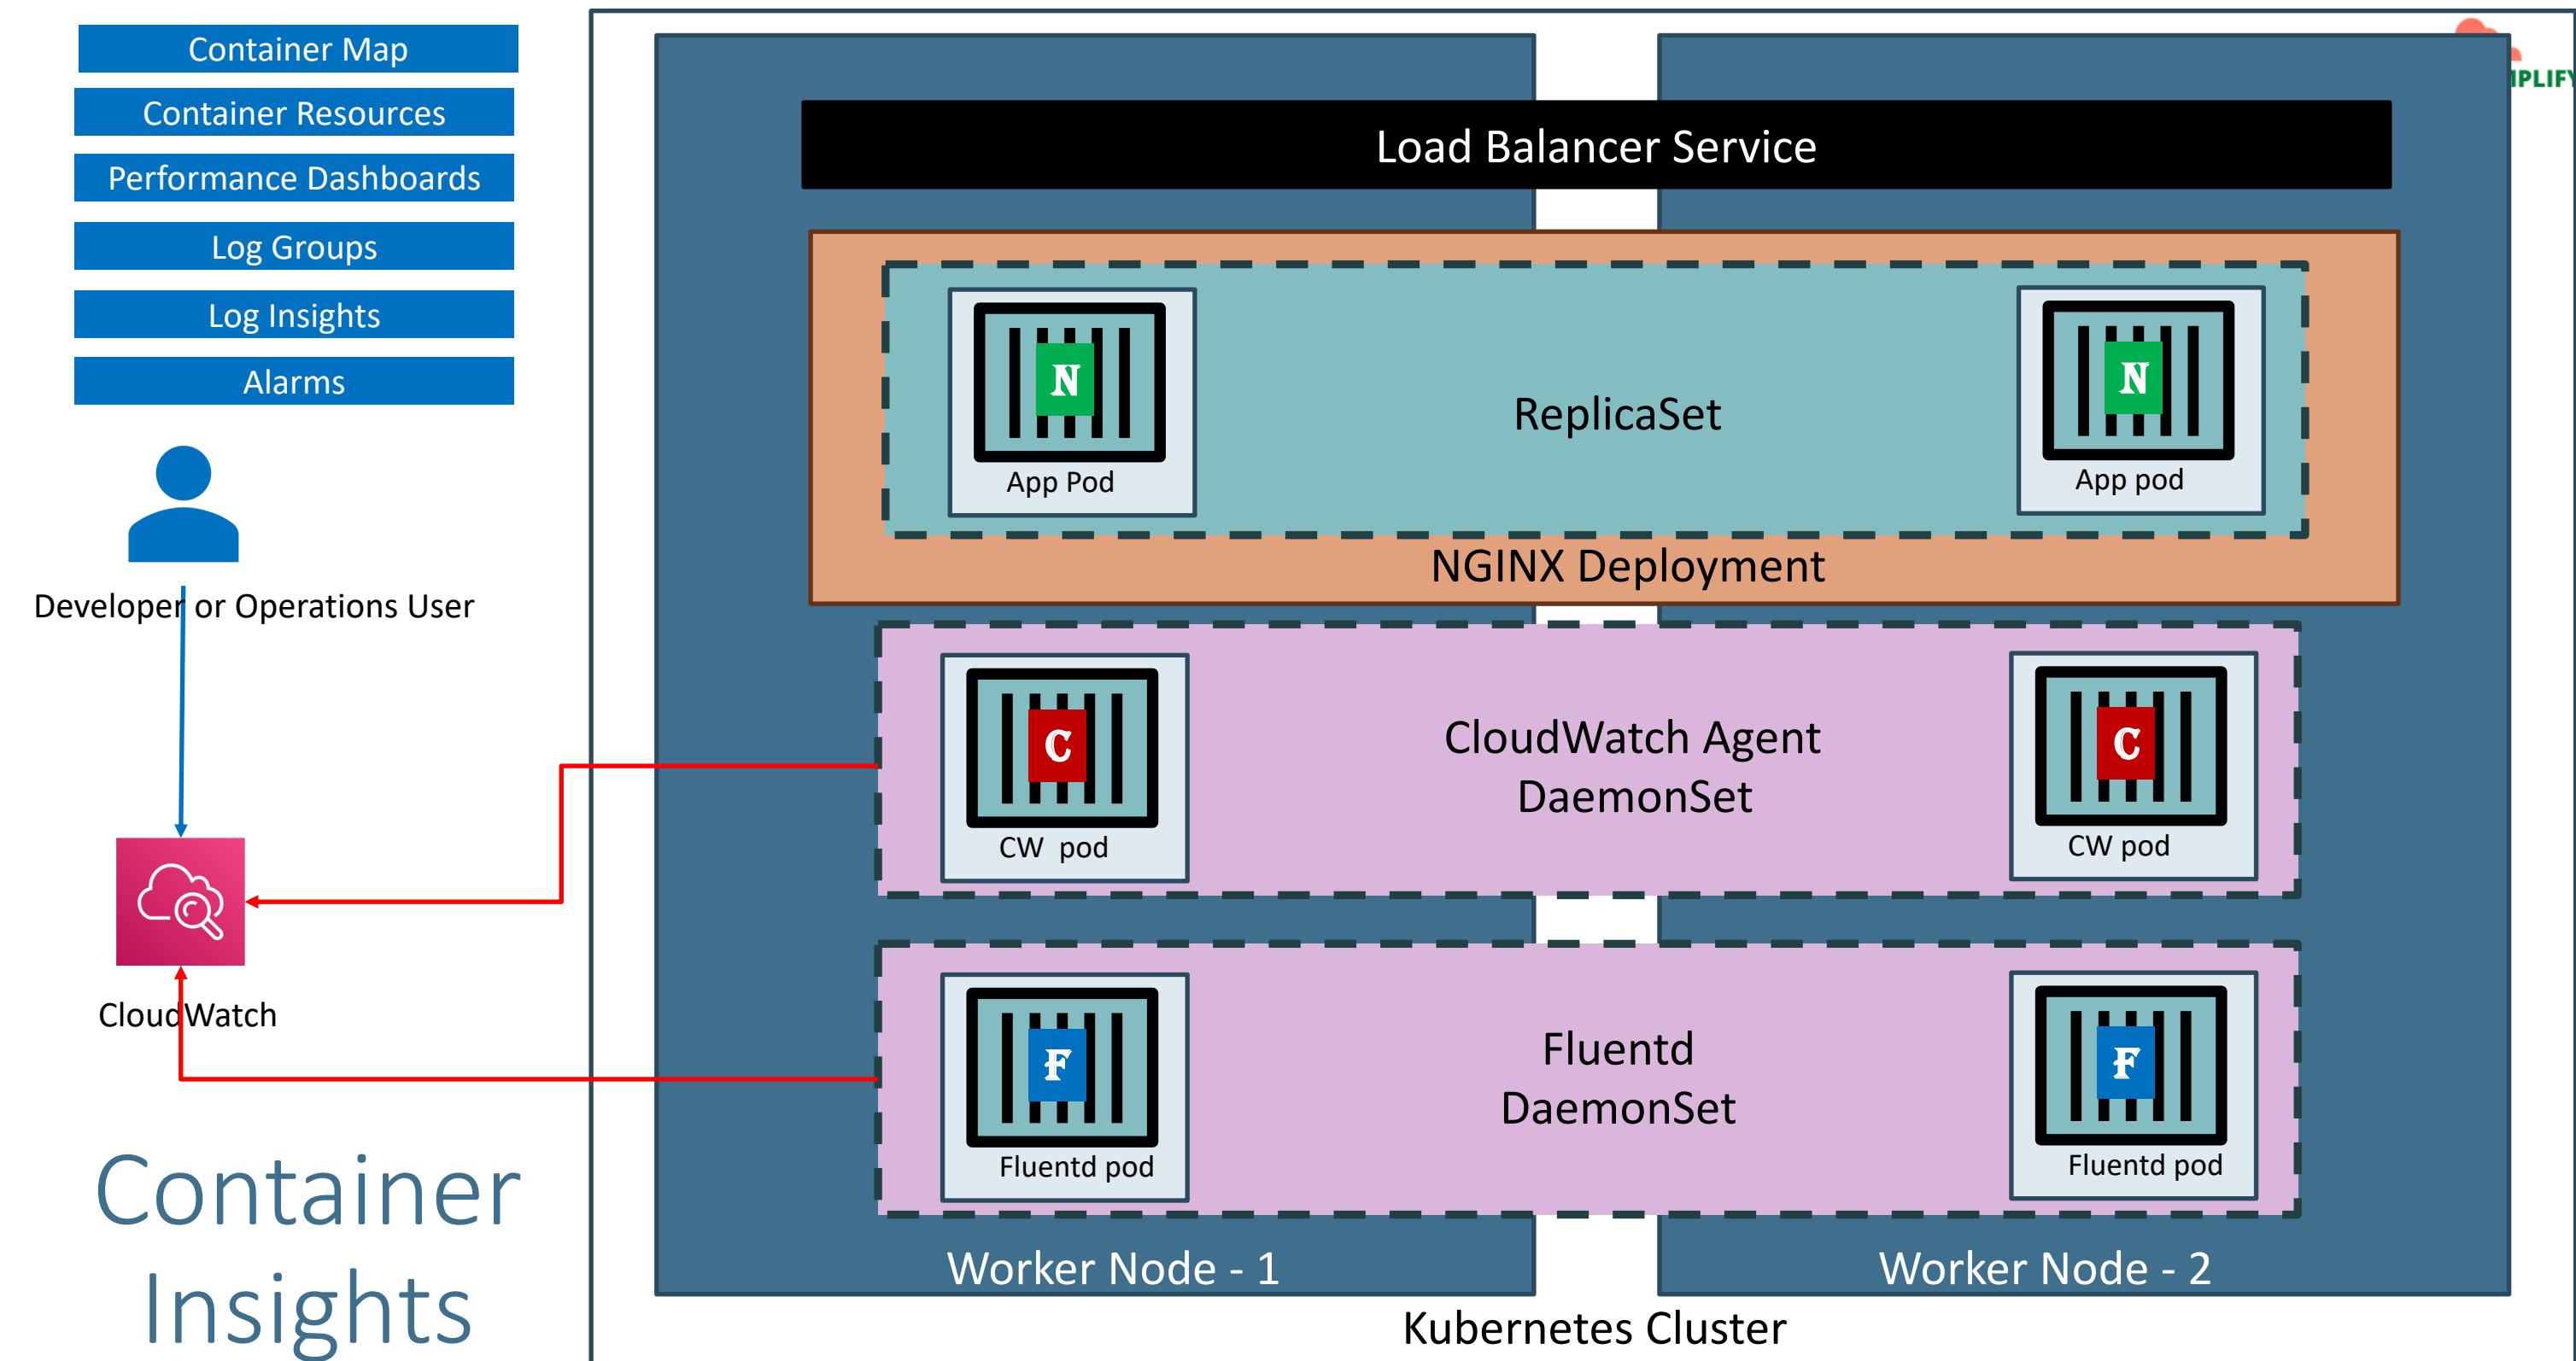
CloudWatch  
Metrics



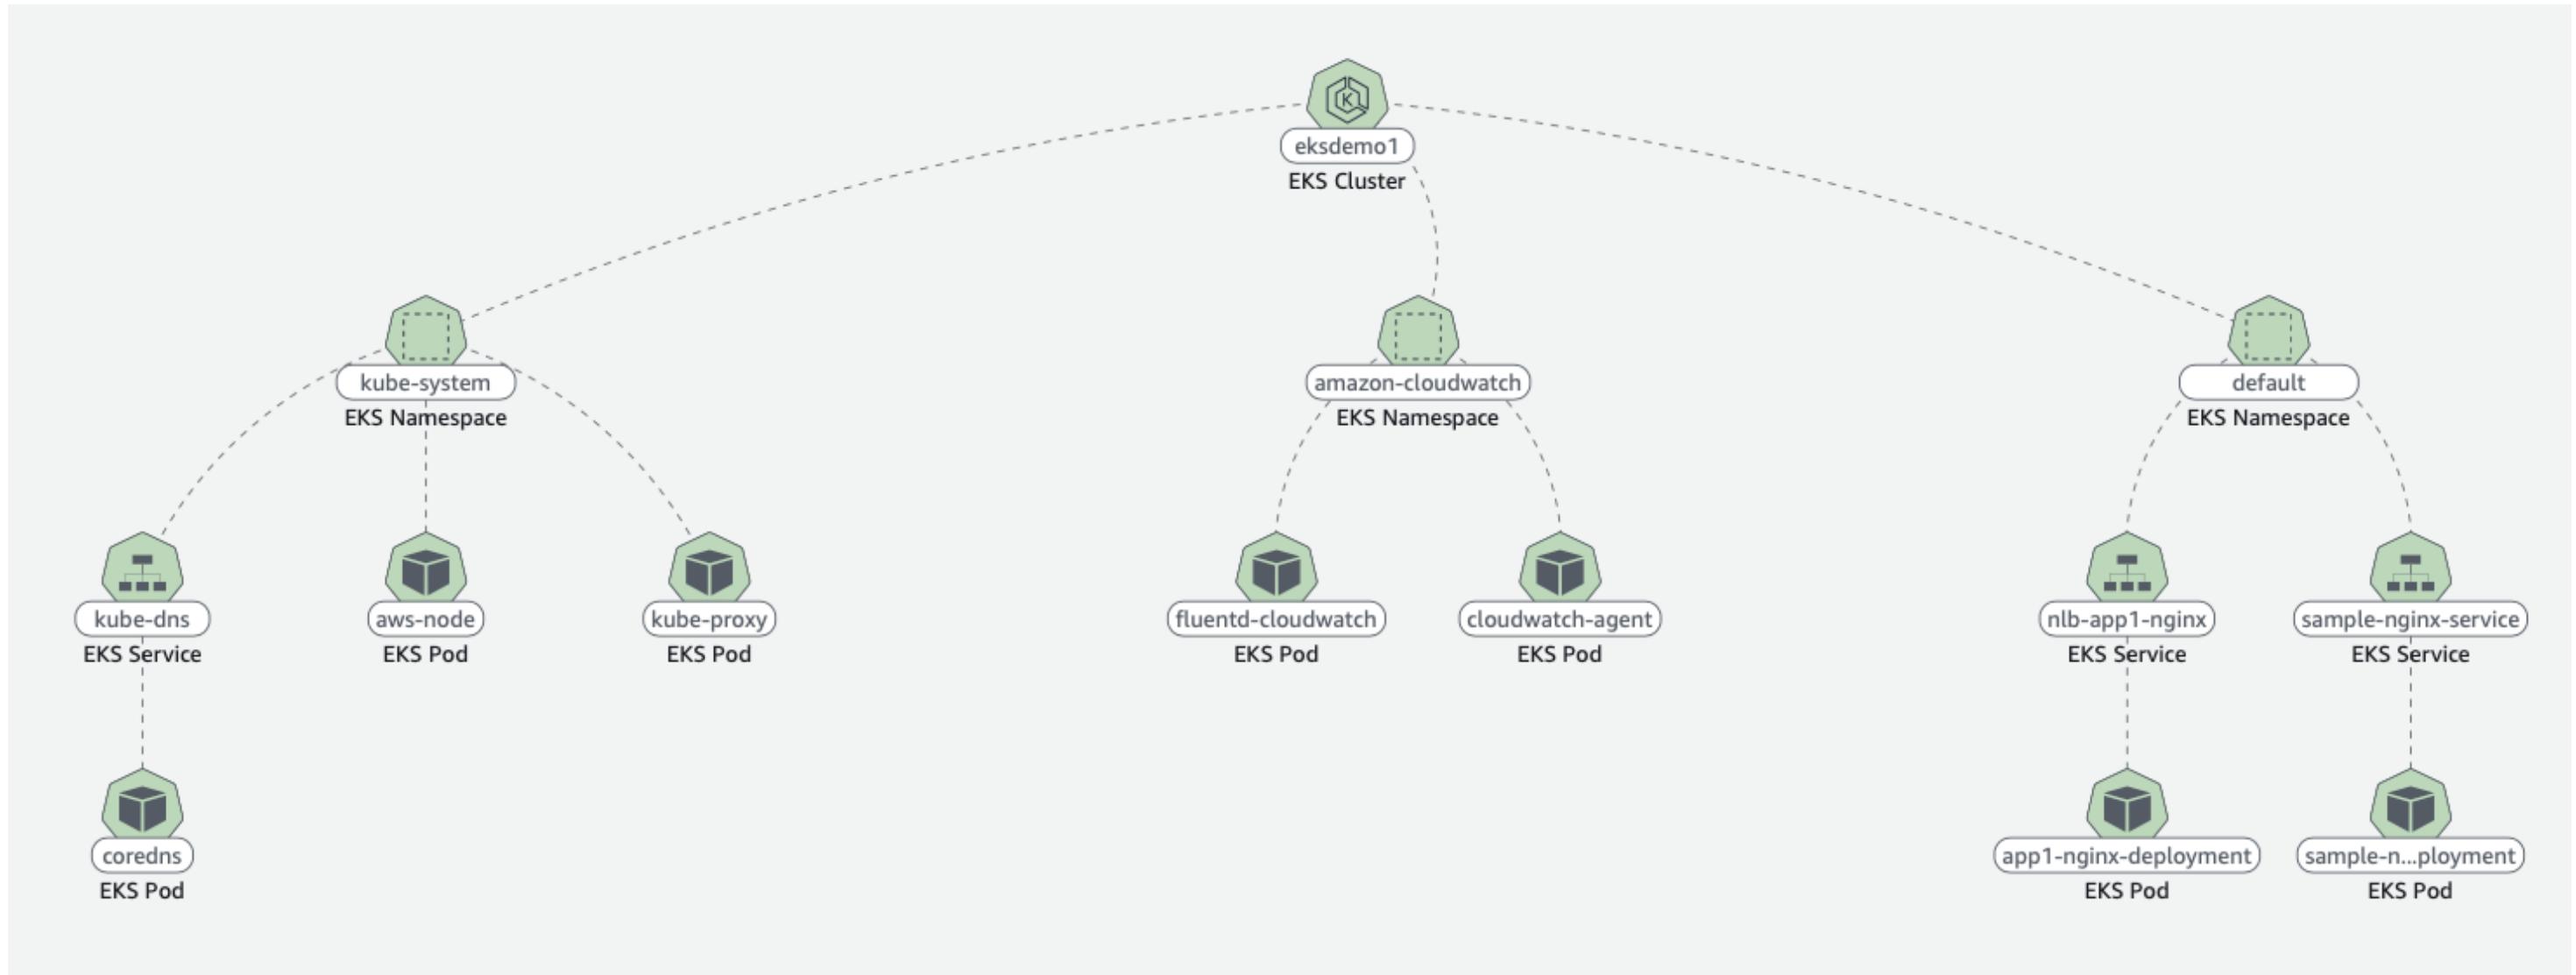
Simple  
Notification Service

# Container Insights

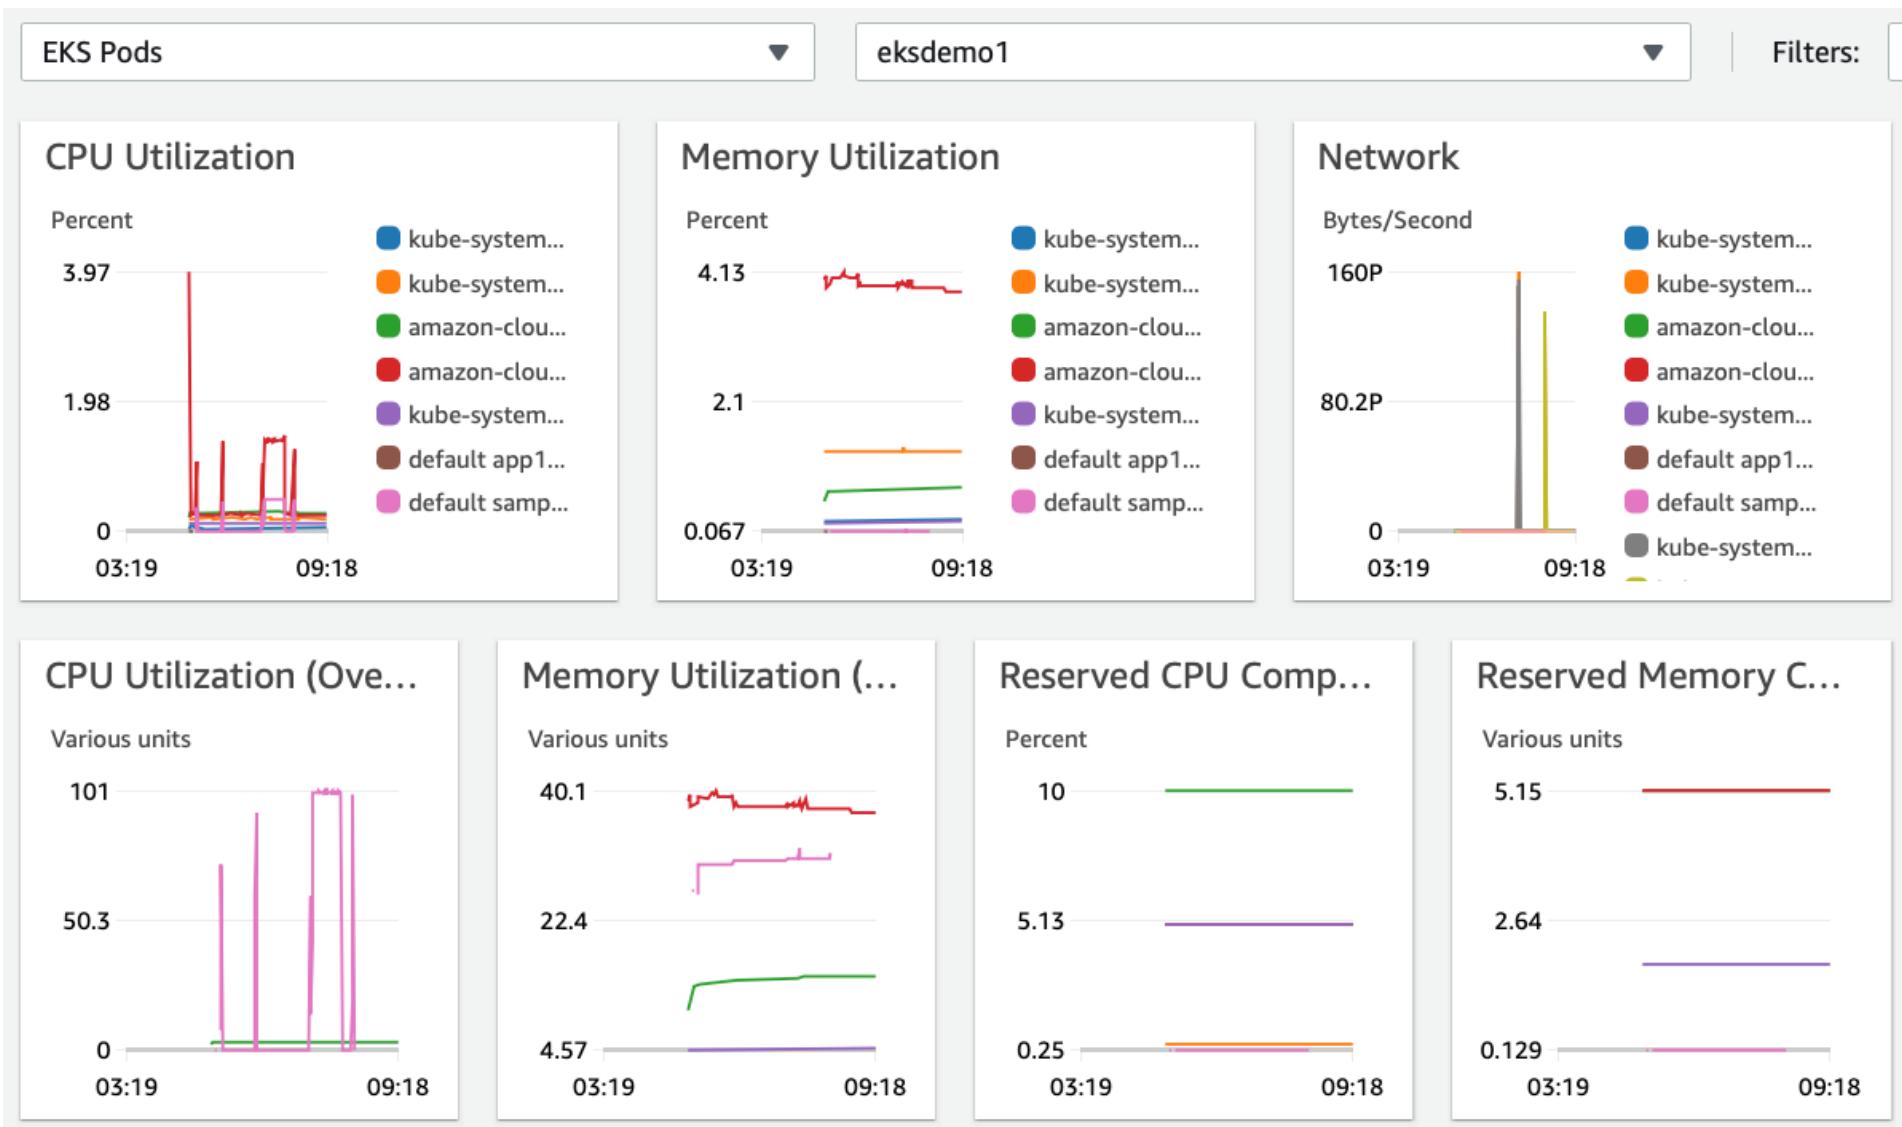
- A **fully managed observability service** for monitoring, troubleshooting and alarming on our containerized applications.
- Container Insights to **collect, aggregate, and summarize metrics and logs** from our containerized applications and microservices.
- The metrics include **utilization for resources** such as CPU, memory, disk, and network.
- It also provides **diagnostic information**, such as container restart failures, to help us isolate issues and resolve them quickly.
- We can also set **CloudWatch alarms** on metrics that Container Insights collects.
- The metrics that Container Insights collects are available in **CloudWatch automatic dashboards**.
- We can analyze and troubleshoot container performance and logs data with **CloudWatch Logs Insights**.



# CloudWatch Container Insights Map

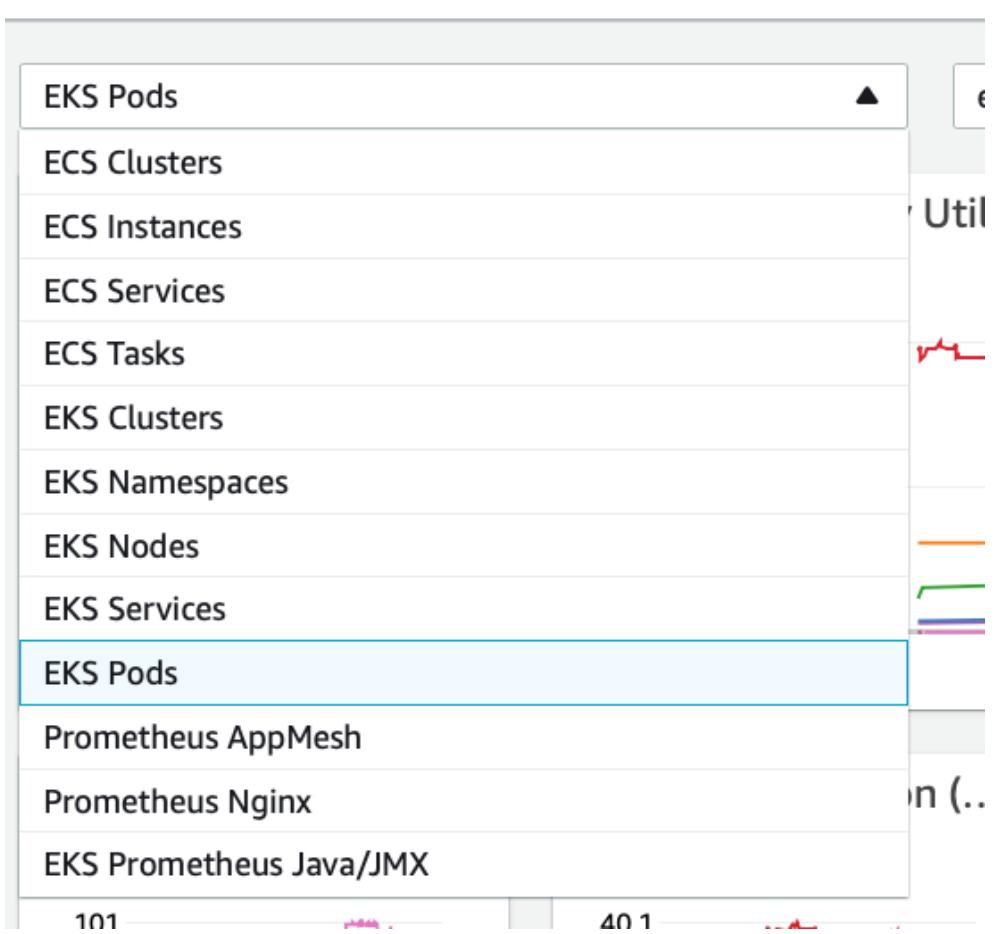


# Automatic Performance Dashboard



[CloudWatch](#) > [Container Insights](#) > [Performance monitoring](#)

## Performance monitoring



# THANK YOU

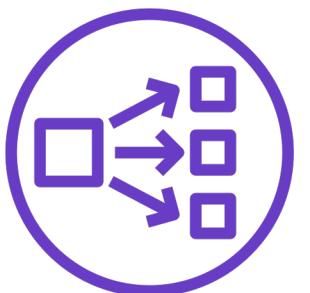
# ALB Ingress Controller v1.x DEPRECATED Slides Continues



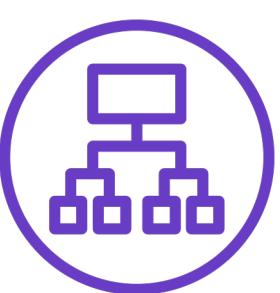
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



Application  
Load Balancer



Amazon RDS

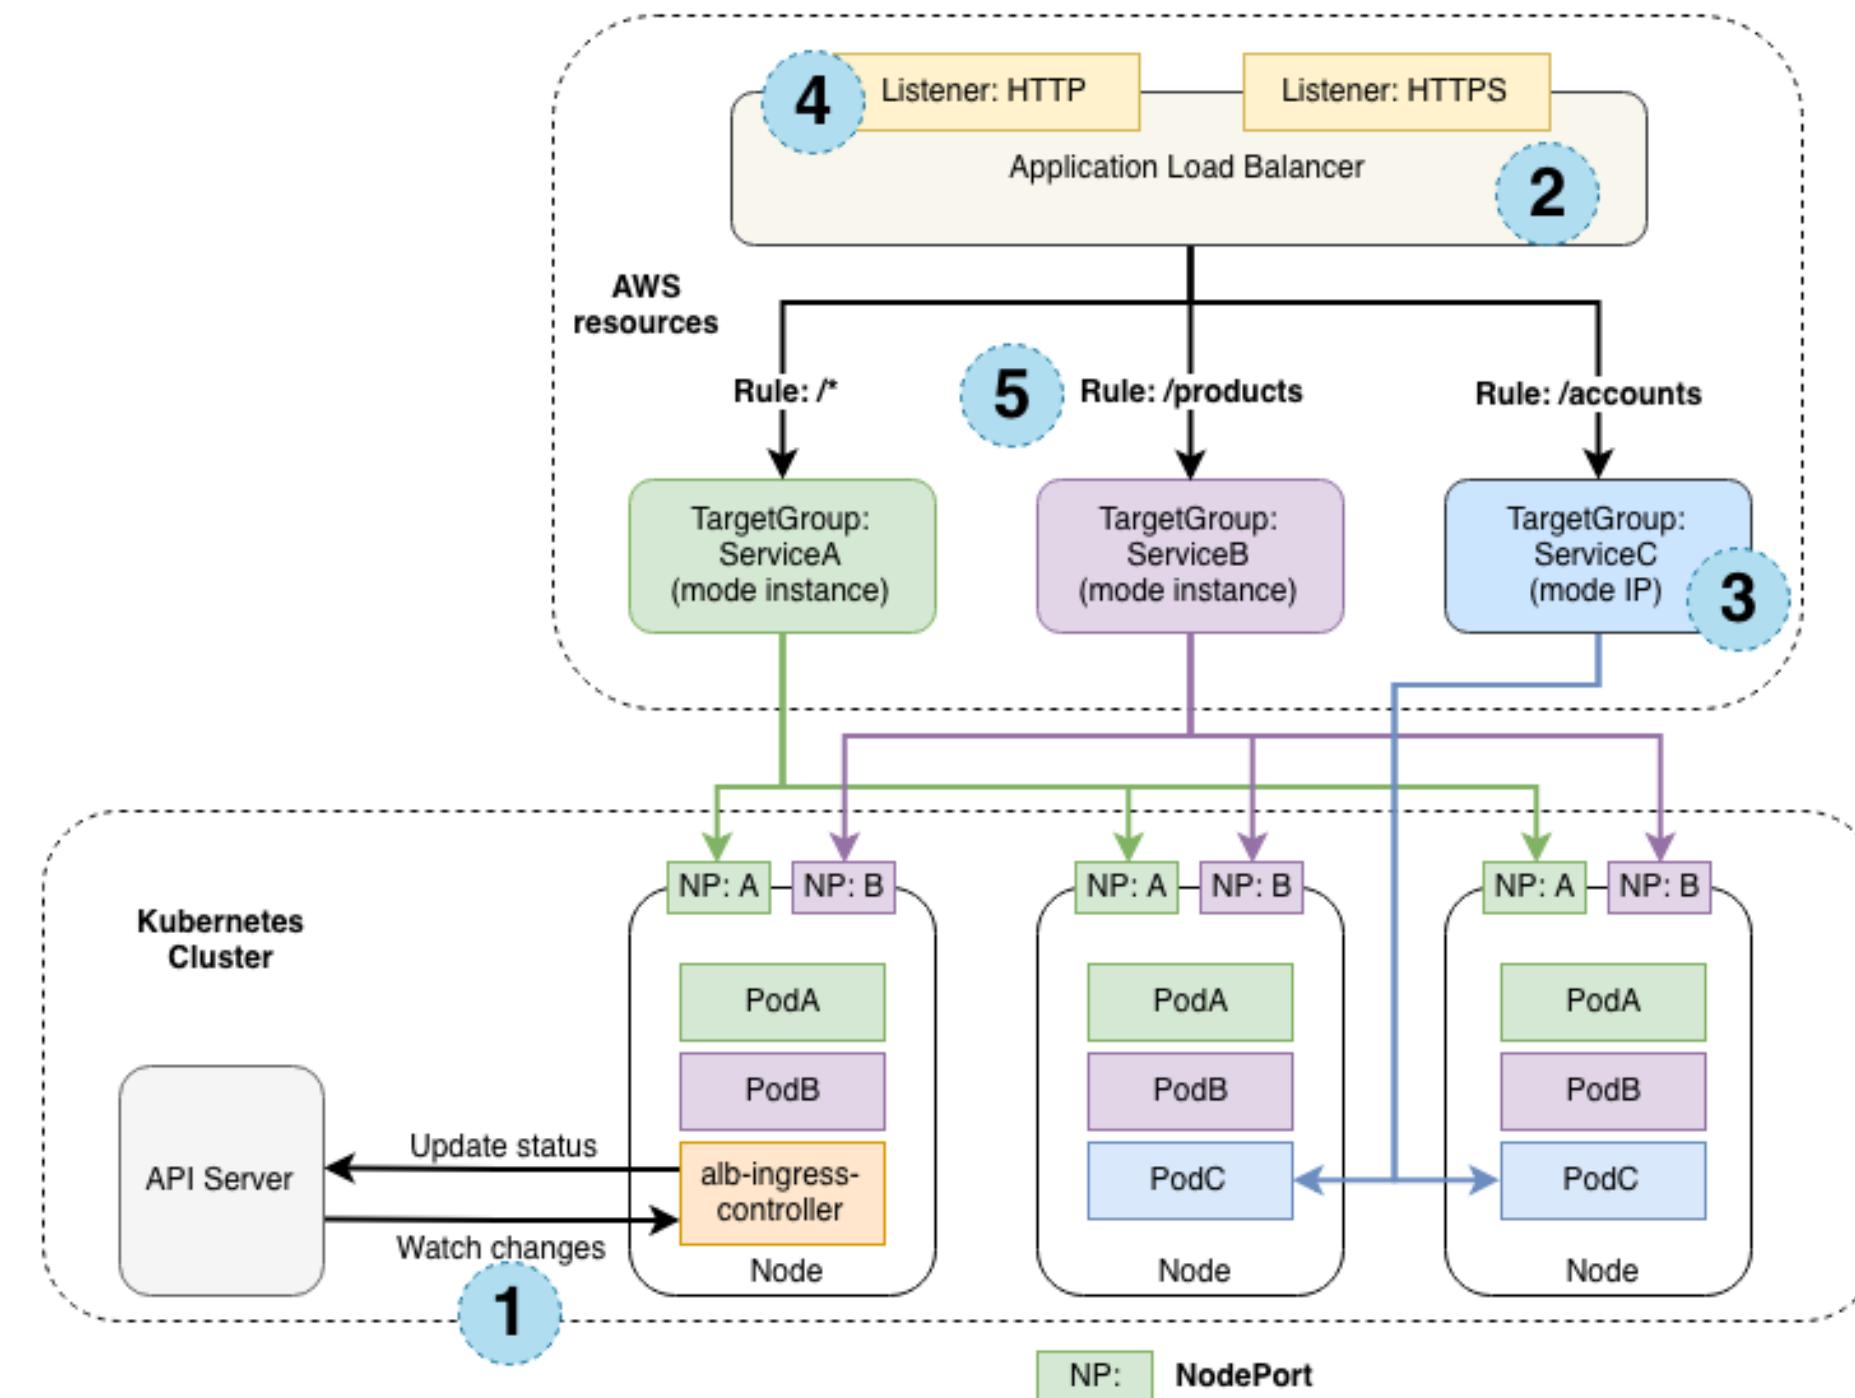


AWS EKS  
&  
RDS & ELB



Application Load Balancer

# How Ingress Works?



## k8s ClusterRole

```

apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  labels:
    app.kubernetes.io/name: alb-ingress-controller
  name: alb-ingress-controller
rules:
- apiGroups:
  - ""
  - extensions
  resources:
  - configmaps
  - endpoints
  - events
  - ingresses
  - ingresses/status
  - services
  - pods/status
  verbs:
  - create
  - get
  - list
  - update
  - watch
  - patch
- apiGroups:
  - ""
  - extensions
  resources:
  - nodes
  - pods
  - secrets
  - services
  - namespaces
  verbs:
  - get
  - list
  - watch

```

## k8s ServiceAccount

```

apiVersion: v1
kind: ServiceAccount
metadata:
  labels:
    app.kubernetes.io/name: alb-ingress-controller
  name: alb-ingress-controller
  namespace: kube-system

```

## k8s ClusterRoleBinding

```

apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  labels:
    app.kubernetes.io/name: alb-ingress-controller
  name: alb-ingress-controller
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: alb-ingress-controller
subjects:
- kind: ServiceAccount
  name: alb-ingress-controller
  namespace: kube-system

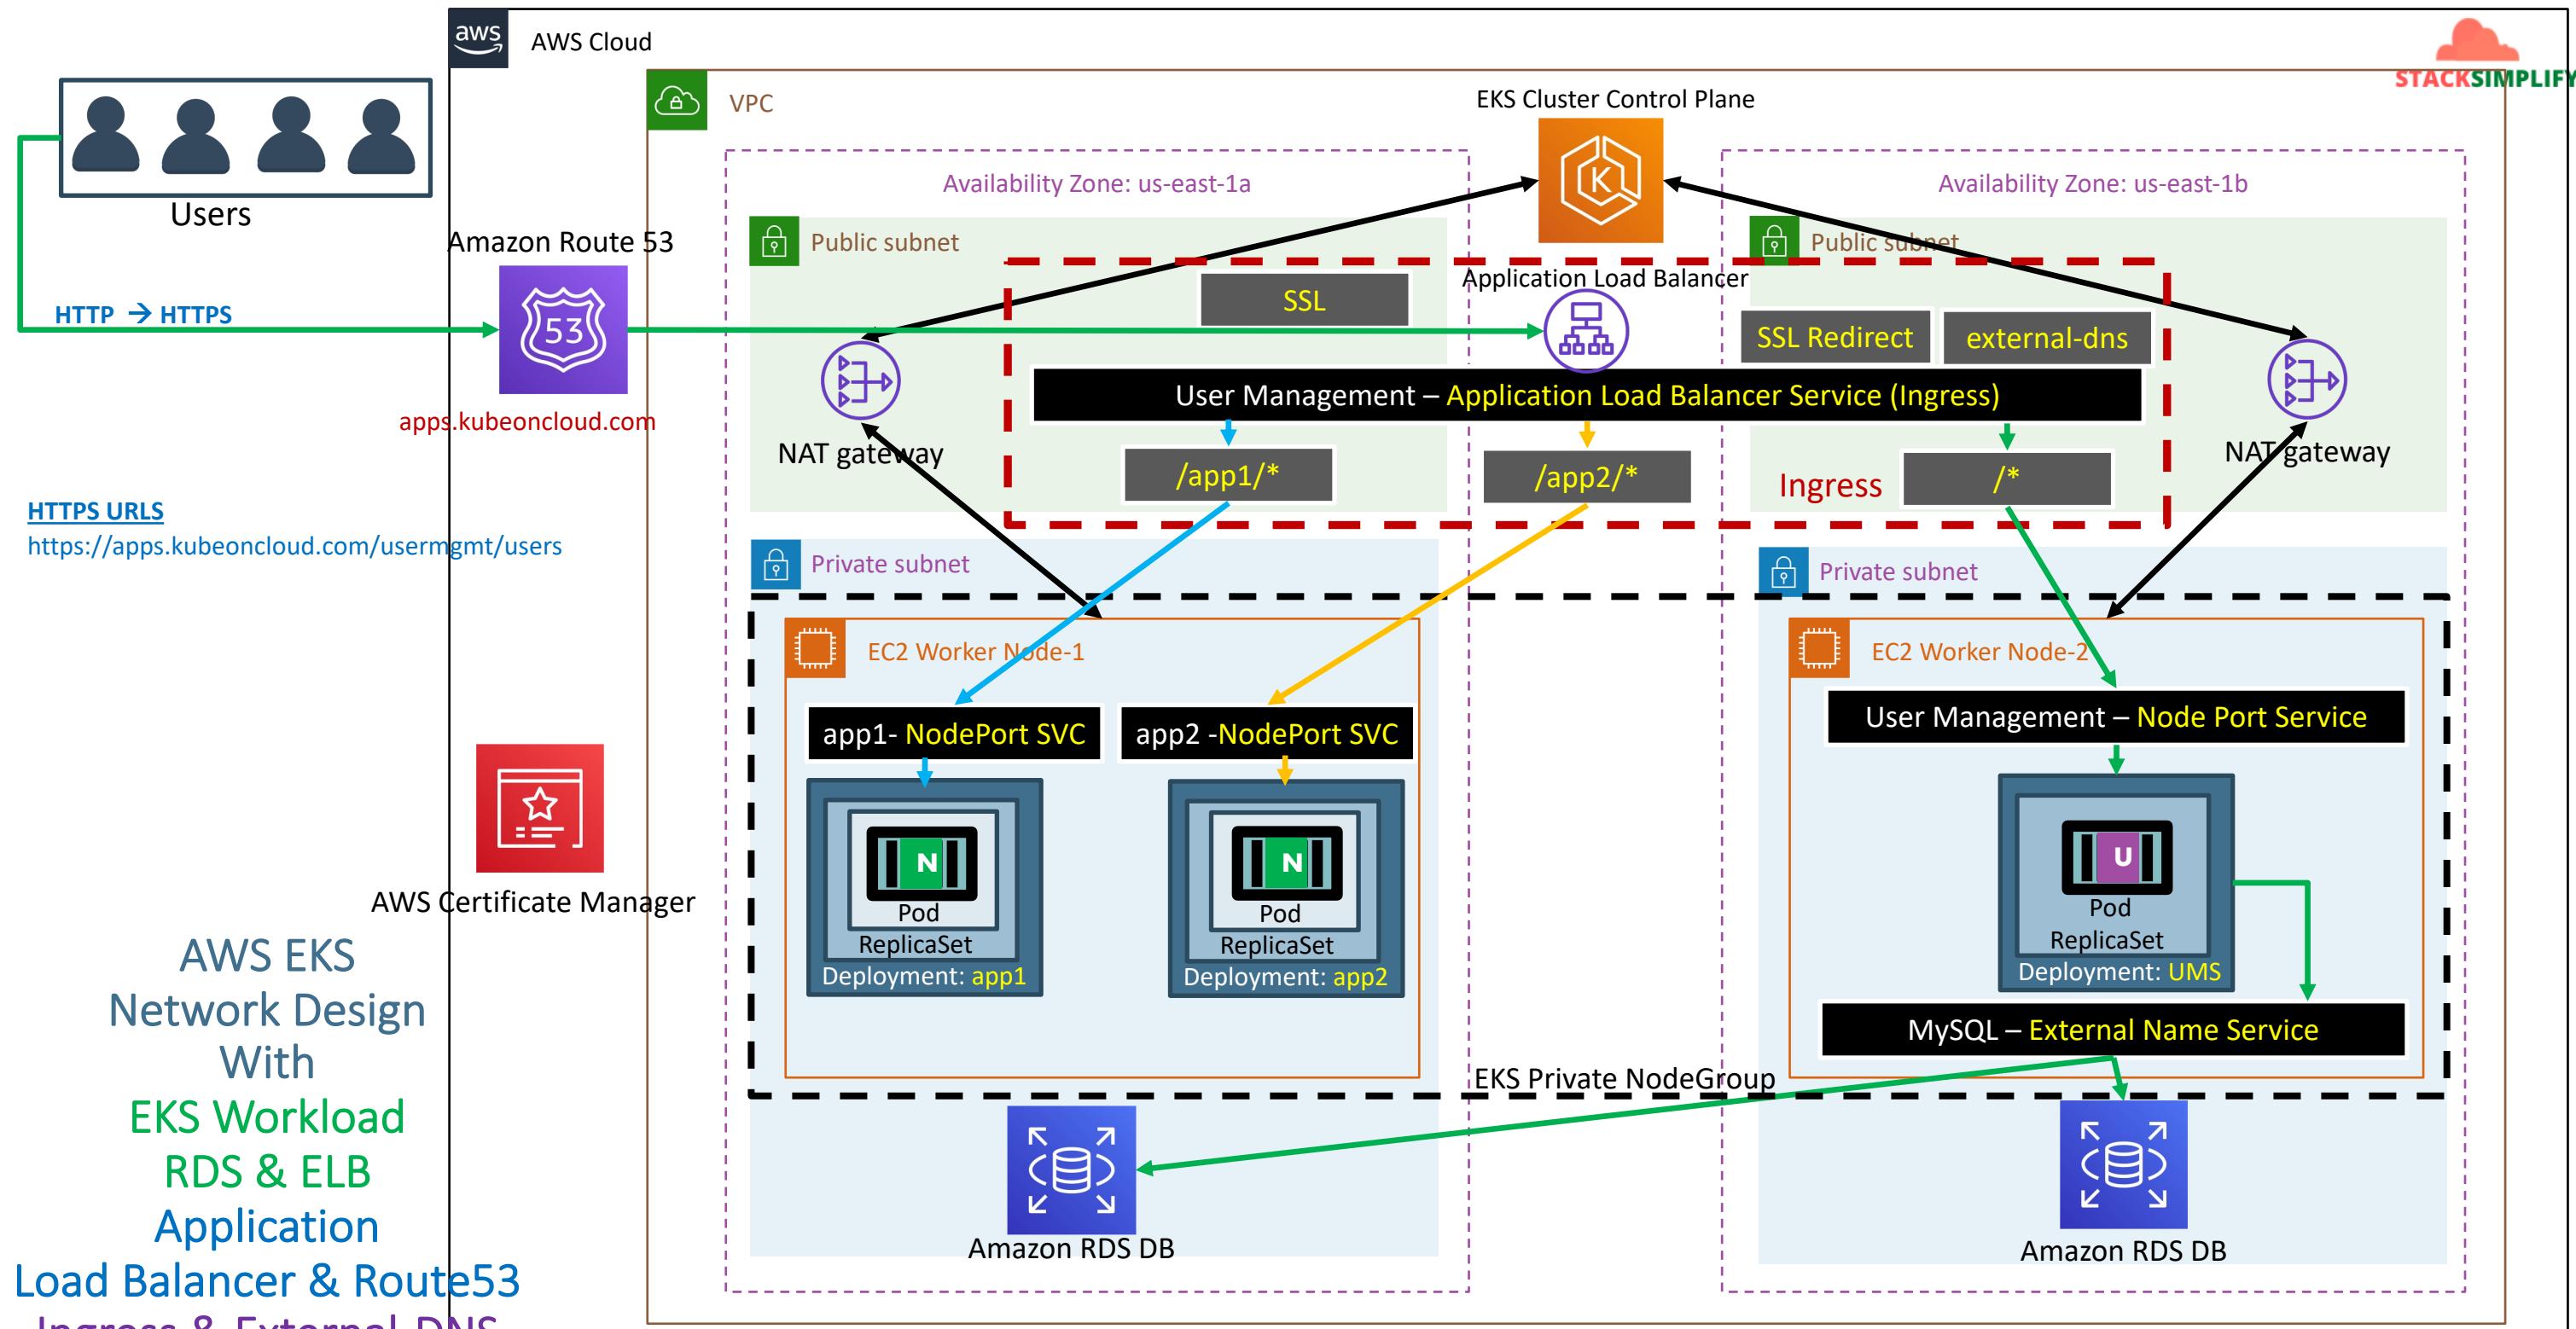
```

## AWS IAM Policy

<https://raw.githubusercontent.com/kubernetes-sigs/aws-alb-ingress-controller/master/docs/examples/iam-policy.json>



STACKSIMPLIFY

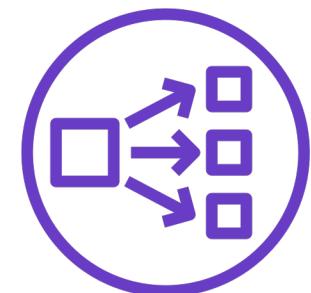




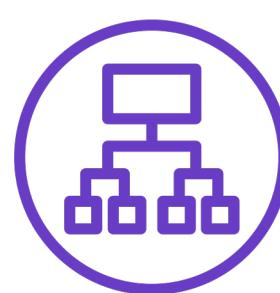
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



Application  
Load Balancer

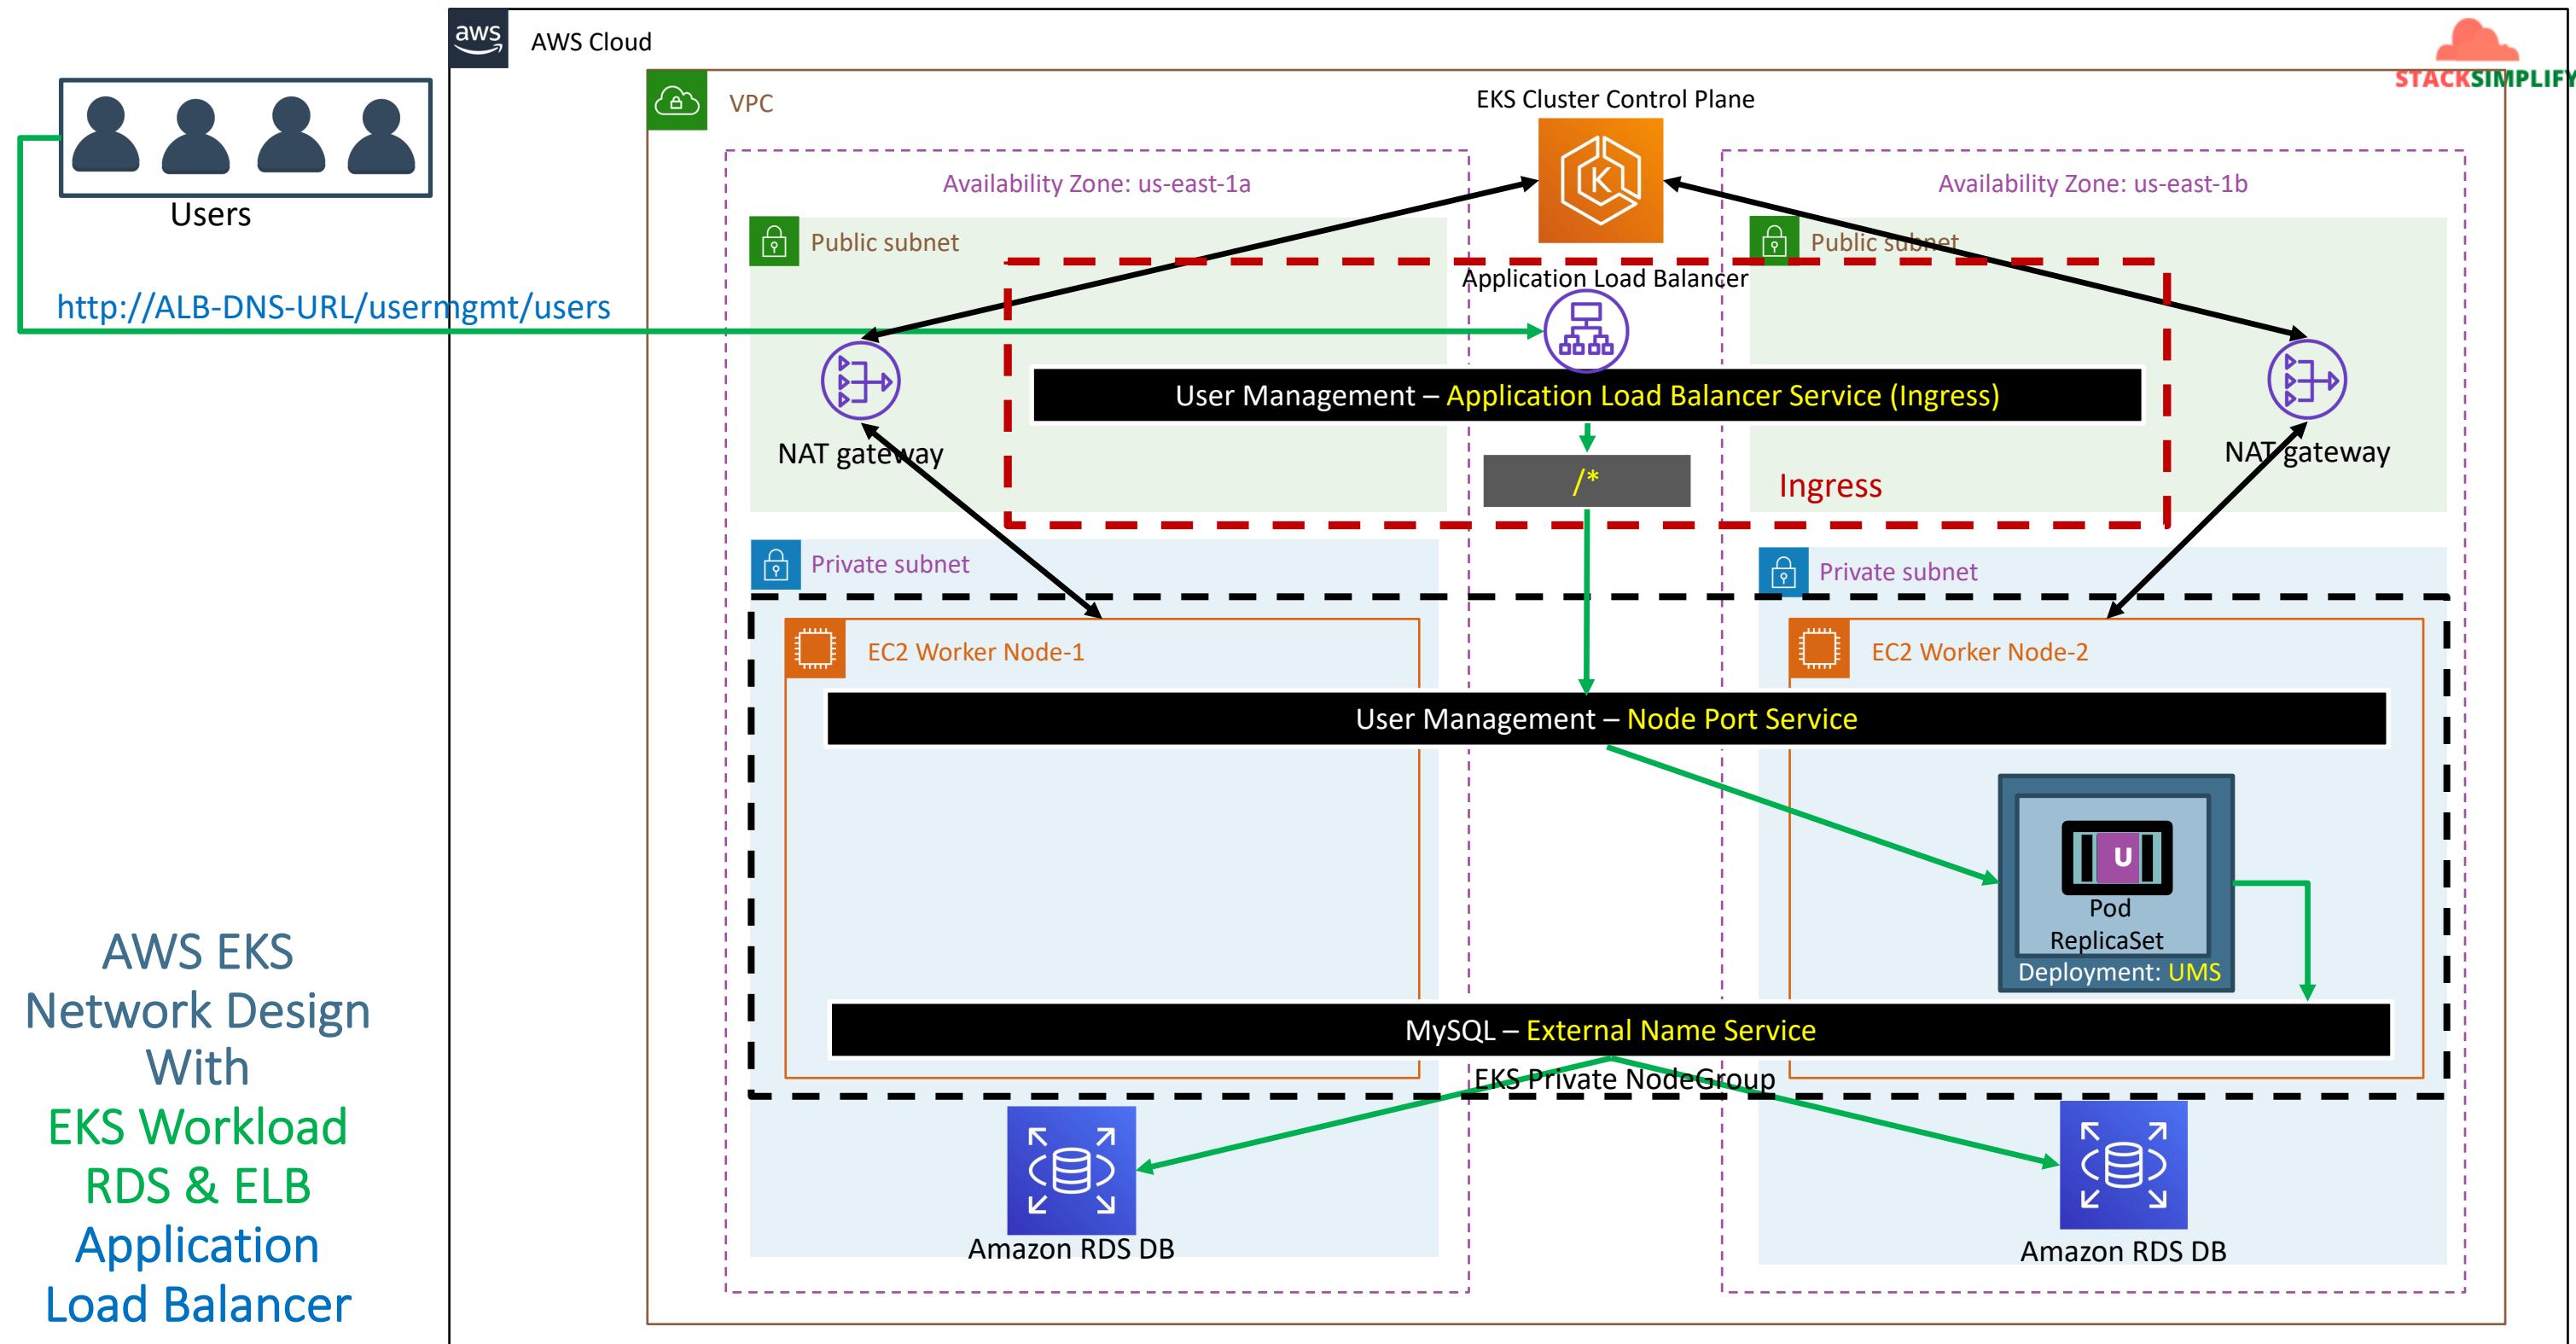


Amazon RDS



# AWS EKS RDS & ELB Application Load Balancer Ingress Controller Basics



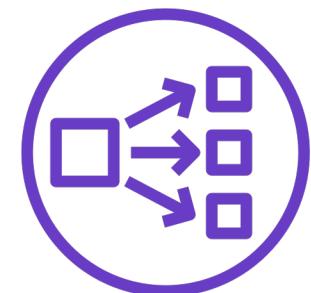




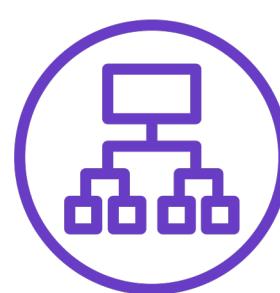
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



Application  
Load Balancer



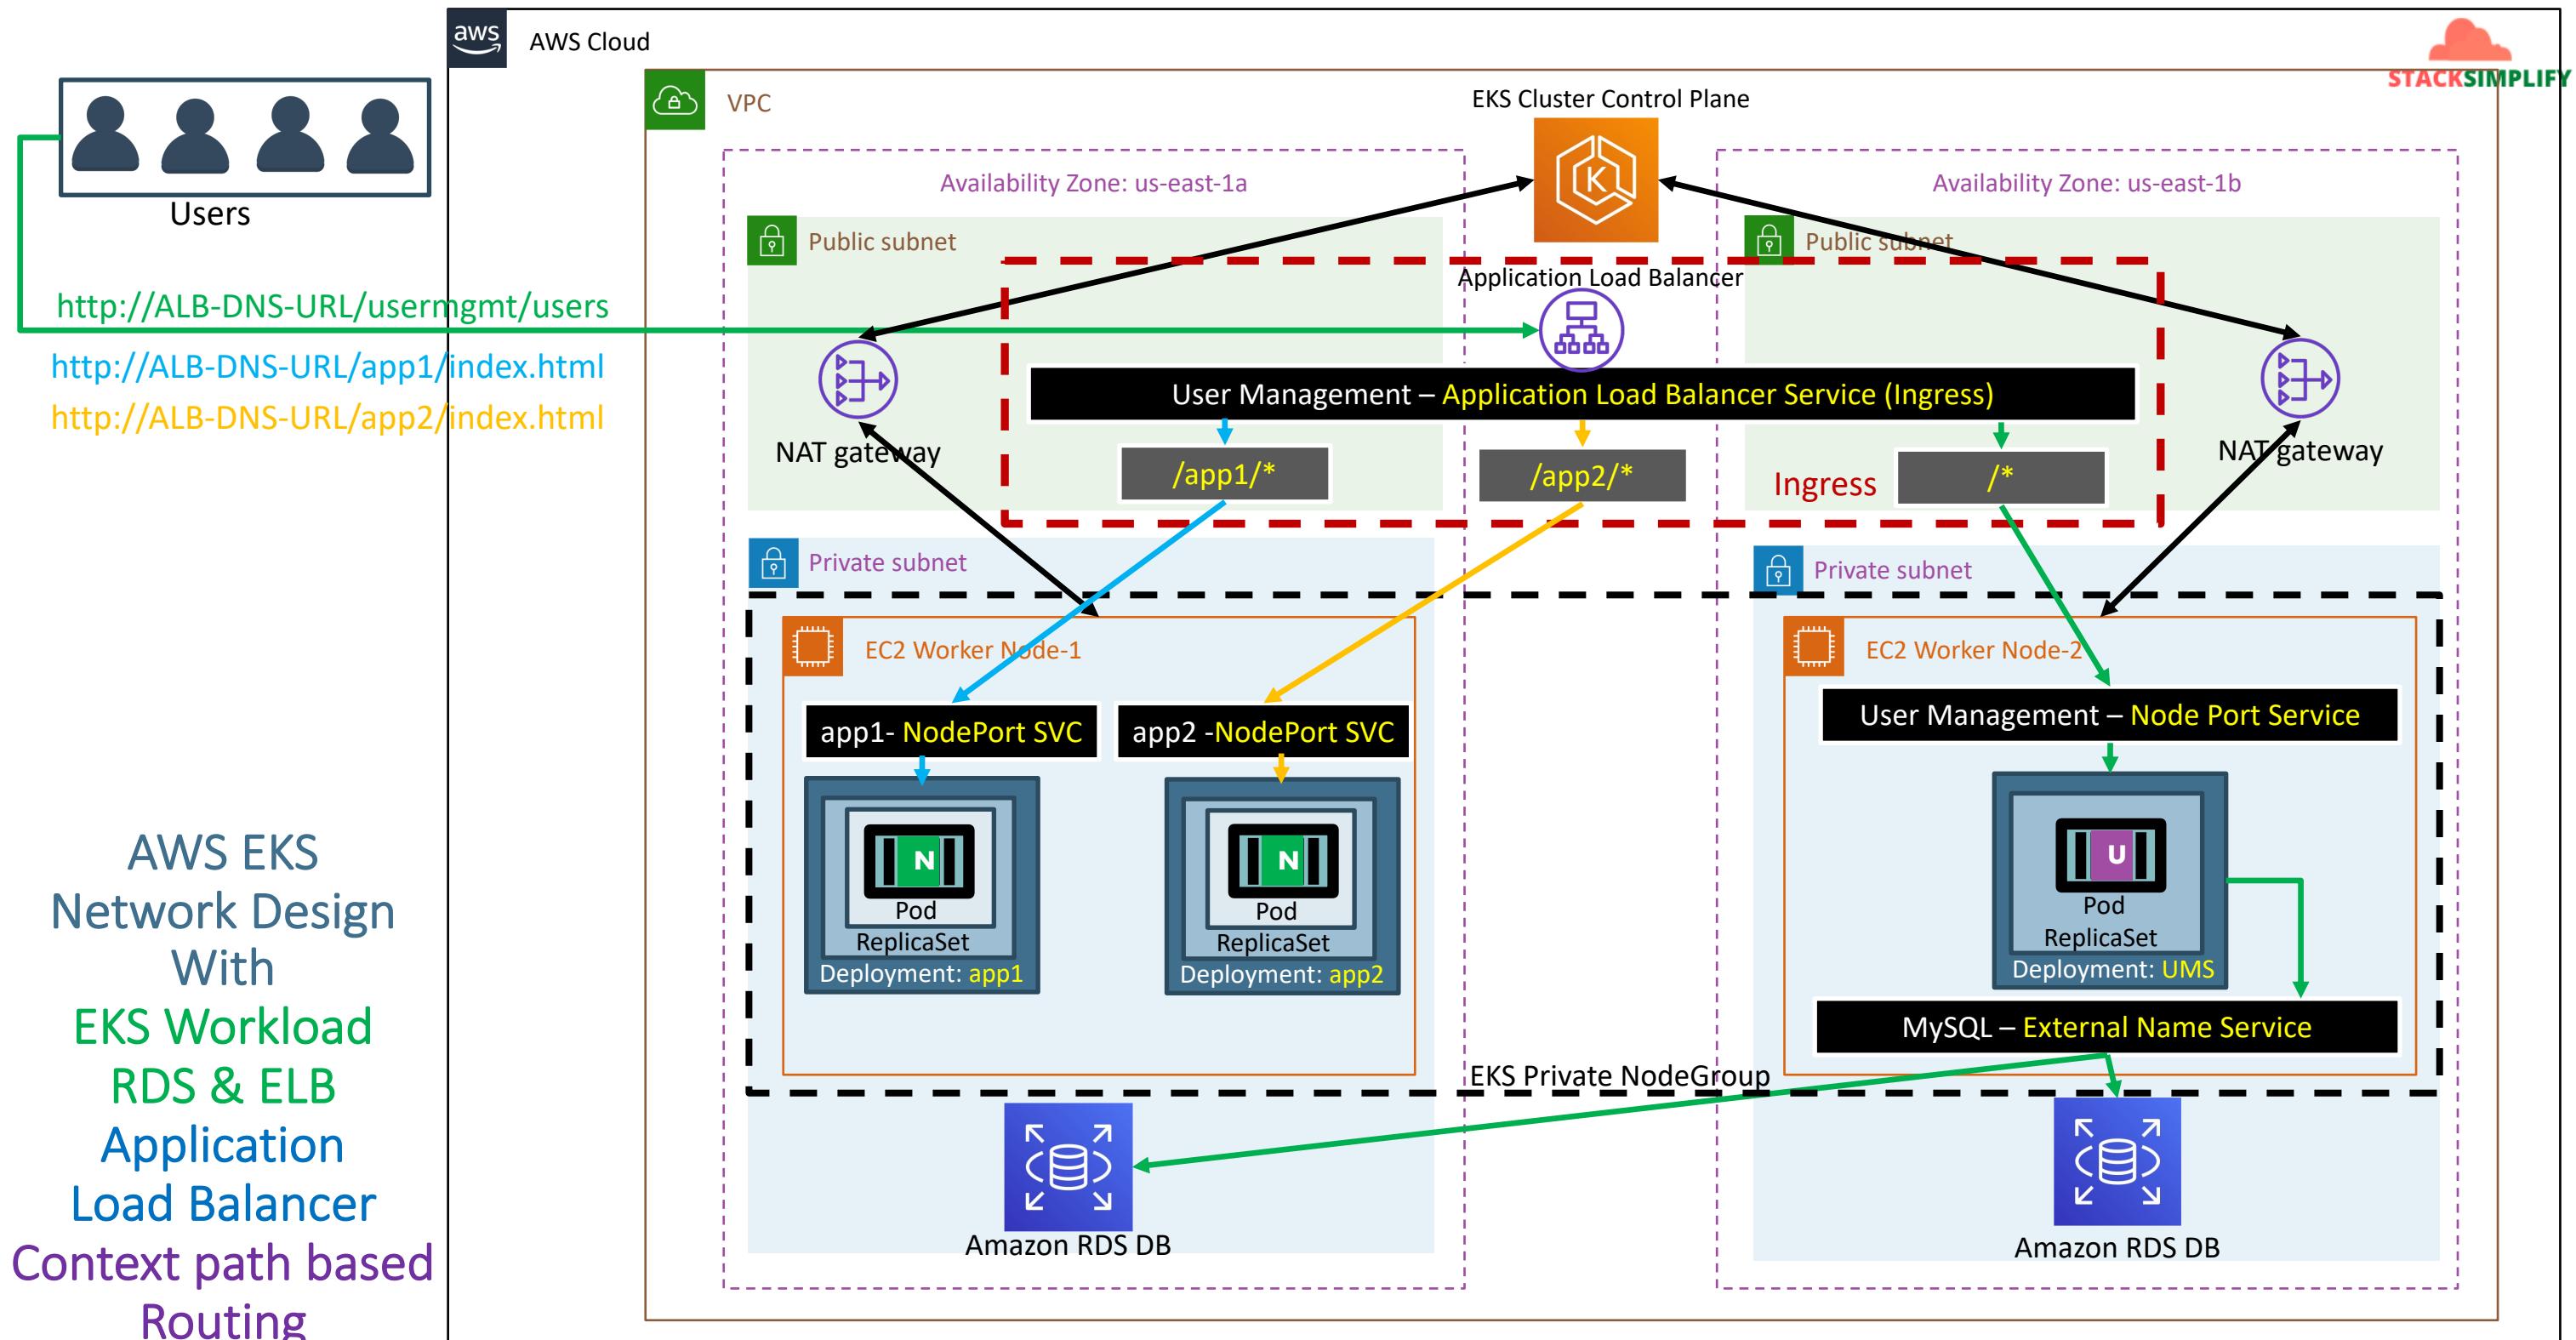
Amazon RDS

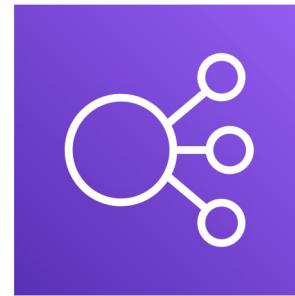


# AWS EKS RDS & ELB Application Load Balancer



## Ingress Context Path based Routing

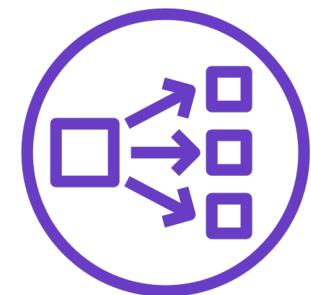




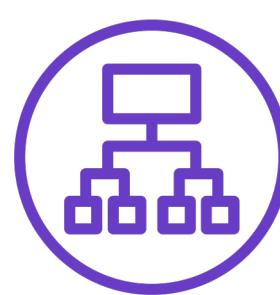
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



Application  
Load Balancer

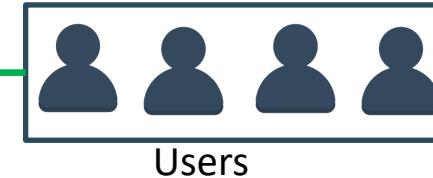


Amazon RDS



# AWS EKS RDS & ELB Application Load Balancer Ingress SSL





AWS Cloud

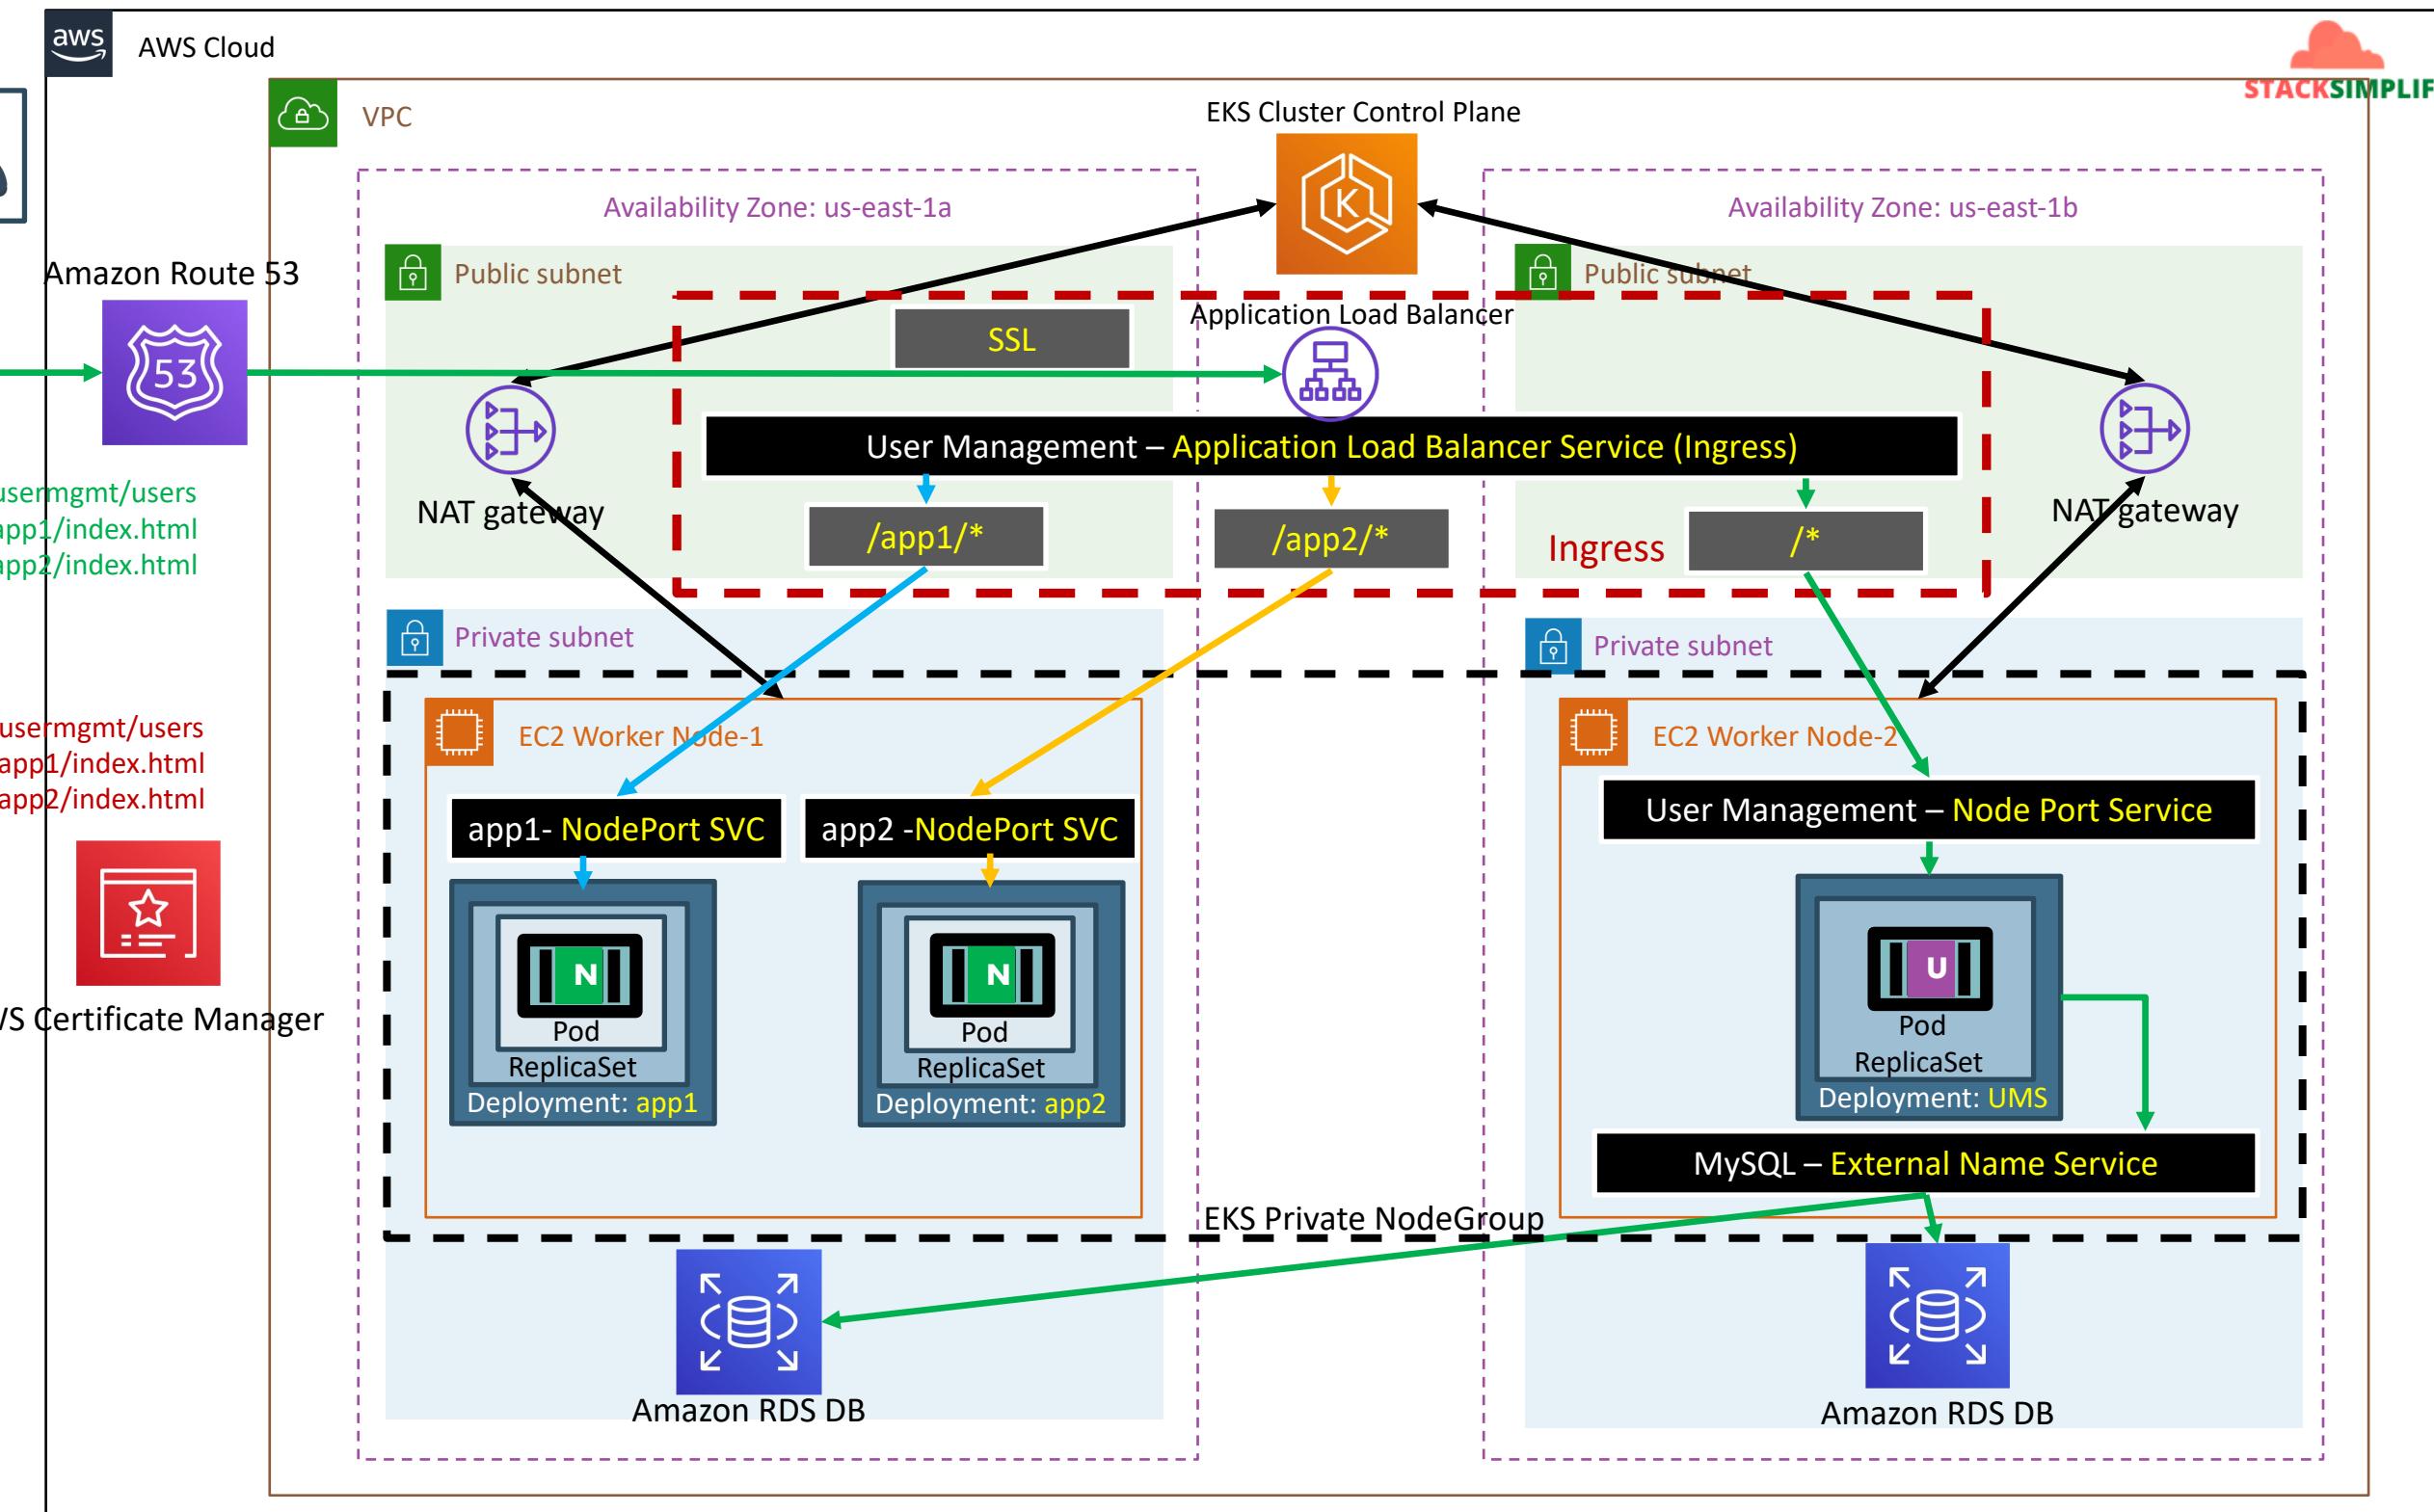
#### HTTP URLs

<http://ssldemo.kubeoncloud.com/usermgmt/users>  
<http://ssldemo.kubeoncloud.com/app1/index.html>  
<http://ssldemo.kubeoncloud.com/app2/index.html>

#### HTTPS URLs

<https://ssldemo.kubeoncloud.com/usermgmt/users>  
<https://ssldemo.kubeoncloud.com/app1/index.html>  
<https://ssldemo.kubeoncloud.com/app2/index.html>

## AWS EKS Network Design With EKS Workload RDS & ELB Application Load Balancer Ingress SSL

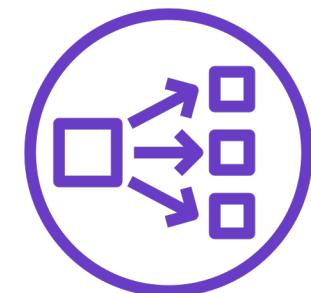




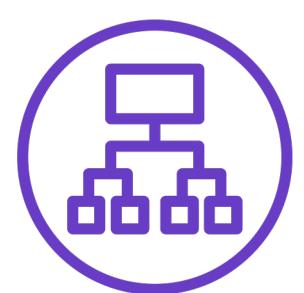
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



Application  
Load Balancer



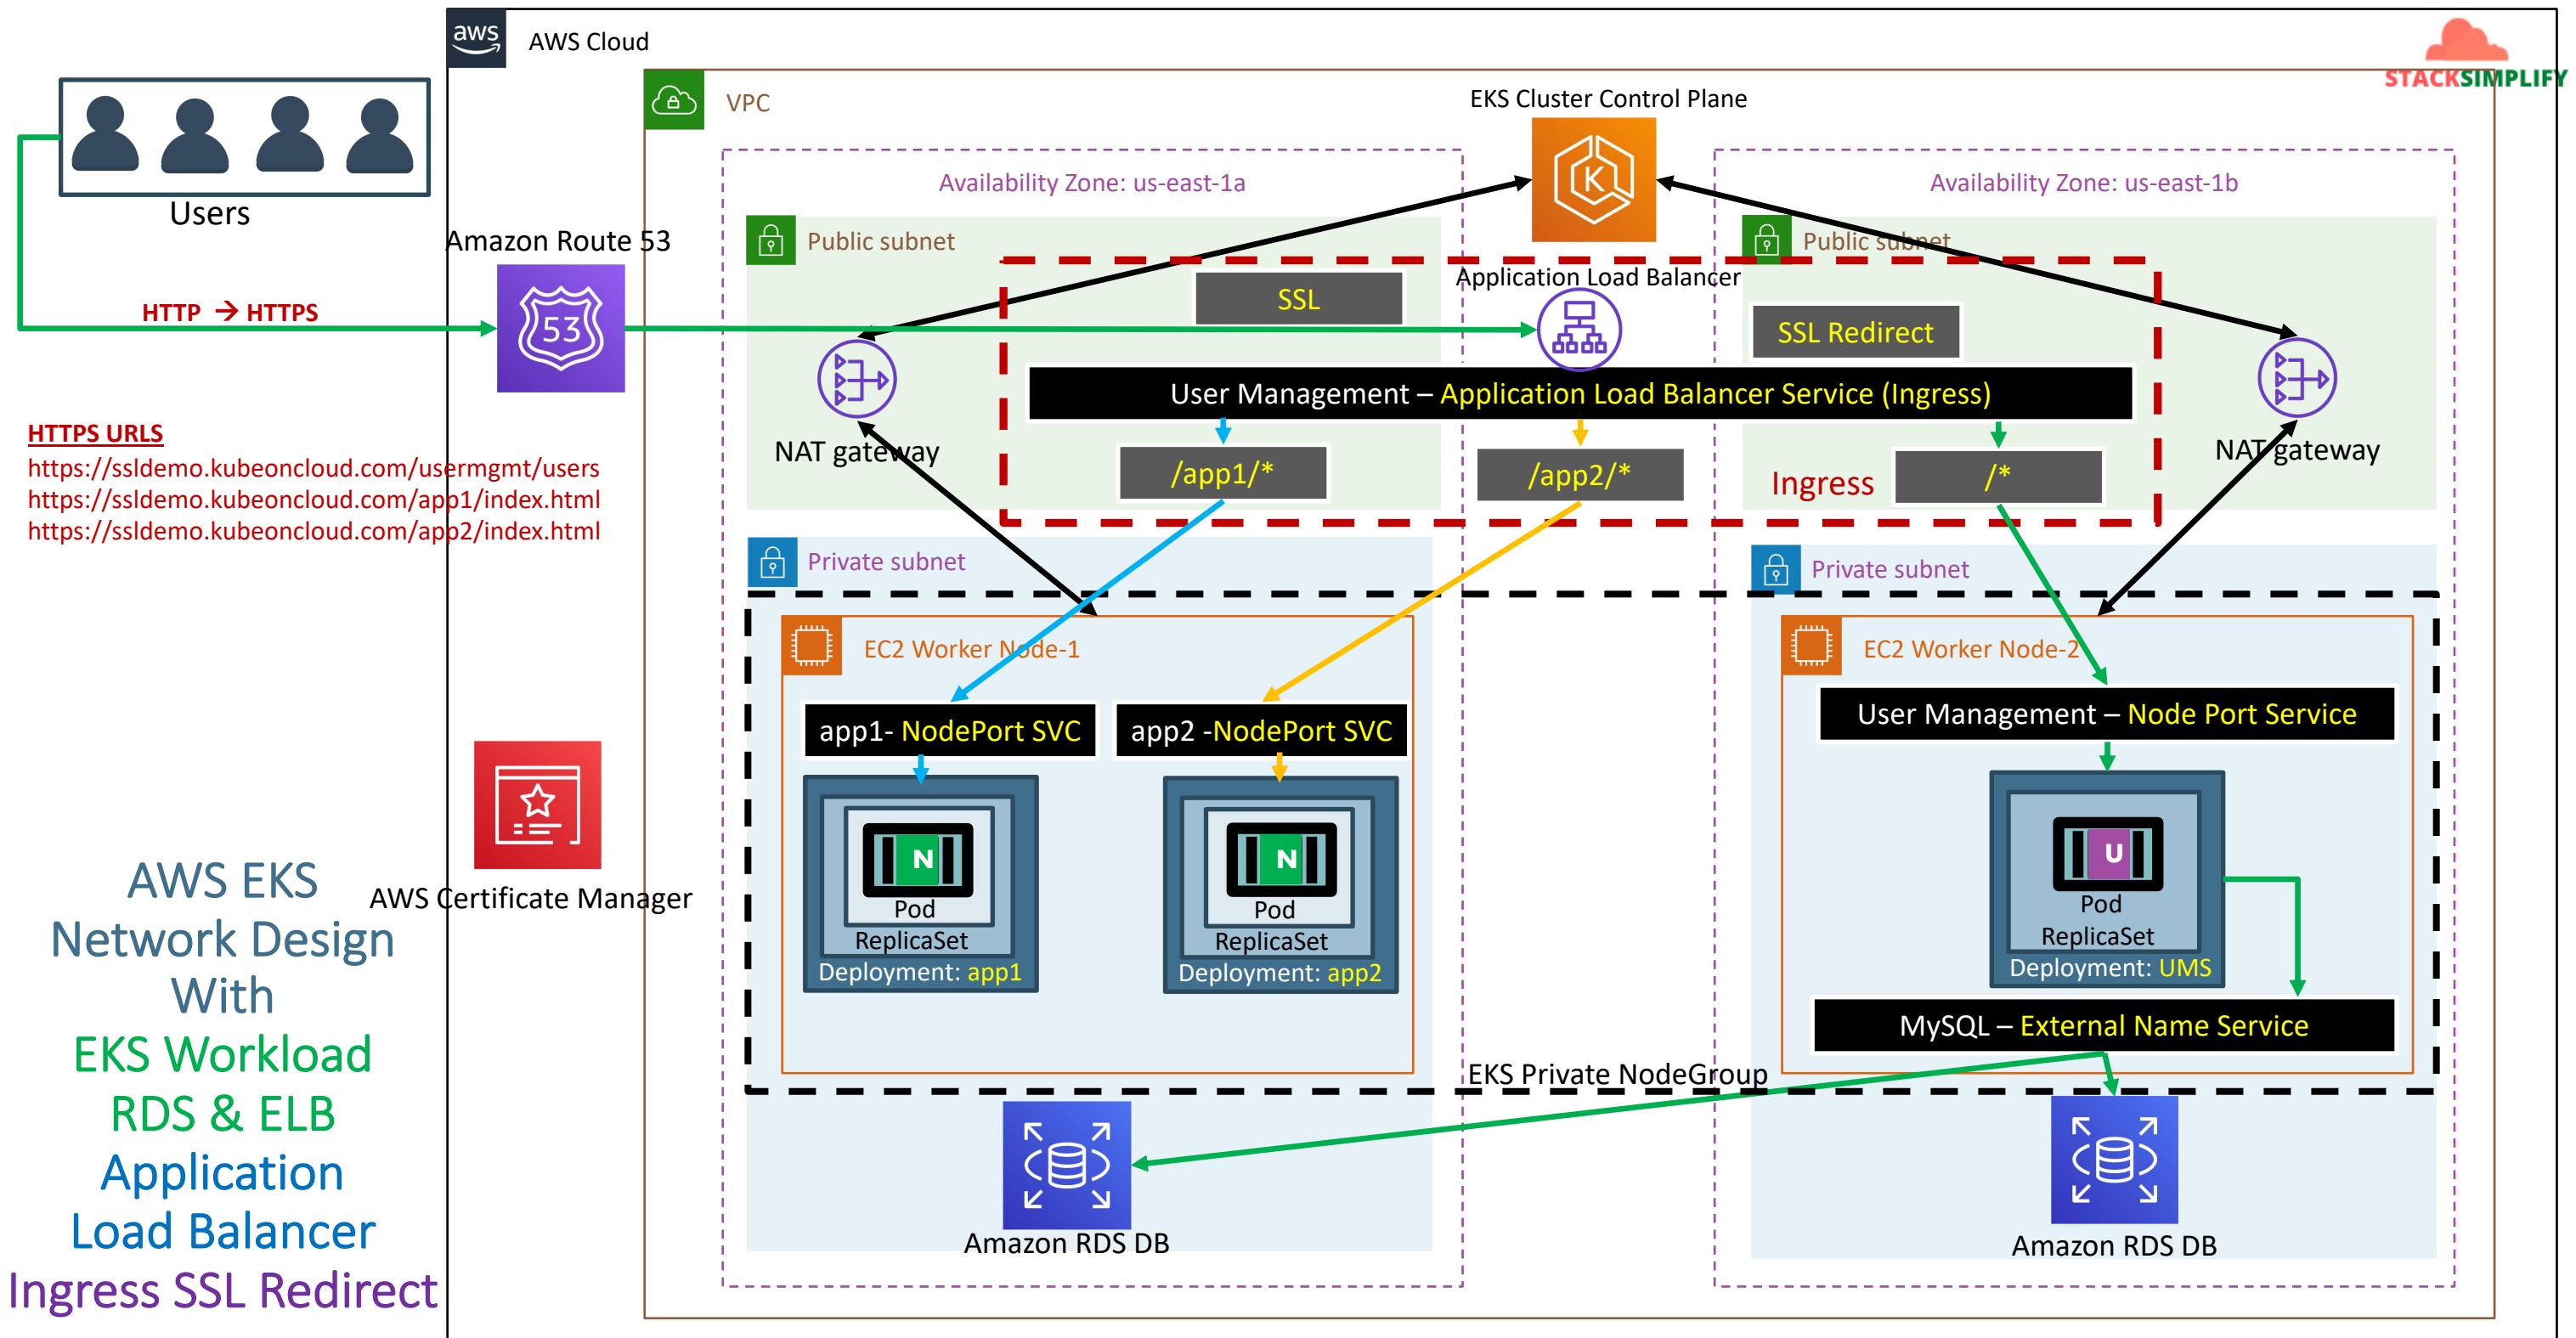
Amazon RDS



**AWS EKS**  
**RDS & ELB**  
**Application Load Balancer**



**Ingress SSL Redirect**

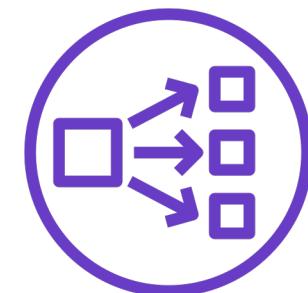




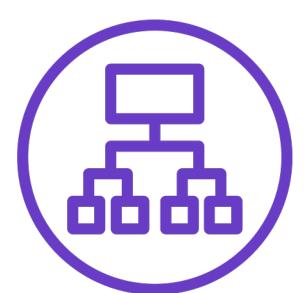
Elastic Load  
Balancing



Classic  
Load Balancer



Network  
Load Balancer



Application  
Load Balancer



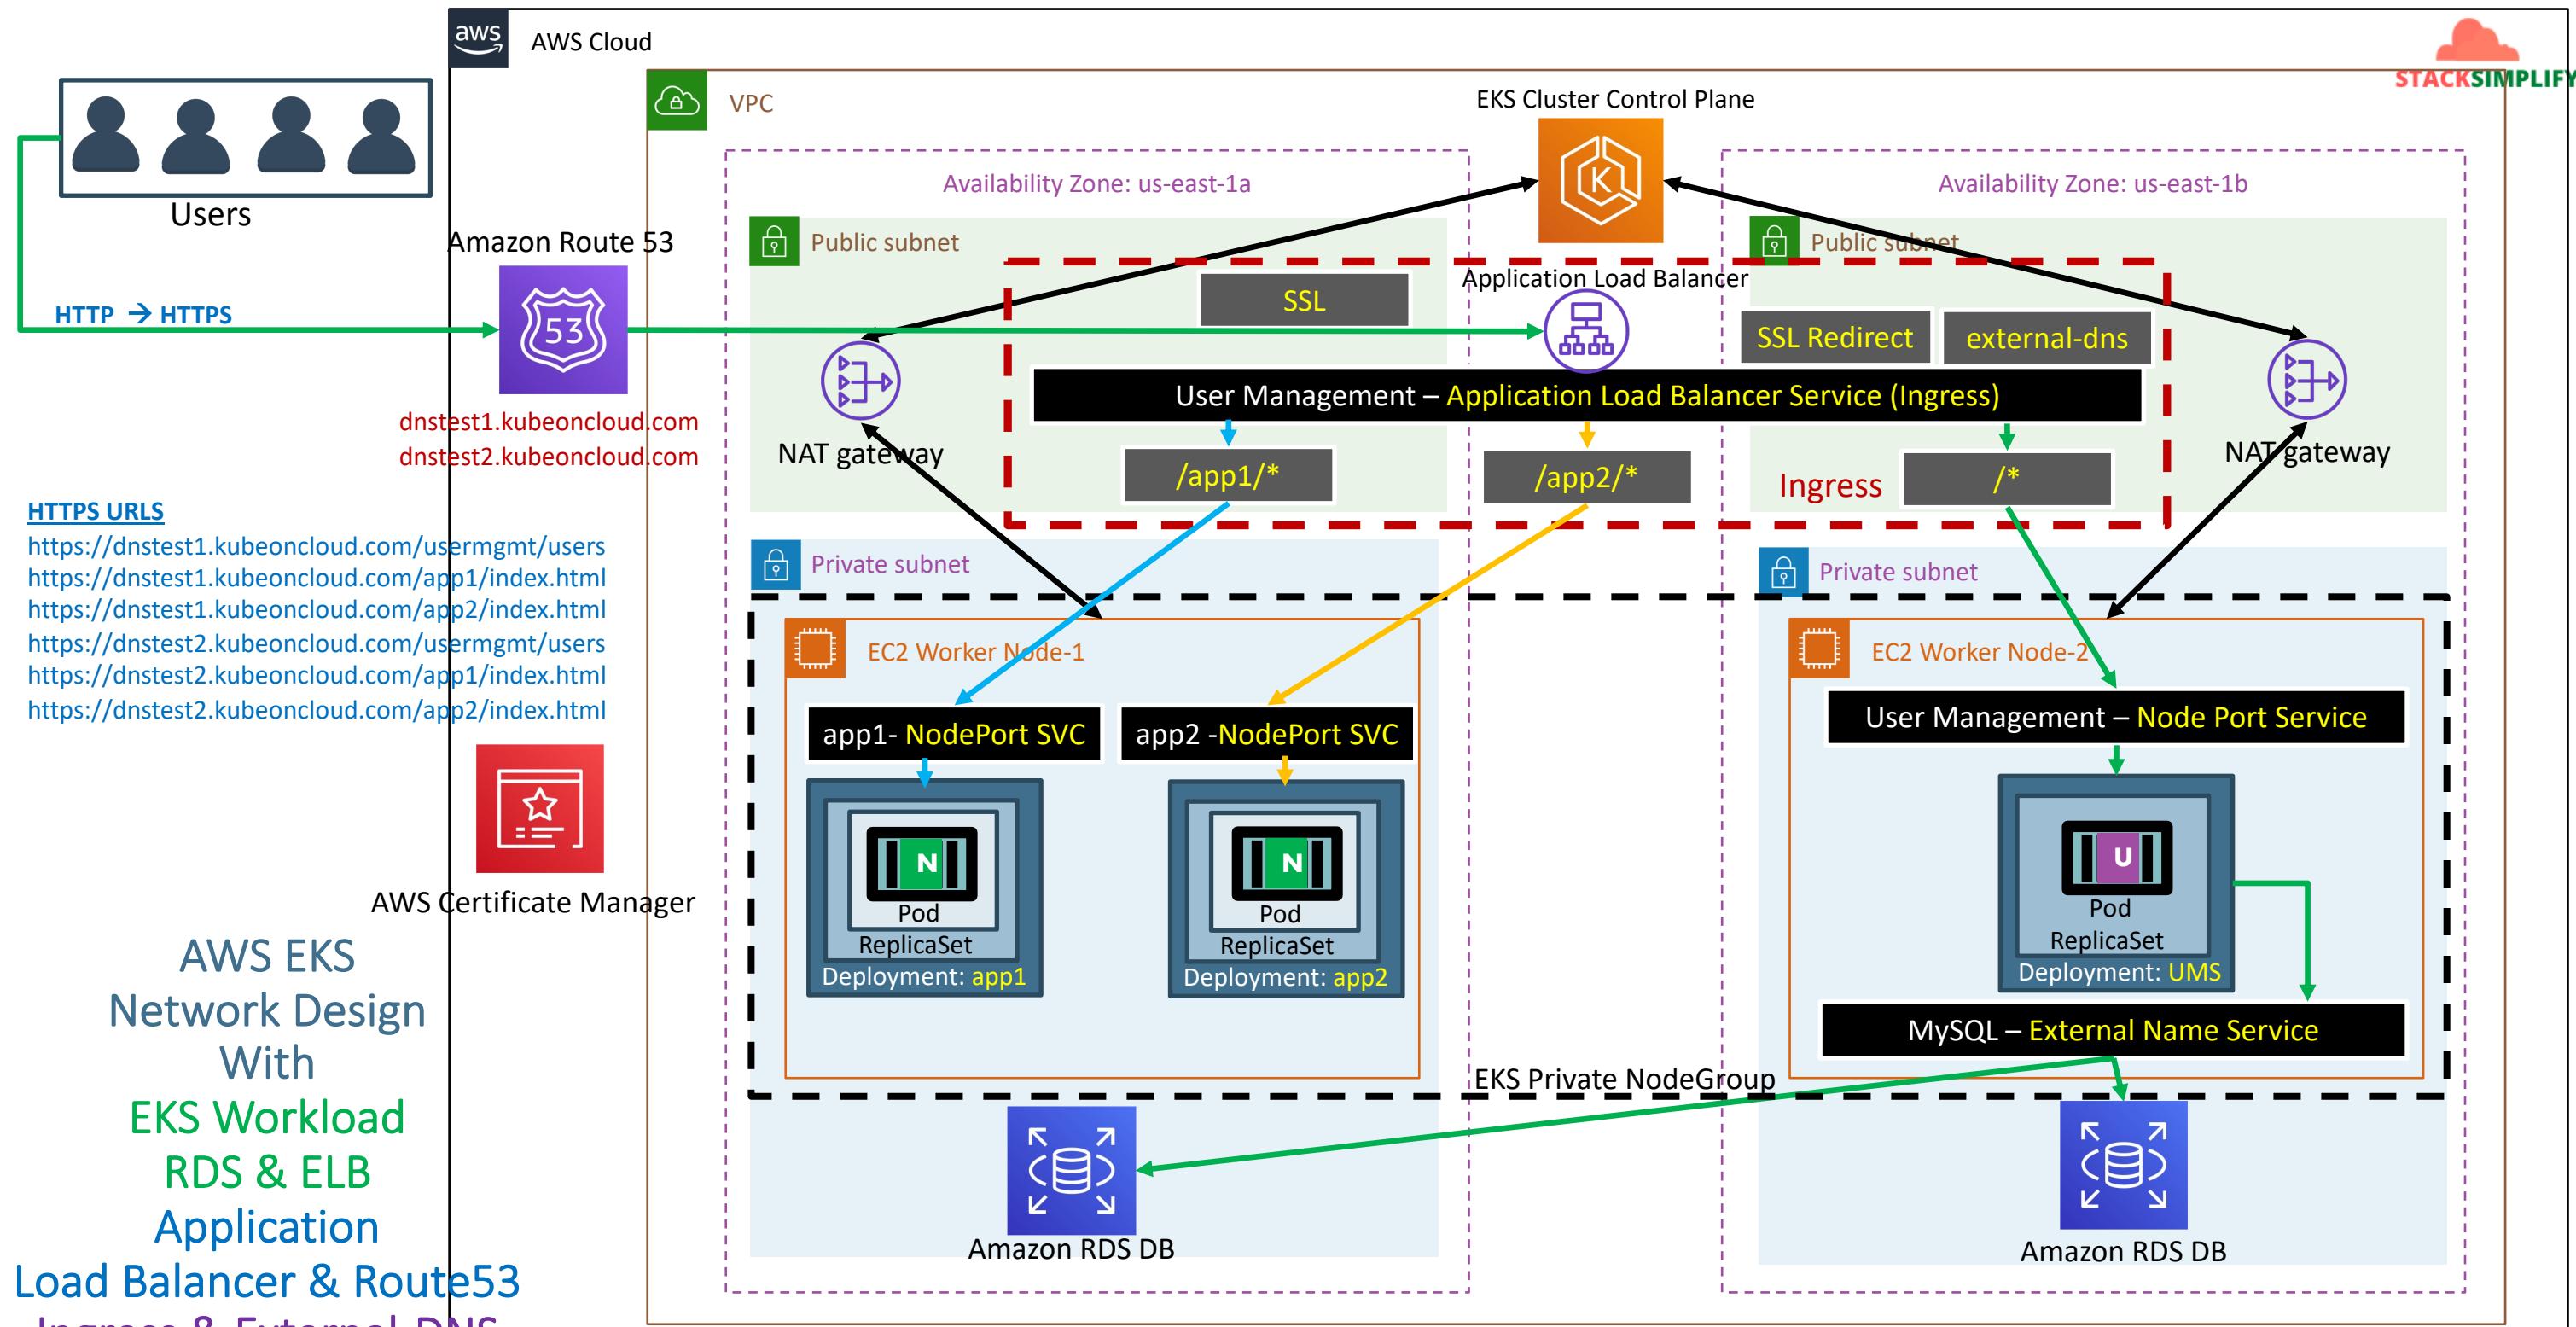
Amazon RDS



# AWS EKS RDS & ELB Application Load Balancer & Route53



## Ingress & External-DNS



**AWS EKS Network Design With EKS Workload RDS & ELB Application Load Balancer & Route53 Ingress & External-DNS**